

Robust Verification of Stochastic Systems

**Guarantees in the
Presence of Uncertainty**

Thom Badings

Author: Thom Badings
Title: Robust Verification of Stochastic Systems: Guarantees in the Presence of Uncertainty

Radboud Dissertations Series
ISSN: 2950-2772 (Online); 2950-2780 (Print)

Published by RABDOUD UNIVERSITY PRESS
Postbus 9100, 6500 HA Nijmegen, The Netherlands
www.radbouduniversitypress.nl

Design: Thom Badings
Cover: Thom Badings
Printing: DPN Rikken/Pumbo

ISBN: 9789493296909
DOI: <https://doi.org/10.54195/9789493296909>
Free download at: www.boekenbestellen.nl/radboud-university-press/dissertations

©2025 Thom Badings

This Ph.D. research has been funded by the Dutch Research Council (NWO) under the grant PrimaVera (NWA.1160.18.238) and has been carried out under the auspices of the research school IPA (Institute for Programming research and Algorithmics).

IPA Dissertation Series 2025-03

**RADBOUD
UNIVERSITY
PRESS**



This is an Open Access book published under the terms of Creative Commons Attribution-Noncommercial-NoDerivatives International license (CC BY-NC-ND 4.0). This license allows reusers to copy and distribute the material in any medium or format in unadapted form only, for noncommercial purposes only, and only so long as attribution is given to the creator, see <http://creativecommons.org/licenses/by-nc-nd/4.0/>.

Robust Verification of Stochastic Systems: Guarantees in the Presence of Uncertainty

Proefschrift ter verkrijging van de graad van doctor
aan de Radboud Universiteit Nijmegen
op gezag van de rector magnificus prof. dr. J.M. Sanders,
volgens besluit van het college voor promoties
in het openbaar te verdedigen op

donderdag 27 maart 2025
om 16:30 uur precies

door

Thomas Sebastiaan Badings

geboren te Zwolle

Promotoren

Prof. dr. Nils H. Jansen (Ruhr-Universität Bochum, Duitsland)

Prof. dr. Mariëlle I.A. Stoelinga

Manuscriptcommissie

Prof. dr. Frits W. Vaandrager

(Universität des Saarlandes, Duitsland)

Prof. dr. Jan Křetínský

(Masarykova univerzita, Tsjechië)

Dr. Frans A. Oliehoek

(Technische Universiteit Delft)

Dr. Jana Tůmová

(Kungliga Tekniska högskolan, Zweden)

Abstract

Verifying that systems are safe and reliable is crucial in today's world. For example, we want to prove that an autonomous drone will safely reach its target, or that a manufacturing system will not break down. Classical algorithms for verifying such properties often rely on a precise mathematical model of the system, for example, in the form of a *Markov chain* or a *Markov decision process (MDP)*. Such *Markov models* are probabilistic transition systems and are ubiquitous in many areas, including control theory, artificial intelligence (AI), formal methods, and operations research.

However, as systems become increasingly complex with more cyber-physical and AI components, *uncertainty* about the system's behavior is *inevitable*. As a result, transition probabilities in Markov models are subject to uncertainty, rendering many existing analysis algorithms inapplicable.

In this thesis, we fill this gap by developing novel verification methods for Markov models in the presence of uncertainty. To capture this uncertainty, we use *parametric* Markov models, where probabilities are described as functions over parameters. We study two perspectives on rendering verification methods *robust* against uncertainty: (1) *set-bounded uncertainty*, where only the set of possible parameter values is known and one can be robust in a *worst-case sense*, and (2) *stochastic uncertainty*, where the parameter values are described by a probability distribution and one can be robust in a *probabilistic sense*. By combining techniques from formal methods, AI, and control theory, our contributions span the following general problem settings:

1. We develop robust *abstraction techniques* for solving control problems for Markov models with continuous state and action spaces, and with set-bounded uncertain parameters. Based on the notion of *probabilistic simulation relations*, we show that such continuous control problems can be solved by only analyzing a finite-state abstraction, formalized as an MDP with sets of transition probabilities.
2. We present novel and scalable verification techniques for parametric Markov models with a *prior distribution* over the parameters. Our approaches are *sampling-based*, do not require any assumptions on the parameter distribution, and provide *probably approximately correct (PAC)* guarantees on the verification results.
3. We show that parametric models can be used to improve the *sample efficiency* of data-driven learning methods. We leverage tools from *convex optimization* to perform a sensitivity analysis, where we measure sensitivity in terms of *partial derivatives* of the polynomial function that describes a measure of interest.
4. We study *continuous-time Markov chains* where the initial state must be inferred from a set of (possibly uncertain) state observations. This setting is particularly relevant in *runtime monitoring*. We compute upper and lower bounds on reachability probabilities, *conditioned* on these observations. Our approach is based on a robust abstraction into an MDP with intervals of transition probabilities.

In conclusion, we develop verification methods that reason over uncertainty without sacrificing guarantees. While dealing with uncertainty can be computationally expensive, providing such guarantees is crucial for designing safe and reliable systems with cyber-physical and AI components. The aim of this thesis is to contribute to a better understanding of dealing with uncertainty in stochastic verification problems.

Samenvatting

Het verifiëren dat systemen veilig en betrouwbaar zijn is cruciaal. We willen bijvoorbeeld bewijzen dat een autonome drone veilig zijn bestemming bereikt, of dat een productiemachine niet kapot zal gaan. Klassieke algoritmes voor het verifiëren van zulke eigenschappen vereisen een exact wiskundig model van het systeem, bijvoorbeeld als een *Markov chain* of een *Markov decision process (MDP)*. Zulke *Markov modellen* zijn probabilistische transitiesystemen die in vele gebieden voorkomen, waaronder de regeltechniek, kunstmatige intelligentie (AI), formele methoden en besliskunde.

Doordat systemen steeds meer cyber-fysieke en AI componenten hebben neemt hun complexiteit toe. Hierdoor is *onzekerheid* over het gedrag en de dynamica van deze systemen *onvermijdelijk*. Dit betekent dat de transitiekansen in Markov modellen onzeker zullen zijn, waardoor veel klassieke algoritmes niet toepasbaar zijn.

In deze thesis ontwikkelen we nieuwe verificatiemethodes voor Markov modellen met onzekerheid. We modelleren onzekerheid met *parametrische* Markov modellen, waarin transitiekansen als functies over parameters zijn gedefinieerd. We onderzoeken twee perspectieven voor het modelleren en *robust* analyseren van onzekerheden: (1) *set-begrensde onzekerheid*, waarbij alleen de set van mogelijke parameterwaarden bekend is en een *worst-case* robuuste analyse mogelijk is, en (2) *stochastische onzekerheid*, waarbij de parameterwaarden worden beschreven door een kansverdeling en robuustheid een *probabilistische* opvatting heeft. We combineren technieken uit de formele methoden, AI, en de regeltechniek in de volgende problemstellingen:

1. We ontwikkelen robuuste *abstractie-technieken* voor planningsproblemen in Markov modellen met continue toestands- en actieruimtes, en met set-begrensde onzekere parameters. We gebruiken de notie van *probabilistische simulatierelaties* om zulke problemen op te lossen door enkel een eindige abstractie te analyseren. Deze abstractie formaliseren we als een MDP met sets van transitiekansen.
2. We presenteren nieuwe en schaalbare verificatietechnieken voor parametrische Markov modellen met een *kansverdeling* over de parameters. Onze technieken zijn *data-gedreven*, vereisen geen aannames op deze kansverdeling, en geven *probably approximately correct (PAC)* garanties op de verificatieresultaten.
3. We gebruiken parametrische modellen om de *sample efficiëntie* van data-gedreven leermethodes te verbeteren. We gebruiken technieken uit *convexe optimalisatie* om sensitiviteitsanalyses te doen, waarbij we de sensitiviteit uitdrukken in termen van de *partiële afgeleiden* van de polynomiale functie die het model beschrijft.
4. We onderzoeken *continuous-time Markov chains* waar de initiële toestand moet worden afgeleid uit (mogelijk onzekere) observaties. Deze setting is in het bijzonder relevant voor *systeemonitoring*. We berekenen boven- en ondergrenzen op analyses, *geconditioneerd* op de gegeven observaties. Onze methode is gebaseerd op een robuuste abstractie als een MDP met intervallen van transitiekansen.

In conclusie, we ontwikkelen verificatiemethodes die wiskundige garanties geven ondanks onzekerheden. Hoewel het geven van zulke garanties rekenintensief kan zijn, is dit cruciaal voor het ontwerpen van veilige en betrouwbare systemen met cyber-fysieke en AI componenten. Met deze thesis hopen we bij te dragen aan een beter begrip van het omgaan met onzekerheid in stochastische verificatieproblemen.

Acknowledgements

I can still remember if it were yesterday that I sent an email to Ingrid, the secretary of the Software Science Department at Radboud University, asking whether there were any openings for doing a Ph.D. in Nijmegen. Ingrid forwarded me to Nils, who by luck was soon going to hire three Ph.D. students. That ended up being Christoph, Marnix, and me, and it was the start of an amazing 4-year journey.

I can only say that I have been incredibly lucky to have Nils as my Ph.D. advisor. Nils, you give me so much freedom to explore my academic path. You guided me when needed, but more often than not, you simply asked what I would like to do the most. From you, I learned the great attitude of: “*If something makes you happier doing your job, then just do it!*” Moreover, your (copyrighted) quote that “*uncertainty is inevitable*” has been part of the fundamental basis for my Ph.D. research. By always involving us in the workshops and events you have organized, you have not just taught me how to do research, but you have also been my guide into the academic world.

If Nils’ enthusiasm wasn’t enough yet, then I could always still count on Mariëlle. Mariëlle, you have been a great advisor and supervisor who always managed to make things exciting. You apply your saying “*no risk, no fun*” to many different situations, and (to some extent) that saying is also absolutely true! My Ph.D. project is part of PrimaVera, a consortium on predictive maintenance that Mariëlle has tirelessly led over the past years. Being part of the PrimaVera team, I have met many amazing people from different universities and companies in the Netherlands.

Next, I would like to thank Frits Vaandrager, Holger Hermanns, Jan Křetínský, Frans Oliehoek, and Jana Tůmová for acting as members of the manuscript committee, and for providing sharp and valuable feedback on my thesis. Special thanks go out to Matthijs for reading the introduction of my thesis in an early stage.

All my fellow LAVA-lab colleagues (Dennis, Christoph, Marnix, Thiago, Eline, Merlijn, Maris, and Wietze), a.k.a. Nils’ paper factory, I thank you for making my Ph.D. journey a lively experience. Later in my Ph.D., when Nils started his position in Bochum, our group was expanded with more amazing people (Joshua, Jule, Marcel, Markel, Miriam, Özner, and Verena). The same holds for all others in the Software Science Department, with whom I had the pleasure of having countless important lunch discussions. Sebastian, not only do we (almost) share a (middle) name, but we also share authorship on several papers. Matthias, I always enjoyed our collaborations, and I love how you can put basically anything in perspective by saying it was still more reliable than the German trains. Also, shoutout to Sebastian and Matthias for being willing to share their thesis template with me—You absolutely saved me from many frustrating hours of stumbling around with L^AT_EX templates!

The interaction with many incredible people, both within and outside the Netherlands, is what made my Ph.D. journey an amazing experience. I don’t think there’s any other type of job in which people are so passionate about their work and eager to collaborate. Steve, you were the best possible host during Thiago’s and my trip to Austin, making it a truly awesome experience (the “chicken shit bingo” must still be the most Texan thing I’ve ever seen). Licio, I’m writing part of these acknowledgments while sitting in the Californian sun—You made my visit to Stanford a great trip. You have been a wonderful collaborator, and I’ve always enjoyed our intense discussions. Similar thanks go to the

many people I've met over the past years at conferences, workshops, and other events.

I am grateful to the PrimaVera project for sponsoring my Ph.D. research. Special thanks to my fellow PrimaVera Ph.D.'s and postdocs Bas, Lisandro, Luc, Matthias, Natália, Núbia, Ragnar, Roel, Thiago, Zaharah, and others from the PrimaVera project who I am now inevitably forgetting.

Thanks to all my collaborators and co-authors before and during my Ph.D., namely: Alessandro Abate, Eline Bovy, Murat Cubuktepe, Arnd Hartmanns, Nils Jansen, Sebastian Junges, Joost-Pieter Katoen, Wietze Koops, Ahmadreza Marandi, Mahdi Nazeri, Dave Parker, Hasan Poonawala, Dennis van Putten, Luke Rickard, Licio Romao, Vahab Rostampour, Jacquelien Scherpen, Thiago Dias Simão, Sadegh Soudjani, Francisco Souza, Mariëlle Stoelinga, Marnix Suilen, Ufuk Topeu, and Matthias Volk. This is a subtle way of me saying: Thanks for the great collaborations, and let's continue to do more!

My first hands-on experience with academic research actually goes back further than my Ph.D. During my master's thesis at the University of Groningen, on which I was supervised by Jacquelien Scherpen and Vahab Rostampour, I got the opportunity to publish my results. While I was not sure at that point whether I wanted to pursue a career in academia, I am very sure that the seed for my passion for research was already planted back then!

Outside of academia, I owe a lot of thanks to all "Hendige Heeren" (Sake, Balazs, Cees, Colin, Eelco, Herman, Jippe, Stefan, Koen, Roland, Sebastian, and Yorick), who are my former water polo team from my life as a student in Groningen. We may not see each other that often anymore, but the events that we do, such as the sailing weekend and the yearly winter sports, are among the best events of the year. In particular, Herman and Sake, you're the best for being my paranympths (a.k.a. bodyguards) at my defense. Similar shoutouts go to my friends from high school in Zwolle (Roy, Daan, Gerjan, Jan, Justin, Mathieu, Menno, Rico, Rutger, Sander, Stephan, and Thomas), with whom I had many great times and adventures as well.

To my mom (Pauline), dad (Erik), sister (Nienke), and brother (Matthijs), thanks for the unconditional support, and for always being there for me. It was in 2018 that my dad defended his Ph.D. at Radboud University. Now, it is up to me to do the same very soon. Finally, Marjolein, I am so grateful for how amazing, kind, and supportive you are. Living with someone who has spent the past four years on making pictures, irregularly traveling across the world to talk to people, and responding to paper reviews on vacation is not always easiest. I'm thankful that you've always supported me, even now I've decided to do a postdoc in Oxford. But in the end, maybe these acknowledgments are not the best place to express how much you truly mean to me.

*Thom Badings
November 15, 2024
Lent, The Netherlands*

Contents

1	Introduction	1
1.1	Stochastic Systems	1
1.2	Verification	3
1.3	Uncertainty is Inevitable	5
1.4	Robustness	8
1.5	Challenges and Contributions	9
1.6	Overview of Key Techniques	11
1.7	Navigating This Thesis	15
1.8	Overview of Publications	18
I	Foundations	21
2	Preliminaries	23
2.1	Basic Notation	23
2.2	Optimization Problems	24
2.3	Probability Theory	25
2.3.1	Probability distributions	25
2.3.2	Random variables	26
2.3.3	Stochastic processes	26
3	A Primer on Markov Decision Processes	29
3.1	Markov Decision Processes	29
3.1.1	Paths and sets of paths	31
3.1.2	Schedulers	31
3.2	Analyzing MDPs	34
3.2.1	Probabilistic computation tree logic	34
3.2.2	Measures	35
3.2.3	Value iteration for MDPs	37
3.3	Robust Markov Decision Processes	40
3.3.1	Interval MDPs	41
3.3.2	Nature	42
3.3.3	Robust measures	43
3.3.4	Optimal schedulers for RMDPs	44
3.3.5	Connection to other models	45
	Summary	46

II Discrete-Time Stochastic Systems	47
4 Foundations of DTSSs	49
4.1 Introduction	49
4.2 Discrete-Time Stochastic Systems	50
4.2.1 Markov policy	53
4.2.2 Stochastic kernel	54
4.3 Reach-Avoid Probability	54
4.3.1 Computing satisfaction probabilities	55
4.3.2 Extension to PCTL	56
Summary	56
5 Probabilistic Simulation Relations	59
5.1 Introduction	59
5.1.1 Policy evaluation	60
5.1.2 Optimal control	60
5.1.3 Lower bound control	61
5.2 The DTSS Policy Synthesis Problem	62
5.2.1 Approaches to DTSS policy synthesis	62
5.2.2 Shortcomings of abstraction-based control	63
5.2.3 An overview of our approach	64
5.3 Probabilistic Simulation Relations	64
5.3.1 Relating reach-avoid specifications	65
5.3.2 Relating DTSSs and MDPs	66
5.3.3 Comparison to other behavioral relations	68
5.4 Correct-by-Construction Markov Policy Synthesis	69
5.5 DTSS Relations With Robust MDPs	74
5.5.1 Probabilistic alternating simulation relation	74
5.5.2 Lower bounding satisfaction probabilities	76
5.5.3 Markov policy synthesis with RMDPs	77
Summary	78
6 Reach-Avoid Control of Linear DTSSs	81
6.1 Linear DTSS	81
6.1.1 Assumptions	82
6.1.2 Problem statement	84
6.2 MDP Abstraction of Linear DTSS	85
6.2.1 Relation induced by the abstract MDP	89
6.3 Sampling-Based Probability Intervals	90
6.3.1 Bounds for the transition probabilities	91
6.3.2 *The scenario approach	93
6.3.3 *Proof of Theorem 6.19	96
6.3.4 Tightness of probability intervals	99
6.4 Abstraction-Based Control Algorithm	100
6.4.1 Interval MDP abstraction	100
6.4.2 Solving Problem 6.6 with high probability	102

6.5	Exploiting Stability for Smaller Abstractions	104
6.5.1	Backward reachable sets	105
6.5.2	Constructing smaller abstractions	106
6.6	Experimental Evaluation	107
6.6.1	UAV motion planning	107
6.6.2	Spacecraft docking	110
6.7	Related Work	112
6.8	Discussion	113
	Summary	114
7	DTSSs With Uncertain Parameters	117
7.1	Parameter Uncertainty in DTSSs	117
7.1.1	Assumptions	118
7.1.2	Problem statement	120
7.1.3	Overview of our abstraction technique	120
7.2	Parameter Robustness in IMDP Abstractions	121
7.2.1	Nominal dynamics model	121
7.2.2	IMDP abstraction of the nominal model	121
7.2.3	PAC probability intervals	127
7.3	Abstraction Algorithm	131
7.3.1	Solving Problem 7.5 with high probability	132
7.4	Experimental Evaluation	133
7.4.1	Longitudinal drone dynamics	133
7.4.2	Building temperature control	135
7.5	Related Work	137
7.6	Discussion	138
	Summary	138
III	Parametric Markov Decision Processes	141
8	Foundations of Parametric MDPs	143
8.1	Introduction	143
8.2	Parametric MDPs	144
8.2.1	Parameter instantiation	145
8.3	Verifying Parametric MDPs	146
8.3.1	Solution function	146
8.3.2	Parameter synthesis	148
8.4	Challenges	149
	Summary	150
9	The Scenario Approach for Parametric MDPs	151
9.1	Introduction	151
9.2	Motivating Example	152
9.3	Problem Statement	153
9.4	Bounding the Satisfaction Probability	155
9.4.1	Chance-constrained problem	155

9.4.2	Scenario problem	156
9.4.3	Sample complexity	159
9.5	Improving Bounds by Discarding Samples	159
9.5.1	Scenario problem with discarded samples	160
9.5.2	Problem 9.6 solved	161
9.6	Experimental Evaluation	163
9.6.1	UAV Motion Planning	164
9.6.2	Parameter Synthesis Benchmarks	166
9.7	Discussion	169
	Summary	170
10	Sensitivity Analysis for Parametric Markov Chains	171
10.1	Introduction	171
10.2	Overview	173
10.3	Problem Statement	175
10.3.1	Parametric robust Markov chains	176
10.3.2	Problem statement	178
10.4	Differentiating Solution Functions for pMCs	178
10.4.1	Computing derivatives explicitly	179
10.4.2	Computing k highest derivatives	179
10.5	Differentiating Solution Functions for pRMCs	181
10.5.1	Computing derivatives via pMCs	182
10.5.2	Computing derivatives explicitly	183
10.5.3	Computing k highest derivatives	185
10.6	Numerical Experiments	186
10.7	Related Work	191
10.8	Discussion	192
	Summary	193
IV	Continuous-Time Markov Chains	195
11	Foundations of CTMCs	197
11.1	Introduction	197
11.2	Continuous-Time Markov Chains	198
11.3	Verifying CTMCs	201
11.3.1	Continuous stochastic logic	201
11.3.2	Measures	203
11.3.3	Algorithms	204
11.4	Parametric Continuous-Time Markov Chains	205
11.4.1	Parameter instantiation	205
11.4.2	Verifying pCTMCs	206
11.5	Challenges	207
	Summary	208
12	CTMCs With Uncertain Rates	209
12.1	Introduction	209

12.2	CTMCs With Uncertain Rates	211
12.2.1	Measures and solution functions	212
12.2.2	Problem statement	212
12.2.3	Illustrative example	214
12.2.4	Our approach	215
12.3	Precise Sampling-Based Prediction Regions	216
12.3.1	Constructing prediction regions	216
12.3.2	Bounding the containment probability	219
12.3.3	Algorithm for computing prediction regions	221
12.4	Imprecise Sampling-Based Prediction Regions	222
12.4.1	Prediction regions on imprecise solutions	223
12.4.2	*Proof of Theorem 12.21	224
12.4.3	Computing the complexity	226
12.4.4	Solution refinement scheme	227
12.5	Batch Verification for CTMCs	227
12.6	Numerical Experiments	228
12.6.1	Converting pCTMCs into upCTMCs	229
12.6.2	Applicability	230
12.6.3	Scalability	232
12.6.4	Comparison to baselines	232
12.7	Related Work	233
	Summary	234
13	CTMCs With Imprecisely Timed Observations	235
13.1	Introduction	235
13.2	The CTMC Monitoring Problem	236
13.2.1	Problem statement	238
13.2.2	Our approach	239
13.3	Conditional Reachability With Imprecise Evidence	240
13.3.1	Unfolding the CTMC into an MDP	240
13.3.2	Computing conditional probabilities in MDPs	242
13.3.3	Computing evidence probability	243
13.4	Abstraction of Conditioned MDPs	244
13.4.1	Abstracting evidence times	245
13.4.2	Abstraction refinement	247
13.5	Bounding the Conditional Reachability	248
13.6	Numerical Experiments	249
13.6.1	Feasibility	250
13.6.2	Scalability	251
13.7	Related Work	252
13.8	*Proofs	253
13.8.1	Proof of Theorem 13.11	253
13.8.2	Proof of Theorem 13.17	256
13.9	Discussion	256
	Summary	257

V Outlook	259
14 Tool Support	261
14.1 Probabilistic model checkers	261
14.2 DynAbs	261
14.3 Scenario Approach for pMDPs	262
14.4 Differentiation of pRMCs	263
14.5 SLURF	263
14.6 Conditional Reachability in CTMCs	264
15 Conclusion and Future Work	265
15.1 Summary of Contributions	265
15.2 A Guide to Robust Verification Under Uncertainty	267
15.3 Limitations, Challenges, and Perspectives	268
15.3.1 Combining learning and verification	269
15.3.2 Integration with reinforcement learning	269
15.3.3 Partial observability	270
15.3.4 Exploiting structure in AI	270
15.3.5 Continuous-time models with nondeterminism	271
15.3.6 Mature tool support	271
15.4 Final Remarks	271
VI Back Matter	273
A Bibliography	275
B Index	307
C Research Data Management	309
D About the Author	311

1 Introduction

Motivation | Making predictions about complex systems is at the core of many decision-making problems. With the rise of artificial intelligence (AI), more and more of the predictions involved in these decision-making problems are automated and autonomous [RN10; ABRD⁺20]. For example, we want to predict whether an autonomous drone will safely reach its destination [GKM10], whether a manufacturing system will break down within the next month [TBBB⁺20; Mob02], whether a warehouse inventory will not run empty [GMT14], or whether a financial portfolio will yield a profit [CBSN⁺16]. Most of these problems are of a *sequential* nature, meaning that we need to repetitively make predictions and decisions that depend on each other. To make such predictions accurately, we need to create *mathematical models* that describe how systems behave. However, accurately modeling systems at the right level of abstraction is highly challenging [BK08], and no model will ever capture the behavior of the system it represents exactly. As the British statistician George E. P. Box once said: “*All models are wrong, but some are useful*” [Box76]. In this thesis, we postulate that eliminating all uncertainty from a model is generally impossible. Instead, we need to *embrace uncertainty in mathematical models* and develop methods that *rigorously account for it*.

Outline | This thesis is about analyzing mathematical models that are subject to uncertainty. In this introductory chapter, we discuss the motivation, scope, and contributions of the thesis. In Sect. 1.1, we introduce stochastic systems, and in particular Markov models, as the main models that we study. Then, in Sect. 1.2, we describe common questions we can ask about the behavior of a Markov model. In Sect. 1.3, we discuss what we mean by uncertainty and how to incorporate it into Markov models. Thereafter, in Sect. 1.4, we introduce the concept of robustness in analyzing models. In Sect. 1.5, we present the main research goals of this thesis, and in Sect. 1.6, we give an overview of some of the key techniques we use to achieve these goals. Finally, we present a reading guide for this thesis in Sect. 1.7 and list the origins of each chapter in Sect. 1.8.

1.1 Stochastic Systems

Mathematical modeling classically assumes that the current status of a system can be described by a set of variables, jointly called the *state* of the system. In our models, we wish to capture how the state of the system changes over time. As a simple example, consider an autonomous delivery drone flying straight towards a target at a constant speed of 100 km/h. In this simplified example, the state of the drone is described by just one variable, namely the total distance traveled. Using the laws of physics, we can derive a model that describes how the distance traveled changes over time, as illustrated by Fig. 1.1a. Using this model, we can, for example, analyze the distance the drone travels in one hour (which is, not surprisingly, 100 km).

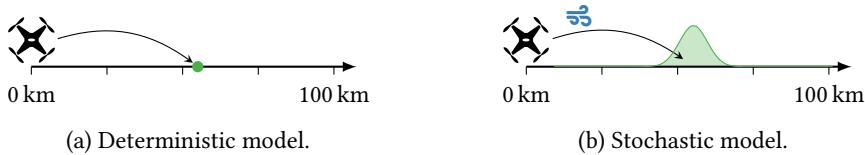


Figure 1.1: Simplified model of a drone flying in a straight line. In the left plot, the model is *deterministic*, so the distance traveled after 30 minutes is precisely known (green dot). In the right plot, the model is *stochastic*, so the distance traveled after 30 minutes is given as a probability distribution (green area).

Stochasticity | What if some variables influencing the drone's behavior depend on random events and are *stochastic*? A stochastic variable is one whose value is not deterministic but instead described by a *probability distribution*. For example, suppose that the wind acts as a disturbance on the speed of the drone, thus affecting the distance traveled within an hour. If each value for the wind speed occurs with a particular probability (for example, weak, medium, and strong wind all occurring with probability 1/3), then we can use these probabilities to derive a *probability distribution* over traveled distances. In practice, these likelihoods can, for example, be derived from historical weather data [WW99; PK08]. As a result, our model becomes a so-called *stochastic process*, and each time we simulate the model, we may obtain a different distance traveled in an hour. This idea is illustrated by Fig. 1.1b.

stochastic
process

Markov models | Slightly more formally, the mathematical model for the drone with stochastic wind can be formalized as a *Markov model*. A Markov model is a particular type of stochastic process where the future evolution of the state only depends on the current state, and not on the previous states (also known as the *Markov property* [Put94]). Why is this such a nice property to have? Because when analyzing a Markov model, we only have to look at the *current state* of the system and not at all the *previous states* in which the system has been.

Markov models appear in various areas, including control theory [Ast12], operations research [Dav18], artificial intelligence [RN10], systems biology [All10], and many more. Well-known examples of Markov models that we study in this thesis are Markov chains, Markov decision processes (MDPs), and dynamical control systems. The following example is the first instance of a Markov model—namely a discrete-time Markov chain (DTMC)—that we will encounter in this thesis.

Example 1.1 (Hello, Markov chain!) We simplify the drone example even further by discretizing time into steps of 30 minutes. Suppose that if the wind is weak, the drone travels 55 km per 30 minutes. Similarly, the travel distance per 30 minutes is 50 km for medium wind and 45 km for strong wind. Assuming the same probabilistic model for the wind as above (namely that weak, medium, and strong wind all occur with probability 1/3), we can model the evolution of the cumulative distance traveled. Specifically, starting from an initial distance of 0 km, after 30 minutes the drone has traveled either 45, 50, or 55 km, each with probability 1/3. After 60

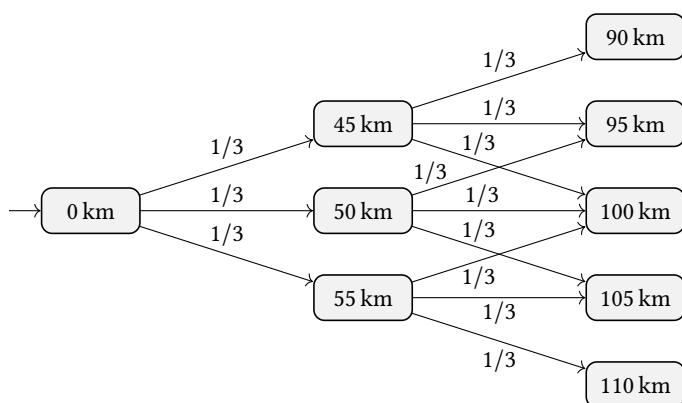


Figure 1.2: Discrete-time Markov chain for the distance traveled by the drone. Each box is a possible state (the total distance traveled), each edge represents a time duration of 30 minutes, and the numbers on the edges represent the probabilities for transitions between states to occur.

minutes, the drone has traveled 90, 95, 100, 105, or 110 km, but now the probabilities are no longer uniform. This model is depicted in Fig. 1.2 and is a simple example of a (discrete-time) Markov chain—a type of Markov model we formally define in Chapter 3. Loosely speaking, this Markov chain consists of:

- a set of *states* $S = \{0, 45, 50, 55, 90, 95, 100, 105, 110\}$, each of which represents a total distance traveled,
- an *initial state* $s_I = 0$, and
- a *transition function*, which describes the probability of transitioning between any two states (visualized by the numbers on the edges in Fig. 1.2).

Action choices | Often, Markov models describe *controllable elements* present in the underlying system. A controllable element in the drone example can be that a human at a control center can accelerate or decelerate the drone, thus influencing the drone's behavior. These controllable elements are modeled as *actions* (or *control inputs*) and render the Markov model *nondeterministic*. The best-known example of a Markov model with nondeterminism is the Markov decision process (MDP). Nondeterminism means that the action choices are not determined by the model itself, but by an external decision rule called the *policy* (also called *controller* or *scheduler*).

1.2 Verification

Research in areas such as probabilistic model checking, operations research, and control theory has led to the development of a wealth of approaches for analyzing Markov models. These approaches can roughly be divided into two categories. First, *qualitative* verification considers the question: “*Does a system satisfy a certain set of requirements?*” On the other hand, *quantitative* verification asks: “*How well does a system satisfy a certain set of requirements?*” Let us discuss both types of verification in more detail.

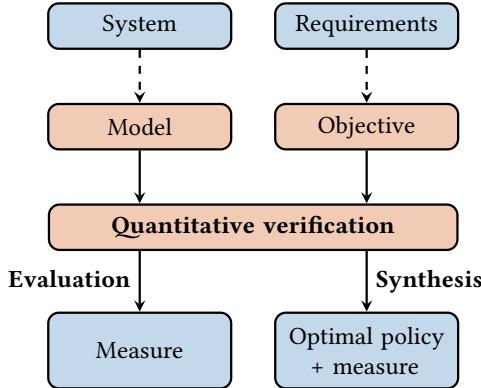


Figure 1.3: Ingredients of the two types of quantitative verification problems, namely evaluation and synthesis problems.

Qualitative verification | A classical qualitative verification problem has two ingredients: a Markov model representing an underlying system, and a logical specification representing the requirements for this system [BK08]. The qualitative verification (also called model checking) problem is then to check whether the model satisfies the specification. This problem leads to a yes/no answer and is thus of a qualitative nature. For instance, for the Markov chain from Example 1.1, we can consider the specification: “*The distance traveled within one hour never exceeds 110 km.*” This specification is satisfied as the maximum distance traveled is precisely 110 km. We can also pose a probabilistic variant of this specification, such as “*with a probability of at least 0.9, the distance traveled within one hour is at least 100 km.*” This specification is not satisfied: The probability of traveling 90 km is $\frac{1}{3} \cdot \frac{1}{3} = \frac{1}{9}$, and the probability of traveling 95 km is $\frac{1}{3} \cdot \frac{1}{3} \cdot 2 = \frac{2}{9}$, so already with probability $\frac{1}{3}$, the distance traveled is less than 100 km. Algorithms for solving qualitative verification problems essentially explore all possible states of the model in a brute-force manner. In this way, we can show that a given system formally satisfies the modeled system requirements.

Quantitative verification | Often, we do not just care about whether a set of requirements is satisfied, but we also care about *to what extent* a specific set of requirements (or rather, a quantitative objective) is satisfied. For example, rather than asking whether the distance traveled never exceeds 110 km, we may ask what the expected travel distance within one hour actually is. This leads to a quantitative verification problem, where the answer is a *number* rather than a *yes/no* statement. Depending on whether or not the Markov model has controllable elements, we roughly distinguish between the two types of quantitative verification problems also shown in Fig. 1.3:

- **Evaluation:** If the Markov model has no controllable elements, then the quantitative verification problem becomes an *evaluation problem*. Intuitively, we *evaluate* the quantitative objective on the Markov model, resulting in a concrete number.
- **Synthesis:** If the Markov model has controllable elements, we obtain a so-called *synthesis problem*. A typical synthesis problem is to determine what policy should

be used, such that the quantitative objective is maximized (or minimized). Synthesis problems are analogous to control problems (in control theory) or planning problems (in AI and operations research).

In this thesis, we study quantitative verification problems for Markov models, focusing on both evaluation and synthesis problems.

1.3 Uncertainty is Inevitable

Ultimately, we want the solution to a verification problem to perform well on the actual system that the Markov model represents. Needless to say, this performance depends on how well the Markov model represents the underlying system. Unfortunately, as systems become increasingly more complex, it becomes harder to model them accurately. As a result, eliminating all uncertainty about the behavior of the underlying system is practically impossible. Instead, we need to develop methods that explicitly capture and rigorously account for uncertainty in models.

However, the above-mentioned verification approaches for Markov models classically assume that the model is precisely specified, without any uncertainty. For example, computing the probability that the travel time is at least 100 km requires knowing the exact transition probabilities of the Markov chain. This issue raises a fundamental research challenge that forms the motivation for this thesis:

How can we analyze Markov models that are subject to uncertainty?

What is uncertainty? | Before discussing how we can analyze models subject to uncertainty, we first need to define what we mean by uncertainty. Because Markov models are stochastic by definition, simulating them multiple times will lead to possibly different outcomes (which we call *executions*). For example, simulating the Markov chain from Example 1.1 can lead to the execution (0 km, 45 km, 95 km), but also to (0 km, 50 km, 105 km). Thus, we can say that the *execution* of a Markov model is uncertain. However, does this also mean that the Markov model itself is uncertain?

We argue that a *stochastic* model is not the same as an *uncertain* model. For example, the execution of the Markov chain in Fig. 1.2 is stochastic, but the model itself is not uncertain, as each of the transition probabilities of 1/3 is precisely known. Thus, we distinguish between uncertainty about the execution of a model and uncertainty about the model itself. For example, uncertainty about the execution of a model can be caused by uncertainty in:

- the current position of the drone (for example, due to sensor imprecision), or
- the future position of the drone (for example, due to the stochastic wind).

On the other hand, uncertainty about a Markov model includes uncertainty in:

- the probability for each wind speed (for example, due to limited weather data), or
- the physical parameters of the drone influencing the dynamics (for example, the mass of the drone, its drag coefficient, and so on).

Our focus in this thesis is mainly on dealing with uncertainty at the level of the model.

In this thesis, we primarily consider uncertainty about the *transition probabilities* of Markov models, and sometimes also about other elements of the model, such as the *initial state*.

How do we model uncertainty? | Now that we have defined what we mean by uncertainty, we can discuss how to incorporate uncertainty into Markov models. The general idea is to replace elements of the Markov model that are uncertain with *parameters* that can take on different values, as illustrated by the next example.

Example 1.2 (Parametric model) In the drone example, we assumed that the probability for each wind speed to occur is $1/3$. However, since we assumed that these weather conditions are approximated from historical weather data, these probabilities are, realistically speaking, only estimates. To acknowledge that the true probabilities may differ from our estimates, we can replace the precise transition probabilities of $1/3$ in Fig. 1.2 with *parameters* that can take on different values. Doing so results in the model depicted in Fig. 1.4a, which is a so-called *parametric Markov chain*. The parameter x is the probability of strong wind, y is the probability of weak wind, and thus, the probability of medium wind is automatically determined as $1 - x - y$ (as probabilities need to sum to one).

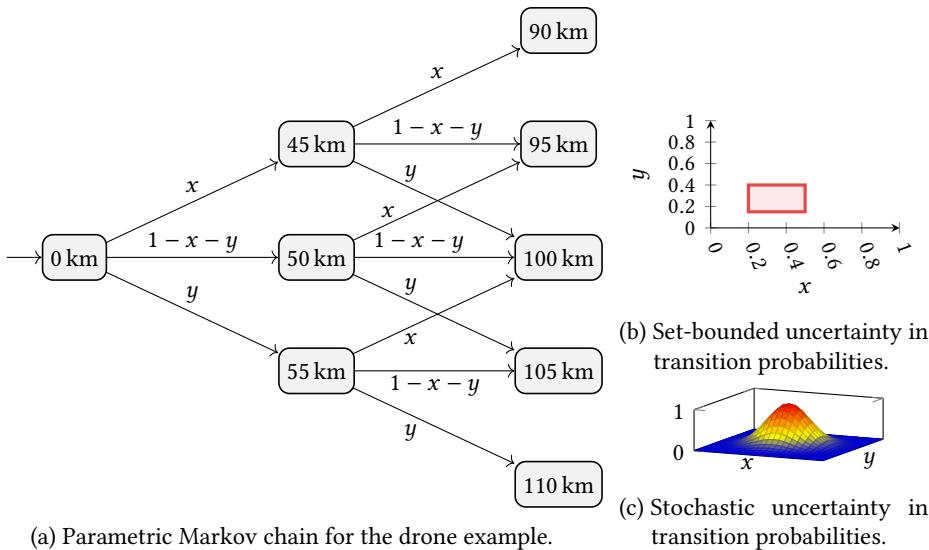


Figure 1.4: Capturing uncertainty about the probabilities for each wind speed in the drone example by introducing two parameters x and y . We can either consider the parameters (x, y) to be contained in the set in (b) or be drawn according to the distribution in (c).

Introducing parameters in a Markov model allows us to distinguish between two flavors when it comes to modeling uncertainty:

1. **Set-bounded uncertainty:** The first option is to consider all possible values the uncertain element can take, but without assigning likelihoods to each possible value. We achieve this by constraining each of the parameters to a *set of possible values*. We call the set of all allowed values for the parameters the *uncertainty set*. This perspective leads to a *set of possible models* but without likelihoods for each model in this set. In fact, this set of models is a simple example of a *robust Markov chain*; a type of Markov model we formally define in Chapter 3.
2. **Stochastic uncertainty:** The second option is to consider a probability distribution over the values of the parameter modeling the uncertainty. We achieve this by assigning a probability distribution (such as a normal distribution or a uniform distribution) to the parameters introduced in the Markov model. This perspective leads to a *distribution over possible models*, i.e., a set of possible models together with a likelihood for each model in this set. This type of model, which we call an *uncertain parametric Markov chain*, is central to Chapter 9 of this thesis.

uncertainty set

The set-bounded perspective can especially be useful to model parameters that can be *controlled*. By contrast, in the stochastic perspective, the parameter values are purely governed by the distribution. Let us illustrate both perspectives on uncertainty with the following example for the parametric Markov chain from Fig. 1.4a.

Example 1.3 (Modeling uncertainty) Using the set-bounded perspective, we can define the *set of possible values* (that is, the uncertainty set) for the parameters (x, y) shown in Fig. 1.4b. In this example, x (the probability of strong wind) is at least 0.2 and at most 0.5, such that $0.2 \leq x \leq 0.5$. Similarly, y (the probability of weak wind) is constrained to $0.15 \leq y \leq 0.4$. Setting a value for x and y automatically determines the probability of medium wind as $1 - x - y$.

By contrast, using the stochastic perspective, we define a *probability distribution over possible values* for the parameters (x, y) , as shown in Fig. 1.4c. In this example, the parameters are normally distributed. Thus, parameter values close to the mean of this distribution are more likely to occur, but other values are still possible.

We study both the set-bounded and stochastic perspectives on uncertainty. As we will see, both of these perspectives enable a different perspective on decision-making.

Where does uncertainty come from? | A popular classification of uncertainty is to distinguish between *aleatoric* and *epistemic* uncertainty [FU11; Sul15; CSKG22]. Aleatoric uncertainty is caused by randomness, whereas epistemic uncertainty is caused by a lack of knowledge of, for example, system parameters [Smi14]. While aleatoric is *irreducible*, epistemic uncertainty is *reducible* by collecting more data.

We argue that uncertainty classification (aleatoric vs. epistemic) is largely independent of how to model the uncertainty (set-bounded vs. stochastic). That is, aleatoric and epistemic uncertainty can both be modeled as either set-bounded or stochastic uncertainty in a Markov model. In this thesis, we are more interested in how to model and deal with uncertainty in Markov models, than in the source of the uncertainty. Thus, we do not further investigate aleatoric versus epistemic uncertainty.

1.4 Robustness

Now we have defined what uncertainty is and how to capture it in Markov models, we can discuss how to actually deal with the uncertainty. In general, we want to analyze the behavior of a Markov model with uncertainty, such that the analysis outcome is *robust against different outcomes of the uncertainty*. The term “*robustness*” has been used a lot in the literature on control theory [ZD98], operations research [BGN09; BH22], and artificial intelligence [Die17; Mar20]. Informally, a robust solution does not just perform well under a single scenario, but under a set of different scenarios. As we will see next, robustness against set-bounded and stochastic uncertainty both have a different interpretation.

Robustness for set-bounded uncertainty | Generally, a solution to a problem is robust (against a particular source of set-bounded uncertainty) if that solution performs properly for *all possible outcomes* of the uncertainty. For example, in robust optimization [BH22; BGN09], the idea is to constrain certain parameters to a compact set (called an *uncertainty set*) and compute a solution to the optimization problem that is feasible for all values of these parameters in the uncertainty set. This perspective is similar to the setting in Fig. 1.4, where we constrained the parameters (x, y) of the Markov chain to the set in Fig. 1.4b. For each possible value of the parameters (x, y) , we could compute the average distance traveled within one hour. A robust verification problem is then to compute the minimum (or maximum) average distance traveled over all possible values of (x, y) in the uncertainty set.

Example 1.4 (Robustness to set-bounded uncertainty) We again consider Fig. 1.4 with the set-bounded uncertainty in (x, y) described by the set in Fig. 1.4b. In this example, we have that $0.2 \leq x \leq 0.5$ and $0.15 \leq y \leq 0.4$ (and thus, we implicitly have that $0.1 \leq 1 - x - y \leq 0.65$). We ask the question: “*What is the minimum average distance traveled within one hour, over all possible values of (x, y) in the uncertainty set?*” This minimum average distance is attained when we choose x as high as possible (which is 0.5) and y as low as possible (which is 0.15), which automatically means that $1 - x - y$ is 0.35. For these parameter values, we find that the minimum average distance traveled is 96.5 km, which is computed as the weighted average of the distances (weighted by their probabilities):

$$\begin{aligned} & 90 \cdot (0.5 \cdot 0.5) + 95 \cdot (0.5 \cdot 0.35 + 0.35 \cdot 0.5) \\ & + 100 \cdot (0.5 \cdot 0.15 + 0.35 \cdot 0.35 + 0.15 \cdot 0.5) \\ & + 105 \cdot (0.35 \cdot 0.15 + 0.15 \cdot 0.35) + 110 \cdot (0.15 \cdot 0.15) = 96.5. \end{aligned}$$

This answer of 96.5 km is robust against the set-bounded uncertainty: No matter what values for the parameters (x, y) we choose within their allowed set, the average distance traveled will always be at least 96.5 km.

Robustness for stochastic uncertainty | For the case of stochastic uncertainty, we need to incorporate the likelihood of each possible realization of the uncertainty. Rather than being robust against *all* possible values of the uncertainty, we want to be robust

against a set of values of the uncertainty that occur with a certain probability, such as 99%, as illustrated by the next example.

Example 1.5 (Robustness to stochastic uncertainty) The distribution over the parameters (x, y) shown in Fig. 1.4c consists of two independent normal distributions for x and y .^a Suppose the distribution for x has a mean of 0.4 and a standard deviation of 0.1, and the distribution for y has a mean of 0.2 and a standard deviation of 0.05. We want to compute the minimum average distance traveled, robust against 99% probability mass of the joint distribution over the parameters (x, y) . We can obtain a solution to this problem in two steps. First, we compute a set that, with probability 0.99, contains the values for the parameters (x, y) when sampled from the joint distribution. Since x and y are independently distributed, we can compute a set with probability mass of $\sqrt{0.99} \approx 0.995$ for both parameters, and combine them to obtain a box set similar to the one in Fig. 1.4b. Second, once we have obtained this set, we can proceed as in Example 1.4.

^aIn fact, the distributions are truncated at 0 and 1 to avoid probabilities that are impossible.

In general, the set-bounded perspective leads to more conservative solutions than the stochastic perspective. Thus, the set-bounded perspective is typically more appropriate for safety-critical settings, where little or no risk is acceptable, whereas the stochastic perspective is better suited for settings where less conservative solutions are preferred. However, computing the set that contains a certain probability mass of the distribution over parameters is often challenging, especially because this set is not unique. In Example 1.5, we can either determine the set by taking the 99% probability mass around the mean of the distribution, but we can also take an asymmetric approach and instead shift the set to one of the tails. Thus, which of the two perspectives is most appropriate depends on the specific problem at hand.

1.5 Challenges and Contributions

The previous examples about uncertainty and robustness in Markov models help to illustrate the problem but are still very simplistic. In more realistic applications, models can consist of millions or even infinitely many states and actions. Moreover, sets and distributions over uncertain parameters can be very complex and high-dimensional, or may not even be known at all. In such cases, classical algorithms for the quantitative verification of Markov models are no longer applicable. Thus, the overall research goal of this thesis is as follows.

The goal of this thesis is to develop novel methods for quantitative verification of Markov models with uncertainty, which can be used to provide rigorous guarantees that are robust against the model's uncertainty.

We now break down this overall research goal into four key challenges in robust decision-making under uncertainty for Markov models. Thereafter, we discuss how we aim to address each challenge in this thesis.

Challenge 1: Robust policy synthesis for uncertain Markov models with continuous state and action spaces (Part II of this thesis).

Classical model checking algorithms rely on exhaustively exploring all states of a Markov model and are thus restricted to models with finitely many states and actions. However, dealing with real-world physical systems requires models with *continuous* state and action spaces. While the field of stochastic control deals with such continuous models, these approaches can typically only handle simple specifications of, for example, stability and convergence. More realistic applications often require more complex specifications, which reason over the *temporal evolution* of the system. Furthermore, existing approaches for continuous-state/action models often assume that the model is precisely known, without any uncertainty. Thus, there is a lack of approaches for robust control in uncertain Markov models with continuous state and action spaces, and with more complex specifications.

We address this first challenge for a particular class of continuous-state/action Markov models, called *discrete-time stochastic systems* (DTSSs). We focus on objectives that go beyond the simple tasks usually considered in control theory, such as stability and convergence. Instead, we consider richer objectives that are formalized in so-called temporal logics [Pnu77]. We, in particular, focus on *reach-avoid control tasks*, where the goal is to reach a set of target states (within a given time horizon) while always avoiding a set of unsafe states. As our first contribution, we develop a novel framework for computing policies for DTSSs that provably satisfy complex temporal specifications, and that are robust against set-bounded parameter uncertainty.

Challenge 2: Data-driven verification of uncertain Markov models with prior knowledge (Part III of this thesis).

Most approaches for the verification of Markov models are *model-based*. That is, these approaches rely on an explicit description of the model and its uncertainty. However, in many realistic applications, the model or some of its elements may not be available in explicit form. In Example 1.5, we assumed that we knew the distribution over the parameters (x, y) , but what if this distribution is unknown? Similarly, what if the transition probabilities of a Markov model are unknown and instead must be learned by interacting with the system? In reality, assuming no prior knowledge about a system at all may be too strong. Instead, we often have some prior knowledge about the system. We can encode this prior knowledge in a Markov model with a fixed parametric structure, such as in Fig. 1.4a. With this parametric structure, we can encode dependencies between different elements of the model. However, data-driven approaches for uncertain Markov models whose parametric structure is known are severely lacking.

In addressing this second challenge, we develop novel data-driven approaches for the verification of uncertain Markov models whose parametric structure is known. In particular, our approaches leverage the parametric structure of the model as prior knowledge, which leads to more efficient and less conservative solutions than neglecting these parametric structures completely.

Challenge 3: Robust verification for continuous-time Markov chains with uncertainty (*Part IV of this thesis*).

Over the past decades, various versions of Markov models with uncertainty in the transition probabilities have been developed, such as interval MDPs (IMDPs) and robust MDPs (RMDPs). These models inherently evolve over *discrete* time steps; however, accurately modeling physical systems often requires a *continuous-time* perspective. In this thesis, we specifically focus on continuous-time Markov chains (CTMCs), which are the continuous-time variant of the discrete-time Markov chains discussed earlier in this chapter. Approaches for CTMCs with uncertainty are significantly less developed than their discrete-time counterparts. Thus, the verification of CTMCs with uncertainty, for example in the transition probabilities (also called transition rates), remains a challenging task.

To address this third challenge, our next main contribution is to develop novel approaches for verifying CTMCs with uncertainty. We consider two distinct settings of uncertainty: (1) CTMCs with distributions over the transition rates, and (2) CTMCs where the initial state is unknown and must instead be estimated from previous observations. For both settings, we develop an algorithm that solves the corresponding verification problems, while providing rigorous mathematical guarantees.

Challenge 4: Tool support for analyzing Markov models with uncertainty (*all parts of this thesis*).

Mature tool support exists for analyzing Markov models without uncertainty, with Storm [HJKQ⁺22] and PRISM [KNP11] being two notable examples that we prominently use throughout this thesis. Most of the algorithms are based on variants of *dynamic programming* [Bel66], such as value iteration algorithms [SB98]. However, these tools are largely restricted to Markov models without uncertainty, with the exception of interval MDPs, for which support has recently been added. Thus, in addressing the previous three challenges, we also aim to develop prototypical tools that implement the algorithms proposed in this thesis and can be used as a basis for further research.

Our contribution to addressing this fourth challenge is to develop a set of tools that implement the algorithms proposed in this thesis. Our goal is not to develop full-fledged commercial tools but instead to develop prototypical tools that other researchers can use as a basis for further research. All of our implementations are open-source, archived in public repositories, and documented. By doing so, we strive to maximize the reusability of our implementations by other researchers who aim to advance the field of robust verification of stochastic systems.

1.6 Overview of Key Techniques

We now provide a brief introduction to four key techniques that we use prominently throughout this thesis. Our goal here is to provide high-level and accessible intuition about these techniques while postponing technical details to the relevant chapters.

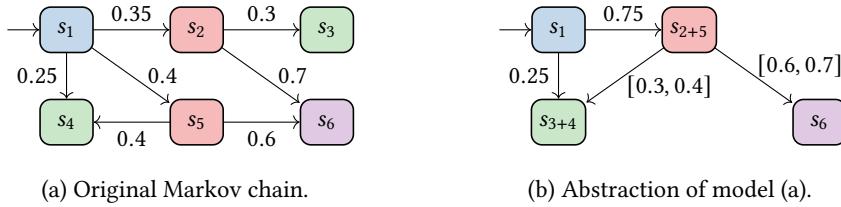


Figure 1.5: A Markov chain and its abstraction in the form of an interval Markov chain, obtained by merging states with the same color.

Probabilistic model checking

As already discussed, the verification of Markov models involves checking whether a model satisfies a given specification, or to what extent a specification is satisfied [BK08]. Probabilistic model checkers are tools that implement such verification methods. Throughout this thesis, we prominently use both the model checkers Storm [HJKQ⁺22] and PRISM [KNP11] to analyze (discrete-time and continuous-time) Markov chains, and Markov decision processes and their extension with parametric and interval-valued probabilities. Since we already discussed the quantitative verification problems that we will consider in this thesis (namely, evaluation and synthesis problems), we will not go into more detail here.

Model abstraction

abstraction

Abstraction is the process of generalizing details away from a model under study to focus attention on details of greater importance [Kra07]. In the same spirit, we can also say that one model is an *abstraction* of another model. Abstraction is a key concept in computer science [CS07] and, in particular, in model checking [BK08; Her02] and model learning [Vaa17]. In this thesis, we prominently use abstraction to simplify Markov models whose analysis is intractable or infeasible into Markov models that we can analyze. For example, the discrete-state Markov chain from Fig. 1.2 is an abstraction of the continuous-state representation for the drone dynamics in Fig. 1.1. Abstraction is particularly useful for Markov models with continuous state and action spaces: Analyzing continuous Markov models is generally intractable, while their finite abstractions can be analyzed efficiently with standard methods from probabilistic model checking.

Abstraction errors | However, the discrete-state Markov chain from Fig. 1.2 does not capture the original continuous dynamics *exactly*. For example, in the continuous model, we could ask for the distance traveled at any time $t \geq 0$, but in the discretized Markov chain, we only modeled time in steps of 30 minutes. In other words, we incurred an *abstraction error* by converting the continuous model into a discrete one. Due to abstraction errors, we cannot just expect the analysis outcomes on the abstract model to carry over to the original, continuous system.

Sound abstraction | In practice, we often do not want an abstraction that simply *approximates* the behavior of a model. Instead, we want that analyzing the abstraction provides a *sound* result for the original model, often in the form of an upper or lower bound on the analysis outcome. To achieve this, we typically want an abstraction that

either *overapproximates* or *underapproximates* the behavior of a model. For example, consider the Markov chain in Fig. 1.5a and its abstraction in Fig. 1.5b with intervals of transition probabilities, where we merged the states with the same colors. This abstraction is a simple example of a so-called *interval Markov chain*, which is a variant of a Markov chain where the transition probabilities are given as intervals.

Now consider the verification problem of computing the probability to reach state s_6 . In the original Markov chain, this probability is $0.35 \cdot 0.7 + 0.4 \cdot 0.6 = 0.485$. For the abstraction, we can perform a *robust analysis* similar to the analysis in Example 1.4, to compute the minimal probability of reaching s_6 for any choice of probabilities. This minimal probability of reaching s_6 is obtained by choosing the lower bound probability of 0.6 in the interval $[0.6, 0.7]$. Thus, we obtain a robust reachability probability of $0.75 \cdot 0.6 = 0.45$, which is a lower bound on the outcome of 0.485 for the original Markov chain. Hence, we say that, for this verification problem, the interval Markov chain in Fig. 1.5b is a *sound abstraction* of the original Markov chain in Fig. 1.5a.

Robust optimization

Robust optimization is a branch of mathematical optimization that deals with problems where certain parameters are only known up to a given uncertainty set [BH22; RM19; BBC11; BGN09]. Let Δ denote this uncertainty set. A typical robust optimization problem has the following form:

$$\begin{aligned} & \underset{v \in \mathbb{R}^n}{\text{minimize}} \quad f(v) \\ & \text{subject to} \quad g(v, \delta) \leq 0 \quad \forall \delta \in \Delta, \end{aligned} \tag{1.1}$$

where $f(v)$ is a cost function (or objective function), which is a function of the n decision variables $v \in \mathbb{R}^n$, and $g(v, \delta)$ is a constraint that depends on both v and the uncertainty variable δ . Intuitively, solving Eq. (1.1) amounts to finding values for v (which is a vector of n decision variables) such that $f(v)$ is minimized, while at the same time ensuring that the constraint $g(v, \delta)$ is satisfied for this choice of v and for all values of the uncertainty variable $\delta \in \Delta$. A more graphical interpretation of this optimization problem can be found in [CG18a, Chapter 1].

Solving robust optimization problems | The optimization problem in Eq. (1.1) is called a robust optimization problem because the constraint is enforced *for all* values of the uncertainty variable $\delta \in \Delta$. If f and g , as well as the uncertainty set Δ , have a “nice” structure,¹ then the robust optimization problem can be solved efficiently.

Application in verification | Robust optimization is tightly connected to verifying Markov models with uncertain transition probabilities [GLD00; NG05; XM10; WKR13]. In particular, several common verification problems for robust Markov chains and robust MDPs (which we formally introduce in Sect. 3.3) can be formulated as robust optimization problems [PLSS13]. Intuitively, the uncertainty set Δ of the robust optimization problem represents the set of possible transition probabilities of the Markov model. Then, solving this robust optimization problem corresponds to finding the minimal (or maximal) value

¹For example, when both f and g are affine and Δ is a convex polytope, then the problem reduces to a linear program, which can be solved efficiently; see, e.g., the books [BH22; BV14] for details.

of, for example, the probability of reaching a target state, over all possible transition probabilities in the uncertainty set.

The scenario approach

Robust optimization does not assume any structure over the uncertainty set Δ and always assumes that $\delta \in \Delta$ takes the *worst possible value*. This assumption can be too conservative in practice. Moreover, in many cases, we actually have some knowledge about the likelihoods of different $\delta \in \Delta$. We can capture this knowledge in a *probability distribution* \mathbb{P} over the set Δ , which assigns a probability for each $\delta \in \Delta$ to occur.² Instead of enforcing the constraint $g(v, \delta)$ for all $\delta \in \Delta$ (as we did in the robust optimization problem), we will allow for a *small probability* that the constraint is violated. Let ε be the maximal probability by which we allow the constraint $g(v, \delta)$ to be violated. Then, we consider the following optimization problem, which is commonly called a *chance-constrained problem*:

$$\begin{aligned} & \underset{v \in \mathbb{R}^n}{\text{minimize}} \quad f(v) \\ & \text{subject to } \mathbb{P}(\delta \in \Delta : g(v, \delta) \leq 0) \geq 1 - \varepsilon. \end{aligned} \tag{1.2}$$

Solving this chance-constrained problem can be interpreted as follows. We want to find a value $v \in \mathbb{R}^n$ for the decision variables that minimizes $f(v)$. However, at the same time, we must choose v such that, for at least a $1 - \varepsilon$ probability of the values $\delta \in \Delta$, the constraint $g(v, \delta)$ is satisfied. A graphical interpretation of this type of solution can be found in [CG18a, Chapter 1].

Performance vs. risk | In the chance-constrained problem, the parameter ε is a tuning knob that allows for a trade-off between performance (in terms of a low cost $f(v)$) and risk (in terms of the probability that the constraint $g(v, \delta)$ is violated). In other words, the parameter ε represents the level of risk we are willing to take. If we pick $\varepsilon = 1$, then we completely neglect the constraint. If, on the other extremum, we pick $\varepsilon = 0$, then the chance-constrained reduces to a robust problem as in Eq. (1.1).³

Solving chance-constrained problems | In general, solving chance-constrained optimization problems is very difficult.⁴ This is where the so-called *scenario approach* comes in, a methodology that became popular due to [CC05; CG08]. Since then, many important extensions of the theory have been developed, such as [CGP09; CG11; ESL15; GZMG⁺16; CG18b; GC22; RPM23]. Intuitively, the scenario approach assumes that we have access to a set of N samples of the uncertainty variable, $\{\delta_1, \dots, \delta_N\}$, drawn independently from the probability distribution \mathbb{P} . Instead of the chance-constrained problem in Eq. (1.2), we then formulate the following *scenario (optimization) problem*:

$$\begin{aligned} & \underset{v \in \mathbb{R}^n}{\text{minimize}} \quad f(v) \\ & \text{subject to } g(v, \delta_i) \leq 0 \quad \forall i = 1, \dots, N. \end{aligned} \tag{1.3}$$

²We formalize such a setting further in Chapters 9 and 12.

³To be more precise, this reduction is only up to values of $\delta \in \Delta$ with a nonzero probability.

⁴A detailed discussion of why solving chance-constrained problems is difficult is beyond our scope, and we refer the interested reader to the papers [Pré03; RS03] or the book [Pré13] instead.

scenario
approach

scenario
problem

Observe that the scenario optimization problem only enforces the constraint $g(v, \delta_i)$ for each sample δ_i , where i ranges from 1 to N . Thus, given that f and g are convex functions in the decision variables v , the scenario optimization problem can be solved efficiently with convex optimization techniques [BV14]. Now, the main beauty of the scenario approach is that the solution to the scenario optimization problem is, with at least a certain probability $1 - \beta$, also a feasible solution to the chance-constrained optimization problem in Eq. (1.3). The parameter β is called the *confidence level* and depends on the number of samples N we used, and on the violation probability ε we considered. The theory of the scenario approach gives us a formula that relates the sample size N , the violation probability ε , and the confidence level β [CG08; CG11].

The scenario approach in this thesis | We use the scenario approach in several parts of this thesis. First, we use the scenario approach in Chapters 6 and 7 for estimating upper and lower bounds on transition probabilities of abstractions of continuous-state Markov models. Second, we use the scenario approach in Chapter 9 for verifying parametric MDPs where we have access to a probability distribution over the parameter values. Finally, we consider a similar setting in Chapter 12 for parametric continuous-time Markov chains with distributions over parameter values.

1.7 Navigating This Thesis

This thesis consists of the five main parts shown in Fig. 1.6. Some (sub)sections of the chapters are preceded with an asterisk (*), meaning that these sections contain mathematical details or proofs that can be safely skipped by the reader.

Of course, you are most welcome to read everything from this thesis. However, if you are interested in a specific topic, you can also read most of the main parts independently. The overall structure and dependencies between the parts and chapters of this thesis are sketched in Fig. 1.6.

We briefly summarize each of the chapters. Where applicable, we list the publications (see Sect. 1.8 for the full references) on which each of the chapters is based.

Part I: Foundations

- **Chapter 2. Preliminaries**

We introduce the mathematical notation used throughout this thesis, and we recap several important concepts and definitions from probability theory.

- **Chapter 3. A Primer on Markov Decision Processes**

We give a primer on MDPs and discuss how to perform common analyses for these models. Furthermore, we discuss generalizations of MDPs to uncertain transition probabilities, particularly as so-called robust and interval MDPs.

Part II: Discrete-time stochastic systems

In this first main part of the thesis, we study discrete-time stochastic systems (DTSSs)—a particular type of Markov model with continuous state and action spaces. We focus on *reach-avoid control tasks*: Reach a desirable target state (within a given time limit), while always avoiding unsafe states in the meantime.

- **Chapter 4. Foundations of DTSSs**

We introduce the fundamentals of DTSSs and discuss probabilistic reach-avoid control tasks as the main control objectives that we study. Moreover, we show that, due to the continuous and stochastic nature of DTSSs, computing the probability that a given policy satisfies a reach-avoid task is intractable in general.

- **Chapter 5. Probabilistic Simulation Relations** (based on [1; 6; 7])

We solve control tasks for DTSS based on an abstraction into a (discrete) MDP. We discuss the requirements for this MDP to be a *sound abstraction* of the DTSS (in the sense discussed in Sect. 1.6). These requirements are captured in a mathematical relation between the DTSS and the MDP, called a *probabilistic simulation relation*. We use the MDP abstraction to compute a guaranteed lower bound on the probability that a reach-avoid task is satisfied by the DTSS.

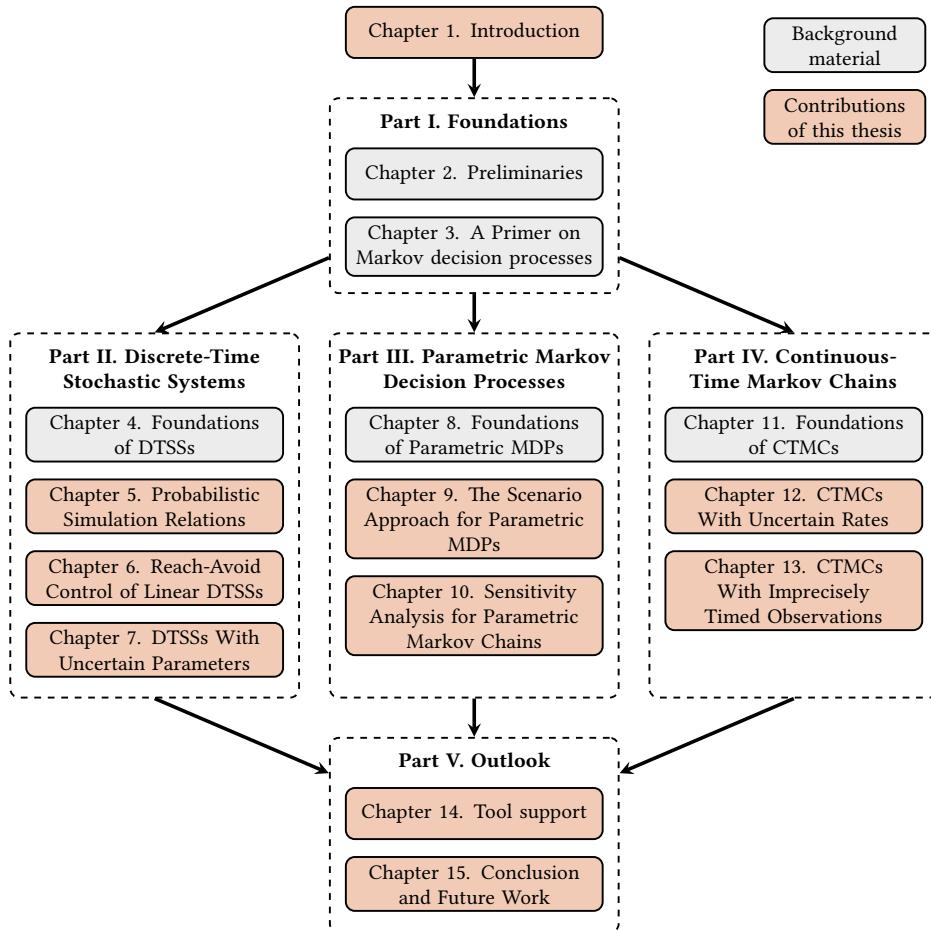


Figure 1.6: The overall structure of this thesis, highlighting its five main parts. Chapters that contain novel contributions are shown in red, while the chapters that contain background material are shown in gray.

- **Chapter 6. Reach-Avoid Control of Linear DTSSs** (based on [1; 7; 5; 10])
To use the results from Chapter 5 to solve a reach-avoid control problem for a DTSS, we need to construct a finite-state abstraction in the form of an MDP that creates a probabilistic simulation relation. In this chapter, we present a tractable algorithm to find such an MDP abstraction. We focus on a particular class of DTSS, where the dynamics are linear in the state and control input.
- **Chapter 7. DTSSs With Uncertain Parameters** (based on [6])
We study a variant of DTSS with set-bounded uncertain parameters. Specifically, these uncertain parameters are assumed to lie in a convex uncertainty set, but we do not assume any probability distribution over these parameters. In this chapter, we extend our abstraction framework and the notion of probabilistic simulation relations to such DTSS with uncertain parameters.

Part III: Parametric Markov decision processes

In this part, we focus on verification problems for parametric Markov decision processes (pMDPs), which are a generalization of MDPs to transition probabilities given as polynomial functions over parameters.

- **Chapter 8. Foundations of Parametric MDPs**
We describe the fundamentals of pMDPs and discuss how, by fixing a value for the parameters of the pMDP, we obtain a standard MDP that we can analyze using the methods from Chapter 3.
- **Chapter 9. The Scenario Approach for Parametric MDPs** (based on [2])
We study a setting similar to Example 1.5, where we are given a pMDP and a distribution over the parameters. However, in contrast with this example, we do not assume that this distribution over the parameters is known. The problem is to compute the probability that the pMDP satisfies a given specification when the parameter values are sampled according to this unknown distribution. We present a data-driven method based on the *scenario approach* to obtain a rigorous solution to this problem.
- **Chapter 10. Sensitivity Analysis for Parametric Markov Chains** (based on [4])
We present a novel and efficient method for sensitivity analysis of parametric robust Markov chains. These models incorporate parameters and sets of probability distributions to alleviate the often unrealistic assumption that precise probabilities are available. Using techniques from convex optimization, we perform a sensitivity analysis in terms of partial derivatives with respect to the uncertain transition probabilities. Our approach can be used to reduce the sample complexity in a model-based learning setting.

Part IV: Continuous-time Markov chains

In this part, we study verification problems for CTMC, which are the continuous-time counterpart of standard (discrete-time) Markov chains.

- **Chapter 11. Foundations of CTMCs**

We describe the fundamentals of CTMCs, and we discuss how to express and compute common measures of interest. Furthermore, we introduce parametric CTMC (pCTMC) as a parametric extension similar to pMDP.

- **Chapter 12. CTMCs With Uncertain Rates** (based on [3])

We study pCTMCs with a probability distribution over the parameters. This probability distribution encodes uncertainty about the transition rates of the CTMC. We consider the following verification problem: From a finite set of parameter samples and a user-specified confidence level, compute prediction regions on the reachability probabilities. We provide a principled solution to this problem based on techniques from the scenario approach.

- **Chapter 13. CTMCs With Imprecisely Timed Observations** (based on [11])

We consider runtime monitoring for CTMCs. In such applications, we must incorporate past observations, whose timings may be uncertain. Thus, we consider a setting in which we are given a sequence of imprecisely timed labels called the evidence. We provide a principled approach to compute reachability probabilities, which we condition on this evidence. In other words, we compute a method for computing conditional probabilities for CTMCs with imprecisely timed observations.

Part V: Outlook

- **Chapter 14. Tool Support**

In this chapter, we provide a brief overview of our prototypical Python tools that implement the algorithms presented in this thesis.

- **Chapter 15. Conclusion and Future Work**

In this final chapter, we reflect on the four main challenges that we posed above. We discuss to what extent our contributions in this thesis have addressed each of these challenges. Furthermore, we present a guide to robust verification under uncertainty, with several questions that practitioners can ask when aiming to verify a Markov model with uncertainty. Finally, we present an outlook on future research directions which, we believe, are promising to pursue.

1.8 Overview of Publications

The following academic publications form the basis of the contributions presented in this thesis. The reference numbers listed here are also those used throughout the thesis to refer to these publications.

The papers below are the result of great team effort and dedication from all authors involved. Thus, attributing specific contributions of papers to the individual authors is generally impossible. Nevertheless, the Radboud University doctorate regulations request a specification of the contributions to the core papers of the thesis. For all papers where I am listed as the first author, I have been involved as a leading author in all stages of the research, including developing theory and methodology, literature research, writing, implementation. The only exception is our position paper [8], which is rather a joint survey in which we describe our group's research vision.

Cited publications

- [1] T. Badings, A. Abate, N. Jansen, D. Parker, H. A. Poonawala and M. Stoelinga. ‘Sampling-Based Robust Control of Autonomous Systems with Non-Gaussian Noise’. AAAI. AAAI Press, 2022, pages 9669–9678. doi: [10.1609/AAAI.V36I9.21201](https://doi.org/10.1609/AAAI.V36I9.21201).
- [2] T. Badings, M. Cubuktepe, N. Jansen, S. Junges, J. Katoen and U. Topcu. ‘Scenario-Based Verification of Uncertain Parametric MDPs’. *Int. J. Softw. Tools Technol. Transf.* 24.5 (2022), pages 803–819. doi: [10.1007/S10009-022-00673-Z](https://doi.org/10.1007/S10009-022-00673-Z).
- [3] T. Badings, N. Jansen, S. Junges, M. Stoelinga and M. Volk. ‘Sampling-Based Verification of CTMCs with Uncertain Rates’. CAV (2). Volume 13372. Lecture Notes in Computer Science. Springer, 2022, pages 26–47. doi: [10.1007/978-3-031-13188-2_2](https://doi.org/10.1007/978-3-031-13188-2_2).
- [4] T. Badings, S. Junges, A. Marandi, U. Topcu and N. Jansen. ‘Efficient Sensitivity Analysis for Parametric Robust Markov Chains’. CAV (3). Volume 13966. Lecture Notes in Computer Science. Springer, 2023, pages 62–85. doi: [10.1007/978-3-031-37709-9_4](https://doi.org/10.1007/978-3-031-37709-9_4).
- [5] T. Badings, H. A. Poonawala, M. Stoelinga and N. Jansen. ‘Correct-by-Construction Reach-Avoid Control of Partially Observable Linear Stochastic Systems’. ArXiv preprint (2023). doi: [10.48550/arXiv.2103.02398](https://doi.org/10.48550/arXiv.2103.02398).
- [6] T. Badings, L. Romao, A. Abate and N. Jansen. ‘Probabilities Are Not Enough: Formal Controller Synthesis for Stochastic Dynamical Models with Epistemic Uncertainty’. AAAI. AAAI Press, 2023, pages 14701–14710. doi: [10.1609/AAAI.V37I12.26718](https://doi.org/10.1609/AAAI.V37I12.26718).
- [7] T. Badings, L. Romao, A. Abate, D. Parker, H. A. Poonawala, M. Stoelinga and N. Jansen. ‘Robust Control for Dynamical Systems with Non-Gaussian Noise via Formal Abstractions’. *J. Artif. Intell. Res.* 76 (2023), pages 341–391. doi: [10.1613/JAIR.1.14253](https://doi.org/10.1613/JAIR.1.14253).
- [8] T. Badings, T. D. Simão, M. Suilen and N. Jansen. ‘Decision-Making Under Uncertainty: Beyond Probabilities’. *Int. J. Softw. Tools Technol. Transf.* 25.3 (2023), pages 375–391. doi: [10.1007/S10009-023-00704-3](https://doi.org/10.1007/S10009-023-00704-3).
- [9] L. Rickard, T. Badings, L. Romao and A. Abate. ‘Formal Controller Synthesis for Markov Jump Linear Systems with Uncertain Dynamics’. QEST. Volume 14287. Lecture Notes in Computer Science. Springer, 2023, pages 10–29. doi: [10.1007/978-3-031-43835-6_2](https://doi.org/10.1007/978-3-031-43835-6_2).
- [10] T. Badings, L. Romao, A. Abate and N. Jansen. ‘A Stability-Based Abstraction Framework for Reach-Avoid Control of Stochastic Dynamical Systems with Unknown Noise Distributions’. ECC. 2024. doi: [10.48550/arXiv.2404.01726](https://doi.org/10.48550/arXiv.2404.01726).
- [11] T. Badings, M. Volk, S. Junges, M. Stoelinga and N. Jansen. ‘CTMCs with Imprecisely Timed Observations’. TACAS (2). Volume 14571. Lecture Notes in Computer Science. Springer, 2024, pages 258–278. doi: [10.1007/978-3-031-57249-4_13](https://doi.org/10.1007/978-3-031-57249-4_13).

Further publications

Beyond the publications directly cited in this thesis, I have contributed to the following papers before and during my Ph.D.

- T. Badings, V. Rostampour and J. M. Scherpen. ‘*Distributed Building Energy Storage Units for Frequency Control Service in Power Systems*’. *IFAC-PapersOnLine* 52.4 (2019). CSGRES 2019, pages 228–233. doi: [10.1016/j.ifacol.2019.08.190](https://doi.org/10.1016/j.ifacol.2019.08.190).
- V. Rostampour, T. Badings and J. M. A. Scherpen. ‘*Buildings-to-Grid Integration with High Wind Power Penetration*’. *CDC*. IEEE, 2019, pages 2976–2981. doi: [10.1109/CDC40024.2019.9030242](https://doi.org/10.1109/CDC40024.2019.9030242).
- T. Badings and D. S. van Putten. ‘*Data Validation and Reconciliation for Error Correction and Gross Error Detection in Multiphase Allocation Systems*’. *Journal of Petroleum Science and Engineering* 195 (2020), page 107567. doi: [10.1016/j.petrol.2020.107567](https://doi.org/10.1016/j.petrol.2020.107567).
- V. Rostampour, T. Badings and J. M. A. Scherpen. ‘*Demand Flexibility Management for Buildings-to-Grid Integration with Uncertain Generation*’. *ENERGIES* 13.24 (2020). doi: [10.3390/en13246532](https://doi.org/10.3390/en13246532).
- T. Badings, A. Hartmanns, N. Jansen and M. Suilen. ‘*Balancing Wind and Batteries: Towards Predictive Verification of Smart Grids*’. *NFM*. Volume 12673. Lecture Notes in Computer Science. Springer, 2021, pages 1–18. doi: [10.1007/978-3-030-76384-8_1](https://doi.org/10.1007/978-3-030-76384-8_1).
- T. Badings, N. Jansen, L. Romao and A. Abate. ‘*Correct-by-Construction Control for Stochastic and Uncertain Dynamical Models via Formal Abstractions*’. *FMAS@iFM*. Volume 395. EPTCS. 2023, pages 144–152. doi: [10.4204/EPTCS.395.10](https://doi.org/10.4204/EPTCS.395.10).
- T. Badings, W. Koops, S. Junges and N. Jansen. ‘*Learning-Based Verification of Stochastic Dynamical Systems with Neural Network Policies*’. *CoRR* abs/2406.00826 (2024). doi: [10.48550/ARXIV.2406.00826](https://doi.org/10.48550/ARXIV.2406.00826).
- M. Nazeri, T. Badings, S. Soudjani and A. Abate. ‘*Data-Driven Yet Formal Policy Synthesis for Stochastic Nonlinear Dynamical Systems*’. Under submission (2024). URL: <https://arxiv.org/abs/2501.01191>.
- M. Suilen, T. Badings, E. M. Bovy, P. David and N. Jansen. ‘*Robust Markov Decision Processes: A Place Where AI and Formal Methods Meet*’. *Principles of Verification: Cycling the Probabilistic Landscape : Essays Dedicated to Joost-Pieter Katoen on the Occasion of His 60th Birthday, Part III*. Springer Nature Switzerland, 2024, pages 126–154. doi: [10.1007/978-3-031-75778-5_7](https://doi.org/10.1007/978-3-031-75778-5_7).
- F. Souza, T. Badings, G. Postma and J. Jansen. ‘*Integrating Expert and Physics Knowledge for Modeling Heat Load in District Heating Systems*’. Accepted to *IEEE Transactions on Industrial Informatics* (2025).

Part I

Foundations

2 Preliminaries

In this short chapter, we define the basic notation that we use throughout the thesis. Furthermore, we describe the basic concepts from convex optimization and probability theory that we use in this thesis.

2.1 Basic Notation

We start by defining the mathematical notation that we use in the thesis. As a starting point, we assume that the reader is familiar with basic concepts from real analysis [Rud⁺64] and linear algebra [Str23], referring to these books for details on definitions and notation.

Sets | As is standard, we use \mathbb{R} , \mathbb{Q} , and \mathbb{N} to denote the real, rational, and natural numbers (with 0 as a natural number), respectively. We use subscripts $\mathbb{R}_{\geq 0}$, $\mathbb{R}_{\leq 0}$, $\mathbb{R}_{> 0}$, and $\mathbb{R}_{< 0}$ to restrict the real numbers to nonnegative, nonpositive, strictly positive, and strictly negative numbers, respectively. We use the same subscript-notation to constrain the rational numbers \mathbb{Q} and real numbers \mathbb{N} . The cardinality of a discrete set X is $|X|$. We write closed and open intervals as $[a, b] \subseteq \mathbb{R}$ and $(a, b) \subseteq \mathbb{R}$, respectively, where $a, b \in \mathbb{R}$ and $a \leq b$.

For sets $Y_1, \dots, Y_n, n \in \mathbb{N}_{\geq 0}$, the generalized Cartesian product is written as $\times_{i=1}^n Y_i = Y_1 \times Y_2 \times \dots \times Y_n$. We interchangeably use the notations $i = 1, \dots, p$ and $i \in \{1, \dots, p\}$ to denote ranging i over the integers 1 to p .

Properties of sets | Let $B_\varepsilon(x) \subset \mathbb{R}^n$ be the open ball of size $\varepsilon > 0$ and centered at $x \in \mathbb{R}^n$ in the Euclidean space:

$$B_\varepsilon(x) = \{x' \in \mathbb{R}^n : |x - x'| < \varepsilon\}.$$

Given a set V , we can define several important related sets. The *interior* of a set $V \subset \mathbb{R}^n$, denoted by $\text{int}(V)$, is defined as the set of all interior points of V :

$$\text{int}(V) = \{v \in V : \exists \varepsilon > 0. B_\varepsilon(v) \subset V\}.$$

The *closure* of a set $V \subset \mathbb{R}^n$, denoted by $\text{closure}(V)$, is the set of all points $v \in V$ such that every open ball centered at v contains a point in V :

$$\text{closure}(V) = \{v \in \mathbb{R}^n : \forall \varepsilon > 0. B_\varepsilon(v) \cap V \neq \emptyset\}.$$

Finally, the *boundary* of a set $V \subset \mathbb{R}^n$, denoted by δV , is defined as the closure of V minus the interior of V :

$$\delta V = \text{closure}(V) \setminus \text{int}(V).$$

As a simple example, consider the set $W = (0, 1] \subset \mathbb{R}$. The interior of W is $\text{int}(W) = (0, 1)$, its closure is $\text{closure}(W) = [0, 1]$, and its boundary is $\delta W = \{0, 1\}$.

In this thesis, we will only consider the interior, boundary, and closure of sets defined on Euclidean spaces. Note that there are alternative yet equivalent ways to define these sets; see [Rud⁺64] for more details.

Vectors and matrices | We use $x \in \mathbb{R}^n$ to denote a real-valued column vector of size $n \in \mathbb{N}$. All vectors are column vectors unless stated otherwise. Similarly, we use $A \in \mathbb{R}^{n \times m}$ to denote a real-valued matrix with $n \in \mathbb{N}$ rows and $m \in \mathbb{N}$ columns.

The *identity matrix* of size $n \in \mathbb{N}$, i.e., the matrix with ones on the diagonal and zeros elsewhere, is written as $I_n \in \mathbb{R}^{n \times n}$.

The *diagonal matrix* $\text{diag}(x) \in \mathbb{R}^{n \times n}$ has vector $x \in \mathbb{R}^n$ on the diagonal and is zero elsewhere.

The *inverse* of a real-valued square matrix $A \in \mathbb{R}^{n \times n}$ is denoted by A^{-1} . The (Moore-Penrose) *pseudoinverse* B^\dagger of a real-valued matrix $B \in \mathbb{R}^{n \times m}$ satisfies four conditions: (1) $BB^\dagger B = B$, (2) $B^\dagger BB^\dagger = B^\dagger$, (3) $(BB^\dagger)^\top = BB^\dagger$, and (4) $(B^\dagger B)^\top = B^\dagger B$. While not all matrices have an inverse, the pseudoinverse always exists. For a square, nonsingular matrix $A \in \mathbb{R}^{n \times n}$, the inverse and the pseudoinverse are the same, i.e., $A^{-1} = A^\dagger$.

Functions | We write a function f from a set X to a set Y as $f: X \rightarrow Y$. We may also use the notation $f: x \mapsto f(x)$ to denote a mapping from $x \in X$ to $f(x) \in Y$. Note that we use different symbols \rightarrow and \mapsto for these two ways of defining a mapping.

A *partial map* (or *partial function*) f from a set X to a set Y is a function from a subset $S \subseteq X$ to Y and is written as $f: S \rightarrow Y$. For a (partial) function f , we denote the domain and support as $\text{dom}(f)$ and $\text{supp}(f)$, respectively. We prominently use partial maps to define Markov decision processes (MDPs) (see Def. 3.1) where only a subset of actions may be enabled in each state.

The *indicator function* $\mathbb{1}_Z: X \rightarrow \{0, 1\}$ over a set $Z \subseteq X$ is defined for all $x \in X$ as $\mathbb{1}_Z(x) = 1$ if $x \in Z$ and $\mathbb{1}_Z(x) = 0$ otherwise.

Convex polytopes | A (bounded) *convex polytope* $T \subset \mathbb{R}^n$ is defined as the intersection of a finite number of half-spaces, such that

$$T = \{x \in \mathbb{R}^n : Hx \leq b, \quad H \in \mathbb{R}^{m \times n}, b \in \mathbb{R}^m\}.$$

This half-space representation of a convex polytope is also known as the *H-representation*. A convex polytope T is equivalently described by its vertices v_1, \dots, v_q , $q \in \mathbb{N}_{>0}$, referred to as the *V-representation*. Concretely, the convex polytope T is the *convex hull* of its vertices $\{v_1, \dots, v_m\}$, which we denote by $\text{conv}(v_1, \dots, v_n) \subset \mathbb{R}^n$. Algorithms for switching between convex polytopes in H-representation and V-representation exist, such as the double description method [MRTT53], which is implemented in the C-library cddlib [Fuk21].

2.2 Optimization Problems

We prominently use (convex) optimization throughout this thesis. We briefly introduce the notation for optimization problems and discuss convexity, while referring to the book [BV14] for a comprehensive introduction. A typical *optimization problem* (also

called *optimization program*) is defined as

$$\begin{aligned} & \underset{v \in \mathbb{R}^n}{\text{minimize}} \quad f_0(v) \\ & \text{subject to} \quad f_i(v) \leq 0 \quad \forall i = 1, \dots, m, \end{aligned} \tag{2.1}$$

where $v \in \mathbb{R}^n$ are called the *decision variables*, $f_0: \mathbb{R}^n \rightarrow \mathbb{R}$ is the *objective function*, and $f_i: \mathbb{R}^n \rightarrow \mathbb{R}$ for $i = 1, \dots, m$ are the *constraint functions*.

Optimal solution | The *value* of the optimization problem in Eq. (2.1) for $v \in \mathbb{R}^n$ is $f_0(v)$. A choice $v \in \mathbb{R}^n$ is *feasible* if all constraints are satisfied, i.e., if $f_i(v) \leq 0$ for all $i = 1, \dots, m$. An optimal value of the optimization problem is the smallest value over all feasible $v \in \mathbb{R}^n$:

$$f_0^* = \inf \{f_0(v) : v \in \mathbb{R}^N, f_i(v) \leq 0 \forall i = 1, \dots, m, \}.$$

While the optimization problem in Eq. (2.1) minimizes the objective function, we can also maximize a function by minimizing its negation $-f_0(v)$.

Convexity | Convexity is a crucial property of optimization problems because it allows for efficient solution methods [BV14]. A *convex function* is defined as follows.

Definition 2.1 A function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is *convex* if

$$f(\alpha x + \beta y) \leq \alpha f(x) + \beta f(y)$$

for all $x, y \in \mathbb{R}^n$ and for all $\alpha, \beta \in \mathbb{R}$ such that $\alpha + \beta = 1, \alpha \geq 0, \beta \geq 0$.

convex function

Intuitively, a function is convex if the line segment between any two points on the graph of the function lies above the graph itself. An optimization problem is *convex* if the objective function f_0 and all constraint functions $f_i, i = 1, \dots, m$ are convex. Most of the optimization problems we encounter in this thesis are convex and can thus be solved efficiently with convex optimization solvers, such as Gurobi [Gur23] or ECOS [DCB13].

convex optimization problem

2.3 Probability Theory

In this thesis, we reason about stochastic uncertainty in a probability-theoretic framework. While we expect you (the reader) to obtain a deeper understanding of our results if you are familiar with the basics of probability theory, it is not a big problem if you do not have this background knowledge.

We assume familiarity with basic concepts such as σ -algebras and probability measures. We refer to [BR07] for a comprehensive introduction of measure theory, and to books such as [Dur10; Wil91; Kal02] for an introduction of probability theory.

2.3.1 Probability distributions

We start our discussion with the concept of a *measurable space*. For the definition of a σ -algebra, we refer to any of the books above.

measurable space

Definition 2.2 (Measurable space) A *measurable space* is a pair (Ω, \mathcal{F}) consisting of a set Ω and a σ -algebra \mathcal{F} of subsets of Ω .

probability space

When we equip a measurable space with a probability measure $\mathbb{P}: \mathcal{F} \rightarrow [0, 1]$, we obtain a *probability space*.

Definition 2.3 (Probability space) A *probability space* is a triple $(\Omega, \mathcal{F}, \mathbb{P})$, where Ω is a set of outcomes called the *sample space*, \mathcal{F} is a σ -algebra over Ω , and $\mathbb{P}: \mathcal{F} \rightarrow [0, 1]$ is a probability measure.

Probability spaces will form the basis for reasoning about stochastic uncertainty throughout this thesis.

Discrete distributions | Measurability concerns do not apply to distributions over discrete sets. Thus, we simplify the notation for (*discrete*) *probability distribution* over finite or countably infinite sets as follows.

discrete distribution

Definition 2.4 (Discrete probability distribution) Let X be a finite or countably infinite set. A *discrete probability distribution* over X is a function $\mu: X \rightarrow \mathbb{R}_{\geq 0}$ such that $\sum_{x \in X} \mu(x) = 1$.

In this definition, the set X is implicitly endowed with a σ -algebra $\mathcal{F} = 2^X$ being the power set of X . For a probability distribution μ over a set X , we call $\mu(x)$ the probability of $x \in X$. The support of a distribution μ over X , written as $\text{supp}(\mu)$, is the subset of X that has positive probability, i.e., $\text{supp}(\mu) = \{x \in X : \mu(x) > 0\}$. For a set X , we denote the set of all distributions over X by $\text{Distr}(X)$. One special distribution if the *Dirac distribution* $\delta_y: X \rightarrow \mathbb{R}_{\geq 0}$ over a set X , with $y \in X$, which is defined as $\delta_y(x) = 1$ for $x = y$ and $\delta_y(x) = 0$ otherwise.

Dirac distribution

2.3.2 Random variables

Random variables formalize the concept of a quantity that depends on random events. By definition, a random variable is a *measurable function*.

measurable function

Definition 2.5 (Measurable function) Let (Ω, \mathcal{F}) and (Ω', \mathcal{G}) be measurable spaces. A function $f: \Omega \rightarrow \Omega'$ is *measurable* (with respect to the pair $(\mathcal{F}, \mathcal{G})$) if $f^{-1}(G) \in \mathcal{F}$ for all $G \in \mathcal{G}$.

A *random variable* is a measurable function between two measurable spaces.

random variable

Definition 2.6 (Random variable) Let (Ω, \mathcal{F}) and (Ω', \mathcal{G}) be measurable spaces. A function $X: \Omega \rightarrow \Omega'$ is called a *random variable* if X is measurable with respect to the pair $(\mathcal{F}, \mathcal{G})$.

2.3.3 Stochastic processes

By ordering multiple random variables over an index set, we obtain a *stochastic process*. Often, this index set represents time, such that a stochastic process models the evolution

of a random process over time. In this thesis, we only consider stochastic processes evolving over discrete time steps, which are the result of a countable collection of random variables.

Definition 2.7 (Stochastic process) A sequence $(X_k)_{k \in \mathbb{N}} := \{X_k : k = 0, 1, \dots\}$ of random variables $X_k : \Omega \rightarrow \Omega'$ from the common probability space $(\Omega, \mathcal{F}, \mathbb{P})$ to the measurable space (Ω', \mathcal{G}) is called a (discrete-time) *stochastic process*.

stochastic process

Without loss of generality, Def. 2.7 assumes that each random variable X_k of the stochastic process maps to the same measurable space. The next important concept is that of a *filtration*, which intuitively can be used to encode the information contained in the history of a stochastic process.

Definition 2.8 (Filtration) A collection $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$ such that $\mathcal{F}_0 \subseteq \mathcal{F}_1 \subseteq \dots \subseteq \mathcal{F}$ is called a *filtration* on the measurable space (Ω, \mathcal{F}) . If (Ω, \mathcal{F}) is equipped with a probability measure \mathbb{P} , then the triple $(\Omega, \mathcal{F}, \{\mathcal{F}_k\}_{k \in \mathbb{N}}, \mathbb{P})$ is called a filtered probability space.

filtration

Particularly useful is the so-called *natural filtration* of a stochastic process. For this definition, let $\sigma(\{X_1, X_2, \dots, X_n\})$ denote the σ -algebra generated by the collection of $n \in \mathbb{N}$ random variables X_1, X_2, \dots, X_n , as defined in the standard way; see, e.g., [CE15] for a definition.

natural filtration

Definition 2.9 (Natural filtration) Let $(X_k)_{k \in \mathbb{N}}$ be a stochastic process defined on the probability space $(\Omega, \mathcal{F}, \mathbb{P})$. The *natural filtration* $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$ is defined for all $k \in \mathbb{N}$ as $\mathcal{F}_k = \sigma(\cup_{i=0}^k X_i)$.

Consider a concrete outcome $\omega \in \Omega$ which determines the values for a stochastic process $(X_k)_{k \in \mathbb{N}}$. Intuitively, the natural filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$ models the information about ω at every time step $k \in \mathbb{N}$. The more fine-grained the σ -algebra \mathcal{F}_k at time k , the more information we have about ω . Furthermore, observe that, by definition of the filtration, the information about ω cannot decrease over time.

adapted stochastic process

Definition 2.10 A stochastic process $(X_k)_{k \in \mathbb{N}}$ from the probability space $(\Omega, \mathcal{F}, \mathbb{P})$ to the measurable space (Ω', \mathcal{G}) is called *adapted* (to the filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$) if for each $k \in \mathbb{N}$, the random variable X_k is measurable with respect to the pair $(\mathcal{F}_k, \mathcal{G})$, i.e., if $X_k^{-1}(G) \in \mathcal{F}_k$ for all $G \in \mathcal{G}$.

If a stochastic process is adapted, the value $X_k(\omega)$ is known to us at time k . Note that any stochastic process is an adapted process with respect to its natural filtration.

3 A Primer on Markov Decision Processes

Summary | Discrete-time Markov chains (DTMCs) are probabilistic models that exhibit the so-called Markov property: In any state, the probability distribution over the next state is independent of previous states. Markov decision processes (MDPs) extend DTMCs with nondeterministic action choices and are widely used for modeling sequential decision-making problems. In this chapter, we give a general overview of DTMCs and MDPs and discuss how to formulate common measures for these models, such as reachability probabilities and cumulative expected rewards. Finally, we introduce robust Markov decision processes (RMDPs) and interval Markov decision processes (IMDPs) as extensions of MDPs with uncertain transition probabilities.

Origins | None of the content in this chapter is novel. For comprehensive introductions to DTMCs and MDPs, we refer to [BK08; Put94]. For details on RMDPs, we refer to the seminal papers [GLD00; NG05; WKR13].

Background | While we introduce the relevant Markov models from their basic definitions, some level of acquaintance with transition systems is desirable; see, for example, [BK08] for a comprehensive introduction.



3.1 Markov Decision Processes

We define Markov decision processes (MDPs) and discrete-time Markov chains (DTMCs). Thereafter, we discuss paths, schedulers, and how to perform common analyses (such as computing reachability probabilities and expected rewards).

Atomic propositions | Throughout this thesis, let AP be a finite set of *atomic propositions*. In practice, these propositions express fundamental statements that are either true or false in each state of a model. For example, an atomic proposition $a \in AP$ may model being in an unsafe state and is thus true in a subset of unsafe states. Multiple atomic propositions can be true in a single state.

Definition 3.1 (MDP) A *Markov decision process (MDP)* is a tuple $\mathcal{M} := (S, Act, s_I, P, L, r)$ where S is a finite set of *states*, Act is a finite set of *actions*, $s_I \in S$ is an *initial state*, $P: S \times Act \rightarrow \text{Distr}(S)$ is a *transition probability function*, $L: S \rightarrow 2^{AP}$ is a *labeling function*, and $r: S \rightarrow \mathbb{R}_{\geq 0}$ is a *state reward function*.

The MDP transition function is *partial* in general (denoted by the symbol \rightarrow), meaning that not all state-action pairs $(s, a) \in S \times Act$ are in the domain of P . By using a partial transition function, we can model that only a subset of actions may be enabled in each state. The set of actions enabled in state $s \in S$ is $Act(s) := \{a \in Act \mid (s, a) \in \text{dom}(P)\} \subseteq$

atomic
proposition

Markov
decision
process

partial
transition
function

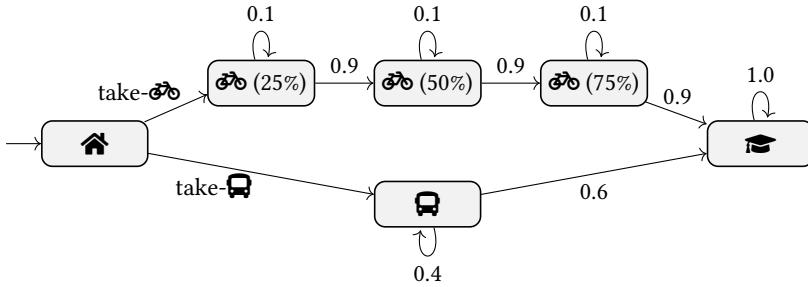


Figure 3.1: Going to university (🎓) by bicycle (🚲) or bus (🚌), modeled as an MDP with two actions in the initial states. The other states have only a single action enabled (and thus have nondeterminism).

Act. Without loss of generality, we assume that at least one action is enabled in every deadlock state, i.e., $|Act(s)| \geq 1$ for all $s \in S$, also referred to as the MDP being free of *deadlocks*.

Example 3.2 Suppose that you have to go to the university for a lecture. You can choose to travel by bike or by bus—which option will get you the fastest at the university? We model this scenario as an MDP, where each step represents 10 min. If you go by bus, there is a 0.6 probability that the bus shows up, and a 0.4 probability that it does not. However, once the bus shows up, you are at university in one time step. By contrast, taking the bike takes multiple time steps, but the probability of a delay (e.g., due to a red traffic light) is lower. The resulting MDP is shown in Fig. 3.1 and will be our running example throughout this chapter.

Remark 3.3 (Omitting rewards and labels) For some analyses (such as reachability probabilities; see Sect. 3.2.2), the reward function r is irrelevant. In such cases, we omit the reward function and write an MDP as $\mathcal{M} := (S, Act, s_I, P, L)$. Similarly, we often omit the labeling function L from the MDP definition and, instead, work with specifications defined directly over subsets of states $S' \subset S$.

transition
self-loop
absorbing state

We call the triple (s, a, s') with probability $P(s, a)(s') > 0$ a *transition*. Each transition (s, a, s') represents an edge in the graph structure of the MDP, where s is the outgoing state and s' is the ingoing state. A transition (s, a, s) where the outgoing and ingoing states are the same is called a *self-loop*. A state $s \in S$ for which the only outgoing transition (s, a, s) is a self-loop is called *absorbing*.

We interpret a DTMC as a special case of MDP with a single action in every state.

discrete-time Markov chain

Definition 3.4 (DTMC) A *(discrete-time) Markov chain* is an MDP with $|Act(s)| = 1$ for all $s \in S$. We write a DTMC as a tuple $\mathcal{D} := (S, s_I, P, L, r)$, where S , s_I , L , and r are defined as in Def. 3.1, and where the transition probability function is defined as $P: S \rightarrow \text{Distr}(S)$.

For DTMCs, the transition probability function can equivalently be seen as a matrix

$P \in \mathbb{R}^{|S| \times |S|}$, such that $P_{i,j} \geq 0$ for all $i, j \in \{1, \dots, |S|\}$ and $\sum_{j=1}^{|S|} P_{i,j} = 1$ for all $i \in \{1, \dots, |S|\}$. In this thesis, however, we will interpret the transition function P as a map from states to distributions rather than a matrix.

3.1.1 Paths and sets of paths

Let $\mathcal{M} = (S, Act, s_I, P, L)$ be an MDP. A *path* of MDP \mathcal{M} is an (in)finite sequence $\pi := s_0 a_0 s_1 a_1 \dots$ of states and actions, such that $a_i \in Act(s_i)$ for all $i \in \mathbb{N}$, and $P(s_i, a_i)(s_{i+1}) > 0$. The $(i+1)^{\text{th}}$ state in an (in)finite path $\pi = s_0 a_0 \dots s_i a_i \dots$ is denoted by $\pi(i) := s_i$. Similarly, we denote a path fragment as $\pi(i, \dots, j) := s_i a_i s_{i+1} a_{i+1} \dots a_{j-1} s_j$, where $i, j \in \mathbb{N}$ and $i \leq j$. For an (in)finite path $\pi = s_0 a_0 s_1 a_1 \dots$, we write the *first state* of π as $\text{first}(\pi) := s_0$. Furthermore, for a finite path $\pi = s_0 a_0 s_1 a_1 \dots s_n$, we denote the *last state* by $\text{last}(\pi) := s_n$ and we define the *length* of π as $|\pi| := n + 1$. Note that a path of length $|\pi| = 1$ is just a single state, $\pi = s_0$.

We write the set of all finite paths as $\Pi_{\text{fin}}^{\mathcal{M}}$ and the set of all infinite paths as $\Pi_{\text{inf}}^{\mathcal{M}}$. The set of all paths $\Pi^{\mathcal{M}}$ is the union of the finite and infinite paths, i.e., $\Pi^{\mathcal{M}} := \Pi_{\text{fin}}^{\mathcal{M}} \cup \Pi_{\text{inf}}^{\mathcal{M}}$. For a subset of states $S' \subseteq S$, we define the set of all paths starting in a state in S' as

$$\Pi^{\mathcal{M}}(S') := \left\{ \pi \in \Pi^{\mathcal{M}} : \text{first}(\pi) \in S' \right\},$$

and we define the sets $\Pi_{\text{fin}}^{\mathcal{M}}(S')$ and $\Pi_{\text{inf}}^{\mathcal{M}}(S')$ analogously. For a singleton set $S' = \{s'\}$, we omit the brackets and write $\Pi^{\mathcal{M}}(s')$.

Remark 3.5 (Paths in DTMCs) For DTMCs, the actions of a path are trivial and thus omitted. Thus, for a DTMC \mathcal{D} , a path is defined as $\pi = s_0 s_1 s_2 \dots$, and the sets of all, all finite, and all infinite paths defined above are denoted by $\Pi^{\mathcal{D}}$, $\Pi_{\text{fin}}^{\mathcal{D}}$, and $\Pi_{\text{inf}}^{\mathcal{D}}$, respectively.

3.1.2 Schedulers

Because the action choices are unspecified, MDPs are called *nondeterministic*. A *scheduler* (also called *policy* or *strategy*) is a decision rule that resolves these nondeterministic action choices.

scheduler

Remark 3.6 (Schedulers vs. policies) In the literature, the term *policy* is arguably more popular than *scheduler*. However, the word *policy* may also relate to the continuous-state/action systems we will discuss in Part II [BS78]. For clarity, we thus prefer using the word *scheduler* for (finite) MDPs, whereas we will use the word *policy* for the (continuous) systems in Part II.

Schedulers for MDPs are commonly defined as a function from paths to *distributions over actions* [BK08; Put94]. Such schedulers are called *randomized*. On the other hand, a scheduler that maps every path to a Dirac distribution (or equivalently, maps every path to a single action $a \in Act$) is called *deterministic*. For the measures and specifications for MDPs that we consider in this thesis, randomized and deterministic schedulers lead to the same optimal values. Thus, randomized schedulers do not provide any extra power despite being more general than deterministic schedulers. As a result, we restrict ourselves to deterministic schedulers in this thesis.

randomized scheduler

deterministic scheduler

Definition 3.7 (Scheduler) A *(deterministic) scheduler* for MDP \mathcal{M} is a function $\sigma: \Pi_{\text{fin}}^{\mathcal{M}} \rightarrow \text{Act}$ mapping finite paths to actions, with $\sigma(\pi) \in \text{Act}(\text{last}(\pi))$ for all paths $\pi \in \Pi_{\text{fin}}^{\mathcal{M}}$. The set of all schedulers for an MDP \mathcal{M} is denoted by $\mathfrak{S}^{\mathcal{M}}$.

induced DTMC

Applying a scheduler to an MDP resolves the nondeterministic action choices, resulting in a DTMC that we call the induced Markov chain or *induced DTMC*. As we shall see, for most objectives in MDPs, the optimal actions depend only on the current state and not the history of states. *Stationary* schedulers (also known as *memoryless* schedulers [BK08]) formalize this concept.

stationary scheduler

Definition 3.8 (Stationary scheduler) A scheduler σ is *stationary* if for all paths $\pi, \pi' \in \Pi_{\text{fin}}^{\mathcal{M}}$, it holds that

$$\text{last}(\pi) = \text{last}(\pi') \implies \sigma(\pi) = \sigma(\pi').$$

We simplify notation and write stationary schedulers as $\sigma: S \rightarrow \text{Act}$. The set of all stationary schedulers for an MDP \mathcal{M} is denoted by $\mathfrak{S}_{\text{stat}}^{\mathcal{M}}$.

For measures with no explicit time constraints (i.e., measures over an infinite time horizon), we shall see that we can restrict the class of schedulers to the set of stationary schedulers when analyzing an MDP.

Example 3.9 Consider again the example MDP in Fig. 3.1. This MDP has two stationary schedulers. The first scheduler corresponds with taking the bike, such that $\sigma(\text{🏠}) = \text{take-bike}$ (the other states have only a single action enabled, so we omit them from the scheduler definition). The second scheduler corresponds with taking the bus, such that $\sigma'(\text{🏠}) = \text{take-bus}$.

However, when dealing with measures over a finite time horizon, stationary schedulers do not suffice. As an intuitive example, consider the problem of reaching a set of target states $S_T \subset S$ within ten steps. If we have nine more steps to go (i.e., we can still choose nine actions), then the optimal action can be different from when we only have one more step to go. Following the nomenclature from [Put94, Section 2.1], we formalize this idea in the notion of a *Markov scheduler*.

Markov scheduler

Definition 3.10 (Markov scheduler) A scheduler σ is *Markov* (or *Markovian*) if for all paths $\pi, \pi' \in \Pi_{\text{fin}}^{\mathcal{M}}$, it holds that

$$\text{last}(\pi) = \text{last}(\pi') \wedge |\pi| = |\pi'| \implies \sigma(\pi) = \sigma(\pi').$$

We simplify notation and write Markov schedulers as a sequence of stationary schedulers, $\sigma = (\sigma_0, \sigma_1, \dots)$, where $\sigma_k: S \rightarrow \text{Act}$ for $k \in \mathbb{N}$. The set of all Markov schedulers is denoted by $\mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$.

Observe that a Markov scheduler only depends on the last state $\text{last}(\pi) \in S$ and the length $|\pi| \in \mathbb{N}_{>0}$ of a path $\pi \in \Pi_{\text{fin}}^{\mathcal{M}}$. This observation motivates the notation $\sigma = (\sigma_0, \sigma_1, \dots)$ for a Markov scheduler, which is also used by [Put94] and makes explicit

that such a scheduler only depends on the last state and the length of the path (and not on the full path). Markov schedulers are less general than the standard definition of a scheduler in Def. 3.7 but more general than the stationary schedulers in Def. 3.8.

Remark 3.11 (Encoding time in schedulers) Instead of using a Markov scheduler $\sigma = (\sigma_0, \sigma_1, \dots)$ as in Def. 3.10, another way to make schedulers dependent on time is to use a stationary scheduler $\sigma': S \rightarrow \text{Act}$ and create a copy of the MDP states S for every time step $k \in \mathbb{N}$ (the initial states, enabled actions, and transition function should be modified accordingly as well; however, we omit these definitions for brevity). Intuitively, we thus obtain a larger MDP with a factored state space $S \times \mathbb{N}$, in which actions are chosen as

$$\sigma'((s, k)) = \sigma_k(s) \quad \forall s \in S, k \in \mathbb{N}.$$

Both of these perspectives are equivalent. In this thesis, we prefer using Markov schedulers because it avoids having to define time in MDPs explicitly.

Probability measure | Let $\mathcal{M} = (S, \text{Act}, s_I, P, L)$ be an MDP with a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$. We define the (standard) **probability measure** $\Pr_{\sigma}^{\mathcal{M}}: \Pi_{\text{fin}}^{\mathcal{M}} \rightarrow [0, 1]$ over finite paths $\pi = s_0 s_1 \cdots s_n \in \Pi_{\text{fin}}^{\mathcal{M}}$ as

$$\Pr_{\sigma}^{\mathcal{M}}(\pi) := \prod_{i=0}^{n-1} P(s_i, \sigma(\pi(0, \dots, i)))(s_{i+1}). \quad (3.1)$$

probability measure (MDP)

For stationary and Markov schedulers, the scheduler $\sigma(\pi(0, \dots, i))$ in Eq. (3.1) is replaced by the appropriate forms. Thus, for schedulers $\sigma' \in \mathfrak{S}_{\text{stat}}$ and $\sigma'' \in \mathfrak{S}_{\text{Markov}}$, we obtain the probability measures

$$\Pr_{\sigma'}^{\mathcal{M}}(\pi) := \prod_{i=0}^{n-1} P(s_i, \sigma'(s_i))(s_{i+1}), \text{ and } \Pr_{\sigma''}^{\mathcal{M}}(\pi) := \prod_{i=0}^{n-1} P(s_i, \sigma''(s_i))(s_{i+1}). \quad (3.2)$$

Intuitively, $\Pr_{\sigma}^{\mathcal{M}}(\pi)$ is the probability that the Markov chain induced by applying the scheduler σ to MDP \mathcal{M} generates the path $\pi \in \Pi_{\text{fin}}^{\mathcal{M}}$. The probability measure $\Pr_{\sigma}^{\mathcal{M}}$ also defines the probability $\Pr_{\sigma}^{\mathcal{M}}(\Psi)$ for any set of finite paths $\Psi \subseteq \Pi_{\text{fin}}^{\mathcal{M}}$. For infinite paths, we use the same notation and define $\Pr_{\sigma}^{\mathcal{M}}$ as the (unique) probability measure over the smallest σ -algebra containing the cylinder sets of all finite paths. This cylinder set construction is standard, so we refer to [BK08, Def. 10.10] for details.

Example 3.12 For our example MDP in Fig. 3.1 of traveling to the university, consider the path $\pi = (\text{House}, \text{Bus}, \text{Bus}, \text{Graduation Cap})$. The probability $\Pr_{\sigma'}^{\mathcal{M}}(\pi)$ for this path (under the stationary scheduler σ' that always chooses the bus) is obtained by multiplying the transition probabilities of the corresponding transitions:

$$\Pr_{\sigma'}^{\mathcal{M}}(\pi) = 0.4 \cdot 0.6 = 0.24.$$

Remark 3.13 (Probability measure for DTMCs) Recall that applying a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$ to an MDP \mathcal{M} induces a DTMC. In fact, the probability measure $\text{Pr}_{\sigma}^{\mathcal{M}}$ is defined on this induced DTMC. As such, we may define the probability measure for a DTMC, denoted by $\text{Pr}^{\mathcal{D}}$, in the same way.

3.2 Analyzing MDPs

In this section, we introduce the analyses of MDPs we use in this thesis. For MDPs, we, in particular, compute so-called measures of reachability and expected rewards.

3.2.1 Probabilistic computation tree logic

PCTL In this section, we introduce *probabilistic computation tree logic (PCTL)* [HJ94], which is a logic widely used for MDPs, and which extends the classical computation tree logic (CTL) with the probabilistic operator \mathbb{P} . Most of the specifications for MDP and DTMCs considered in this thesis can be expressed in PCTL.

Definition 3.14 (Syntax of PCTL) A PCTL formula over a set of atomic propositions AP is built using the following grammar, starting with symbol Φ :

$$\begin{aligned}\Phi &::= \text{True} \mid a \mid \neg\Phi \mid \Phi \wedge \Phi \mid \mathbb{P}_{\triangleleft\lambda}(\varphi) \\ \varphi &::= \bigcirc\Phi \mid \Phi \cup \Phi \mid \Phi \cup^{\leq h} \Phi,\end{aligned}$$

where $a \in AP$ is an atomic proposition, $\triangleleft \in \{<, \leq, \geq, >\}$ is a comparison operator, $\lambda \in [0, 1]$ is a threshold probability, and $h \in \mathbb{N}$ is a time bound.

state/path formula

The symbols \neg and \wedge are the logical negation and conjunction, and the temporal operators \bigcirc (*next*), \cup (*until*), and $\cup^{\leq h}$ (*bounded until*) are defined in the usual way, e.g., as also used in linear temporal logic (LTL) [Pnu77]. In Def. 3.14, the symbol Φ is called a *state formula* and is interpreted on the state of a DTMC. Similarly, φ is called a *path formula* and is interpreted on the infinite paths of a DTMC. The interpretation is Boolean, i.e., a state (or a path) either satisfies a state (or path) formula or not.

We can derive several other useful operators from the basic PCTL syntax.

Definition 3.15 (Derived operators) From the PCTL syntax in Def. 3.14, we derive the *eventually* operator \diamond and the *globally* (or *always*) operator \square as follows:

$$\mathbb{P}_{\triangleleft\lambda}(\diamond\Phi) := \mathbb{P}_{\triangleleft\lambda}(\text{True} \cup \Phi) \quad \text{and} \quad \mathbb{P}_{\triangleleft\lambda}(\square\Phi) := \mathbb{P}_{\triangleright(1-\lambda)}(\text{True} \cup \neg\Phi),$$

where $\triangleleft \in \{<, \leq, \geq, >\}$ is a comparison operator and \triangleright is its opposite,^a and $\lambda \in [0, 1]$ is a threshold probability. Similarly, we define the *step-bounded eventually* and *globally* operators as

$$\mathbb{P}_{\triangleleft\lambda}(\diamond^{\leq h}\Phi) := \mathbb{P}_{\triangleleft\lambda}(\text{True} \cup^{\leq h} \Phi) \quad \text{and} \quad \mathbb{P}_{\triangleleft\lambda}(\square^{\leq h}\Phi) := \mathbb{P}_{\triangleright(1-\lambda)}(\text{True} \cup^{\leq h} \neg\Phi).$$

^aThis opposite is defined such that if \triangleleft is $<$ then \triangleright is $>$, if \triangleleft is \leq then \triangleright is \geq , and so on.

Satisfaction relation | We often ask for the probability that an MDP, starting in a fixed state $s \in S$, generates a path satisfying a given PCTL path formula. We call this probability the *satisfaction probability*. In the following, we define the satisfaction probability for the DTMC induced by an MDP \mathcal{M} and a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$. For a DTMC $\mathcal{D} = (S, s_I, P, L)$, the satisfaction probability is defined equivalently by replacing the measure $\text{Pr}_{\sigma}^{\mathcal{M}}$ with $\text{Pr}^{\mathcal{D}}$ and the path set $\Pi^{\mathcal{M}}$ by $\Pi^{\mathcal{D}}$.

3

satisfaction probability

Definition 3.16 (Satisfaction probability) Let \mathcal{M} be an MDP with a state $s \in S$ and a scheduler $\mathfrak{S}^{\mathcal{M}}$. The *satisfaction probability* $\text{Pr}_{\sigma}^{\mathcal{M}}(s \models \varphi)$ of a PCTL path formula φ is defined as

$$\text{Pr}_{\sigma}^{\mathcal{M}}(s \models \varphi) := \text{Pr}_{\sigma}^{\mathcal{M}}(\pi \in \Pi^{\mathcal{M}}(s) : \pi \models \varphi).$$

When no state s is specified, we use the initial MDP state s_I , i.e.,

$$\text{Pr}_{\sigma}^{\mathcal{M}}(\varphi) := \text{Pr}_{\sigma}^{\mathcal{M}}(\pi \in \Pi^{\mathcal{M}}(s_I) : \pi \models \varphi).$$

In Def. 3.16, the term $\pi \models \varphi$ is the satisfaction relation of a PCTL formula in the standard sense, which we revisit below [BK08, Chapter 10]. For state formulae, the satisfaction relation \models is a relation between the states S of a DTMC and state formulae. For path formulae, \models is a relation between infinite paths and path formulae.

Definition 3.17 (Satisfaction of PCTL) Let $\mathcal{M} = (S, \text{Act}, s_I, P, L)$ be an MDP with a state $s \in S$ and a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$, let $a \in AP$ be an atomic proposition, let Φ, Φ' be PCTL state formulae, and let φ be a PCTL path formula. The satisfaction relation \models is defined for state formulae as

$$\begin{aligned} s \models a &\quad \text{iff } a \in L(s) \\ s \models \neg\Phi &\quad \text{iff } s \not\models \Phi \\ s \models \Phi \wedge \Phi' &\quad \text{iff } s \models \Phi \text{ and } s \models \Phi' \\ s \models \mathbb{P}_{\lambda}(\varphi) &\quad \text{iff } \text{Pr}_{\sigma}^{\mathcal{M}}(s \models \varphi) \triangleleft \lambda. \end{aligned}$$

The satisfaction relation for a path π in \mathcal{M} is defined as

$$\begin{aligned} \pi \models \bigcirc\Phi &\quad \text{iff } \pi(1) \models \Phi \\ \pi \models \Phi \cup \Phi' &\quad \text{iff } \exists j \geq 0. (\pi(j) \models \Phi' \wedge (\forall 0 \leq k < j. \pi(k) \models \Phi)) \\ \pi \models \Phi \cup^{\leq h} \Phi' &\quad \text{iff } \exists 0 \leq j \leq h. (\pi(j) \models \Phi' \wedge (\forall 0 \leq k < j. \pi(k) \models \Phi)). \end{aligned}$$

It has been shown that the events specified by PCTL formulae are measurable, and thus, satisfaction probabilities are ensured to be well-defined [BK08, Chapter 10].

3.2.2 Measures

The satisfaction probability in Def. 3.16 provides the basis for various common *measures* of interest. In this section, we present the measures of reachability and expected rewards that we use throughout this thesis. For MDPs, these measures are defined on the induced

measure
(for MDP)

Markov chain for a fixed scheduler. Thus, in our notation, we will often specify both an MDP \mathcal{M} and a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$. For DTMCs, the same definitions apply while omitting any subscripts related to the scheduler.

3.2.2.1 Reachability

Reachability measures express the probability of eventually reaching a subset of states $S_T \subseteq S$. A standard reachability measure is expressed by the PCTL path formula $\varphi = \diamond S_T$, and its satisfaction probability is defined as follows.

Definition 3.18 (Reachability probability) Let \mathcal{M} be an MDP with a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$ and a state $s \in S$, and let $S_T \subseteq S$ be a set of target states.

- The probability $\Pr_{\sigma}^{\mathcal{M}}(s \models \diamond S_T)$ is called a *reachability probability*;
- For a time bound $h \in \mathbb{N}$, the probability $\Pr_{\sigma}^{\mathcal{M}}(s \models \diamond^{\leq h} S_T)$ is called a *step-bounded reachability probability*.

A *reach-avoid measure* extends a reachability measure with a set of states $S_U \subseteq S$ that must be avoided until the target states S_T are reached.

Definition 3.19 (Reach-avoid probability) Let \mathcal{M} be an MDP with a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$ and a state $s \in S$, and let $S_T, S_U \subseteq S$ be sets of states.

- The probability $\Pr_{\sigma}^{\mathcal{M}}(s \models \neg S_U \cup S_T)$ is called a *reach-avoid probability*;
- For a time bound $h \in \mathbb{N}$, the probability $\Pr_{\sigma}^{\mathcal{M}}(s \models \neg S_U \cup^{\leq h} S_T)$ is called a *step-bounded reach-avoid probability*.

3.2.2.2 Expected rewards

Besides measures based on PCTL, we also use measures on expected rewards. Recall from Def. 3.1 that we consider MDPs equipped with a reward function $r: S \rightarrow \mathbb{R}_{\geq 0}$. For a path $\pi = s_0 a_0 s_1 \dots \in \Pi^{\mathcal{M}}$, we define the cumulative reward $\text{rew}(\pi)$ as

$$\text{rew}(\pi) := \sum_{i=0}^{|\pi|} \text{rew}(s_i) = \text{rew}(s_0) + \text{rew}(s_1) + \dots + \text{rew}(s_{|\pi|}).$$

The cumulative reward is finite for all finite paths but can be infinite for infinite paths. We circumvent issues with infinite rewards by considering the expected cumulative reward *before* reaching a given set of states. Formally, let $B \subseteq S$ be a subset of states of MDP \mathcal{M} and suppose that $\Pr_{\sigma}^{\mathcal{M}}(s \models B) = 1$, that is, S_T is reached with probability one from state $s \in S$. Then, the *expected cumulative reward* before reaching B is characterized by a summation over finite paths only.

Definition 3.20 (Expected cumulative reward) Let $\mathcal{M} = (S, \text{Act}, s_I, P, L, r)$ be an MDP with a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}}$, and let $B \subseteq S$. The *expected cumulative reward* before reaching B is defined as

$$\text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B) := \begin{cases} \sum_{\pi \in \Pi_{\text{fin}}^{\mathcal{M}}(s, \diamond B)} \Pr_{\sigma}^{\mathcal{M}}(\pi) \cdot \text{rew}(\pi) & \text{if } \Pr_{\sigma}^{\mathcal{M}}(s \models \diamond B) = 1, \\ +\infty & \text{otherwise,} \end{cases}$$

reachability probability

reach-avoid probability

expected cumulative reward

where $\Pi_{\text{fin}}^{\mathcal{M}}(s, \diamond B) \subseteq \Pi_{\text{fin}}^{\mathcal{M}}$ is the set of paths $s_0 a_0 s_1 a_1 \cdots a_{n-1} s_n$ with $s_0 = s$, $s_n \in B$, and $s_i \notin B$ for all $i = 0, \dots, n-1$. When no initial state is specified, we use the initial MDP state:

$$\text{ExpRew}_{\sigma}^{\mathcal{M}}(\diamond B) := \text{ExpRew}_{\sigma}^{\mathcal{M}}(s_I \models \diamond B).$$

Observe that $\Pi_{\text{fin}}^{\mathcal{M}}(s, \diamond B)$ in Def. 3.20 is the set of paths starting in $s \in S$, ending in B , and not having reached B before. The intuitive argument allowing us to neglect infinite paths in Def. 3.20 is that $\Pr_{\sigma}^{\mathcal{M}}(s \models \diamond B) = 1$ implies that all infinite paths that never reach B have probability zero. As a result, the expected cumulative reward is finite and can be computed as the sum over all finite paths [BK08, Sec. 10.5.1].

Note that we define $\text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B)$ as $+\infty$ if B is not reached with probability one, as also done by [BK08, Sec. 10.5.1]. In practice, however, we will only consider cases where $\Pr_{\sigma}^{\mathcal{M}}(s \models \diamond B) = 1$, thus avoiding this corner case anyway.

Remark 3.21 (Discount factor) Our definition of expected cumulative rewards in Def. 3.20 is common in formal methods but differs from the standard definition used in artificial intelligence [SB98] and operations research [Put94]. In these areas, it is common not to specify a set of goal states (as we did in Def. 3.20) and instead incorporate a *discount factor* $\gamma \in (0, 1)$. The discount factor promotes immediate rewards over future rewards and is another way to ensure that the expected cumulative reward is finite (also over infinite paths). Then, the standard goal is to compute a Markov (or even stationary) scheduler $\sigma = (\sigma_0, \sigma_1, \dots) \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ that maximizes the expected cumulative discounted reward defined as

$$\mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) \right],$$

where $s_{t+1} \sim P(s_t, \sigma_t(s_t))$ for all $t \in \mathbb{N}$, and the initial state $s_0 = s_I$ is specified by the MDP. While we refrain from using such a discount factor in this thesis and instead follow Def. 3.20, our methods can readily be adapted to incorporate a discount factor (e.g., value iteration is immediately compatible with a discount factor [SB98]).

discount factor

3.2.3 Value iteration for MDPs

Often, we are interested in computing a scheduler that *maximizes* or *minimizes* a given measure. For example, for a cumulative expected reward, we want to compute

$$\sigma_{\max}^* \in \operatorname{argmax}_{\sigma \in \mathfrak{S}^{\mathcal{M}}} \text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B), \quad \text{or} \quad \sigma_{\min}^* \in \operatorname{argmin}_{\sigma \in \mathfrak{S}^{\mathcal{M}}} \text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B).$$

We write such *optimal schedulers* with a \star -script. In practice, whether we maximize or minimize will be clear from the context, and we simply write σ^* . Note that optimal schedulers may not be unique.

optimal scheduler
(for MDP)

The following is a classical result for MDPs, which shows that *deterministic stationary* schedulers are sufficient to maximize (or minimize) cumulative expected rewards.

Proposition 3.22 (Stationary schedulers suffice [Put94]) (Deterministic) stationary schedulers $\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}$ suffice to obtain maximal and minimal values for cumulative expected rewards:

$$\max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}} \text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B) = \max_{\sigma' \in \mathfrak{S}^{\mathcal{M}}} \text{ExpRew}_{\sigma'}^{\mathcal{M}}(s \models \diamond B), \quad \text{and}$$

$$\min_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}} \text{ExpRew}_{\sigma}^{\mathcal{M}}(s \models \diamond B) = \min_{\sigma' \in \mathfrak{S}^{\mathcal{M}}} \text{ExpRew}_{\sigma'}^{\mathcal{M}}(s \models \diamond B).$$

Proposition 3.22 motivates searching the *finite* set $\mathfrak{S}_{\text{stat}}^{\mathcal{M}}$ of stationary schedulers only (rather than the set $\mathfrak{S}^{\mathcal{M}}$ of *infinitely many* schedulers; see Def. 3.7).

value iteration

Value iteration, described in Algorithm 3.1, is a classical algorithm to estimate optimal stationary schedulers for expected cumulative rewards in MDPs. The algorithm stores a *value* $V(s)$ for each MDP state $s \in S$, which represents an estimate of a given measure in that state. Starting from an arbitrary value $V_0(s)$ for all states $s \in S$ (which can simply be zero), the algorithm iteratively updates the values by successively applying the Bellman backup operator (i.e., line 6 of Algorithm 3.1). The algorithm can loop over the states $s \in S$ in any order, but some choices lead to faster convergence than others due to the dependencies between $V_i(s)$ for different states $s \in S$ (see [SB98, Chapter 4] for details). Algorithm 3.1 terminates when the values in two consecutive iterations differ less than some predefined $\epsilon \geq 0$. Upon termination, the algorithm returns the (stationary) scheduler that maximizes the value estimates in every state.

Remark 3.23 (Minimizing schedulers) To minimize a cumulative expected reward, we replace the max and argmax operators in Algorithm 3.1 by min and argmin, respectively.

Algorithm 3.1 Value iteration for maximizing expected cumulative rewards in MDPs.

Input: MDP $\mathcal{M} = (S, \text{Act}, s_I, P, L, r)$

Params: Precision parameter $\epsilon \geq 0$

Output: Memoryless deterministic scheduler $\sigma: S \rightarrow \text{Act}$

-
- 1: $V_0(s) \leftarrow 0, \quad \forall s \in S$
 - 2: $i \leftarrow 0$
 - 3: **repeat**
 - 4: $i \leftarrow i + 1$
 - 5: **for** $s \in S$ **do**
 - 6: $V_i(s) \leftarrow \max_{a \in \text{Act}(s)} [r(s) + \sum_{s' \in S} P(s, a)(s')V_{i-1}(s')]$
 - 7: **until** $|V_i(s) - V_{i-1}| \leq \epsilon$
 - 8: **for** $s \in S$ **do**
 - 9: $\sigma(s) \leftarrow \arg\max_{a \in \text{Act}(s)} [r(s) + \sum_{s' \in S} P(s, a)(s')V_{i-1}(s')]$
 - 10: **return** σ
-

Convergence guarantees | To analyze the convergence guarantees of the value iteration algorithm in Algorithm 3.1, let $V^\sigma(s)$ denote the value in state $s \in S$ under a (fixed) scheduler $\sigma \in \mathfrak{S}_{\text{stat}}$:

$$V^\sigma(s) = r(s) + \sum_{s' \in S} P(s, a)(s')V^\sigma(s').$$

Using this notation, $V_{\max}^{\sigma^*}(s)$ is the value in state $s \in S$ under an optimal scheduler σ^* . The convergence of value iteration can then be stated as follows.

Theorem 3.24 (Correctness of value iteration [Put94, Thm. 6.3.1]) For every $\epsilon > 0$, Algorithm 3.1 terminates in finitely many iterations. Furthermore, for $\epsilon = 0$, the values $V_i(s)$ for all $s \in S$ in Algorithm 3.1 converge to $V_{\max}^{\sigma^*}(s)$ as $i \rightarrow \infty$.

Thus, in the limit of infinitely many iterations, value iteration converges to an optimal scheduler. However, no lower bound on the required number of iterations to optimality can be given a-priori. Thus, practical implementations of value iteration use a strictly positive precision parameter $\epsilon > 0$ to ensure the termination of the algorithm.

Other methods | Variants of value iteration exists for many different objectives [CH08], such as long-run average rewards [ACDK⁺17]. Besides value iteration, other classical methods to compute (near) optimal schedulers for MDPs include policy iteration and linear programming. For brevity, we do not discuss these methods in any detail. A more complete introduction of policy iteration is given in [Put94; SB98], and linear programming formulations for MDPs are provided in [BK08; Put94].

Tool support | Mature tools exist that implement algorithms such as value iteration for MDPs. In this thesis, we specifically use the probabilistic model checkers PRISM [KNP11] and Storm [HJKQ⁺22]. Due to the maturity and efficiency of these tools, implementing algorithms such as value iteration ourselves is unnecessary. Hence, we do not discuss these algorithms in more technical detail.

3.2.3.1 Unbounded reach-avoid probabilities

An unbounded reach-avoid probability $\Pr_\sigma^{\mathcal{M}}(s \models \neg S_U \cup S_T)$ for MDP \mathcal{M} can be computed using the value iteration algorithm in Algorithm 3.1 as well. Concretely, we define a modified MDP \mathcal{M}' in which all states in S_T and S_U transition to a dedicated absorbing state $s' \notin S_T \cup S_U$, and which has the reward function

$$r(s) = \begin{cases} 1 & \text{if } s \in S_T \setminus S_U, \\ 0 & \text{otherwise.} \end{cases}$$

Then, running value iteration on the modified MDP \mathcal{M}' converges (as $\epsilon \rightarrow 0$) to a (stationary) scheduler $\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}'}$ that attains the maximum reach-avoid probability in the original MDP \mathcal{M} , i.e., for all $s \in S$,

$$V_i(s) \rightarrow \max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}} \Pr_\sigma^{\mathcal{M}}(s \models \neg S_U \cup S_T).$$

Algorithm 3.2 Value iteration for maximizing over a finite horizon in MDPs.

Input: MDP $\mathcal{M} = (S, Act, s_I, P, L, r)$; time horizon $h \in \mathbb{N}$
Output: Markov scheduler $\sigma = (\sigma_0, \dots, \sigma_h)$, $\sigma_k : S \rightarrow Act$, $k = 0, \dots, h$

```

1:  $V_h(s) = r(s)$ ,  $\forall s \in S$ 
2: for  $i = h - 1, h - 2, \dots, 1, 0$  do
3:   for  $s \in S$  do
4:      $V_i(s) \leftarrow \max_{a \in Act(s)} [r(s) + \sum_{s' \in S} P(s, a)(s')V_{i+1}(s')]$ 
5:      $\sigma_i(s) \leftarrow \max_{a \in Act(s)} [r(s) + \sum_{s' \in S} P(s, a)(s')V_{i+1}(s')]$ 
6: return  $\sigma$ 

```

Thus, we can use value iteration to compute optimal schedulers for unbounded reach-avoid probabilities. In addition, Proposition 3.22 for the optimality of stationary schedulers also carries over. Unbounded reachability probabilities can be computed analogously by setting $S_U = \emptyset$.

3.2.3.2 Bounded reach-avoid probabilities

As already discussed, optimal schedulers for step-bounded PCTL formula may not be stationary and instead depend on the number of steps until the horizon is reached. Thus, stationary schedulers do not suffice to attain optimal values, and we need to consider the (larger) set $\mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ of Markov schedulers instead (as defined in Def. 3.10).

A version of value iteration for finite horizons is described by Algorithm 3.2, which is based on a backward recursion on the value function. Intuitively, $V_i(s)$ and $\sigma_i(s)$ are, respectively, the value and the optimal action in state s at time i . The algorithm computes the values at all previous times $h - 1, h - 2, \dots, 0$, such that $V_0(s)$ is the optimal value at time 0. While Algorithm 3.1 converges to the optimum as $\epsilon \rightarrow 0$, Algorithm 3.2 is *exact*. In other words, the scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ returned by Algorithm 3.2 satisfies

$$\Pr_{\sigma}^{\mathcal{M}}(s \models \neg S_U \cup^{\leq h} S_T) = \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}} \Pr_{\sigma}^{\mathcal{M}}(s \models \neg S_U \cup^{\leq h} S_T).$$

The intuitive argument is that Algorithm 3.2 terminates in exactly h iterations, whereas Algorithm 3.1 may require infinitely many iterations for $\epsilon \rightarrow 0$. Moreover, Algorithm 3.2 extracts the (Markov) scheduler *within* the loop, while Algorithm 3.1 extracts the (stationary) scheduler *after* terminating the loop.

3.3 Robust Markov Decision Processes

An MDP specifies a precise probability distribution over states for every pair $s \in S$ and $a \in Act(s)$. Thus, MDPs cannot capture uncertainty about the probabilistic behavior of a model. Robust Markov decision processes (RMDPs) are a natural extension of MDPs with *sets* of transition probabilities [GLD00; NG05; XM10]. In the literature, RMDPs have also been called *bounded-parameter MDPs* [GLD00] and *uncertain MDPs* [WTM12]. However, RMDP is arguably the most commonly used name. Formally, an RMDP is defined as follows.

Definition 3.25 (RMDP) A *robust MDP* is a tuple $\mathcal{M}_R := (S, \text{Act}, s_I, \mathcal{P}, L, r)$ where (as for MDPs) S is a set of *states*, Act is a set of *actions*, s_I is the *initial state*, $L: S \rightarrow 2^{\text{AP}}$ is a *labeling function*, and $r: S \rightarrow \mathbb{R}_{\geq 0}$ is the *state reward function*. Different from MDPs, the *uncertain probabilistic transition function* $\mathcal{P}: S \times \text{Act} \rightarrow 2^{\text{Distr}(S)}$ maps to sets of probability distributions over S .

Observe that the only element of an RMDP different from an MDP is the uncertain transition function \mathcal{P} . As for MDPs, the transition function for an RMDP is partial in general, as only a subset $\text{Act}(s) \subseteq \text{Act}$ of actions may be enabled in every state $s \in S$. For each $s \in S$ and $a \in \text{Act}(s)$, we call $\mathcal{P}(s, a) \subseteq \text{Distr}(S)$ an *uncertainty set*. For simplicity, we assume that for all $s \in S$ and $a \in \text{Act}(s)$, the uncertainty set $\mathcal{P}(s, a)$ in Def. 3.25 is non-empty. Analogous to Def. 3.4, a *robust Markov chain* is an RMDP with exactly one enabled action in every state.

Remark 3.26 (Geometry of uncertainty sets) Our definition of an RMDP in Def. 3.25 assumes that the uncertainty sets for all state-action pairs are independent. This assumption is known as the *rectangularity* assumption, which is common in the literature on RMDPs [WKR13] and robust optimization [BBC11], and is necessary to guarantee computational tractability (in fact, most problems for non-rectangular RMDPs are proven to be NP-hard [WKR13]). Moreover, in practice, we will use RMDPs in which every uncertainty set $\mathcal{P}(s, a)$ is a convex polytope, which is another common assumption in the literature.

3.3.1 Interval MDPs

An *interval Markov decision process (IMDP)* is a special case of RMDP where the uncertain transition functions are given as intervals [BGN09].

Definition 3.27 (IMDP) An *interval MDP* is an RMDP $\mathcal{M}_R = (S, \text{Act}, s_I, \mathcal{P}, L, r)$ where the transition function \mathcal{P} is defined such that for all $s \in S$ and $a \in \text{Act}(s)$,

$$\begin{aligned} \mathcal{P}(s, a) = \{ & \mu \in \text{Distr}(S) : \forall s' \in S, \check{p}(s, a, s') \leq \mu(s') \leq \hat{p}(s, a, s') \\ & 0 \leq \check{p}(s, a, s') \leq \hat{p}(s, a, s') \leq 1 \}. \end{aligned}$$

For clarity, we identify interval Markov decision processes (IMDPs) by writing \mathcal{M}_I . Paths, sets of paths, and schedulers for RMDPs (and IMDPs) are defined and denoted analogously to those for MDPs.

Convex uncertainty sets | Intuitively, for each transition (s, a, s') , the uncertain transition function of an IMDP is the interval $[\check{p}(s, a, s'), \hat{p}(s, a, s')]$, plus the condition that each set $P(s, a)$ contains only valid probability distributions. Geometrically speaking, for each state-action pair (s, a) , the set of distributions $\mathcal{P}(s, a)$ can be thought of as the intersection of the probability simplex over S and the hyperrectangle induced by the intervals for each $s' \in S$. Thus, IMDPs have convex uncertainty sets.

robust MDP

3

uncertainty set

robust MC

rectangularity

interval MDP

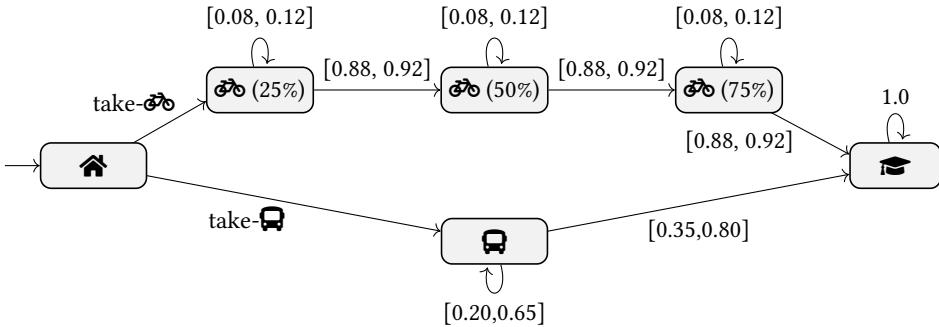


Figure 3.2: The scenario of traveling from home (🏠) to university (🎓), with uncertainty in the probabilities of delays when going by bike and by bus.

Example 3.28 In the bike-versus-bus example MDP from Fig. 3.1, we assumed that we know the probability for a bus or bike delay *exactly*. In practice, however, these probabilities will typically be uncertain, and the MDP in Fig. 3.1 will only be an approximation of the scenario. By contrast, we can also take a more risk-averse stance and instead say that we only know these probabilities up to a given interval. For example, suppose that we have enough data to say that the probability for a bus to be canceled is in the interval $[0.20, 0.65]$, and similarly, the probability for a bike delay is $[0.08, 0.12]$. Then, we can model our scenario as the interval Markov decision process (IMDP) in Fig. 3.2.

3.3.2 Nature

Loosely speaking, an RMDP \mathcal{M}_R can be interpreted as a set of MDPs varying only in their transition function. By fixing a precise transition function within the set of transition functions \mathcal{P} , an RMDP reduces to a standard MDP. We adopt the common terminology that this choice of transition function from the uncertainty set \mathcal{P} is made by (the scheduler of) *nature* [NG05; BSJJ24].

nature

Definition 3.29 (Nature) A (scheduler of) *nature* for an RMDP $\mathcal{M}_R = (S, \text{Act}, s_I, \mathcal{P}, L, r)$ is a function $\tau: \Pi_{\text{fin}}^{\mathcal{M}_R} \times \text{Act} \rightarrow \mathcal{P}$ mapping each finite path and an action to a possible transition function $P \in \mathcal{P}$. The set of all schedulers of nature for an MDP \mathcal{M}_R is denoted by $\mathfrak{T}^{\mathcal{M}_R}$.

Semantics | The semantics of applying a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}_R}$ together with a nature $\tau \in \mathfrak{T}^{\mathcal{M}_R}$ on an RMDP $\mathcal{M}_R = (S, \text{Act}, s_I, \mathcal{P}, L)$ are as follows. Given a finite path $\pi \in \Pi_{\text{fin}}^{\mathcal{M}_R}$ of the RMDP execution thus far, the scheduler picks an action $\tilde{a} = \sigma(\pi)$, and nature picks a choice of transition function $\tilde{P} = \tau(\pi, \tilde{a})$. Then, the successor state is sampled according to the probability distribution $\tilde{P}(\text{last}(\pi), \tilde{a})$. As with regular MDPs, this process is repeated indefinitely.

Restricting nature | As with schedulers, it is unnecessary to consider the full set $\mathfrak{T}_{\text{stat}}^{\mathcal{M}_R}$ of all (infinitely many) schedulers of nature [Iye05]. Analogous to stationary schedulers, we can choose to restrict the choice of transition function to be fixed *once and for all* for every state-action pair (called the *static uncertainty model*).

Definition 3.30 (Stationary nature) A nature τ is *stationary* if for all paths $\pi, \pi' \in \Pi_{\text{fin}}^{\mathcal{M}}$ and for all $a \in \text{Act}(\text{last}(\pi))$, it holds that

$$\text{last}(\pi) = \text{last}(\pi') \implies \tau(\pi, a) = \tau(\pi', a).$$

We simplify notation and write stationary natures as $\tau: S \times \text{Act} \rightarrow \mathcal{P}$. The set of all stationary natures for an RMDP \mathcal{M}_R is denoted by $\mathfrak{T}_{\text{stat}}^{\mathcal{M}_R}$.

stationary
nature

Similarly, we can weaken this requirement by restricting the choice of transition function to be fixed *once per time step* for every state-action pair (called the *dynamic uncertainty model*).

Definition 3.31 (Markov nature) A nature τ is *Markov* (or *Markovian*) if for all paths $\pi, \pi' \in \Pi_{\text{fin}}^{\mathcal{M}}$ and for all $a \in \text{Act}(\text{last}(\pi))$, it holds that

$$\text{last}(\pi) = \text{last}(\pi') \wedge |\pi| = |\pi'| \implies \tau(\pi, a) = \tau(\pi', a).$$

Markov
nature

We simplify notation and write Markov natures as a sequence of stationary natures, $\tau = (\tau_0, \tau_1, \dots)$, where $\tau: S \times \text{Act} \rightarrow \mathcal{P}$ for $k \in \mathbb{N}$. The set of all Markov natures for an RMDP \mathcal{M}_R is denoted by $\mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$.

As we shall see, and equivalent to schedulers, stationary natures are sufficient to attain optimal values for unbounded measures, whereas Markov natures are needed for step-bounded measures.

3.3.3 Robust measures

Applying a scheduler and a nature to an RMDP induces a DTMC with a well-defined probability measure $\Pr_{\sigma, \tau}^{\mathcal{M}_R}: \Pi_{\text{fin}}^{\mathcal{M}_R} \rightarrow [0, 1]$. This probability measure is defined equivalently as for MDPs, but we replace the transition function P in Eq. (3.2) with the transition function $\tau(s, a)$ chosen by a stationary nature, or by the transition function $\tau_k(s, a)$ chosen by a Markov nature. We can reason over all of the measures introduced in Sect. 3.2.2 on the resulting induced DTMC.

Often, we are interested in the maximal or minimal value of a measure *under any* choice of nature. For example, given a scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}_R}$ for RMDP \mathcal{M}_R , we want to compute the *robust values* for a PCTL path formula φ :

$$\min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi), \quad \text{and} \quad \max_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi).$$

robust value

Similarly, for expected cumulative rewards, we want to compute the robust values

$$\min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \text{ExpRew}_{\sigma, \tau}^{\mathcal{M}_R}(s \models \diamond B), \quad \text{and} \quad \max_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \text{ExpRew}_{\sigma, \tau}^{\mathcal{M}_R}(s \models \diamond B).$$

These robust values are, by definition, upper and lower bounds on the values for any fixed nature $\tau' \in \mathfrak{T}^{\mathcal{M}_R}$. We state this trivial yet useful result without proof in the following “*sandwich lemma*.¹

Lemma 3.32 (RMDP sandwich lemma) Let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P}, L, r)$ be an RMDP with scheduler $\sigma \in \mathfrak{S}^{\mathcal{M}_R}$ and a state $s \in S$, let $\tilde{\tau} \in \mathfrak{T}^{\mathcal{M}_R}$ be a nature, and let φ be a PCTL path formula. It holds that

$$\min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi) \leq \Pr_{\sigma, \tilde{\tau}}^{\mathcal{M}_R}(s \models \varphi) \leq \max_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi).$$

In addition, let $B \subseteq S$ be a set of states. It holds that

$$\begin{aligned} \min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \text{ExpRew}_{\sigma, \tau}^{\mathcal{M}_R}(s \models \Diamond B) &\leq \text{ExpRew}_{\sigma, \tilde{\tau}}^{\mathcal{M}_R}(s \models \Diamond B) \\ &\leq \max_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \text{ExpRew}_{\sigma, \tau}^{\mathcal{M}_R}(s \models \Diamond B). \end{aligned}$$

3.3.4 Optimal schedulers for RMDPs

We now have two axes to maximize or minimize over: the scheduler and the nature. This yields four cases for maximizing or minimizing over the schedulers and maximizing or minimizing over the natures. Suppose that we have a PCTL path formula φ for which we compute a maximizing or minimizing scheduler. Asking for this probability under the most *optimistic nature*, we obtain

$$\hat{\sigma}_{\max}^* \in \operatorname{argmax}_{\sigma \in \mathfrak{S}^{\mathcal{M}_R}} \max_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi), \text{ or } \check{\sigma}_{\min}^* \in \operatorname{argmin}_{\sigma \in \mathfrak{S}^{\mathcal{M}_R}} \min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi). \quad (3.3)$$

pessimistic nature

Similarly, we can also ask for this probability under the most *pessimistic nature*:

$$\check{\sigma}_{\max}^* \in \operatorname{argmax}_{\sigma \in \mathfrak{S}^{\mathcal{M}_R}} \min_{\tau \in \mathfrak{T}_{\text{stat}}}^{\mathcal{M}_R} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi), \text{ or } \hat{\sigma}_{\min}^* \in \operatorname{argmin}_{\sigma \in \mathfrak{S}^{\mathcal{M}_R}} \max_{\tau \in \mathfrak{T}_{\text{stat}}}^{\mathcal{M}_R} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi). \quad (3.4)$$

Pessimistic values are commonly called *robust*.

Proposition 3.33 ([NG05, Theorem 1]) The minimize/maximize operators over nature and the scheduler in Eqs. (3.3) and (3.4) are interchangeable.

In practice, whether we maximize or minimize over the schedulers will be clear from the context.

Lemma 3.34 (Static vs. dynamic uncertainty [Iye05]) For stationary schedulers $\sigma \in \mathfrak{S}_{\text{stat}}$, the optimal values for general, stationary, and Markov natures

¹In fact, this result holds for any measure, but we state it for the measures used in this thesis only.

coincide. For example, it holds that

$$\begin{aligned} \max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_R}} \min_{\tau \in \mathfrak{T}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi) &= \max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_R}} \min_{\tau \in \mathfrak{T}_{\text{stat}}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi) \\ &= \max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_R}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s \models \varphi), \end{aligned}$$

and the same holds for the other formulations in Eqs. (3.3) and (3.4).

Lemma 3.34 states that, when dealing with stationary schedulers, it suffices to work with stationary natures. By contrast, for Markov schedulers, we instead need to use Markov natures $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$. The following statement formalizes these observations.

Proposition 3.35 ([NG05, Corollary 2]) For unbounded PCTL formulae and expected cumulative rewards, stationary (deterministic) schedulers $\sigma \in \mathfrak{S}_{\text{stat}}$ suffice to attain optimal values. For step-bounded formulae, Markov (deterministic) schedulers $\sigma \in \mathfrak{S}_{\text{Markov}}$ suffice to attain optimal values.

Robust value iteration | The value iteration algorithms in Algorithms 3.1 and 3.2 can be adapted to compute optimal policies under the minimizing or maximizing nature [WKR13; Iye05; NG05]. For IMDPs, these algorithms are implemented in the probabilistic model checkers PRISM [KNP11] and Storm [HJKQ⁺22]. We use these model checkers when analyzing IMDPs in this thesis. For brevity, we omit further details on robust value iteration and instead refer to the references above.

3.3.5 Connection to other models

RMDPs are closely related to several other modeling formalisms. In this section, we briefly discuss the most prominent connections. However, further exploration of these connections is beyond the scope of this thesis.

Parametric MDPs | RMDPs are closely related to parametric Markov decision processes (pMDPs), which we formally define in Chapter 8. Intuitively, a pMDP is an MDP where the transition probabilities are given as polynomial functions over a given set of parameters [JJK22; Jun20]. The same parameters can be shared between transitions, making pMDPs a powerful but often expensive-to-analyze modeling formalism. By assigning a lower and upper bound to every parameter, we can effectively define an RMDP with a stationary nature.

Stochastic games | An RMDP can also be interpreted as a stochastic game between the scheduler (player one) and nature (player two) [Iye05; WKR13; XM12; NG05; GG23]. Stochastic games are usually defined with finite action spaces for both players, whereas in an RMDP, nature generally has *infinitely many* choices (e.g., the continuous probability intervals an IMDP). However, for the convex uncertainty sets we consider in this thesis (see Remark 3.26), optimistic and pessimistic optimal values are attained at the *vertices of the uncertainty sets* [NG05]. Intuitively, an RMDP (with convex uncertainty set) can thus be defined as a stochastic game, with (in every state) an action for player two for every vertex of the uncertainty set.

Probabilistic automata | The connection with probabilistic automata is similar to the connection with stochastic games. Probabilistic automata combine nondeterminism with probabilistic behavior and subsume MDPs [Seg95; SL95; Sto02]. The transition function of a probabilistic automaton is usually defined as $\Delta \subseteq S \times Act \times \text{Distr}(S)$, where S is the set of states, and Act is the set of actions. Thus, multiple distributions over states may be enabled for the same state-action pair, which is analogous to the choice of nature in an RMDP. Like stochastic games, probabilistic automata are usually defined with a finite action space and thus capture discrete sets of probability distributions (whereas RMDPs usually consist of continuous sets of distributions).

Summary

- ⇒ Markov decision processes (MDPs) are probabilistic models with non-deterministic action choices.
- ⇒ Schedulers resolve the nondeterministic action choices in MDPs.
- ⇒ Common measures of interest can be specified in probabilistic computation tree logic (PCTL).
- ⇒ Robust MDPs (RMDPs) extend MDPs with sets of transition probabilities.

Part II

Discrete-Time Stochastic Systems

4 Foundations of DTSSs

Summary | We study a version of Markov decision processes (MDPs) with continuous state and action spaces. A common formalism for such models is a discrete-time stochastic system (DTSS). In this background chapter, we introduce the foundations of DTSSs and define schedulers for them (which we call policies to distinguish from finite MDPs). We discuss common control problems for DTSSs, focusing on reach-avoid control tasks. We show how to characterize the probability that a given reach-avoid task is satisfied when applying a given policy to the DTSS. Finally, we discuss that, due to the continuous and stochastic nature of DTSSs, computing this satisfaction probability exactly is intractable in general.

Origins | This chapter does not contain novel content and instead provides the foundations for the subsequent chapters. For a much more in-depth treatment of discrete-time stochastic control, we refer to [BS78].

Background | We assume the reader is familiar with MDPs (Def. 3.1) and robust MDPs (Def. 3.25), and with computing optimal (robust) schedulers for reach-avoid probabilities. Moreover, to capture stochastic uncertainty in models, we build upon the measure- and probability-theoretic definitions from Sect. 2.3.



4.1 Introduction

In Chapter 3, we have discussed Markov decision processes (MDPs) and their robust variants, which we defined to have *finite* state and action spaces. However, in many situations, these finite models do not suffice, especially when modeling physical systems that are inherently continuous. Instead, we need stochastic models with *continuous* state and action spaces. In this chapter, we introduce *discrete-time stochastic systems* (DTSSs) as models for complex controlled systems with continuous state and action spaces [KGS07]. DTSSs are used in many areas around control theory, including forecasting, controller design, and filtering techniques [Söd02].

UAV motion planning | As an example, an unmanned aerial vehicle (UAV) can be modeled as a discrete-time dynamical system, where the continuous state reflects variables such as the current position and velocity, and the continuous actions (which we will call *control inputs*) reflect actuator inputs that change the state over time. Simply discretizing the continuous states and actions means we incur a discretization error, which is unacceptable in safety-critical settings. Furthermore, factors such as wind gusts and actuator imprecision cause uncertainty about the evolution of the system's state [BOBW10]. We model such uncertainty as a *stochastic disturbance* (also called

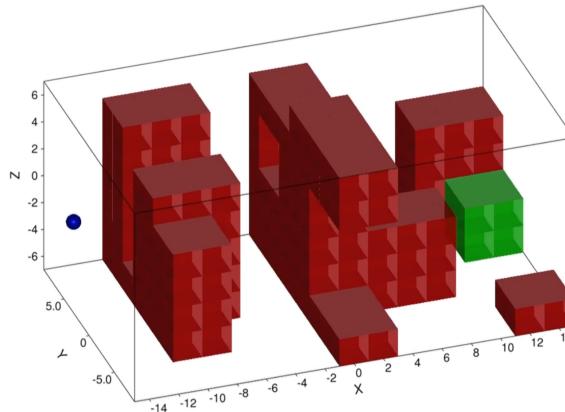


Figure 4.1: A reach-avoid problem for a UAV, which needs to navigate from the blue point to the green target while avoiding the red obstacles. In fact, we solve this particular task for a DTSS with linear dynamics in Chapter 6.

process noise) that affects the state transitions of the system. As a result, we obtain a model of the UAV and its environment in the form of a DTSS.

Reach-avoid control | A common task for a DTSS is to *reach* a desired target region while *avoiding* collision with certain obstacles [KIKF22]. Reach-avoid tasks are ubiquitous in controller design [PAQM18; SKLT11; FCTS15; ECL11; YTCB⁺12], e.g., in motion planning and process control. For example, the UAV in Fig. 4.1 must reach the target region (green) within a desired time horizon, while avoiding crashing into the obstacles (red). Often, a reach-avoid task needs to be satisfied with at least a certain probability. Such a control problem is similar to the probabilistic computation tree logic (PCTL) *reach-avoid formulae* for MDPs that we considered in Chapter 3 (see Def. 3.19). While our methods can readily be extended to general PCTL formulae for DTSS, we specifically focus on reach-avoid tasks in this thesis for brevity.

4.2 Discrete-Time Stochastic Systems

In this section, we provide the foundations of DTSSs and discuss how to express reach-avoid control problems. The material in this section largely follows the definitions in [BS78; APLS08; SL10], although the latter two references consider discrete-time stochastic *hybrid* systems, which extend DTSSs with discrete behavior. Here, we neglect partial observability and assume that the state of the system is fully observable.

To ensure that the probabilistic behavior of a DTSS is well-defined, we need to constrain the state and action spaces to so-called *Borel spaces* (loosely speaking, the inexperienced reader may simply interpret a Borel space as any “nice” subset of an n -dimensional Euclidean space \mathbb{R}^n , or \mathbb{R}^n itself).

Definition 4.1 (Borel space) Let (Y, d) be a complete and separable metric space^a and let $X \in \mathcal{B}(Y)$, where $\mathcal{B}(Y)$ are the Borel sets of Y induced by the topology given

by the metric d . The measurable space $(X, \mathcal{B}(X))$ is called a *Borel space*, where the σ -algebra $\mathcal{B}(X)$ is defined as

$$\mathcal{B}(X) := \{X \cap B : B \in \mathcal{B}(Y)\} = X \cap \mathcal{B}(Y).$$

When clear from the context, we omit the σ -algebra $\mathcal{B}(X)$ and simply write X for a Borel space $(X, \mathcal{B}(X))$.

^aFor a formal introduction of complete and separable metric spaces, we refer to [BS78].

A simple example of a Borel space is $(\mathbb{R}^n, \mathcal{B}(\mathbb{R}^n))$ for any dimension $n \in \mathbb{N}$, i.e., the set \mathbb{R}^n endowed with the Borel σ -algebra over \mathbb{R}^n . Moreover, any Borel subset of a Borel space is itself a Borel space.

The continuous state of a DTSS evolves dynamically over discrete time steps according to some fixed update rule (a.k.a. transition function) that depends on a chosen action. Like in an MDP, this update rule is stochastic, so repeating the same action twice in the same state may lead to a different successor state. We formally define a DTSS as follows.

Definition 4.2 (DTSS) A *discrete-time stochastic system (DTSS)* is a tuple $\mathcal{S} := (X, U, x_I, \zeta, f)$, where

- $X \subseteq \mathbb{R}^n$ is a Borel space, called the state space of the system;
- $U \subseteq \mathbb{R}^m$ is a Borel space, called the (control) input space of the system;
- $x_I \in X$ is the initial state;
- $\zeta = (\zeta_k)_{k \in \mathbb{N}}$ is a discrete-time stochastic process, called the process noise, which is defined on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ with its natural filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$ (see Def. 2.10), where for all $k \in \mathbb{N}$, the random variable $\zeta_k: \Omega \rightarrow \mathcal{V}_\zeta$ maps to the same measurable space $(\mathcal{V}_\zeta, \mathcal{F}_\zeta)$;
- $f: X \times U \times \mathcal{V}_\zeta \rightarrow X$ is a measurable function, called the state transition function of the system, which characterizes the state evolution of the system.

discrete-time stochastic system

In practice, we only consider DTSSs defined with *relatively simple* Borel spaces, such as \mathbb{R}^n or compact subsets of \mathbb{R}^n . Thus, in Def. 4.2, a Borel space can be loosely interpreted as any “nice subset of the Euclidean space” (such as a compact or polytopic subset of \mathbb{R}^n), endowed with its Borel σ -algebra. The DTSS in Def. 4.2 is called *time-invariant* because the transition function f is independent of the time step $k \in \mathbb{N}$.

time-invariance

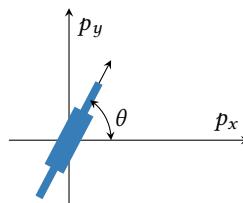


Figure 4.2: Top-down view of the unicycle model.

Example 4.3 (Adapted from [LaV06]) A simple unicycle model (also called a Dubins vehicle) is described by its x and y coordinates (denoted by p_x and p_y) and its orientation θ (see Fig. 4.2). In this simplified model, there are two input variables: the driver can set the driving speed (v) and set the change in steering angle (r), both of which are bounded by a minimum and maximum value denoted by v_{\min} , v_{\max} , r_{\min} , and r_{\max} , respectively. However, the actual change in steering angle is noisy, which we model by an additive Gaussian random variable. The unicycle can be modeled as a DTSS, where

- The state space is $X = \mathbb{R}^3$, such that $x = [p_x, p_y, \theta] \in X$;
- The input space is $U = [v_{\min}, v_{\max}] \times [r_{\min}, r_{\max}] \subset \mathbb{R}^2$;
- The initial state is $x_I \in X$;
- The stochastic process $\zeta = (\zeta_k)_{k \in \mathbb{N}}$ is defined such that each $\zeta_k = \mathcal{N}(0, \Sigma)$ is a zero-mean Gaussian random variable with covariance matrix Σ ;
- The transition function f is defined, for all $[p_x, p_y, r] \in X$ and $[v, r] \in U$, as

$$f([p_x, p_y, r]^T, [v, r]^T, \zeta_k) = \begin{bmatrix} p_x + \delta_t \cdot v \cdot \cos \theta \\ p_y + \delta_t \cdot v \cdot \sin \theta \\ \theta + \delta_t \cdot (r + \zeta_k) \end{bmatrix},$$

where $\delta_t \in \mathbb{R}_{>0}$ is a discretization time.

process noise

The stochastic process ζ is also called the *process noise* and is a stochastic disturbance affecting the state transitions. Specifically, each random variable ζ_k for $k \in \mathbb{N}$ is a mapping from the measurable space (Ω, \mathcal{F}) to the measurable space $(\mathcal{V}_\zeta, \mathcal{F}_\zeta)$. As such, each ζ_k also induces a mapping of the probability measure \mathbb{P} on (Ω, \mathcal{F}) to a probability measure \mathbb{P}_{ζ_k} on $(\mathcal{V}_\zeta, \mathcal{F}_\zeta)$. We make the following (standard) assumptions on the distribution of the process noise.

Assumption 4.4 (Noise i.i.d.) For all $k \in \mathbb{N}$, let $\mathbb{P}_{\zeta_k} : \mathcal{F}_\zeta \rightarrow [0, 1]$ be the induced measure by ζ_k on \mathbb{P} . We assume that

- (*independent*): For every finite collection of sets $V_1, \dots, V_m \in \mathcal{F}_\zeta$, with $m \in \mathbb{N}$, it holds that $\mathbb{P}_{(\zeta_1, \dots, \zeta_m)}(\cap_{i=1}^m V_i) = \prod_{i=1}^m \mathbb{P}_{\zeta_i}(V_i)$, where $\mathbb{P}_{(\zeta_1, \dots, \zeta_m)}$ is the product probability measure over $(\zeta_1, \dots, \zeta_m)$;
- (*identically distributed*): For all $k, l \in \mathbb{N}$ and for all $V \in \mathcal{F}_\zeta$, it holds that $\mathbb{P}_{\zeta_k}(V) = \mathbb{P}_{\zeta_l}(V)$.

Remark 4.5 (DTSS as a continuous MDP) The DTSS in Def. 4.2 can equivalently be represented as an MDP with an uncountably infinite number of states S (representing the state space $X \subset \mathbb{R}^n$ of the DTSS), and an uncountably infinite number of actions Act (representing the input space $U \subset \mathbb{R}^m$ of the DTSS). While interesting, this connection does not provide more insights that add to the contributions of this thesis. Thus, we do not explore this connection in more detail in this thesis and instead refer to [Put94, Section 2.3.2] for further details.

4.2.1 Markov policy

The inputs to a DTSS are chosen based on a *Markov policy*,¹ which is a sequence of decision rules for each time step $k \in \mathbb{N}$. Each decision rule is memoryless (i.e., it only uses the current state x_k to determine the input u_k) and is a measurable function from X to U (needed to make the closed-loop system ‘well-behaved’ [BS78]).

Definition 4.6 (Markov policy) A *Markov policy* for a DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ is a sequence $\mu := (\mu_0, \mu_1, \mu_2, \dots)$ of measurable maps $\mu_k: X \rightarrow U, k \in \mathbb{N}$ from states $x \in X$ to control inputs $u \in U$.

The semantics for executing a Markov policy $\mu = (\mu_k)_{k \in \mathbb{N}}$ on a DTSS are as follows. The initial state $x_0 := x_I$ is given by the DTSS, and for all discrete steps $k \in \mathbb{N}$, the next state x_{k+1} is determined recursively as

$$x_{k+1} := f(x_k, \mu_k(x_k), \varsigma_k). \quad (4.1)$$

Thus, each state x_{k+1} for $k \in \mathbb{N}$ is determined by the Markov policy μ and the values $\varsigma_0, \varsigma_1, \dots, \varsigma_k$ of the process noise up to time k .

Example 4.7 For the unicycle model from Example 4.3, consider the control problem of tracking the p_x axis while avoiding y-coordinates p_y for which $|p_y| \geq 1$, starting from the origin $x_I = [0, 0, 0]$. Our task is to design a Markov policy μ as in Def. 4.6 that achieves this *tracking task*. For simplicity, we use a constant speed of $v = \bar{v} > 0$ so that the Markov policy only sets the change to the steering angle r . Intuitively, if $p_y > 0$, we want to reduce the steering angle, while if $p_y < 0$, we want to increase the steering angle. Thus, a naïve *proportional* Markov policy $\mu = (\mu_0, \mu_1, \mu_2, \dots)$ could be linear in the y-coordinate, such that

$$\mu_k = \begin{bmatrix} \bar{v} \\ -\alpha \cdot p_y \end{bmatrix} \quad \forall k \in \mathbb{N},$$

where $\alpha \in \mathbb{R}_{>0}$ is a coefficient. However, while this simple Markov policy may work decently in most cases, it is unclear whether the closed-loop system satisfies the tracking task *always*, or, for example, with a probability of at least 99%.

The example above illustrates the need for more powerful methods that can formally verify the performance of a closed-loop system consisting of a DTSS and a Markov policy. In subsequent chapters of Part II of this thesis, we will present such a method, which is based on formally relating the system with a simpler, finite-state abstraction.

Execution of DTSS | Let us investigate the probabilistic behavior of a DTSS in more detail. Each $\varsigma_k: \Omega \rightarrow \mathcal{V}_\varsigma$ is a random variable measurable with respect to the k^{th} element \mathcal{F}_k of the natural filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$. Since the state transition function f and the Markov policy μ are measurable functions, each $x_{k+1}: \Omega \rightarrow X$ is a random variable measurable with respect to the pair $(\mathcal{F}_k, \mathcal{B}(X))$. Hence, for a fixed initial state $x_I \in X$ and

¹In control theory, Markov policies are better known as (time-varying) feedback controllers or control laws. For clarity, we use the term *policy* for DTSSs, while using *scheduler* for discrete MDPs.

execution sample path a Markov policy μ , the *execution* $(x_k)_{k \in \mathbb{N}}$ is a discrete-time stochastic process defined on the probability space $(\Omega, \mathcal{F}, \mathbb{P})$, adapted to the filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$. By fixing an element $\omega \in \Omega$ and a Markov policy μ , we thus obtain the *sample path* $x_I(\omega), x_1(\omega), x_2(\omega), \dots$

4.2.2 Stochastic kernel

We will now introduce an alternative yet equivalent representation of the DTSS that often leads to more concise definitions and derivations. In particular, we replace the transition function f and the process noise ς with a *stochastic kernel* $T: \mathcal{B}(X) \times X \times U \rightarrow [0, 1]$, which is a collection of probability measures on X , parameterized by X and U [Kal02]. For each $V \in \mathcal{B}(X)$, we define the kernel T as the probability for x_{k+1} to be contained in V , conditioned on the current state x and input u :

$$T(V | x, u) = \mathbb{P}\{\omega \in \Omega : f(x, u, \varsigma_k(\omega)) \in V\}. \quad (4.2)$$

The vertical bar indicates that $T(V | x, u)$ is the probability for $V \in \mathcal{B}(X)$, *conditioned* on $(x, u) \in X \times U$. Using the kernel notation, Eq. (4.1) is equivalently expressed as

$$x_{k+1} \sim T(\cdot | x_k, \mu_k(x_k)), \quad (4.3)$$

where the symbol \sim means that $x_{k+1} \in X$ is sampled according to $T(\cdot | x_k, \mu_k(x_k))$. Thus, in both Eqs. (4.1) and (4.3), $x_{k+1}: \Omega \rightarrow X$ is a random variable, which takes on a concrete value $x_{k+1}(\omega) \in X$ for a given value $\omega \in \Omega$. Throughout the thesis, we will equivalently use the transition function and kernel notations where convenient.

It has been shown [APLS08; SL10; BS78] that the DTSS execution induces a probability measure over paths x_I, x_1, x_2, \dots , which is uniquely defined by the stochastic kernel T , the initial state $x_I \in X$, and the Markov policy μ .

Proposition 4.8 (DTSS probability measure) Let $\mathcal{S} = (X, U, x_I, \varsigma, f)$ be a DTSS and let μ be a Markov policy. The execution $(x_k)_{k \in \mathbb{N}}$ is defined on the measurable space $(\Omega', \mathcal{B}(\Omega'))$, with the sample space $\Omega' = X \times X \times \dots$ endowed with its product σ -algebra $\mathcal{B}(\Omega')$. The execution $(x_k)_{k \in \mathbb{N}}$ induces a *probability measure* $\mathbb{P}_\mu^{\mathcal{S}}: \mathcal{B}(\Omega') \rightarrow [0, 1]$ that is uniquely defined by the stochastic kernel T , the initial state $x_I \in X$, and the Markov policy μ .

Proof. The proposition directly follows from applying [BS78, Proposition 7.45] to the stochastic process $(x_k)_{k \in \mathbb{N}}$ for the execution of the DTSS. ■

As a consequence of Proposition 4.8, we can use $\mathbb{P}_\mu^{\mathcal{S}}$ to define the probability that the execution $(x_k)_{k \in \mathbb{N}}$ is contained in *any Borel set* $V \in \mathcal{B}(\Omega')$. In the next section, we will use this machinery to formally express the control objective that we consider.

4.3 Reach-Avoid Probability

Our goal is to find a Markov policy μ such that the DTSS execution $(x_k)_{k \in \mathbb{N}}$ satisfies some objective (with high probability). In this thesis, we focus on *reach-avoid* objectives, which are analogous to the PCTL reach-avoid formulae introduced in Chapter 3 for MDPs. Specifically, we consider the objective of reaching a desired set of states $X_T \subset X$ in at most $h \in \mathbb{N} \cup \{\infty\}$ steps, while always avoiding unsafe states $X_U \subset X$.

Definition 4.9 (Reach-avoid specification) A *reach-avoid specification* for a DTSS \mathcal{S} is a tuple $\varphi := (X_T, X_U, h)$, where $X_T \subset X$ is a Borel set of target states, $X_U \subset X$ is a Borel set of unsafe states, and $h \in \mathbb{N} \cup \{\infty\}$ is a horizon.

Observe that the horizon is allowed to be infinite. Without loss of generality, we assume that the target states X_T and the unsafe states X_U are disjoint.²

Assumption 4.10 (Target and unsafe states disjoint) The set of target and unsafe states are assumed to be disjoint, i.e., $X_T \cap X_U = \emptyset$.

Intuitively, a sample path $x_I(\omega), x_1(\omega), x_2(\omega)$ for $\omega \in \Omega$ satisfies a given reach-avoid specification φ if there exists a $k \in \{0, \dots, h\}$ such that $x_k(\omega) \in X_T$ and $x_{k'}(\omega) \notin X_U$ for all $k' \in \{0, \dots, k\}$. Recall from Proposition 4.8 that the execution $(x_k)_{k \in \mathbb{N}}$ is defined on the probability space $(\Omega', \mathcal{B}(\Omega'), \mathbb{P}_\mu^S)$, where the sample space is the (countably infinite) Cartesian product $\Omega' = X \times X \times \dots$. Thus, we can use this probability space to reason over the probability of satisfying a reach-avoid specification.

Definition 4.11 (Satisfaction probability) Let $\mathcal{S} = (X, U, x_I, \varsigma, f)$ be a DTSS, let $\mu = (\mu_0, \mu_1, \dots)$, $\mu_k: \mathbb{R}^n \rightarrow U$ for all $k \in \mathbb{N}$, be a Markov policy, and let $\varphi = (X_T, X_U, h)$ be a reach-avoid specification. The satisfaction probability of φ under μ is defined as

$$\Pr_\mu^S(x_I \models \varphi) := \mathbb{P}_\mu^S \left\{ \exists k \in \{0, \dots, h\} : x_k \in X_T \wedge (x_{k'} \notin X_U \forall k' \in \{0, \dots, k\}) \right\}. \quad (4.4)$$

We remark that in our definition, we do not care about what happens *after* reaching the target states X_T . Moreover, reaching the target states may happen at any step within the horizon (not just at the end of the horizon), which is also referred to as the *first hitting time* reach-avoid problem by [SL10].

4.3.1 Computing satisfaction probabilities

It has been shown by [SL10] that the satisfaction probability $\Pr_\mu^S(x_I \models \varphi)$ can be characterized as a value function over the state space. However, as we will show next, actually *computing* the satisfaction probability using this value function is often intractable, even for a fixed Markov policy μ .

We sketch the characterization of the reach-avoid probability as proposed in [SL10]. Consider a reach-avoid specification $\varphi = (X_T, X_U, h)$ with a finite horizon $h \in \mathbb{N}$. The main idea is to define a *backward recursion*, which is initialized as the satisfaction probability (in terms of a value function over the state space) at the very end of the horizon, i.e., at time step h . From this final time step h , we then work *backward* until we reach the initial time step of zero.

First, suppose that we are at the final time step h , which means that we are at the very end of the horizon and cannot choose any action anymore. Thus, the satisfaction probability at this step is one for all states $x \in X_T$ and zero elsewhere. Consequently,

²If the target and unsafe states would not be disjoint, we shrink the target states to $X_T \setminus X_U$.

the value function $V_h^\mu(x) : X \rightarrow [0, 1]$ at time step k is defined for all $x \in X$ as

$$V_h^\mu(x) = \mathbb{1}_{X_T}(x).$$

For previous steps $k < h$, the value in the state $x \in X$ is

- one if $x \in X_T$,
- zero if $x \in X_U$, and
- determined by the value at the next time step $k + 1$ if $x \in X \setminus (X_T \cup X_U)$.

For brevity, define $X_S = X \setminus (X_T \cup X_U)$. Taking into account the distribution over states at the next step $k + 1$, we obtain the value function

$$V_k^\mu(x) = \mathbb{1}_{X_T}(x) + \mathbb{1}_{X_S}(x) \int_{\mathbb{R}^n} \mathbb{1}_X(\xi) \cdot V_{k+1}^\mu(\xi) \cdot T(d\xi | x, \mu_k(x)), \quad \forall x \in X.$$

In other words, the value function at step k is defined by the probability of reaching a target state $x \in X_T$, combined with the value function at step $k + 1$. By repeating this recursive definition until step $k = 0$, we obtain the satisfaction probability at the initial time step:

$$\Pr_\mu^S(x_I \models \varphi) = V_0^\mu(x_I).$$

The main challenge in actually computing $V_0^\mu(x_I)$ lies in the integration over the stochastic kernel (which represents a potentially complex distribution over the state space). In particular, we either have to represent the value functions V_k^μ for $k = 0, \dots, h-1$ explicitly, or we have to compute a nesting of h integrals directly.

An alternative is to approximate the value function, for example, by numerical integration. However, we argue that such approximative methods are undesirable for safety-critical settings. Thus, in this thesis, we focus on computing *sound lower bounds* on the satisfaction probability instead.

4.3.2 Extension to PCTL

All of the results for DTSS that we present in this thesis can be extended to general PCTL specifications by equipping the DTSS with a labeling $L : X \rightarrow 2^{AP}$. In fact, the reach-avoid specification that we defined in Def. 4.9 can be expressed as a PCTL formula $\text{unsafe} \cup^{\leq h} \text{goal}$ over two atomic propositions $AP = \{\text{goal}, \text{unsafe}\}$: one proposition goal corresponding to the set of target states X_T and another proposition unsafe for the set of unsafe states X_U . Then, the labeling function is defined such that:

$$\begin{aligned} \text{goal} \in L(x) &\iff x \in X_T, \text{ and} \\ \text{unsafe} \in L(x) &\iff x \in X_U. \end{aligned}$$

However, considering control problems given as general PCTL formulae quickly gets cumbersome while not providing much deeper theoretical insights. For the sake of readability, we, therefore, focus on reach-avoid specifications in this thesis, while referring to our paper [9] for more details on the extension of our methods to PCTL.

Summary

- ▷ Discrete-time stochastic systems (DTSSs) are continuous-state/action systems that evolve stochastically over discrete time steps.
- ▷ The nondeterministic choices in a DTSS are resolved by a Markov policy.
- ▷ The probability that a given reach-avoid specification is satisfied under a given policy can be characterized using a backward recursion on a value function over the state space.
- ▷ Computing this satisfaction probability exactly is intractable in general.

5 Probabilistic Simulation Relations

Summary | As discussed in Chapter 4, exactly computing satisfaction probabilities for discrete-time stochastic systems (DTSSs) is intractable in general. In this chapter, we present an abstraction-based framework to synthesize a Markov policy for a given DTSSs, together with a lower bound on the satisfaction probability under that policy. Our framework hinges on a formal relation between the DTSS and a finite Markov decision process (MDP) abstraction, which is called a probabilistic simulation relation. We show that our framework leads to Markov policies for DTSSs that provably satisfy reach-avoid specifications with at least the obtained lower bound probability.

Origins | The results in this chapter combine our contributions from several papers:

- [1] Badings, Abate, Jansen, Parker, Poonawala and Stoelinga (2022) ‘Sampling-Based Robust Control of Autonomous Systems with Non-Gaussian Noise’. AAAI.
- [6] Badings, Romao, Abate, and Jansen (2023) ‘Probabilities Are Not Enough: Formal Controller Synthesis for Stochastic Dynamical Models with Epistemic Uncertainty’. AAAI.
- [7] Badings, Romao, Abate, Parker, Poonawala, Stoelinga and Jansen (2023) ‘Robust Control for Dynamical Systems with Non-Gaussian Noise via Formal Abstractions’. J. Artif. Intell. Res.
- [9] Rickard, Badings, Romao and Abate (2023). ‘Formal Controller Synthesis for Markov Jump Linear Systems with Uncertain Dynamics’. QEST.
- [10] Badings, Romao, Abate and Jansen (2024) ‘Exploiting Stability for Abstractions of Stochastic Dynamical Systems’. ECC.

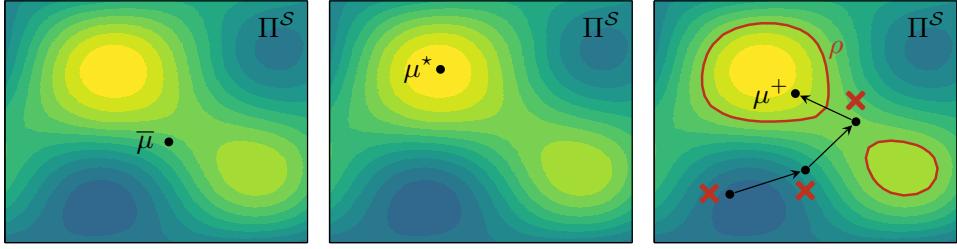
This chapter generalizes results from these papers into an overarching theoretical framework. Thus, this chapter cannot directly be traced to one of these papers.

Background | This chapter builds directly on the definitions for DTSSs from Chapter 4. Moreover, the reader should be familiar with MDPs and robust Markov decision processes (RMDPs), and with computing optimal (robust) schedulers for them (see Chapter 3).



5.1 Introduction

In Chapter 4, we have discussed reach-avoid control problems for discrete-time stochastic systems (DTSSs). Recall that a DTSS is a tuple $\mathcal{S} = (X, U, x_I, \varsigma, f)$, and that a Markov policy $\mu = (\mu_0, \mu_1, \mu_2, \dots)$ is a sequence of measurable maps $\mu_k : X \rightarrow U$, $k \in \mathbb{N}$, from states $x \in X$ to control inputs $u \in U$. Let $\Pi^{\mathcal{S}}$ denote the set of all Markov policies for DTSS \mathcal{S} . A reach-avoid control task can be formalized as a reach-avoid specification, which we defined in Def. 4.9 as a tuple $\varphi = (X_T, X_U, h)$ consisting of a set of target states X_T , a set of unsafe states X_U , and a time horizon $h \in \mathbb{N} \cup \{\infty\}$.



(a) **Policy evaluation:** given a fixed policy $\bar{\mu} \in \Pi^S$, compute $\Pr_{\bar{\mu}}^S(x_I \models \varphi)$. (b) **Optimal control:** compute a policy $\mu^* \in \Pi^S$ that maximizes $\Pr_{\mu^*}^S(x_I \models \varphi)$. (c) **LB control:** compute $\mu^+ \in \Pi^S$ such that $\Pr_{\mu^+}^S(x_I \models \varphi) \geq \rho$ or return False.

Figure 5.1: Three DTSS verification problems. The color gradients visualize level sets of the satisfaction probability. In this thesis, we consider problem (c), which we solve by iteratively searching for Markov policies, until one with a (lower bound) satisfaction probability above the threshold of ρ is found.

Certifying that a DTSS satisfies a desired reach-avoid specification is paramount, especially in safety-critical applications. In this section, we discuss the three verification problems for DTSSs depicted in Fig. 5.1, which we respectively call the policy evaluation, optimal control, and lower bound (LB) control problem.

5.1.1 Policy evaluation

In Chapter 4, we already considered the problem of computing the satisfaction probability for a given policy, which we call the *policy evaluation* problem.

Policy evaluation problem: For a fixed Markov policy $\bar{\mu}$ and a reach-avoid specification φ , compute the probability $\Pr_{\bar{\mu}}^S(x_I \models \varphi)$ of satisfying the specification.

As we have seen in Sect. 4.3.1, computing the satisfaction probability involves integrating the probability measure over (in)finite paths of the DTSS. For nontrivial Markov policies and distributions for the stochastic noise, computing these integrals exactly becomes infeasible in practice [BS78].

5.1.2 Optimal control

Another common problem is to compute a policy that *maximizes* the probability of satisfying the reach-avoid specification, resulting in an *optimal control* problem.

Optimal control problem: For a fixed reach-avoid specification φ , compute a Markov policy μ^* and its satisfaction probability $p^* = \Pr_{\mu^*}^S(x_I \models \varphi)$, such that

$$\mu^* \in \operatorname{argmax}_{\mu \in \Pi^S} \Pr_{\mu}^S(x_I \models \varphi).$$

One typically fixes a template (such as the set of linear Markov policies, i.e., policies with the structure $\mu_k(x) = Kx$ for some fixed matrix K) for the set of policies to

optimize over. The resulting stochastic optimal control problem can be formulated as a nonconvex stochastic optimization problem [BS78], where the objective is to maximize the probability of satisfying the reach-avoid specification, and where the constraints encode the (stochastic) dynamics and the value function for the reach-avoid probability (as described in Sect. 4.3.1). However, solving this optimization problem is intractable in general, due to the feedback structure of the Markov policy in combination with the nonconvexity and stochasticity of the dynamics.

5.1.3 Lower bound control

Because the previous two problems are intractable in general, we study a third type of verification problem in this thesis. Instead of asking for a policy that maximizes the satisfaction probability, we aim to find one for which the satisfaction probability meets a desired threshold. However, in general we cannot guarantee that we actually find such a policy. Thus, in this thesis, we consider the following problem.

Lower bound (LB) control problem: For a fixed reach-avoid specification φ and a threshold $\rho \in [0, 1]$, compute a Markov policy $\mu^+ \in \hat{\Pi}^S$ such that

$$\Pr_{\mu^+}^S(x_I \models \varphi) \geq \rho,$$

or return `False` if no such policy could be found.

The lower bound control problem asks for a *sound but not complete* solution method. If the method returns a Markov policy, then the satisfaction probability of this policy is at least ρ . However, if the method returns `False`, then we cannot conclude anything about the existence of a policy with satisfaction probability at least ρ .

An iterative solution method | Before jumping into details, let us sketch how we will solve the lower bound control problem in this thesis. As also visualized in Fig. 5.1c, the main idea is to iteratively search for Markov policies until we find a policy whose satisfaction probability meets the desired threshold of ρ . However, to determine whether the threshold is met, we still need to solve a policy evaluation problem in each iteration (which we have shown to be intractable).

Nevertheless, we shall see that, for policies with a particular structure, we can instead compute a *lower bound* on the satisfaction probability. Thus, let us for now assume that we have an *oracle* that, given a Markov policy with such a structure, returns a lower bound on the satisfaction probability. If this lower bound meets the threshold, then the actual satisfaction probability surely meets the threshold as well, so we have solved the lower bound control problem. If the lower bound does not meet the threshold, then we keep searching for a policy with a higher lower bound (until some termination criterion is reached). This general solution method is expressed by Algorithm 5.1, where the oracle is captured by the subroutine COMPUTEPOLICY.

Remark 5.1 (Nontrivial solutions) The lower bound control problem admits trivial and uninformative solutions, because always returning `False` is a valid solution. However, as we will see in the experiments of Chapters 6 and 7, the algorithms we develop to solve the problem generally lead to nontrivial solutions in practice.

Algorithm 5.1 High-level algorithm for solving the lower bound control problem.

Input: DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$; reach-avoid specification $\varphi = (X_T, X_U, h)$; satisfaction probability threshold $p \in [0, 1]$

Params: Termination criterion TERMINATE

Output: Markov policy μ^+ for \mathcal{S} , or False

```

1:  $p^0 = 0$ 
2: for  $i = 1, 2, \dots$  do
3:    $(\mu^i, p^i) \leftarrow \text{COMPUTEPOLICY}(\mathcal{S}, \varphi)$  such that  $p^i \geq \Pr_{\mu^i}^{\mathcal{S}}(x_I \models \varphi)$ 
4:   if  $p^i \geq p$  then
5:     return  $\mu^+ \leftarrow \mu^i$                                  $\triangleright$  Threshold probability met
6:   else if TERMINATE is True then
7:     return False                                          $\triangleright$  Inconclusive result

```

5.2 The DTSS Policy Synthesis Problem

In this chapter, we zoom in on the lower bound control problem and develop a theoretical basis for the subroutine COMPUTEPOLICY in Algorithm 5.1 which, given a DTSS and a reach-avoid specification, computes a policy and a sound lower bound on its satisfaction probability. Formally, we solve the following problem.

Problem 5.2 (DTSS policy synthesis) Given a DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ and a reach-avoid specification $\varphi = (X_T, X_U, h)$, compute a Markov policy μ and a lower bound p on the satisfaction probability, i.e., such that

$$\Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi) \geq p.$$

Throughout this chapter, we consider again Assumption 4.4, which states that the stochastic noise ς of the DTSS is independent and identically distributed (i.i.d.).

5.2.1 Approaches to DTSS policy synthesis

Synthesizing Markov policies for DTSSs that provably satisfy reach-avoid specifications (and, more generally, temporal logic specifications) is an active research area [KG02; BYG17]. Traditional methods from control theory largely focus on satisfying simple specifications and are thus insufficient for solving Problem 5.2. For example, the linear quadratic regulator (LQR) can be used to compute a linear feedback control law that minimizes a cost function that is quadratic in the state and control input. However, as we have seen in Sect. 4.3.1, the cost function for a reach-avoid specifications is not quadratic. Similarly, Lyapunov and barrier functions (and their stochastic variants) can be used to verify the (asymptotic) stability and safety of a system [DDNZ00; TSYA20], respectively. However, these traditional methods do not provide formal guarantees about temporal specifications [BK08; FQMN⁺22; STBR11; YTCB⁺12].

Formal policy synthesis | Consequently, various approaches have been developed over the past decades to compute Markov policies for reach-avoid problems. These approaches primarily focus on applications in safety-critical control engineering [APLS08;

[LSAZ22](#)] and can be divided into several categories. First, there are approaches that compute reachable sets directly in the continuous domain, e.g., using Hamilton-Jacobi reachability analysis [[BCHT17](#); [HCHB⁺17](#)] or optimization [[RSA22](#)]. Second, there are approaches that use variants of Lyapunov methods that are applicable to richer specifications, such as [[MCL23](#); [AGR24](#); [ZLHC23b](#)]. Finally, various approaches are based on a formal abstraction of the continuous system into a simpler (typically discrete) model [[AHLPO0](#); [LAB15](#); [SA13](#)].

Abstraction-based control | In this thesis, we focus on the latter abstraction-based paradigm, which is well-studied [[APLS08](#); [AHLPO0](#)] and has applications to, for example, stochastic hybrid [[CLLA⁺19](#); [LSAZ22](#); [ZSRH⁺12](#); [FHHW⁺11](#)], switched [[LAB15](#)], linear [[SKCC⁺15](#)], and partially observable systems [[HNWT⁺18](#)]. In the control literature, abstractions are also referred to as *symbolic models* [[APLS08](#); [LAB15](#); [Tab09](#)]. As we discuss in more detail later in this chapter, under an appropriate behavioral relation (e.g., a simulation relation [[Tab09](#)]) between both models, trajectories of the abstraction are related to those of the dynamical system. Thus, a scheduler for the abstraction can, by construction, be refined to a Markov policy for the original system. Various tools exist, such as StochHy [[CA19](#)], ProbReach [[SZ15](#)], and SReachTools [[VGO19](#)].

5.2.2 Shortcomings of abstraction-based control

Despite the active research efforts on abstraction-based policy synthesis, existing approaches still have significant shortcomings.

First of all, discretizing a continuous-state model inherently leads to finite-state models with exponentially many states in the number of state variables (also known as the “*curse of dimensionality*”, a term already coined by Richard Bellman in the sixties [[Bel66](#)]). The number of transitions¹ of the abstraction grows even faster: For every discrete state, several discrete actions are available, each of which has multiple possible successor states. Hence, scalability remains a general concern.

Second, existing abstraction methods rely on *full knowledge* of the underlying DTSS and are not robust against uncertainty in the dynamics. As a result, system parameters (such as the mass of the aforementioned unmanned aerial vehicle (UAV), or its friction coefficient) must be known precisely. Similarly, most methods require an explicit representation of the distribution of the stochastic noise, for example, as a Gaussian [[PSQ13](#)]. Hence, existing approaches are difficult to apply (or inapplicable at all) when there is uncertainty about the dynamics of the DTSS.

Jumping ahead, in Chapters 6 and 7, we zoom in on two specific settings, each of which considers DTSSs with a different form of uncertainty. As such, we address the shortcoming that existing abstraction methods require full knowledge of the dynamics. For each setting, we develop a tractable abstraction algorithm that aims to maximize practical scalability, thus addressing (to some extent) the shortcoming of limited scalability. In the remainder of this chapter, we lay the theoretical foundation that underpins these algorithms.

¹The number of transitions of a finite-state model is the number of edges in the underlying graph.

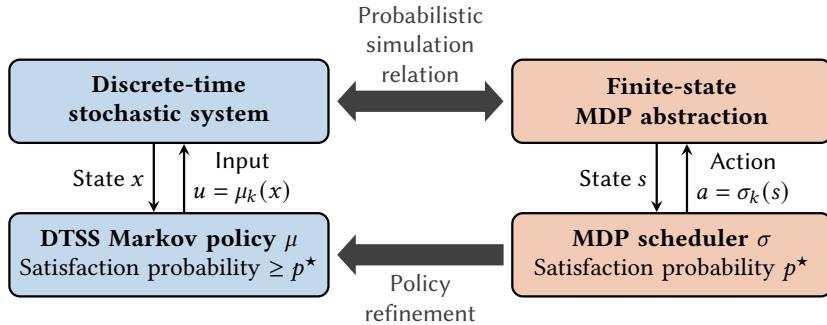


Figure 5.2: Our approach to solving Problem 5.2 is based on a probabilistic simulation relation between the DTSS and an MDP abstraction.

5.2.3 An overview of our approach

A high-level overview of our approach to solving Problem 5.2 is shown in Fig. 5.2. In a nutshell, we relate the DTSS (which has *infinitely* many states and actions) to an abstraction with *finitely* many states and actions, which we will formalize as a Markov decision process (MDP).² We require that the DTSS *simulates* this abstract MDP, which, loosely speaking, means that any possible probabilistic behavior of the abstract MDP can be mimicked by the DTSS (under some Markov policy). As a main contribution, we present sufficient requirements for the DTSS to simulate the abstract MDP in the form of a so-called *probabilistic simulation relation*.

Importantly, the existence of a probabilistic simulation relation between the DTSS and the MDP means that we can transfer satisfaction probabilities from the MDP to the DTSS. Specifically, any scheduler for the MDP with a corresponding satisfaction probability $p^* \in [0, 1]$ can, by construction, be translated (or *refined*) into a Markov policy for the DTSS with a satisfaction probability of *at least* p^* . This result allows us to reduce Problem 5.2 to computing an (optimal) scheduler for the MDP abstraction, which can be done using the value iteration algorithms described in Chapter 3.

Outline | In the remainder of this chapter, we introduce the behavioral relation that underpins our abstraction-based approach to solving Problem 5.2. In Sect. 5.3, we first present the particular type of probabilistic simulation relation that we use. To put our approach in perspective, in Sect. 5.3.3 we compare our particular relation to others classical relations from the verification and AI literature. We show in Sect. 5.4 how the existence of a probabilistic simulation relation between a DTSS and an MDP abstraction can be used to solve Problem 5.2. Finally, in Sect. 5.5 we extend these results to a relation between a DTSS and an RMDP.

5.3 Probabilistic Simulation Relations

At the core of many behavioral relations, including ours, is a (*binary*) *relation* between the states of two models. Such a relation is a set that contains all pairs of states that are, according to some criterion, related (or even equivalent).

²Later on, we will also use robust Markov decision process (RMDP) abstractions.

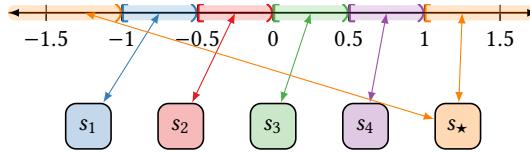


Figure 5.3: Binary relation R between the real number line \mathbb{R} and five discrete states $\{s_1, s_2, s_3, s_4, s_\star\}$, where states with the same color are related under R .

binary relation

5

Definition 5.3 (Binary relation) A set $R \subseteq X \times Y$ is called a *binary relation* between sets X and Y . We use the notation $R(x) := \{y \in Y : (x, y) \in R\}$ and $R^{-1}(y) := \{x \in X : (x, y) \in R\}$. Moreover:

- If $R(x) \neq \emptyset$ for all $x \in X$, then R is *strict* (in X);
- If $R(x) \neq \emptyset \implies |R(x)| = 1$ for all $x \in X$, then R is *single-valued* (in X) and we write $R(x) = y$ for $(x, y) \in R$ (i.e., we omit the fact that $R(x)$ is a set).

Thus, R is *strict single-valued* if $|R(x)| = 1$ for all $x \in X$.

partition

Intuitively, a strict single-valued binary relation $R \subseteq X \times Y$ relates each item $x \in X$ with exactly one element $y \in Y$. Thus, such a relation creates a *partition* of X into subsets that are related to the same element of Y , as illustrated by the next example. Equivalently, we can represent a strict single-valued binary relation $R \subseteq X \times Y$ as a function $f: X \rightarrow Y$, which is defined as $f(x) = y$ for all $(x, y) \in R$.

Example 5.4 (Partition as a relation) Consider a DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ whose state space is defined as $X = \mathbb{R}$, and consider an MDP $\mathcal{M} = (S, Act, s_I, P)$ whose state space is defined as $S = \{s_1, s_2, s_3, s_4, s_\star\}$. As shown in Fig. 5.3, we partition X into five elements defined as $X_1 = [-1, -0.5]$, $X_2 = [-0.5, 0]$, $X_3 = [0, 0.5]$, $X_4 = [0.5, 1]$, and $X_\star = (-\infty, -1) \cup [1, \infty)$. This partition can be captured in a strict binary relation R between X and S , which is defined as

$$\begin{aligned} R = & \{(x, s_1) : x \in [-1, -0.5]\} \cup \{(x, s_2) : x \in [-0.5, 0]\} \cup \\ & \{(x, s_3) : x \in [0, 0.5]\} \cup \{(x, s_4) : x \in [0.5, 1]\} \cup \\ & \{(x, s_\star) : x \in (-\infty, -1) \cup [1, \infty)\}\,. \end{aligned}$$

In this and the next chapters, we exclusively use binary relations defined between the states X of a DTSS and the states S of an MDP or RMDP.

5.3.1 Relating reach-avoid specifications

To tailor our relation to reach-avoid specifications, we first need to define how we relate these specifications between the DTSS and the MDP. We do so using the same binary relation R that we later use to relate the states of the DTSS and the MDP.

Definition 5.5 (Relation between specifications) Let $\mathcal{S} = (X, U, x_I, \varsigma, f)$ be a DTSS with a reach-avoid specification $\varphi = (X_T, U_T, h)$, and let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP with sets of target and unsafe states $S_T, S_U \subseteq S$. Consider a relation

consistent specifications

$R \subseteq X \times S$ for \mathcal{S} and \mathcal{M} . We call these specifications *consistent* under R , denoted by $(S_T, S_U) \leq_R (X_T, X_U)$, if

$$\mathbb{1}_{S_T}(s) \leq \mathbb{1}_{X_T}(x) \quad \text{and} \quad \mathbb{1}_{S_U}(s) \geq \mathbb{1}_{X_U}(x) \quad \forall (x, s) \in R.$$

equivalent specifications

Furthermore, we call these specifications *equivalent* under R , denoted by $(S_T, S_U) \equiv_R (X_T, X_U)$, if the above holds with equality, i.e.,

$$\mathbb{1}_{S_T}(s) = \mathbb{1}_{X_T}(x) \quad \text{and} \quad \mathbb{1}_{S_U}(s) = \mathbb{1}_{X_U}(x) \quad \forall (x, s) \in R.$$

Note that every pair of *equivalent* specifications is also *consistent*:

$$(S_T, S_U) \equiv_R (X_T, X_U) \implies (S_T, S_U) \leq_R (X_T, X_U)$$

A pair of reach-avoid specifications is *consistent* under R , if the MDP target states S_T *underapproximate* the DTSS target states X_T , and if the MDP unsafe states S_U *overapproximate* the DTSS unsafe states X_U . Similarly, *equivalence* means that the target and unsafe states of the DTSS and MDP correspond exactly under R . Thus, equivalence is a stronger condition than consistency: All pairs of equivalent specifications are also consistent, but not vice versa.

Example 5.6 We continue Example 5.4 and consider the sets $S_T = \{s_1\}$ and $S_U = \{s_4\}$ for the MDP. Consider the following cases for the DTSS:

- For the sets $X_T = [-1, -0.5]$ and $X_U = [0.5, 1]$, we have that $(S_T, S_U) \equiv_R (X_T, X_U)$, i.e., the specifications are equivalent under R ;
- For the sets $X_T = [-1.1, -0.4]$ and $X_U = [0.6, 0.9]$, we have that $(S_T, S_U) \leq_R (X_T, X_U)$, i.e., the specifications are consistent under R ;
- For the sets $X_T = [-1, -0.5]$ and $X_U = [0.4, 1.1]$, the specifications are not equivalent nor consistent. For example, for $(x, s) = (1.05, s_\star) \in R$, we have that $\mathbb{1}_{X_U}(x) = 1 > \mathbb{1}_{S_U}(s) = 0$.

5.3.2 Relating DTSSs and MDPs

We are now ready to define the particular variant of a probabilistic simulation relation we employ in this thesis.

Remark 5.7 (Generality of the relation) For clarity, we tailor the probabilistic simulation relation to the DTSS (reach-avoid) policy synthesis problem in Problem 5.2. As such, we specifically define a probabilistic simulation for the DTSS and an MDP abstraction, and we focus on simulation with respect to reach-avoid specifications. Nevertheless, the relation we propose can be generalized to other specifications and models. In particular, as we discussed in Sect. 4.3.2, we may consider more general probabilistic computation tree logic (PCTL) specifications. Yet, for the sake of readability, we tailor our definitions and results to the simpler yet more concrete reach-avoid policy synthesis problem in Problem 5.2.

Definition 5.8 (Relation between distributions) Let $\mathcal{S} = (X, U, x_I, \xi, f)$ be a DTSS with a stochastic kernel $T : \mathcal{B}(X) \times X \times U \rightarrow [0, 1]$, and let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP. Consider a strict binary relation $R \subseteq X \times S$. The relation R defines another binary relation $\bar{R} \subseteq \text{Distr}(X) \times \text{Distr}(S)$, called the *lifted relation*, defined as

$$\bar{R} = \left\{ (\nu, p) \in \text{Distr}(X) \times \text{Distr}(S) : \forall s' \in S. p(s') = \int_X \mathbb{1}_{R^{-1}(s')}(\xi) \cdot \nu(d\xi) \right\}.$$

The relation \bar{R} is called the *lifted* relation because it lifts R from the states of \mathcal{S} and \mathcal{M} to distributions over the states. In fact, Def. 5.8 can be seen as a concretization of the lifting of relations defined by [HSA17, Def. 5], tailored to relations between a DTSS and an MDP. The lifting of relations over continuous sets using measures (like we have for the DTSS) is analogous to the lifting of relations over countable or finite sets using *weight functions*, as done in [SL95; Sto02].

Definition 5.9 (Probabilistic simulation relation) Let $\mathcal{S} = (X, U, x_I, \xi, f)$ be a DTSS with a stochastic kernel $T : \mathcal{B}(X) \times X \times U \rightarrow [0, 1]$, and let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP. A strict binary relation $R \subseteq X \times S$ is a *probabilistic simulation relation* from MDP \mathcal{M} to DTSS \mathcal{S} if the following holds:

1. (*initial states are related*): For the initial states, we have $(x_I, s_I) \in R$;
2. (*next states are related*): For all $(x, s) \in R$, we have that

$$\forall a \in Act(s), \exists u \in U : (T(\cdot | x, u), P(s, a)) \in \bar{R},$$

where $\bar{R} \subseteq \text{Distr}(X) \times \text{Distr}(S)$ is the lifted relation for R as defined by Def. 5.8. We denote such a probabilistic simulation relation R by $\mathcal{M} \leq_R \mathcal{S}$.^a

^aIntuitively, all possible behaviors of the MDP are *contained* in the behaviors of the DTSS.

Informally, the second requirement in Def. 5.9 states that every pair of related MDP and DTSS states leads to related distributions in \bar{R} over successor states (under some Markov policy). More precisely, for every related pair $(x, s) \in R$ and for every MDP action $a \in Act(s)$, there exists a DTSS input $u \in U$ such that the probability $P(s, a)(s')$ that the MDP transitions to state $s' \in S$ is the same as the probability that the DTSS transitions to a state $x \in R^{-1}(s') \subseteq X$ related to s' . This interpretation of a probabilistic simulation relation is visualized by Fig. 5.4 and will enable us to relate satisfaction probabilities between the MDP and the DTSS.

Example 5.10 Consider again the DTSS \mathcal{S} , the MDP \mathcal{M} , and the relation R from Example 5.4. Suppose that, for a given state-input pair $(x, u) \in X \times U$, the DTSS kernel $T(\cdot | x, u)$ corresponds with a uniform distribution over $[-1.5, 1.5]$. We want to find the MDP distribution $p \in \text{Distr}(S)$ such that $(T(\cdot | x, u), p) \in \bar{R}$, i.e., the kernel and MDP distribution are related under R . From Fig. 5.3, we observe that this MDP distribution p is defined as

$$p(s_i) = \frac{1}{6}, \quad i = 1, \dots, 4, \quad \text{and} \quad p(s_\star) = \frac{1}{3}.$$

lifted relation

5

probabilistic simulation relation

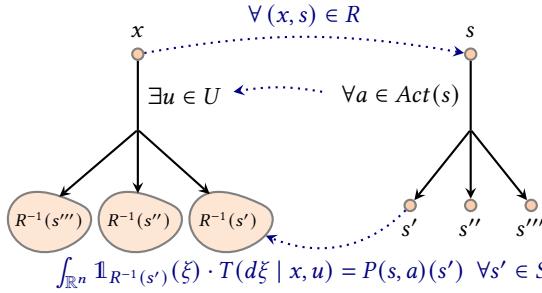


Figure 5.4: Visualization of a probabilistic simulation relation. For every pair $(x, s) \in R$ of related states and every $a \in \text{Act}(s)$, there must exist $u \in U$ such that we transition to a related successor state with equal probability.

5.3.3 Comparison to other behavioral relations

We briefly zoom out and make a comparison to other behavioral relations that have been proposed in the literature. Here, we focus only on behavioral relations for stochastic systems (and thus skip the classical notions of simulation [Mil71], alternating simulation [AHKV98], and forward-backward simulation [LV95], as well as feedback refinement relations [RWR17] for nonstochastic systems).

Probabilistic (bi)simulation | Probabilistic extensions of simulation and bisimulation³ date back to the early 90s [LS91; SL95]. Most of these early works originate from process calculi, which led to a slightly different perspective on probabilistic simulation than the one we use [BK08]. For example, for two states (s, s') to be bisimilar, [LS91] requires that, for all actions in state s , the resulting distribution over states is related to the distribution over states upon choosing *that same action* in state s' . By contrast, Def. 5.9 allows states to be related by comparing their behavior *under different actions*.

Probabilistic (bi)simulation has been studied for various models and from a variety of perspectives [BH97; DEP02; HPSW⁺11; HKK14]. Because exact probabilistic (bi)simulation can pose too strict conditions, various metrics and approximate notions have been developed, especially for MDPs [FPK14; GDG03; DGJP04].

As already mentioned, closest to our probabilistic simulation relation is [HSA17]. However, our approach differs from [HSA17] in the way we will deal with approximations of the relation in Def. 5.9. While [HSA17] considers approximations by allowing for a distance between the kernels of the two models (thus enlarging the lifted relation \bar{R}), we keep the exact relation \bar{R} and instead overapproximate the abstract MDP as an RMDP (see Sect. 5.5). In doing so, we establish a relation between the DTSS and an RMDP, which is closer to the recent works on probabilistic (bi)simulation for interval Markov decision processes (IMDPs) [HHSS⁺16; HHHT16].

MDP homomorphism | MDP homomorphisms were proposed by [RB01; RB03] as a relation between state-action pairs of two MDPs. One of the main motivations for MDP homomorphisms was to exploit symmetries in MDPs to improve the sample efficiency of, for example, deep reinforcement learning [PKOW20; PWHO⁺20]. Because MDP

³Two systems are bisimilar if one is simulated by the other and vice versa [BK08].

homomorphisms originate from the AI community, they are typically defined for MDPs with a state-action reward function $r: S \times Act \rightarrow \mathbb{R}_{\geq 0}$ and a discount factor $\gamma \in [0, 1]$. Formally, an MDP homomorphism is defined as follows.

Definition 5.11 (MDP homomorphism [RB03]) Consider two MDPs $\mathcal{M}_1 = (S_1, Act_1, s_{I1}, P_1, r_1, \gamma)$ and $\mathcal{M}_2 = (S_2, Act_2, s_{I2}, P_2, r_2, \gamma)$. An *MDP homomorphism* \mathcal{H} is defined by a tuple of surjective maps $(\beta, \{\alpha_s : s \in S\})$, where $\beta: S_1 \rightarrow S_2$ is the state map, $\alpha_s: Act_1(s) \rightarrow Act_2(\beta(s))$ is the action map for state s , and where the following conditions hold:

1. (*rewards match*): For all $s \in S_1, a \in Act_1(s)$, we have $r_2(\beta(s), \alpha_s(a)) = r_1(s, a)$;
2. (*transition functions match*): For all $s, s' \in S_1, a \in Act_1(s)$, we have $P_2(\beta(s), \alpha_s(a))(\beta(s')) = \sum_{s'' \in \beta^{-1}(s')} P_1(s, a)(s'')$.

An MDP homomorphism implies that the values of related state-action pairs are equal [LWL06; SLO22]. Much like bisimulation (and unlike simulation) relations, MDP homomorphisms are bidirectional: If \mathcal{H} is an MDP homomorphism from \mathcal{M}_1 to \mathcal{M}_2 , then we can invert the maps to obtain an MDP homomorphism from \mathcal{M}_2 to \mathcal{M}_1 .

Informally, the relation in Def. 5.9 can also be interpreted as a *unidirectional homomorphism*. To understand this connection, replace the reach-avoid specification with the equivalent reward function (as we discussed in Sect. 3.2.3). The homomorphism from Def. 5.11 requires *all* state-action pairs from the DTSS to be represented in the abstract MDP. By contrast, Def. 5.9 only requires that *some* state-action pairs from the DTSS are represented in the abstract MDP. Loosely speaking, we can thus interpret a probabilistic simulation relation as an homomorphism between the abstract MDP and a version of the DTSS where we removed part of the state-action pairs.

Remark 5.12 (Homomorphism vs. bisimulation) A key motivation for using MDP homomorphisms was that bisimulations used in process calculi were usually formulated as relations between states, and not between state-action pairs [RB01]. Bisimulations used in process calculi, e.g., [Mil89; LS91], indeed only consider the behavior between states for the same actions. Other definitions, such as e.g., [BK08], instead abstract away from actions completely and purely look at the possible transitions between two states. That is, two states $(x_1, x_2) \in R$ are bisimilar if for any successor of x_1 , there is a related successor of x_2 and vice versa (i.e., we are comparing behavior under different actions). Thus, we believe that probabilistic (bi)simulation and homomorphism for MDPs are more closely related than often argued in the literature. However, a more detailed discussion is beyond our scope.

5.4 Correct-by-Construction Markov Policy Synthesis

The following theorem, which is the main result of this section, defines a solution to Problem 5.2 based on a finite MDP abstraction of the DTSS. We use a probabilistic simulation relation R to relate the satisfaction probability of a scheduler for the MDP to that of a Markov policy for the DTSS. This theorem only guarantees *the existence of such a Markov policy*; we discuss how to compute this Markov policy in the next section.

Theorem 5.13 (Equivalence of satisfaction probabilities) Consider a DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP with $S_T, S_U \subseteq S$. Suppose there exists a strict single-valued binary relation $R \subseteq X \times S$ such that $\mathcal{M} \leq_R \mathcal{S}$ and $(S_T, S_U) \equiv_R (X_T, X_U)$. Then, for every Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ for MDP \mathcal{M} , there exists a Markov policy μ as in Def. 4.6 such that

$$\Pr_{\sigma}^{\mathcal{M}}(s_I \models \neg S_U \cup^{< h} S_T) = \Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi). \quad (5.1)$$

Proof. Recall from Sect. 4.3.1 that, as shown in [SL10, Lemma 4], the satisfaction probability for a given Markov policy μ can be computed recursively. For brevity, define $X_S = X \setminus (X_T \cup X_U)$ and define the backward recursion $V_k^{\mu}: X \rightarrow [0, 1]$ for $k = 0, \dots, h - 1$ as

$$V_k^{\mu}(x) = \mathbb{1}_{X_T}(x) + \mathbb{1}_{X_S}(x) \int_{\mathbb{R}^n} \mathbb{1}_X(\xi) \cdot V_{k+1}^{\mu}(\xi) \cdot T(d\xi | x, \mu_k(x)), \quad \forall x \in X, \quad (5.2)$$

initialized with $V_h^{\mu}(x) = \mathbb{1}_{X_T}(x)$ for all $x \in X$. Then, the satisfaction probability is $\Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi) = V_0^{\mu}(x_I)$. Similarly, for a scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$, the satisfaction probability for MDP \mathcal{M} is computed as $\Pr_{\sigma}^{\mathcal{M}}(s_I \models \neg S_U \cup^{< h} S_T) = W_0^{\sigma}(s_I)$, where the backward recursion $W_k^{\sigma}: S \rightarrow [0, 1]$, $k = 0, \dots, h - 1$ is defined as

$$W_k^{\sigma}(s) = \mathbb{1}_{S_T}(s) + \mathbb{1}_{S_S}(s) \sum_{s' \in S} W_{k+1}^{\sigma}(s') \cdot P(s, \sigma_k(s))(s'), \quad \forall s \in S, \quad (5.3)$$

initialized with $W_h^{\sigma}(s) = \mathbb{1}_{S_T}(s)$ for all $s \in S$, and where $S_S = S \setminus (S_T \cup S_U)$. We will show by induction that for all schedulers $\sigma \in \mathfrak{S}_{\text{Markov}}$, there exists a Markov policy μ such that $V_0^{\mu}(x) = W_0^{\sigma}(s)$ for all $(x, s) \in R$.

Base case | First, suppose $k = h$ and let $s = R(x)$ for any $x \in X$. Recall from Def. 5.5 that $(S_T, S_U) \equiv_R (X_T, X_U)$ implies that $\mathbb{1}_{X_T}(x) = \mathbb{1}_{S_T}(s)$ and $\mathbb{1}_{X_U}(x) = \mathbb{1}_{S_U}(s)$ for all $(x, s) \in R$. Thus, we have that

$$W_h^{\sigma}(s) = \mathbb{1}_{S_T}(s) = \mathbb{1}_{X_T}(x) = V_h^{\mu}(x), \quad (x, s) \in R.$$

This concludes the base case.

Inductive step | Next, suppose $k < h$. Because R is a strict single-valued relation, it partitions X into sets $R^{-1}(s)$, $s \in S$. Thus, we rewrite Eq. (5.2) as

$$V_k^{\mu}(x) = \mathbb{1}_{X_T}(x) + \mathbb{1}_{X_S}(x) \sum_{s' \in S} \left[\int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot V_{k+1}^{\mu}(\xi) \cdot T(d\xi | x, \mu_k(x)) \right].$$

For $s' \in S$ and $k = h - 1$, observe that for all $(x, s') \in R$, it holds that $V_{k+1}^{\mu}(x) = W_{k+1}^{\sigma}(s')$, i.e., the values $V_{k+1}^{\mu}(x)$ and $V_{k+1}^{\mu}(x')$ for DTSS states $R(x) = R(x')$ related to the same

MDP states are the same. Thus, we obtain

$$\begin{aligned} V_{h-1}^\mu(x) &= \mathbb{1}_{X_T}(x) + \mathbb{1}_{X_S}(x) \sum_{s' \in S} \left[W_h^\sigma(s') \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(x) \cdot T(d\xi | x, \mu_{h-1}(x)) \right] \\ &= \mathbb{1}_{X_T}(x) + \mathbb{1}_{X_S}(x) \sum_{s' \in S} \left[W_h^\sigma(s') \cdot P(R(x), \sigma_{h-1}(R(x)))(s') \right], \end{aligned}$$

where the second equality follows from the second requirement in Def. 5.9. Writing the difference between the value functions for $k = h - 1$:

$$\begin{aligned} V_{h-1}^\mu(x) - W_{h-1}^\sigma(R(x)) &= [\mathbb{1}_{X_T}(x) - \mathbb{1}_{S_T}(R(x))] + [\mathbb{1}_{X_S}(x) - \mathbb{1}_{S_S}(R(x))] \\ &\quad \times \sum_{s' \in S} \{W_h^\sigma(s') \cdot P(R(x), \sigma_{h-1}(R(x)))(s')\} \\ &= [\mathbb{1}_{X_T}(x) - \mathbb{1}_{R^{-1}(S_T)}(x)] + [\mathbb{1}_{X_S}(x) - \mathbb{1}_{R^{-1}(S_S)}(x)] \\ &\quad \times \sum_{s' \in S} \{W_h^\sigma(s') \cdot P(R(x), \sigma_{h-1}(R(x)))(s')\}. \end{aligned} \tag{5}$$

Again, due to $(S_T, S_U) \equiv_R (X_T, X_U)$, we have that $\mathbb{1}_{X_T}(x) = \mathbb{1}_{S_T}(s)$ and $\mathbb{1}_{X_U}(x) = \mathbb{1}_{S_U}(s)$ for all $(x, s) \in R$. Thus, it holds that $V_{h-1}^\mu(x) = W_{h-1}^\sigma(s)$ for all $(x, s) \in R$, which we use to obtain the general result for $k < h - 1$:

$$\begin{aligned} V_k^\mu(x) - W_k^\sigma(R(x)) &= [\mathbb{1}_{X_T}(x) - \mathbb{1}_{R^{-1}(S_T)}(x)] + [\mathbb{1}_{X_S}(x) - \mathbb{1}_{R^{-1}(S_S)}(x)] \\ &\quad \times \sum_{s' \in S} \{W_{k+1}^\sigma(s') \cdot P(R(x), \sigma_k(R(x)))(s')\} = 0. \end{aligned} \tag{5.4}$$

Using $k = 0$ in Eq. (5.4) gives $V_0^\mu(x) = W_0^\sigma(R(x))$, which concludes the proof. ■

Theorem 5.13 shows that the existence of a probabilistic simulation relation implies that, for any MDP scheduler, there exists a Markov policy for the DTSS under which the probabilities of satisfying the reach-avoid specifications in both models are equal. As a result, if we are able to construct an MDP abstraction of the DTSS, then we can compute an optimal policy on the MDP and translate back the results to the DTSS.

However, Theorem 5.13 requires that $(S_T, S_U) \equiv_R (X_T, X_U)$, i.e., the specifications are equivalent under R , which is a strong requirement in general (especially if X_T and S_T have a complex shape). The following result relaxes this requirement by only assuming that $(S_T, S_U) \leq_R (X_T, X_U)$, i.e., the specifications are consistent (but not necessarily equivalent) under R . In that case, the MDP satisfaction probability is still a *lower bound* on the DTSS satisfaction probability.

Corollary 5.14 (Lower bound on satisfaction probability) Consider again the assumptions from Theorem 5.13, but now consider that $(S_T, S_U) \leq_R (X_T, X_U)$ instead of $(S_T, S_U) \equiv_R (X_T, X_U)$. Then, for every Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ for MDP \mathcal{M} , there exists a Markov policy μ as in Def. 4.6 such that

$$\Pr_\sigma^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_\mu^S(x_I \models \varphi). \tag{5.5}$$

Proof. Def. 5.5 states that $(S_T, S_U) \leq_R (X_T, X_U)$ implies that

$$\mathbb{1}_{X_T}(x) \geq \mathbb{1}_{S_T}(s) \text{ and } \mathbb{1}_{X_U}(x) \leq \mathbb{1}_{S_U}(s) \quad \forall (x, s) \in R.$$

Using these expressions to modify the proof of Theorem 5.13 leads to the value functions satisfying $W_k^\sigma(R(x)) \leq V_k^\mu(x)$ for all $k = 0, \dots, h - 1$, which in turn leads to the claim in Eq. (5.5). ■

We have seen that the existence of a probabilistic simulation relation implies *the existence* of a Markov policy such that the DTSS satisfies a reach-avoid specification with at least the same probability as the MDP does. In this section, we show how to find this Markov policy based on a given MDP scheduler.

Intuitively, this Markov policy needs to pick inputs $u \in U$ that preserve the probabilistic simulation relation R from the MDP to the DTSS. Recall from Def. 5.9 that, to preserve the relation, we require for all $(x, s) \in R$ and all $s' \in S$ that

$$P(s, a)(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u).$$

Thus, if we are given a fixed Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ for the MDP, such that $a = \sigma_k(s)$, we obtain

$$P(s, \sigma_k(s))(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u). \quad (5.6)$$

Choosing any control input $u \in U$ that satisfies Eq. (5.6) is guaranteed to preserve the probabilistic simulation relation. This intuition is formalized in the notion of an *interface function*, which defines the set of all control inputs such that the DTSS \mathcal{S} has the same probabilistic behavior as the MDP \mathcal{M} (under scheduler σ) and thus the probabilistic simulation relation is preserved [GP09]. Since $h \in \mathbb{N}$ is the horizon of the reach-avoid specification, we only consider actions up to time $h - 1$.⁴

interface
function

Definition 5.15 (Interface function) An *interface function* for a probabilistic simulation relation $R \subseteq X \times S$ (between MDP \mathcal{M} and DTSS \mathcal{S}) and a Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ is a set-valued map $I_R^\sigma: X \times \{0, \dots, h - 1\} \rightarrow 2^U$, which is defined for all $x \in X$ and $k \in \{0, \dots, h - 1\}$ as

$$I_R^\sigma(x, k) = \left\{ u \in U : \forall s' \in S, P(s, \sigma_k(s))(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u) \right\},$$

where $s \in S$ is such that $(x, s) \in R$, and T is the stochastic kernel of DTSS \mathcal{S} .

Thus, the interface function describes precisely the Markov policies for the DTSS for which Theorem 5.13 holds. Embedding the interface function in a control loop leads to the scheme shown in Fig. 5.5. As shown in this figure, the interface function uses the scheduler as a *look-up table*, which, based on the current state x_k , returns a control input u_k that preserves the probabilistic simulation relation.

⁴If $h = \infty$, then a stationary MDP scheduler suffices to attain optimal reach-avoid probabilities. In that case, the interface function is independent of time and is defined as a function $I_R^\sigma: X \rightarrow 2^U$.

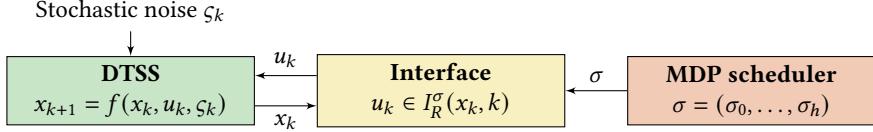


Figure 5.5: The interface function I_R^σ translates an MDP scheduler σ into a Markov policy for the DTSS for the probabilistic simulation relation R is preserved.

Lemma 5.16 (Nonemptyiness of the interface) The interface function $I_R^\sigma(x, k)$ is nonempty for all $x \in X$ and $k \in \{0, \dots, h - 1\}$.

Proof. A probabilistic simulation relation R is strict by definition. Thus, for all $x \in X$, there exists an $s \in S$ such that $(x, s) \in R$. By Def. 5.9, for all $(x, s) \in R$ and all $a \in Act(s)$ (which includes $\sigma_k(s) \in Act(s)$), there exists a $u \in U$ such that

$$P(s, a)(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u), \quad \forall s' \in S, \quad (5.7)$$

which implies that $I_R^\sigma(x, k) \neq \emptyset$ for all $x \in X$ and $k \in \{0, \dots, h - 1\}$. ■

The following theorem shows that, by restricting the Markov policy μ to the control inputs in an interface function I_R^σ , the satisfaction probability $\Pr_\mu^S(x_I \models \varphi)$ on the DTSS is at least as high as the satisfaction probability on the MDP under scheduler σ .

Theorem 5.17 (Markov policy synthesis) Let $\mathcal{S} = (X, U, x_I, \zeta, f)$ be a DTSS with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP with $S_T, S_U \subseteq S$. Suppose there exists a strict single-valued binary relation $R \subseteq X \times S$ such that $\mathcal{M} \leq_R \mathcal{S}$ and $(S_T, S_U) \leq_R (X_T, X_U)$. Let $\mu = (\mu_0, \mu_1, \mu_2, \dots)$ be the Markov policy defined for all $x \in X$ and $k \in \{0, \dots, h - 1\}$ as

$$\mu_k(x) \in I_R^\sigma(x, k),$$

where I_R^σ is an interface function for R and $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$. Then, it holds that

$$\Pr_\sigma^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_\mu^S(x_I \models \varphi). \quad (5.8)$$

Furthermore, if $(S_T, S_U) \equiv_R (X_T, X_U)$, then the above holds with equality, i.e.,

$$\Pr_\sigma^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T) = \Pr_\mu^S(x_I \models \varphi). \quad (5.9)$$

Proof. We prove the theorem by showing that Theorem 5.13 holds for any Markov policy that satisfies $\mu_k(x) \in I_R^\sigma(x, k)$. That is, we need to show that

$$\mu_k(x) \in I_R^\sigma(x, k) \implies \forall s' \in S, \quad P(s, a)(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, \mu_k(x)),$$

which is satisfied by construction, because $I_R^\sigma(x, k) \subseteq U$ contains precisely the control inputs $u \in U$ such that

$$\forall s' \in S, \quad P(s, a)(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid x, u).$$

The remainder of the proof is then equivalent to the proof of Theorem 5.13. ■

Finally, we have all the ingredients to claim that Theorem 5.17 can be used to compute Markov policies that solve Problem 5.2 in three general steps:

1. For a given DTSS, find an MDP abstraction \mathcal{M} which yields a probabilistic simulation relation R as per Def. 5.9, i.e., $\mathcal{M} \leq_R S$;
2. Pick any Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ for the MDP and define the corresponding interface function I_R^σ as per Def. 5.15;
3. Define a Markov policy that chooses inputs contained in the interface, i.e., $\mu_k(x) \in I_R^\sigma(x, k)$ for all x and k . Due to Theorem 5.17, the satisfaction probability under this Markov policy is at least the satisfaction probability on the MDP.

In practice, we are interested in maximizing the satisfaction probability on the DTSS. Thus, instead of picking *any* scheduler for the MDP in step two, we can compute an *optimal* scheduler that maximizes the satisfaction probability on the MDP. While we cannot provide any guarantees on the value of the satisfaction probability *a-priori*, we can still use this general approach to compute a Markov policy for the DTSS with a guaranteed lower bound on the satisfaction probability.

Remark 5.18 (Computing interface functions) One may wonder how to compute the interface function I_R^σ from Def. 5.15 in practice, especially due to the integral over the state space of the DTSS. While this integral may indeed be intractable to compute, it turns out that by carefully constructing the MDP abstraction, we can often compute the interface function without needing to compute any integrals. For example, in the next chapter, we will consider DTSS with additive stochastic noise, and we show that interface functions can be computed efficiently and exactly without computing any integrals.

5.5 DTSS Relations With Robust MDPs

Def. 5.9 of a probabilistic simulation relation requires generating an MDP abstraction with the *exact* same probabilistic behavior as the DTSS. However, ensuring the exact equivalent probabilistic behavior is often impossible, e.g., when transition probabilities must be estimated from data (as we shall see in, Chapter 6). In such cases, it is instead often possible to estimate the transition probabilities up to a given set and formalize the resulting abstraction as an RMDP (or, in the specific case that these sets are intervals, as an IMDP). We extend Theorem 5.13 to RMDP abstractions to facilitate this setting.

5.5.1 Probabilistic alternating simulation relation

The first step is to define a more general notion of probabilistic simulation. Inspired by the notion of *alternating simulation* for nonstochastic systems [AHKV98; Tab09], we

define the following relation between a DTSS and an RMDP.

Definition 5.19 (Probabilistic alternating simulation relation) Let $\mathcal{S} = (X, U, x_I, \varsigma, f)$ be a DTSS with a stochastic kernel $T: \mathcal{B}(X) \times X \times U \rightarrow [0, 1]$, and let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P})$ be an RMDP. A strict binary relation $R \subseteq X \times S$ is a *probabilistic alternating simulation relation* from RMDP \mathcal{M} to DTSS \mathcal{S} if:

1. (*initial states are related*): For the initial states, we have $(x_I, s_I) \in R$;
2. (*next states are related*): For all $(x, s) \in R$, we have that

$$\forall a \in Act(s), \exists u \in U, \exists P(s, a) \in \mathcal{P}(s, a) : (T(\cdot | x, u), P(s, a)) \in \bar{R},$$

where $\bar{R} \subseteq \text{Distr}(X) \times \text{Distr}(S)$ is the lifted relation for R as defined by Def. 5.8. We denote such a probabilistic simulation relation R by $\mathcal{M}_R \leq_R^{\text{alt}} \mathcal{S}$.

probabilistic
alternating
simulation
relation

Note that a probabilistic alternating simulation relation is almost identical to the non-alternating version from Def. 5.9, with the exception of an additional existential quantifier over the transition function in the second condition. From a game-based perspective, this additional quantifier means that the DTSS can, besides the input $u \in U$, now also “choose” the transition function $P(s, a) \in \mathcal{P}(s, a)$.

5

Remark 5.20 (Comparison to classical alternating simulation) The classical notion of alternating simulation reasons over two layers of nondeterminism in *both* models. That is, alternating simulation is used to relate two models with both action choices and nondeterminism in the outcomes of actions. In our setting, the RMDP has two layers of nondeterminism (the choice of action and the choice of transition function), whereas DTSS only has a single layer of nondeterminism (after choosing a control input, the stochastic kernel is deterministic). Nevertheless, we use the term *alternating* to emphasize the similarity to the classical notion from [AHKV98].

Recall that choosing $P(s, a) \in \mathcal{P}(s, a)$ for all state $s \in S$ and actions $a \in Act(s)$ is equivalent to reducing the RMDP to an MDP. As such, the existence of a probabilistic alternating simulation relation from RMDP \mathcal{M}_R to DTSS \mathcal{S} implies that there exists a nature $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$ such that there is a probabilistic (non-alternating) simulation relation from the induced MDP \mathcal{M} to DTSS \mathcal{S} .⁵ This argument directly leads to the following result, which we state without further proof.

Lemma 5.21 Let $\mathcal{S} = (X, U, x_I, \varsigma, f)$ be a DTSS and let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P})$ be an RMDP. If $\mathcal{M}_R \leq_R^{\text{alt}} \mathcal{S}$, then there exists a nature $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$ such that $\mathcal{M} \leq_R \mathcal{S}$, where $\mathcal{M} = (S, Act, s_I, P)$ is the MDP induced by \mathcal{M}_R and τ .

Conversely, suppose that we have a probabilistic simulation relation R between an MDP \mathcal{M} and a DTSS \mathcal{S} . Moreover, fix any RMDP \mathcal{M}_R that can induce MDP \mathcal{M} (under some nature). Then, the same relation R is also a probabilistic alternating simulation relation between \mathcal{M}_R and \mathcal{S} . We formalize this result in the following lemma.

⁵Recall from Sect. 3.3 that the MDP induced by applying a nature τ to an RMDP \mathcal{M}_R has the same states S , actions Act , and initial states s_I as \mathcal{M}_R , but the transition function is defined as $P(s, a) = \tau(s, a) \in \mathcal{P}(s, a)$. For a Markov nature, we additionally allow for a different $P(s, a)$ in every step of the execution; see the RMDP semantics in Sect. 3.3 for details.

Lemma 5.22 Let $\mathcal{S} = (X, U, x_I, \zeta, f)$ be a DTSS, let $\mathcal{M} = (S, Act, s_I, P)$ be an MDP, and let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P})$ be an RMDP with the same states, actions, and initial state as \mathcal{M} , and such that the uncertain transition function \mathcal{P} satisfies

$$P(s, a) \in \mathcal{P}(s, a) \quad \forall s \in S, \forall a \in Act(s). \quad (5.10)$$

Furthermore, let $R \subseteq X \times S$. Then, if $\mathcal{M} \leq_R \mathcal{S}$, it holds that $\mathcal{M}_R \leq_R^{\text{alt}} \mathcal{S}$.

Proof. To prove the lemma, we show that satisfying the conditions of a probabilistic simulation relation (Def. 5.9) implies the satisfaction of the conditions of a probabilistic alternating simulation relation (Def. 5.19). First, note that the states, actions, and initial state of \mathcal{M}_R and \mathcal{M} are the same. Thus, the first condition (i.e., initial states are related) is trivially satisfied.

The second condition for a probabilistic simulation relation in Def. 5.9 states that for all $(x, s) \in R$, it holds that

$$\forall a \in Act(s), \exists u \in U : (T(\cdot | x, u), P(s, a)) \in \bar{R}.$$

where \bar{R} is the lifted relation for R as per Def. 5.8. Now, we use the fact from Eq. (5.10) that $P(s, a) \in \mathcal{P}(s, a)$. Thus, we obtain that for all $(x, s) \in R$, it holds that

$$\exists P(s, a) \in \mathcal{P}(s, a), \forall a \in Act(s), \exists u \in U : (T(\cdot | x, u), P(s, a)) \in \bar{R}.$$

By pulling the outer existential quantifier inside the universal quantifier, we obtain

$$\forall a \in Act(s), \exists u \in U, \exists P(s, a) \in \mathcal{P}(s, a) : (T(\cdot | x, u), P(s, a)) \in \bar{R},$$

which is precisely the second condition for a probabilistic alternating simulation relation in Def. 5.19. Thus, we conclude the proof. ■

We will use Lemma 5.22 in Chapters 6 and 7 to solve the DTSS policy synthesis problem using IMDP abstractions, rather than glsMDP abstractions.⁶ In particular, we shall see that computing an MDP abstraction (with precise transition probabilities) is often infeasible. As a solution, we instead compute an interval around each of these transition probabilities, resulting in an IMDP abstraction for which we can apply the results presented in section.

5.5.2 Lower bounding satisfaction probabilities

We have seen that $\mathcal{M} \leq_R \mathcal{S}$ implies we can lower bound the satisfaction probability on \mathcal{S} by analyzing \mathcal{M} . At the same time, recall from the RMDP sandwich lemma (Lemma 3.32) that the robust (i.e., pessimistic) satisfaction probability on an RMDP is defined as the minimum over all natures. In other words, the satisfaction probability on RMDP \mathcal{M}_R under the most pessimistic nature is a lower bound on the satisfaction probability on any MDP \mathcal{M} it can induce, which thus immediately leads to a lower bound on the satisfaction probability on \mathcal{S} . Formalizing this intuition leads to the following result.

⁶Recall from Def. 3.27 that an IMDP is a particular type of RMDP and hence, all results for RMDPs naturally carry over to IMDPs.

Theorem 5.23 (Robust lower bound on satisfaction probability) Consider a DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P})$ be an RMDP with $S_T, S_U \subseteq S$. Suppose there exists a strict single-valued binary relation $R \subseteq X \times S$ such that $\mathcal{M}_R \leq_R^{\text{alt}} \mathcal{S}$ and $(S_T, S_U) \leq_R (X_T, X_U)$. Then, for every Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_R}$ for RMDP \mathcal{M}_R , there exists a Markov policy μ as in Def. 4.6 such that

$$\min_{\tau \in \mathfrak{T}_{\text{Markov}}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi). \quad (5.11)$$

Proof. Fix any nature $\bar{\tau} \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$ for the RMDP. From Lemma 3.32 (the RMDP sandwich lemma), it follows that for all RMDP schedulers $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_R} \subseteq \mathfrak{S}^{\mathcal{M}_R}$, it holds that

$$\begin{aligned} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s_I \models \neg S_U \cup^{\leq h} S_T) &\leq \Pr_{\sigma, \bar{\tau}}^{\mathcal{M}_R}(s_I \models \neg S_U \cup^{\leq h} S_T) \\ &= \Pr_{\sigma}^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T), \end{aligned} \quad (5.12)$$

where \mathcal{M} is the MDP induced by applying nature $\bar{\tau}$ to RMDP \mathcal{M}_R . From Lemma 5.21, we know that there exists a nature $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_R}$ such that the induced MDP satisfies $\mathcal{M} \leq_R \mathcal{S}$. Thus, we can choose $\bar{\tau}$ to be the nature such that $\mathcal{M} \leq_R \mathcal{S}$ holds. From Corollary 5.14 we know that $\mathcal{M} \leq_R \mathcal{S}$ and $(S_T, S_U) \leq_R (X_T, X_U)$ implies that there exists a Markov policy μ such that

$$\Pr_{\sigma}^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi). \quad (5.13)$$

By combining Eq. (5.12) with Eq. (5.13), we obtain

$$\begin{aligned} \min_{\tau \in \mathfrak{T}_{\text{Markov}}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s_I \models \neg S_U \cup^{\leq h} S_T) &\leq \Pr_{\sigma}^{\mathcal{M}}(s_I \models \neg S_U \cup^{\leq h} S_T) \\ &\leq \Pr_{\mu}^{\mathcal{S}}(x_I \models \varphi), \end{aligned}$$

which equals Eq. (5.11), so we conclude the proof. ■

While the proof above is straightforward, Theorem 5.23 will prove to be a powerful tool in later chapters to construct RMDP (or in fact IMDP) abstractions of DTSSs.

5.5.3 Markov policy synthesis with RMDPs

All that is left to solve the DTSS synthesis problem in Problem 5.2 is to derive the Markov policy for which the probabilistic alternating simulation relation is preserved. First, we extend the interface function from Def. 5.15 to match the conditions of the probabilistic alternating simulation relation.

Definition 5.24 (Robust interface function) A *robust interface function* for a probabilistic alternating simulation relation $R \subseteq X \times S$ (between RMDP \mathcal{M}_R and DTSS \mathcal{S}) and a Markov scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ is a set-valued map $\tilde{I}_R^\sigma: X \times$

robust
interface
function

$\{0, \dots, h-1\} \rightarrow 2^U$, which is defined for all $x \in X$ and $k \in \{0, \dots, h-1\}$ as

$$\tilde{I}_R^\sigma(x, k) = \left\{ u \in U : \forall s' \in S, \mathcal{P}(s, \sigma_k(s))(s') \ni \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u) \right\},$$

where $s \in S$ is such that $(x, s) \in R$, and T is the stochastic kernel of DTSS \mathcal{S} .

Finally, we show that, by defining the robust interface function for the RMDP, we can compute a Markov policy for the DTSS for which the satisfaction probability is at least the (robust) satisfaction probability of the RMDP.

Corollary 5.25 (Robust Markov policy synthesis) Let $\mathcal{S} = (X, U, x_I, \zeta, f)$ be a DTSS with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\mathcal{M}_R = (S, Act, s_I, \mathcal{P})$ be an RMDP with $S_T, S_U \subseteq S$. Suppose there exists a strict single-valued binary relation $R \subseteq X \times S$ such that $\mathcal{M}_R \leq_R^{\text{alt}} \mathcal{S}$ and $(S_T, S_U) \leq_R (X_T, X_U)$. Let $\mu = (\mu_0, \mu_1, \mu_2, \dots)$ be the Markov policy defined for all $x \in X$ and $k \in \{0, \dots, h-1\}$ as

$$\mu_k(x) \in \tilde{I}_R^\sigma(x, k),$$

where \tilde{I}_R^σ is a robust interface function for R and $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_R}$. Then, it holds that

$$\min_{\tau \in \mathfrak{T}_{\text{Markov}}} \Pr_{\sigma, \tau}^{\mathcal{M}_R}(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_\mu^{\mathcal{S}}(x_I \models \varphi). \quad (5.14)$$

Proof. We prove the theorem by showing that Theorem 5.23 holds for any policy that satisfies $\mu_k(x) \in \tilde{I}_R^\sigma(x, k)$. Specifically, Theorem 5.23 holds for any policy that preserves the probabilistic alternating simulation relation. Thus, we need to show that

$$\mu_k(x) \in \tilde{I}_R^\sigma(x, k) \implies \forall s' \in S, \mathcal{P}(s, \sigma_k(s))(s') \ni \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi | x, u),$$

which is satisfied by construction of the robust interface function in Def. 5.24. ■

We will use Theorem 5.23 and Corollary 5.25 to derive several main results in the next two chapters. In particular, by defining an RMDP such that Theorem 5.23 holds, we can use Corollary 5.25 to define an appropriate robust interface function for any scheduler for the RMDP, such that the lower bound on the satisfaction probability carries over.

Summary

- ⇒ The lower bound control problem for a DTSS asks for a Markov policy such that the reach-avoid specification is satisfied with at least a certain probability (or return that no such policy could be found).
- ⇒ We have developed a framework to compute a Markov policy together with a lower bound on the probability of satisfying the specification.
- ⇒ Probabilistic simulation relations allow us to compute such DTSS Markov policies based on a finite MDP abstraction.
- ⇒ Similarly, by defining the notion of a probabilistic alternating simulation relation, we can weaken the requirements on this finite-state abstraction and instead consider RMDPs abstractions.

6 Reach-Avoid Control of Linear DTSSs

Summary | In this chapter, we zoom in on a class of discrete-time stochastic systems (DTSSs) where the transition function is linear in the state, control input, and stochastic noise. For this class of linear DTSSs, we present a tractable Markov decision process (MDP) abstraction method that induces a probabilistic simulation relation. However, computing the transition probabilities of this MDP exactly is not possible in general. Moreover, the common assumptions that the noise distributions of DTSSs are known and/or Gaussian are unrealistic in practice. Thus, we drop these assumptions and instead use sampling-based techniques from the scenario approach to compute probably approximately correct (PAC) bounds on the transition probabilities of the abstract MDP. We use these bounds to formalize the abstraction as an interval Markov decision process (IMDP), for which the results from Chapter 5 for robust Markov decision processes (RMDPs) naturally carry over. Based on our abstraction method, we present an algorithm to solve the lower bound control problem introduced in Chapter 5. Furthermore, we investigate how the stability of a DTSS may be exploited to reduce the size of abstract models while retaining the correctness guarantees.

Origins | This chapter is based on the following publications:

- [1] Badings, Abate, Jansen, Parker, Poonawala and Stoelinga (2022) ‘Sampling-Based Robust Control of Autonomous Systems with Non-Gaussian Noise’. AAAI.
- [7] Badings, Romao, Abate, Parker, Poonawala, Stoelinga and Jansen (2023) ‘Robust Control for Dynamical Systems with Non-Gaussian Noise via Formal Abstractions’. J. Artif. Intell. Res.
- [10] Badings, Romao, Abate, and Parker (2024) ‘Exploiting Stability for Abstractions of Stochastic Dynamical Systems’. ECC.

Specifically, most content originates from [1] and the later journal publication [7]. The idea of exploiting stability to generate smaller abstractions originates from [10].

Background | We assume the reader is familiar with MDPs (Def. 3.1) and IMDPs (Def. 3.27), and with computing optimal schedulers for reach-avoid probabilities. To capture stochastic uncertainty in models, we build upon the probability-theoretic definitions from Sect. 2.3. We also use the results from Chapter 5 on probabilistic (alternating) simulations to prove the correctness of (I)MDP abstractions.



6.1 Linear DTSS

In Chapter 5, we have established that the *existence* of a probabilistic simulation relation between a discrete-time stochastic system (DTSS) and a Markov decision process (MDP) leads to a Markov policy together with a lower bound on the probability of satisfying a

reach-avoid specification (and thus solving Problem 5.2). Furthermore, we have discussed how to use this result in a *sound but not complete* method for computing a Markov policy that satisfies a reach-avoid specification with a desired threshold probability. We called this problem the *lower bound control problem*.

However, the dynamics of a DTSS as defined by Def. 4.2 may, in general, be highly nonlinear or even discontinuous, in which case it is intractable to actually find an MDP that induces a probabilistic simulation relation with the DTSS. In this chapter, we thus zoom in on a particular class of DTSSs for which solving the lower bound control problem is tractable. Specifically, we consider DTSSs with a transition function that is linear in the state, control input, and stochastic noise. Such DTSSs are commonly called *linear*. Recall from Def. 4.2 that x_k , u_k , and ζ_k are the state, control input, and stochastic noise at time step k , respectively. Then, we define a linear DTSS as follows.

linear DTSS

Definition 6.1 (Linear DTSS) A DTSS $\mathcal{S} = (X, U, x_I, \zeta, f)$ is *linear* if the state transition function f is linear, i.e., can be written in the form

$$f(x, u, \zeta) = Ax + Bu + q + \zeta, \quad (6.1)$$

where $A \in \mathbb{R}^{n \times n}$, $B \in \mathbb{R}^{n \times m}$, and $q \in \mathbb{R}^n$ are of appropriate size.

The matrix A is called the *system matrix* and models how the state x_{k+1} at step $k+1$ depends (linearly) on the state x_k . Similarly, the matrix B is called the *input matrix* and models how x_{k+1} depends (linearly) on the control input u_k .

To ensure that the set of reachable states is closed under the state transition function f , we take $X := \mathbb{R}^n$ as the state space in this chapter and write \mathbb{R}^n instead of X where convenient. Recall from Sect. 4.2.2 that T denotes the stochastic kernel associated with the DTSS, which in the linear case is defined for each $V \in \mathcal{B}(\mathbb{R}^n)$ as¹.

$$T(V | x, u) = \mathbb{P}\{\omega \in \Omega : Ax + Bu + q + \zeta_k(\omega) \in V\}. \quad (6.2)$$

6.1.1 Assumptions

In addition to Assumption 4.4 (the stochastic noise is i.i.d.) and Assumption 4.10 (the target and unsafe sets of the reach-avoid specification are disjoint), we make the next assumptions on the noise ζ .

Assumption 6.2 (Noise has density) The Radon-Nikodym derivative of the probability measure \mathbb{P} of the process noise $\zeta = (\zeta_k)_{k \in \mathbb{N}}$ with respect to the Lebesgue measure exists.

Assuming the existence of the Radon-Nikodym derivative is quite standard in probability theory and means that the process noise ζ has density [Dur10]. Intuitively, Assumption 6.2 implies that the probability $T(V | x_k, u_k)$ for x_{k+1} to be contained in the Borel set $V \in \mathcal{B}(X)$ can be nonzero only if V has nonzero volume. Consequently, the probability for x_{k+1} to lie on any subspace of X is zero (as any subspace has volume zero). In practice, this assumption rules out anomalous cases in which there is a nonzero probability for

¹Recall from Assumption 4.10 that the stochastic noise is assumed independent and identically distributed (i.i.d.), so the kernel T is time-invariant and thus equal for all time steps k

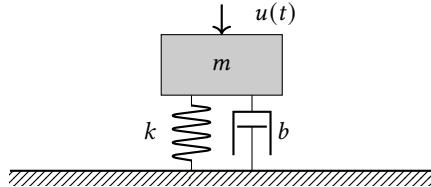


Figure 6.1: Mass-spring-damper circuit with mass m , spring stiffness k , damping coefficient b , and control input $u(t)$ perturbed by a disturbance $v(t)$.

drawing two values $\omega, \omega' \in \Omega$ according to \mathbb{P} for which $\varsigma_k(\omega) = \varsigma_k(\omega')$, i.e., the noise has exactly the same value (we will need this assumption in Sect. 6.3).

Assumption 6.3 (Polytopic input space) The control input space U is a convex polytope $U = \{u \in \mathbb{R}^m : Gu \leq g\} \subset \mathbb{R}^m$, where $G \in \mathbb{R}^{q \times m}$ and $g \in \mathbb{R}^q$.

Finally, we also make the following assumption on the dynamics of the linear DTSS.

Assumption 6.4 (Non-singular and controllable) The matrix $A \in \mathbb{R}^{n \times n}$ is non-singular. Furthermore, the matrix pair (A, B) is controllable, i.e., the matrix

$$C = [B \ AB \ \cdots \ A^{n-1}B]$$

has full row rank, i.e., $\text{rank}(C) = n$.

The controllability assumption requires that each column of the matrix C is linearly independent, which is needed for our backward reachability analysis. Controllability is a standard assumption in control theory, which roughly ensures that the system's state can be steered to any desired state in a finite number of steps.

Matrix A being non-singular implies that its inverse A^{-1} exists, which we need in our abstraction procedure (to perform backward reachability analysis on the DTSS; see Sect. 6.2). This non-singularity assumption is less standard and can be seen as “*the price we pay*” to be able to use backward reachability analysis to construct abstractions.

The mass-spring-damper model | Many physical systems can be modeled as a linear DTSS, where the transition function can often be derived from differential equations, as illustrated by the following example.

Example 6.5 Fig. 6.1 shows a classical mass-spring-damper model consisting of a point-mass connected to a spring and a damper. The spring force is proportional (with stiffness $k > 0$) to the displacement $p \in \mathbb{R}$ of the mass, and the viscous damping force is proportional (with coefficient $b > 0$) to the velocity $v = \dot{p} \in \mathbb{R}$ of the mass. At time $t \in \mathbb{R}$, an actuator exerts a force of $u(t) \in U \subset \mathbb{R}$ on the mass. The evolution of the state variables $x = [p, v]^\top \in \mathbb{R}^2$ for the mass-spring-damper is described by the following system of first-order stochastic differential equations:

$$\dot{x}(t) = \begin{bmatrix} \dot{p} \\ \dot{v} \end{bmatrix} = \begin{bmatrix} \dot{p} \\ \dot{v} \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{b}{m} \end{bmatrix} x(t) + \begin{bmatrix} 0 \\ \frac{1}{m} \end{bmatrix} u(t),$$

A simple discretization of this continuous-time model is given by the (forward) Euler method. Moreover, we want to account for the fact that the desired control input cannot be realized exactly (due to actuator noise). We model this noise as an additive stochastic disturbance to the input, such that the actual force exerted at time step k is $u_k + \zeta_k$, where we assume that ζ_k is i.i.d. and normally distributed (thus satisfying Assumptions 4.4 and 6.2). As such, for a step size of $\delta_t \in \mathbb{R}_{>0}$, we obtain a DTSS with the linear transition function

$$\begin{aligned} x_{k+1} &= x_k + \delta_t \left(\begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{b}{m} \end{bmatrix} x_k + \begin{bmatrix} 0 \\ \frac{1}{m} \end{bmatrix} u_k + \zeta_k \right) \\ &= \begin{bmatrix} 1 & \delta_t \\ -\frac{\delta_t k}{m} & 1 - \frac{\delta_t b}{m} \end{bmatrix} x_k + \begin{bmatrix} 0 \\ \frac{\delta_t}{m} \end{bmatrix} u_k + \delta_t \zeta_k = Ax_k + Bu_k + \tilde{\zeta}_k, \end{aligned}$$

where $\tilde{\zeta}_k$ is normally distributed with appropriate covariance, and A and B are matrices of appropriate size as defined in Def. 6.1.

6.1.2 Problem statement

We formalize the problem that we solve in this chapter, which deviates from Problem 5.2 only in that it is tailored to linear DTSSs. Recall from Def. 4.11 that $\text{Pr}_\mu^S(x_I \models \varphi)$ denotes the probability that the execution of DTSS S satisfies the reach-avoid specification φ (see Def. 4.9) under the Markov policy μ (see Def. 4.6).

Problem 6.6 (Lower bound control for linear DTSS) Given a linear DTSS $S = (X, U, x_I, \zeta, f)$, a reach-avoid specification $\varphi = (X_T, X_U, h)$, and a desired threshold probability $\rho \in [0, 1]$, compute a Markov policy μ such that

$$\text{Pr}_\mu^S(x_I \models \varphi) \geq \rho,$$

or return `False` if no such policy could be found.

Outline | We will use the framework presented in Chapter 5 to solve Problem 6.6 in the following way. First, in Sect. 6.2, we present an algorithm to generate an abstraction of a linear DTSS in the form of an MDP. This abstraction leads to a probabilistic simulation relation (as per Def. 5.9) from the MDP to the DTSS, which (as discussed in Sect. 5.4) leads to a Markov policy that solves Problem 6.6.

However, computing the transition probabilities of this MDP requires integrating over the stochastic process noise, which is often intractable. Thus, in Sect. 6.3, we estimate these transition probabilities as intervals and represent the abstraction as an interval Markov decision process (IMDP), leading to a probabilistic alternating simulation relation (as per Def. 5.19). Each of these intervals contains the precise transition probability with a user-specified confidence probability.

In Sect. 6.4, we present an algorithm that uses this IMDP abstraction to solve Problem 6.6 with high confidence. In Sect. 6.5, we propose a method to reduce the size of abstractions significantly, and in Sect. 6.6, we demonstrate the applicability of our approach on several benchmarks. Finally, we survey related work in Sect. 6.7 and discuss open challenges in Sect. 6.8.

6.2 MDP Abstraction of Linear DTSS

We present an algorithm to generate an MDP *abstraction* $\mathcal{M} = (S, Act, s_I, P)$ of a linear DTSS. First, we partition a compact subset \mathcal{Z} of the state space $X := \mathbb{R}^n$ into a finite set of convex polytopes. We add an additional element to the partition representing $\mathbb{R}^n \setminus \mathcal{Z}$ to ensure that the partition covers the entire state space \mathbb{R}^n . Thereafter, we define the states, actions, and transition probabilities of the MDP abstraction.

abstraction

Definition 6.7 (Partition) A polyhedral *partition* $\Psi := \{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus \mathcal{Z}\}$ of a compact subset $\mathcal{Z} \subset \mathbb{R}^n$ is a finite collection of sets such that

1. Each \mathcal{V}_i is a convex polytope, i.e., $\mathcal{V}_i = \{x \in \mathbb{R}^n : H_i x \leq b_i\}$ for $H_i \in \mathbb{R}^{\xi_i \times n}$, $b_i \in \mathbb{R}^{\xi_i}$, and $\xi_i \in \mathbb{N}$;
2. The union of the regions covers \mathcal{Z} , i.e., $\mathcal{Z} = \bigcup_{i=1}^L \mathcal{V}_i$;
3. The interiors of all regions are disjoint, i.e., $\text{int}(\mathcal{V}_i) \cap \text{int}(\mathcal{V}_j) = \emptyset, \forall i, j \in \{1, \dots, L\}, i \neq j$.

partition

A partition Ψ induces the following binary relation.

6

Definition 6.8 (Induced relation) A partition Ψ induces a strict single-valued binary relation $R \subseteq \mathbb{R}^n \times \Psi$ such that $R(x) \in \{\mathcal{V} \in \Psi : x \in \mathcal{V}\}$ for all $x \in \mathbb{R}^n$.

Formally, Def. 6.7 is not a proper partition because the boundaries between elements can be contained in multiple sets. As a result, the induced binary relation is not unique for states $x \in \mathbb{R}^n$ precisely on the boundary of a partition element $\mathcal{V} \in \Psi$, i.e., for $x \in \text{closure}(\mathcal{V}) \setminus \text{int}(\mathcal{V})$. For such a state x that is precisely on the boundary between multiple partition elements, we arbitrarily choose one of these regions as being related to state x . However, due to Assumption 6.2, the probability for the DTSS state x to lie exactly on the boundary of any partition element $\mathcal{V} \in \Psi$ is zero. Thus, this arbitrary choice will not affect the correctness of our algorithm.

Remark 6.9 (Equivalence classes) Another common way to denote the partition element $R(x)$ related to a state $x \in \mathbb{R}^n$ is to use the common *equivalence class* notation $[x]_R$. Here, we instead use the notation based on R and R^{-1} . However, both notations are equivalent.

States | The states of the MDP abstraction are the elements of the partition Ψ , that is, $S := \Psi = \{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus \mathcal{Z}\}$. This definition of the MDP states is also depicted by Fig. 6.2. Thus, each state $s \in S$ is also a subset of \mathbb{R}^n . The initial MDP state $s_I := R(x_I) \in S$ is the partition element that contains x_I . We call the MDP state associated with the partition element $\mathbb{R}^n \setminus \mathcal{Z}$ the *absorbing state*.

absorbing state

Remark 6.10 (Relation with discrete states) Since, $S := \Psi$, we can replace the partition Ψ in Def. 6.8 by the MDP states S , which results in the binary relation $R \subseteq \mathbb{R}^n \times S$. As per Def. 5.3, $R(x) \in S$ is the (unique) MDP state that contains $x \in \mathbb{R}^n$, whereas $R^{-1}(s) \subset \mathbb{R}^n$ is the set of DTSS states related to $s \in S$.

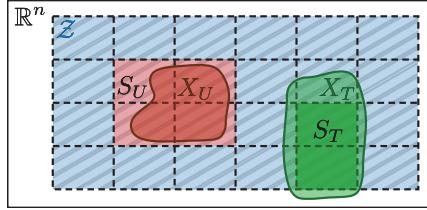


Figure 6.2: We partition a compact subset $Z \subset \mathbb{R}^n$ into a finite set of convex polytopes (dashed boxes). Furthermore, we define the MDP target states S_T as an underapproximation of X_T (green boxes), and we define the MDP unsafe states S_U as an overapproximation of X_U (red boxes).

Target and unsafe states | Let $X_T \subset \mathbb{R}^n$ and $X_U \subset \mathbb{R}^n$ be the target and unsafe sets of the reach-avoid specification for the DTSS under study. We choose S_T and S_U such that the resulting MDP reach-avoid specification is consistent with the specification for the DTSS, as defined in Def. 5.5. More precisely, recall from Def. 5.5 that these specifications are consistent, denoted by $(S_T, S_U) \leq_R (X_T, X_U)$, if

$$\mathbb{1}_{S_T}(s) \leq \mathbb{1}_{X_T}(x) \quad \text{and} \quad \mathbb{1}_{S_U}(s) \geq \mathbb{1}_{X_U}(x) \quad \forall (x, s) \in R. \quad (6.3)$$

Thus, we choose S_T and S_U exactly such that Eq. (6.3) holds, i.e.,

$$\begin{aligned} S_T &= \{s \in S : \forall (x, s) \in R, x \in X_T\} \quad \text{and} \\ S_U &= \{s \in S : \exists (x, s) \in R, x \in X_U\}. \end{aligned} \quad (6.4)$$

Essentially, Eq. (6.4) means that the MDP *underapproximates* the target states and *overapproximates* the unsafe states of the DTSS, such that $(S_T, S_U) \leq_R (X_T, X_U)$ holds as per Def. 5.5. This procedure is also depicted by Fig. 6.2. Note that, if the target states X_T and unsafe states X_U are exactly described by a union of partition elements $\{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus Z\}$, then the resulting specifications are even equivalent as per Def. 5.5, i.e., $(S_T, S_U) \equiv_R (X_T, X_U)$.

Actions | Actions in the abstraction correspond to executing control inputs $u_k \in U$ in the DTSS. We define $q \in \mathbb{N}$ MDP actions in total, so $Act := \{a_1, \dots, a_q\}$, $q \in \mathbb{N}$. Every action a is associated with a fixed continuous *target point* $d_a \in X$ defined on the state space of the DTSS. While not a restriction of our approach, it is often convenient to define one action for every MDP state $s \in S$ (except for the absorbing state) and choose the target point to be the center of the corresponding partition element.²

Choosing an MDP action $a \in Act$ corresponds with choosing a control input $u \in U$ in the DTSS such that, when we neglect the stochastic noise ζ , the successor state of the DTSS is *exactly* the target point d_a . To capture this intuition, we define the *backward reachable set* $\text{reach}^{-1}(d_a)$ for every target point d_a , $a \in Act$, as the set of all DTSS states

²If this choice results in an MDP that is too large, we may reduce the number of actions. If, on the other hand, a more refined abstraction is required, we may define additional actions.

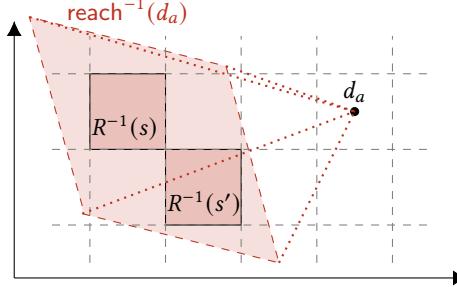


Figure 6.3: A part of a partition of $\mathcal{Z} \subset \mathbb{R}^2$, showing the backward reachable set $\text{reach}^{-1}(d_a)$ of an action $a \in \text{Act}$ with target point d_a . Action a is enabled in states s and s' , since $R^{-1}(s) \subseteq \text{reach}^{-1}(d_a)$ and $R^{-1}(s') \subseteq \text{reach}^{-1}(d_a)$.

from which d_a can be reached in one step:

$$\text{reach}^{-1}(d_a) = \{x \in \mathbb{R}^n : \exists u \in U. d_a = Ax + Bu + q\},$$

Now, let the control input space $U = \text{conv}(v^1, \dots, v^q)$, $q \in \mathbb{N}$, be given in its vertex representation, which is possible because U is assumed to be a convex polytope (see Assumption 6.3). Due to the linearity of the dynamics and Assumption 6.4 (invertibility of A), we rewrite the backward reachable set as

$$\text{reach}^{-1}(d_a) = \text{conv}(A^{-1}(d_a - Bv^i - q) : i = 1, \dots, q). \quad (6.5)$$

Intuitively, we enable an action $a \in \text{Act}$ in MDP state $s \in S$ if and only if the set $R^{-1}(s)$ is contained in $\text{reach}^{-1}(d_a)$. Thus, the subset $\text{Act}(s) \subseteq \text{Act}$ of actions enabled in the MDP state $s \in S$ is

$$\text{Act}(s) = \{a \in \text{Act} \mid R^{-1}(s) \subseteq \text{reach}^{-1}(d_a)\}. \quad (6.6)$$

As an example, Fig. 6.3 shows the backward reachable set of an action $a \in \text{Act}$, which contains regions $R^{-1}(s)$ and $R^{-1}(s')$. Hence, this action is enabled in states s and s' .

Transition probabilities | From Eq. (6.5), we observe that each MDP action $a \in \text{Act}$ is defined such that $d_a = Ax_k + Bu_k + q$, which holds for $u_k \in B^\dagger(d_a - Ax_k - q)$.³ Due to Eq. (6.6), we have $u_k \in B^\dagger(d_a - Ax_k - q) \subseteq U$ by construction for any DTSS state $x_k \in X$ for which $a \in \text{Act}(R(x_k))$. Since the noise is additive, the actual successor state is then $x_{k+1} = d_a + \xi_k$, which is a random variable with distribution $T(\cdot \mid x_k, B^\dagger(d_a - Ax_k - q))$, where T is the stochastic kernel as introduced in Sect. 4.2.2. As a result, for every $s \in S$, $a \in \text{Act}(s)$, and $s' \in S$, we define the transition probability $P(s, a)(s')$ as

$$\begin{aligned} P(s, a)(s') &= \int_{\mathbb{R}^n} \mathbf{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid x, B^\dagger(d_a - Ax - q)) \\ &= \mathbb{P}\{\omega \in \Omega : d_a + \xi_k(\omega) \in R^{-1}(s')\}. \end{aligned} \quad (6.7)$$

³Recall that B^\dagger denotes the pseudoinverse of B .

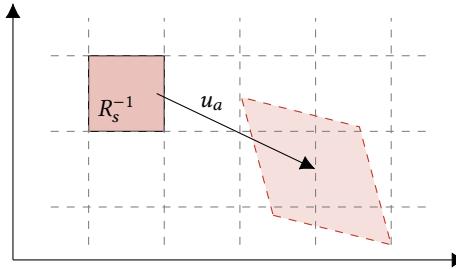


Figure 6.4: The forward abstraction method fixes a control input $u_a \in \text{Act}$ for each action $a \in \text{Act}$ and computes forward reachable sets by propagating $R^{-1}(s)$, $s \in S$ through the dynamics under the input u_a .

Remark 6.11 (Backward vs. forward methods) Most abstraction methods associate each abstract action $a \in \text{Act}$ with a *fixed control input* $u_a \in U$ and compute the *forward reachable set* associated with this input. We call this type of abstraction procedures *forward methods*. By contrast, we call our method based on backward reachability analysis a *backward method*. Using backward reachability is a key novelty of our approach that makes it distinct from other abstraction techniques.

As illustrated in Fig. 6.4, using the forward method, the distribution over DTSS states x_{k+1} associated with an MDP action $a \in \text{Act}$ depends on the precise state x_k where the action is chosen. By contrast, for every action $a \in \text{Act}$, our backward method leads to the same distribution $T(\cdot | x_k, B^\dagger(d_a - Ax_k - q))$ over DTSS states, independent of where the action is chosen.

However, computing backward reachable sets (especially for nonlinear systems) is more challenging than computing forward reachable sets. In the linear case, our backward method requires the system matrix A to be non-singular and the pair (A, B) to be controllable (Assumption 6.4), which is not needed for computing forward reachable sets. Thus, which of the two methods is better suited depends on the situation at hand.

Complete MDP | We now put all elements from the preceding paragraphs together. Given a partition Ψ and its induced binary relation $R \subseteq \mathbb{R}^n \times \Psi$, we define the MDP abstraction $\mathcal{M} = (S, \text{Act}, s_I, P)$ with:

- Set of states $S := \Psi = \{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus \mathcal{Z}\}$, with initial state $s_I := R(x_I) \in S$, and the set of target states S_T and unsafe states S_U as per Eq. (6.4);
- Set of actions $\text{Act} := \{a_1, \dots, a_q\}$ for some $q \in \mathbb{N}$, with the enabled actions $\text{Act}(s)$ defined by Eq. (6.6) for all $s \in S$;
- For all $s, s' \in S$ and $a \in \text{Act}(s)$, the probability $P(s, a)(s')$ is defined by Eq. (6.7).

Remark 6.12 (Number of transition probabilities of the MDP) For a given action $a \in \text{Act}$ and successor state $s' \in S$, the transition probability $P(s, a)(s')$ is equal for all MDP states s for which $a \in \text{Act}(s)$. In other words, for any two states $s, \tilde{s} \in S$ and an action a for which $a \in \text{Act}(s)$ and $a \in \text{Act}(\tilde{s})$, we have that

$P(s, a)(s') = P(\tilde{s}, a)(s')$ for all $s' \in S$. Thus, the MDP has at most $|Act| \cdot |S|$ unique transition probabilities (rather than at most $|S| \cdot |Act| \cdot |S|$ probabilities), which reduces the complexity of generating abstractions.

6.2.1 Relation induced by the abstract MDP

We close this section by showing that the abstract MDP \mathcal{M} induces a *probabilistic simulation relation*, which means that we can use the MDP abstraction to solve Problem 6.6 using the framework presented in Chapter 5.

probabilistic simulation relation

Proposition 6.13 (Probabilistic simulation relation) The induced binary relation $R \subseteq \mathbb{R}^n \times S$ is a probabilistic simulation relation from the abstract MDP \mathcal{M} to the DTSS \mathcal{S} , i.e., $\mathcal{M} \leq_R \mathcal{S}$.

Proof. First, observe that the initial states x_I and s_I are related by definition. Second, observe that Eq. (6.7) satisfies the requirements on the relation $D \subseteq \text{Distr}(X) \times \text{Distr}(S)$ in Def. 5.9. Thus, both requirements for a probabilistic simulation relation in Def. 5.9 are satisfied. ■

Moreover, the following lemma shows how to find an interface function (as defined by Def. 5.15) for the induced relation R between \mathcal{M} and \mathcal{S} . The intuition is that we restrict the DTSS control inputs to precisely those $u \in U$ that preserve the probabilistic simulation relation, which is the case if $d_a = Ax + Bu + q$.

6

Lemma 6.14 (Interface function) Let $R \subseteq \mathbb{R}^n \times S$ be the probabilistic simulation relation from abstract MDP \mathcal{M} to DTSS \mathcal{S} , and let $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}}$ be a Markov scheduler. The function $I_R^\sigma: X \times \{0, \dots, h-1\} \rightarrow 2^U$ defined for all $x \in X$ and for all $k \in \{0, \dots, h-1\}$ as

$$I_R^\sigma(x, k) := \{u \in U : d_a = Ax + Bu + q, a = \sigma_k(s), (x, s) \in R\}$$

is an interface function for the relation R as defined by Def. 5.15.

Proof. Let $\bar{x} \in X$ and $\bar{k} \in \{0, \dots, h-1\}$. To prove that I_R^σ is an interface function, we must show that, for all control inputs $u \in I_R^\sigma(\bar{x}, \bar{k})$ in the interface, it holds that

$$\forall s' \in S, \quad P(s, \sigma_{\bar{k}}(s))(s') = \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid \bar{x}, u).$$

From the definition of the MDP transition probabilities in Eq. (6.7), we have that

$$\begin{aligned} \forall s' \in S, \quad P(s, \sigma_{\bar{k}}(s))(s') &= \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid \bar{x}, B^\dagger(d_{\sigma_{\bar{k}}(s)} - Ax - q)) \\ &= \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid \bar{x}, \bar{u}), \end{aligned}$$

for any $\bar{u} \in I_R^\sigma(\bar{x}, \bar{k})$. Thus, it follows that $I_R^\sigma(x, k)$ is an interface function. ■

Remark 6.15 (The interface is independent from the noise) Observe that the interface function I_R^σ is *independent of the stochastic noise* ς . This independence is a crucial property of the interface function because it will allow us to solve Problem 6.6 even if the distribution of the noise is unknown.

Recall from Theorem 5.17 that by restricting the DTSS Markov policy to the interface function $I_R^\sigma(x, k)$, the satisfaction probability of σ on the MDP carries over as a lower bound to the DTSS. In mathematical terms, Theorem 5.17 implies that, for all MDP schedulers $\sigma \in \mathfrak{S}_{\text{Markov}}^M$ and for $\mu_k(x) \in I_R^\sigma(x, k)$, $k \in \{0, \dots, h - 1\}$, it holds that

$$\Pr_\sigma^M(s_I \models \neg S_U \cup^{\leq h} S_T) \leq \Pr_\mu^S(x_I \models \varphi).$$

Now, recall that Problem 6.6 asks for a Markov policy μ such that $\Pr_\mu^S(x_I \models \varphi) \geq \rho$. Thus, if we find that an optimal⁴ MDP scheduler has a satisfaction probability of at least ρ , then we can, by construction, compute a DTSS Markov policy which also has a satisfaction probability of at least ρ . If, on the other hand, the satisfaction probability on the MDP is below ρ , then we either need to improve the abstraction⁵ or conclude that we could not solve Problem 6.6 (and thus return False). Thus, our MDP abstraction yields a solution to the lower bound control problem for linear DTSS described in Problem 6.6. We present a more concrete algorithm in Sect. 6.4.

Remark 6.16 (Integration of the stochastic kernel) Even though the interface function is independent of the stochastic noise, the MDP's transition probabilities in Eq. (6.7) are defined in terms of integrals of the stochastic kernel T . Computing these integrals exactly is possible for simple (e.g., triangular) distributions but is often infeasible for more complex distributions. Moreover, distributions may not be precisely known (or even unknown) at all. Thus, whether we can compute the MDP abstraction in practice (and thus use Lemma 6.14 to solve Problem 6.6) depends on the noise distribution of the DTSS at hand.

6.3 Sampling-Based Probability Intervals

A common assumption to achieve computational tractability is that the process noise is Gaussian [PSQ13], e.g., as is classically assumed in linear-quadratic-Gaussian control [AM90]. However, in realistic problems, such as a UAV operating under turbulence, this assumption yields a poor approximation of the uncertainty [BOBW10]. Distributions may even be *unknown*, meaning that one cannot derive a set-bounded or a precise probabilistic representation of the noise. Motivated by these examples, we drop the common assumption that the distribution of the process noise ς is known.

Assumption 6.17 (Noise distribution unknown) The probability measure \mathbb{P} of the stochastic process $\varsigma = (\varsigma_k)_{k \in \mathbb{N}}$ of the DTSS $\mathcal{S} = (X, U, x_I, \varsigma, f)$ is unknown.

Assumption 6.17 prevents us from computing the MDP transition probabilities via Eq. (6.7). In this section, we describe our sampling-based method proposed in [1; 7]

⁴In fact, this result holds for any MDP scheduler, but in practice, we will work with optimal schedulers.

⁵We will discuss ways to improve the MDP abstraction in Sect. 6.4.

to estimate the transition probabilities using intervals. In a nutshell, we will leverage Assumptions 4.4 and 6.2, which respectively state that the noise is independent and identically distributed, and has a probability density function. Then, we assume access to a finite set of $N \in \mathbb{N}$ i.i.d. observations of the process noise, $\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}$, each of which is a realization $\varsigma_k(\omega) \in \mathcal{V}_\varsigma$ of the noise for some $\omega \in \Omega$ drawn according to \mathbb{P} .

With slight abuse of notation, we can regard the set of samples $\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}$ to be an element from the probability space Ω^N equipped with the product probability \mathbb{P}^N and the product Borel σ -algebra $\mathcal{B}(\Omega^N)$.⁶ For a fixed action $a \in \text{Act}$, each sample $\varsigma_k^{(i)}$, $i = 1, \dots, N$, is associated with a possible successor state $x_{k+1}^{(i)} = d_a + \varsigma_k^{(i)}$. As such, we can generate these samples by inferring the process noise from state trajectories of the DTSS, or we may sample from the noise distribution directly (e.g., if a simulator is available). Thus, in our setting, noise samples can be obtained at a relatively low cost.

A frequentist approach to estimation | As an example, we want to estimate the probability $P(s, a)(s')$ that choosing action $a \in \text{Act}(s)$ in state $s \in S$ leads to a transition to state $s' \in S$. Naturally, we can approximate this probability as the number of samples $N_{a,s}^{\text{in}}$, $\leq N$ leading to a transition to state s' , divided by the total of N samples. This approach is known as a *frequentist approach* and is statistically justified by the strong law of large numbers. Concretely, the value of $N_{a,s'}^{\text{in}}$ is obtained as follows.

Definition 6.18 (Sample counts) The number of samples $N_{a,s'}^{\text{in}} \in \{0, \dots, N\}$ leading to a successor state $x_{k+1} \in R^{-1}(s')$ associated with the MDP state $s' \in S$ is

$$N_{a,s'}^{\text{in}} := \left| \left\{ i \in \{1, \dots, N\} : (d_a + \varsigma_k^{(i)}) \in R^{-1}(s') \right\} \right|. \quad (6.8)$$

Similarly, $N_{a,s'}^{\text{out}} = N - N_{a,s'}^{\text{in}}$ is the number of samples for which $d_a + \varsigma_k^{(i)} \notin R^{-1}(s')$.

The frequentist approach is simple but may lead to estimates that deviate critically from their true values if the number of samples is limited (we illustrate this issue in the UAV experiment in Sect. 6.6.1). In what follows, we discuss how to render our method robust against such estimation errors.

6.3.1 Bounds for the transition probabilities

We present a method based on the *scenario approach* [CG18a] to compute *intervals of probabilities* instead of precise estimates. Specifically, for every transition (s, a, s') , we compute an upper and lower bound, i.e., an interval, that contains $P(s, a)(s')$ defined by Eq. (6.7) with a user-specified (high) confidence probability. We represent the resulting abstraction as an IMDP, where these intervals enter the uncertain transition function.

We first state the main contribution of this section, which is a non-trivial variant of [RPM23, Theorem 5], adapted to our context. Specifically, for a given transition

⁶More formally, recall from Def. 4.2 that each sample $\varsigma_k^{(i)}$, $i = 1, \dots, N$, is an element of \mathcal{V}_ς , so the set of samples $\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}$ is an element of \mathcal{V}_ς^N , whose probability measure is defined uniquely by the pushforward of the product measure \mathbb{P}^N under ς . This interpretation coincides with ours in practice, so we simplify notation and say that $\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}$ is an element of $(\Omega^N, \mathcal{B}(\Omega^N), \mathbb{P}^N)$.

(s, a, s') and the resulting number of samples $N_{a,s'}^{\text{out}}$ outside of region $R^{-1}(s')$ (as per Def. 6.18), Theorem 6.19 returns an interval $[\check{p}, \hat{p}]$ that contains $P(s, a)(s')$ with at least a pre-defined confidence probability $1 - \beta \in (0, 1)$.

Theorem 6.19 (PAC probability intervals) Let $\zeta_k^{(1)}, \dots, \zeta_k^{(N)}$ be a set of $N \in \mathbb{N}$ samples of the noise ζ_k , and let $\beta \in (0, 1)$ be a confidence parameter. Furthermore, let $a \in \text{Act}$ and $s' \in S$ be a state and action of the abstract MDP $\mathcal{M} = (S, \text{Act}, s_I, P)$, respectively, and determine the value of $N_{a,s'}^{\text{out}}$. Then, for any $s \in S$ such that $a \in \text{Act}(s)$, the transition probability $P(s, a)(s')$ is bounded by

$$\mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : \check{p} \leq P(s, a)(s') \leq \hat{p} \right\} \geq 1 - \beta, \quad (6.9)$$

where $\check{p} = 0$ if $N_{a,s'}^{\text{out}} = N$, and otherwise \check{p} is the solution of

$$\frac{\beta}{2N} = \sum_{i=0}^{N_{a,s'}^{\text{out}}} \binom{N}{i} (1 - \check{p})^i \check{p}^{N-i}, \quad (6.10)$$

and $\hat{p} = 1$ if $N_{a,s'}^{\text{out}} = 0$, and otherwise \hat{p} is the solution of

$$\frac{\beta}{2N} = 1 - \sum_{i=0}^{N_{a,s'}^{\text{out}}-1} \binom{N}{i} (1 - \hat{p})^i \hat{p}^{N-i}. \quad (6.11)$$

Theorem 6.19 states that, with a probability of at least $1 - \beta$, the transition probability $P(s, a)(s')$ is bounded by the obtained interval $[\check{p}, \hat{p}]$. Importantly, this claim holds for *any* probability space $(\Omega, \mathcal{F}, \mathbb{P})$ for the process noise ζ_k , $k \in \mathbb{N}$, that satisfies the previous assumptions, so we can bound the probability in Eq. (6.7), even when the probability distribution of the noise is unknown.

Remark 6.20 (Beta distribution) Eqs. (6.10) and (6.11) are cumulative distribution functions of a beta distribution with parameters $N_{a,s'}^{\text{out}} + 1$ (or $N_{a,s'}^{\text{out}}$) and $N - N_{a,s'}^{\text{out}}$ (or $N - N_{a,s'}^{\text{out}} - 1$), respectively [CG18a], which can be solved numerically for \check{p} or \hat{p} up to arbitrary precision. Thus, we can speed up computations by tabulating the intervals for all relevant values of N , β , and $N_{a,s'}^{\text{out}}$ upfront.

The proof of Theorem 6.19, which we provide in Sect. 6.3.3, relies on recasting the estimation of transition probabilities into a set of $2N$ optimization problems with different numbers of constraints. To obtain probably approximately correct (PAC) intervals on these probabilities, we use results from [RPM23] on the probability of *constraint violation* on the solutions to these optimization problems. Interestingly, these optimization problems can be solved analytically based on their geometry. As a result, Theorem 6.19 only depends on the sample count $N_{a,s'}^{\text{out}}$, the total number of samples N , and the confidence parameter β . We can thus compute PAC probability intervals without explicitly solving optimization programs. In practice, our method is as simple as the frequentist approach but has the notable advantage that we obtain robust intervals of probabilities.

Remark 6.21 (Difference to [RPM23]) The factor $\frac{1}{2N}$ on the left side in Eqs. (6.10) and (6.11) is not present in the original result in [RPM23, Theorem 5]. As we will show in Sect. 6.3.3, we obtain this extra term because we are implicitly considering $2N$ optimization problems to compute a probability interval. Dividing the confidence β by $2N$ in Eqs. (6.10) and (6.11) essentially means that we force the result from [RPM23] to hold for all of these $2N$ problems.

6.3.2 *The scenario approach

Toward the proof of Theorem 6.19, we introduce several core concepts related to the scenario approach, largely following the exposition from [CG18a; CCG21]. We remark that these concepts are only used for the derivation of the proof in Sect. 6.3.3, so the reader may decide to skip directly to Sect. 6.3.4.

*Section with details that can be skipped safely

6

Scaled polytopes | Recall from Def. 6.7 that each element $\mathcal{V} \in \Psi$ is a convex polytope defined as $\mathcal{V} = \{x \in \mathbb{R}^n : Hx \leq b\}$, where the matrix $H \in \mathbb{R}^{\xi \times n}$ and vector $b \in \mathbb{R}^\xi$ define the $\xi \in \mathbb{N}$ halfspace constraints (note that H , b , and ξ can be different for every element $\mathcal{V} \in \Psi$). With slight abuse of notation, we define $\mathcal{V}(\lambda)$ as a version of $\mathcal{V} \in \Psi$ which is *scaled* by a factor $\lambda \geq 0$ relative to a Chebyshev center⁷ $\tilde{v} \in \mathbb{R}^n$ of \mathcal{V} :

$$\mathcal{V}(\lambda) := \{x \in \mathbb{R}^n : Hx \leq \lambda(b - H\tilde{v}) + H\tilde{v}\}. \quad (6.12)$$

Note that $\mathcal{V}(1) = \mathcal{V}$ and that shifting by \tilde{v} ensures we scale around a point contained in \mathcal{V} , such that $\mathcal{V}(\lambda_1) \subset \mathcal{V}(\lambda_2)$ for all $0 \leq \lambda_1 < \lambda_2$. Using this notation, $R^{-1}(s')(\lambda)$ denotes a version of the partition element associated with MDP state $s' \in S$ scaled by a factor λ . A visualization of the scaling for an arbitrary region $R^{-1}(s')$ associated with a discrete successor state $s' \in S$ is shown in Fig. 6.5. In this example, the Chebyshev center is not unique since the circle can be shifted while remaining within $R^{-1}(s')$.

⁷Informally, a Chebyshev center of a bounded set Y is the center of the largest inscribed ball of Y ; see [BV14, Section 4.3.1] for details.

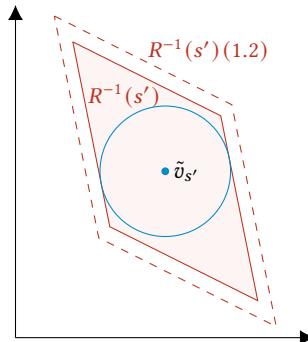


Figure 6.5: Polytope $R^{-1}(s')$ has a Chebyshev center $\tilde{v}_{s'}$ (which is not unique, as the circle can be shifted while remaining within $R^{-1}(s')$). Polytope $R^{-1}(s')(1.2)$ is scaled by a factor $\lambda = 1.2$ and is computed using Eq. (6.12).

Remark 6.22 (Rectification) In our papers [1; 7], we defined a scaled polytope as

$$\mathcal{V}(\lambda) := \{x \in \mathbb{R}^n : Hx \leq \lambda(b - \tilde{v}) + \tilde{v}\}. \quad (6.13)$$

Unfortunately, this original definition is incorrect (note that $b \in \mathbb{R}^\xi$ and $\tilde{v} \in \mathbb{R}^n$ have different dimensions). As such, Eq. (6.12) is a correction of the definition in Eq. (6.13). Nevertheless, all further results from [1; 7] remain valid.

Scenario optimization problem | For a fixed pair $a \in \text{Act}$ and $s' \in S$ and a given set of N noise samples $\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}$, we formulate the following linear optimization problem $\mathfrak{L}_{a,s'}^Q$, with a scalar decision variable $\lambda \geq 0$:

$$\begin{aligned} \mathfrak{L}_{a,s'}^Q : & \underset{\lambda \geq 0}{\text{minimize}} \quad \lambda \\ & \text{subject to } d_a + \varsigma_k^{(i)} \in R^{-1}(s')(\lambda) \quad \forall i \in \{1, \dots, N\} \setminus Q, \end{aligned} \quad (6.14)$$

where $Q \subset \{1, \dots, N\}$. Roughly speaking, solving $\mathfrak{L}_{a,s'}^Q$ amounts to finding the smallest λ such that the scaled polytope $R^{-1}(s')(\lambda)$ contains $d_a + \varsigma_k^{(i)}$ for all noise samples with index in $\{1, \dots, N\} \setminus Q$. Intuitively, the set Q is thus a subset of samples whose constraints have been *discarded* from the optimization problem.

Definition 6.23 (Active constraints) Let $\lambda^* \geq 0$ be an optimal solution to problem $\mathfrak{L}_{a,s'}^Q$. We say that the constraint for noise sample $\varsigma_k^{(i)}$ is *active* if $d_a + \varsigma_k^{(i)}$ is on the boundary of the scaled polytope $R^{-1}(s')(\lambda^*)$, i.e., if

$$d_a + \varsigma_k^{(i)} \in \text{closure}(R^{-1}(s')(\lambda^*)) \setminus \text{int}(R^{-1}(s')(\lambda^*)).$$

We denote the set of all active constraints by $\text{active}(\mathfrak{L}_{a,s'}^Q) \subset \{1, \dots, N\} \setminus Q$.

The next lemma states that, in our setting, there is always a single active constraint.

Lemma 6.24 (Uniqueness of active constraint) For all $Q \subset \{1, \dots, N\}$, $a \in \text{Act}$, $s' \in S$, and $\lambda^* \geq 0$, the number $|\text{active}(\mathfrak{L}_{a,s'}^Q)|$ of active constraints is one almost surely (i.e., with probability one).

Proof. First, as $\{1, \dots, N\} \setminus Q \neq \emptyset$, there is at least one active constraint (see [CG08] for details). Second, suppose there are multiple active constraints. Then, $d_a + \varsigma_k^{(i)}$ lies on the boundary of $R^{-1}(s')(\lambda^*)$ for multiple values $i \in \{1, \dots, N\} \setminus Q$. However, due to Assumption 6.2, the probability for this to occur is zero, so the claim follows. ■

Set of discarded samples | We use Lemma 6.24 to construct a sequence of strictly increasing subsets $Q_0, Q_1, \dots, Q_N \subset \{1, \dots, N\}$ of discarded samples, where at each step, we add the unique active constraint from the previous solution.

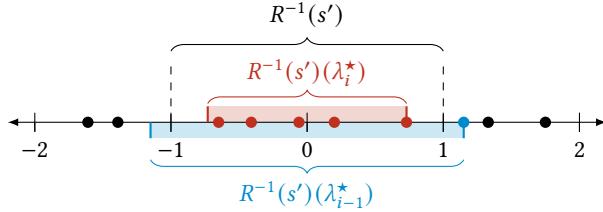


Figure 6.6: Bounding the region $R^{-1}(s') = [-1, 1]$ using $N = 10$ samples $d_a + \zeta_k^{(i)}$, $i = 1, \dots, 10$. Discarding $i := N_{a,s'}^{\text{out}} = 5$ samples defines the red region $R^{-1}(s')(\lambda_i^*) \subseteq R^{-1}(s')$, whereas discarding $i - 1$ samples defines the blue region $R^{-1}(s')(\lambda_{i-1}^*) \supset R^{-1}(s')$.

Definition 6.25 (Discarded samples) The sequence of subsets $Q_0, Q_1, \dots, Q_N \subset \{1, \dots, N\}$ is defined recursively, where:

1. The initial set $Q_0 = \emptyset$ is empty;
2. For all $\ell = 1, \dots, N$, we define $Q_\ell = Q_{\ell-1} \cup \text{active}(\mathfrak{L}_{a,s'}^{Q_{\ell-1}})$.

Thus, we have that $Q_0 \subset Q_1 \subset \dots \subset Q_N$. Furthermore, let λ_ℓ^* denote the optimal solution to problem $\mathfrak{L}_{a,s'}^{Q_\ell}$. From Lemma 6.24, it follows that $\lambda_0^* < \lambda_1^* < \dots < \lambda_N^*$ with probability one.

Under/overapproximating regions | Recall from Def. 6.18 that $N_{a,s'}^{\text{out}}$ is the number of samples leading to a successor state outside of region $R^{-1}(s')$ for state $s' \in S$. The following lemma uses $N_{a,s'}^{\text{out}}$ to define an under- or overapproximation of $R^{-1}(s')$.

Lemma 6.26 (Over/underapproximating regions) Let $i := N_{s'}^{\text{out}} \leq N$ be shorthand for the number of successor state samples outside of $R^{-1}(s')$. Then, it holds that $R^{-1}(s')(\lambda_i^*) \subseteq R^{-1}(s') \subset R^{-1}(s')(\lambda_{i-1}^*)$.

Proof. First consider optimization problem $\mathfrak{L}_{a,s'}^{Q_i}$ with discarded samples Q_i . In this case, exactly those samples for which $d_a + \zeta_k^{(i)}$ is outside of $R^{-1}(s')$ have been discarded, so $\lambda_i^* \leq 1$ and thus, $R^{-1}(s')(\lambda_i^*) \subseteq R^{-1}(s')$. Similarly, for problem $\mathfrak{L}_{a,s'}^{Q_{i-1}}$, with discarded samples Q_{i-1} , we must have one sample outside of $R^{-1}(s')$, so $\lambda_{i-1}^* > 1$ and thus, $R^{-1}(s') \subset R^{-1}(s')(\lambda_{i-1}^*)$. This concludes the proof. ■

Intuitively, for $i := N_{s'}^{\text{out}} \leq N$, the scaled polytope $R^{-1}(s')(\lambda_i^*)$ contains *exactly* those samples in $R^{-1}(s')$, while $R^{-1}(s')(\lambda_{i-1}^*)$ additionally contains *the sample closest outside of* $R^{-1}(s')$, as visualized in Fig. 6.6 for a 1-dimensional example.

Risk of violation | As a final ingredient, we introduce the concept of *risk* (or the *violation probability*), which is the probability that the DTSS state x_{k+1} is *not* in a given subset of the state space upon choosing action $a \in \text{Act}(s)$ in state $s \in S$ [CG08].

Definition 6.27 (Risk of violation) The *risk* $\mathbb{P}_a(x_{k+1} \notin V)$ that $x_{k+1} = d_a + \varsigma_k$ is not in a Borel set $V \in \mathcal{B}(X)$ upon choosing abstract action $a \in \text{Act}$ is defined as the shorthand notation

$$\mathbb{P}_a(x_{k+1} \notin V) := \mathbb{P}\{\omega \in \Omega : d_a + \varsigma_k(\omega) \notin V\}.$$

Crucially, observe from Eq. (6.7) that the transition probability $P(s, a)(s')$ we aim to estimate is the complement of the violation probability over $R^{-1}(s')$, i.e.,

$$P(s, a)(s') = 1 - \mathbb{P}_a(x_{k+1} \notin V).$$

6.3.3 *Proof of Theorem 6.19

We use the ingredients from Sect. 6.3.2 to prove Theorem 6.19. The proof is adapted from [RPM23, Theorem 5], which requires three key assumptions: (1) the scenario problem belongs to the class of so-called *fully-supported problems*,⁸ (2) its solution is unique, and (3) discarded samples violate the optimal solution with probability one. In our case, (1) is satisfied as $\mathfrak{L}_{a,s'}^Q$ has one decision variable and one active constraint with probability one (see Lemma 6.24). Furthermore, (2) is implied by Assumption 6.2, and (3) is satisfied by our choice of the discarded samples via Def. 6.25.

For the remainder of the proof, we choose a fixed action $a \in \text{Act}$ and a state $s' \in S$. Under the three assumptions above, [RPM23, Theorem 5] states that, for a fixed $\ell \in \{0, \dots, N\}$, the risk associated with an optimal solution λ_ℓ^\star to Eq. (6.14) for discarded samples Q_ℓ satisfies the following expression:

$$\mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \notin R^{-1}(s')(\lambda_\ell^\star)\} \leq \epsilon \right\} = 1 - \sum_{i=0}^{\ell} \binom{N}{i} \epsilon^i (1-\epsilon)^{N-i}. \quad (6.15)$$

Eq. (6.15) is the cumulative distribution function (CDF) $F_\ell: [0, 1] \rightarrow [0, 1]$ of a beta distribution with parameters $\ell + 1$ and $N - \ell$, which is defined for all $\epsilon \in [0, 1]$ as

$$F_\ell(\epsilon) = \mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \notin R^{-1}(s')(\lambda_\ell^\star)\} \leq \epsilon \right\} = \tilde{\beta}. \quad (6.16)$$

Thus, for any $\ell \in \{0, \dots, N\}$ and $\epsilon \in [0, 1]$, Eq. (6.16) returns the confidence probability $\tilde{\beta} \in (0, 1)$ by which the probability $\mathbb{P}_a \{x_{k+1} \notin R^{-1}(s')(\lambda_\ell^\star)\}$ is upper bounded by ϵ . Conversely, we can compute the value of ϵ that results in a confidence probability of $\tilde{\beta}$, using the percent point function (PPF) $G_\ell(\tilde{\beta})$:

$$\mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \notin R^{-1}(s')(\lambda_\ell^\star)\} \leq G_\ell(\tilde{\beta}) \right\} = \tilde{\beta}. \quad (6.17)$$

The PPF is the inverse of the CDF, so by definition, we have

$$\epsilon = G_\ell(\tilde{\beta}) = G_\ell(F_\ell(\epsilon)). \quad (6.18)$$

⁸An optimization problem is fully supported if the number of active constraints is equal to the number of decision variables with probability one. See [CG08] for a more formal definition.

Since $\mathbb{P}_a(x \in \mathcal{V}) + \mathbb{P}_a(x \notin \mathcal{V}) = 1$ for any x and \mathcal{V} , we rewrite Eq. (6.16) as

$$F_\ell(\epsilon) = \mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_\ell^\star)\} \geq 1 - \epsilon \right\} = \tilde{\beta}. \quad (6.19)$$

By defining $p = 1 - \epsilon$, Eqs. (6.18) and (6.19) are combined as

$$\mathbb{P}^N \left\{ 1 - G_\ell(\tilde{\beta}) \leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_\ell^\star)\} \right\} = 1 - \sum_{i=0}^{\ell} \binom{N}{i} (1-p)^i p^{N-i} = \tilde{\beta}. \quad (6.20)$$

Lower bound | We first prove the lower bound. There are N possible values for ℓ , ranging from 0 to $N - 1$. The case $\ell = N$ (i.e., all samples are discarded) is treated as a special case in Theorem 6.19. We fix $\tilde{\beta} = 1 - \frac{\beta}{2N}$ in Eq. (6.20), yielding the equations

$$\begin{aligned} \mathbb{P}^N \left\{ 1 - G_0 \left(1 - \frac{\beta}{2N} \right) \leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_0^\star)\} \right\} &= 1 - \frac{\beta}{2N} \\ \vdots &\vdots \\ \mathbb{P}^N \left\{ 1 - G_{N-1} \left(1 - \frac{\beta}{2N} \right) \leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N-1}^\star)\} \right\} &= 1 - \frac{\beta}{2N}. \end{aligned} \quad (6.21)$$

Denote the event that $1 - G_n \left(1 - \frac{\beta}{2N} \right) \leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_n^\star)\}$ for $n = 0, \dots, N - 1$ by \mathcal{A}_n . Regardless of n , this event has a probability of $\mathbb{P}^N \{\mathcal{A}_n\} = 1 - \frac{\beta}{2N}$, and its complement \mathcal{A}'_n of $\mathbb{P}^N \{\mathcal{A}'_n\} = \frac{\beta}{2N}$. Via Boole's inequality, we know that

$$\mathbb{P}^N \left\{ \bigcup_{i=0}^{N-1} \mathcal{A}'_n \right\} \leq \sum_{i=0}^{N-1} \mathbb{P}^N \{\mathcal{A}'_n\} = \frac{\beta}{2N} N = \frac{\beta}{2}. \quad (6.22)$$

Thus, for the intersection of all events in Eq. (6.21) we have

$$\mathbb{P}^N \left\{ \bigcap_{i=0}^{N-1} \mathcal{A}_n \right\} = 1 - \mathbb{P}^N \left\{ \bigcup_{i=0}^{N-1} \mathcal{A}'_n \right\} \geq 1 - \frac{\beta}{2}. \quad (6.23)$$

After observing the samples at hand, we replace ℓ by the value of $N_{a,s'}^{\text{out}}$ (as per Def. 6.18), resulting in one of the expressions in Eq. (6.21). This expression holds with *at least* the probability of the intersection of all events in Eq. (6.23). Thus, we obtain

$$\mathbb{P}^N \left\{ \check{p} \leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N_{a,s'}^{\text{out}}}^\star)\} \right\} \geq 1 - \frac{\beta}{2}, \quad (6.24)$$

where $\check{p} = 0$ if $N_{a,s'}^{\text{out}} = N$ (which is a trivial lower bound), and otherwise $\check{p} = 1 - G_{N_{a,s'}^{\text{out}}} \left(1 - \frac{\beta}{2N} \right)$ is the solution for p to Eq. (6.20), with $\ell = N_{a,s'}^{\text{out}}$ and $\tilde{\beta} = 1 - \frac{\beta}{2N}$:

$$1 - \frac{\beta}{2N} = 1 - \sum_{i=0}^{N_{a,s'}^{\text{out}}} \binom{N}{i} (1-p)^i p^{N-i}, \quad (6.25)$$

which is equivalent to Eq. (6.10).

Upper bound | Eq. (6.20) is rewritten as an upper bound as

$$\mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_\ell^\star)\} < 1 - G_\ell(\tilde{\beta}) \right\} = 1 - \tilde{\beta}, \quad (6.26)$$

where ℓ ranges from 0 to $N - 1$. However, to obtain high-confidence guarantees on the upper bound, we now fix $\tilde{\beta} = \frac{\beta}{2N}$, which yields the series of equations

$$\begin{aligned} \mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_0^\star)\} < 1 - G_0\left(\frac{\beta}{2N}\right) \right\} &= 1 - \frac{\beta}{2N} \\ \vdots &\quad \vdots \\ \mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N-1}^\star)\} < 1 - G_{N-1}\left(\frac{\beta}{2N}\right) \right\} &= 1 - \frac{\beta}{2N}. \end{aligned} \quad (6.27)$$

Analogous to the lower bound case, Boole's inequality implies that the intersection of all expressions in Eq. (6.27) has a probability of at least $1 - \frac{\beta}{2}$. After observing the samples at hand, we replace ℓ by $N_{a,s'}^{\text{out}} - 1$, yielding one of the expressions in Eq. (6.27). For this expression, it holds that

$$\mathbb{P}^N \left\{ \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N_{a,s'}^{\text{out}}-1}^\star)\} \leq \hat{p} \right\} \geq 1 - \frac{\beta}{2}, \quad (6.28)$$

where $\hat{p} = 1$ if $N_{a,s'}^{\text{out}} = 0$ (which is a trivial upper bound), and otherwise $\hat{p} = 1 - G_{N_{a,s'}^{\text{out}}-1}\left(\frac{\beta}{2N}\right)$ is the solution for p to Eq. (6.20), with $\ell = N_{a,s'}^{\text{out}} - 1$ and $\tilde{\beta} = \frac{\beta}{2N}$:

$$\frac{\beta}{2N} = 1 - \sum_{i=0}^{N_{a,s'}^{\text{out}}-1} \binom{N}{i} (1-p)^i p^{N-i}, \quad (6.29)$$

which is equivalent to Eq. (6.11).

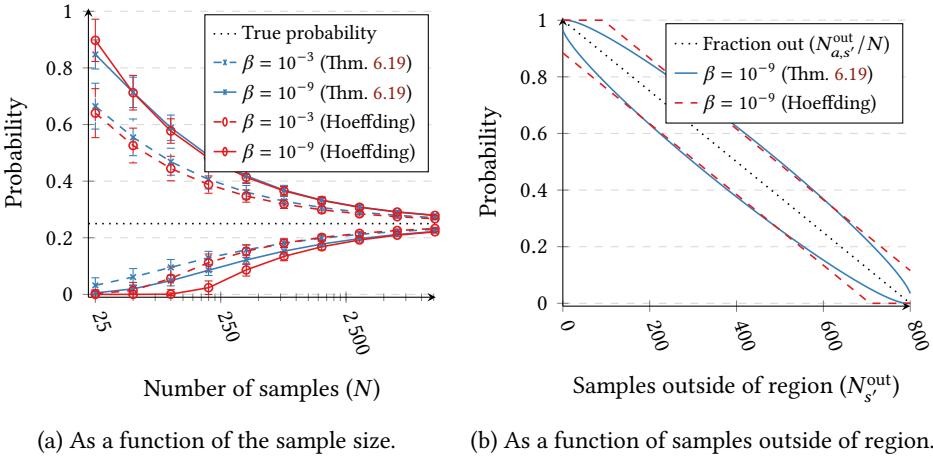
Probability interval | We invoke Lemma 6.26, which states that $R^{-1}(s')(\lambda_i^\star) \subseteq R^{-1}(s') \subset R^{-1}(s')(\lambda_{i-1}^\star)$. Thus, for all $s \in S$ for which $a \in \text{Act}(s)$, we have that

$$\begin{aligned} \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N_{a,s'}^{\text{out}}}^\star)\} &\leq \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')\} = P(s, a)(s') \\ &< \mathbb{P}_a \{x_{k+1} \in R^{-1}(s')(\lambda_{N_{a,s'}^{\text{out}}-1}^\star)\}. \end{aligned} \quad (6.30)$$

We use Eq. (6.30) to write Eqs. (6.24) and (6.28) in terms of the transition probability $P(s, a)(s')$. By applying Boole's inequality, we combine Eqs. (6.24) and (6.28) as follows:

$$\begin{aligned} \mathbb{P}^N \left\{ \check{p} \leq P(s, a)(s') \cap P(s, a)(s') \leq \hat{p} \right\} &= \mathbb{P}^N \left\{ \check{p} \leq P(s, a)(s') \leq \hat{p} \right\} \\ &\geq 1 - \beta, \end{aligned} \quad (6.31)$$

which is equal to Eq. (6.9), so we conclude the proof.



(a) As a function of the sample size. (b) As a function of samples outside of region.

Figure 6.7: Probability intervals from Theorem 6.19 vs. Hoeffding's inequality.

6

6.3.4 Tightness of probability intervals

We investigate the tightness of the probability intervals obtained from Theorem 6.19. We can interpret a transition probability $P(s, a)(s')$ as the probability of a Bernoulli random variable. In this context, we may use *Hoeffding's inequality*, a well-known concentration inequality, to infer PAC bounds on this probability [BLM13]. In particular, given N successor state samples of which $N_{a,s'}^{\text{in}}$ are contained in region $R^{-1}(s')$, Hoeffding's inequality states that, for any $\beta \in (0, 1)$, it holds that

$$\mathbb{P}^N \left\{ \frac{N_{a,s'}^{\text{in}}}{N} - \varepsilon \leq P(s, a)(s') \leq \frac{N_{a,s'}^{\text{in}}}{N} + \varepsilon \right\} \geq 1 - \beta,$$

where $\varepsilon = \sqrt{\frac{1}{2N} \log(\frac{2}{\beta})}$. In what follows, we evaluate the tightness of our intervals obtained from Theorem 6.19, versus those obtained from Hoeffding's inequality.

Hoeffding's inequality

Number of noise samples N | To illustrate how the choice for the number of samples N affects the tightness of the intervals, consider a system with a 1-dimensional state $x_k \in \mathbb{R}$, where the distribution over successor states for a specific action $a \in \text{Act}$ with target point d_a is given by a uniform distribution over the domain $[-4, 4]$. For a given region $R^{-1}(s') = [-1, 1]$ (also shown in Fig. 6.6), we want to evaluate the probability that $x_{k+1} \in R^{-1}(s')$, which is 0.25. To this end, we apply Theorem 6.19 for different numbers of samples $25 \leq N \leq 12800$ and a confidence level of $\beta = 10^{-3}$ or 10^{-9} . The obtained probability bounds are random variables through their dependence on the samples, so we repeat each experiment with 100 000 sets of N noise samples, resulting in the probability intervals shown in Fig. 6.7a. We observe that increasing the number of samples reduces the uncertainty in the transition probability. Moreover, the lower bounds obtained from Theorem 6.19 are better than those obtained from Hoeffding's inequality, while the converse holds for the upper bounds.

Sample count $N_{a,s'}^{\text{out}}$ | To explain why Theorem 6.19 yields tighter lower bounds, while Hoeffding's inequality yields tighter upper bounds, we plot in Fig. 6.7b the probability intervals for $N = 800$ samples and different values of $0 \leq N_{a,s'}^{\text{out}} \leq N$ (recall that $N = N_{a,s'}^{\text{out}} + N_{a,s'}^{\text{in}}$). Our scenario-based approach results in *significantly tighter* intervals for values of $N_{a,s'}^{\text{out}}$ close to 0 and N . On the other hand, Hoeffding's inequality leads to *slightly better* intervals for moderate values of $N_{a,s'}^{\text{out}}$ around $\frac{N}{2}$. Hoeffding's inequality yields bounds given by the *sample mean* $N_{a,s'}^{\text{in}}$, plus or minus a *fixed* value of ϵ which is *independent of the sample mean*. Thus, if the probability $P(s, a)(s')$ is close to one, Hoeffding's inequality leads to intervals whose upper bound is above one (or whose lower bound is below one if $P(s, a)(s')$ close to zero).

6.4 Abstraction-Based Control Algorithm

Theorem 6.19 enables us to compute a PAC interval for each transition probability of the MDP abstraction from Sect. 6.2. In this section, we present an iterative algorithm that uses this result to solve Problem 6.6 with a predefined confidence probability. Our scheme is summarized in Fig. 6.8 and consists of an *offline planning phase* to generate an IMDP abstraction with PAC guarantees, and an *online control phase* to derive a Markov policy for the DTSS on the fly. We discuss both phases in more detail.

6.4.1 Interval MDP abstraction

The offline phase is presented in Algorithm 6.1 and follows the same general structure as Algorithm 5.1 from Chapter 5. We first generate an MDP abstraction by defining its states S , actions Act , initial state s_I , and for all $s \in S$ the set $Act(s)$ of enabled actions (lines 1–5). However, since the process noise has an unknown distribution (as per Assumption 6.17), we cannot compute the transition probabilities of the MDP abstraction exactly.

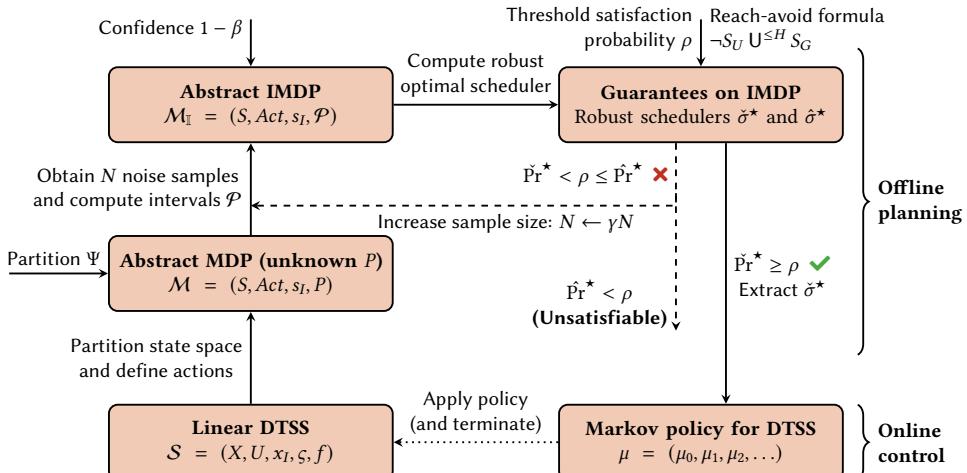


Figure 6.8: Our iterative IMDP abstraction scheme for linear DTSSs, where N is the number of noise samples to compute probability intervals, ρ is the threshold reach-avoid probability, and $\check{P}r^*$, $\hat{P}r^*$, and σ^* are defined as in Algorithm 6.1.

Algorithm 6.1 Iterative IMDP abstraction with PAC guarantees.

Input: Linear DTSS $\mathcal{S} = (X, U, x_I, \xi, f)$; reach-avoid specification $\varphi = (X_T, X_U, h)$; satisfaction probability threshold $\rho \in [0, 1]$

Params: Partition Ψ ; confidence lvl. β ; increment factor $\gamma > 1$; initial size $N_0 \in \mathbb{N}$; maximum sample size $N_{\max} \in \mathbb{N}$

Output: Optimal robust IMDP scheduler $\check{\sigma}^*$

```

1:  $S \leftarrow \Psi = \{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus \mathcal{Z}\}$                                 ▷ Abstract states
2:  $S_T, S_U \subseteq S$  as per Def. 5.5 for  $\varphi$                                 ▷ Target/unsafe states
3:  $s_I \leftarrow R(x_I) \in S$                                               ▷ Initial state
4:  $Act \leftarrow \{a_1, \dots, a_q\}, q \in \mathbb{N}$                                 ▷ Abstract actions
5:  $Act(s) \leftarrow \{a \in Act : R^{-1}(s) \subseteq \text{reach}^{-1}(d_a)\}, \quad \forall s \in S$     ▷ Enabled actions
6:  $N \leftarrow N_0; \quad \check{Pr}^* \leftarrow 0; \quad \hat{Pr}^* \leftarrow 1$ 
7: while  $\check{Pr}^* < \rho \leq \hat{Pr}^*$  and  $N \leq N_{\max}$  do
8:    $\zeta_k^{(1)}, \dots, \zeta_k^{(N)} \in \Omega$                                 ▷ Obtain  $N$  noise samples
9:   for all  $a$  in  $Act$  do
10:    for all  $s'$  in  $S$  do
11:      Compute  $[\check{p}, \hat{p}]$  for  $N, N_{s'}^{out}$ , and  $\beta$  (via Theorem 6.19)
12:       $\mathcal{P}(s, a, s') \leftarrow [\check{p}, \hat{p}], \quad \forall s \in S$  such that  $a \in Act(s)$     ▷ Prob. intervals
13:       $\mathcal{M}_{\mathbb{I}} := (S, Act, s_I, \mathcal{P})$                                               ▷ Store IMDP
14:       $\check{Pr}^* \leftarrow \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s_I \models \neg S_U \cup^{\leq h} S_G)$ 
15:       $\hat{Pr}^* \leftarrow \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \max_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s_I \models \neg S_U \cup^{\leq h} S_G)$     ▷ Reach-avoid probs.
16:       $N \leftarrow \gamma N$ 
17:    if  $\rho > \check{Pr}^*$  then
18:      return False
19:    else
20:      return  $\check{\sigma}^* \leftarrow \operatorname{argmax}_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s_I \models \neg S_U \cup^{\leq h} S_G)$ 

```

Thus, we instead propose the following iterative IMDP abstraction scheme. In each iteration, we first obtain N samples of the noise (line 8). Recall from Sect. 6.3 that we can obtain these samples by inferring the process noise from previously generated state trajectories of the DTSS or by sampling the noise distribution directly using a simulator. For every action $a \in Act$ and successor state $s' \in S$, we compute a PAC transition probability interval $\mathcal{P}(\cdot, a, s') := [\check{p}, \hat{p}]$ using Theorem 6.19 (lines 9–12). We store the resulting IMDP as $\mathcal{M}_{\mathbb{I}} = (S, Act, s_I, \mathcal{P})$ (line 13).

For the IMDP, we compute the optimal reach-avoid probabilities under the minimizing and maximizing choices $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}$ of nature (lines 14–15):

$$\check{Pr}^* := \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s_I \models \neg S_U \cup^{\leq h} S_G), \quad \text{and}$$

$$\hat{Pr}^* := \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \max_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s_I \models \neg S_U \cup^{\leq h} S_G),$$

which can be computed using robust value iteration (see Sect. 3.3.4). We then proceed in one of the following ways:

1. If $\check{\Pr}^* \geq \rho$, the algorithm terminates, returning an optimal robust scheduler $\check{\sigma}^*$, which is computed via robust value iteration as
$$\check{\sigma}^* := \operatorname{argmax}_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_I}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_I}} \Pr_{\sigma, \tau}^{\mathcal{M}_I}(s_I \models \neg S_U \cup^{\leq h} S_G).$$
2. If $\check{\Pr}^* < \rho$ (i.e., the reach-avoid probability under the most pessimistic nature $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_I}$ violates the threshold ρ) but $\check{\Pr}^* \geq \rho$, we increase the number of samples to $N \leftarrow \gamma N$ for a fixed $\gamma > 1$ and repeat the loop (until the maximum size N_{\max} is exceeded).
3. If $\hat{\Pr}^* < \rho$ (i.e., the reach-avoid probability under the most optimistic nature $\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_I}$ violates the threshold ρ), the problem is (with high confidence) unsatisfiable on the current partition. In this case, increasing the number of samples does not help, so we terminate the scheme.

6.4.2 Solving Problem 6.6 with high probability

Upon termination, Sect. 6.4.1 returns an optimal robust scheduler $\check{\sigma}^*$ for the IMDP abstraction (or returns `False` in case no satisfactory scheduler was found). What remains is to show that we can compute a Markov policy via a suitable interface function, such that the satisfaction probability on the IMDP carries over (as a lower bound) to the DTSS. We will obtain such an interface function by establishing the existence of a *probabilistic alternating simulation relation*.

Recall from Def. 5.19 that a probabilistic alternating simulation relation $R \subseteq X \times S$ is the natural extension of a relation between an MDP and a DTSS, to an IMDP (or, more generally, a robust Markov decision process (RMDP)) and a DTSS. The first step is to show that, with a certain (high) probability, Algorithm 6.1 returns an IMDP \mathcal{M}_I that induces a probabilistic alternating simulation relation with the DTSS.

Lemma 6.28 Suppose that Algorithm 6.1 terminates with an abstract IMDP $\mathcal{M}_I = (S, \text{Act}, s_I, \mathcal{P})$. Then, the following holds:

$$\mathbb{P}^N \{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : \mathcal{M}_I \leq_R^{\text{alt}} \mathcal{S} \} \geq 1 - \beta \cdot |\text{Act}| \cdot |S|.$$

Proof. The lemma is almost a direct application of Lemma 5.22 to the IMDP \mathcal{M}_I . Concretely, this lemma states that

$$\mathcal{M} \leq_R \mathcal{S} \text{ and } \left[P(s, a) \in \mathcal{P}(s, a) \quad \forall s \in S, \forall a \in \text{Act}(s) \right] \implies \mathcal{M}_I \leq_R^{\text{alt}} \mathcal{S}, \quad (6.32)$$

where $\mathcal{M} = (S, \text{Act}, s_I, P)$. For all $a \in \text{Act}$ and $s' \in S$, we have that

$$\mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : P(s, a)(s') \in \mathcal{P}(s, a)(s'), \quad \forall s \in S \right\} \geq 1 - \beta.$$

That is, every MDP distribution $P(s, a)$ is contained in the set of distributions $\mathcal{P}(s, a)$ of the IMDP with a probability of at least $1 - \beta$. Recall from Remark 6.12 that the

abstraction has at most $|Act| \cdot |S|$ unique transition probabilities, such that

$$\begin{aligned} \mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : P(s, a)(s') \in \mathcal{P}(s, a)(s'), \forall s, s' \in S, a \in Act \right\} \\ \geq 1 - \beta \cdot |Act| \cdot |S|. \end{aligned}$$

Thus, the second condition in Eq. (6.32) holds with probability at least $1 - \beta \cdot |Act| \cdot |S|$, whereas the first condition ($\mathcal{M} \leq_R \mathcal{S}$) is always satisfied (as we established in Proposition 6.13). Hence, we conclude that $\mathcal{M}_I \leq_R^{\text{alt}} \mathcal{S}$ holds with a probability of at least $1 - \beta \cdot |Act| \cdot |S|$. ■

As the final step, we provide an explicit definition of the desired robust interface function. To distinguish from the MDP interface function I_R^σ defined in Lemma 6.14, let us denote the IMDP interface function as J_R^σ . Recall from Def. 5.24 that this robust interface function must satisfy

$$J_R^\sigma(x, k) = \left\{ u \in U : \forall s' \in S, \mathcal{P}(s, \sigma_k(s))(s') \ni \int_{\mathbb{R}^n} \mathbb{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid x, u) \right\}, \quad (6.33)$$

where $s \in S$ is the MDP state related to $x \in X$, i.e., $(s, x) \in R$, and T is the stochastic kernel of the DTSS. Thus, in general, the interface function involves an integral over the stochastic kernel T , which (as discussed before) we cannot compute exactly. Luckily, and as before (see Remark 6.15), the interface function for the particular IMDP abstraction we generate *does not depend on the stochastic kernel*. In fact, the following result shows that we can restrict the Markov policy to the same interface function as defined in Lemma 6.14 for the MDP abstraction.

Theorem 6.29 (Solution to Problem 6.6) Let $\mathcal{S} = (X, U, x_I, \zeta, f)$ be a DTSS with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\rho \in [0, 1]$. Suppose that Algorithm 6.1 terminates and returns an optimal robust scheduler $\check{\sigma}^* \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_I}$. Then, it holds that

$$\mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : \Pr_\mu^{\mathcal{S}}(x_I \models \varphi) \geq \rho \right\} \geq 1 - \beta \cdot |Act| \cdot |S|,$$

where the policy is chosen as $\mu_k(x) \in I_R^{\check{\sigma}^*}(x, k)$ for all $x \in X$ and $k \in \{0, \dots, h-1\}$, with the same interface function $I_R^{\check{\sigma}^*}(x, k)$ as for the MDP abstraction:

$$I_R^{\check{\sigma}^*}(x, k) := \left\{ u \in U : d_a = Ax + Bu + q, a = \check{\sigma}_k^*(x), (x, s) \in R \right\}.$$

Proof. For now, let us assume that $\mathcal{M}_I \leq_R^{\text{alt}} \mathcal{S}$ holds. Then, we know from Corollary 5.25 that

$$\Pr_\mu^{\mathcal{S}}(x_I \models \varphi) \geq \min_{\tau \in \mathfrak{T}_{\text{Markov}}} \Pr_{\check{\sigma}^*, \tau}^{\mathcal{M}_I}(s_I \models \neg S_U \cup^{\leq h} S_T) \quad (6.34)$$

where $\mu_k(x) \in J_R^{\check{\sigma}^*}(x, k)$ for all $x \in X$ and $k \in \{0, \dots, h-1\}$. We need to show that any Markov policy $\mu_k(x) \in I_R^{\check{\sigma}^*}(x, k)$ in the MDP's interface function is also

contained in the IMDP's interface function, i.e., $\mu_k(x) \in J_R^{\sigma^*}(x, k)$ for all $x \in X$ and $k \in \{0, \dots, h-1\}$. Observe that $P(s, \check{\sigma}_k^*(s))(s') \in \mathcal{P}(s, \check{\sigma}_k^*(s))(s')$, which implies

$$\mathcal{P}(s, \check{\sigma}_k^*(s))(s') \ni P(s, \check{\sigma}_k^*(s))(s') = \int_{\mathbb{R}^n} \mathbf{1}_{R^{-1}(s')}(\xi) \cdot T(d\xi \mid x, u).$$

Thus, the MDP's interface function $I_R^{\sigma^*}$ is a subset of the IMDP's interface function $J_R^{\sigma^*}$, such that $I_R^{\sigma^*}(x, k) \subseteq J_R^{\sigma^*}(x, k)$. It follows that

$$\mu_k(x) \in I_R^{\sigma^*}(x, k) \implies \mu_k(x) \in J_R^{\sigma^*}(x, k),$$

which shows that $I_R^{\sigma^*}$ is a valid interface function for the DTSS.

Finally, we assumed in this proof that $\mathcal{M}_{\mathbb{I}} \leq_R^{\text{alt}} \mathcal{S}$ but, in fact, this is only true with a probability of at least $1 - \beta \cdot |\text{Act}| \cdot |S|$. By incorporating this probability into the statements in this proof, we obtain the desired claim and thus conclude the proof. ■

Theorem 6.29 shows that Algorithm 6.1 yields, with probability at least $1 - \beta \cdot |\text{Act}| \cdot |S|$ (with respect to the noise samples $(\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N$), a solution to Problem 6.6. In other words, termination of Sect. 6.4.1 implies that (with a probability of at least $1 - \beta \cdot |\text{Act}| \cdot |S|$) we have found a solution to Problem 6.6. Interestingly, we concluded that the interface function that determines the Markov policy is the same as for the MDP abstraction and is thus (as discussed earlier on) independent of the stochastic kernel T . This observation allows for simpler control design, in which only the IMDP scheduler (but not the abstraction itself) is needed within the online control loop.

6.5 Exploiting Stability for Smaller Abstractions

The size of IMDPs can be expressed by the number of transitions in the underlying graph. Inspired by the fact that only a small subset of the transitions are taken under an optimal scheduler, we adapt our abstraction framework presented thus far to create smaller abstractions. We leverage the two-layer control design framework shown in Fig. 6.9, which first stabilizes the dynamics and then creates an abstraction of the closed-loop dynamics. Specifically, we use the composite feedback control given by

$$u_k := -Kx_k + u'_k, \quad (6.35)$$

where the gain matrix $K \in \mathbb{R}^{m \times n}$ represents a stabilizing Markov policy, and u'_k is a control input obtained from a Markov policy $u'_k := \mu_k(x_k)$ as we did before. We may obtain the feedback gain matrix by solving an instance of a linear quadratic regulator (LQR) control problem [FPE21].

Applying the composite feedback policy in Eq. (6.35) to the linear DTSS yields the closed-loop dynamics given by

$$\begin{aligned} x_{k+1} &= Ax_k + B(-Kx_k + u'_k) + q + \zeta_k \\ &= A_{\text{cl}}x_k + Bu'_k + q + \zeta_k, \end{aligned}$$

where $A_{\text{cl}} = A - BK$. We assume that the feedback gain K satisfies the input constraints in the following way.

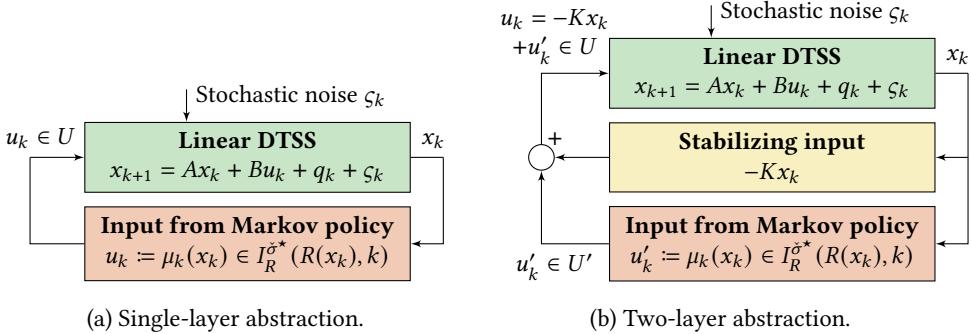


Figure 6.9: The feedback design framework from Sect. 6.4 (a), versus a two-layer feedback design framework, which combines a linear feedback gain with the Markov policy obtained from the abstraction.

Assumption 6.30 (Non-singular closed-loop dynamics) The gain matrix $K \in \mathbb{R}^{m \times n}$ is such that $-Kx \in U$ for all $x \in \mathcal{Z}$ and the matrix A_{cl} is non-singular.

6.5.1 Backward reachable sets

We show how our IMDP abstraction algorithm from Sect. 6.4 can be employed together with the composite Markov policy in Eq. (6.35). Recall from Assumption 6.3 that the control input space U is a convex polytope $U = \{u \in \mathbb{R}^m : Gu \leq g\} \subset \mathbb{R}^m$. Let us define a second polytopic control input space $U' = \{u \in \mathbb{R}^m : G'u \leq g'\} \subset \mathbb{R}^m$, where $G' \in \mathbb{R}^{q \times m}$ and $g' \in \mathbb{R}^q$. Then, we replace the backward reachable set computation in Eq. (6.5) with one that depends explicitly on this newly defined set U' :

$$\begin{aligned} \text{reach}_{\text{cl}}^{-1}(d_a, U') = \{x \in \mathbb{R}^n : d_a &= A_{\text{cl}}x + Bu + q, \\ &-Kx + u' \in U, \quad u' \in U'\}. \end{aligned} \quad (6.36)$$

The constraint $-Kx + u' \in U$ enforces that the total input is admissible, and the constraint $u' \in U'$ will enable us to control the size of the abstraction.

Assumption 6.31 (Origin contained in U') The set $U' = \{u \in \mathbb{R}^m : G'u \leq h'\}$ contains the origin: $0 \in U'$.

Observe that Eq. (6.36) is of the same form as Eq. (6.5) (despite imposing additional constraints) and can thus be computed similarly, as shown by the following lemma.

Lemma 6.32 (Backward reachable set properties) Under Assumptions 6.30 and 6.31, the following holds:

- i) For any $d_a \in \mathbb{R}^n$, the set $\text{reach}_{\text{cl}}^{-1}(d_a, U')$ is non-empty;
- ii) $\text{reach}_{\text{cl}}^{-1}(d_a, U') = \{x \in \mathbb{R}^n \mid d_a = A_{\text{cl}}x + Bu' + q, u' \in \tilde{U}\}$, where $\tilde{U} \subset \mathbb{R}^m$ is a

convex polytope defined as

$$\tilde{U} = \{u \in \mathbb{R}^m : G(\alpha + \beta u) \leq g, G'u \leq g'\}, \quad (6.37)$$

with $\alpha = -K(A_{\text{cl}}^{-1}d_a - q)$ and $\beta = I + KA_{\text{cl}}^{-1}B$.

Proof. Item *i*): We will show that the point $\tilde{x} = A_{\text{cl}}^{-1}d_a \in \text{reach}_{\text{cl}}^{-1}(d_a, U')$. Note that this point \tilde{x} is obtained for $u' = 0$ in Eq. (6.36), which is an admissible input due to Assumption 6.31. Moreover, due to Assumption 6.30, we have that $-Kx \in U$, and thus, it holds that $\tilde{x} \in \text{reach}_{\text{cl}}^{-1}(d_a, U')$, which concludes the proof of item *i*).

Item *ii*): Solving the equality constraint in Eq. (6.36) for x yields $x = A_{\text{cl}}^{-1}(d_a - Bu - q)$, so the input constraint $-Kx + u' \in U$ can be written as

$$-K(A_{\text{cl}}^{-1}d_a - q) + (I + KA_{\text{cl}}^{-1}B)u' = \alpha + \beta u' \in U.$$

Thus, we have two convex polyhedral constraints on u , given by $\alpha + \beta u' \in U = \{u \in \mathbb{R}^m : Gu \leq g\}$ and $u' \in U' = \{u \in \mathbb{R}^m : G'u \leq g'\}$. The intersection of the feasible sets for u is the set \tilde{U} in Eq. (6.37), which concludes the proof of item *ii*). ■

Recall from Eq. (6.5) in Sect. 6.2 that the backward reachable set $\text{reach}^{-1}(d_a)$ can be computed by enumerating the vertices of U . Lemma 6.32 shows that $\text{reach}^{-1}(d_a)$ is of the same form as $\text{reach}^{-1}(d_a)$ and can thus be computed by enumerating the vertices of \tilde{U} defined in Eq. (6.37). However, the number of vertices to consider is generally higher due to the additional input constraint $u' \in U'$.

6.5.2 Constructing smaller abstractions

We can use the modified backward reachable set in Eq. (6.36) to construct abstractions with fewer enabled actions, as shown by the following lemma.

Lemma 6.33 (Smaller backward reachable sets) Compare $\text{reach}^{-1}(d_a)$ obtained from Eq. (6.5) with $\text{reach}_{\text{cl}}^{-1}(d_a, U'')$ obtained from Eq. (6.36). If $U' = \mathbb{R}^p$ in Eq. (6.36), then we have that $\text{reach}_{\text{cl}}^{-1}(d_a, U'') = \text{reach}^{-1}(d_a)$. Moreover, for any $U'' \subset U' \subseteq \mathbb{R}^p$, it holds that $\text{reach}_{\text{cl}}^{-1}(d_a, U'') \subseteq \text{reach}_{\text{cl}}^{-1}(d_a, U')$.

Proof. Letting $U' = \mathbb{R}^p$ and $A_{\text{cl}} = A - BK$ in Eq. (6.36) gives

$$\begin{aligned} \text{reach}_{\text{cl}}^{-1}(d_a, U') &= \{x \in \mathbb{R}^n : d_a = Ax + B(-Kx + u') + q, -Kx + u' \in U\} \\ &= \{x \in \mathbb{R}^n : d_a = Ax + Bu + q, u \in U\} \\ &= \text{reach}^{-1}(d_a), \end{aligned}$$

which concludes the proof of the first part of the lemma. Second, for $U'' \subset U'$, observe from Eq. (6.36) that $u' \in U'' \subset U'$. Thus, we obtain that $\text{reach}_{\text{cl}}^{-1}(d_a, U'') \subseteq \text{reach}_{\text{cl}}^{-1}(d_a, U')$, which concludes the second part of the proof. ■

Recall from Eq. (6.6) in Sect. 6.2 that an IMDP action $a \in \text{Act}$ is enabled in state $s \in S$ if and only if the set $R^{-1}(s)$ is contained in the backward reachable set $\text{reach}^{-1}(d_a)$. We replace $\text{reach}^{-1}(d_a)$ with the modified version $\text{reach}_{\text{cl}}^{-1}(d_a, U')$, such that for all $s \in S$,

the subset $\text{Act}(s) \subseteq \text{Act}$ of enabled actions becomes

$$\text{Act}(s) = \{a \in \text{Act} : R^{-1}(s) \subseteq \text{reach}_{\text{cl}}^{-1}(d_a, U')\}. \quad (6.38)$$

By controlling the size of U' , we can control the size of the modified backward reachable set $\text{reach}_{\text{cl}}^{-1}(d_a)$ and thus the size of the graph of the IMDP (through the number of enabled actions). In Sect. 6.6, we show how suitable choices for the feedback gain K and input constraint U' may lead to significantly smaller IMDP abstractions.

6.6 Experimental Evaluation

We perform numerical experiments to demonstrate the applicability of Algorithm 6.1 for solving Problem 6.6. We implement Algorithm 6.1 in a Python tool called DynAbs, which we present in more detail in Chapter 14. Given a linear DTSS, a reach-avoid specification, a satisfaction probability threshold, and the required parameters, DynAbs applies the iterative abstraction scheme from Algorithm 6.1. When DynAbs returns an optimal robust IMDP scheduler σ^* , then it is guaranteed (with the specified confidence probability) to solve Problem 6.6 as per Theorem 6.29.

Reproducibility | The Python code to reproduce the experimental results is provided with DynAbs; see Chapter 14 for details. All experiments in this chapter are run on a computer with 32 3.7GHz cores and 64 GB of RAM.

Overview of experiments | In this chapter, we report the performance on (1) a UAV motion control problem, and (2) a spacecraft docking problem. In particular, we will consider the following questions about our approach:

1. Is it, in practice, necessary to use IMDP instead of MDP abstractions to obtain safe Markov policies for the DTSS?
2. How does the sample size N used to compute the probability intervals of the IMDP affect the performance of our approach?
3. Can the two-layer feedback design framework presented in Sect. 6.5 be used to construct smaller IMDP abstractions while retaining performance?

To limit the length of this thesis, we selected the experimental results which we believe are most representative of our approach. As a result, several experiments that investigate other questions about our approach are omitted. We refer to [7] for more details on the scalability of our approach and comparisons against two state-of-the-art tools, namely StocHy [CA19] and SReachTools [VGO19]. Similarly, we refer to [10] for more experiments on the two-layer feedback design framework from Sect. 6.5.

6.6.1 UAV motion planning

We consider Problem 6.6 for a reach-avoid problem for a UAV operating under turbulence. The horizon for the reach-avoid problem is $h = 64$ time steps, and the target and unsafe sets are displayed in Fig. 6.10 in green and red, respectively.

DTSS formulation | We use a simplified model for the UAV as a system of 3 double integrators, such that the 6D state vector is $x = [p_x, v_x, p_y, v_y, p_z, v_z]^T \in \mathbb{R}^6$, where p_i and v_i are the position and velocity in the direction i . Control inputs $u_k \in U \subset \mathbb{R}^3$ model

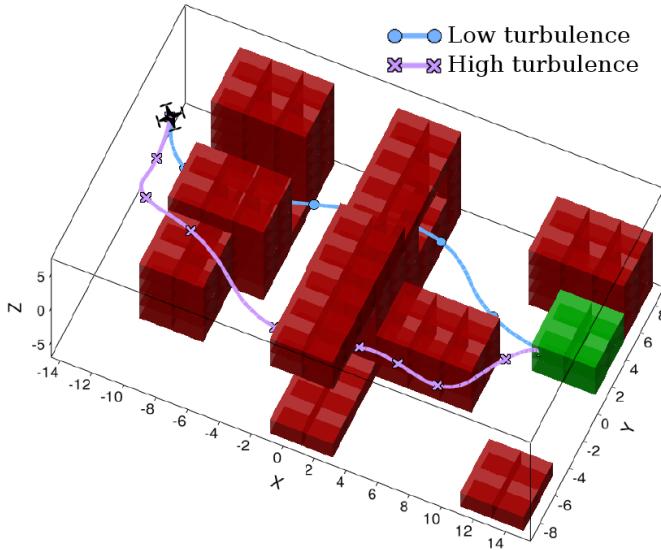


Figure 6.10: UAV reach-avoid problem (target in green; unsafe states in red), plus trajectories under the optimal IMDP-based Markov policy from initial state $x_0 = [-14, 0, 6, 0, -6, 0]^\top$, under high and low turbulence.

actuators that change the velocity, where $U = [-4, 4]^3$. The discrete-time dynamics are then defined as follows:

$$x_{k+1} = \begin{bmatrix} 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix} x_k + \begin{bmatrix} 0.5 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0.5 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0.5 \\ 0 & 0 & 1 \end{bmatrix} u_k + \xi_k, \quad (6.39)$$

where $\xi_k \in \mathbb{R}^6$ is the effect of turbulence on the dynamics. Note that the drift term $q = 0$ and is thus omitted from Eq. (6.39). We group every two discrete time steps together to render the DTSS fully actuated (referring to [7] for details).

The process noise distribution comes from a Dryden gust model [BCMJ19; Dry43]. Importantly, this noise model is non-Gaussian but enables us to cheaply obtain samples, which we use to compute probability intervals for the IMDP abstraction.

Reach-avoid problem | We consider Problem 6.6 with a probability threshold of $\rho = 0.75$. Thus, the problem is to compute a Markov policy μ for the DTSS such that $\Pr_\mu^S(x_I \models \varphi) \geq 0.75$. The set of target states is defined as

$$X_T = \{x \in \mathbb{R}^6 \mid 11 \leq p_x \leq 15, 1 \leq p_y \leq 5, -7 \leq p_z \leq -3\}.$$

For brevity, we omit an explicit definition of the critical regions.

IMDP abstraction | We apply Algorithm 6.1 with $\beta = 0.01$, $\gamma = 2$, $N_0 = 25$, with an upper bound of $N_{\max} = 12\,800$ samples. We use a partition Ψ of the 6-dimensional state space into 25 515 hyperrectangular regions. The resulting reach-avoid⁹ probabilities \check{Pr}^* obtained from Algorithm 6.1 are shown in Fig. 6.11. Observe that the threshold probability of $\rho = 0.75$ is satisfied for $N = 3\,200$ samples and higher.

Validating the satisfaction probability | To validate whether the reach-avoid probabilities \check{Pr}^* for the IMDP are indeed a lower bound of the reach-avoid probability on the DTSS, we perform Monte Carlo simulations. Specifically, for every value of N , we use Theorem 6.29 to obtain a Markov policy μ defined as $\mu_k(x) \in I_R^{\check{\sigma}^*}(x, k)$ contained in the interface for the optimal IMDP scheduler $\check{\sigma}^*$. For each Markov policy, we simulate the DTSS 10 000 times and compute the empirical reach-avoid probability (as the fraction of executions that satisfy the reach-avoid specification), which are shown in Fig. 6.11 by the dashed line. As expected, the IMDP reach-avoid probability \check{Pr}^* (solid line) consistently lower bounds the fraction of satisfying DTSS executions.

Accounting for noise matters for safety | In Fig. 6.10, we show state trajectories under the Markov policy $\mu_k(x) \in I_R^{\check{\sigma}^*}(x, k)$ derived from the optimal IMDP scheduler, under two different strengths (high and low) for the turbulence (i.e., process noise). Under low process noise, the UAV takes the short but narrow path. On the other hand, with the high process noise level, the longer but safer path is preferred since the risk of colliding with an obstacle is too high. Thus, accounting for process noise is important to obtain Markov policies that are safe.

Robust abstractions give safer guarantees | To show the importance of using robust abstractions, we compare our robust IMDP approach against a naïve MDP abstraction. This MDP has the same states and actions as the IMDP but uses *precise transition probabilities*, which are computed using the frequentist approach in Def. 6.18. For the MDP abstraction, we compute (using value iteration; see Sect. 3.2.3) an optimal scheduler σ^* and the corresponding maximal reach-avoid probability $Pr_{\sigma^*}^M$:

$$\sigma^* \in \operatorname{argmax}_{\sigma \in \mathfrak{S}_{\text{Markov}}^M} Pr_{\sigma}^M(s \models \varphi_{\text{ra}}), \quad \text{and} \quad Pr^+ := \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^M} Pr_{\sigma}^M(s \models \varphi_{\text{ra}}).$$

The values of Pr^+ for different values of N are also shown in Fig. 6.11, as well as the fraction of satisfying DTSS executions (in Monte Carlo simulations). Observe that the (non-robust) MDP abstractions yield *invalid satisfaction guarantees*: The actual reach-avoid probability of the Markov policy on the DTSS is much lower than the reach-avoid probability Pr^* on the MDP. The intuitive reason for this unsafe behavior is that the MDP consists of millions of transitions, so even though the estimation error in each individual transition probability is small, the overall error in the satisfaction probability is still significant. By contrast, despite being more conservative, our robust IMDP-based approach consistently yields safe lower-bound guarantees on the actual satisfaction probability observed in simulations.

⁹Since we only consider reach-avoid problems, we use the words *reach-avoid probability* and *satisfaction probability* interchangeably.

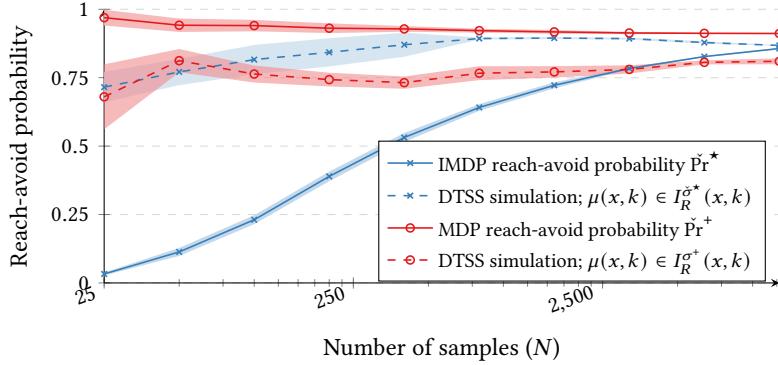


Figure 6.11: The maximal reach-avoid probabilities on the IMDP (blue) and MDP (orange), versus the fraction of satisfying DTSS executions in simulations (under the resulting Markov policies). The shaded areas show the standard deviation across 10 iterations.

6.6.2 Spacecraft docking

In the second experiment, we demonstrate how the two-layer feedback design framework from Sect. 6.5 can be used to construct smaller IMDP abstractions. We consider the spacecraft docking problem from [VGO19], which has a 4D state $x \in \mathbb{R}^4$ modeling the position and velocity in two dimensions, and which evolves as

$$x_{k+1} = \begin{bmatrix} 1.0006 & 0.0000 & 19.9986 & 0.4100 \\ 8.6200 & 1.0000 & -0.4100 & 19.9944 \\ 6.3000 & 0.0000 & 0.9998 & 0.0410 \\ -1.2900 & 0.0000 & -0.0410 & 0.9992 \end{bmatrix} x_k + \begin{bmatrix} 0.33332 & 0.00456 \\ -0.00456 & 0.33329 \\ 0.33331 & 0.00068 \\ -0.00069 & 0.33324 \end{bmatrix} u_k + \zeta_k, \quad (6.40)$$

where the process noise $\zeta_k \sim \mathcal{N}(0, \text{diag}(10^{-4}, 10^{-4}, 5 \times 10^{-8}, 5 \times 10^{-8}))$ has a Gaussian distribution, and $u_k \in U = [-0.1, 0.1]^2 \subset \mathbb{R}^2$. We partition the state space into 3 200 rectangular regions and consider the reach-avoid problem depicted in Fig. 6.12 with a horizon of $h = 8$ steps.

Comparison to SReachTools | SReachTools [VGO19] is an optimization-based toolbox for probabilistic reach-avoid problems. While we use samples to generate a model abstraction, SReachTools employs sample-based methods over the specifications directly. Distinctively, SReachTools does not create abstractions (as in our case) and is thus generally faster than our method. However, its complexity is exponential in the number of samples (versus the linear complexity for our method). Importantly, we derive Markov policies, which use state *feedback*, while the sampling-based methods of SReachTools compute *open-loop* controllers. Open-loop controllers do not consider any feedback, making such controllers unsafe in settings with significant noise levels [AM10]. Markov policies use state feedback and are, therefore, more robust against strong disturbances from noise, as also shown in Fig. 6.12.

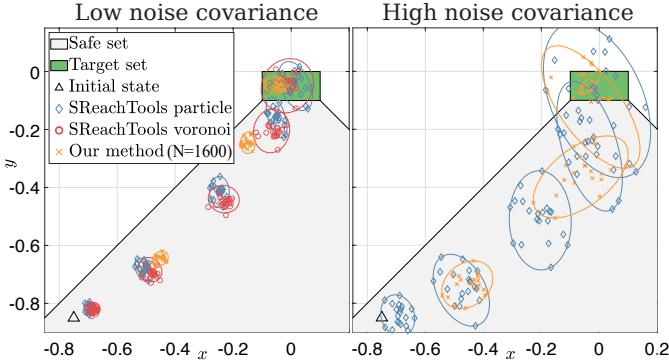


Figure 6.12: Simulated state trajectories for the spacecraft docking problem, under process noise with low and high covariance, respectively. Our Markov policies use state feedback and are thus more robust than open-loop controllers, as shown by the smaller error in the state trajectories over time (the Voronoi method under high covariance failed to generate a solution).

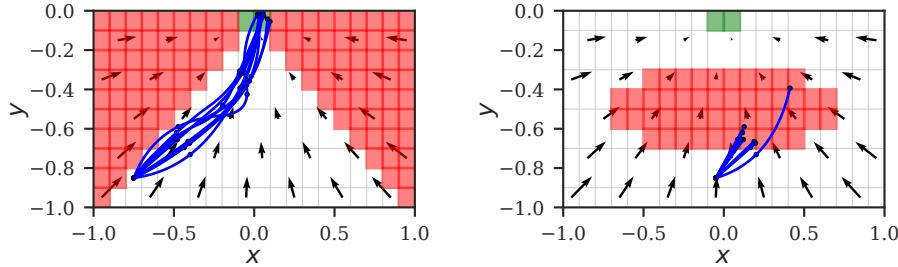
Exploiting stability for smaller abstractions | Suppose that we choose an overall confidence of $1 - \beta \cdot |\text{Act}| \cdot S = 0.99$ in Theorem 6.29. The resulting IMDP has 1.6 million transitions (i.e., the underlying graph has 1.6 million edges) and leads to a reach-avoid probability of $\Pr^* = 0.80$, such that we can guarantee that $\Pr_\mu^S(x_I \models \varphi) \geq 0.80$, where μ is the Markov policy obtained from Theorem 6.29. We can use the two-layer feedback design framework from Sect. 6.5 to obtain a significantly smaller abstraction at the cost of a small reduction in the reach-avoid probability \Pr^* . Specifically, before generating the abstraction, we stabilize the dynamics in Eq. (6.40) with an LQR, such that we obtain closed-loop dynamics of the form

$$x_{k+1} = (A - BK)x_k + Bu_k + \zeta_k,$$

where $K \in \mathbb{R}^{2 \times 4}$ is the feedback gain. Then, we replace the backward reachable set computations in Algorithm 6.1 with the computation in Eq. (6.36), where we fix $U' = [-0.08, 0.08]^2$, which is slightly smaller than the original input space $U = [-0.1, 0.1]^2$. The resulting IMDP abstraction has 280 thousand transitions (a reduction of 79%) and leads to a satisfaction probability of $\Pr^* = 0.79$ (only 0.01 below the baseline).

To explain this significant improvement, we plot the vector field $(A - BK)x$ under the stabilizing controller for this reach-avoid problem in Fig. 6.13a. Observe that the stabilizing controller steers the system toward the target states X_T and away from the unsafe states X_U . Thus, the Markov policy derived from the abstraction only has to deviate slightly from this natural vector field. In such cases, we can thus construct significantly smaller abstractions at negligible performance loss.

Alignment of the stabilizing controller is important | However, if the reach-avoid specification and the stabilizing controller are not well-aligned, performance may drop significantly. For example, in Fig. 6.13b, we consider a different reach-avoid problem. In this case, when we abstract the stabilizing dynamics, the satisfaction probability drops



(a) Specification and stabilization aligned. (b) Specification and stabilization misaligned.

Figure 6.13: Simulated trajectories and stabilized vector fields $(A - BK)x$ for the reach-avoid problems on the spacecraft docking benchmark (only the two position state variables are shown). In case (a), the specification and vector fields are aligned, while this is not the case for case (b).

to $\bar{P}_r = 0.0072$, i.e., almost to zero. This severe performance loss is clearly observed from the simulated trajectories in Fig. 6.13b: The stabilizing controller directly steers the system towards the unsafe states, and the abstraction cannot counteract this action due to the reduced control limits $U' \subset U$.

6.7 Related Work

In Sect. 5.2, we discussed related work on abstraction-based policy synthesis for DTSSs. The majority of these methods rely on full knowledge of the probabilistic models. In other words, the distribution of the stochastic noise must be known. Our approach breaks away from this literature and can be used to generate formal abstractions *without requiring any knowledge of the noise distribution*. We leverage results from the scenario approach [CC05; CG18a] to deal with these unknown noise distributions in a sampling-based fashion.

Backward reachability computations | A defining characteristic of the abstraction method presented in this chapter is that each abstract action is associated with a *fixed distribution* over (continuous) successor states (which we call the *backward method*). By contrast, other abstractions typically associate each abstract action with a *fixed control input*, such that the distribution over successor states depends on the precise continuous state where the abstract action is chosen (which we call the *forward method*). We demonstrate the difference between the backward and forward methods in more detail in Remark 6.11. With our approach, each abstract action leads to a single distribution over successor states that is independent of the precise continuous state where the abstract action is chosen. Doing so comes at the cost of more restrictive assumptions on the dynamics, namely Assumption 6.4.

Other sampling techniques | The controller synthesis tool SReachTools [VGO19] also exhibits a sampling-based method but relies on Hoeffding's inequality to obtain confidence guarantees [SVAO19], so the noise is assumed to be sub-Gaussian [BLM13].

By contrast, the scenario approach is *completely distribution-free* [CG18a]. Moreover, as we will show in this chapter (specifically, in Sect. 6.3.4), the scenario approach may lead to abstract models with significantly better probability intervals compared to Hoeffding's inequality. In addition, SReachTools is limited to problems with convex safe sets (a restrictive assumption in many problems), and its sampling-based methods can only synthesize open-loop controllers, which may undermine the robustness of the overall approach. Another body of relevant literature entails sampling-based feedback motion planning algorithms, e.g., LQR-Trees [TMT10]. However, sampling in LQR-Trees relates to random exploration of the state space and not to stochastic noise affecting the dynamics as in our setting [RPT16]. Finally, Monte Carlo methods (e.g., particle methods) have also been used to solve stochastic reach-avoid problems [BOBW10; LOE13]. These methods simulate the system via many samples of the uncertain variable [Smi13]. Monte Carlo methods *approximate* stochastic problems but do not provide rigorous *bounds* with a desired *confidence level* on the obtained results as our approach does.

Distributionally robust optimization | In distributionally robust optimization (DRO), decisions are robust with respect to *ambiguity sets* of distributions [EK18; GS10; WKS14]. While the scenario approach uses samples of the uncertain variable, DRO works on the domain of uncertainty directly, thus involving potentially complex ambiguity sets [GC22]. Designing robust schedulers for IMDPs with known uncertainty sets was studied by [PLSS13], and [WTM12]. Furthermore, hybrid methods between the scenario approach and robust optimization also exist [MGL14].

PAC literature | The term probably approximately correct (PAC) refers to obtaining, with high probability, a hypothesis that is a good approximation of some unknown phenomenon [HW93]. PAC learning methods for discrete-state MDPs are developed in [BT02], [FT14], and [KS02]. Furthermore, PAC statistical model checking for MDPs is studied by [AKW19; MWW24].

6.8 Discussion

We finish this chapter by discussing the open challenges and limitations of our abstraction-based policy synthesis scheme for linear DTSS.

Tighter probability intervals | While our experiments illustrate that our scenario-based method to compute probability intervals leads to reasonable lower bounds on the satisfaction probability, the tightness of the intervals can still be improved. Recently, [MWW24] pointed out that our probability intervals based on the scenario approach can be improved by using the Clopper-Pearson interval. The Clopper-Pearson interval is a classical statistical method for calculating binomial confidence intervals, and can thus be applied to bound the unknown success probability of a binomial random variable [CP34; New98]: Thus, as pointed out in [MWW24, Theorem 2], these results could readily be incorporated into our framework to construct IMDPs with smaller probability intervals and thus obtain a tighter lower bound on the satisfaction probability.

Robust MDP abstractions | In this chapter, we have focused on IMDP abstractions. However, the theoretical results from Chapter 5 are valid for general RMDPs. An interesting direction for future work is to consider abstractions into RMDPs with convex

polytopic uncertainty sets, which allow for modeling the estimates of transition probabilities with tighter uncertainty sets. By contrast, our IMDP abstraction method essentially estimates each transition probability independently, which leads to more conservative abstractions in general. Such an approach for abstractions into RMDPs has recently been explored by [GBLL24].

Curse of dimensionality | While the ideas from Sect. 6.5 can be used to construct significantly smaller abstractions, the scalability of our approach (and any discretization-based method) remains a general concern. One approach to improve the scalability is to better exploit the structure of the DTSS, e.g., as recently done by X for models represented as mixtures of Gaussian processes. Another interesting avenue for future research is to combine our formal abstractions with a learning scheme. The general idea is to first use learning techniques (e.g., reinforcement learning [SB98]) to efficiently learn a candidate Markov policy without any formal guarantees. We can then use our abstraction-based method to provide hard guarantees on the satisfaction probability of the learned candidate policy. More specifically, our stability-based scheme from Sect. 6.5 may readily be extended from a single stabilizing controller to a piecewise affine controller over the state space. Interestingly, a neural network with ReLU activation functions (which are among the most popular in deep reinforcement learning) represents such a piecewise affine controller, thus posing an interesting potential connection between learning and verification.

Partial observability | In this chapter, we considered DTSSs with full state observability. Intuitively, having full state observability means that Markov policies can use the full state information, i.e., they are defined as a mapping from states to actions. However, in many realistic systems, states are rarely fully observable. For example, a UAV may only have access to noisy GPS measurements of its position, and a self-driving car may only have access to camera images or LIDAR to determine the movement of pedestrians. In such applications, the assumption that the full state is observable is unrealistic, and *partially observable* models are needed instead. However, control problems for partially observable systems are notoriously hard in general, and several problems are proven to be undecidable already for discrete-state/action models, i.e., for partially observable Markov decision processes (POMDPs) [MHC99; CCT16].

One exception, which we investigate in [5], is when both the state transition dynamics and the observation model are linear, and when both are subject to additive Gaussian noise. The abstraction algorithm from this chapter can be extended to such *linear Gaussian systems* relatively straightforwardly, by using a Kalman filter [Kal60; WB01] to estimate the state of the system and then construction an abstraction of the *estimated state* (called the *belief*) instead of the *actual state*. Importantly, the Kalman filter represents the belief of the state as a Gaussian distribution, such that the belief is fully characterized by its mean and covariance matrix. For linear Gaussian systems, the Kalman filter is optimal in the minimum mean-square-error sense, meaning its estimate is the least uncertain of any filter given the same history of information [TBF05; HRW12]. We omit a full discussion of this extension in this thesis, and we refer the interested reader to our paper [5] for more details instead.

Summary

- ⇒ We have presented an algorithm for synthesizing Markov policies for linear DTSSs that provably satisfy reach-avoid specifications.
- ⇒ We drop the unrealistic assumption that the distribution of the stochastic noise of the DTSS is known and instead only assume access to a finite set of noise samples.
- ⇒ We generate a finite abstraction of the DTSS in the form of an IMDP, whose probability intervals are computed using data-driven techniques.
- ⇒ The correctness of our algorithm is agnostic to the noise distribution.
- ⇒ The stability of a DTSS can be exploited to reduce the size of abstractions significantly while retaining the correctness guarantees.

7 DTSSs With Uncertain Parameters

Summary | So far, we have considered discrete-time stochastic systems (DTSSs) whose (probabilistic) transition function is precisely known. As a result, these systems exclusively capture stochastic uncertainty and require that model parameters be known precisely. In this chapter, we study the lower bound control problem (introduced in Chapter 5) for DTSSs with uncertain parameters. These uncertain parameters are assumed to lie in a convex uncertainty set, but we do not assume any probability distribution over these parameters. By sampling techniques and robust analysis, we capture both the stochastic and parameter uncertainty, with a user-specified confidence level, in the transition probability intervals of an interval Markov decision process (IMDP). We show that the probabilistic simulation relations from Chapter 5 extend naturally to DTSSs with uncertain parameters, and that the IMDP abstraction induces such a relation. Our experiments show that our approach leads to Markov policies that are more robust against variations in parameter values.

Origins | The results of this chapter are based on the conference paper:

[6] Badings, Romao, Abate, Jansen (2023) ‘Probabilities Are Not Enough: Formal Controller Synthesis for Stochastic Dynamical Models with Epistemic Uncertainty’. AAAI.

In this paper, we present the first abstraction-based, formal policy synthesis method that simultaneously captures set-bounded parameter uncertainty and stochastic uncertainty for models with continuous state and action spaces.

Background | We assume the reader is familiar with Markov decision processes (MDPs) (Def. 3.1) and IMDPs (Def. 3.27), and with computing optimal schedulers for reach-avoid probabilities. Moreover, to capture stochastic uncertainty in models, we build upon the probability-theoretic definitions from Sect. 2.3.



7.1 Parameter Uncertainty in DTSSs

Recall from Chapter 4 that the state evolution of a discrete-time stochastic system (DTSS) is characterized by its state transition function $f: X \times U \times \mathcal{V}_S \rightarrow X$, or (equivalently) by the stochastic kernel $T: \mathcal{B}(X) \times X \times U \rightarrow [0, 1]$, which assigns to every state-control pair (x_k, u_k) a distribution over successor states x_{k+1} . Thus, the DTSS captures uncertainty about its execution of a *purely stochastic nature*. For example, in Sect. 6.6.1, we modeled the effect of turbulence on a drone as such a (non-Gaussian) stochastic disturbance.

However, what if certain parameters of the DTSS are only known up to a limited precision? For example, the mass of a drone may only be known up to a given interval, e.g., 0.75–1.25 kg. Assuming that we do not have any information about the likelihood of each value for the mass, it is unrealistic to employ a probabilistic model to capture the

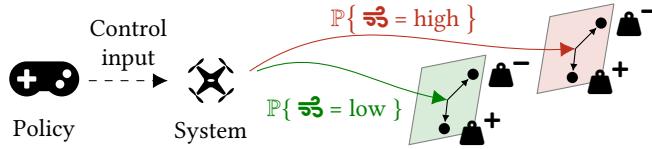


Figure 7.1: Stochastic uncertainty in the wind (\bar{w}) causes probability distributions over the outcomes of controls, while set-bounded uncertainty in the mass (\bar{m}) of the drone causes state transitions to be nondeterministic.

uncertainty [HW21]. Instead, we can only assume that every value for the mass in the interval is a *possibility*, and thus, we obtain a set of possible state transition dynamics. This perspective is illustrated in Fig. 7.1, showing that stochastic uncertainty leads to *distributions over outcomes* of controls, whereas set-bounded parameter uncertainty leads to *sets of possible outcomes*. In other words, parameter uncertainty leads to *nondeterminism* in the outcomes of actions.

Formally, we extend the definition of a DTSS in Def. 4.2 as follows to capture parameter uncertainty. We will refer to such a model as a *robust DTSS (RDTSS)*.

robust
DTSS

Definition 7.1 (RDTSS) A *robust DTSS (RDTSS)* is a tuple $S_R := (X, U, \Gamma, x_I, \zeta, F)$, where

- $X \subseteq \mathbb{R}^n$ is a Borel space, called the state space of the system;
- $U \subseteq \mathbb{R}^m$ is a Borel space, called the (control) input space of the system;
- $\Gamma \subset \mathbb{R}^r$ is a Borel space, called the parameter space of the system;
- $x_I \in X$ is the initial state;
- $\zeta = (\zeta_k)_{k \in \mathbb{N}}$ is a discrete-time stochastic process defined on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$ with its natural filtration $\{\mathcal{F}_k\}_{k \in \mathbb{N}}$ (see Def. 2.10), where each $\zeta_k: \Omega \rightarrow \mathcal{V}_\zeta$ maps to a common measurable space $(\mathcal{V}_\zeta, \mathcal{F}_\zeta)$;
- $F: X \times U \times \mathcal{V}_\zeta \times \Gamma \rightarrow X$ is a parametric state transition function.

Note that the difference between Defs. 4.2 and 7.1 is the parametric transition function F . In particular, fixing a parameter value $\gamma \in \Gamma$ for an robust DTSS (RDTSS) yields a standard DTSS, which we call the *induced DTSS*.

induced
DTSS

Definition 7.2 (Induced DTSS) Fixing a parameter value $\gamma \in \Gamma$ for an RDTSS $S_R = (X, U, \Gamma, x_I, \zeta, F)$ *induces* a standard DTSS, denoted by $S_R[\gamma] = (X, U, x_I, \zeta, f)$, where the transition function f is defined as

$$f(x, u, \bar{\zeta}) = F(x, u, \zeta, \gamma), \quad \forall x \in X, u \in U, \bar{\zeta} \in \mathcal{V}_\zeta.$$

7.1.1 Assumptions

In this chapter, and analogous to Chapter 6, we consider RDTSSs with a linear transition function. Specifically, we assume that the transition function F of the RDTSS is defined

for all $x \in X$, $u \in U$, $\zeta \in \mathcal{V}_\zeta$, and $\gamma \in \Gamma$ as

$$F(x, u, \zeta, \gamma) = A(\gamma)x + B(\gamma)u + q(\gamma) + \zeta, \quad (7.1)$$

where the dynamics matrix $A(\gamma) \in \mathbb{R}^{n \times n}$, control matrix $B(\gamma) \in \mathbb{R}^{n \times m}$, and disturbance $q(\gamma) \in \mathbb{R}^n$ are convex combinations of a finite set of $r \in \mathbb{N}$ known elements:

$$A(\gamma) = \sum_{i=1}^r \gamma_i A_i, \quad B(\gamma) = \sum_{i=1}^r \gamma_i B_i, \quad \text{and} \quad q(\gamma) = \sum_{i=1}^r \gamma_i q_i,$$

where the (unknown) parameter $\gamma \in \Gamma$ can be any point in the unit simplex $\Gamma \subset \mathbb{R}^r$:

$$\Gamma = \left\{ \gamma \in \mathbb{R}^r : \gamma_i \geq 0, \forall i \in \{1, \dots, r\}, \sum_{i=1}^r \gamma_i = 1 \right\}.$$

The model in Eq. (7.1) has *set-bounded uncertain parameters* through $\gamma \in \Gamma$ (against which we want to be robust) and is stochastic due to the process $(\zeta_k)_{k \in \mathbb{N}}$ (which we want to reason over probabilistically). Importantly, we assume that the (unknown) parameter $\gamma \in \Gamma$ is not observable to the Markov policy. Thus, we will aim to synthesize a *single* Markov policy that is robust against all possible values of the parameter $\gamma \in \Gamma$.

Suppose that we are given an RDTSS with a linear transition function. For this RDTSS, we again consider Assumptions 4.4 and 6.2 (the process noise is i.i.d. and has density), Assumption 6.3 (the set U is a convex polytope), and Assumption 6.17 (the noise distribution is unknown). However, as an important difference to the previous chapter, instead of imposing Assumption 6.4, we only make the following assumption.

Assumption 7.3 (Non-singular dynamics) The matrix $A(\hat{\gamma})$ in Eq. (7.1) is non-singular for some $\hat{\gamma} \in \Gamma$.

Thus, compared to Chapter 6, we drop the strong assumption that the pair (A, B) is controllable. Assumption 7.3 is milder as it only requires the existence of a non-singular matrix in the convex hull $\text{conv}\{A_1, \dots, A_r\}$ of the dynamics matrix.

Modeling uncertain parameters | The model defined in Eq. (7.1) can be used to capture parameters known up to an interval, as illustrated by the following example.

Example 7.4 Consider again the drone depicted in Fig. 7.1. The drone's longitudinal position p_k and velocity v_k are modeled as

$$x_{k+1} = \begin{bmatrix} p_{k+1} \\ v_{k+1} \end{bmatrix} = \begin{bmatrix} 1 & \delta_t \\ 0 & \frac{m-0.1\delta_t}{m} \end{bmatrix} x_k + \begin{bmatrix} \frac{\delta_t^2}{2m} \\ \frac{\delta_t}{m} \end{bmatrix} u_k + \zeta_k,$$

with δ_t the discretization time, and $U = [-5, 5]$. Assume that the mass m is only known to lie within $[0.75, 1.25]$. Then, we obtain an RDTSS with a linear transition function as in Eq. (7.1), with $r = 2$ vertices where A_1, B_1 are obtained for $m := 0.75$, and A_2, B_2 for $m := 1.25$ (and $q_1 = q_2 = 0$).

set-bounded parameter uncertainty

7.1.2 Problem statement

In this chapter, we consider a robust variant of the lower bound control problem (see Chapter 5) tailored to RDTSSs with a linear transition function as in Eq. (7.1). Intuitively, our goal is to synthesize a Markov policy that (1) is *robust against nondeterminism* due to parameter uncertainty and (2) *reasons probabilistically* over the stochastic noise. Using the terminology from [WKR13], our goal is to be *risk-averse* against the uncertain parameters, while being *risk-neutral* against the stochastic noise. More specifically, we aim to find a Markov policy for which the reach-avoid probability is above a certain threshold $\rho \in [0, 1]$ for every possible value of the unknown parameter $\gamma \in \Gamma$.

Recall from Def. 7.2 that fixing a parameter value $\gamma \in \Gamma$ for an RDTSS induces a standard DTSS denoted by $S_R[\gamma]$. Furthermore, recall from Def. 4.11 that $\Pr_{\mu}^{S_R[\gamma]}(x_I \models \varphi)$ is the probability of satisfying a reach-avoid specification φ under a Markov policy μ , starting from the initial state x_I . We then consider the following problem.

Problem 7.5 (Lower bound control for linear RDTSS) Given an RDTSS $S_R = (X, U, \Gamma, x_I, \varsigma, F)$ with a linear transition function as in Eq. (7.1), a reach-avoid specification $\varphi = (X_T, X_U, h)$, and a desired threshold probability $\rho \in [0, 1]$, compute a Markov policy μ such that

$$\min_{\gamma \in \Gamma} \Pr_{\mu}^{S_R[\gamma]}(x_I \models \varphi) \geq \rho,$$

or return False if no such policy could be found.

Thus, we want to find a *single* Markov policy μ with a lower bound on the satisfaction probability that is robust against *all possible values* of the (unknown) parameter $\gamma \in \Gamma$. In doing so, we assume that the unknown parameter $\gamma \in \Gamma$ is *unobservable* to the Markov policy. As in Chapter 6, we will solve Problem 7.5 up to a user-specified confidence probability due to the noise of unknown distribution.

7.1.3 Overview of our abstraction technique

The main contribution that allows us to be robust against parameter uncertainty, is that we reason over *sets* of potential transitions (as shown by the boxes in Fig. 7.1), rather than *precise* transitions as in Chapter 6. Intuitively, for a given action, the stochastic uncertainty creates a probability distribution over *sets of possible outcomes*. To ensure robustness against parameter uncertainty, we consider *all possible outcomes* within these sets. We show that, for RDTSSs with linear transition functions, computing these sets of all possible outcomes is computationally tractable. Crucially, the correctness guarantees from Chapters 5 and 6 carry over to the setting in this chapter.

Outline | This chapter is organized as follows. In Sect. 7.2, we present our method for abstracting an RDTSS into a finite interval Markov decision process (IMDP) that soundly captures both the stochastic and parameter uncertainty. Then, in Sect. 7.3, we present a concrete algorithm to solve the RDTSS policy synthesis problem in Problem 7.5. In Sect. 7.4, we present numerical experiments that demonstrate the effectiveness of our approach. Finally, we discuss related work in Sect. 7.5, and we discuss the open challenges and promising directions for future research in Sect. 7.6.

7.2 Parameter Robustness in IMDP Abstractions

Similar to Chapter 6, we solve Problem 7.5 using an IMDP abstraction. However, adding the robustness against any parameter value $\gamma \in \Gamma$ means that the abstraction algorithm changes significantly. In particular, we will introduce a so-called *nominal model* that neglects both the stochastic noise and parameter uncertainty. We will then define the actions of this IMDP via backward reachability computations on this nominal model. We compensate for the error caused by this modeling simplification in the IMDP's transition probability intervals.

7.2.1 Nominal dynamics model

To build our abstraction, we rely on a *nominal model* that neglects both the stochastic noise and parameter uncertainty in Eq. (7.1), and is thus deterministic. Concretely, we choose any value $\hat{\gamma} \in \Gamma$ (called the *nominal parameter value*) for which Assumption 7.3 holds and define the nominal model dynamics as

$$\hat{x}_{k+1} = A(\hat{\gamma})x_k + B(\hat{\gamma})u_k + q(\hat{\gamma}). \quad (7.2)$$

Due to the linearity of the dynamics, we can now express the successor state x_{k+1} *with full uncertainty*, from Eq. (7.1), as

$$x_{k+1} = \hat{x}_{k+1} + \delta(\gamma, x_k, u_k) + \zeta_k, \quad (7.3)$$

with $\delta(\gamma, x_k, u_k)$ being a new term, called the *epistemic error*, encompassing the error caused by parameter uncertainty:

$$\delta(\gamma, x_k, u_k) = [A(\gamma) - A(\hat{\gamma})]x_k + [B(\gamma) - B(\hat{\gamma})]u_k + [q(\gamma) - q(\hat{\gamma})]. \quad (7.4)$$

In other words, the successor state x_{k+1} is the nominal one, plus the epistemic error, and plus the stochastic noise. Note that for $\gamma = \hat{\gamma}$ (i.e., the true model parameters equal their nominal values), we obtain $\delta(\gamma, x_k, u_k) = 0$.

7.2.2 IMDP abstraction of the nominal model

For the nominal model in Eq. (7.2), we generate a finite abstraction in a similar vein as in Chapter 6. However, to capture the parameter uncertainty in the abstraction, we immediately obtain an IMDP (whereas we initially obtained an MDP in Chapter 6).

States | The states of the IMDP abstraction are defined equivalently to Chapter 6. That is, the states of the IMDP abstraction $S := \Psi = \{\mathcal{V}_1, \dots, \mathcal{V}_L, \mathbb{R}^n \setminus \mathcal{Z}\}$ are the elements of a polyhedral partition Ψ (see Def. 6.7). Recall from Remark 6.10 that $R(x) \in S$ denotes the (unique) IMDP state containing $x \in \mathbb{R}^n$, whereas $R^{-1}(s) \subset \mathbb{R}^n$ denotes the set of DTSS states related to $s \in S$.

Actions | We define the IMDP actions via backward reachability computations on the nominal model in Eq. (7.2). Let $\mathcal{T} = \{\mathcal{T}_1, \dots, \mathcal{T}_q\}$ be a finite collection of $q \in \mathbb{N}$ *target sets*, each of which is a convex polytope, $\mathcal{T}_\ell = \text{conv}(t_\ell^1, \dots, t_\ell^d) \subset \mathbb{R}^n$, consisting of $d \in \mathbb{N}$ vertices $t_\ell^1, \dots, t_\ell^d \in \mathbb{R}^n$. Every target set corresponds to an IMDP action, yielding the set $Act = \{a_\ell : \ell = 1, \dots, q\}$ of actions. We denote the target set for action $a \in Act$ by \mathcal{T}_a .

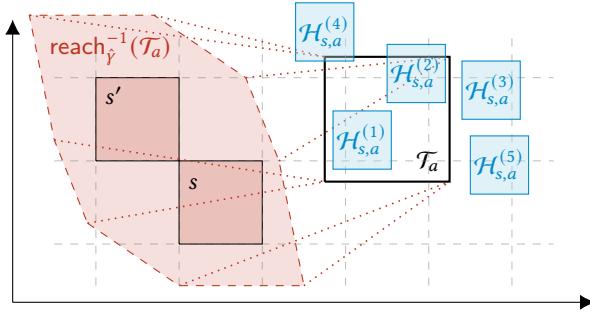


Figure 7.2: Action $a \in \text{Act}$ to target set \mathcal{T}_a is enabled in states $s, s' \in S$, as their regions are contained in the backward reachable set $\text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a)$. The successor state sets $\mathcal{H}_{s,a}^{(1)}, \dots, \mathcal{H}_{s,a}^{(5)}$ for five noise samples (see Sect. 7.2.3), overapproximated as boxes, are shown in blue.

Intuitively, action $a \in \text{Act}$ represents a transition to $\hat{x}_{k+1} \in \mathcal{T}_a$ feasible under the backward reachable set

$$\text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a) = \{x \in \mathbb{R}^n : \exists u \in U, A(\hat{y})x + B(\hat{y})u + q(\hat{y}) \in \mathcal{T}_a\}.$$

As in Chapter 6, we enable action $a \in \text{Act}$ in state $s \in S$ if and only if $\hat{x}_{k+1} \in \mathcal{T}_a$ can be realized from any $x_k \in R^{-1}(s)$ in the associated region, or in other words, the partition element $R^{-1}(s)$ must be contained in the backward reachable set. As such, the set $\text{Act}_{\hat{y}}(s)$ of enabled actions in state $s \in S$ under $\hat{y} \in \Gamma$ is defined as

$$\text{Act}_{\hat{y}}(s) = \{a \in \text{Act} : R^{-1}(s) \subseteq \text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a)\}. \quad (7.5)$$

Remark 7.6 (Target points vs. target sets) Observe the differences between Figs. 6.3 and 7.2. Whereas we used a *precise* target point $d_a \in \mathbb{R}^n$ in Chapter 6, we here use target *sets* to accommodate the parameter uncertainty and the milder Assumption 7.3. Moreover, observe that the backward reachable set shown in Fig. 7.2 has more vertices, which is generally also the case. Finally, as we will discuss later in Sect. 7.2.3, Fig. 7.2 shows that each sample of the process noise leads to a set of possible successor states, rather than the precise successor state samples we obtained in Chapter 6.

To compute the set of enabled actions using Eq. (7.5), we must compute $\text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a)$ for each action $a \in \text{Act}$ with associated target set \mathcal{T}_a . The following lemma shows that $\text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a)$ is a polytope characterized by the vertices of U and \mathcal{T}_a , which is computationally tractable to compute. The intuition is that \mathcal{T}_a and U are convex polytopes, so the inverse of the dynamics of the nominal model in Eq. (7.2) is also a convex polytope, which is exactly characterized by the vertices of U and \mathcal{T}_a .

Lemma 7.7 (Backward reachable set as convex polytope) Let the control space $U = \text{conv}(u^1, \dots, u^q)$, $q \in \mathbb{N}$, and the target set $\mathcal{T}_a = \text{conv}(t^1, \dots, t^d)$, $d \in \mathbb{N}$, of action $a \in \text{Act}$ be given in their vertex representations. Under Assumption 7.3, we have that

$$\text{reach}_{\hat{\gamma}}^{-1}(\mathcal{T}_a) = \text{conv}(\bar{x}_{ij} : i = 1, \dots, d, j = 1, \dots, q), \quad (7.6)$$

where \bar{x}_{ij} is the unique solution of the linear equation system

$$A(\hat{\gamma})\bar{x}_{ij} + B(\hat{\gamma})u^j + q(\hat{\gamma}) = t^i. \quad (7.7)$$

Proof. We first prove that Eq. (7.6) holds with inclusion, i.e.,

$$\text{conv}(\bar{x}_{ij} : i = 1, \dots, d, j = 1, \dots, q) \subseteq \text{reach}_{\hat{\gamma}}^{-1}(\mathcal{T}_a). \quad (7.8)$$

Let z be any element belonging to the right-hand side of Eq. (7.6), i.e. $z \in \text{conv}(\bar{x}_{ij} : i = 1, \dots, d, j = 1, \dots, q)$. Then, there exists θ_{ij} , $i = 1, \dots, d$, $j = 1, \dots, q$, such that

$$\theta_{ij} \geq 0, \quad \sum_{i,j=1}^{d,q} \theta_{ij} = 1, \quad z = \sum_{i,j=1}^{d,q} \theta_{ij} \bar{x}_{ij}.$$

For each vertex u^j of U , $j = 1, \dots, q$, let $\xi_j = \sum_{i=1}^d \theta_{ij}$ and write the control input u corresponding to point z :

$$u = \sum_{j=1}^q \xi_j u^j, \quad u \in U,$$

which is admissible by construction. Now, note that the mapping of the pair (z, u) under the dynamics satisfies

$$\begin{aligned} A(\hat{\gamma})z + B(\hat{\gamma})u + q(\hat{\gamma}) &= \sum_{i,j=1}^{d,q} \theta_{ij} A(\hat{\gamma}) \bar{x}_{ij} + \sum_{j=1}^q \xi_j B(\hat{\gamma}) u^j + q(\hat{\gamma}) \\ &= \sum_{i,j=1}^{d,q} \theta_{ij} (A(\hat{\gamma}) \bar{x}_{ij} + B(\hat{\gamma}) u^j + q(\hat{\gamma})) \\ &= \sum_{i=1}^d \bar{\theta}_i t^i \in \mathcal{T}_a, \end{aligned}$$

where the first equality follows from the definition of z and u , the second by the definition of $\xi_j = \sum_{i=1}^d \theta_{ij}$, and the third by letting $\bar{\theta}_i = \sum_{j=1}^q \theta_{ij}$ and noting that $\sum_{i=1}^d \bar{\theta}_i = 1$ and $\bar{\theta}_i \geq 0$ for all $i = 1, \dots, d$. In other words, the mapping of the pair (z, u) belongs to the target set \mathcal{T}_a , which implies that Eq. (7.8) holds by construction. This concludes the first part of the lemma.

To show the opposite direction in Eq. (7.6), let z be any element in $\text{reach}_{\hat{\gamma}}^{-1}(\mathcal{T}_a)$. By the definition of $\text{reach}_{\hat{\gamma}}^{-1}(\mathcal{T}_a)$, this means that there exist $\xi_j \geq 0$ and $\theta_i \geq 0$,

$i = 1, \dots, d$, $j = 1, \dots, q$, with $\sum_{j=1}^q \xi_j = 1$ and $\sum_{i=1}^d \theta_i = 1$, such that

$$A(\hat{y})z + B(\hat{y}) \left(\sum_{j=1}^q \xi_j u^j \right) + q(\hat{y}) = \sum_{i=1}^d \theta_i t^i. \quad (7.9)$$

In other words, if $z \in \text{reach}_{\hat{y}}^{-1}(\mathcal{T}_a)$ then there exists an input $u \in U$ such that $A(\hat{y})z + B(\hat{y})u + q(\hat{y}) \in \mathcal{T}_a$. By substituting (7.7) into (7.9), we obtain

$$\begin{aligned} A(\hat{y})z + B(\hat{y}) \sum_{j=1}^q \xi_j u^j + q(\hat{y}) &= \sum_{i=1}^d \theta_i \left(A(\hat{y})\bar{x}_{ik} + B(\hat{y})u^k + q(\hat{y}) \right) \\ A(\hat{y})z &= \sum_{i=1}^d \theta_i \left(A(\hat{y})\bar{x}_{ik} + B(\hat{y})u^k \right) - B(\hat{y}) \sum_{j=1}^q \xi_j u^j, \end{aligned} \quad (7.10)$$

for all $k = 1, \dots, q$. Multiplying both sides of (7.10) by $A(\hat{y})^{-1}$, which is allowed due to Assumption 7.3, yields

$$z = \sum_{i=1}^d \theta_i (\bar{x}_{ik} + A(\hat{y})^{-1}B(\hat{y})u^k) - A(\hat{y})^{-1}B(\hat{y}) \sum_{j=1}^q \xi_j u^j. \quad (7.11)$$

Since Eq. (7.11) holds for all $k = 1, \dots, q$, we can multiply both sides by ξ_k for each $k = 1, \dots, q$, and sum up the resulting expression. As $\sum_{k=1}^q \xi_k = 1$, we obtain

$$z = \sum_{i,k=1}^{d,q} \bar{\theta}_{ik} \bar{x}_{ik} + A(\hat{y})^{-1}B(\hat{y}) \sum_{k=1}^q \xi_k u^k - A(\hat{y})^{-1}B(\hat{y}) \sum_{j=1}^q \xi_j u^j, \quad (7.12)$$

where $\bar{\theta}_{ik} = \xi_k \theta_i$, which is larger than or equal to zero for all $i = 1, \dots, d$, $k = 1, \dots, q$. Since the last two terms on the right-hand side of (7.12) cancel out and $\sum_{i,k=1}^{d,q} \bar{\theta}_{ik} = 1$, we conclude that $z \in \text{conv}(\bar{x}_{ij} : i = 1, \dots, d, j = 1, \dots, q)$, thus proving the opposite inclusion and concluding the proof of the lemma. ■

Transition probabilities | So far, we have established that taking the abstract action $a \in \text{Act}(s)$ enabled in state $s \in S$ corresponds to transitioning to a continuous successor state that satisfies $\hat{x}_{k+1} = A(\hat{y})x_k + B(\hat{y})u_k + q(\hat{y}) \in \mathcal{T}_a$. Recall from Eq. (7.3) that the actual continuous successor state is $x_{k+1} = \hat{x}_{k+1} + \delta(y, x_k, u_k) + \varsigma_k$, which is, for all $\hat{x}_{k+1} \in \mathcal{T}_a$ and $y \in \Gamma$, a random variable. In other words, for a given action $a \in \text{Act}$, fixing a value for $\hat{x}_{k+1} \in \mathcal{T}_a$ and $y \in \Gamma$ leads to a distribution over continuous successor states in \mathbb{R}^n . By integrating this distribution to the set $R^{-1}(s') \subset \mathbb{R}^n$ related to each abstract state $s' \in S$, we obtain the transition probabilities of our abstraction. However, the precise values of $\hat{x}_{k+1} \in \mathcal{T}_a$ and $y \in \Gamma$ are not determined, so we instead define a set of possible transition probabilities $\mathcal{P}(s, a)(s')$ for all $s, s' \in S$ and $a \in \text{Act}(s)$:

$$\mathcal{P}(s, a)(s') = \left\{ \mathbb{P}\{\omega \in \Omega : \hat{x}_{k+1} + \delta(y, x_k, u_k) + \varsigma_k \in R^{-1}(s')\} : \hat{x}_{k+1} \in \mathcal{T}_a, y \in \Gamma \right\}.$$

Two factors prevent us from computing this set of probabilities $\mathcal{P}(s, a)(s')$:

1. the nominal successor state \hat{x}_{k+1} and the epistemic error $\delta(\gamma, x_k, u_k)$ are non-deterministic, and
2. the distribution of the noise ς_k is unknown.

We deal with the nondeterminism of the nominal successor state in the following paragraph while addressing the stochastic noise in Sect. 7.2.3.

Capturing nondeterminism | We capture the nondeterminism in the nominal successor state \hat{x}_{k+1} and the epistemic error $\delta(\gamma, x_k, u_k)$ in a single set that we reason robustly against. First, we define the set Δ_s of all possible epistemic errors for any $x_k \in R^{-1}(s)$ in Eq. (7.4) as

$$\Delta_s = \{\delta(\gamma, x_k, u_k) : \gamma \in \Gamma, x_k \in R^{-1}(s), u_k \in U\}.$$

Then, the successor state x_{k+1} (upon choosing action $a \in \text{Act}$ in state $x_k \in R^{-1}(s)$, $s \in S$) is an element of a set that we denote by $\mathcal{H}_{s,a}$:

$$x_{k+1} \in \mathcal{T}_a + \Delta_s + \varsigma_k = \mathcal{H}_{s,a}. \quad (7.13)$$

Observe that $\mathcal{H}_{s,a}$ is a random variable defined by ς_k , such that for all $\omega \in \Omega$, we have $\mathcal{T}_a + \Delta_s + \varsigma_k(\omega) = \mathcal{H}_{s,a}(\omega)$. Based on the set $\mathcal{H}_{s,a}$, we have for all $p \in \mathcal{P}(s, a)(s')$ that

$$\mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \subseteq R^{-1}(s')\} \leq p \leq \mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \cap R^{-1}(s') \neq \emptyset\}. \quad (7.14)$$

In other words, Eq. (7.14) defines a *lower and upper bound* on the set of transition probabilities $\mathcal{P}(s, a)(s')$ for all $s, s' \in S$ and $a \in \text{Act}(s)$. The lower bound follows from Eq. (7.13), since if $\mathcal{H}_{s,a} \subseteq R^{-1}(s')$, then $x_{k+1} \in R^{-1}(s')$ for any $x_{k+1} \in \mathcal{H}_{s,a}$. The upper bound holds, since by Eq. (7.13) we have that $x_{k+1} \in \mathcal{H}_{s,a}$, and thus, if $x_{k+1} \in R^{-1}(s')$, then the intersection $\mathcal{H}_{s,a} \cap R^{-1}(s')$ must be nonempty.

Before describing our sampling-based method to bound Eq. (7.14), we show with the following lemma that Δ_s is a subset of a convex polytope, which is characterized by the region $R^{-1}(s)$, the feasible control space U , and the model dynamics.

Lemma 7.8 (Representation of the epistemic error) Given the vertex representations of sets $R^{-1}(s) = \text{conv}(v^1, \dots, v^p)$ and $U = \text{conv}(u^1, \dots, u^q)$ for $p, q \in \mathbb{N}$, it holds that

$$\begin{aligned} \Delta_s \subseteq \text{conv}\Big(& (A_\iota - A(\hat{y}))v^j + (B_\iota - B(\hat{y}))u^\ell + (q_\iota - q(\hat{y})) \\ & : \iota = 1, \dots, r, \ j = 1, \dots, p, \ \ell = 1, \dots, q \Big), \end{aligned} \quad (7.15)$$

where $A_1, \dots, A_r, B_1, \dots, B_r$, and q_1, \dots, q_r are as defined in Sect. 7.1.1.

Proof. First, let us fix any $\gamma \in \Gamma$, and observe that Δ_s evaluated at γ is written as

$$\begin{aligned}\Delta_s(\gamma) &= \left\{ \delta(\gamma, x_k, u_k) : x_k \in R^{-1}(s), u_k \in U \right\} \\ &= \left\{ (A(\gamma) - A(\hat{\gamma}))x_k + (B(\gamma) - B(\hat{\gamma}))u_k + q(\gamma) - q(\hat{\gamma}) \right. \\ &\quad \left. : x_k \in R^{-1}(s), u_k \in U \right\}.\end{aligned}\tag{7.16}$$

We observe that the sets $\{(A(\gamma) - A(\hat{\gamma}))x_k : x_k \in R^{-1}(s)\}$ and $\{(B(\gamma) - B(\hat{\gamma}))u_k : u_k \in U\}$ are both convex polytopes characterized by the vertices of $R^{-1}(s)$ and U , respectively. Thus, we rewrite Eq. (7.16) as

$$\Delta_s(\gamma) = \text{conv} \left((A(\gamma) - A(\hat{\gamma}))v^j + (B(\gamma) - B(\hat{\gamma}))u^\ell + q(\gamma) - q(\hat{\gamma}) \right. \\ \left. : j = 1, \dots, p, \ell = 1, \dots, q \right).$$

Note that the full set Δ_s is the union of $\Delta_s(\gamma)$ over all $\gamma \in \Gamma$:

$$\Delta_s = \bigcup_{\gamma \in \Gamma} \Delta_s(\gamma) \subseteq \text{conv} \left((A(\gamma) - A(\hat{\gamma}))v^j + (B(\gamma) - B(\hat{\gamma}))u^\ell + q(\gamma) - q(\hat{\gamma}) \right. \\ \left. : j = 1, \dots, p, \ell = 1, \dots, q, \gamma \in \Gamma \right).\tag{7.17}$$

Crucially, observe that for any fixed pair of vertices $\bar{v} := v^j$ and $\bar{u} := u^\ell$, where $j = 1, \dots, p$ and $\ell = 1, \dots, q$, we can write the convex hull in Eq. (7.17) in terms of only the matrices A_ι, B_ι for $\iota = 1, \dots, r$ of which $A(\gamma)$ and $B(\gamma)$ are a convex combination (as in Sect. 7.1.1):

$$\begin{aligned}&\text{conv}((A(\gamma) - A(\hat{\gamma}))\bar{v} + (B(\gamma) - B(\hat{\gamma}))\bar{u} + q(\gamma) - q(\hat{\gamma}) : \gamma \in \Gamma) \\ &= \text{conv}((A(\gamma) - A(\hat{\gamma}))\bar{v} + (B(\gamma) - B(\hat{\gamma}))\bar{u} + q(\gamma) - q(\hat{\gamma}) : \gamma = e_1, \dots, e_r) \\ &= \text{conv}((A_\iota - A(\hat{\gamma}))\bar{v} + (B_\iota - B(\hat{\gamma}))\bar{u} + q(\gamma) - q(\hat{\gamma}) : \iota = 1, \dots, r),\end{aligned}$$

where $e_i \in \mathbb{R}^r$ is the vector with all components equal to 0, except the i^{th} , which is 1. The last equality holds, since $A(e_\iota) = A_\iota$, $B(e_\iota) = B_\iota$, and $q(e_\iota) = q_\iota$ for any $\iota = 1, \dots, r$. In other words, considering the values $\gamma \in \Gamma \setminus \{e_1, \dots, e_r\}$ in Eq. (7.17) is redundant since these values can be expressed as a convex combination of $\gamma \in \{e_1, \dots, e_r\}$. As a result, we simplify Eq. (7.17) as

$$\Delta_s \subseteq \text{conv} \left((A_\iota - A(\hat{\gamma}))v^j + (B_\iota - B(\hat{\gamma}))u^\ell + (q_\iota - q(\hat{\gamma})) \right. \\ \left. : \iota = 1, \dots, r, j = 1, \dots, p, \ell = 1, \dots, q \right),$$

which equals Eq. (7.15). This concludes the proof of Lemma 7.8. ■

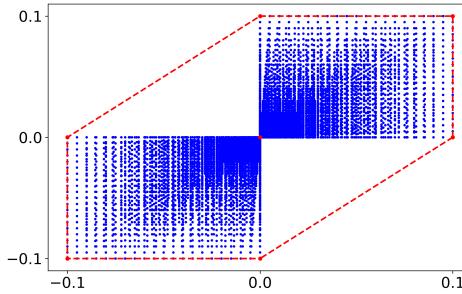


Figure 7.3: Overapproximation of Δ_s obtained from Lemma 7.8, shown as the red hull, versus sampled (blue) points in Δ_s for different $x_k \in R^{-1}(s)$ and $u_k \in U$.

Remark 7.9 (Eq. (7.15) does not hold with equality) It may be tempting to conclude that Eq. (7.15) holds with equality. However, we show with a simple example that this is not the case. Specifically, we apply Lemma 7.8 to a model as in Eq. (7.1) with the matrices

$$A_1 = \begin{bmatrix} 0.9 & 1 \\ 0 & 0.9 \end{bmatrix}, \quad A_2 = \begin{bmatrix} 1.1 & 1 \\ 0 & 1.1 \end{bmatrix}, \quad A(\hat{\gamma}) = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix},$$

$$B_1 = B_2 = B(\hat{\gamma}) = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \quad R^{-1}(s) = [0, 1]^2, \quad U = [-5, 5],$$

and with $q_1 = q_2 = 0$. The resulting right-hand side of Eq. (7.15) is shown by the dashed hull in Fig. 7.3. Moreover, to approximate Δ_s , we compute $\delta(\gamma, x_k, u_k)$ for many linearly spaced points $\gamma \in \Gamma$, $x_k \in R^{-1}(s)$, and $u_k \in U$, which are shown by the blue dots in Fig. 7.3. While the convex hull is a sound overapproximation of the set Δ_s , *the opposite is clearly not the case* (there are points in the convex hull that are not included in Δ_s). This result empirically shows that Eq. (7.15) in Lemma 7.8 does not hold with equality.

7.2.3 PAC probability intervals

The probability interval in Eq. (7.14) depends on the noise ζ_k , whose density function is unknown. We show how to compute PAC bounds on this interval by sampling a set of $N \in \mathbb{N}$ samples of the noise, denoted by $\zeta_k^{(1)}, \dots, \zeta_k^{(N)}$. As in Chapter 6, this set of N noise samples is an element from the probability space Ω^N equipped with the product probability \mathbb{P}^N and the product σ -algebra. Observe from Eq. (7.13) that each sample $\zeta_k^{(i)}$, $i = 1, \dots, N$, yields a set $\mathcal{H}_{s,a}^{(i)}$ (as shown in Fig. 7.2) that contains the successor state under that value of the noise, i.e.,

$$x_{k+1}^{(i)} \in \mathcal{T}_a + \Delta_s + \zeta_k^{(i)} = \mathcal{H}_{s,a}^{(i)}.$$

For reasons of computational performance, we *overapproximate* each set $\mathcal{H}_{s,a}^{(i)}$ as the smallest *hyperrectangle* in \mathbb{R}^n , by taking the point-wise minimum and maximum over the vertices of $\mathcal{H}_{s,a}^{(i)}$.

7.2.3.1 Lower bounds from the scenario approach

scenario approach We interpret the lower bound in Eq. (7.14) within the *sampling-and-discarding scenario approach* [CG11]. This interpretation is very similar to that in Sect. 6.3; however, each sample is a *set* rather than a *point* in the state space \mathbb{R}^n .

Analogous to the scenario optimization program presented in Sect. 6.3, let $Q \subseteq \{1, \dots, N\}$ be a subset of the noise samples that will be discarded from the scenario optimization program. Then, for a fixed triple $(s, a, s') \in S \times Act \times S$, consider the following convex program:

$$\begin{aligned} \mathfrak{L}_{s,a,s'}^Q : & \underset{\lambda \geq 0}{\text{minimize}} \quad \lambda \\ & \text{subject to } \mathcal{H}_{s,a}^{(i)} \subseteq R^{-1}(s')(\lambda) \quad \forall i \in \{1, \dots, N\} \setminus Q, \end{aligned} \tag{7.18}$$

where $R^{-1}(s')(\lambda)$ is a version of $R^{-1}(s')$ scaled by a factor λ around a Chebyshev center of $R^{-1}(s')$ (as defined in Sect. 6.3). The optimal point λ^* to $\mathfrak{L}_{s,a,s'}^Q$ defines a region $R^{-1}(s')(\lambda^*)$ such that, for all $i \in \{1, \dots, N\} \setminus Q$, the set $\mathcal{H}_{s,a}^{(i)}$ for sample $\zeta_k^{(i)}$ is contained in $R^{-1}(s')(\lambda^*)$.

We again use Def. 6.25 to construct a sequence of strictly increasing subsets $Q_0, Q_1, \dots, Q_N \subseteq \{1, \dots, N\}$ of discarded samples (i.e., such that $Q_0 \subset Q_1 \subset \dots \subset Q_N$), where at each step, we add the unique active constraint from the previous solution. Let us denote λ_ℓ^* as the optimal solution to problem $\mathfrak{L}_{s,a,s'}^Q$. As a result, it follows that $\lambda_0^* < \lambda_1^* < \dots < \lambda_N^*$ with probability one.

For every transition (s, a, s') , let us define $N_{s,a,s'}^{\text{out}} \leq N$ as the number of samples $\mathcal{H}_{s,a}^{(i)}$ that are *not* fully contained in $R^{-1}(s')$:

$$N_{s,a,s'}^{\text{out}} := |\{i \in \{1, \dots, N\} : \mathcal{H}_{s,a}^{(i)} \not\subseteq R^{-1}(s')\}|.$$

We adapt Theorem 6.19 to lower bound the probability that an $\omega \in \Omega$ drawn according to \mathbb{P} yields $\mathcal{H}_{s,a}(\omega) \subseteq R^{-1}(s')(\lambda^*)$. This leads to the following lower bound on the transition probability.¹

Lemma 7.10 (Lower bound probability) Let $\zeta_k^{(1)}, \dots, \zeta_k^{(N)}$ be a set of $N \in \mathbb{N}$ samples of the noise ζ_k , and let $\beta \in (0, 1)$ be a confidence parameter. For fixed $s, s' \in S$ and $a \in Act(s)$, define $\mathcal{H}_{s,a}^{(i)} = \mathcal{T}_a + \Delta_s + \zeta_k^{(i)}$ and determine the value of $N_{s,a,s'}^{\text{out}} \leq N$. Then, it holds that

$$\mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : \mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \subseteq R^{-1}(s')\} \geq \check{p} \right\} \geq 1 - \frac{\beta}{2},$$

where $\check{p} = 0$ if $N_{s,a,s'}^{\text{out}} = 0$, and otherwise, \check{p} is the solution to

$$\frac{\beta}{2N} = \sum_{i=0}^{N_{s,a,s'}^{\text{out}}} \binom{N}{i} (1 - \check{p})^i \check{p}^{N-i}.$$

¹Analogous to Sect. 6.3, one can readily show that the technical requirements needed to apply the results from [RPM23] are satisfied for the scenario program Eq. (7.18).

Proof. The proof follows analogous to that of Theorem 6.19, adapted to the scenario optimization program in Eq. (7.18). Recall that this program considers a constraint $\mathcal{H}_{s,a}^{(i)} \subseteq R^{-1}(s')(\lambda)$ for all $i \in \{1, \dots, N\} \setminus Q$, whereas the program in Eq. (6.14) instead has constraints $d_a + \zeta_k^{(i)} \in R^{-1}(s')(\lambda)$. Since the proof is so similar, we omit it here and refer to [6] for the complete proof. ■

7.2.3.2 Upper Bounds from Hoeffding's Inequality

As we show in [6, Appendix A], the scenario approach might lead to conservative estimates of the upper bound in Eq. (7.14). Thus, we instead apply *Hoeffding's inequality* [BLM13] to infer an upper bound \hat{p} of the probability $\mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \cap R^{-1}(s') \neq \emptyset\}$ in Eq. (7.14). Concretely, this probability describes the parameter of a Bernoulli random variable, which has value 1 if $\mathcal{H}_{s,a}(\omega) \cap R^{-1}(s') \neq \emptyset$ and 0 otherwise. The sample sum $N_{s,a,s'}^\cap$ of this random variable is given by the number of sets $\mathcal{H}_{s,a}^{(i)}$ that intersect with region $R^{-1}(s')$, i.e.,

$$N_{s,a,s'}^\cap := |\{i \in \{1, \dots, N\} : \mathcal{H}_{s,a}^{(i)} \cap R^{-1}(s') \neq \emptyset\}|.$$

Using Hoeffding's inequality, we obtain the following upper bound on the transition probability in Eq. (7.14).

Lemma 7.11 (Upper bound probability) Let $\zeta_k^{(1)}, \dots, \zeta_k^{(N)}$ be a set of $N \in \mathbb{N}$ samples of the noise ζ_k , and let $\beta \in (0, 1)$ be a confidence parameter. For fixed $s, s' \in S$ and $a \in \text{Act}(s)$, define $\mathcal{H}_{s,a}^{(i)} = \mathcal{T}_a + \Delta_s + \zeta_k^{(i)}$ and determine the value of $N_{s,a,s'}^\cap \leq N$. Then, it holds that

$$\begin{aligned} \mathbb{P}^N \left\{ (\zeta_k^{(1)}, \dots, \zeta_k^{(N)}) \in \Omega^N : \mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \cap R^{-1}(s') \neq \emptyset\} \leq \hat{p} \right\} \\ \geq 1 - \frac{\beta}{2}, \end{aligned}$$

where the upper bound \hat{p} is computed as

$$\hat{p} = \min \left\{ 1, \frac{N_{s,a,s'}^\cap}{N} + \sqrt{\frac{1}{2N} \log \left(\frac{2}{\beta} \right)} \right\}.$$

Proof. Assume we are given N samples of a Bernoulli random variable with unknown probability p , and with sample sum denoted by $\bar{N} \in \{0, \dots, N\}$. For this Bernoulli random variable, Hoeffding's inequality [BLM13] states that

$$\mathbb{P}^N \left\{ pN \leq \bar{N} + \varepsilon N \right\} \geq 1 - e^{-2\varepsilon^2 N}, \quad (7.19)$$

for all $\varepsilon > 0$. Thus, the expected value pN over N samples is upper bounded by the sample sum \bar{N} plus the value of ε . We are interested in the unknown probability p

Hoeffding's inequality

instead of the sum over N samples, so we rewrite Eq. (7.19) as

$$\mathbb{P}^N \left\{ p \leq \frac{\bar{N}}{N} + \varepsilon \right\} \geq 1 - e^{-2\varepsilon^2 N}. \quad (7.20)$$

Moreover, let $\frac{\beta}{2} = e^{-2\varepsilon^2 N}$ and rewrite Eq. (7.20) as

$$\mathbb{P}^N \left\{ p \leq \frac{\bar{N}}{N} + \sqrt{\frac{1}{2N} \log\left(\frac{2}{\beta}\right)} \right\} \geq 1 - \frac{\beta}{2}. \quad (7.21)$$

In Lemma 7.11, the unknown probability p is the probability that for a random $\omega \in \Omega$, the set $\mathcal{H}_{s,a}(\omega)$ intersects with the region $R^{-1}(s')$:

$$p = \mathbb{P}\{\omega \in \Omega : \mathcal{H}_{s,a}(\omega) \cap R^{-1}(s') \neq \emptyset\},$$

such that its sum \bar{N} over N samples becomes

$$\bar{N} = N_{s,a,s'}^\cap.$$

Note that probability p in Eq. (7.21) cannot exceed 1, so we obtain

$$\mathbb{P}^N \left\{ p \leq \min \left\{ 1, \frac{N_{s,a,s'}^\cap}{N} + \sqrt{\frac{1}{2N} \log\left(\frac{2}{\beta}\right)} \right\} \right\} \geq 1 - \frac{\beta}{2},$$

which equals the desired expression. This concludes the proof of Lemma 7.11. ■

7.2.3.3 Probability Intervals with PAC Guarantees

By combining Lemmas 7.10 and 7.11, we obtain the following PAC probability interval for each individual transition $(s, a, s') \in S \times \text{Act} \times S$ of the IMDP abstraction.

Theorem 7.12 (PAC probability interval) For fixed $s, s' \in S$ and $a \in \text{Act}(s)$, let $\xi_k^{(1)}, \dots, \xi_k^{(N)}$ be a set of $N \in \mathbb{N}$ samples of the noise ζ_k . For the collection $(\mathcal{H}_{s,a}^{(i)})_{i=1}^N$, compute \check{p} and \hat{p} using Lemmas 7.10 and 7.11. Then, the transition probability $P(s, a)(s')$ is bounded by

$$\mathbb{P}^N \left\{ (\xi_k^{(1)}, \dots, \xi_k^{(N)}) \in \Omega^N : \check{p} \leq P(s, a)(s') \leq \hat{p} \right\} \geq 1 - \beta. \quad (7.22)$$

Proof. Theorem 7.12 follows directly by combining Lemmas 7.10 and 7.11 via the union bound^a with the probability interval in Eq. (7.14), which asserts that these bounds are both correct with a probability of at least $1 - \beta$. ■

^aThe union bound (Boole's inequality) states that the probability that at least one of a finite set of events happens, is upper bounded by the sum of the probabilities of these events [CB21].

Observe that the statistical guarantee in Theorem 7.12 is for an individual transition probability interval. In Sect. 7.3, we will lift this guarantee on individual transitions to a guarantee on the whole IMDP abstraction.

7.2.3.4 Counting samples

Theorem 7.12 requires two ingredients:

1. the sample counts $N_{s,a,s'}^{\text{out}}$ (fully outside of $R^{-1}(s')$) and $N_{s,a,s'}^{\cap}$ (at least partially contained in $R^{-1}(s')$), and
2. the confidence probability β .

Thus, the problem of computing probability intervals reduces to a *counting problem* on the samples.

Merging samples | To reduce the complexity of this counting procedure, we merge samples that are very similar into a single, overapproximating sample. Formally, let $\xi > 0$ be a tuning parameter that reflects the maximum distance for two samples to be merged. We merge two samples $\mathcal{H}_{s,a}$ and $\mathcal{H}'_{s,a}$ if their centers, denoted by $h \in \mathcal{H}_{s,a}$ and $h' \in \mathcal{H}'_{s,a}$ are at most a ξ -distance apart, i.e.,

$$\|h - h'\|_2 \leq \xi. \quad (7.23)$$

If Eq. (7.23) holds, we define one larger set $\mathcal{H}''_{s,a} \supseteq \mathcal{H}_{s,a} \cup \mathcal{H}'_{s,a}$ (without loss of generality, we define $\mathcal{H}''_{s,a}$ as a hyperrectangle for simplicity). Then, to determine sets $N_{s,a,s'}^{\text{out}}$ and $N_{s,a,s'}^{\cap}$, a merged sample set is associated with *the number of samples that it represents*. For example, if $\mathcal{H}''_{s,a} \subseteq R^{-1}(s')$ (i.e., the merged sample that represents $\mathcal{H}_{s,a}$ and $\mathcal{H}'_{s,a}$ is contained in $R^{-1}(s')$), we add 2 to the value of $N_{s,a,s'}^{\text{out}}$. Doing so yields more conservative yet sound estimates of the counts $N_{s,a,s'}^{\text{out}}$ and $N_{s,a,s'}^{\cap}$, and thus also of the PAC probability intervals. It is easily verified that as $\xi \rightarrow 0$, the number of merged samples converges to zero as well.

Heuristic for merging samples | Determining the best way to merge samples is a problem of combinatorial complexity. In our implementation, we thus use a heuristic in which we greedily try to merge samples by following three steps:

1. Select a sample, denoted by $\mathcal{H}_{s,a}$, that has not been merged yet.
2. Merge this sample with all other (non-merged) samples for which Eq. (7.23) holds, and we remove these samples from the list of non-merged samples. Mark $\mathcal{H}_{s,a}$ as a merged sample, even if no samples are within a ξ -distance of $\mathcal{H}_{s,a}$.
3. Repeat until no non-merged samples remain.

Because we mark a sample as merged even if no samples are within a ξ -distance of it, this heuristic terminates in at most N iterations.

While the improvement in computational complexity strongly depends on the model at hand, we observe significant improvements in the experiments in Sect. 7.4. For example, for the numerical experiments in Sect. 7.4, we used 20 000 samples to compute probability intervals, but using the proposed merging procedure with $\xi = 0.01$, we reduced this to around 1 000 merged samples.

7.3 Abstraction Algorithm

We present a variant of Algorithm 6.1 to solve Problem 7.5 based on the IMDP abstraction proposed in this chapter. Recall that Algorithm 6.1 consists of an *offline planning* phase,

in which we create the IMDP and compute a robust optimal policy, and an *online control* phase in which we automatically derive a provably-correct Markov policy for the continuous model. Upon termination, the algorithm either returns a scheduler that can be used to solve Problem 7.5, or that no such satisfactory scheduler could be found. For conciseness, we only describe the changes compared to Algorithm 6.1.

Enabled IMDP actions | Recall that we define the IMDP actions via backward reachability computations under the nominal model in Eq. (7.2), which is defined for a fixed nominal parameter value $\hat{\gamma} \in \Gamma$. Hence, we replace line 5 of Algorithm 6.1 by Eq. (7.5) for the definition of the enabled actions $Act_{\hat{\gamma}}(s)$ in IMDP state $s \in S$.

Computing probability intervals | Similar to Algorithm 6.1, we compute PAC probability intervals for the IMDP abstraction. However, we now use Theorem 7.12 to compute these intervals, thus modifying line 11 of the original algorithm.

Interface function | Recall from Theorem 5.17 and Corollary 5.25 that by restricting the Markov policy μ to an interface function (for the relation induced by the IMDP abstraction), the lower bound on the reach-avoid probability on the IMDP carries over to the DTSS (under that Markov policy). We again leverage this result to derive a control for the RDTSS based on an optimal robust scheduler for the IMDP abstraction. Specifically, let $I_R^\sigma: X \times \{0, \dots, h-1\} \rightarrow 2^U$ be the interface function for IMDP scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_I}$, which is defined for all $x \in X$ and $k \in \{0, \dots, h-1\}$ as

$$I_R^\sigma(x, k) := \{u \in U : A(\hat{\gamma})x + B(\hat{\gamma})u + q(\hat{\gamma}) \in \mathcal{T}_a, a = \sigma_k(s)\}. \quad (7.24)$$

As we shall see, restricting the Markov policy to this interface function leads to a solution to Problem 7.5.

7.3.1 Solving Problem 7.5 with high probability

Analogous to Theorem 6.29, we show that our modified version of Algorithm 6.1 yields, with a probability of at least $1 - \beta \cdot |Act| \cdot |S|$, a solution to Problem 6.6. In other words, termination of the algorithm implies that, with a probability of at least $1 - \beta \cdot |Act| \cdot |S|$, we have found a solution to Problem 7.5.

Theorem 7.13 (Solution to Problem 7.5) Let $\mathcal{S}_R = (X, U, \Gamma, x_I, \varsigma, F)$ be an RDTSS with a reach-avoid specification $\varphi = (X_T, X_U, h)$, and let $\rho \in [0, 1]$. Suppose that Algorithm 6.1 (with the modifications above) terminates and returns the optimal robust scheduler $\dot{\sigma}^* \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_I}$. Then, it holds that

$$\mathbb{P}^N \left\{ (\varsigma_k^{(1)}, \dots, \varsigma_k^{(N)}) \in \Omega^N : \min_{\gamma \in \Gamma} \Pr_{\mu}^{\mathcal{S}_R[\gamma]}(x_I \models \varphi) \geq \rho \right\} \geq 1 - \beta \cdot |Act| \cdot |S|,$$

where $\mu_k(x) \in I_R^{\dot{\sigma}^*}(x, k)$ for all $x \in X$ and $k \in \{0, \dots, h-1\}$, and where the interface function $I_R^{\dot{\sigma}^*}$ is defined by Eq. (7.24).

Proof. The proof is analogous to that of Theorem 6.29, so we only provide a sketch here, while referring to [6] for details. Termination of the algorithm implies that the

optimal reach-avoid probability on the IMDP is at least ρ . Since the abstract IMDP is “correct” with a probability of at least $1 - \beta \cdot |Act| \cdot |S|$, this reach-avoid probability carries over to the RDTSS with this same probability. Thus, with a probability of at least $1 - \beta \cdot |Act| \cdot |S|$, it holds that

$$\min_{\gamma \in \Gamma} \Pr_{\mu}^{S_R[\gamma]}(x_I \models \varphi) \geq \rho,$$

which concludes the proof. ■

7.4 Experimental Evaluation

We perform experiments on three benchmarks to answer the question: “*Can our method synthesize Markov policies that are robust against set-bounded uncertainty in parameters?*” We implement our approach in our Python tool DynAbs, which we present in more detail in Chapter 14. When DynAbs returns an optimal robust IMDP scheduler $\check{\sigma}^*$, then it is guaranteed (with the specified confidence probability) to solve Problem 7.5 as per Theorem 7.13. For all experiments, we use a confidence probability in Theorem 7.13 of $\beta = 10^{-8}$ on every unique probability interval, which (depending on the size of the IMDP) leads to an overall confidence of $1 - \beta \cdot |Act| \cdot |S|$.

Reproducibility | The Python code to reproduce the experimental results is provided with DynAbs; see Chapter 14 for details. All experiments in this chapter are run on a computer with 32 3.7GHz cores and 64 GB of RAM.

7.4.1 Longitudinal drone dynamics

Consider again Example 7.4 of a drone with an uncertain mass $m \in [0.75, 1.25]$. We fix the nominal value of the mass as $\hat{m} = 1$ (which can be interpreted as a maximum likelihood estimate of the mass). To purely show the effect of set-bounded parameter uncertainty, we set the covariance of the stochastic uncertainty in ζ_k (being a Gaussian distribution) negligibly small. The control task is to reach a position of $p_k \geq 8$ before time $h = 12$, while avoiding speeds of $|v_k| \geq 10$. We create a partition covering $\mathcal{Z} = [-10, 14] \times [-10, 10]$ into 24×20 regions. We use Theorem 7.12 with $N = 20\,000$ samples of the stochastic noise to compute the probability intervals of the IMDP. We compare against a baseline that builds an IMDP for the nominal model only, thus neglecting parameter uncertainty.

Neglecting parameter uncertainty is unsafe | The run time for generating the abstract IMDP $\mathcal{M}_{\mathbb{I}}$ and computing an optimal robust scheduler $\check{\sigma}^*$ is around 3 seconds. Let $V^*(s'_I)$ denote the robust reach-avoid probability on the IMDP under the maximizing scheduler $\check{\sigma}^*$ from initial state $s'_I \in S$:

$$V^*(s'_I) := \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \min_{\tau \in \mathfrak{T}_{\text{Markov}}^{\mathcal{M}_{\mathbb{I}}}} \Pr_{\sigma, \tau}^{\mathcal{M}_{\mathbb{I}}}(s'_I \models \neg S_U \cup^{\leq h} S_G).$$

To validate that Theorem 7.13 holds in practice, we estimate (using Monte Carlo simulations) the reach-avoid probability $\Pr_{\mu}^{S_R[\gamma]}(x_I \models \varphi)$ on the RDTSS for different parameter values $\gamma \in \Gamma$. We say that the Markov policy μ is safe for parameter $\gamma \in \Gamma$ and initial state

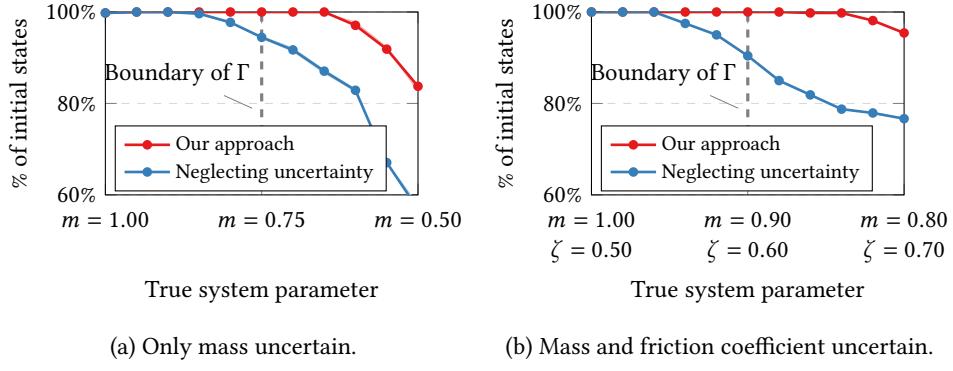


Figure 7.4: Percentage of initial states x_I where $\bar{V}(x_I, \gamma, \mu) \geq \rho$ holds (% of safe initial states) for different values of γ . Figures (a) and (b) show the results for the drone model with one and two uncertain parameters, respectively. Our approach that accounts for parameter uncertainty leads to correct results for all parameter values $\gamma \in \Gamma$ in the uncertainty set, whereas neglecting uncertainty does not.

x_I if the estimated reach-avoid probability, which we denote by $\bar{V}(x_I, \gamma, \mu)$, is indeed above the required threshold ρ , i.e.,

$$\bar{V}(x_I, \gamma, \mu) \geq \rho. \quad (7.25)$$

As per Theorem 7.13, for each triple (x_I, γ, μ) , we expect this inequality to hold with probability at least $1 - \beta \cdot |Act| \cdot |S|$.

In Fig. 7.4a, we show the percentage of initial states x_I where Eq. (7.25) holds, for different values of the uncertain parameter γ . With our approach, Eq. (7.25) holds for all parameter values $\gamma \in \Gamma$ (and even beyond this set). By contrast, neglecting the parameter uncertainty (by reasoning over the nominal mass $\hat{m} = 1$ only) leads to incorrect results. We show simulated trajectories under an actual mass $m = 0.75$ in Fig. 7.5. These trajectories confirm that our approach safely reaches the goal region while the baseline does not, as it neglects parameter uncertainty, which is equivalent to assuming $\Delta_i = 0$ in Eq. (7.13).

Multiple uncertain parameters | To show that our approach is applicable to models with multiple uncertain parameters, we extend the drone model with an uncertain spring coefficient $\zeta \in \mathbb{R}$, yielding the following model:

$$x_{k+1} = \begin{bmatrix} p_{k+1} \\ v_{k+1} \end{bmatrix} = \begin{bmatrix} 1 & \delta_t \\ -\frac{\zeta}{m} & 1 - \frac{0.1\delta_t}{m} \end{bmatrix} x_k + \begin{bmatrix} \frac{\delta_t^2}{2m} \\ \frac{\delta_t}{m} \end{bmatrix} u_k + \xi_k.$$

To write this model in the form of Eq. (7.1), we need four matrices A_1, \dots, A_4 and B_1, \dots, B_4 , which are defined for the combinations of the minimum/maximum mass and spring coefficient. We constrain the mass to $0.9 \leq m \leq 1.1$ and the spring coefficient to $0.4 \leq \zeta \leq 0.6$. We fix their nominal values as $\hat{m} = 1$ and $\hat{\zeta} = 0.5$. Fig. 7.4b shows that our

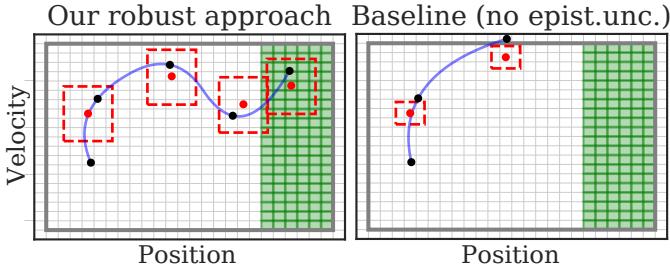


Figure 7.5: With our approach, the system safely reaches the goal (in green). By contrast, the baseline neglecting set-bounded parameter uncertainty leaves the safe set (gray box), as it underestimates the successor state sets $\mathcal{H}_{s,a}$ defined in Eq. (7.13) (red boxes).

approach again results in correct results (well beyond the uncertainty set Γ), whereas neglecting parameter uncertainty does not. This result confirms that our approach can be used to solve Problem 7.5, also in the presence of multiple uncertain parameters.

7.4.2 Building temperature control

We consider a temperature control problem for a 5-room building, each with a dedicated radiator that has an uncertain power output of $\pm 10\%$ around its nominal value. The 10D state of this model captures the temperatures of 5 rooms and 5 radiators. The goal is to maintain a temperature within $21 \pm 2.5^\circ\text{C}$ for 15 steps of 20 min.

Dynamics | Each room $i = 1, \dots, 5$ is modeled by its (zone) temperature $T_i^z \in \mathbb{R}$ and radiator temperature $T_i^r \in \mathbb{R}$. Each room has a scalar control input $T_i^{\text{ac}} \in \mathbb{R}$ reflecting the air conditioning (ac) temperature, which is constrained to $15 \leq T_i^{\text{ac}} \leq 30$. The change in the temperature of zone i depends on the temperatures in the subset of neighboring rooms, denoted by $\mathcal{J} \subseteq \{1, \dots, 5\} \setminus \{i\}$. Thus, the thermodynamics of the room temperature T_i^z and radiator temperature T_i^r of room i are

$$\begin{aligned}\dot{T}_i^z &= \frac{1}{C_i} \left[\sum_{j \in \mathcal{J}} \frac{T_j^z - T_i^z}{R_{i,j}} + \frac{T_{\text{wall}} - T_i^z}{R_{i,\text{wall}}} + m C_{pa} (T_i^{\text{ac}} - T_i^z) + P_i (T_i^r - T_i^z) \right] \\ \dot{T}_i^r &= k_1 (T_i^z - T_i^r) + k_0 w (T_i^{\text{boil}} - T_i^r),\end{aligned}$$

where C_i is the thermal capacitance of zone i , $R_{i,j}$ is the resistance between zones i and j , T_{wall} is the wall temperature, m is the air mass flow, C_{pa} is the specific heat capacity of air, and P_i is the rated output of radiator i . Moreover, k_0 and k_1 are constants, and w is the water mass flow from the boiler. In our experiments, we use a discrete-time representation of the thermodynamics above with additive Gaussian process noise (we omit an explicit definition for brevity).

Parameter uncertainty | We assume the rated output of each radiator $i = 1, \dots, 5$ to be uncertain, within an interval of $0.8 \leq P_i \leq 1.2$. Thus, we can model the building thermodynamics as a linear RDTSS with dynamics of the form in Eq. (7.1). We fix its nominal value to be $\hat{P}_i = 1$, so the uncertainty is $\pm 20\%$ around the nominal value.

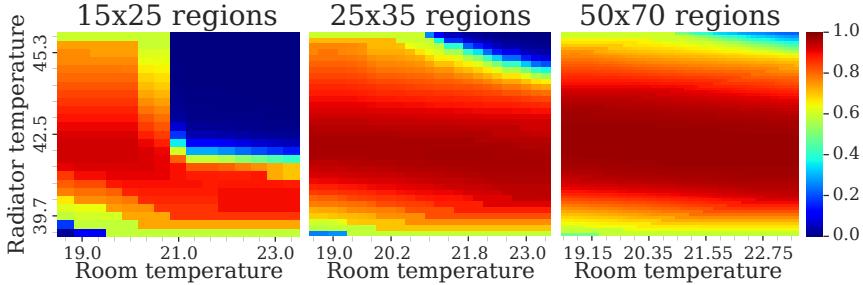


Figure 7.6: Heatmaps of the satisfaction probabilities on the IMDP for the temperature control problem (shown for a single room).

Interactions between rooms as nondeterminism | Since a direct partitioning of the 10D state space is infeasible, we capture *any possible thermodynamic interaction* between rooms in the uncertain additive term $q_k \in Q$, which lies within a convex set $Q \subset \mathbb{R}^n$. Specifically, the set Q_i affecting the thermodynamics of room $i \in \{1, \dots, 5\}$ is defined as follows (recall that X_S denotes the safe set):

$$Q_i = \left\{ \sum_{j \in \mathcal{J}} \frac{T_j^z - T_i^z}{R_{i,j}} : T^z \in X_S \right\}.$$

In other words, the uncertainty set Q is characterized by the maximal difference between T_j^z and T_i^z within the safe set, for all $j \in \mathcal{J}$, which is 5°C for this specific reach-avoid problem (which was defined in Sect. 7.4). Thus, depending on the other parameters, we can easily derive a set-bounded representation of Q .

Discrete-time dynamics | We discretize the thermodynamics of a single room i by a forward Euler method at a time resolution of 20 min. Moreover, we consider an additive Gaussian process noise ξ_k on the room temperature of distribution $\mathcal{N}(0, 0.002)$, and on the radiator temperature of distribution $\mathcal{N}(0, 0.01)$. As the model for room i has only one uncertain parameter (the radiator power output P_i), we obtain a model in the form of Eq. (7.1) with $r = 2$ matrices (we omit the explicit matrices for brevity).

Partition refinement | We apply our method with an increasingly more fine-grained state-space partition. In Fig. 7.6, we present, for three different partitions, the maximum satisfaction probability under the robust optimal IMDP scheduler (for every initial IMDP state $s_I \in S$). These results confirm the idea that partition refinement can lead to Markov policies with better performance guarantees. A more fine-grained partition leads to more actions enabled in the abstraction, which in turn improves the robust lower bound on the reach-avoid probability.

Scalability | We investigate how considering parameter uncertainty affects the scalability of our approach. To this end, we apply our method with different partition sizes, and we compare two cases: (1) with the parameter uncertainty, and (2) without the parameter uncertainty, in which case we assume that the rated power output of each radiator is $P_i = \hat{P}_i = 1$. We present the sizes of the obtained IMDPs (measured in terms

Table 7.1: IMDP sizes and run times for different partitions on the temperature control problem (considering a single room, decoupled from the others).

Epist.unc.	Partition	States	Transitions	Run time [s]
	15×25	378	84 539	5.64
	25×35	878	308 845	10.47
	35×45	1578	2 103 986	25.05
	50×70	3503	6 149 432	62.50
	70×100	7003	51 742 285	308.08
✓	15×25	378	100 994	5.68
✓	25×35	878	463 173	11.49
✓	35×45	1578	2 932 224	30.06
✓	50×70	3505	9 520 698	80.49
✓	70×100	7003	81 763 143	475.35

of the number of edges in the underlying graph) and the run times in Table 7.1. The run times vary between 5.6 s and 8 min for the smallest (15×25) and largest (70×100) partitions, respectively. We observe that accounting for parameter uncertainty yields IMDPs with more transitions and slightly higher run times (for the largest partition: 82 instead of 52 million transitions and 8 instead of 5 min). This is due to larger successor state sets $\mathcal{H}_{s,a}^{(i)}$ in Eq. (7.13) caused by the epistemic error Δ_i .

7.5 Related Work

Set-bounded parameter uncertainty is well-suited for modeling epistemic uncertainty (i.e., uncertainty caused by a lack of knowledge [FÜ11; Sul15]) in the absence of a prior distribution. Distinguishing aleatoric from epistemic uncertainty is a key challenge towards trustworthy AI [TLS21], and has been considered in reinforcement learning [CSKG22], Bayesian neural networks [DHDU18; LSS20], and systems modeling [Smi14]. Dynamical models with set-bounded parameter uncertainty (but *without aleatoric uncertainty*) are considered by [Yed14] and [GC06]. Control of models similar to ours is studied by [Mod22], albeit only for stability specifications.

Model-based approaches | The abstraction-based policy synthesis methods surveyed in Sect. 5.2 only consider stochastic uncertainty (and not set-bounded parameter uncertainty). Thus, we develop the first abstraction-based formal policy synthesis method that simultaneously captures set-bounded parameter uncertainty and stochastic uncertainty for models with continuous state and action spaces. Other methods include, for example, using optimization for reach-avoid control of linear but *non-stochastic* models with bounded disturbances [FQMN⁺22]. Furthermore, so-called *funnel libraries* are used by [MT17] for robust feedback motion planning under set-bounded disturbances.

Data-driven approaches | Models with (partly) unknown dynamics can be used to express epistemic uncertainty about the underlying system. Verification of such models based on data has been done using Bayesian inference [HHA17], optimization [KBJT19; VIT22], and Gaussian process regression [JLFL20; BLD^T21]. Moreover, [KCOB21], and [COB21] use neural networks for feedback motion planning of nonlinear determ-

inistic systems with probabilistic safety and reachability guarantees. Data-driven abstractions have been developed for linear [CPM23], monotone [MGF21], event-triggered [PM23], and also nonlinear systems [GLMA⁺24]. By contrast to our setting, these approaches consider models with *non-stochastic dynamics*. A few recent exceptions exist [SZ23; LSFZ23]; however, these approaches require more strict assumptions (e.g., discrete input sets) than our model-based approach.

7.6 Discussion

We conclude this chapter by discussing the benefits, limitations, and open challenges of our robust IMDP abstraction approach.

Generality of our approach | We stress that the RDTSS in Def. 7.1 can capture many common sources of uncertainty. As we have shown, our approach simultaneously deals with set-bounded uncertainty in one or multiple parameters, as well as stochastic uncertainty due to process noise of an unknown distribution. For example, the additive term q_k enables us to generate abstractions that faithfully capture any error term represented by a bounded set (as we have done for the multi-room temperature control problem). This idea could be used to deal with *nonlinear systems*, such as non-holonomic robots [TBF05]. Concretely, we may apply our abstraction method on a linearized version of the system while treating linearization errors as nondeterministic disturbances $q_k \in Q$ in Eq. (7.1). The main challenge is then to compute this set-bounded representation Q of the linearization error, either exactly or as a tight overapproximation.

Scalability | Enforcing robustness against set-bounded parameter uncertainty hampers the scalability of our approach, especially compared to the setting in Chapter 6 (which only considers stochastic noise). For example, the approach presented in Chapter 6 scaled to systems with a 6D state, whereas we could not scale the approach presented in this chapter beyond 3D systems.² The main reason is that our sampling-based approach for computing probability intervals leads to very wide and conservative intervals, which increases the number of transitions in the IMDP significantly. Thus, an interesting direction for future research is to develop an abstraction method that leads to tighter sets of transition probabilities. One idea (which is in line with the discussion from Sect. 6.8) is to generate convex polytopic uncertainty sets (rather than intervals) and formalize the resulting abstraction as an robust Markov decision process (RMDP) (rather than an IMDP).

Safe learning | Finally, our abstraction-based approach could potentially be integrated into a *safe learning framework* [BGHY⁺22; GF15]. In such a setting, our approach may synthesize policies that guarantee safe interactions with the system, while techniques from, for example, reinforcement learning [BTSK17; ZG21] or stochastic system identification [TP19] can reduce the uncertainty sets (for the uncertain parameters) based on state observations.

²We omitted this experiment from this thesis and instead refer to our paper [6] for details on this *automated anesthesia delivery* benchmark.

Summary

- ⇒ We have presented a novel abstraction-based policy synthesis method for discrete-time dynamical models with both stochastic and set-bounded parameter uncertainty.
- ⇒ Our approach reasons probabilistically over stochastic noise while reasoning robustly over parameter uncertainty.
- ⇒ Capturing parameter uncertainty yields more robust Markov policies.
- ⇒ However, capturing parameter uncertainty also leads to significantly larger abstractions, thus currently limiting the practical scalability to (at most) 3D-state systems.

Part III

Parametric Markov Decision Processes

8 Foundations of Parametric MDPs

Summary | Parametric Markov decision processes (pMDPs) extend standard MDPs with transition probabilities described by polynomials over parameters. Thus, pMDPs encode both uncertainty and dependencies between transition probabilities. By substituting each parameter with a concrete value, a pMDP reduces to a standard MDP which we can analyze using the methods from Chapter 3. In this background chapter, we introduce pMDPs and describe how to analyze them. We highlight challenges related to pMDPs and discuss how we aim to address them in this thesis.

Origins | This chapter does not contain novel content and instead presents a brief introduction to pMDPs. For a comprehensive treatment of pMDPs, we refer to [Jun20].

Background | The reader is assumed to be familiar with standard MDPs (Chapter 3).



8.1 Introduction

In Chapter 3, we have introduced robust Markov decision processes (RMDPs) as extensions of Markov decision processes (MDPs) with sets of transition probabilities. We have heavily used interval Markov decision processes (IMDPs) in Chapters 6 and 7, which are a special case of RMDP, as models for finite-state abstractions of discrete-time stochastic systems (DTSSs). However, to make analyzing RMDPs and IMDPs tractable, we assumed that the actual choice of transition probabilities in the uncertainty sets is made independently for all states and actions (called the *rectangularity assumption*; see Remark 3.26). This assumption is often unrealistic and leads to too pessimistic verification results.

Dependencies between transition probabilities | As an example, consider a simple motion planning scenario where an unmanned aerial vehicle (UAV) is tasked to transport a certain payload to a target location. External factors such as the wind direction may affect the dynamics of the UAV. The assumption that such weather conditions are independent between the different states of the UAV is unrealistic and may yield pessimistic verification results. Similarly, in the verification of network protocols, we typically do not precisely know the channel quality (i.e., the loss rate). However, the loss rate is independent of whether we are either probing or actually sending useful data over the network. A typical verification task would be to show that the protocol yields a sufficiently high quality of service. Assuming that the channel quality depends on the protocol state may be too pessimistic to establish that the protocol provides the required quality of service.

Outline | In this chapter, we introduce parametric Markov decision processes (pMDPs), which add dependencies (or couplings) between the probabilities of different transitions [LMT07; HHZ11a; Daw04; JJK22]. We first introduce the formal definition of a pMDP in Sect. 8.2. Then, we discuss in Sect. 8.3 how to analyze pMDPs. Finally, we list challenges related to pMDPs in Sect. 8.4 and discuss how we aim to address these challenges in Part III of this thesis.

8.2 Parametric MDPs

Formally, let V be a set of parameters. We denote the set of polynomials over parameters V with rational coefficients by $\mathbb{Q}[V]$. A *parametric Markov decision process (pMDP)* extends a standard MDP with a set of parameters and a parametric transition function. The transition probabilities of a pMDP are not expressed as values in $[0, 1]$, but as polynomial functions over the parameters.¹

parametric
MDP

Definition 8.1 (pMDP) A *parametric Markov decision process (pMDP)* is a tuple $\mathcal{M}_V := (S, \text{Act}, s_I, V, P, L, r)$ where S is a finite set of *states*, Act is a finite set of *actions*, $s_I \in S$ is an *initial state*, V is an (ordered) set of *parameters*, $P: S \times \text{Act} \times S \rightarrow \mathbb{Q}[V]$ is a *parametric transition probability function*, $L: S \rightarrow 2^{A^P}$ is a *labeling function*, and $r: S \rightarrow \mathbb{R}_{\geq 0}$ is a *state reward function*.

Because pMDPs are direct extensions of MDPs, the concepts of paths, schedulers, etc. carry over directly, for which we use the same notation as in Chapter 3.

Remark 8.2 Notice that the (parametric) transition function in Def. 8.1 differs slightly from the definition for standard MDPs (apart from incorporating parameters). Specifically, we defined the transition function of an MDP as $P: S \times \text{Act} \rightarrow \text{Distr}(S)$, i.e., as a mapping to distributions over states. By contrast, the transition function of a pMDP maps every state-action-state triple to a function over the parameters. Thus, our definition of a pMDP does not yet encode that the transition function must express valid distributions; instead, we will take care of this later in this section.

We define a parametric Markov chain (pMC) as a pMDP with exactly one action in every state.² Analogous to discrete-time Markov chain (DTMC), we omit the actions of a pMC at all from the definition.

parametric
Markov
chain

Definition 8.3 (pMC) A *parametric Markov chain (pMC)* is a tuple $\mathcal{D}_V := (S, s_I, V, P, L, r)$ where S is a finite set of *states*, $s_I \in S$ is an *initial state*, V is an (ordered) set of *parameters*, $P: S \times S \rightarrow \mathbb{Q}[V]$ is a *parametric transition probability function*, $L: S \rightarrow 2^{A^P}$ is a *labeling function*, and $r: S \rightarrow \mathbb{R}_{\geq 0}$ is a *state reward function*.

If the labeling function L and/or reward function r are irrelevant, we may omit them from the pMDP or pMC tuple. We mainly tailor definitions and notation to pMDPs. However, being a special case of a pMDP, all concepts naturally also apply to pMCs.

¹As is common (see, for example, [Jun20]) we exclude non-rational probabilities for technical reasons.

²For consistency, we should call pMC a parametric DTMC (pDTMC). However, the term pMC is more common in the literature and is thus preferred here.

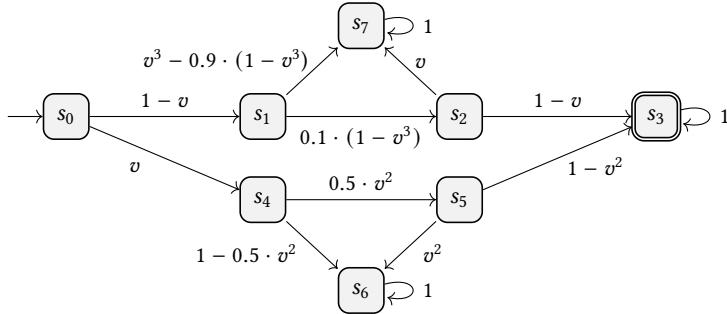


Figure 8.1: A simple pMC with a singleton parameter set $V = \{v\}$.

Example 8.4 Fig. 8.1 shows a simple pMC with a singleton parameter set $V = \{v\}$. The initial state is s_0 and we consider state s_3 as a target set. The transitions in the pMC are annotated with polynomial functions over the parameter v . For example, the transition from s_0 to s_4 is annotated with the polynomial $P(s_0, s_4) = 1 - 0.5 \cdot v^2 \in \mathbb{Q}[V]$.

8.2.1 Parameter instantiation

Given a pMDP, we can substitute each parameter $v \in V$ by a concrete value. We call such an assignment of a value to each of the parameters a *parameter instantiation*.

Formally, a parameter instantiation is a function $u: V \rightarrow \mathbb{Q}$ that maps parameters to concrete values. Because we assumed the parameters $V = \{v_1, \dots, v_n\}$, $n \in \mathbb{N}_{>0}$ to be ordered, we overload notation and often write an instantiation as the vector $u \in \mathbb{Q}^{|V|}$, which is defined as $[u(v_1), \dots, u(v_n)]^\top \in \mathbb{Q}^{|V|}$.

Applying the instantiation u to a polynomial $g \in \mathbb{Q}[V]$ yields $g[u] \in \mathbb{Q}$, which is obtained by substituting every parameter $v \in V$ in g with $u(v)$. Thus, for every $s, s' \in S$ and $a \in Act(s)$ of a pMDP, we can write $P(s, a, s')[u]$ to denote substituting the parameters V in the transition function P of a pMDP with the concrete values $u(v_1), \dots, u(v_n)$ given by the instantiation. With slight abuse of notation, we use the more convenient notation $P[u](s, a, s') := P(s, a, s')[u]$, such that $P[u]$ denotes the (parameter-free) transition function obtained by applying u to the polynomial $P(s, a, s')$ for every transition (s, a, s') .

parameter instantiation

Well-defined instantiations | In general, not all parameter instantiations lead to a valid MDP. For example, in the pMC in Fig. 8.1, assigning a value of 1.5 to the parameter v yields probabilities outside of $[0, 1]$, thus violating the definition of the transition function of a Markov chain. To exclude such invalid parameter valuations, we define a subset of instantiations that we call the *parameter space*. The parameter space is precisely the subset of instantiations that lead to valid MDPs.

Definition 8.5 (Parameter space) The *parameter space* $\mathcal{V}_{\mathcal{M}_V}$ for a pMDP $\mathcal{M}_V = (S, Act, s_I, V, P)$ is a subset of all instantiations $u: V \rightarrow \mathbb{Q}$ defined as

parameter space

$$\mathcal{V}_{\mathcal{M}_V} = \{u: V \rightarrow \mathbb{Q} : P[u](s, a, \cdot) \in \text{Distr}(S) \quad \forall s \in S, a \in Act(s)\}.$$

For any pMDP \mathcal{M}_V , we only consider instantiations $u \in \mathcal{V}_{\mathcal{M}_V}$ that belong to the parameter space $\mathcal{V}_{\mathcal{M}_V}$.

Example 8.6 In the pMC in Fig. 8.1, all values $v \in [0, 1]$ result in a DTMC with valid transition probabilities. Thus, the parameter space $\mathcal{V}_{\mathcal{M}_V}$ for this pMC is defined as the set of instantiations that map v to the interval $[0, 1]$, i.e.,

$$\mathcal{V}_{\mathcal{M}_V} = \{u: V \rightarrow \mathbb{Q} : u(v) \in [0, 1]\}.$$

Induced MDP | Applying a parameter instantiation (from the parameter space) to a pMDP yields an induced MDP, whose transition function is obtained by substituting the parameters with their concrete values. For example, we can assign a value of 0.5 to the parameter v in the pMC in Fig. 8.1. Because this instantiation is in the parameter space $\mathcal{V}_{\mathcal{D}_V}$ of the pMC, applying this instantiation results in a valid DTMC.

Formally, applying the instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$ to pMDP $\mathcal{M}_V = (S, Act, s_I, V, P)$ induces the MDP $\mathcal{M}_V[u] := (S, Act, s_I, P[u])$, where the transition function is defined as $P[u](s, a, s') := P(s, a, s')[u]$ for all $s, s' \in S$ and $a \in Act(s)$. Similarly, applying an instantiation $u \in \mathcal{V}_{\mathcal{D}_V}$ to a pMC $\mathcal{D}_V = (S, s_I, V, P)$ induces the DTMC $\mathcal{D}_V[u] := (S, s_I, P[u])$, where the transition function is defined analogously. For the induced MDP or DTMC, we can perform any of the verification tasks discussed in Chapter 3, such as computing reachability probabilities, computing cumulative expected rewards, and model checking probabilistic computation tree logic (PCTL) formulae.

8.3 Verifying Parametric MDPs

We now discuss how to analyze pMDPs. We first discuss the role of the parameters in analyzing a pMDP. Thereafter, we describe common verification problems for pMDPs.

8.3.1 Solution function

Consider the problem of computing the probability of reaching the state s_3 in the pMC in Fig. 8.1. This reachability probability depends on the values of parameter v . Thus, the result of the verification query can naturally be expressed as a function of the parameter instantiation. This function is formalized by the *solution function*.

Definition 8.7 (Solution function for fixed scheduler) Consider a pMDP $\mathcal{M}_V = (S, Act, s_I, V, P, L)$ with parameter space $\mathcal{V}_{\mathcal{M}_V}$ and a scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_V}$. The *solution function* $\text{sol}_{\varphi, \sigma}^{\mathcal{M}_V}: \mathcal{V}_{\mathcal{M}_V} \rightarrow [0, 1]$ for the satisfaction probability of a PCTL path formulae φ on the pMDP C_V with scheduler σ is defined as

$$\text{sol}_{\varphi, \sigma}^{\mathcal{M}_V}: u \mapsto \Pr_{\sigma}^{\mathcal{M}_V[u]}(s_I \models \varphi).$$

Solution functions for other measures, such as cumulative expected rewards, are defined analogously. However, we omit an explicit definition for brevity.

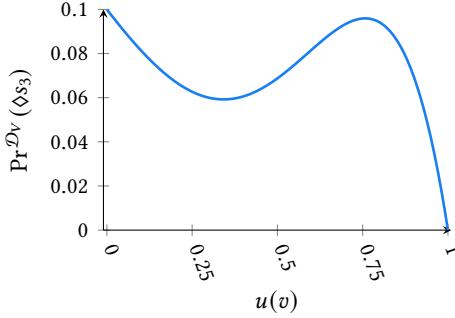


Figure 8.2: Solution function for the probability of reaching s_3 of the pMC in Fig. 8.1.

Example 8.8 For the pMC in Fig. 8.1, consider the PCTL path formula $\varphi = \diamond \{s_3\}$, i.e., eventually reach state s_3 . The solution function $\text{sol}_{\varphi}^{\mathcal{D}_V} : \mathcal{V}_{\mathcal{D}_V} \rightarrow [0, 1]$ for this formula is a polynomial function over the value of parameter v :

$$\begin{aligned}\text{sol}_{\varphi}^{\mathcal{D}_V}(u) &= (1 - u(v)) \cdot 0.1 \cdot (1 - u(v)^3) \cdot (1 - u(v)) \\ &\quad + u(v) \cdot 0.5u(v)^2 \cdot (1 - u(v)^2) \\ &= (1 - u(v))^2 \cdot 0.1 \cdot (1 - u(v)^3) + 0.5u(v)^3 \cdot (1 - u(v)^2).\end{aligned}$$

This solution function is plotted in Fig. 8.2.

One may wonder why we write the solution function in Example 8.8 using $u(v)$ and not directly as a function of v . However, v is a (symbolic) parameter rather than a variable, and thus, writing the solution function as a function of v would be incorrect. For mathematical correctness, we thus insist on the more formal definition and define the solution function as a function of the instantiation u . Then, the actual value of the parameter v is written as $u(v)$, just as we did in Example 8.8.

Maximizing and minimizing schedulers | The solution function in Def. 8.7 is defined for the satisfaction probability on the DTMC obtained by applying a fixed scheduler to the pMDP. We can also define solution functions for the (Markov) schedulers that maximize or minimize a given measure.

Definition 8.9 (Maximizing/minimizing solution function) Consider a pMDP $\mathcal{M}_V = (S, Act, s_I, V, P, L)$ with parameter space $\mathcal{V}_{\mathcal{M}_V}$. The *maximizing solution function* $\text{sol}_{\varphi, \text{max}}^{\mathcal{M}_V} : \mathcal{V}_{\mathcal{M}_V} \rightarrow [0, 1]$ for the satisfaction probability of a PCTL path formulae φ on the pMDP C_V is defined as

$$\text{sol}_{\varphi, \text{max}}^{\mathcal{M}_V} : u \mapsto \max_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_V}} \Pr_{\sigma}^{\mathcal{M}_V[u]}(s_I \models \varphi).$$

Similarly, the *minimizing solution function* $\text{sol}_{\varphi, \text{min}}^{\mathcal{M}_V} : \mathcal{V}_{\mathcal{M}_V} \rightarrow [0, 1]$ is defined as

$$\text{sol}_{\varphi, \text{min}}^{\mathcal{M}_V} : u \mapsto \min_{\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_V}} \Pr_{\sigma}^{\mathcal{M}_V[u]}(s_I \models \varphi).$$

Again, maximizing and minimizing solution functions for other measures, such as cumulative expected rewards, can be defined analogously.

Maximizing (or minimizing) the solution function implicitly assumes that a different maximizing (or minimizing) scheduler can be chosen for every instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$. The maximizing (minimizing) solution function in Def. 8.9 can alternatively be obtained by first computing $\text{sol}_{\varphi, \sigma}^{\mathcal{M}_V}$ for every scheduler $\sigma \in \mathfrak{S}_{\text{Markov}}^{\mathcal{M}_V}$, and then taking the maximum (minimum) over these functions for all instantiations $u \in \mathcal{V}_{\mathcal{M}_V}$.

Properties of solution functions | We describe some key properties of solution functions for pMDPs and pMCs. The exposition here is informal, and we refer to the relevant literature for further details.

1. For the measures in this thesis, the solution function for an pMDP is continuous over the subset of *graph-preserving*³ parameter instantiations [Jun20, Lemma 5.16];
2. For acyclic pMDPs, the solution function for an pMDP is continuous over the entire parameter space $\mathcal{V}_{\mathcal{M}_V}$ [Jun20, Corollary 5.19];
3. Solution functions for pMCs are rational functions over the subset of graph-preserving instantiations, and over the entire parameter space if the pMC is acyclic [HBK17; BHJ⁺20];
4. The size of the solution function is exponential in the number of parameters [HBK17; BHJ⁺20].

These properties indicate that, while solution functions are continuous functions over the (graph-preserving) parameter instantiations, they are computationally expensive to compute when the number of parameters is large.

8.3.2 Parameter synthesis

parameter synthesis

The so-called *parameter synthesis* problem for pMDPs considers computing parameter values such that the induced MDP satisfies a logical specification (for example in PCTL, see Def. 3.14). Several variants of this problem exist. For example, given a pMC and a specification φ , typical parameter synthesis queries include [JJK22]:

1. Do all parameter instantiations in (a subset of) the parameter space satisfy φ ?
2. Which parameter instantiations satisfy φ and which ones do not?
3. What parameter instantiation maximizes the probability of satisfying φ ?

Variants of these queries for pMDPs additionally reason over (typically the minimizing or maximizing) schedulers.

Example 8.10 Consider again the solution function shown in Fig. 8.2. We ask the question: “*For which parameter instantiations is this solution function at most 0.08?*” This question is an instance of the second parameter synthesis problem above. Answering this question leads to the partition of the parameter space shown in Fig. 8.3. The parameter instantiations that satisfy this threshold of 0.08 are colored green, and the ones that do not are colored red.

³A set of instantiations $U \subseteq \mathcal{V}_{\mathcal{M}_V}$ is graph-preserving if all $u, u' \in U$ lead to MDPs $\mathcal{M}[u]$ and $\mathcal{M}[u']$ with the same underlying graph (i.e., transitions cannot vanish under some instantiations).

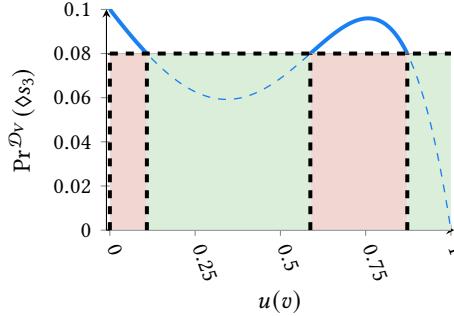


Figure 8.3: Partition of the parameter space $\mathcal{V}_{\mathcal{D}_V}$ into a satisfying (green) and unsatisfying (red) region for the reachability PCTL state formula $\Phi = \mathbb{P}_{\leq 0.13}(\diamond s_3)$ for the pMC in Fig. 8.1.

Solving parameter synthesis problems is intrinsically challenging because parameters may appear in multiple parts of the model and can take uncountably many values. Exact solutions to parameter synthesis problems typically reason over the solution function directly [Daw04; HHZ11b; GHS18; BHJH⁺20]. However, the size of the solution function is exponential in the number of parameters, and thus, exact methods are typically limited to a couple of parameters only. The problem of whether there exists a value in the parameter space that satisfies a reachability specification is ETR-complete,⁴ and finding a satisfying parameter value is exponential in the number of parameters [WJPK19].

Approximating methods for parameter synthesis have also been proposed, for example, utilizing convex optimization [CJJK⁺17; CJJK⁺18] and sampling-based methods [CHHK⁺13; MMAG14]. Another technique is the parameter lifting algorithm (PLA) developed by [QDJH⁺16], which drops any dependencies between parameters and replaces the parametric probabilities with nondeterministic choices. As such, PLA turns a pMC into an MDP, leading to upper and lower bounds on solutions to parameter synthesis problems. By iteratively partitioning the parameter space into regions that (conservatively) satisfy or violate the specification, PLA can provide an approximate solution to parameter synthesis problems [DJJC⁺15; JÁHJ⁺24]. However, due to the partitioning of the parameter space, such approaches still scale poorly with the number of parameters: Most benchmarks reported in [JÁHJ⁺24] contain one or two parameters only, and the only exception with more (namely eight) parameters only has 52 states and 133 transitions.

8.4 Challenges

The following two chapters of this thesis focus on two challenges related to pMDPs.

Chapter 9: Scaling to more parameters via distributions over parameters | As we discussed above, analyzing pMDPs is intrinsically challenging and is typically limited to a couple of parameters only. In Chapter 9, we aim to overcome this challenge by

⁴The ETR satisfiability problem is to decide if there exists a satisfying assignment to the real variables in a Boolean combination of a set of polynomial inequalities. It is known that $\text{NP} \subseteq \text{ETR} \subseteq \text{PSPACE}$.

jointly considering two research questions related to MDPs:

1. How can we scale the verification of pMDPs to higher numbers of parameters?
2. How can we incorporate prior knowledge about the parameter values of a pMDP?

We address these questions by considering a variation of the pMDP verification problem. Specifically, we consider a setting where the parameters of a pMDP are distributed according to a (possibly unknown) probability distribution. This distribution encodes prior knowledge about the parameter values. We only assume sampling access to the distribution, but we do not require the distribution to be known explicitly. As our main contribution, we present a method that, given a set of samples for the parameter values, provides statistical guarantees on the probability of satisfying a given PCTL specification. We show that our method scales to pMDPs with (in the order of) a thousand parameters, which is way beyond the limit for exact methods.

Chapter 10: Using pMCs for more efficient model learning | In Chapter 10, we ask the question: “*How can we leverage knowledge about the parametric structure of a model to make data-driven methods for analyzing this model more efficient?*” To answer this question, we consider the problem of learning a pMC from data. We represent the learned model as a Markov chain whose transition probabilities are both parametric and set-valued. Parametric probabilities allow modeling dependencies, and sets of probabilities reflect uncertainty due to limited data. As our main contribution, we develop novel and efficient methods for performing sensitivity analysis of such *parametric robust Markov chains (prMCs)*. More specifically, we measure sensitivity in terms of partial derivatives of the solution function with respect to the underlying parameters. We show that such a sensitivity analysis can improve the data efficiency of learning algorithms.

Summary

- ⇒ Parametric Markov decision processes (pMDPs) extend standard MDPs with transition probabilities described by polynomials over parameters.
- ⇒ A parameter instantiation substitutes each parameter with a concrete value, yielding an induced MDP.
- ⇒ Solution functions can be used to analyze measures on pMDPs; however, their size is exponential in the number of parameters.

9 The Scenario Approach for Parametric MDPs

Summary | We consider parametric Markov decision processes (pMDPs) with an unknown probability distribution over the parameter instantiations. This setting essentially defines a distribution over (parameter-free) MDPs. The problem is to compute the probability that an MDP drawn according to this distribution satisfies a given temporal logic specification. Solving this problem precisely is infeasible. In this chapter, we propose a sampling-based method based on the scenario approach to compute a solution to this problem with statistical guarantees. Specifically, based on a finite number of samples of the parameters, our method yields high-confidence bounds on the probability of satisfying the specification. The number of samples required to obtain a high confidence on these bounds is independent of the number of states and the number of random parameters. Our experiments show that several thousand samples suffice to obtain tight and high-confidence bounds on the satisfaction probability.

Origins | This chapter is based on the following journal article:

- [2] Badings, Cubuktepe, Jansen, Junges, Katoen and Topcu (2022) ‘Scenario-Based Verification of Uncertain Parametric MDPs’. STTT.

Some experiments presented in [2] are left out of this thesis for brevity.

9

Background | The reader is assumed to be familiar with pMDPs and their analysis, as discussed in Chapter 8.



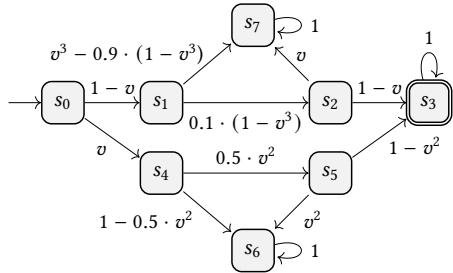
9.1 Introduction

In Chapter 8, we introduced parametric Markov decision processes (pMDPs) as versions of Markov decision processes (MDPs) where the transition probabilities are described by polynomial functions over parameters. While pMDPs allow modeling uncertainty and dependencies between transition probabilities, they do not encode prior knowledge about the values of the parameters. For example, pMDPs cannot model that certain values for the parameters are more likely than others.

We want to analyze a pMDP while considering such additional information about the *typical* (that is, the most likely) parameter values. Similar to [SBHH17], we, therefore, assume that the parameters are random variables. For instance, consider our recurring example of flying an unmanned aerial vehicle (UAV) in windy conditions. The wind speed is uncertain, but from historical weather data, we may derive a probability distribution over wind speeds [PK08]. Thus, modeling parameters as random variables allows for encoding *prior knowledge* about the values of the parameters.

9.2 Motivating Example

To illustrate our problem setting more concretely, consider again the parametric Markov chain (pMC) \mathcal{D}_V with a single parameter $V = \{v\}$ shown in Fig. 8.1 (copied here for convenience). Recall that the parameter space $\mathcal{V}_{\mathcal{D}_V}$ for this pMC is the set of all instantiations $u: V \rightarrow \mathbb{Q}$ such that $u(v) \in [0, 1]$. We again consider the solution function $\text{sol}_{\varphi, \max}^{M_V}$ for the maximal satisfaction probability of the probabilistic computation tree logic (PCTL) path formula $\varphi = \diamond \{s_3\}$ (i.e., eventually reaching state s_3). However, we now also consider a probability measure \mathbb{P} over the parameter space $\mathcal{V}_{\mathcal{D}_V}$, which encodes prior knowledge about the value of parameter v . We do not assume that \mathbb{P} is known explicitly, but we do assume that we can sample from the distribution (or that we have access to a set of independent and identically distributed (i.i.d.) samples). In this example, we consider the Gaussian distribution over the parameter space shown in Fig. 9.1 (truncated at 0 and 1 to avoid invalid instantiations).



Verification problem | Since every parameter instantiation u maps (deterministically) to a solution $\text{sol}_{\varphi, \max}^{M_V}(u)$, we can now ask the question: “*If we draw a parameter instantiation u from the parameter space $\mathcal{V}_{\mathcal{D}_V}$ according to \mathbb{P} , what is the probability that the solution $\text{sol}_{\varphi, \max}^{M_V}(u)$ is below a certain threshold of, let’s say, 0.08?*” This is precisely the type of verification problem that we aim to solve in this chapter. However, solving this problem means integrating the solution function with respect to the probability measure \mathbb{P} , which causes two problems. First, as discussed in Chapter 8, the solution function is often prohibitively large, such that even representing the solution function is infeasible. Second, we assumed that the probability measure \mathbb{P} is not known explicitly, so even if we can represent the solution function, we cannot integrate over it analytically. Thus, solving this verification problem exactly is infeasible.

An approximate solution | By contrast, obtaining an approximate answer to the verification problem above is relatively straightforward. Let us assume that the truncated Gaussian distribution in Fig. 9.1 has a mean of 0.5 and a standard deviation of 0.2. While we assumed we do not know this distribution explicitly, we can obtain i.i.d. samples instead. Thus, we can compute an approximate answer to the verification problem by sampling many instantiations u_1, \dots, u_N from $\mathcal{V}_{\mathcal{D}_V}$ according to \mathbb{P} . For example, for $N = 1000$ samples and a threshold of 0.08, we obtain that the probability of a $\text{sol}_{\varphi, \max}^{M_V}(u)$ below this threshold is approximately 0.708. However, this answer is only an approximation, and repeating the analysis with another set of samples yields a different result of 0.675.

Statistical guarantees | We argue that such an approximate answer is often undesirable, especially for applications in safety-critical settings. Instead, we aim to provide *statistical guarantees* on the approximation quality. For example, we want to say that the probability for an instantiation $u \in \mathcal{V}_{\mathcal{D}_V}$ such that $\text{sol}_{\varphi, \max}^{M_V}(u)$ is below the threshold is

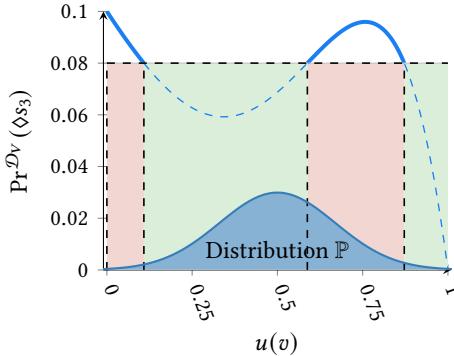


Figure 9.1: Distribution with probability measure \mathbb{P} over the parameter instantiations of the pMC in Fig. 8.1.

at least (or at most) $x \in [0, 1]$, and that claim holds with at least a confidence probability of $y \in [0, 1]$. Intuitively, such a statistical guarantee accounts for the fact that we used a finite number of samples to obtain the result. The larger the set of samples used, the more confident we can be in the result. However, no matter how large the set of samples, we can always have *bad luck* and obtain an *incorrect* result. In the following section, we will make this intuition more precise.

9.3 Problem Statement

More formally, we consider that we are given a pMDP $\mathcal{M}_V = (S, Act, s_I, V, P)$ and a probability measure $\mathbb{P}: \mathcal{B}(\mathcal{V}_{\mathcal{M}_V}) \rightarrow [0, 1]$ over the parameter space $\mathcal{V}_{\mathcal{M}_V}$.¹ We call the resulting model an *uncertain parametric MDP (upMDP)*.

Definition 9.1 (upMDP) An *uncertain parametric MDP (upMDP)* is a tuple $(\mathcal{M}_V, \mathbb{P})$, where $\mathcal{M}_V = (S, Act, s_I, V, P)$ is a pMDP and \mathbb{P} is a probability distribution over the parameter space $\mathcal{V}_{\mathcal{M}_V}$ of \mathcal{M}_V .

9
uncertain parametric MDP

If \mathcal{M}_V is a pMC, then we call the tuple $(\mathcal{M}_V, \mathbb{P})$ an *uncertain parametric MC (upMC)*.

Remark 9.2 (Probability measure \mathbb{P}) The probability measure \mathbb{P} used throughout this chapter should not be confused with the probabilistic operator in PCTL (see Def. 3.14). In this chapter, \mathbb{P} always denotes a probability measure over the parameter space of a pMDP.

Intuitively, a upMDP is a pMDP with an associated distribution over the parameter space. Sampling from the parameter space $\mathcal{V}_{\mathcal{M}_V}$ according to \mathbb{P} yields concrete MDPs $\mathcal{M}_V[u]$ with instantiations $u \in \mathcal{V}_{\mathcal{M}_V}$ (and $\mathbb{P}(u) > 0$).

Assumptions | Let us first make one key assumption about the distribution \mathbb{P} over the parameter instantiations. This assumption states that, although it is possible to draw

¹Recall from Sect. 2.3 that $\mathcal{B}(X)$ is the Borel σ -algebra over a set X . However, we will gloss over details about measurability in this chapter for brevity.

the same instantiation twice, the probability for this to happen is zero. For example, while it is possible to sample the exact same value twice from a Gaussian distribution, the probability for this to happen is zero.

Assumption 9.3 (Parameter distribution) Let $u, u' \in \mathcal{V}_{\mathcal{M}_V}$ be i.i.d. samples extracted according to \mathbb{P} . The probability that $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) = \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u')$ is zero.

Assumption 9.3 is needed for the theoretical results of the scenario approach to hold, which is the key technique we use to solve the problem in this chapter.

Problem | Let us formalize the problem that we aim to solve. Suppose we are given a upMDP $(\mathcal{M}_V, \mathbb{P})$ and a solution function $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$. We aim to compute the probability of sampling a parameter instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$, such that the solution $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u)$ is below a given threshold $\lambda \in [0, 1]$. We call this probability the *satisfaction probability*:

Definition 9.4 (Satisfaction probability) Let $(\mathcal{M}_V, \mathbb{P})$ be a upMDP, $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$ be a solution function, and $\lambda \in [0, 1]$ be a threshold for the solution. The *satisfaction probability* of $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$ with respect to the upper bound λ is

$$\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda\}.$$

Remark 9.5 (PCTL specification) Imposing a threshold λ on a (real-valued) reachability measure defines a specification whose interpretation is Boolean, analogous to the PCTL formulae introduced for MDPs in Chapter 3.

We will compute a *lower bound* on the satisfaction probability based on a finite set of i.i.d. samples from the parameter space $\mathcal{V}_{\mathcal{M}_V}$. Mathematically, this set of samples, denoted by $\mathcal{U}_N = \{u_1, \dots, u_N\}$, is an element of the product probability space $(\mathcal{V}_{\mathcal{M}_V}^N, \mathcal{B}(\mathcal{V}_{\mathcal{M}_V}^N), \mathbb{P}^N)$, where $\mathcal{V}_{\mathcal{M}_V}^N = \times_{i=1}^N \mathcal{V}_{\mathcal{M}_V}$ is the N -times Cartesian product of the parameter space, $\mathcal{B}(\mathcal{V}_{\mathcal{M}_V}^N)$ is its Borel σ -algebra, and \mathbb{P}^N is the product probability measure. We aim to construct an algorithm that, for at least a $\beta \in (0, 1)$ probability of \mathbb{P}^N , returns a valid lower bound on the satisfaction probability. In other words, for at least a $\beta \in (0, 1)$ probability of the sets \mathcal{U}_N we could sample from $(\mathcal{V}_{\mathcal{M}_V}^N, \mathcal{B}(\mathcal{V}_{\mathcal{M}_V}^N), \mathbb{P}^N)$, we obtain a lower bound $\eta \in [0, 1]$ on the satisfaction probability that is correct. Formally, this probably approximately correct (PAC) [HW93] lower bound on the satisfaction probability is defined as follows.

Problem 9.6 (Lower bound satisfaction probability) Let $(\mathcal{M}_V, \mathbb{P})$ be a upMDP, $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$ be a solution function, $\lambda \in [0, 1]$ be a threshold for the solution, and $\beta \in (0, 1)$ be a confidence probability. Compute a lower bound η on the satisfaction probability such that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{\mathcal{M}_V}^N : \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda\} \geq \eta \right\} \geq \beta.$$

Typically, we want a lower bound η on the satisfaction probability that is correct with

high probability, which we can achieve by choosing β close to one.

Example 9.7 Let us reconsider Fig. 9.1, which shows the solution function for the pMC (S, s_I, V, P, L) from Fig. 8.1 with an (unknown) Gaussian distribution \mathbb{P} over the parameter space. Consider again a threshold of $\lambda = 0.08$ on the satisfaction probability. For this upMC $(\mathcal{D}_V, \mathbb{P})$, Problem 9.6 asks to compute a lower bound η on the probability that the solution $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u)$ for a random $u \in \mathcal{V}_{\mathcal{M}_V}$ is below the threshold of $\lambda = 0.08$. We compute this lower bound η based on a finite set of N sampled parameter instantiations, and we aim to provide a confidence guarantee of at least β . That is, only for less than a $1 - \beta$ fraction of such sets of instantiations, our solution is allowed to be incorrect.

Variations | For brevity, this chapter focuses on the particular solution function $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$. However, our approach can directly be applied to any other solution function, such as with measures for minimizing schedulers ($\text{sol}_{\varphi, \min}^{\mathcal{M}_V}$) or cumulative expected rewards. Similarly, while Problem 9.6 considers an *upper bound* λ on the solution, we can easily consider *lower bounds* as well. We also discuss these variations in [2].

9.4 Bounding the Satisfaction Probability

In this section, we present a first step toward solving Problem 9.6. As we shall see, the approach in this section does not quite solve Problem 9.6, because we cannot control the threshold λ on the solution. Instead, the value of λ is determined by the solution to a convex optimization problem, which in turn depends on the parameter instantiation samples at hand. We will see a more appropriate solution to Problem 9.6 in Sect. 9.5.

9.4.1 Chance-constrained problem

First of all, let us recast the computation of the satisfaction probability as a chance-constrained optimization problem. Intuitively, we want to find the smallest value for λ , such that the probability of the solution being below this threshold is at least η . Such a pair of (λ, η) can be found by solving the following *chance-constrained optimization problem*, denoted by CCP:

$$\text{CCP : } \underset{\lambda \geq 0}{\text{minimize}} \quad \lambda \tag{9.1a}$$

$$\text{subject to } \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda\} \geq \eta. \tag{9.1b}$$

9

chance-constrained optimization problem

Observe that this optimization problem has a single decision variable, $\lambda \geq 0$. The constraint in Eq. (9.1b) is commonly called a *chance constraint*.

Let λ^* be an optimal solution to problem CCP for a fixed value of η . The satisfaction probability of the solution function $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}$ is at least λ^* with probability at least η , i.e.,

$$\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda^*\} \geq \eta.$$

Solving Eq. (9.1b) is very challenging in general, even if the probability measure \mathbb{P} is known. The main challenge is that the chance constraint allows us to neglect a $1 - \eta$ portion of the probability mass under \mathbb{P} , but there are many ways in which we can do

so. For example, for the Gaussian distribution in Fig. 9.1, we could try removing a $1 - \eta$ probability from the tails of the distribution. However, it is quite common that such a choice for removing probability mass leads to sub-optimal solutions, especially if the uncertainty domain is high-dimensional [LB02].

9.4.2 Scenario problem

Instead of trying to solve the chance-constrained problem CCP in Eq. (9.1), we propose a reformulation of the problem as a so-called *scenario optimization problem* [CG18a]. Intuitively, the idea is to replace the chance constraint in Eq. (9.1b) by the realization of this constraint for only $N \in \mathbb{N}_{>0}$ samples $u \in \mathcal{V}_{\mathcal{M}_V}$ for the parameter instantiation. Recall that we denote this set of samples by $\mathcal{U}_N = \{u_1, \dots, u_N\}$. Thus, we obtain the following optimization problem.

$$\text{SP}_N : \underset{\lambda \geq 0}{\text{minimize}} \quad \lambda \quad (9.2a)$$

$$\text{subject to } \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u_i) \leq \lambda \quad \forall i = 1, \dots, N. \quad (9.2b)$$

While solving the chance-constrained problem CCP is intractable, the scenario optimization problem SP_N is a simple linear program that can be solved readily with any standard convex optimization solver. The main question is, however, what guarantees a solution to problem SP_N provides on the probability that yet another parameter instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$ satisfies the constraints.

Let us denote an optimal point of SP_N by λ_N^* . Observe that the value of λ_N^* is a random variable whose value is determined by the samples \mathcal{U}_N . In fact, due to the structure of SP_N , this optimal point can be computed analytically as the maximum of the solutions for all sampled instantiations $u \in \mathcal{U}_N$, i.e.,

$$\lambda_N^* = \max_{i=1, \dots, N} \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u_i).$$

Because λ_N^* depends on the set of samples at hand, we can certainly not expect to obtain a perfect lower bound η on the satisfaction probability by solving SP_N . Instead, the theory of the *scenario approach* provides us with a *statistical guarantee* on the quality of the lower bound η in combination with the value of λ_N^* . For this theory to hold, we require the optimum to problem SP_N to exist and be unique, which is automatically satisfied due to Assumption 9.3.

Proposition 9.8 For every $N \in \mathbb{N}_{>0}$, the optimal point to problem SP_N exists and is unique.

Then, the scenario approach provides us with the following important result.

Theorem 9.9 ([CG08, Theorem 2.4]) Let λ_N^* be an optimal point to problem SP_N . Then, it holds that

$$\mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^* \right\} \geq \varepsilon \right\} = (1 - \varepsilon)^N. \quad (9.3)$$

Proof. [CG08, Theorem 2.4] states that under certain assumptions (which we discuss below), it holds that

$$\mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^* \right\} \geq \varepsilon \right\} = \sum_{i=0}^{d-1} \binom{N}{i} \varepsilon^i (1-\varepsilon)^{N-i},$$

where $d \in \mathbb{N}$ is the number of decision variables of the optimization problem. Since $d = 1$ for problem SP_N , we simplify the expression to Eq. (9.3). Next, the assumptions for this expression to hold are:

1. The optimal point to problem SP_N exists and is unique,
2. The scenario problem belongs to the class of *fully supported* optimization problems (see [CG08] for a formal definition).

Assumption (1) is satisfied due to Proposition 9.8. For our particular problem at hand, Assumption (2) boils down to requiring that, with probability 1, none of the constraints in problem SP_N are exactly the same, which is satisfied due to Assumption 9.3. This concludes the proof. ■

Intuitively, Theorem 9.9 can be explained as follows. Let λ_N^* be the optimal point to SP_N for a given set of samples \mathcal{U}_N . This optimal point λ_N^* is a random variable whose distribution is determined by \mathbb{P}^N . For at least a $(1 - \varepsilon)^N$ probability mass of the values λ_N^* we can obtain, we have that the probability for $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^*$ is at most ε , where $u \in \mathcal{V}_{\mathcal{M}_V}$ is sampled according to \mathbb{P} .

We slightly rewrite Theorem 9.9 to be more consistent with Problem 9.6.

Corollary 9.10 (Satisfaction prob. with sample-dependent threshold) Let λ_N^* be an optimal point to problem SP_N , and let $\beta \in (0, 1)$ be a confidence probability. Then, it holds that

$$\mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \right\} \geq (1 - \beta)^{1/N} \right\} = \beta. \quad (9.4)$$

Proof. Observe that $\mathbb{P} \{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^* \} + \mathbb{P} \{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \} = 1$. Thus, rewriting Eq. (9.3) yields

$$\begin{aligned} & \mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^* \right\} \geq \varepsilon \right\} \\ &= \mathbb{P}^N \left\{ 1 - \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \right\} \geq \varepsilon \right\} \\ &= \mathbb{P}^N \left\{ -\mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \right\} \geq \varepsilon - 1 \right\} \\ &= \mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \right\} \leq 1 - \varepsilon \right\} = (1 - \varepsilon)^N \end{aligned}$$

Flipping the sign of the outer inequality gives

$$\mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^* \right\} \geq 1 - \varepsilon \right\} = 1 - (1 - \varepsilon)^N$$

Let $\beta = 1 - (1 - \varepsilon)^N$, which implies $\varepsilon = 1 - (1 - \beta)^{1/N}$. Thus,

$$\mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^*\} \geq (1 - \beta)^{1/N} \right\} = \beta,$$

which concludes the proof. ■

Corollary 9.10 does not solve Problem 9.6 | It is tempting to conclude that Corollary 9.10 solves Problem 9.6. However, the subtle yet key difference is that in Problem 9.6, the threshold λ is fixed, while in Corollary 9.10, λ_N^* is a random variable, whose value depends on the sample set \mathcal{U}_N at hand. Thus, we cannot choose λ_N^* ourselves; instead, its value is determined by the sampling mechanism. We illustrate this important difference with the following example.

Example 9.11 We continue our motivating example, where we considered a pMC with the satisfaction probability for reaching state s_3 with probability at most $\lambda = 0.08$. Thus, Problem 9.6 asks for a lower bound $\eta \in [0, 1]$ such that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{\mathcal{M}_V}^N : \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq 0.08\} \geq \eta \right\} \geq \beta.$$

Now consider solving this problem using Corollary 9.10 with $N = 10$ parameter instantiation samples. Thus, we obtain that

$$\begin{aligned} & \mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \lambda_N^*\} \geq (1 - \beta)^{1/N} \right\} \\ &= \mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) \leq \max_{u \in \mathcal{U}_N} \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u)\} \geq (1 - \beta)^{1/N} \right\} = \beta. \end{aligned}$$

In other words, with probability of at least $(1 - \beta)^{1/N}$, the N^{th} solution is below the highest solution among all sampled instantiations $u \in \mathcal{U}_N$, and that claim holds

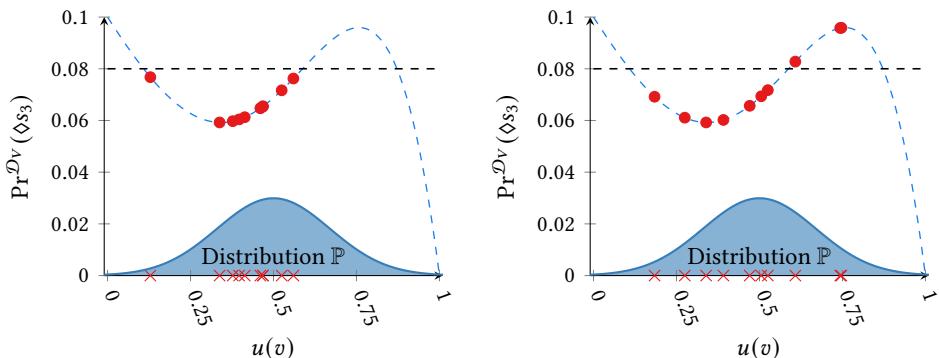


Figure 9.2: Two sets of $N = 10$ parameter instantiations $\mathcal{U}_N = \{u_1, \dots, u_{10}\}$ (shown as red crosses) and the corresponding solutions $\text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u)$ (shown as red points). The optimal point λ_N^* differs for these sets of samples.

for at least a $\beta \in (0, 1)$ probability among all sample sets $\mathcal{U}_N \in \mathcal{V}_{M_V}^N$. Two sets of $N = 10$ together with their corresponding values of λ_N^* are shown in Fig. 9.2. With the result from Corollary 9.10, we cannot control the value λ_N^* ourselves, which means that we cannot solve Problem 9.6 because it requires fixing λ upfront.

9.4.3 Sample complexity

Corollary 9.10 reveals an intricate relationship between the number of samples N , the confidence β , and the lower bound satisfaction probability $\eta = (1 - \beta)^{1/N}$. Given that β is close to one, we can thus obtain a high confidence in the lower bound η on the satisfaction probability. This high confidence is easily achieved for a sufficiently large number of samples N , as seen from the following corollary.

Corollary 9.12 (Required sample size) The sample size N necessary to obtain a desired lower bound $\eta \in (0, 1)$ on the satisfaction probability with at least a confidence of $\beta \in (0, 1)$ is

$$N = \left\lceil \frac{\log(1 - \beta)}{\log \eta} \right\rceil,$$

where $\lceil x \rceil$ denotes the ceil function, i.e., the function which rounds its argument $x \in \mathbb{R}$ upwards to the nearest integer.

Corollary 9.12 states that the sample size N is logarithmic in the confidence probability β . Thus, a significant improvement in β (i.e., closer to one) only requires a marginal increase in N . Similarly, increasing the sample size N improves the lower bound on the satisfaction probability η . For example, applying Corollary 9.10 with $N = 10$ samples yields that the satisfaction probability is lower bounded by 0.794 (with a confidence of at least $\beta = 0.9$) and by 0.631 (with a confidence of at least $\beta = 0.99$). When increasing the number of samples to $N = 100$, we obtain improved lower bounds of 0.977 (for $\beta = 0.9$) and 0.955 (for $\beta = 0.99$).

Next, consider the extreme case of β infinitely close to one. Observe from Corollary 9.12 that such a confidence probability can only be obtained for $N = \infty$. Intuitively, this observation makes sense: We can only be absolutely certain of our lower bound on the satisfaction probability if we have based this estimate on infinitely many samples. In practice, our sample set is finite, and a typical confidence probability is $\beta = 1 - 10^{-3}$.

9.5 Improving Bounds by Discarding Samples

We return to a setting with a fixed threshold λ on the value of the solution function, as we used to formulate Problem 9.6. We again formulate a scenario optimization problem similar to Eq. (9.2). However, instead of enforcing the constraint $\text{sol}_{\varphi, \max}^{M_V}(u) \leq \lambda$ for all samples $u \in \mathcal{U}_N$, we only enforce this constraint for a *subset of samples*. In other words, we *discard* a portion of the samples, which allows us to obtain a less conservative lower bound on the satisfaction probability [CG11]. In what follows, we first formulate the scenario optimization problem and thereafter discuss how we solve Problem 9.6 based on the scenario approach theory.

9.5.1 Scenario problem with discarded samples

Again, we will use a set of N i.i.d. samples from the parameter space $\mathcal{V}_{\mathcal{M}_V}$. In this section, we make the additional assumption that these samples are sorted in *ascending order*. We denote the resulting set of samples by \mathcal{U}_N^\uparrow to indicate that the samples are in ascending order.

Assumption 9.13 (Solutions in ascending order) The set of samples $\mathcal{U}_N^\uparrow = \{u_1, \dots, u_N\}$ is sorted in ascending order, i.e., $u_1 \leq u_2 \leq \dots \leq u_N$.

In practice, after obtaining the N i.i.d. samples from the parameter space, we sort them in ascending order. Sorting the samples will be beneficial for notational purposes but does not affect the theoretical aspects of our approach.

For a fixed value of $k < N$ and the set of samples \mathcal{U}_N^\uparrow (in ascending order), consider the following optimization problem, denoted by $\text{SP}_{N,k}$:

$$\text{SP}_{N,k} : \underset{\lambda \geq 0}{\text{minimize}} \quad \lambda \quad (9.5a)$$

$$\text{subject to } \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u_i) \leq \lambda \quad \forall i = 1, \dots, N - k. \quad (9.5b)$$

We denote the optimal solution to problem $\text{SP}_{N,k}$ by $\lambda_{N,k}^*$. Observe that problem $\text{SP}_{N,k}$ is a relaxation of problem SP_N in Eq. (9.2), where the constraints for k samples with the highest solutions are discarded. Because the samples are sorted, the optimal point to problem $\text{SP}_{N,k}$ is obtained analytically as

$$\lambda_{N,k}^* = \max_{i=1, \dots, N-k} \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u_i).$$

The following result from [CG11] forms the basis of how we solve Problem 9.6 based on the solution to the scenario problem $\text{SP}_{N,k}$ with k discarded constraints.

Theorem 9.14 ([CG11, Theorem 2.1]) Fix $k < N$, let $\lambda_{N,k}^*$ be an optimal point to problem $\text{SP}_{N,k}$, and let β be a confidence probability. Then, it holds that

$$\mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^*\} > \varepsilon \right\} \leq \beta. \quad (9.6)$$

where ε is such that

$$\beta = \sum_{i=0}^k \binom{N}{i} (\varepsilon)^i (1 - \varepsilon)^{N-i}.$$

Proof. [CG11, Theorem 2.1] states that under certain assumptions (which we discuss below), it holds that

$$\mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi, \max}^{\mathcal{M}_V}(u) > \lambda_N^*\} > \varepsilon \right\} \leq \binom{k+d-1}{k} \sum_{i=0}^{k+d-1} \binom{N}{i} \varepsilon^i (1 - \varepsilon)^{N-i},$$

where $d \in \mathbb{N}$ is the number of decision variables of the optimization problem. Plugging in $d = 1$, we obtain the expression in Eq. (9.6). Next, the main assumption (in addition

to those in the proof of Theorem 9.9) for this expression to hold is that all k discarded samples are violated by the optimal point $\lambda_{N,k}^*$, i.e.,

$$\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u_i) > \lambda_{N,k}^* \quad \forall i = N - k + 1, \dots, N.$$

Observe that the solutions $\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u_i)$ for $k = 1, \dots, N - k$ still left in problem $\text{SP}_{N,k}$ are, with probability 1, strictly smaller than the discarded solutions $\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u_i)$ for $k = N - k + 1, \dots, N$. Thus, this assumption is satisfied by construction, which concludes the proof. ■

Theorem 9.14 asserts that, for any fixed $k = 0, \dots, N - 1$, with a probability of at most $\sum_{i=0}^k \binom{N}{i} \varepsilon^i (1 - \varepsilon)^{N-i}$, the probability that $\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) > \lambda_{N,k}^*$ for the next sampled instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$ is at least ε . We emphasize that the value of k in Theorem 9.14 is chosen *a priori*, i.e., before observing the sampled parameter instantiations. Thus, it would *not* be correct to choose k based on the values of the solutions for the samples at hand. Instead, we must fix k independent of these solutions.

9.5.2 Problem 9.6 solved

We discuss how we can leverage Theorem 9.14 to solve Problem 9.6. For this, we need to modify the theorem such that we can choose the number of samples to discard *a posteriori*, i.e., after observing the sampled parameter instantiations. The trick is that we apply Theorem 9.14 not for a single value of k , but *for all values of $k = 0, \dots, N - 1$ simultaneously*. Accordingly, we must also modify the overall confidence probability to acknowledge that we apply the theorem for N different values of k . Formally, we obtain the following result.

Corollary 9.15 (Satisfaction prob. with fixed threshold) Let $\lambda \in [0, 1]$ be a desired threshold and let $\hat{\beta} \in (0, 1)$ be a confidence probability. For the set of samples $\mathcal{U}_N^\uparrow = \{u_1, \dots, u_N\}$ sorted in ascending order, let $M \leq N$ be the number of samples from \mathcal{U}_N^\uparrow that violate this threshold:

$$M := \left| \{u \in \mathcal{U}_N^\uparrow : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) > \lambda\} \right|.$$

Then, it holds that

$$\mathbb{P}^N \left\{ \mathbb{P} \left\{ u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda \right\} \geq 1 - \varepsilon_M \right\} \geq \hat{\beta}, \quad (9.7)$$

where $\varepsilon_N = 1$ (all N samples are violated), and otherwise (if $0 \leq M < N$), ε_M is the solution of

$$\frac{1 - \hat{\beta}}{N} = \sum_{i=0}^M \binom{N}{i} (\varepsilon_M)^i (1 - \varepsilon_M)^{N-i}. \quad (9.8)$$

Proof. Observe that $\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) > \lambda_{N,k}^*\} + \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} = 1$. Fix $\beta = \frac{1-\hat{\beta}}{N}$ in Eq. (9.6). Then, we have for all $k = 0, \dots, N-1$ that

$$\begin{aligned} \mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) > \lambda_{N,k}^*\} > \varepsilon_k \right\} &\leq \frac{1-\hat{\beta}}{N} \\ \mathbb{P}^N \left\{ 1 - \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} > \varepsilon_k \right\} &\leq \frac{1-\hat{\beta}}{N} \\ \mathbb{P}^N \left\{ -\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} > \varepsilon_k - 1 \right\} &\leq \frac{1-\hat{\beta}}{N} \\ \mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} < 1 - \varepsilon_k \right\} &\leq \frac{1-\hat{\beta}}{N} \\ \mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} \geq 1 - \varepsilon_k \right\} &\geq 1 - \frac{1-\hat{\beta}}{N}, \end{aligned}$$

where each ε_k , $k = 0, \dots, N-1$ is such that Eq. (9.8) holds. Via Boole's inequality (i.e., the union bound), we obtain

$$\mathbb{P}^N \left\{ \bigcap_{i=0}^{N-1} \left[\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda_{N,k}^*\} \geq 1 - \varepsilon_k \right] \right\} \geq 1 - \frac{N(1-\hat{\beta})}{N} = \hat{\beta},$$

In other words, with probability at least $\hat{\beta} \in (0, 1)$, the lower bounds on the satisfaction probability with respect to $\lambda_{N,k}^*$ for all $k = 0, \dots, N-1$ hold simultaneously.

Now consider solving Problem 9.6 for a desired threshold of λ , i.e., we want to compute a lower bound on the probability that

$$\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda\}.$$

We distinguish two cases for the number of samples M violating the threshold λ :

1. If $M = N$, then all samples violate the threshold λ , and we invoke the trivial solution of $\varepsilon_M = 1$, which is true by definition;
2. If $0 \leq M < N$, then we have that $\lambda_{N,M}^* \leq \lambda$, so we obtain

$$\mathbb{P}^N \left\{ \mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \leq \lambda\} \geq 1 - \varepsilon_M \right\} \geq \hat{\beta},$$

where ε_M is such that Eq. (9.8) holds.

These cases cover the cases from the corollary, so we conclude the proof. ■

Observe that Corollary 9.15 provides (with a confidence probability of at least $\hat{\beta}$) a lower bound on the satisfaction probability with respect to a fixed threshold λ . Hence, we can use Corollary 9.15 to solve Problem 9.6.

Beta distribution | We note that Eq. (9.8) is the cumulative distribution function of a beta distribution, which can easily be solved numerically for ε , $M = 0, \dots, N-1$; see, e.g., [CG18a] for details. Moreover, we can speed up the computations at run-time by tabulating the solutions to Eq. (9.8) for all relevant values of N , β , and M up front.

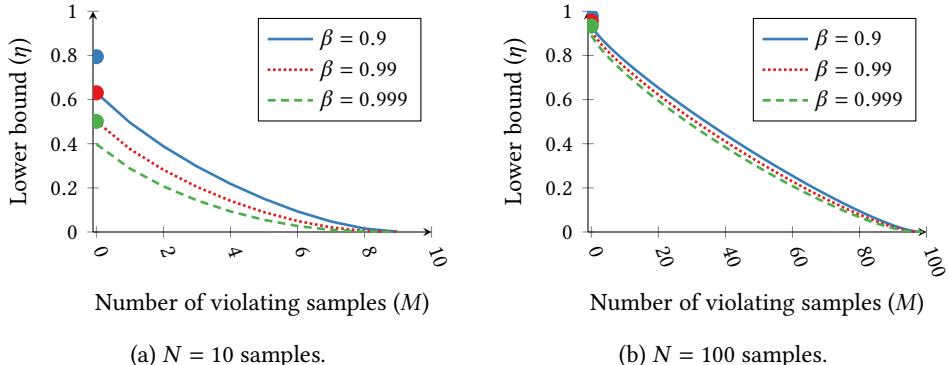


Figure 9.3: Lower bounds on the satisfaction probability obtained from Corollary 9.10 (shown as points at $M = 0$) and Corollary 9.15 for $M = 0, \dots, N$ (lines).

Tightness of lower bounds | Fig. 9.3 illustrates how the number of violating samples, M , influences the quality of the lower bound ε_M on the satisfaction probability. The points at $M = 0$ are the bounds returned by Corollary 9.10, while the lines correspond to Corollary 9.15. Intuitively, the lower bound on the satisfaction probability computed by Corollary 9.15 decreases with an increased number of violating samples. Moreover, Corollary 9.10 yields a better lower bound than Corollary 9.15 (points versus the lines in Fig. 9.3), at the cost of not using a fixed threshold on the solution function, and not being able to deal with violating samples.

In Fig. 9.4, we fix the fraction of violating samples M/N and plot the lower bounds on the satisfaction probability obtained using Corollary 9.15 for different values of N and β . Note that the lower bounds grow toward the fraction of violation for increased sample sizes. As also shown with Corollary 9.12, the confidence probability β only has a marginal effect on the obtained lower bounds.

Finally, we make the following remark with respect to the sample complexity of Corollaries 9.10 and 9.15.

Remark 9.16 (Independence to model size) The number of samples needed to obtain a certain confidence probability in Corollaries 9.10 and 9.15 is independent of the number of states, transitions, or parameters of the upMDP. Despite this independence, note that the time to compute solutions via model checking still depends on the number of states and transitions of the instantiated MDP.

9.6 Experimental Evaluation

We implemented our approach in Python using the model checker Storm [DJKV17] to construct and analyze samples of MDPs. To demonstrate the effectiveness of our approach, we perform experiments on a variety of benchmarks with different numbers of states and parameters. We use our results to compute PAC lower bounds on the satisfaction probabilities of several specifications.²

²Recall from Remark 9.5 that a specification combines a measure of, e.g., reachability, with a threshold λ on this measure.

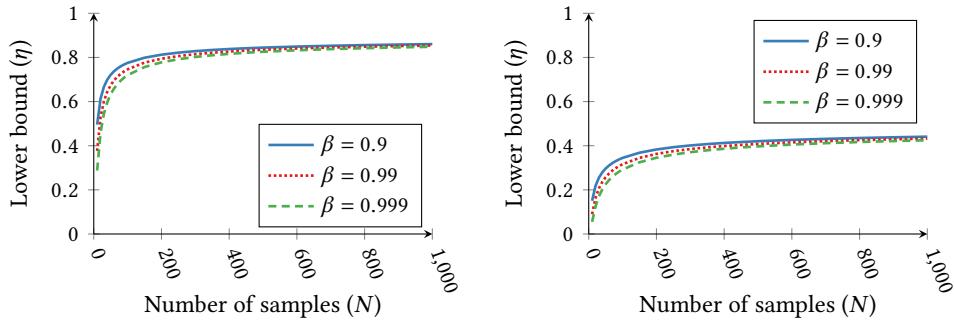


Figure 9.4: Lower bounds on the satisfaction probability obtained from Corollary 9.15 for fixed fractions M/N of violated samples.

Reproducibility | Our implementation is available at <https://doi.org/10.5281/zenodo.6674059>. All of our experiments ran on a computer with 32 3.7 GHz cores and 64 GB of RAM.

Overview of experiments | First, we apply our method to a dedicated UAV motion planning benchmark. Thereafter, we report on a set of well-known benchmarks used in parameter synthesis [JÁHJ⁺24]. These benchmarks are, for instance, available on the website of the tool PARAM [HHZ11b] or part of the PRISM benchmark suite [KNP12]. In this chapter, we focus on computing a lower bound η on the satisfaction probability for a fixed confidence probability β . In [2], we show that we can also do the opposite, i.e., specify a lower bound η and compute the corresponding confidence probability β .

9.6.1 UAV Motion Planning

We apply our method to a UAV motion planning problem. This benchmark originates from [CJJK⁺20] and the later journal version [2] on which this chapter is based.

Model | Consider the UAV shown in Fig. 9.5 that needs to transport a payload from one end of a valley to the green target region, while avoiding the red obstacles. We model the dynamics of the UAV as an MDP, whose state space represents a discretized version of the valley. The states encode the position of the UAV, the current weather conditions (sunny, stormy), and the general wind direction in the valley. The actions model the possible movements of the UAV between the cells of the grid, and the outcome of each action is stochastic due to the influences of the environment. We thus obtain a *grid world* model in which the UAV can decide to fly in either of the six cardinal directions (N, W, S, E, up, down). Upon executing an action, the wind moves the UAV one cell in the wind direction with a probability p , which depends on the wind speed. Furthermore, we assume that the weather and wind conditions change during the day and are described by a stochastic process.

Distribution over weather conditions | The transition probabilities of the MDP are not fixed but rather a function of the weather. Thus, the model is an upMDP whose transition probabilities depend on the weather. Concretely, parameters describe how the

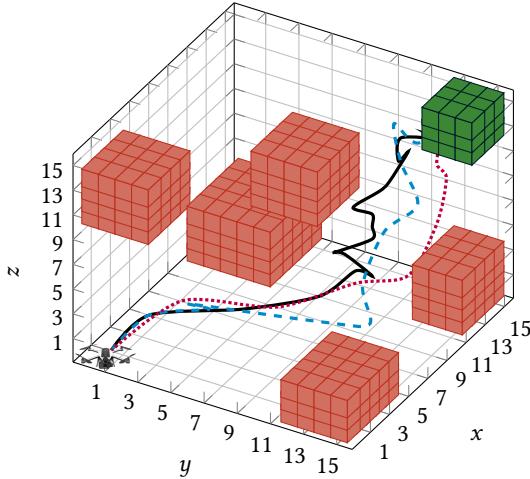


Figure 9.5: The UAV benchmark with obstacles (red boxes) and a target area (green box).

The three trajectories are for a uniform wind distribution (solid line), stronger northbound wind (dashed line), and stronger westbound wind (dotted line).

weather affects the UAV in different zones of the valley, and how the weather/wind may change during the day. In total, the model has 900 parameters. From historical weather data, we can derive a distribution over these parameters. Specifically, we consider the following three cases for the weather conditions:

1. a uniform distribution over the different weather conditions in each zone;
2. the probability for a weather condition inducing a wind direction that pushes the UAV northbound (i.e., into the positive y -direction) is twice as likely as in other directions;
3. it is twice as likely to push the UAV westbound (i.e., into the negative x -direction).

We consider the solution function $\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) : u \mapsto \Pr_{\sigma}^{\mathcal{M}_V[u]}(s_I \models \diamond T)$, i.e., the probability of reaching the target set $T \subset S$ (while avoiding the obstacles, which are modeled as sink states). We consider a lower bound on this solution function of $\lambda = 0.9$. Intuitively, $\text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \geq 0.9$ for a parameter instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$ means that reachability probability is at least 90%. Thus, the satisfaction probability $\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \geq 0.9\}$ is the probability of drawing a parameter instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$, such that the resulting MDP has a reachability probability of at least 90%. Hence, Problem 5.2 asks, with probability at least β , for a lower bound η on the satisfaction probability $\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) \geq 0.9\}$. i.e., the probability that a parameter instantiation $u \in \mathcal{V}_{\mathcal{M}_V}$ sampled according to the distribution \mathbb{P} leads to a solution of at least 0.9.

Satisfaction vs. unsatisfaction | Recall that we can also compute a lower bound on the probability of *not* satisfying the specification. To do so, we consider λ as a *lower bound* (instead of an upper bound) on the solution function. Thus, the *unsatisfaction probability* for this benchmark is $\mathbb{P}\{u \in \mathcal{V}_{\mathcal{M}_V} : \text{sol}_{\varphi,\max}^{\mathcal{M}_V}(u) < 0.9\}$. To compute this

Table 9.1: Lower bounds η on the (un)satisfaction probability for the UAV benchmark with $N = 5\,000$ samples.

Weather condition	$\beta = 0.9$		$\beta = 0.99$		$\beta = 0.999$		$\beta = 0.9999$	
	η , sat	η , unsat	η , sat	η , unsat	η , sat	η , unsat	η , sat	η , unsat
1. Uniform wind	0.91138	0.05830	0.90928	0.05670	0.90735	0.05528	0.90555	0.05398
2. Northbound wind	0.77878	0.17483	0.77577	0.17217	0.77302	0.16978	0.77048	0.16760
3. Westbound wind	0.77680	0.17664	0.77378	0.17397	0.77103	0.17157	0.76847	0.16938

unsatisfaction probability, we basically substitute all maximizations with minimizations and flip all inequalities throughout our approach. For details, we refer to our paper [2].

Trajectories | We depict example trajectories of the UAV under the optimal policies for these three cases in Fig. 9.5. The trajectory depicted by the black line represents a simulated trajectory for the first case (uniform distribution), taking a direct route to reach the target area. Similarly, the trajectories shown by the dashed blue and dotted purple lines are simulated trajectories for the second (stronger northbound wind, i.e., positive x -direction) and third cases (stronger westbound wind, i.e., positive y -direction), respectively. Under these two weather conditions, the UAV takes different paths toward the goal to account for the stronger wind. In particular, in the case of westbound wind (case 3), we observe that the UAV flies close to the obstacle at the right bottom. By contrast, in the case of northbound wind (case 2), the wind may push the UAV into the obstacles, so we observe that the UAV avoids getting close to the obstacles while flying toward the target area.

Bounds on satisfaction probabilities | We sample $N = 1\,000$ parameter values for each case and consider different confidence probabilities β between 0.9 and 0.9999. For all three weather conditions, we use Corollary 9.15 to compute the lower bounds η on (a) the probability of satisfying the reachability specification defined above, and (b) the probability of not satisfying the specification. Specifically, this lower bound η is computed as the value of $(1 - \varepsilon_M)$ in Corollary 9.15, where M is the number of samples that violate the specification.

The results are presented in Table 9.1. The highest lower bound on the satisfaction probability is obtained for the first weather condition, and is $\eta = 0.911$ (for $\beta = 0.9$) and $\eta = 0.906$ (for $\beta = 0.9999$). In other words, under a uniform distribution over the weather conditions, the UAV will (with a confidence of at least $\beta = 0.9999$) satisfy the specification on at least 90.6% of the days. The second and third weather conditions lead to lower bounds of $\eta = 0.770$ and $\eta = 0.768$ (for $\beta = 0.9999$), respectively, showing that it is harder to navigate around the obstacles with non-uniform probability distributions over the parameters. The average time to run our approach on this upMDP with 900 parameters and around 10 000 states (i.e., performing the sampling, model checking, and computing the lower bounds η) with 5 000 parameter samples is 9.5 minutes.

9.6.2 Parameter Synthesis Benchmarks

Setup | In our second set of benchmarks, we adopt pMDPs and pMCs from [QDJJ⁺16]. We adapt the *Consensus* protocol [AH90] and the *Bounded Retransmission Protocol* (BRP) [HSV93; DJJL01] to upMDPs; the *Crowds Protocol* (CROWDS) [Shm04] and the *NAND*

Table 9.2: Information for the benchmark instances and the approximate (un)satisfaction probabilities taken from [QDJJ⁺16].

Benchmark	Measure	λ	Specification			Model size		Approximation	
			Pars.	States	Trans.	Sat. prob.	Unsat. prob.		
BRP	(256,5) Pr($\Diamond T$)	≤ 0.5	2	19 720	26 627	0.055	0.898		
	(16,5) ExpRew($\Diamond T$)	≤ 3	4	1 304	1 731	0.275	0.676		
	(32,5) ExpRew($\Diamond T$)	≤ 3	4	2 600	3 459	0.232	0.718		
CROWDS	(10,5) Pr($\Diamond T$)	≤ 0.9	2	104 512	246 082	0.537	0.413		
	(15,7) Pr($\Diamond T$)	≤ 0.9	2	8 364 409	25 108 729	0.411	0.539		
NAND	(10,5) Pr($\Diamond T$)	≥ 0.5	2	35 112	52 647	0.218	0.733		
	(25,5) Pr($\Diamond T$)	≥ 0.5	2	865 592	1 347 047	0.206	0.744		
CONSENSUS	(2,2) Pr($\Diamond S_T$)	≥ 0.25	2	272	492	0.280	0.669		
	(4,2) Pr($\Diamond S_T$)	≥ 0.25	4	22 656	75 232	0.063	0.888		

Multiplexing benchmark (NAND) [HJ02] become uMCs. Table 9.2 lists the specifications checked and the number of parameters, states, and transitions for all benchmarks. For most benchmarks, we consider reachability measures $\text{Pr}(\Diamond T)$, but for BRP (16,5) and (32,5) we consider expected cumulative reward measures $\text{ExpRew}(\Diamond T)$. Depending on the benchmark, we consider the threshold λ as an upper or lower bound on the measure of interest. For all benchmarks, we assume a uniform distribution over the parameters.

Approximative baseline | Essentially, the parameter lifting algorithm (PLA) from [QDJJ⁺16] allows approximating the size of the satisfying (or unsatisfying) regions of the parameter space. Thus, PLA approximates the satisfaction probability for a uniform distribution over the parameter space. We use PLA as a baseline to approximate the satisfaction probability (or the unsatisfaction probability) of the specifications, for which the results are shown in Table 9.2. We provide these numbers as a baseline only: the computation via PLA cannot scale to more than tens of parameters [QDJJ⁺16] and cannot cope with unknown distributions.

Specifications with variable thresholds λ | For benchmark BRP (16,5), we demonstrate how Corollary 9.10 can be used to compute a lower bound on the satisfaction probability with respect to the expected reward measure $\text{ExpRew}(\Diamond T)$ and the variable threshold λ_N^* . More intuitively, this means we compute a lower bound on the probability that the $(N + 1)^{\text{th}}$ solution is at most the maximum of the first N solutions, which is $\lambda_N^* = \max_{i=1,\dots,N} \text{sol}_{\varphi, \text{max}}^{M_V}(u_i)$. We fix a confidence probability of $\beta = 0.99$ and use either $N = 1\,000$ or $10\,000$ parameter instantiation samples. The resulting lower bounds on the satisfaction probability (which only depend on N and β) are $\eta = 0.9954$ (for $N = 1\,000$) and $\eta = 0.9995$ (for $N = 10\,000$). However, the corresponding value of λ_N^* is a random variable that depends on the set of samples at hand. In the histogram in Fig. 9.6, we plot the distribution over the values of λ_N^* when we repeat this experiment many times. Our results confirm that Corollary 9.10 can be used to lower bound the satisfaction probability when the threshold on the specification does not need to be fixed.

Solving Problem 9.6 | We show how to solve Problem 9.6 for the considered benchmarks. Recall that this problem assumes a fixed threshold $\lambda \in [0, 1]$ on the solution

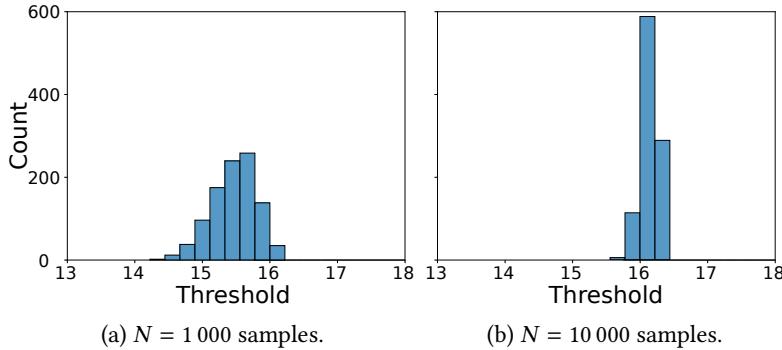


Figure 9.6: Histograms of the obtained thresholds λ_N^* on BRP (16,5).

function and a confidence probability $\beta \in (0, 1)$, and asks to compute a lower bound η on the satisfaction probability such that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{M_V}^N : \mathbb{P}\{u \in \mathcal{V}_{M_V} : \text{sol}_{\varphi, \max}^{M_V}(u) \leq \lambda\} \geq \eta \right\} \geq \beta.$$

We apply Corollary 9.15 to compute such a lower bound η for a fixed λ and β . Concretely, the lower bound η is obtained as $1 - \varepsilon_M$ in Corollary 9.15, where M is the number of samples in $\mathcal{U}_N^\uparrow = \{u_1, \dots, u_N\}$ that violate the threshold λ . For each benchmark, we sample $N = 25\,000$ values for the parameters and apply Corollary 9.15 for increasing confidence probabilities β . We use the values for λ reported in Table 9.2.

We report the resulting bounds η in Table 9.3. We observe that the obtained values of η are slightly more conservative (i.e., lower) for higher values of β . This observation is indeed intuitive: To reduce the $1 - \beta$ probability of obtaining an incorrect bound on the (un)satisfaction probability, the value of η must be more conservative. Moreover, increasing the confidence probability β only marginally reduces the obtained lower bound η . For example, the obtained lower bound on the satisfaction probability for BRP (256,5) with $\beta = 0.9$ is $\eta = 0.07244$, while for $\beta = 0.9999$, it is only reduced to $\eta = 0.07036$ (a reduction of only 0.21%). This observation confirms the important result of Corollary 9.12: A high confidence probability β can typically be obtained without sacrificing the tightness of the obtained lower bound η .

Finally, we remark that the confidence probability β can often be improved significantly by only marginally reducing the lower bound η . Consider, for example, the results for NAND (10,5) in Table 9.3. For $\beta = 0.99$, we obtain a lower bound of 0.23783 on the probability of satisfying the specification, while for a (significantly higher) confidence of $\beta = 0.9999$, we obtain a lower bound of 0.23561. In other words, by weakening our lower bound η by an almost negligible amount of 0.002, we increase the confidence in the results from 99% to a remarkable 99.99%. This result highlights that a very high confidence probability β can be achieved in practice, without hurting the lower bound η significantly—an observation that has been made often in the scenario approach literature [CG18a; CCG21].

Table 9.3: Lower bounds η on the (un)satisfaction probability for $N = 25\,000$ samples.

Benchmark	$\beta = 0.9$		$\beta = 0.99$		$\beta = 0.999$		$\beta = 0.9999$		
	η , sat	η , unsat	η , sat	η , unsat	η , sat	η , unsat	η , sat	η , unsat	
BRP	(256,5)	0.07244	0.91221	0.07168	0.91135	0.07099	0.91056	0.07036	0.90982
	(16,5)	0.28787	0.68619	0.28653	0.68481	0.28531	0.68353	0.28417	0.68234
	(32,5)	0.24356	0.73176	0.24229	0.73044	0.24113	0.72922	0.24005	0.72808
CROWDS	(10,5)	0.55106	0.42091	0.54957	0.41945	0.54821	0.41810	0.54695	0.41685
	(15,7)	0.42397	0.54798	0.42250	0.54650	0.42115	0.54514	0.41990	0.54387
NAND	(10,5)	0.23909	0.73637	0.23783	0.73506	0.23668	0.73384	0.23561	0.73271
	(25,5)	0.20979	0.76673	0.20858	0.76546	0.20748	0.76430	0.20647	0.76321
CONSENSUS	(2,2)	0.29383	0.68009	0.29248	0.67870	0.29125	0.67742	0.29010	0.67622
	(4,2)	0.07367	0.91086	0.07291	0.91000	0.07221	0.90921	0.07157	0.90846

Computing β for a given η | Conversely, we can also compute the confidence probability that can be guaranteed for a fixed lower bound η on the satisfaction probability. However, we omit these results from this thesis for brevity, and we instead refer the interested reader to [2].

9.7 Discussion

We close this chapter with a brief discussion of related work.

Uncertainty in MDPs | MDPs with uncertain transition probabilities, such as the RMDPs and IMDPs discussed in Chapter 3, have received a lot of attention in the artificial intelligence and planning literature [PLSS13; GLD00; WTM12]. However, as already discussed in Chapter 3, most tractable solutions assume that these uncertain transition probabilities are independent across different states and actions of the MDP. Some exceptions that relax this independence assumption exist; however, state-of-the-art exact solution methods are limited to only a few hundred states [HPW18]. Multi-model MDPs [SKD21] treat distributions over probability and cost parameters and aim at finding a single strategy maximizing a weighted value function. This problem is NP-hard for deterministic strategies and PSPACE-hard for history-dependent strategies.

Distributions over parameters | Considering distributions over parameters of pMDPs has been proposed a couple of times in the literature. The authors of [BHL19] consider the analysis of Markov models in the presence of uncertain rewards, utilizing statistical methods to reason about the probability of a pMDP satisfying an expected cost specification. This approach is restricted to reward parameters and does not explicitly compute confidence bounds. The work in [PWHA16] obtains data-driven bounds on the parameter ranges and then uses parameter synthesis techniques to validate specifications for all parameter values in this range. Paper [LBBS⁺18] computes bounds on the long-run probability of satisfying a specification with probabilistic uncertainty for Markov chains. Other related techniques include multi-objective model checking to maximize the average performance with probabilistic uncertainty sets [SBHH17], and sampling-based methods which minimize the *regret* with uncertainty sets [AVLA⁺17]. Finally, the approach in [ABCK⁺18] considers a variant of the problem in this paper where parameter values cannot be observed and thus must be learned. The latter pa-

per formulates the strategy synthesis problem as a computationally harder partially observable Markov decision process (POMDP) synthesis problem and uses off-the-shelf point-based POMDP methods [PGT03; CLZ97].

Observability of parameters | In this chapter, we considered a setting where policies can depend on the values of the parameters. Thus, we made the implicit assumption that the parameter values are *observed* before computing a policy. Whether this assumption is appropriate depends on the context. Other authors have investigated the setting where the parameter values *cannot be observed*, such that a *single policy* must be computed that is robust against the distribution. This setting has recently been considered by [RAM23; SAP24] for upMDPs, and also relates to so-called Bayes-adaptive (PO)MDPs which are commonly used in reinforcement learning [CRLH23; RLH21; GSD12; ÇOK22]. A variant where parameter values cannot be observed and thus must be learned has been studied by [ABCK⁺18]. The setting with unobservable parameters requires a single policy for all parameters and is, thus, arguably more challenging. In that setting, sampled data from the solution function is not i.i.d., which is a requirement for our approach. On the other hand, however, our approach scales significantly better than those that consider unobservable parameters.

Variants for other models | While we focused on pMDPs in this chapter, similar approaches have been proposed for continuous-time Markov chains (CTMCs) with a distribution over the (parametric) transition rates. For example, the authors of [BS18] use Bayesian reasoning to compute parameter values that satisfy a metric temporal logic specification on a CTMC. CTMCs with a distribution over the transition rates is the topic of Chapter 12 of this thesis. In that chapter, which is based on [3], we develop a sampling-based method similar to the one proposed in this chapter for verifying CTMCs with distributions over the transition rates.

Future research | An interesting direction for future work is to adapt our approach for more involved models such as Markov automata [HH12]. Another potential line of future work is to integrate our approach within a parameter synthesis framework for parametric Markov models [JJK22].

Summary

- ⇒ We studied parametric MDPs with a distribution over parameter instantiations, which we call an uncertain parametric MDP (upMDP).
- ⇒ The distribution of the parameter instantiations is unknown, and instead, we only assumed access to a finite set of samples of the parameter.
- ⇒ We have developed a sampling-based method based on the scenario approach to compute PAC-style lower bounds on the satisfaction probability of a given temporal logic specification.
- ⇒ In practice, several thousand samples suffice to obtain tight and high-confidence verification results.

10 Sensitivity Analysis for Parametric Markov Chains

Summary | We introduce parametric robust MCs (pRMCs) as the unification of parametric and robust Markov chains. These models incorporate both parameters and sets of probability distributions to alleviate the often unrealistic assumption that precise probabilities are available. In this chapter, we present an efficient method to compute partial derivatives of the solution function for a pRMC, with respect to its parameters. These partial derivatives measure the sensitivity of the solution function to changes in the parameters. Our method is based on linear programming and differentiating these programs around a given value for the parameters. We show that our approach scales to models with over a million states and thousands of parameters. Moreover, using these partial derivatives to guide sampling in an iterative learning scheme improves the sample efficiency over other sampling strategies.

Origins | This chapter is based on the following publication:

[4] Badings, Junges, Marandi, Topcu and Jansen ‘Efficient Sensitivity Analysis for Parametric Robust Markov Chains’. CAV.

In this paper, we (1) present the first algorithm to compute partial derivatives for pRMCs, (2) develop an efficient method to determine a subset of parameters with the highest derivatives, and (3) apply our methods in an iterative learning scheme.

Background | We assume familiarity with Markov chains and their parametric variant, which we discussed in Chapters 3 and 8, respectively. In particular, we prominently use solution functions for parametric Markov chains (pMCs).

10.1 Introduction

Classical Markov chains assume that all probabilities are precisely known—an assumption that is difficult, if not impossible, to satisfy in many applications [8]. In Sect. 3.3, we have seen how robust MCs (RMCs) alleviate this assumption by using *sets of probability distributions*, e.g., intervals of probabilities in the simplest case [JL91; BGN09]. A typical verification problem for RMCs is to compute upper or lower bounds on measures of interest, such as the expected cumulative reward, under *worst-case realizations* of these probabilities in the set of distributions [WTM12; PLSS13]. Thus, these values are *robust* against any selection of probabilities in these sets.

Where to improve my model? | As an example, consider a ground vehicle navigating toward a target location in an environment with different terrain types. On each terrain type, there is some probability that the vehicle will slip and fail to move. Suppose

we obtain a sufficient number of *samples* to infer upper and lower bounds (i.e., intervals) on the slipping probability on each terrain, which we use to model the environment as an RMC. However, from the RMC, it is unclear how our model (and thus the measure of interest) will change if we obtain more samples. For instance, if we take *one more sample* for a particular terrain, some intervals of the RMC will change, but how can we expect the verification result to change? And if the verification result is unsatisfactory, for which terrain type should we obtain more samples?

Parametric robust MCs | To reason about how additional samples will change our model and thus the verification result, we employ a sensitivity analysis [FTG16]. To that end, we use *parametric robust MCs* (*pRMCs*), which are RMCs whose sets of probability distributions are defined as a function of a set of *parameters* [Del15], e.g., intervals with parametric upper/lower bounds. With these functions over the parameters, we can describe dependencies between the model's states. Applying an instantiation (i.e., assigning a value to each of the parameters; see Chapter 8) to a pRMC induces an RMC by replacing each occurrence of the parameters with their assigned values. For this induced RMC, we compute a (robust) value for a given measure, and we call this verification result the *solution* for this instantiation. Analogous to parametric Markov decision processes (*pMDPs*), we can associate a pRMC with a *solution function* that maps parameter instantiations to solutions.

Differentiation for pRMCs | For the ground vehicle example above, we choose the parameters to represent the number of samples we have obtained for each terrain. Naturally, the *derivative of this solution function* with respect to each parameter (a.k.a. sample size) then corresponds to the expected change in the solution upon obtaining more samples. Such differentiation for parametric Markov chains (*pMCs*), where parameter instantiations yield one precise probability distribution, has been studied in [HSJM⁺22]. For pRMCs, however, it is unclear how to compute derivatives and under what conditions the derivative exists. We thus consider the following problem:

Problem 1 (Compute all derivatives): Given a pRMC and a parameter instantiation, compute the partial derivative of the solution function (evaluated at this instantiation) with respect to each of the parameters.

Our approach | We compute derivatives for pRMCs by solving a parameterized linear optimization problem. We build upon results from convex optimization theory for differentiating the optimal solution of this optimization problem [BV14; Bar18]. We also present sufficient conditions for the derivative to exist.

Improving efficiency | However, computing the derivative for every parameter explicitly does not scale to more realistic models with thousands of parameters. Instead, we observe that to determine for which parameter we should obtain more samples, we do not need to know *all partial derivatives explicitly*. Instead, it may suffice to know which parameters have *the highest* (or lowest, depending on the application) derivative. Thus, we also solve the following (related) problem:

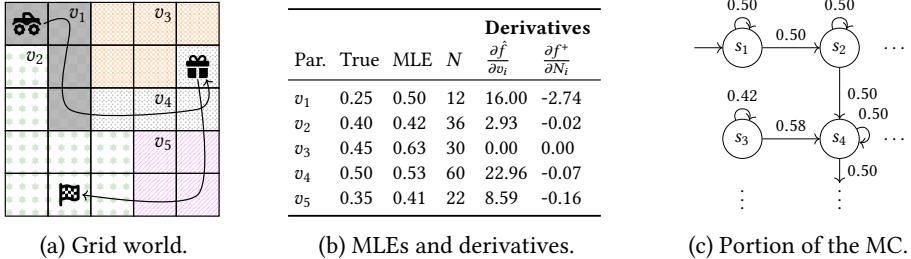


Figure 10.1: Grid world environment (a). The vehicle () must deliver the package () to the warehouse () We obtain the maximum likelihood estimates (MLEs) in (b), leading to the MC in (c).

Problem 2 (k highest derivatives): Given a pRMC with $|V|$ parameters, determine the $k < |V|$ parameters with the highest (or lowest) partial derivative.

We develop novel and efficient methods for solving Problem 2. Concretely, we design a linear program (LP) that finds the k parameters with the highest (or lowest) partial derivative without computing all derivatives explicitly. This LP constitutes a polynomial-time algorithm for Problem 2 and is, in practice, *orders of magnitude faster* than computing all derivatives explicitly, especially if the number of parameters is high. Moreover, if the concrete values for the partial derivatives are required, one can additionally solve Problem 1 for only the resulting k parameters. In our experiments, we show that we can compute derivatives for models with over a million states and thousands of parameters.

Learning framework | Learning in stochastic environments is very data-intensive in general, and millions of samples may be required to obtain sufficiently tight bounds on measures of interest [Kak03; MBPJ23]. Several methods exist to obtain intervals on probabilities based on sampling, including statistical methods such as Hoeffding's inequality [BLM13] and Bayesian methods that iteratively update intervals [SSPJ22]. Motivated by this challenge of reducing the sample complexity of learning algorithms, we embed our methods in an iterative learning scheme that profits from having access to sensitivity values for the parameters. In our experiments, we show that derivative information can be used effectively to guide sampling when learning an unknown Markov chain with hundreds of parameters.

Outline | We give an overview of our approach in Sect. 10.2 and formalize the problem statement in Sect. 10.3. In Sect. 10.4, we solve Problems (1) and (2) for pMCs, and in Sect. 10.5 for pRMCs. We present the embedding in a learning scheme and the experiments in Sect. 10.6. Finally, we survey related work in Sect. 10.7, and we discuss open challenges for further research in Sect. 10.8.

10.2 Overview

We expand the example from Sect. 10.1 to illustrate our approach more concretely. The environment, shown in Fig. 10.1a, is partitioned into five regions of the same terrain

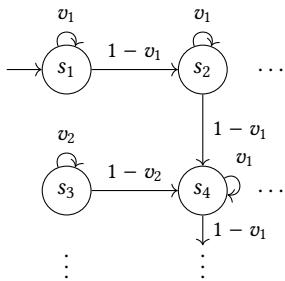


Figure 10.2: Parametric MC.

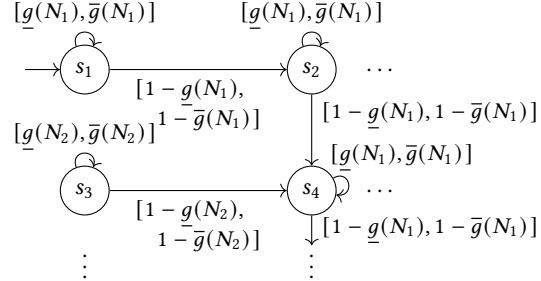


Figure 10.3: Parametric robust MC.

type. The vehicle can move in the four cardinal directions. Recall that the slipping probabilities are the same for all states with the same terrain. The vehicle follows a dedicated route to collect and deliver a package to a warehouse. Our goal is to estimate the expected number of steps f^* to complete the mission.

Estimating probabilities | Classically, we would derive maximum likelihood estimates (MLEs) of the probabilities by sampling. Consider that, using N samples per slipping probability, we obtained the rough MLEs shown in Fig. 10.1b and thus the MC in Fig. 10.1c. Verifying this Markov chain (MC) shows that the expected travel time (called the solution) under these estimates is $\hat{f} = 25.51$ steps, which is far from the travel time of $f^* = 21.62$ steps under the true slipping probabilities. We want to close this *verification-to-real gap* by taking more samples for one of the terrain types. For which of the five terrain types should we obtain more samples?

Parametric model | We can model the grid world as a parametric Markov chain (pMC), with a solution function representing the travel time \hat{f} for different parameter instantiations. We sketch four states of this pMC in Fig. 10.2. The most relevant parameter is then naturally defined as the parameter with the *largest partial derivative of the solution function*. As shown in Fig. 10.1b, parameter v_4 has the highest partial derivative of $\frac{\partial \hat{f}}{\partial v_4} = 22.96$, while the derivative of v_3 is zero as no states related to this parameter are ever visited.

Parametric robust model | The approach above does not account for the (level of) uncertainty in each MLE. Terrain type v_4 has the highest derivative but also the largest sample size, so sampling v_4 once more has likely less impact than for, e.g., v_1 . So, is v_4 actually the best choice to obtain additional samples for? The parametric robust MC (pRMC) that allows us to answer this question is shown in Fig. 10.3, where we use (parametric) intervals as uncertainty sets. The parameters are the sample sizes N_1, \dots, N_5 for all terrain types (contrary to the pMC, where parameters represent slipping probabilities). Now, if we obtain one additional sample for a particular terrain type, how can we expect the uncertainty sets to change?

Derivatives for pRMCs | We use the pRMC to compute an upper bound f^+ on the true solution f^* . Obtaining one more sample for terrain type v_i (i.e., increasing N_i by one) shrinks the interval $[g(N_i), \bar{g}(N_i)]$ on expectation, which in turn decreases our

upper bound f^+ . Here, \underline{g} and \bar{g} are functions mapping sample sizes to interval bounds. The partial derivatives $\frac{\partial f^+}{\partial N_i}$ for the pRMC are also shown in Fig. 10.1b and give a very different outcome than the derivatives for the pMC. In fact, sampling v_1 yields the biggest decrease in the upper bound f^+ , so we ultimately decide to sample for terrain type v_1 instead of v_4 .

Efficient differentiation | We remark that we do not need to know all derivatives explicitly to determine where to obtain samples. Instead, it suffices to know *which parameter has the highest (or lowest) derivative*. In the rest of the paper, we develop efficient methods for computing either all, or only the $k \in \mathbb{N}_{>0}$ highest, partial derivatives of the solution functions for pMCs and pRMCs.

Supported extensions | Our approaches are applicable to general pMCs and pRMCs whose parameters can be shared between distributions (and thus capture dependencies, being a common advantage of parametric models in general [JÁHJ⁺24]). Besides parameters in transition probabilities, we can handle parametric initial states, rewards, and policies. We could, e.g., use parameters to model the policy of a surveillance drone in our example and compute derivatives for these parameters.

10.3 Problem Statement

Recap | Recall from Chapter 8 that a pMC is defined as $\mathcal{D}_V = (S, s_I, V, P, r)$, where $V = \{v_1, \dots, v_{|V|}\}$, is a finite and ordered set of parameters.¹ The set of polynomials over parameters V with rational coefficients is $\mathbb{Q}[V]$. A parameter instantiation is a function $u: V \rightarrow \mathbb{Q}$ mapping parameters to real valuations. We again overload notation and write $u \in \mathbb{Q}^{|V|}$ for an instantiation of all parameters in V , which is defined as $[u(v_1), \dots, u(v_V)]^\top \in \mathbb{R}^{|V|}$. Applying an instantiation u to a pMC yields an MC $\mathcal{D}_V[u]$ by replacing each transition probability $f \in \mathbb{Q}[V]$ by $f[u]$. Finally, recall from Def. 8.5 that the parameter space $\mathcal{V}_{\mathcal{D}_V}$ is the subset of instantiations that lead to valid MCs (i.e., with valid probability distribution).

Remark 10.1 (Initial state distribution) In this chapter, we consider Markov chains where the initial state $s_I \in \text{Distr}(S)$ is defined as a *distribution over states*, instead of a fixed element of S (as we do in most other chapters). However, this difference is without loss of generality because we can always express an initial state distribution by expanding the Markov chain. Similarly, we can always express a fixed initial state $s_I \in S$ as a Dirac distribution over that state.

Expected rewards | In Chapter 8, we focused on solution functions for satisfaction probabilities, such as reachability probabilities. In this chapter, we shift focus and consider expected reward measures based on the state-reward function $r: S \rightarrow \mathbb{R}_{\geq 0}$ of a pMC. The next assumption ensures that the expected cumulative reward is finite.

¹We do not need the labeling function $L: S \rightarrow 2^{AP}$ in this chapter, so we omit it from the definition.

Assumption 10.2 (Sink states) Let $B \subseteq S$ be a set of absorbing (sink) states of pMC \mathcal{D}_V with zero reward, i.e., $r(s') = 0 \forall s' \in B$. From every state $s \in S$ and for every parameter instantiation $u \in \mathcal{V}_{\mathcal{D}_V}$, the pMC reaches the set B with probability one, i.e., $\Pr^{\mathcal{D}_V[u]}(s \models \diamond B) = 1$.

We remark that Assumption 10.2 is quite standard in probabilistic model checking to ensure finiteness of expected cumulative rewards [BK08].

Remark 10.3 (Discount factor) We may drop Assumption 10.2 by instead considering expected cumulative rewards with a discount factor $\gamma \in (0, 1)$. In this chapter, we purely focus on the case without a discount factor for brevity. Nevertheless, our methods can readily be adapted for the discounted case. For details on using discount factors, we refer back to Remark 3.21.

Recall from Chapter 3 that $\Pi^{\mathcal{D}_V[u]}(s)$ is the set of paths starting in $s \in S$, that the probability for a path $\pi \in \Pi^{\mathcal{D}_V[u]}(s)$ is denoted by $\Pr^{\mathcal{D}_V[u]}(\pi)$, and that the (undiscounted) cumulative reward over the path $\pi = s_0 s_1 \dots$ is $\text{rew}(\pi) = R(s_0) + R(s_1) + \dots$. Using the notation from Def. 3.20, the solution function $\text{sol}_{\mathbb{E}}: \mathcal{V}_{\mathcal{D}_V} \rightarrow \mathbb{R}_{\geq 0}$ for the expected cumulative reward in pMC \mathcal{D}_V is defined for all $u \in \mathcal{V}_{\mathcal{D}_V}$ as follows:

$$\begin{aligned}\text{sol}_{\mathbb{E}}(u) &= \sum_{s \in S} \left(s_I(s) \cdot \text{ExpRew}^{\mathcal{D}_V[u]}(s \models \diamond B) \right) \\ &= \sum_{s \in S} \left(s_I(s) \cdot \sum_{\pi \in \Pi_{\text{fin}}^{\mathcal{D}_V[u]}(s, \diamond B)} \Pr^{\mathcal{D}_V[u]}(\pi) \cdot \text{rew}(\pi) \right),\end{aligned}\tag{10.1}$$

where the second equality follows from the definition of the cumulative expected reward in Def. 3.20 (as Assumption 10.2 implies that $B \subseteq S$ is reached with probability one). Recall from Sect. 3.2.2.2 that, as a result, the expected cumulative reward is finite and characterized by a summation over finite paths $\pi \in \Pi_{\text{fin}}^{\mathcal{D}_V[u]}(s, \diamond B)$ only.

Remark 10.4 (Notation of solution function) In Chapter 8, we used the notation $\text{sol}_{\varphi}^{\mathcal{D}_V}$ for the solution function of a probabilistic computation tree logic (PCTL) formula φ in a pMC \mathcal{D}_V . In this chapter, we somewhat simplify notation and drop the superscript for the model for brevity. Thus, the model for which the solution function is defined is kept implicit.

10.3.1 Parametric robust Markov chains

In this section, we introduce parametric robust MCs (pRMCS) as an extension of robust MCs (RMCS) with parametric uncertainty sets. As we define below, these uncertainty sets are convex polytopes whose halfspace inequality constraints are defined by polynomials over a finite set of parameters.

Convex polytopes | Recall from Chapter 2 that the *convex polytope* $T_{A,b} \subseteq \mathbb{R}^n$ defined by matrix $A \in \mathbb{R}^{m \times n}$ and vector $b \in \mathbb{R}^m$ is the set $T_{A,b} = \{p \in \mathbb{R}^n : Ap \leq b\}$. We denote

by \mathbb{T}_n the set of all convex polytopes of dimension $n \in \mathbb{N}_{>0}$, i.e.,

$$\mathbb{T}_n = \{T_{A,b} : A \in \mathbb{R}^{m \times n}, b \in \mathbb{R}^m, m \in \mathbb{N}_{>0}\}.$$

In this chapter, we consider an RMC as a tuple (S, s_I, \mathcal{P}, r) , where the uncertain transition function $\mathcal{P} : S \rightarrow \mathbb{T}_{|S|}$ maps states to convex polytopes $T \in \mathbb{T}_{|S|}$. Compared to the RMC definition from Sect. 3.3, we thus impose the additional assumption that the uncertainty set for each state-action pair is a convex polytope.

Parametric polytopes | We extend RMCs with polytopes whose halfspaces are defined by polynomials $\mathbb{Q}[V]$ over V . To this end, let $\mathbb{T}_n[V]$ be the set of all such *parametric polytopes* of dimension $n \in \mathbb{N}_{>0}$:

$$\mathbb{T}_n[V] = \{T_{A,b} : A \in \mathbb{Q}[V]^{m \times n}, b \in \mathbb{Q}[V]^m, m \in \mathbb{N}_{>0}\}. \quad (10.2)$$

parametric polytope

An element $T \in \mathbb{T}_n[V]$ can be interpreted as a function $T : \mathbb{R}^{|V|} \rightarrow 2^{(\mathbb{R}^n)}$ that maps an instantiation u to a (possibly empty) convex polytopic subset of \mathbb{R}^n . The set $T[u]$ is obtained by substituting each v_i in T by $u(v_i)$ for all $i = 1, \dots, |V|$.

Example 10.5 The uncertainty set for state s_1 of the pRMC in Fig. 10.3 is the parametric polytope $T \in \mathbb{T}_2[V]$ with singleton parameter set $V = \{N_1\}$, such that

$$T = \{[p_{1,1}, p_{1,2}]^\top \in \mathbb{R}^2 : \underline{g}_1(N_1) \leq p_{1,1} \leq \bar{g}_1(N_1), \\ 1 - \bar{g}_1(N_1) \leq p_{1,2} \leq 1 - \underline{g}_1(N_1), \quad p_{1,2} + p_{1,1} = 1\}.$$

parametric robust Markov chain

We use parametric convex polytopes to define a parametric robust MC (pRMC) in the following way:

Definition 10.6 (pRMC) A *parametric robust MC (pRMC)* is a tuple $\mathcal{M}_R^V := (S, s_I, V, \mathcal{P}, r)$, where S is a finite set of states, $s_I \in \text{Distr}(S)$ is an *initial state (distribution)*, V is an (ordered) set of *parameters*, $r : S \rightarrow \mathbb{R}_{\geq 0}$ is a *state reward function*, and (different from RMCs) $\mathcal{P} : S \rightarrow \mathbb{T}_{|S|}[V]$ is a parametric and uncertain transition function that maps states to parametric convex polytopes.

The scripts V and R in \mathcal{M}_R^V indicate that the pRMC is both parametric and robust.

Applying an instantiation u to a pRMC yields an RMC $\mathcal{M}_R^V[u]$ by replacing each parametric polytope $T \in \mathbb{T}_{|S|}[V]$ by $T[u]$, which is a (nonparametric) polytope defined by a concrete matrix $A \in \mathbb{R}^{m \times n}$ and vector $b \in \mathbb{R}^m$. Without loss of generality, we consider that nature (see Def. 3.29) *minimizes* the expected cumulative reward until reaching a set of terminal states $S_T \subseteq S$. Because we consider infinite-horizon rewards, we restrict ourselves to *stationary* natures $\tau \in \mathfrak{T}_{\text{stat}}^{\mathcal{M}_R^V[u]}$ (as defined in Sect. 3.3).

The minimum expected cumulative reward $\text{sol}_{\mathbb{E}, R}(u)$, called the *robust solution* on the instantiated pRMC $\mathcal{M}_R^V[u]$, is defined as

$$\text{sol}_{\mathbb{E}, R}(u) = \sum_{s \in S} \left(s_I(s) \cdot \min_{\tau \in \mathfrak{T}_{\text{stat}}^{\mathcal{M}_R^V[u]}} \sum_{\pi \in \Pi_{\text{fin}}^{\mathcal{D}_V[u]}(s, \diamond B)} \text{Pr}_{\tau}^{\mathcal{M}_R^V[u]}(\pi) \cdot \text{rew}(\pi) \right), \quad (10.3)$$

where the subscript R expresses that the solution function is robust (i.e., computed for the most pessimistic nature; see Sect. 3.3.4). We refer to the function $\text{sol}_{\mathbb{E},R} : \mathcal{V}_{\mathcal{M}_R^V} \rightarrow \mathbb{R}_{>0}$ as the *robust solution function* (for the cumulative expected reward measure).

robust
solution
function

Assumption 10.7 (Graph-preservation of p(R)MCs) In this chapter, we consider Assumption 10.2 for all pMCs and pRMCs. Furthermore, we assume that transitions cannot vanish under any instantiation (graph-preservation). That is, for every $s, s' \in S$, we have that $P(s)[u](s')$ (for pMCs) and $\mathcal{P}(s)[u](s')$ (for pRMCs) are either zero or strictly positive for all instantiations u . Graph-preservation ensures that solution functions are smooth for all pMCs and for most pRMCs (we discuss the exceptions in Sect. 10.5.1).

10.3.2 Problem statement

Let $f(q_1, \dots, q_n) \in \mathbb{R}^m$ be a differentiable multivariate function with $m \in \mathbb{N}_{>0}$. We denote the *partial derivative* of the function f with respect to its arguments $q = [q_1, \dots, q_n]$ by $\frac{\partial f}{\partial q} \in \mathbb{R}^m$. The gradient of f combines all partial derivatives in a single vector as $\nabla_q f = [\frac{\partial f}{\partial q_1}, \dots, \frac{\partial f}{\partial q_n}] \in \mathbb{R}^{m \times n}$. We only use gradients $\nabla_u f$ with respect to the parameter instantiation u , so we simply write ∇f in the remainder.

partial
derivative

The *gradient* of the robust solution function evaluated at the instantiation u is $\nabla \text{sol}_{\mathbb{E},R}[u] = \left[\left(\frac{\partial \text{sol}_{\mathbb{E},R}}{\partial u(v_1)} \right)[u], \dots, \left(\frac{\partial \text{sol}_{\mathbb{E},R}}{\partial u(v_{|V|})} \right)[u] \right]$. We formalize the first problem that we described in Sect. 10.1 as follows.

Problem 10.8 (Computing all $|V|$ derivatives) Given a pRMC \mathcal{M}_R^V and a parameter instantiation u , compute the gradient $\nabla \text{sol}_{\mathbb{E},R}[u]$ of the robust solution function evaluated at u .

Solving Problem 10.8 is linear in the number of parameters, which leads to significant overhead if the number of parameters is large. Typically, it suffices to obtain only the parameters with the highest derivatives, leading to the second problem statement:

Problem 10.9 (Computing $k < |V|$ highest derivatives) Given a pRMC \mathcal{M}_R^V , an instantiation u , and a $k < |V|$, compute a subset V^* of k parameters for which the partial derivatives are maximal.

For both problems, we present polynomial-time algorithms for pMCs (Sect. 10.4) and pRMCs (Sect. 10.5). Sect. 10.6 defines problem variations that we study empirically.

10.4 Differentiating Solution Functions for pMCs

Before solving Problems 10.8 and 10.9 for pRMCs, we first consider the simpler version of these problems for pMCs.

We can compute the solution of an MC $\mathcal{D}_V[u]$ with instantiation u based on a system of $|S|$ linear equations; here for an expected reward measure [BK08]. Let $x = [x_{s_1}, \dots, x_{s_{|S|}}]^T$ and $r = [r_{s_1}, \dots, r_{s_{|S|}}]^T$ be variables for the expected cumulative reward and the instantaneous reward in each state $s \in S$, respectively. Then, for a set of terminal

(sink) states $S_T \subset S$, we obtain the equation system

$$x_s = 0, \quad \forall s \in S_T \quad (10.4a)$$

$$x_s = r_s + P(s)[u]x, \quad \forall s \in S \setminus S_T. \quad (10.4b)$$

Let us set $P(s)[u] = 0$ for all $s \in S_T$ and define the matrix $P[u] \in \mathbb{R}^{|S| \times |S|}$ by stacking the rows $P(s)[u]$ for all $s \in S$. Then, Eq. (10.4) is written in matrix form as $(I_{|S|} - P[u])x = r$. The equation system in Eq. (10.4) can be efficiently solved by, e.g., Gaussian elimination or more advanced iterative equation solvers.

10.4.1 Computing derivatives explicitly

We differentiate the equation system in Eq. (10.4) with respect to an instantiation $u(v_i)$ for parameter $v_i \in V$, similar to, e.g., [HSJM⁺22]. For all $s \in S_T$, the derivative $\frac{\partial x_s}{\partial u(v_i)}$ is trivially zero. For all $s \in S \setminus S_T$, we obtain via the product rule that

$$\frac{\partial x_s}{\partial u(v_i)} = \frac{\partial P(s)x}{\partial u(v_i)}[u] = (x^\star)^\top \frac{\partial P(s)^\top}{\partial u(v_i)}[u] + P(s)[u] \frac{\partial x}{\partial u(v_i)},$$

where $x^\star \in \mathbb{R}^{|S|}$ is the solution to Eq. (10.4). In matrix form for all $s \in S$, this yields

$$(I_{|S|} - P[u]) \frac{\partial x}{\partial u(v_i)} = \frac{\partial Px^\star}{\partial u(v_i)}[u]. \quad (10.5)$$

The solution defined in Eq. (10.1) is computed as $\text{sol}_{\mathbb{E}}[u] = s_I^\top x^\star$. Thus, the partial derivative of the solution function with respect to $u(v_i)$ in closed form is

$$\left(\frac{\partial \text{sol}_{\mathbb{E}}}{\partial u(v_i)} \right) [u] = s_I^\top \frac{\partial x}{\partial u(v_i)} = s_I^\top (I_{|S|} - P[u])^{-1} \frac{\partial Px^\star}{\partial u(v_i)}[u]. \quad (10.6)$$

Algorithm for Problem 10.8 | Let us provide an algorithm to solve Problem 10.8 for pMCs. Eq. (10.6) provides a closed-form expression for the partial derivative of the solution function, which is a function of the vector x^\star in Eq. (10.4). However, due to the inversion of $(I_{|S|} - P[u])$, it is generally more efficient to solve the system of equations in Eq. (10.5). Doing so, the partial derivative of the solution with respect to $u(v_i)$ is obtained by: (1) solving Eq. (10.4) with u to obtain $x^\star \in \mathbb{R}^{|S|}$, and (2) solving the equation system in Eq. (10.5) with $|S|$ unknowns for this vector x^\star . We repeat step 2 for all of the $|V|$ parameters. Thus, we can solve Problem 10.8 by solving $|V| + 1$ linear equation systems with $|S|$ unknowns each.

10.4.2 Computing k highest derivatives

To solve Problem 10.9 for pMCs, we present a method to compute only the $k \leq |V|$ parameters with the highest (or lowest) partial derivative without computing all derivatives explicitly. Without loss of generality, we focus on the highest derivative. We can determine these parameters by solving a combinatorial optimization problem with binary variables $z_i \in \{0, 1\}$ for $i = 1, \dots, |V|$. Our goal is to formulate this optimization problem such that an optimal value of $z_i^\star = 1$ implies that parameter $v_i \in V$ belongs to the set of k highest derivatives. Concretely, we formulate the following *mixed integer*

linear problem (MILP) [Wol20]:

$$\underset{y \in \mathbb{R}^{|S|}, z \in \{0,1\}^{|V|}}{\text{maximize}} \quad s_I^\top y \quad (10.7a)$$

$$\text{subject to } (I_{|S|} - P[u]) y = \sum_{i=1}^{|V|} z_i \frac{\partial P x^\star}{\partial u(v_i)} [u] \quad (10.7b)$$

$$z_1 + \dots + z_{|V|} = k. \quad (10.7c)$$

Constraint (10.7c) ensures that any feasible solution to Eq. (10.7) has exactly k nonzero entries. Since matrix $(I_{|S|} - P[u])$ is invertible by construction (see, e.g., [Put94]), Eq. (10.7) has a unique solution in y for each choice of $z \in \{0, 1\}^{|V|}$. Thus, the objective value $s_I^\top y$ is the sum of the derivatives for the parameters $v_i \in V$ for which $z_i = 1$. Since we maximize this objective, an optimal solution y^\star, z^\star to Eq. (10.7) is guaranteed to correspond to the k parameters that maximize the derivative of the solution in Eq. (10.6). We state this correctness claim for the MILP:

Proposition 10.10 (k highest derivatives for pMC) Let y^\star, z^\star be an optimal solution to Eq. (10.7). Then, the set $V^\star = \{v_i \in V : z_i^\star = 1\}$ is a subset of $k \leq |V|$ parameters with maximal derivatives.

The set V^\star may not be unique. However, to solve Problem 10.9, it suffices to obtain a set of k parameters for which the partial derivatives are maximal. Therefore, the set V^\star provides a solution to Problem 10.9. We remark that, to solve Problem 10.9 for the k lowest derivatives, we change the objective in Eq. (10.7a) to minimize $s_I^\top y$.

Linear relaxation | The MILP in Eq. (10.7) is computationally intractable for high values of $|V|$ and k . Instead, we compute the set V^\star via a *linear relaxation* of the MILP. Specifically, we relax the *binary* variables $z \in \{0, 1\}^{|V|}$ to *continuous* variables $z \in [0, 1]^{|V|}$. As such, we obtain the following LP relaxation of Eq. (10.7):

$$\underset{y \in \mathbb{R}^{|S|}, z \in \mathbb{R}^{|V|}}{\text{maximize}} \quad s_I^\top y \quad (10.8a)$$

$$\text{subject to } (I_{|S|} - P[u]) y = \sum_{i=1}^{|V|} z_i \frac{\partial P x^\star}{\partial u(v_i)} [u] \quad (10.8b)$$

$$0 \leq z_i \leq 1, \quad \forall i = 1, \dots, |V| \quad (10.8c)$$

$$z_1 + \dots + z_{|V|} = k. \quad (10.8d)$$

Denote by y^+, z^+ the solution of the LP relaxation in Eq. (10.8). For details on such linear relaxations of integer problems, we refer to [HK10; MG07]. In our case, every optimal solution y^+, z^+ to the LP relaxation with only binary values $z_i^+ \in \{0, 1\}$ is also optimal for the MILP, resulting in the following theorem.

Theorem 10.11 (Correctness of LP relaxation) The LP relaxation in Eq. (10.8) has an optimal solution y^+, z^+ with $z^+ \in \{0, 1\}^{|V|}$ (i.e., every z_i^+ is binary), and every such a solution is also an optimal solution of the MILP in Eq. (10.7).

Proof. From invertibility of $(I_{|S|} - P[u])$, we know that Eq. (10.7) is equivalent to

$$\underset{z \in \{0,1\}^{|V|}}{\text{maximize}} \sum_{i=1}^{|V|} z_i \left(s_I^\top (I_{|S|} - P[u])^{-1} \frac{\partial P x^*}{\partial u(v_i)}[u] \right) \quad (10.9a)$$

$$\text{subject to } z_1 + \cdots + z_{|V|} = k. \quad (10.9b)$$

The linear relaxation of Eq. (10.9) is an LP whose feasible region has integer vertices (see, e.g., [HK10]). Therefore, both Eq. (10.9) and its relaxation Eq. (10.8) have an integer optimal solution z^+ , which constructs z^* in Eq. (10.7).

The binary solutions $z^+ \in \{0,1\}^{|V|}$ are the vertices of the feasible set of the LP in Eq. (10.8). A simplex-based LP solver can be set to return such a solution.²

Algorithm for Problem 10.9 | We provide a two-step algorithm to solve Problem 10.9 for pMCs. First, for pMC \mathcal{D}_V and parameter instantiation u , we solve the linear equation system in Eq. (10.5) for x^* to obtain the solution $\text{sol}_{\mathbb{E}}[u] = s_I^\top x^*$. Second, we fix a number of parameters $k \leq |V|$ and solve the LP relaxation in Eq. (10.8). The set V^* of parameters with maximal derivatives is then obtained as defined in Proposition 10.10. The parameter set V^* is a solution to Problem 10.9.

10.5 Differentiating Solution Functions for pRMCs

We shift focus to pRMCs. Recall that robust solutions $\text{sol}_{\mathbb{E}, R}[u]$ for pRMCs are computed for the minimizing stationary nature $\tau \in \mathfrak{T}_{\text{stat}}^{M_R^V[u]}$. For each state $s \in S$, such a nature fixes a probability distribution $p \in \mathcal{P}(s)$ in the uncertain transition function of the pRMC. Thus, we derive the following equation system, where, as for pMCs, $x \in \mathbb{R}^{|S|}$ represents the expected cumulative reward in each state.

$$x_s = 0, \quad \forall s \in S_T \quad (10.10a)$$

$$x_s = r_s + \inf_{p \in \mathcal{P}(s)[u]} (p^\top x), \quad \forall s \in S \setminus S_T. \quad (10.10b)$$

Solving Eq. (10.10) directly corresponds to solving a system of nonlinear equations due to the inner infimum in Eq. (10.10b). The standard approach from robust optimization [BGN09] is to leverage the dual problem for each inner infimum, e.g., as is done in [PLSS13; CFRS14]. For each $s \in S$, $\mathcal{P}(s)$ is a parametric convex polytope $T_{A,b}$ as defined in Eq. (10.2). The dimensionality of this polytope depends on the number of successor states, which is typically much lower than the total number of states. To make the number of successor states explicit, we denote by $\text{post}(s) \subseteq S$ the successor states of $s \in S$ and define $T_{A,b} \in \mathbb{T}_{|\text{post}(s)|}[V]$ with $A_s \in \mathbb{Q}^{m_s \times |\text{post}(s)|}$ and $b_s[u] \in \mathbb{Q}^{m_s}$ (recall m_s is the number of halfspaces of the polytope). Then, the infimum in Eq. (10.10b) for

²Even if a non-vertex solution y^+, z^+ is obtained, we can use an arbitrary tie-break rule on z^+ , which forces each z_i^+ binary and preserves the sum in Eq. (10.8d).

each $s \in S \setminus S_T$ is written as follows:

$$\text{minimize } p^\top x \quad (10.11a)$$

$$\text{subject to } A_s[u]p \leq b_s[u] \quad (10.11b)$$

$$\mathbb{1}^\top p = 1, \quad (10.11c)$$

where $\mathbb{1}$ denotes a column vector of ones of appropriate size. Let $x_{\text{post}(s)} = [x_s]_{s \in \text{post}(s)}$ be the vector of decision variables corresponding to the (ordered) successor states in $\text{post}(s)$. The dual problem of Eq. (10.11), with dual variables $\alpha \in \mathbb{R}^{m_s}$ and $\beta \in \mathbb{R}$ (see, e.g., [BJS11] for details), is written as follows:

$$\text{maximize } -b_s[u]^\top \alpha - \beta \quad (10.12a)$$

$$\text{subject to } A_s[u]^\top \alpha + x_{\text{post}(s)} + \beta \mathbb{1} = 0, \quad \alpha \geq 0. \quad (10.12b)$$

By using this dual problem in Eq. (10.10b), we obtain the following LP with decision variables $x \in \mathbb{R}^{|S|}$, and with $\alpha_s \in \mathbb{R}^{m_s}$ and $\beta_s \in \mathbb{R}$ for every $s \in S$:

$$\text{maximize } s_I^\top x \quad (10.13a)$$

$$\text{subject to } x_s = 0, \quad \forall s \in S_T \quad (10.13b)$$

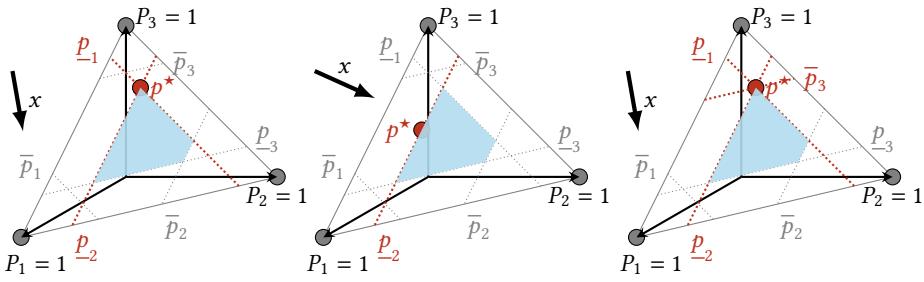
$$x_s = r_s - (b_s[u]^\top \alpha_s + \beta_s), \quad \forall s \in S \setminus S_T \quad (10.13c)$$

$$A_s[u]^\top \alpha_s + x_{\text{post}(s)} + \beta_s \mathbb{1} = 0, \quad \alpha_s \geq 0, \quad \forall s \in S \setminus S_T. \quad (10.13d)$$

The reformulation of Eq. (10.10) to Eq. (10.13) requires that every element of the initial distribution s_I is nonnegative, which is trivially satisfied because s_I is a probability distribution. Denote by x^*, α^*, β^* an optimal point of Eq. (10.13). The x^* element of this optimum is also an optimal solution of Eq. (10.10) [BGN09]. Thus, the robust solution defined in Eq. (10.3) is $\text{sol}_{\mathbb{E}, R}[u] = s_I^\top x^*$.

10.5.1 Computing derivatives via pMCs

Toward solving Problem 10.8, we provide some intuition about computing robust solutions for pRMCS. The infimum in Eq. (10.10b) finds the *worst-case* point p^* in each set



(a) Well-defined optimum. (b) Non-unique optimum. (c) Too many active constraints.

Figure 10.4: Three polytopic uncertainty sets (blue shade), with the vector x , the worst-case points p^* , and the active constraints shown in red.

$\mathcal{P}(s)[u]$ that minimizes $(p^*)^\top x$. This minimization is visualized in Fig. 10.4a for an uncertainty set that captures three probability intervals $\underline{p}_i \leq p_i \leq \bar{p}_i$, $i = 1, 2, 3$. Given the optimization direction x (arrow in Fig. 10.4a), the point p^* (red dot) is attained at the vertex where the constraints $\underline{p}_1 \leq p_1$ and $\underline{p}_2 \leq p_2$ are active.³ Thus, we obtain that the point in the polytope that minimizes $(p^*)^\top x$ is $p^* = [\underline{p}_1, \underline{p}_2, 1 - \underline{p}_1 - \underline{p}_2]^\top$. Using this procedure, we can obtain a worst-case point p_s^* for each state $s \in S$. We can use these points to convert the pRMC into an induced pMC with transition function $P(s) = p_s^*$ for each state $s \in S$.

For small changes in the parameters, the point p^* in Fig. 10.4a changes smoothly, and its closed-form expression (i.e., the functional form) remains the same. As such, it feels intuitive that we could apply the methods from Sect. 10.4 to compute partial derivatives on the induced pMC. However, this approach does not always work, as illustrated by the following two corner cases.

1. Consider Fig. 10.4b, where the optimization direction defined by x is parallel to one of the facets of the uncertainty set. In this case, the worst-case point p^* is not unique, but an infinitesimal change in the optimization direction x will force the point to one of the vertices again. Which point should we choose to obtain the induced pMC (and does this choice affect the derivative)?
2. Consider Fig. 10.4c with more than $|S| - 1$ active constraints at the point p^* . Observe that decreasing \bar{p}_3 changes the point p^* while increasing \bar{p}_3 does not. In fact, the optimal point p^* changes *non-smoothly* with the halfspaces of the polytope. As a result, also the solution changes non-smoothly, and thus, the derivative is not defined. How do we deal with such a situation?

These examples show that computing derivatives via an induced pMC by obtaining p_s^* for every state $s \in S$ can be tricky or is, in some cases, not possible at all. In what follows, we present a method that directly derives a set of linear equations to obtain derivatives for pRMCs (all or only the k highest) based on the solution to the LP in Eq. (10.13). Using this method, we intrinsically identify the corner cases above in which the derivative is not defined.

10.5.2 Computing derivatives explicitly

We now develop a dedicated method for identifying whether the derivative of the solution function for a pRMC exists, and if so, to compute this derivative. Observe from Fig. 10.4 that the point p^* is uniquely defined and has a smooth derivative only in the case of Fig. 10.4a, which has two active constraints. For only one active constraint (Fig. 10.4b), the point is *underdetermined*, while for three active constraints (Fig. 10.4c), the derivative may *not be smooth*. In the general case, having exactly $n - 1$ active constraints (whose facets are nonparallel) is a sufficient condition for obtaining a unique and smoothly changing point p^* in the n -dimensional probability simplex.

Optimal dual variables | The optimal dual variables $\alpha_s^* \geq 0$ for each $s \in S \setminus S_T$ in Eq. (10.13) indicate which constraints of the polytope $A_s[u]p \leq b_s[u]$ are active, i.e., for which rows $a_{s,i}[u]$ of $A_s[u]$ it holds that $a_{s,i}[u]p^* = b_s[u]$. Specifically, a value of $\alpha_{s,i} > 0$ implies that the i^{th} constraint is active, and $\alpha_{s,i} = 0$ indicates a nonactive

³An inequality constraint $gx \leq h$ is active under the optimal solution x^* if $gx^* = h$ [BV14].

constraint [BV14]. We define $E_s = [e_1, \dots, e_{m_s}] \in \{0, 1\}^{m_s}$ as a vector whose binary values $e_i \forall i \in \{1, \dots, m_s\}$ are given as $e_i = [\![\alpha_{s,i}^* > 0]\!]$.⁴ Moreover, denote by $\text{diag}(E_s)$ the matrix with E_s on the diagonal and zeros elsewhere. We reduce the LP in Eq. (10.13) to a system of linear equations that encodes only the constraints that are active under the worst-case point p_s^* for each $s \in S \setminus S_T$:

$$x_s = 0, \quad \forall s \in S_T \quad (10.14a)$$

$$x_s = r_s - (b_s[u]^\top \text{diag}(E_s)\alpha_s + \beta_s), \quad \forall s \in S \setminus S_T \quad (10.14b)$$

$$A_s[u]^\top \text{diag}(E_s)\alpha_s + x_{\text{post}(s)} + \beta_s \mathbb{1} = 0, \quad \alpha_s \geq 0, \quad \forall s \in S \setminus S_T. \quad (10.14c)$$

Differentiation | However, when does Eq. (10.14) have a (unique) optimal solution? To provide some intuition, let us write the equation system in matrix form, i.e., $C[x \ \alpha \ \beta]^\top = d$, where we omit an explicit definition of matrix C and vector d for brevity. It is apparent that if matrix C is nonsingular, then Eq. (10.14) has a unique solution. This requires matrix C to be square, which is achieved if, for each $s \in S \setminus S_T$, we have $|\text{post}(s)| = \sum E_s + 1$. In other words, the number of successor states of s is equal to the number of active constraints of the polytope plus one. This confirms our previous intuition from Sect. 10.5.1 on a polytope for $|\text{post}(s)| = 3$ successor states, which required $\sum_{i=1}^{m_s} E_i = 2$ active constraints.

Let us formalize this intuition about computing derivatives for pRMCs. We can compute the derivative of the solution x^* by differentiating the equation system in Eq. (10.14) through the product rule in a very similar manner to the approach in Sect. 10.4. We state this key result in the following theorem.

Theorem 10.12 (Derivatives for pRMCs) Given a pRMC \mathcal{M}_R^V and an instantiation u , compute x^*, α^*, β^* for Eq. (10.13) and choose a parameter $v_i \in V$. The partial derivatives $\frac{\partial x}{\partial u(v_i)}$, $\frac{\partial \alpha}{\partial u(v_i)}$, and $\frac{\partial \beta}{\partial u(v_i)}$ are obtained as the solution to the linear equation system

$$\frac{\partial x_s}{\partial u(v_i)} = 0, \quad \forall s \in S_T \quad (10.15a)$$

$$\frac{\partial x_s}{\partial u(v_i)} + b_s[u]^\top \text{diag}(E_s) \frac{\partial \alpha_s}{\partial u(v_i)} + \frac{\partial \beta_s}{\partial u(v_i)} = -(\alpha_s^*)^\top \text{diag}(E_s) \frac{\partial b_s[u]}{\partial u(v_i)}, \\ \forall s \in S \setminus S_T \quad (10.15b)$$

$$A_s[u]^\top \text{diag}(E_s) \frac{\partial \alpha_s}{\partial u(v_i)} + \frac{\partial x_{\text{post}(s)}}{\partial u(v_i)} + \frac{\partial \beta_s}{\partial u(v_i)} \mathbb{1} = -(\alpha_s^*)^\top \text{diag}(E_s) \frac{\partial A_s[u]}{\partial u(v_i)}, \\ \forall s \in S \setminus S_T. \quad (10.15c)$$

Proof. The proof follows from applying the product rule to the equation system in Eq. (10.14). The derivative of the right-hand side Eq. (10.14a) (i.e., for all terminal

⁴Here, $[\![x]\!]$ is an Iverson bracket, which is defined as $[\![x]\!] = 1$ if x is true and as $[\![x]\!] = 0$ otherwise.

states $s \in S_T$) is trivially zero. The derivative of Eq. (10.14b) is

$$\begin{aligned}\frac{\partial x_s}{\partial u(v_i)} &= -\frac{\partial (b_s[u]^\top \text{diag}(E_s)\alpha_s)}{\partial u(v_i)} - \frac{\partial \beta_s}{\partial u(v_i)} \\ \frac{\partial x_s}{\partial u(v_i)} &= -(\alpha_s^*)^\top \text{diag}(E_s) \frac{\partial b_s[u]}{\partial u(v_i)} - b_s[u]^\top \text{diag}(E_s) \frac{\partial \alpha}{\partial u(v_i)} - \frac{\partial \beta_s}{\partial u(v_i)},\end{aligned}$$

which, after rearranging, yields Eq. (10.15b). The derivative of Eq. (10.14c) is

$$(\alpha_s^*)^\top \text{diag}(E_s) \frac{\partial A_s[u]}{\partial u(v_i)} + A_s[u]^\top \text{diag}(E_s) \frac{\partial \alpha_s}{\partial u(v_i)} + \frac{\partial x_{\text{post}(s)}}{\partial u(v_i)} + \frac{\partial \beta_s}{\partial u(v_i)} \mathbb{1},$$

which after rearranging yields Eq. (10.15c), so we conclude the proof. ■

To compute the derivative for a parameter $v_i \in V$, we thus solve a system of linear equations of size $|S| + \sum_{s \in S \setminus S_T} |\text{post}(s)|$. Using Theorem 10.12, we obtain sufficient conditions for the solution function to be differentiable.

Lemma 10.13 (Differentiability for pRMCS) Write the linear equation system in Eq. (10.15) in matrix form, i.e.,

$$C \left[\frac{\partial x}{\partial u(v_i)}, \frac{\partial \alpha}{\partial u(v_i)}, \frac{\partial \beta}{\partial u(v_i)} \right]^\top = d, \quad (10.16)$$

for $C \in \mathbb{R}^{q \times q}$ and $d \in \mathbb{R}^q$, $q = |S| + \sum_{s \in S \setminus S_T} |\text{post}(s)|$, which are implicitly given by Eq. (10.15). The solution function $\text{sol}_{\mathbb{B}, R}[u]$ is differentiable at instantiation u if matrix C is nonsingular, in which case we obtain $(\frac{\partial \text{sol}_{\mathbb{B}, R}}{\partial u(v_i)})[u] = s_I^\top \frac{\partial x}{\partial u(v_i)}$.

Proof. The partial derivative of the solution function is $\frac{\partial \text{sol}_{\mathbb{B}, R}}{\partial u(v_i)}[u] = s_I^\top \frac{\partial x^*}{\partial u(v_i)}$, where $\frac{\partial x^*}{\partial u(v_i)}$ is (a part of) the solution to Eq. (10.14). Thus, the solution function is differentiable if there is a (unique) solution to Eq. (10.14), which is guaranteed if matrix C is nonsingular. Thus, the claim in Lemma 10.13 follows.

Algorithm for Problem 10.8 | We use Theorem 10.12 to solve Problem 10.8 for pRMCS, similarly as for pMCs. Given a pRMC \mathcal{M}_R^V and an instantiation u , we first solve Eq. (10.13) to obtain x^*, α^*, β^* . Second, we use α_s^* to compute the vector E_s of active constraints for each $s \in S \setminus S_T$. Third, for every parameter $v \in V$, we solve the equation system in Eq. (10.15). Thus, to compute the gradient of the solution function, we solve one LP and $|V|$ linear equation systems.

10.5.3 Computing k highest derivatives

Analogous to the MILP in Eq. (10.7) formulated for a pMC, we can compute the $k \leq |V|$ highest derivatives of a pRMC based on the solution to a MILP. For brevity, let us define the notations $x'_s = \frac{\partial x_s}{\partial u(v_i)} \in \mathbb{R}$, $\alpha'_s = \frac{\partial \alpha_s}{\partial u(v_i)} \in \mathbb{R}^{m_s}$ and $\beta'_s = \frac{\partial \beta_s}{\partial u(v_i)} \in \mathbb{R}$. Using this

notation, we obtain the following MILP:

$$\underset{x', \alpha', \beta', z \in \{0,1\}^{|V|}}{\text{maximize}} \quad s_I^\top x' \quad (10.17a)$$

$$\text{subject to } x'_s = 0, \quad \forall s \in S_T \quad (10.17b)$$

$$x'_s + b_s[u]^\top \text{diag}(E_s) \alpha'_s + \beta'_s = -(\alpha_s^\star)^\top \text{diag}(E_s) \sum_{i=1}^{|V|} z_i \frac{\partial b_s[u]}{\partial u(v_i)}, \\ \forall s \in S \setminus S_T \quad (10.17c)$$

$$A_s[u]^\top \text{diag}(E_s) \alpha'_s + x'_{\text{post}(s)} + \beta'_s \mathbb{1} = -(\alpha_s^\star)^\top \text{diag}(E_s) \sum_{i=1}^{|V|} z_i \frac{\partial A_s[u]}{\partial u(v_i)}, \\ \forall s \in S \setminus S_T. \quad (10.17d)$$

$$z_1 + \dots + z_{|V|} = k. \quad (10.17e)$$

Observe that the difference between the constraints in Eqs. (10.17c) and (10.17d), versus the equation system in Eq. (10.15) lies in the summation over $i = 1, \dots, |V|$. We derive the same LP relaxation as in Eq. (10.8), i.e., we relax the binary variables $z \in \{0, 1\}^{|V|}$ to continuous variables $z \in [0, 1]^{|V|}$. Since Eq. (10.17) has the exact same characteristics as Eq. (10.7), Theorem 10.11 applies equivalently to the case for pRMCs. In other words, the LP relaxation is exact, and we can use the resulting solution to find the set V^\star of parameters with maximal derivatives using Proposition 10.10. This set V^\star is a solution to Problem 10.9 for pRMCs.

10.6 Numerical Experiments

We perform experiments to answer the following questions about our approach:

1. Is it feasible (in terms of computational complexity and runtimes) to compute all $|V|$ derivatives, in particular compared to computing (robust) solutions?
2. How does computing only the k highest derivatives compare to computing all $|V|$ derivatives?
3. Can we apply our approach to effectively determine for which parameters to sample in a learning framework?

Let us briefly summarize the computations involved in answering these questions. First of all, computing the solution $\text{sol}_{\mathbb{B}}(u)$ for a pMC, which is defined in Eq. (10.1), means solving the linear equation system in Eq. (10.4). Similarly, computing the robust solution $\text{sol}_{\mathbb{B}, R}(u)$ for a pRMC means solving the LP in Eq. (10.13). Then, solving Problem 10.8, i.e., computing all $|V|$ partial derivatives, amounts to solving a linear equation system for each parameter $v \in V$ (namely, Eq. (10.4) for a pRMC and Eq. (10.15) for a pRMC). In contrast, solving Problem 10.9, i.e., computing a subset V^\star of parameters with maximal (or minimal) derivative, means for a pMC that we solve the LP in Eq. (10.8) (or the equivalent LP for a pRMC) and thereafter extract the subset of V^\star parameters using Proposition 10.10.

Problem 3: Computing the k highest derivatives | A solution to Problem 10.9 is a set V^* of k parameters but does not include the computation of the derivatives. However, it is straightforward to also obtain the actual derivatives $\left(\frac{\partial \text{sol}_{\mathbb{E}, R}}{\partial u(v)}\right)[u]$ for each parameter $v \in V^*$. Specifically, we solve Problem 10.8 for the k parameters in V^* , such that we obtain the partial derivatives for all $v \in V^*$. We remark that, for $k = 1$, the derivative follows directly from the optimal value $s_I^\top y^+$ of the LP in Eq. (10.8), so this additional step is not necessary. We will refer to computing the actual values of the k highest derivatives as *Problem 3*.

Reproducibility | We implement our approach in Python 3.10, using the model checker Storm [HJKQ²²] to parse pMCs, Gurobi [Gur23] to solve LPs, and the SciPy sparse solver [VGOH¹⁹] to solve equation systems. All experiments run on a computer with a 4GHz Intel Core i9 CPU and 64 GB RAM, with a timeout of one hour. Our implementation is available at <https://doi.org/10.5281/zenodo.7864260>.

Grid world benchmarks | We use scaled versions of the grid world from the example in Sect. 10.2 with over a million states and up to 10 000 terrain types. The vehicle only moves right or down, both with 50% probability (wrapping around when leaving the grid). Slipping only occurs when moving down and (slightly different from the example in Sect. 10.2) means that the vehicle moves *two cells instead of one*. We obtain between $N = 500$ and 1 000 samples of each slipping probability. For the pMCs, we use maximum likelihood estimation ($\frac{\bar{p}}{N}$, with \bar{p} the sample mean) obtained from these samples as probabilities, whereas, for the pRMCs, we infer probability intervals using Hoeffding's inequality (see Q3 for details).

Benchmarks from literature | We consider several instances of parametric extensions of MCs and Markov decision processes (MDPs) from standard benchmark suits [HKPQ¹⁹; KNP11]. We also use pMC benchmarks from [CJJK²²] and [2] (our paper on which Chapter 9 is based) as these models have more parameters than the traditional benchmarks. We extend these benchmarks to pRMCs by constructing probability intervals around the pMC's probabilities.

Results | The full results for all benchmarks are here omitted for brevity and instead presented in [4, Appendix B, Tab. 2–3]. In this chapter, we highlight the most important results, referring to [4] for the full details.

Q1. Computing solutions vs. derivatives

We investigate whether computing derivatives is feasible on p(R)MCs. In particular, we compare the computation times for computing derivatives on p(R)MCs (Problems 1 and 3) with the times for computing the solution for these models.

In Fig. 10.5, we show for all benchmarks the times for computing the solution (defined in Eqs. (10.1) and (10.3)), versus computing either a single derivative for Problem 10.8 (left) or the highest derivative of all parameters resulting from Problem 3 (right). A point (x, y) in the left plot means that computing a single derivative took x seconds while computing the solution took y seconds. A line above the (center) diagonal means we obtained a speed-up over the time for computing the solution; a point over the upper diagonal indicates a 10× speed-up or larger.

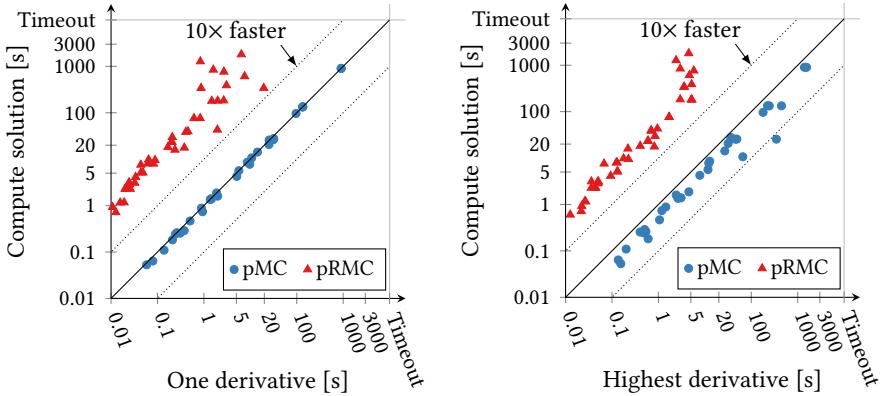


Figure 10.5: Runtimes (log-scale) for computing a single derivative (left, Problem 1) or the highest derivative (right, Problem 3), versus computing the solution $\text{sol}_{\mathbb{E}}[u]$ or $\text{sol}_{\mathbb{E},R}[u]$.

One derivative | The left plot in Fig. 10.5 shows that, for pMCs, the times for computing the solution and a single derivative are approximately the same. This is expected since both problems amount to solving a single equation system with $|S|$ unknowns. Recall that, for pRMCs, computing the solution means solving the LP in Eq. (10.13), while for derivatives, we solve an equation system. Thus, computing a derivative for a pRMC is relatively cheap compared to computing the solution, which is confirmed by the results in Fig. 10.5.

Highest derivative | The right plot in Fig. 10.5 shows that, for pMCs, computing the highest derivative is slightly slower than computing the solution (the LP to compute the highest derivative takes longer than the equation system to compute the solution). On the other hand, computing the highest derivative for a pRMC is still cheap compared to computing the solution. Thus, if we are using a pRMC anyway, computing the derivatives is relatively cheap.

Q2. Runtime improvement of computing only k derivatives

We want to understand the computational benefits of solving Problem 3 over solving Problem 10.8. For Q2, we consider all models with $|V| \geq 10$ parameters.

An excerpt of results for the grid world benchmarks is presented in Table 10.1. Recall that, after obtaining the (robust) solution, solving Problem 10.8 amounts to solving $|V|$ linear equation systems, whereas Problem 3 involves solving a single LP and k equations systems. From Table 10.1, it is clear that computing k derivatives is orders of magnitudes faster than computing all $|V|$ derivatives, especially if the number of parameters is high.

We compare computing all derivatives (Problem 10.8) versus only the $k = 1$ or 10 highest derivatives (Problem 3). The left plot of Fig. 10.6 shows the runtimes for $k = 1$, and the right plot for the $k = 10$ highest derivatives. The interpretation for Fig. 10.6 is the same as for Fig. 10.5. From Fig. 10.6, we observe that computing only the k highest derivatives leads to significant speed-ups, often of more than 10 times (except for very small models). Moreover, the difference between $k = 1$ and $k = 10$ is minor, showing that retrieving the actual derivatives after solving Problem 10.9 is relatively cheap.

Table 10.1: Model sizes, runtimes, and derivatives for selection of grid world models.

Type	XModel statistics			Verifying		Problem 1		Problem 3		Derivatives	
	S	\Gamma	#trans	sol _(R) [u]	Time [s]	All derivs. [s]	k = 1 [s]	k = 10 [s]	Highest	Error %	
pMC	5000	50	14995	5.07	1.39	3.32	2.64	2.69	1.54e+00	0.0	
pMC	5000	100	14995	5.05	1.36	4.17	2.63	2.66	1.28e+00	0.0	
pMC	5000	921	14995	4.93	1.87	19.92	4.52	2.87	1.20e+00	0.0	
pMC	80000	100	239995	8.01	25.54	98.47	45.18	46.87	1.95e+00	0.0	
pMC	80000	1000	239995	8.01	25.64	612.97	48.92	58.20	2.08e+00	0.0	
pMC	80000	9831	239995	7.93	25.52	5,650.25	347.76	1,343.59	2.10e+00	0.0	
pMC	1280000	100	3839995	12.90	902.52	4,747.43	1,396.51	1,507.77	3.32e+00	0.0	
pMC	1280000	1000	3839995	12.79	902.67	37,078.12	1,550.45	1,617.27	3.18e+00	0.0	
pMC	1280000	10000	3839995	Timeout ^b	—	—	—	—	—	—	
pRMC	5000	100	14995	136.07	23.46	3.55	0.60	1.58	-1.26e-02	-0.0	
pRMC	5000	921	14995	138.74	29.82	25.23	0.85	1.09	-4.44e-03	-0.0	
pRMC	20000	100	59995	2,789.77	1,276.43	15.68	2.40	2.70	-4.96e-01	-0.1	
pRMC	20000	1000	59995	2,258.41	339.96	159.70	3.53	4.09	-9.51e-02	-0.0	
pRMC	80000	100	239995	Timeout ^b	—	—	—	—	—	—	

^a Extrapolated from the runtimes for 10 to all $|V|$ parameters.

^b Timeout (1 hour) occurred for verifying the p(R)MC, not for computing derivatives.

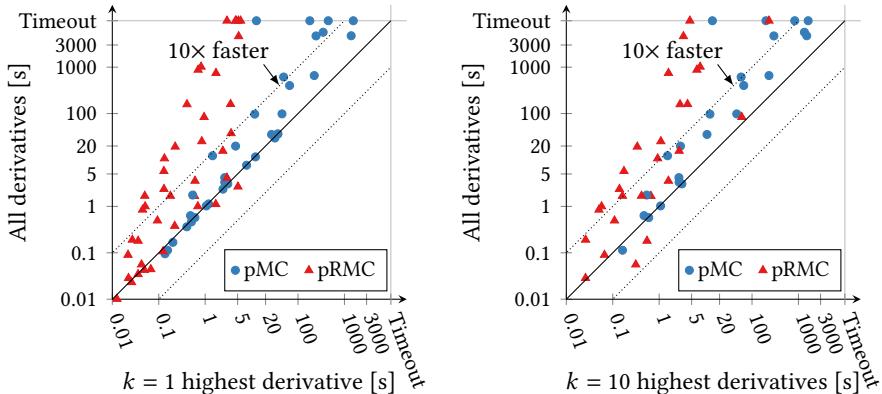


Figure 10.6: Runtimes (log-scale) for computing the highest (left) or 10 highest (right) derivatives (Problem 3), versus computing all derivatives (Problem 10.8).

Numerical stability | While our algorithm is exact, our implementation uses floating-point arithmetic for efficiency. To evaluate the numerical stability, we compare the highest derivatives (solving Problem 3 for $k = 1$) with an empirical approximation of the derivative obtained by perturbing the parameter by 10^{-3} . The difference (column ‘Error. %’ in Table 10.1) between both is marginal, indicating that our implementation is sufficiently numerically stable to return accurate derivatives.

Q3. Application in a learning framework

Reducing the sample complexity is a key challenge in learning under uncertainty [Kak03; MBPJ23]. Especially in stochastic environments, learning is very data-intensive, and realistic applications tend to require millions of samples to provide tight bounds on measures of interest [BHTB⁺18]. Motivated by this challenge, we apply our approach in a learning framework to investigate if derivatives can be used to effectively guide exploration, compared to alternative exploration strategies.

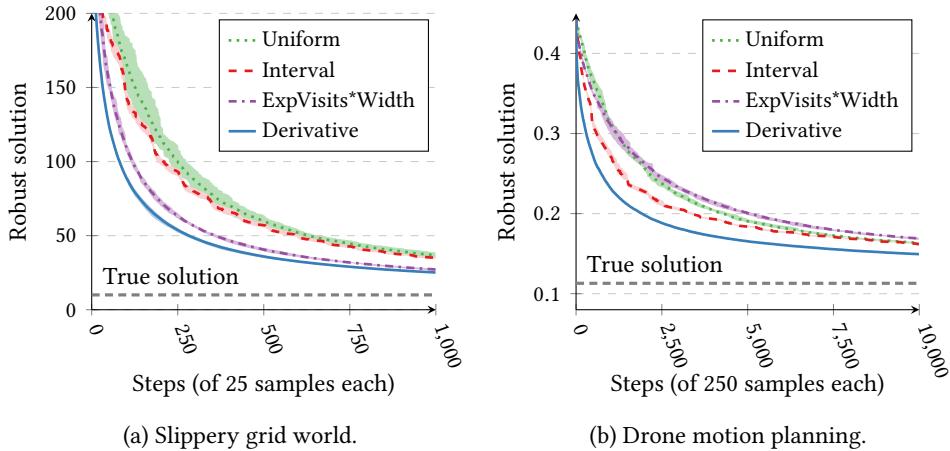


Figure 10.7: Robust solutions for each sampling strategy in the learning framework for the grid world (a) and drone (b) benchmarks. Averages values of 10 (grid world) or 5 (drone) repetitions are shown, with shaded areas showing the min/max of all repetitions.

Models | We consider the problem of where to sample in (1) a slippery grid world with $|S| = 800$ and $|V| = 100$ terrain types, and (2) the drone benchmark from [CJJK⁺22] with $|S| = 4\,179$ and $|V| = 1\,053$ parameters. As in the motivating example in Sect. 10.2, we learn a model of the unknown MC in the form of a pRMC, where the parameters are the sample sizes for each parameter. We assume access to a model that can arbitrarily sample each parameter (i.e., the slipping probability in the case of the grid world). We use an initial sample size of $N_i = 100$ for each parameter $i \in \{1, \dots, |V|\}$, from which we infer a $\beta = 0.9$ (90%) confidence interval using *Hoeffding's inequality*. The interval for parameter i is $[\hat{p}_i - \epsilon_i, \hat{p}_i + \epsilon_i]$, with \hat{p}_i the sample mean and $\epsilon_i = \sqrt{\frac{\log 2 - \log(1-\beta)}{2N}}$ (see, e.g., [BLM13] for details).

Hoeffding's inequality

Learning scheme | We iteratively choose for which parameter $v_i \in V$ to obtain 25 (for the grid world) or 250 (for the drone) additional samples. We compare four strategies for choosing the parameter v_i to sample for: (1) with the highest derivative, i.e., solving Problem 3 for $k = 1$; (2) with the biggest interval width ϵ_i ; (3) uniformly; and (4) sampling according to the expected number of visits times the interval width (see [4, Appendix B.1] for details). After each step, we update the robust upper bound on the solution for the pRMC with the additional samples.

Results | The upper bounds on the solution for each sampling strategy, as well as the solution for the MC with the true parameter values, are shown in Fig. 10.7. For both benchmarks, our derivative-guided sampling strategy converges to the true solution faster than the other strategies. Notably, our derivative-guided strategy accounts for both the uncertainty and importance of each parameter, which leads to a lower sample complexity required to approach the true solution.

10.7 Related Work

We discuss related work in three areas: pMCs, their extension to parametric interval Markov chains (pIMCs), and general sensitivity analysis methods.

Parametric Markov chains | As we also discussed in Chapter 3, pMCs [Daw04; LMT07] have traditionally been studied in terms of computing the solution function [HHZ11b; DJJC⁺15; BMS16; FTG16; FCGA21; JJK22]. For our paper, particularly relevant are [SJK21], which checks whether a derivative is positive (for all parameter valuations), and [HSJM⁺22], which solves parameter synthesis via gradient descent. We note that all these problems are (co-)ETR complete [JKPW21] and that the solution function is exponentially large in the number of parameters [BHHJ⁺20], whereas we consider a polynomial-time algorithm. Furthermore, practical *verification* procedures for uncontrollable parameters (as we do) are limited to less than 10 parameters. Parametric verification is used in [PWHA17] to guide model refinement by detecting for which parameter values a specification is satisfied. In contrast, we consider slightly more conservative RMCs and aim to stepwise optimize an objective. Solution functions also provide an approach to compute and refine confidence intervals [CGJP⁺16]; however, the size of the solution function hampers scalability.

Parametric interval Markov chains (pIMCs) | To the best of our knowledge, pRMCs have not been studied before. Nevertheless, their slightly more restricted version with *interval-valued* transition probabilities with parametric bounds, called parametric interval Markov chains (pIMCs) have been considered by some authors. Work on pIMCs falls into two categories. First, *consistency* [DLP16; PP18] asks whether there exists a parameter instantiation such that the (reachable fragment of the) induced interval MC contains valid probability distributions. Second, parameter synthesis for quantitative and qualitative reachability in pIMCs with up to 12 parameters [BDFL⁺18].

Perturbation analysis | Perturbation analysis considers the change in solution by any perturbation vector X for the parameter instantiation, whose norm is upper bounded by δ , i.e., $\|X\| \leq \delta$ (or conversely, which δ ensures the solution perturbation is below a given maximum). Likewise, [Cho19] uses the distance between two instantiations of a pMC (called augmented interval MC) to bound the change in reachability probability. Similar analyses exist for stationary distributions [ABH16]. These problems are closely related to the verification problem in pMCs and are equally (in)tractable if there are dependencies over multiple parameters. To improve tractability, a follow-up [SFCR16] derives asymptotic bounds based on first or second-order Taylor expansions. Other approaches to perturbation analysis analyze individual paths of a system [FH94; CC97; CW98]. Sensitivity analysis in (parameter-free) imprecise MCs, a variation to RMCs, is thoroughly studied in [CHQ08].

Exploration in learning | Similar to Q3 considered in Sect. 10.6, determining where to sample is relevant in many learning settings. Approaches such as probably approximately correct (PAC) statistical model checking [AKW19; AGKM22] and model-based reinforcement learning [MBPJ23] commonly use optimistic exploration policies [Mun14]. By contrast, we guide exploration based on the sensitivity analysis of the solution function with respect to the parametric model.

10.8 Discussion

We discuss interesting directions for further research along three dimensions: (1) computing derivatives for models with nondeterminism and applying this extension in a learning scheme, and (2) considering higher-order derivatives to obtain explicit bounds on the solution function.

Models with nondeterminism | Computing derivatives for models with nondeterminism, in particular parametric (robust) Markov decision processes, is an interesting yet challenging avenue for further research. The challenge is that each policy of the parametric MDP defines a different solution function, and one is typically interested in the maximum of all these functions. As a result, solution functions for parametric MDPs are nonsmooth, and to compute derivatives, one first has to fix a policy to resolve the nondeterminism. The problem is that different policies may have derivatives pointing in opposite directions. As a result, we may end up switching between different policies that have derivatives pointing in the opposite direction. Hence, we postulate that one falls more easily into local optima.

One idea is to encode the policy explicitly in the optimization problems that we formulated to compute derivatives. However, this renders our optimization problems nonconvex, because policy variables are multiplied with variables representing the parameters V . In the context of partially observable Markov decision processes (POMDPs), such nonconvex optimization problems have been (approximately) solved by convexifying them using different relaxations, such as convex-concave procedures [SJCT20; LB16] and sequential convex programming schemes [CJMJ⁺21; MSXA18]. However, it is questionable how informative derivatives from such a convexification of the problem still are. Nevertheless, we believe such a direction could be interesting for further research.

Embedding our techniques in model-based (reinforcement) learning | Extensions to models with nondeterminism are particularly interesting for model-based reinforcement learning [MBPJ23]. In this context, our methods could be used to guide learning in an MDP under the assumption that the parametric structure of its transition function is known. There is, however, one major problem with such MDP learning applications: The derivatives of parameters relevant to states that are never visited by the current policy are always zero. For example, in the grid world from Fig. 10.1a, the ground vehicle never visits the terrain types related to parameter v_3 , and thus, the derivative with respect to this parameter is zero. However, this derivative of zero does not mean that this parameter is not important at all: Under a different policy, the parameter could suddenly be relevant to sample more. This issue is a typical manifestation of the *exploration-exploitation* trade-off, which has received much attention in the reinforcement learning community [SB98; KLM96]. Dealing with exploration-exploitation trade-offs in an optimal yet tractable manner is a wide-open problem.

Bounds on solution function | In this chapter, we have focused on computing the gradient of the solution function, that is, the vector of all partial derivatives. The gradient provides information on how the solution function changes with (infinitesimal) changes in the parameters, but not with respect to larger changes in the parameters. For example, our methods can not immediately be used to perform a perturbation analysis (see the related work), where the goal is to obtain an upper and lower bound

on the solution function over a region in the parameter space. Still, one could try to construct a Taylor expansion of the solution function, by considering not just the (first-order) but also the higher-order partial derivatives. However, the size of this Taylor expansion is exponential in the number of parameters. For example, for the second-order expansion, one needs to consider differentiating with respect to all combinations of two parameters. Thus, it remains an open research question to what extent such an approach is tractable and whether the resulting bounds on the solution function are worth the high computational expenses.

Summary

- ▷ We introduced parametric robust MCs (pRMCs), which combine uncertainty in transition probabilities with dependencies between states.
- ▷ Partial derivatives of the solution function are an effective measure of how values change with respect to the parameters.
- ▷ Computing only the k highest parameters is orders of magnitudes faster than computing all $|V|$ parameters.
- ▷ Partial derivatives can effectively guide sampling to increase the sample efficiency in a learning framework.

Part IV

Continuous-Time Markov Chains

11 Foundations of CTMCs

Summary | Continuous-time Markov chains (CTMCs) are stochastic processes subject to random timing that can be interpreted as the continuous-time counterpart of discrete-time Markov chains (DTMCs). In this background chapter, we present the fundamentals of CTMCs and how to perform common analyses. We also introduce parametric CTMCs (pCTMCs), which relax the assumption that transition rates are precisely known and instead model transition rates as (polynomials over) parameters.

Origins | None of the material presented in this chapter is novel, and much of the background on CTMCs originates from the seminal papers [BHK03; ASSB00].

Background | While not strictly required to understand this chapter, familiarity with MDPs (Chapter 3) and pMDPs (Chapter 8) may help the reader to understand the basics of CTMCs and pCTMCs.



11.1 Introduction

Thus far, we have considered stochastic models in which the state evolves dynamically over *discrete* time steps. In practice, such *discrete-time models* are often the result of discretizing time in a system of differential equations. For example, an epidemic can abstractly be modeled using a so-called SIR (susceptible-infected-recovered) model [AB12], which is shown in Fig. 11.1 for a population of two. Initially, one person is susceptible (S), and the other is infected (I). From this initial state, there is a certain probability that the infected person infects the susceptible person within one discrete time step (and similarly, with some probability, the infected person recovers before infecting the other).

A practical consideration when using such a model is what time period every discrete step should resemble. If we choose a small time period (e.g., one minute), we obtain a very detailed but potentially also very expensive model to analyze. On the other hand, if we choose a long time period (e.g., a month), we obtain a much coarser model but also lose a lot of detail and potentially incur a high approximation error. This raises the following natural question. Can we avoid discretizing time and directly reason about a continuous-time model where transitions between states can occur at any point in time?

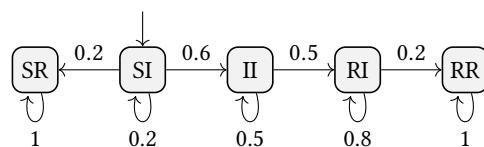


Figure 11.1: DTMC for the SIR model with a population of two.

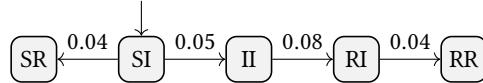


Figure 11.2: CTMC for the SIR model with a population of two, where labels on the edges represent the transition rates.

11.2 Continuous-Time Markov Chains

In this final part of the thesis, we consider such continuous-time models that alleviate the challenges mentioned above. Specifically, we study continuous-time Markov chains (CTMCs), which are stochastic processes subject to random timing. While transitions in a discrete-time Markov chain (DTMC) occur at fixed discrete time steps, transitions in a CTMC can occur at *any continuous point in time*. CTMCs are widely used to model complex probabilistic systems in reliability engineering [RS15], network processes [HHK00; HMS99], systems biology [CDPK⁺17; BS18] and epidemic modeling [All17].

Recall that $\text{Distr}(X)$ is the set of all distributions over a set X . Formally, we define a CTMC as follows. Note that, without loss of generality, we define CTMCs with an initial state *distribution*, rather than a *single* initial state (as we did for Markov decision processes (MDPs) in Chapter 3).¹

Definition 11.1 (CTMC) A (labeled) *continuous-time Markov chain (CTMC)* is a tuple $C := (S, s_I, R, L)$ with a finite set S of *states*, an *initial state distribution* $s_I \in \text{Distr}(S)$, a *transition rate function* $R: S \times S \rightarrow \mathbb{Q}_{\geq 0}$, and a *labeling function* $L: S \rightarrow 2^{AP}$, where AP is a set of atomic propositions.

Intuitively, a rate $R(s, s') > 0$ implies there is a *transition* from state $s \in S$ to $s' \in S$. If $R(s, s') > 0$ for more than one state $s' \in S$, a *race condition* between the transitions exists. Let $E: S \rightarrow \mathbb{Q}_{\geq 0}$ be the so-called *exit rates*, where $E(s) = \sum_{s' \in S} R(s, s')$ is the total rate at which any outgoing transition of state $s \in S$ is taken. Then, the probability $P(s, s', t)$ of transitioning from $s \in S$ to $s' \in S$ within $t \in \mathbb{R}_{\geq 0}$ time units is

$$P(s, s', t) = \frac{R(s, s')}{E(s)} \cdot \left(1 - e^{-E(s) \cdot t}\right).$$

The exit rate $E(s)$ characterizes the *residence time* for state $s \in S$. Specifically, the probability to reside in state s at most $t \in \mathbb{R}_{\geq 0}$ time units is $1 - e^{-E(s) \cdot t}$. Hence, the average residence time of state s is $\frac{1}{E(s)}$. CTMCs exhibit the Markov specification in the following way: The probability of leaving the state $s \in S$ within the next $t \in \mathbb{R}_{\geq 0}$ time units is independent of how many time units ago the model entered state s .

Example 11.2 (Epidemic model as CTMC) In reality, getting infected and recovering from the disease in the epidemic DTMC model in Fig. 11.1 should not be restricted to discrete time steps. Thus, a more realistic model for the epidemic is the

¹Recall from Remark 10.1 that we can always transform a model with an initial state distribution into one with a single initial state and vice versa.

CTMC shown in Fig. 11.2 [All17]. The label on each edge between states s and s' is the transition rate $R(s, s')$ of the CTMC. Observe that the CTMC does not explicitly model self-loops (as is the case in the DTMC in Fig. 11.1 does); instead, self-loops in CTMCs are implicit.

transient distribution

Transient distribution | The *transient probability distribution* of a CTMC specifies the probability distribution over states as a function of time, given some initial distribution $\bar{s} \in \text{Distr}(S)$. Concretely, the transient distribution $P_{\bar{s}}(t) \in \text{Distr}(S)$ from state distribution \bar{s} after time $t \geq 0$ is

$$P_{\bar{s}}(t) = \bar{s} \cdot e^{(R - \text{diag}(E)) \cdot t},$$

where $\text{diag}(E)$ is the diagonal matrix with the exit rates E (interpreted as a vector) on the diagonal. Intuitively, $P_{\bar{s}}(t)(s') \in [0, 1]$ is the probability of being in state $s' \in S$ at time t , when starting from the state distribution \bar{s} at time zero.

Similarly, for a fixed starting state $s \in S$, we can compute the transient distribution for the Dirac distribution δ_s (as defined in Chapter 2) for state s :

$$P_{\delta_s}(t) = \delta_s \cdot e^{(R - \text{diag}(E)) \cdot t}.$$

Example 11.3 (Transient distribution for epidemic CTMC) For the epidemic CTMC in Fig. 11.2, consider a state ordering [SI, SR, II, RI, RR]. Then, the initial state is $s_I = [1, 0, 0, 0, 0]$. We wish to compute the transient distribution from s_I at time $t = 10$, which is $P_{s_I}(10) = s_I \cdot e^{(R - \text{diag}(E)) \cdot 10}$. Based on the state ordering, we can interpret the transition rate function R as a matrix and the exit rates E as a vector:

$$R = \begin{bmatrix} 0 & 0.04 & 0.05 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0.08 & 0 \\ 0 & 0 & 0 & 0 & 0.04 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}, \quad E = \begin{bmatrix} 0.09 \\ 1 \\ 0.08 \\ 0.04 \\ 1 \end{bmatrix},$$

where we choose an (arbitrary) nonzero rate of 1 for the sink states, which is needed later on (for applying the uniformization technique). Thus, we have that

$$R - \text{diag}(E) = \begin{bmatrix} -0.09 & 0.04 & 0.05 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & -0.08 & 0.08 & 0 \\ 0 & 0 & 0 & -0.04 & 0.04 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix},$$

which leads to the transient distribution $P_{s_I}(10) = [0.41, 0.26, 0.21, 0.10, 0.02]$. This transient distribution tells us that after 10 time units, the CTMC is in state SI with probability 0.41, in state SR with probability 0.26, and so on.

The matrix exponent involved in the transient distribution can be computed using a Taylor-Maclaurin expansion such that

$$\mathbf{P}_{\bar{s}}(t) = \bar{s} \cdot e^{(R - \text{diag}(E)) \cdot t} = \bar{s} \cdot \sum_{i=0}^{\infty} \frac{((R - \text{diag}(E)) \cdot t)^i}{i!}. \quad (11.1)$$

However, as discussed in more detail by [ML03; Ste94], using this expansion is impractical for two main reasons:

1. It is difficult to find a proper criterion for truncating the infinite summation;
2. Numerical methods can suffer from instability due to the non-sparsity of $(R - \text{diag}(E)) \cdot t$ and the presence of both positive and negative entries.

Thus, most practical techniques for computing transient distributions resort to a technique called *uniformization* [BHHK03], which we describe in more detail in Def. 11.11.

Paths | An *infinite path* in the CTMC $C = (S, s_I, R, L)$ is an alternating sequence $\pi := s_0 t_0 s_1 t_1 s_2 \dots$ of states and residence times, where $R(s_i, s_{i+1}) > 0$ for all $i \in \mathbb{N}$. Similarly, a *finite* CTMC path is an alternating sequence $\pi := s_0 t_0 s_1 t_1 \dots t_{\ell-1} s_\ell$ such that $R(s_i, s_{i+1}) > 0$ for all $i < \ell$ and s_ℓ is absorbing, i.e., $R(s_\ell, s') = 0$ for all $s' \in S$. The set of all (finite and infinite) paths in the CTMC C is denoted by $\Pi^C = S \times (\mathbb{R}_{\geq 0} \times S)^*$. As for MDPs, we denote the set of all (in)finite CTMC paths starting in a state $s \in S$ by $\Pi^C(s)$.

Intuitively, the finite path $s_0 3 s_1 4 s_2$ means we stayed exactly 3 time units in s_0 , then transitioned to s_1 , where we stayed 4 time units before moving to s_2 . For a path $\pi = s_0 t_0 s_1 t_1 \dots$, let $\pi_i = s_i$ denote the $(i + 1)^{\text{th}}$ state of π , and $\delta_i(\pi) = t_i$ the time spent in s_i . Furthermore, the CTMC state occupied at time $t \in \mathbb{R}_{\geq 0}$ is denoted by $\pi(t) \in S$.

Example 11.4 (Path for epidemic CTMC) Consider the finite CTMC path $\pi = s 3 s' 4 s''$. For this path, we have $\pi_0 = s$, $\delta_0(\pi) = 3$, $\pi_1 = s'$, and $\delta_1(\pi) = 4$. Similarly, we find that the state occupied at time 6.2 is $\pi(6.2) = s'$.

Alternative view | An alternative (and equivalent; see [Kat16]) view of CTMCs is to combine the exit rates $E: S \rightarrow \mathbb{Q}_{\geq 0}$ with a transition matrix $\Delta: S \rightarrow \text{Distr}(S)$, which is defined such that $\Delta(s, s') = \frac{R(s, s')}{E(s)}$ for all $s, s' \in S$. For the absorbing states, the exit rate $E(s)$ is zero, and we also set $\Delta(s, s') = 0$ for all $s' \in S$. For consistency, we primarily represent CTMCs by their transition rate function R (and not by their exit rates E and transition matrix Δ).

Probability measure | We can define a probability measure over infinite timed paths by cylinder sets [Kat16]. Here, we only sketch the construction while referring to, for example, [BHHK03] for formal details. The main idea is to define cylinder sets over infinite paths, which is similar yet more complex than the construction for DTMCs because we need to consider the residence times in states along a path. As such, let $s_0 I_0 s_1 I_1 \dots I_{\ell-1} s_\ell$ be an (interval) path fragment, where for all $j = 0, \dots, \ell-1$, $I_j = [a_j, b_j]$ is an interval with $a_j \leq b_j$, and $a_j, b_j \in \mathbb{R}_{\geq 0}$. For this interval path fragment, we define the (interval) cylinder set as the set of all CTMCs paths that have $s_0 t_0 s_1 t_1 \dots t_{\ell-1} s_\ell$ as a prefix, where each $t_j \in I_j$. Doing so, we can define the *probability measure* \Pr^C on paths in the CTMC C (with initial distribution s_I).

Fault trees | CTMCs are closely related to (dynamic) *fault trees*, which are a common modeling formalism in reliability engineering [RS15]. Fault trees are widely used to predict how component faults lead to system failure in safety and economically critical assets [RRBS19]. A fault tree is a *directed acyclic graph*, whose leaves are *basic events* that represent component faults, and whose root is a *top event* related to system failure. The basic and top events are connected through logical *gates* that represent failure propagation. Fault trees can be represented as a CTMC, and thus, techniques for analyzing CTMCs can also be applied to fault trees. We refer to [Vol22] for a more comprehensive introduction to (dynamic) fault trees.

fault tree

Continuous-time MDPs | CTMCs can be extended with nondeterministic action choices, resulting in a continuous-time Markov decision process (CTMDP) [Kat16; GH09; BHHK05]. Schedulers for CTMDPs not only decide which action to play based on the states previously visited but also on the elapsed time in every state [NSK09; Mil68]. Thus, a CTMDP has uncountably many deterministic schedulers (compared to countably many for MDPs; see Chapter 3 for details), and existing algorithms resort to discretization to obtain ϵ -optimal schedulers in practice [Kat16; BFKK⁺13]. We do not study CTMDPs in this thesis, and we refer to the references above for details.

11.3 Verifying CTMCs

Over the past decades, efficient verification algorithms have been developed for the analysis of CTMCs [BHHK03; ASSB00; KKNP01]. In this section, we provide a brief introduction to specifications for CTMCs and computing measures of interest.

11.3.1 Continuous stochastic logic

Continuous stochastic logic (CSL) is a branching-time temporal logic that is commonly used for specifying CTMC specifications. The logic was developed by [ASSB00] and later extended with a time-bounded until operator by [BHHK03]. While we primarily focus on measures of reachability and reward in this thesis, we still present the syntax of continuous stochastic logic (CSL), as taken from [BHHK03], for completeness.

continuous
stochastic
logic

Definition 11.5 (Syntax of CSL) Let AP be a set of atomic propositions. A CSL formula over AP is built using the following grammar:

$$\begin{aligned}\Phi &::= \text{True} \mid a \mid \neg\Phi \mid \Phi \wedge \Phi \mid \mathbb{L}_{\triangleleft\lambda}(\Phi) \mid \mathbb{P}_{\triangleleft\lambda}(\varphi) \\ \varphi &::= \bigcirc^I \Phi \mid \Phi \cup^I \Phi,\end{aligned}$$

where $a \in AP$ is an atomic proposition, $\triangleleft \in \{<, \leq, \geq, >\}$ is a comparison operator, $\lambda \in [0, 1]$ is a threshold probability, and $I \subseteq \mathbb{R}_{\geq 0}$ is a nonempty interval.

As for probabilistic computation tree logic (PCTL) introduced in Def. 3.14, Φ is called a *state formula* and φ a *path formula*. The symbols in the Def. 11.5 are defined as follows. First, $\mathbb{L}_{\triangleleft\lambda}(\Phi)$ holds if the steady-state (also called *long-run*) probability of occupying a state that satisfies the state formula Φ meets the condition $\triangleleft\lambda$. Analogous to the probabilistic operator in PCTL, $\mathbb{P}_{\triangleleft\lambda}(\varphi)$ holds if the probability measure of the paths satisfying the path formula φ meets the condition $\triangleleft\lambda$. The path formula $\bigcirc^I \Phi$ holds if a

state
and path
formula

transition to a state satisfying Φ is made within the interval I . Finally, the path formula $\Phi \cup^I \Phi'$ holds if Φ' is satisfied at some time in the interval I and Φ holds at all times before that.

Remark 11.6 (Omitting labeling function) For simplicity, we often define state formulae directly over CTMC states $s \in S$ instead of over atomic propositions $a \in AP$. In such a case, we implicitly assume that the states and atomic propositions coincide and that the labeling function L is defined such that $L(s) = \{s\}$ for all $s \in S$.

Example 11.7 (CSL formula for epidemic CTMC) For the CTMC in Fig. 11.2, we can consider the CSL state formula $\mathbb{P}_{\geq 0.9}(\text{True} \cup^{[0,10]} \{\text{SR}, \text{RR}\})$, which is satisfied if, starting from the initial state SI, the probability of reaching a state in which the disease has become extinct (namely, states SR and RR) within the first 10 time units is at least 0.9.

Semantics | We consider the semantics of CSL as presented in [BHHK03]. We use the common notation $s \models \Phi$ to denote that state $s \in S$ satisfies the state formula Φ . Similarly, we write $\pi \models \varphi$ to denote that CTMC path π satisfies the path formula φ .

Definition 11.8 (Satisfaction of CSL) Let $C = (S, s_I, R, L)$ be a CTMC, $s \in S$ be a state, $a \in AP$ be an atomic proposition, Φ, Φ' be CSL state formulae, and φ be a CSL path formula. The satisfaction relation \models is defined for state formulae as

$$\begin{aligned} s \models a &\quad \text{iff } a \in L(s) \\ s \models \neg\Phi &\quad \text{iff } s \not\models \Phi \\ s \models \Phi \wedge \Phi' &\quad \text{iff } s \models \Phi \text{ and } s \models \Phi' \\ s \models \mathbb{L}_{\leq \lambda}(\Phi) &\quad \text{iff } \lim_{t \rightarrow \infty} \Pr^C(\pi \in \Pi^C(s) : \pi(t) \models \Phi) \leq \lambda \\ s \models \mathbb{P}_{\leq \lambda}(\varphi) &\quad \text{iff } \Pr^C(\pi \in \Pi^C(s) : \pi \models \varphi) \leq \lambda. \end{aligned}$$

Similarly, the satisfaction relation for a path π in C is defined as

$$\begin{aligned} \pi \models \bigcirc^I \Phi &\quad \text{iff } \pi_1 \models \Phi \wedge \delta_0(\pi) \in I \\ \pi \models \Phi \cup^I \Phi' &\quad \text{iff } \exists t \in I. (\pi(t) \models \Phi' \wedge (\forall t' \in [0, t). \pi(t') \models \Phi)). \end{aligned}$$

Derived operators | Much like for PCTL, we can derive more operators from the syntax in Def. 11.5. First, we define the standard (untimed) next and until operators as

$$\bigcirc \Phi := \bigcirc^{[0,\infty)} \Phi \quad \text{and} \quad \Phi \cup \Phi' := \Phi \cup^{[0,\infty)} \Phi'.$$

The *eventually* operator \diamond and its step-bounded variant \diamond^I are then defined as

$$\mathbb{P}_{\leq \lambda}(\diamond \Phi) := \mathbb{P}_{\leq \lambda}(\text{True} \cup \Phi) \quad \text{and} \quad \mathbb{P}_{\leq \lambda}(\diamond^I \Phi) := \mathbb{P}_{\leq \lambda}(\text{True} \cup^I \Phi).$$

Similarly, the *always* operator \square and its step-bounded variant \square^I are defined as

$$\mathbb{P}_{\leq \lambda}(\square \Phi) := \mathbb{P}_{\geq (1-\lambda)}(\text{True} \cup \neg \Phi) \quad \text{and} \quad \mathbb{P}_{\leq \lambda}(\square^{\leq I} \Phi) := \mathbb{P}_{\geq (1-\lambda)}(\text{True} \cup^{\leq I} \neg \Phi),$$

where \triangleright is the opposite comparison operator from \triangleleft , e.g., if \triangleleft is \leq then \triangleright is \geq .

11.3.2 Measures

In this thesis, we are primarily interested in *quantitative* aspects of CTMCs, without fixing a threshold probability λ for the satisfaction of a formula upfront. While the interpretation of CSL formulae is boolean (i.e., a formula is either satisfied or not), we can also use the CSL syntax to express common *measures*. In this section, we discuss the measures that we use for CTMCs in this thesis.

measure
(for CTMC)

Transient probabilities | Often, we wish to compute the probability of reaching a subset of CTMC states within a given time interval. We express such *transient (satisfaction) probability* in the following way.²

Definition 11.9 (Transient satisfaction probability) Let $C = (S, s_I, R, L)$ be a CTMC, $s \in S$ be a state, and φ be a CSL path formula. The *transient satisfaction probabilities* of formula φ in state s is defined as

$$\Pr^C(s \models \varphi) := \Pr^C(\pi \in \Pi^C(s) : \pi \models \varphi).$$

transient
probability

An unbounded transient reachability probability with respect to an initial state $s \in S$ and the target states $S_T \subseteq S$ is a particular instance of Def. 11.9 with a path formula of the form $\varphi = \diamond S_T$, i.e.,

$$\Pr(s \models \diamond S_T).$$

Similarly, a step-bounded transient reachability probability with respect to an initial state $s \in S$, the target states $S_T \subseteq S$ and a time interval $I \subseteq \mathbb{R}_{\geq 0}$ is of the form $\Pr(s \models \diamond^I S_T)$. Transient satisfaction probabilities for other path formulae are defined analogously.

Steady-state probabilities | In a very similar manner, we can define *steady-state (satisfaction) probability* with respect to state formulae.

steady-state
probability

Definition 11.10 (Steady-state satisfaction probability) Let $C = (S, s_I, R, L)$ be a CTMC, $s \in S$ be a state, and Φ be a CSL state formula. The *steady-state satisfaction probabilities* of formula Φ in state s is defined as

$$\text{St}^C(s \models \Phi) := \lim_{t \rightarrow \infty} \Pr^C(\pi \in \Pi^C(s) : \pi(t) \models \Phi).$$

11

Intuitively, $\text{St}^C(s \models \Phi)$ is the steady-state probability to be in a state satisfying the state formula Φ , starting from state $s \in S$.

Expected reward | We may augment CTMCs with a reward structure (ρ, ι) , where $\rho: S \rightarrow \mathbb{R}_{\geq 0}$ is a state-based reward function, and $\iota: S \times S \rightarrow \mathbb{R}_{\geq 0}$ is a transition-based reward function. For CTMCs, state rewards are obtained for the time during which a state is visited, while transition rewards are obtained upon entering a state. For example, consider again the example CTMC path $\pi = s3s'4s''$ from Example 11.4. In this path, a state reward of $3\rho(s)$ is obtained while being in state s and of $4\rho(s')$ while being in state s' . By contrast, a transition reward of $\iota(s, s')$ is obtained upon transitioning from

²The word *transient* refers the time interval I in the CSL path formula φ for which the satisfaction probability is defined.

expected cumulative reward s to s' , and of $\iota(s', s'')$ upon transitioning from s' to s'' . While some of the developed methods in this thesis are amenable to *expected reward measures* for CTMCs, we omit them for simplicity and instead refer to [KNP06; DHK15] for details.

11.3.3 Algorithms

Mature and efficient algorithms have been developed for model checking CTMCs against CSL formulae. Especially the seminal papers [ASSB00] and [BHHK03] have laid much of the foundations for several popular CTMC model checking algorithms. Depending on the CSL formula, model checking requires solving a linear equation system (for unbounded until and the steady-state operator) or computing transient probabilities in the CTMC (for time-bounded until formulae).

Uniformization | In practice, transient probabilities in CTMCs can be computed using *uniformization* (also referred to as *randomization*), a technique first proposed by Jensen in 1953 [Jen53]. The main idea is to transform the CTMC into a DTMC that captures the transient behavior of the CTMC. More concretely, the uniformized DTMC for a given CTMC is defined as follows.

uniformiza-
tion

Definition 11.11 (Uniformization) Let $C = (S, s_I, R, L)$ be a CTMC, and fix a *uniformization rate* $q \in \mathbb{R}_{>0}$ such that $q \geq \max_{s \in S} E(s)$. The *uniformized DTMC* $\mathcal{D} = (S, s_I, P, L)$ has the same states S , initial rate s_I and labeling function L , and the transition function $P: S \rightarrow \text{Distr}(S)$ is defined for all $s, s' \in S$ as

$$P(s)(s') = \begin{cases} \frac{E(s)}{q} \cdot \Delta(s, s') & \text{if } s' \neq s \\ \frac{E(s)}{q} \cdot (\Delta(s, s') - 1) + 1 & \text{otherwise,} \end{cases}$$

where $E: S \rightarrow \mathbb{Q}_{\geq 0}$ and $\Delta: S \rightarrow \text{Distr}(S)$ are the exit rates and transition matrix for the CTMC C .

Note that the uniformization rate q must be at least the highest exit rate $E(s)$ over all states $s \in S$. Uniformization is possible for any CTMC and allows computing transient probabilities for a CTMC by analyzing the uniformized DTMC instead. Specifically, the transient distribution at time t for the initial state distribution \bar{s} can be rewritten using the transition function P of the uniformized DTMC as

$$P_{\bar{s}}(t) = \bar{s} \cdot e^{(R - \text{diag}(E)) \cdot t} = \bar{s} \cdot \sum_{i=0}^{\infty} e^{-q \cdot t} \frac{(q \cdot t)^i}{i!} \cdot P^i. \quad (11.2)$$

Observe that the term $\frac{(q \cdot t)^i}{i!}$ decreases rapidly with i . Furthermore, we can compute the required number of terms in Eq. (11.2) such that $P_{\bar{s}}(t)$ is approximated with any desired precision [BHHK03]. As a result, uniformization leads to numerically stable techniques for computing transient probabilities and CSL model checking of CTMCs [KKNP01]. For a more comprehensive treatment of uniformization for CTMCs and the resulting model checking algorithms, we refer to [GM84; DBB18; BHHK03].

Implementation in model checkers | These algorithms, and in particular the uniformization technique, are implemented in probabilistic model checkers, such as

PRISM [KNP11] and Storm [HJKQ⁺22]. In this thesis, we especially use Storm to analyze CTMCs. For a given CTMC and a CSL formula, these model checkers are able to verify whether the formula is satisfied or not. Furthermore, these model checkers can compute common measures, including the transient probabilities, steady-state probabilities, and expected rewards discussed in Sect. 11.3.2. We omit further details about the algorithms used by these model checkers, because the contributions of this thesis are independent of these algorithms. Instead, we regard the model checker as an *oracle* that, given a CSL formula or measure, provides a solution to the verification problem at hand.

11.4 Parametric Continuous-Time Markov Chains

Standard CTMC verification algorithms require that the transition rates are precisely known; an assumption that is often unrealistic in practice [HKM08]. To this end, parametric CTMCs (pCTMCs) extend standard CTMC with transition rates given as polynomials over parameters [HKM08; CCGK⁺18].

Let V be a set of parameters. The set of polynomials over parameters V with rational coefficients is denoted by $\mathbb{Q}[V]$. We formally define a pCTMC as follows.

Definition 11.12 (pCTMC) A (labeled) *pCTMC* is a tuple $C_V := (S, s_I, V, R, L)$, with a finite set S of *states*, an *initial state distribution* $s_I \in \text{Distr}(S)$, an (ordered) set of *parameters* V , a *parametric transition rate function* $R: S \times S \rightarrow \mathbb{Q}[V]$, and a labeling function $L: S \rightarrow 2^{AP}$, where AP is a set of atomic propositions.

parametric
CTMC

Observe that the difference with a standard CTMC is the rate function, which now maps every pair of states $s, s' \in S$ to a polynomial over the set of parameters V . Assuming that the parameters V in Def. 11.12 are ordered will conveniently allow us to reason over values of the parameters as vectors.

11.4.1 Parameter instantiation

Given a pCTMC, we can fix a value for each parameter $v \in V$. Assigning a value to the parameters is modeled by an *instantiation* $u: V \rightarrow \mathbb{Q}$, which is a function that maps parameters to concrete values. Using the fact that the parameters $V = \{v_1, \dots, v_n\}$, $n \in \mathbb{N}_{>0}$, are ordered, we will often denote an instantiation as the vector $u \in \mathbb{Q}^{|V|}$.

parameter
instanti-
ation

Induced CTMC | Recall from Chapter 8 that applying the instantiation u to a polynomial $g \in \mathbb{Q}[V]$ yields $g[u] \in \mathbb{Q}$, which is obtained by substituting every parameter $v \in V$ in g with $u(v)$. Thus, applying an *instantiation* u to a pCTMC C_V yields an *induced CTMC* $C_V[u] := (S, s_I, V, R[u])$, where $R[u](s, s') := R(s, s')[u]$ for all $s, s' \in S$. Intuitively, the transition rate function $R[u]$ is obtained by assigning the value $u(v)$ to each parameter $v \in V$. For the induced CTMC, we can perform any of the CSL model checking queries discussed in Sect. 11.3.

11

Well-defined instantiations | The transition rate function of a CTMC maps pairs of states to nonnegative rational numbers. However, some instantiations may lead to an induced CTMC with *negative* transition rates. To avoid such situations, we restrict ourselves to instantiations that yield a valid transition rate function. We call the set of such instantiations the *parameter space* \mathcal{V}_{C_V} for pCTMC C_V .

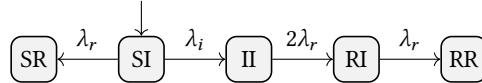


Figure 11.3: PCTMC for the SIR model with a population of two, where λ_i and λ_r model the infection and recovery rate, respectively.

parameter space

Definition 11.13 (Parameter space) The *parameter space* \mathcal{V}_{C_V} for a pCTMC $C_V = (S, s_I, V, R, L)$ is a subset of all instantiations $u: V \rightarrow \mathbb{Q}$ defined as

$$\mathcal{V}_{C_V} = \{u: V \rightarrow \mathbb{Q} : R[u](s, s') \in \mathbb{Q}_{\geq 0} \ \forall s, s' \in S\}.$$

For any pCTMC C_V , we only consider instantiations $u \in \mathcal{V}_{C_V}$ that belong to the parameter space \mathcal{V}_{C_V} .

Example 11.14 (Instantiations for epidemic pCTMC) We continue our example of the SIR epidemic model. A parametric version of this CTMC is shown in Fig. 11.3, where we replaced the concrete transition rates by parameters λ_i and λ_r , modeling the infection and recovery rate, respectively. Thus, the set of parameters of this pCTMC is $V = \{\lambda_i, \lambda_r\}$. Consider now the following two instantiations:

- The instantiation $u: V \rightarrow \mathbb{Q}$ defined as $u(\lambda_i) = 0.5$ and $u(\lambda_r) = 0.4$ induces precisely the (parameter-free) CTMC shown in Fig. 11.2. Observe that this instantiation leads to transition rates that belong to $\mathbb{Q}_{\geq 0}$, i.e., all rates are nonnegative rational numbers. Thus, this instantiation u is valid and belongs to the parameter space \mathcal{V}_{C_V} .
- By contrast, the instantiation $u': V \rightarrow \mathbb{Q}$ defined as $u'(\lambda_i) = -0.2$ and $u'(\lambda_r) = -0.1$ does not lead to a well-defined transition rate (because negative rates are now allowed), and thus, $u' \notin \mathcal{V}_{C_V}$.

11.4.2 Verifying pCTMCs

A standard problem for pCTMCs is to find all parameter instantiations $u \in \mathcal{V}_{C_V}$ that induce a CTMC $C_V[u]$ that satisfies a given CSL specification. A similar query is to determine whether there exists an instantiation $u \in \mathcal{V}_{C_V}$ such that the induced CTMC $C_V[u]$ satisfies the specification. Such model checking queries are similar to those for parametric Markov decision processes (pMDPs) as we discussed in Chapter 8, and for which [JJK22; Jun20] provide more detailed treatments.

Verification of parametric CTMCs is investigated in [HKM08; CDPK⁺17]; however, the resulting algorithms are generally restricted to a few parameters. In general, verifying pCTMCs remains a challenging problem, which typically scales poorly with the number of parameters. In the next chapter, we present a sampling-based verification method that does not suffer from this limited scalability. Instead, our method assumes access to a prior distribution over the parameter values of the pCTMC, which we use to obtain verification results with statistical guarantees.

Solution function | The intuition that every parameter instantiation $u \in \mathcal{V}_{C_V}$ induces a concrete CTMC provides a natural construct to lift measures from CTMC to pCTMC. In particular, we can define a function, called the *solution function*, that maps every instantiation $u \in \mathcal{V}_{C_V}$ to the value of a (fixed) measure on the induced CTMC $C_V[u]$. For example, the solution function for the transient satisfaction probability of a CSL path formula φ is formally defined as follows.

solution
function

Definition 11.15 (Solution function for transient probability) Let

$C_V = (S, s_I, V, R, L)$ be a pCTMC with parameter space \mathcal{V}_{C_V} . The *solution function* $\text{sol}_\varphi^{C_V} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}$ for the transient satisfaction probability of a CSL path formulae φ on the pCTMC C_V is defined as

$$\text{sol}_\varphi^{C_V} : u \mapsto \Pr^{C_V[u]}(s_I \models \varphi).$$

The solution function for the steady-state satisfaction probability of a CSL state formula Φ is defined analogously:

Definition 11.16 (Solution function for steady-state probability) Let $C_V = (S, s_I, V, R, L)$ be a pCTMC with parameter space \mathcal{V}_{C_V} . The *solution function* $\text{sol}_\Phi^{C_V} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}$ for the steady-state satisfaction probability of a CSL state formulae Φ on the pCTMC C_V is defined as

$$\text{sol}_\Phi^{C_V} : u \mapsto \text{St}^C(s_I \models \Phi).$$

Again, the solution function for an expected reward measure can be defined analogously, but we omit an explicit definition for brevity.

11.5 Challenges

Standard verification algorithms for CTMCs assume that the model dynamics, including the initial state distribution, are precisely known. That is, given a CTMC (or a pCTMC with a parameter instantiation) with an initial state distribution $s_I \in \text{Distr}(S)$, we can verify any of the CSL formulae or measures discussed in Sect. 11.3.

When dealing with CTMCs in realistic settings, however, the dynamics and/or initial state are often subject to uncertainty. The verification of CTMCs subject to uncertainty, such as the pCTMCs discussed in the previous section, is a significantly less developed research area. In Chapters 12 and 13, we study two such settings for CTMCs with uncertainty in either the transition rates or the initial state. To close this chapter, we briefly introduce these settings and the respective research questions we consider.

Chapter 12: Uncertain transition rates | In Chapter 12, we study pCTMC with uncertain transition rates. We model this setting as a pCTMC together with a prior distribution over the parameter values that encodes uncertainty about the actual transition rates. Each sample of this prior yields a standard CTMC that we can analyze, e.g., by computing any of the measures discussed in Sect. 11.3. However, the outcome of this analysis may be different for each sample from the prior. Thus, the main question that we aim to answer is as follows: “How can we use the analysis outcomes for previous samples to make predictions about the analysis outcome for yet another sample?”

Chapter 13: Uncertain initial state | In Chapter 13, we study CTMCs where the initial state is unknown and must instead be inferred from a sequence of previously observed labels. This setting is motivated by applications such as runtime monitoring, which involves analyzing an already running system without a static initial state. The main question we consider is thus: “*How can we make predictions about the CTMC, conditioned on a sequence of previously observed labels (and their observation times)?*”

Summary

- ⇒ Continuous-time Markov chains (CTMCs) are stochastic processes subject to random timing and are the continuous-time analog of DTMCs.
- ⇒ Specifications for CTMCs are commonly expressed in continuous stochastic logic (CSL).
- ⇒ We can use the CSL syntax to define common measures of reachability and reward for CTMCs.
- ⇒ Parametric CTMCs extend CTMCs with parametric transition rates.
- ⇒ Assigning a value to each of the parameters of a pCTMC is called an instantiation and yields an induced CTMC.

12 CTMCs With Uncertain Rates

Summary | We study parametric CTMCs (pCTMCs) with a (possibly unknown) prior distribution over the parameters. The prior encodes uncertainty about the actual transition rates, while the parameters allow dependencies between transition rates. Sampling the parameter values from the prior distribution yields a standard CTMC, for which we may compute relevant measures, such as reachability probabilities. We provide a principled solution, based on a technique called scenario optimization, to the following problem: From a finite set of parameter samples and a user-specified confidence level, compute prediction regions on the reachability probabilities. The prediction regions should (with high probability) contain the reachability probabilities of a CTMC induced by any additional sample. To boost the scalability of the approach, we employ standard abstraction techniques and adapt our methodology to support approximate reachability probabilities. Experiments with various well-known benchmarks show the applicability of the approach.

Origins | The contents of this chapter are based on:

[3] Badings, Jansen, Junges, Stoelinga and Volk (2022) ‘Sampling-Based Verification of CTMCs with Uncertain Rates’. CAV.

Additional experimental results left out of this thesis can be found in [3, Appendix C].

Background | The reader is assumed to be familiar with (parametric) continuous-time Markov chains (CTMCs) and their analysis, as discussed in Chapter 11.



12.1 Introduction

In Chapter 11, we have introduced continuous-time Markov chains (CTMCs) as stochastic processes subject to random timing. Standard model checking algorithms for CTMCs require that the transition rates are precisely known; an assumption that is often unrealistic in practice [HKM08]. For example, consider again the SIR (susceptible-infected-recovered) model from Fig. 11.2. This CTMC assumes *fixed and known values* for the infection rate (denoted here as λ_i) and the recovery rate (λ_r). Suppose that we are interested in the probability of the disease becoming extinct over time. The outcome of this analysis for fixed values of λ_i and λ_r may be a *probability curve* like in Fig. 12.1a, where we plot the probability (y-axis) of reaching a target state that corresponds to the epidemic becoming extinct against varying time horizons (x-axis).¹ In fact, the plot is obtained via a smooth interpolation of the results at finitely many horizons, as depicted in Fig. 12.1b. However, what if the rates we used for this analysis turn out to be inaccurate?

¹For visual clarity, we plot the reachability probability *between* time 100 and t_1, \dots, t_N .

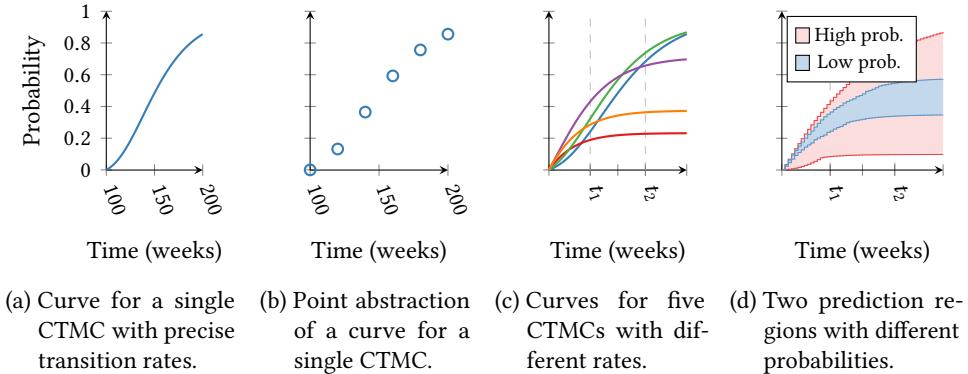


Figure 12.1: The probability of extinction in the SIR (140) model for horizons $[100, t]$.

Parametric models | To acknowledge that the rates λ_i and λ_r are, in fact, uncertain, we may specify the model as the parametric CTMC (pCTMC) in Fig. 11.3. We can then analyze this pCTMC for different values of λ_i and λ_r , resulting in a set of curves as in Fig. 12.1c. These individual curves, however, provide *no guarantees* about the curve obtained from yet another infection and recovery rate. Furthermore, what if some parameter values are more likely than others? In other words, how can we incorporate *prior knowledge* about the parameter values into the analysis?

Prior on parameters | In this chapter, we consider a setting in which we have access to a *prior distribution* over the transition rates of a CTMC, which is similar to assumptions made in, e.g., [BS18; MMAG14]. These priors may result from asking different experts which value they would assume for, e.g., the infection rate. The prior may also be the result of Bayesian reasoning [WA19]. We capture the uncertainty in the rates by an arbitrary and potentially unknown *probability distribution* over the parameter space of the pCTMC. We call this model an *uncertain parametric CTMC* (*upCTMC*); see Sect. 12.2 for a formal definition. The distribution allows drawing independent and identically distributed (i.i.d.) *samples*, each of which yields a (parameter-free) CTMC.

Overall goal | For the epidemic modeled as a pCTMC, we may, for example, assume that the infection and recovery rates λ_i and λ_r are both normally distributed. Each time we sample a pair of parameters, we obtain a concrete CTMC for which we can compute a probability curve as in Fig. 12.1c. Intuitively, we want to use these curves to make a statement about the curve that we would obtain for *yet another sample of the parameters*. More concretely, we aim to construct so-called *prediction regions* around a set of probability curves, as those shown in Fig. 12.1d. Then, with high probability and high confidence, sampling a set of transition rates should induce a probability curve within this prediction region. In this chapter, we develop an efficient *probably approximately correct*, or PAC-style, method that computes these prediction regions.

Outline | This chapter is structured as follows. In Sect. 12.2, we introduce uncertain parametric CTMCs (upCTMCs) as the model that we study in this chapter, and we formally define the problem of interest. In Sect. 12.3, we solve this problem assuming

that we can verify the CTMCs obtained for the different parameter values exactly. In Sect. 12.4, we consider a less restricted setting in which the verification result for every parameter value is imprecise, i.e., only known to lie in a certain interval. We discuss algorithmic improvements for computing verification results in Sect. 12.5, and we evaluate our approach by performing numerical experiments in Sect. 12.6.

12.2 CTMCs With Uncertain Rates

Let us formalize the model that we consider in this chapter. We propose to extend pCTMCs (as defined by Def. 11.12) with distributions over the parameter values. We call the resulting model an *uncertain parametric CTMC (upCTMC)*.

Definition 12.1 (upCTMC) An *uncertain parametric CTMC (upCTMC)* is a tuple (C_V, \mathbb{P}) , where $C_V = (S, s_I, V, R, L)$ is a pCTMC and \mathbb{P} is a probability distribution over the parameter space \mathcal{V}_{C_V} of C_V .

uncertain
parametric
CTMC

In Def. 12.1, we implicitly assume a probability space $(\mathcal{V}_{C_V}, \mathcal{B}(\mathcal{V}_{C_V}), \mathbb{P})$, where the sample space is the parameter space \mathcal{V}_{C_V} , the σ -algebra is the Borel σ -algebra $\mathcal{B}(\mathcal{V}_{C_V})$, and \mathbb{P} is the probability measure. For brevity, we omit the full probability space and instead only refer to the distribution \mathbb{P} over the parameter space.

Remark 12.2 (Probability measure \mathbb{P}) The probability measure \mathbb{P} used throughout this chapter should not be confused with the probabilistic operator in continuous stochastic logic (CSL), defined in Def. 11.5. In this chapter, \mathbb{P} always denotes a probability measure over the parameter space of a pCTMC.

The semantics of a upCTMC are as follows. A upCTMC specifies a probability distribution over the parameter values of pCTMC C_V , whose domain is defined by the parameter space \mathcal{V}_{C_V} . We denote a *sample* from \mathcal{V}_{C_V} drawn according to \mathbb{P} by $u \in \mathcal{V}_{C_V}$. Recall from Sect. 11.4 that applying this instantiation $u \in \mathcal{V}_{C_V}$ to the pCTMC induces the CTMC denoted by $C_V[u]$. Thus, a upCTMC implicitly defines a *probability distribution over CTMCs* with concrete transition rates.

Sampling from \mathbb{P} | We make the following assumption on the distribution \mathbb{P} , which is analogous to Assumption 6.2 made in Chapter 6.

Assumption 12.3 (Distribution unknown) We merely assume i.i.d. sampling access to the probability distribution \mathbb{P} of a upCTMC, and that the Radon-Nikodym derivative of \mathbb{P} with respect to the Lebesgue measure exists. However, the distribution itself can be highly complex or even unknown.

12

The existence of the Radon-Nikodym derivative is a rather standard assumption in probability theory [Dur10]. A consequence of this assumption is that the probability of drawing two samples $u, u' \in \mathcal{V}_{C_V}$ according to the probability measure \mathbb{P} that are exactly the same, $u = u'$, is zero. This assumption is standard in the literature on scenario optimization [CG18a; CCG21], which is the sample-based methodology that we use to solve the problem we consider in this chapter.

12.2.1 Measures and solution functions

Recall from Sect. 11.3 that measures for CTMCs can be specified in CSL. Furthermore, recall from Sect. 11.4 that the satisfaction probability for a CSL path formulae φ (or state formula Φ) can be lifted to a pCTMC C_V using the solution function $\text{sol}_{\varphi}^{C_V} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}$ (or $\text{sol}_{\Phi}^{C_V} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}$). This solution function maps every parameter instantiation $u \in \mathcal{V}_{C_V}$ to the satisfaction probability on the induced CTMC $C_V[u]$.

Remark 12.4 (Notation for solution function) In Sect. 11.4, we denoted a solution function as $\text{sol}_{\varphi}^{C_V}$ or $\text{sol}_{\Phi}^{C_V}$ to emphasize that it is defined for a fixed pCTMC C_V and for the satisfaction probability of a CSL path formula φ or state formulae Φ . For brevity, we simplify notation in this chapter and write a solution function as $\text{sol} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}$, thus making the pCTMC and CSL formula implicit.

In this chapter, we generalize the concept of a solution function to (ordered) sets of measures. To this end, we introduce the notion of a *vector-valued solution function*.

vector-valued solution function

Definition 12.5 (Vector-valued solution function) Let $C_V = (S, s_I, V, R, L)$ be a pCTMC with parameter space \mathcal{V}_{C_V} . A *vector-valued solution function* $\text{sol} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}^m$ for the pCTMC C_V is a function that maps every instantiation $u \in \mathcal{V}_{C_V}$ to a value in \mathbb{R}^m , for some $m \in \mathbb{N}_{>0}$.

The dimension of a vector-valued solution function will be clear from the context and is, thus, often omitted. Furthermore, a standard solution function is a special case of a vector-valued solution function mapping to \mathbb{R}^1 . Thus, we also drop the *vector-valued* and simply refer to *solution functions* in this chapter. For brevity, we also refer to $\text{sol}(u) \in \mathbb{R}^m$ as the *solution vector* of instantiation $u \in \mathcal{V}_{C_V}$.

Example 12.6 (Solution function for epidemic model) Consider again the point abstraction of a probability curve in Fig. 12.1b. This point abstraction consists of six time points t_1, \dots, t_6 at which the probability of the disease becoming extinct is computed. Thus, this setting can be modeled by the (vector-valued) solution function $\text{sol} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}^6$, which is defined for all $u \in \mathcal{V}_{C_V}$ as

$$\text{sol}(u) = \left[\Pr^{C_V[u]}(s_I \models \Diamond^{\leq t_1} S_T), \dots, \Pr^{C_V[u]}(s_I \models \Diamond^{\leq t_6} S_T) \right]^\top,$$

where $S_T \subset S$ are the states corresponding with the disease being extinct, and $t_1, \dots, t_6 \in \mathbb{R}_{>0}$ are the six time points at which to compute the measure.

12.2.2 Problem statement

Let (C_V, \mathbb{P}) be a upCTMC and let $\text{sol} : \mathcal{V}_{C_V} \rightarrow \mathbb{R}^m$ be a solution function modeling $m \in \mathbb{N}_{>0}$ measures. We want to make predictions about the solution vector $\text{sol}(u)$ for a random instantiation $u \in \mathcal{V}_{C_V}$ drawn according to \mathbb{P} . Intuitively, we represent this prediction as a *prediction region*, denoted by R , on the codomain \mathbb{R}^m of the solution function sol , such as those shown in Fig. 12.1d. We consider only prediction regions that are compact subsets of \mathbb{R}^m , yielding the following definition.

Definition 12.7 (Prediction region) An m -dimensional *prediction region* R is a compact subset of \mathbb{R}^m , where $m \in \mathbb{N} \cup \{\infty\}$.

prediction
region

Note that a prediction region can be infinite-dimensional, which we will use in Def. 12.10 to define prediction regions over time (as those shown in Fig. 12.1d).

We define the so-called *containment probability* of a prediction region R , which is the probability that the solution vector $\text{sol}(u)$ for a randomly sampled parameter $u \in \mathcal{V}_{C_V}$ is contained in R .

Definition 12.8 (Containment probability) Let (C_V, \mathbb{P}) be a upCTMC with parameter space \mathcal{V}_{C_V} , let R be a prediction region, and let $\text{sol}: \mathcal{V}_{C_V} \rightarrow \mathbb{R}^m$ be a solution function. The *containment probability* $\Pr_{\mathcal{V}_{C_V}}(R)$ is the probability that $\text{sol}(u)$ is contained in R for u sampled according to \mathbb{P} :

containment
probability

$$\Pr_{\mathcal{V}_{C_V}}(R) := \mathbb{P}\{u \in \mathcal{V}_{C_V} : \text{sol}(u) \in R\}.$$

Sampling parameter values | Our goal is to compute a prediction region R and provide guarantees on the containment probability $\Pr_{\mathcal{V}_{C_V}}(R)$. We will achieve this using a sampling-based approach. In particular, let $\mathcal{U}_N = \{u_1, \dots, u_N\}$ be a set of $N \in \mathbb{N}_{>0}$ i.i.d. samples from \mathcal{V}_{C_V} drawn according to \mathbb{P} . This set of samples is an element from the probability space $\mathcal{V}_{C_V}^N = \times_{i=1}^N \mathcal{V}_{C_V}$ equipped with the product probability measure \mathbb{P}^N and the product σ -algebra. We aim to use the set of samples $\mathcal{U}_N = \{u_1, \dots, u_N\}$ from \mathbb{P} to obtain *statistical guarantees* on the containment probability, i.e., guarantees that hold with a (user-specified) confidence level, denoted by $\beta \in (0, 1)$. Formally, we solve the following problem.

Problem 12.9 (Prediction region for upCTMC) Given a upCTMC (C_V, \mathbb{P}) , a solution function $\text{sol}: \mathcal{V}_{C_V} \rightarrow \mathbb{R}^m$, and a confidence level $\beta \in (0, 1)$, compute a (tight) prediction region R and a (tight) lower bound $\mu \in (0, 1)$ on the containment probability, such that

$$\mathbb{P}^N\left\{\{u_1, \dots, u_N\} \in \mathcal{V}_{C_V}^N : \Pr_{\mathcal{V}_{C_V}}(R) \geq \mu\right\} \geq \beta,$$

where $\Pr_{\mathcal{V}_{C_V}}(R) = \mathbb{P}\{u \in \mathcal{V}_{C_V} : \text{sol}(u) \in R\}$ is the containment probability for the solution function sol and the prediction region R .

To explore the intuition of Problem 12.9, let us neglect the confidence probability β for a moment. The solution function sol implicitly encodes a set of measures expressed as CSL formulae (see Remark 12.4). Thus, we aim to compute a prediction region R on the values of these measures, such that when we draw a parameter value $u \in \mathcal{V}_{C_V}$ according to \mathbb{P} , the value $\text{sol}(u)$ of these measures is contained in R .

12

Role of confidence probability β | Problem 12.9 asks for an algorithm that yields a “correct” guarantee with at least the confidence probability β . That is, the algorithm draws a set of N samples $\{u_1, \dots, u_N\} \in \mathcal{V}_{C_V}^N$ according to the product probability measure \mathbb{P}^N . Then, for at least a β probability of \mathbb{P}^N , the algorithm should produce a prediction region R for which the containment probability $\Pr_{\mathcal{V}_{C_V}}(R)$ is lower bounded

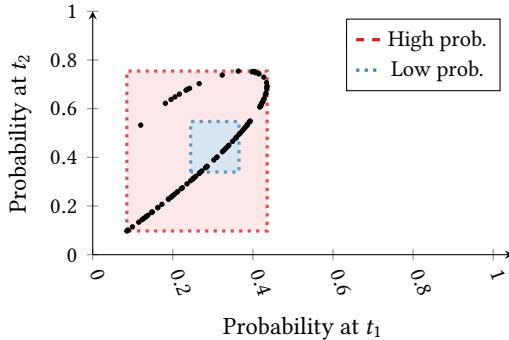


Figure 12.2: Prediction region for the epidemic CTMC, showing the extinction probabilities from Fig. 12.1c at time points t_1 and t_2 , together with two prediction regions around these points.

by μ . Naturally, we expect a trade-off between the size of the prediction region R and the lower bound μ on the containment probability. Two other important parameters are the number of samples N and the confidence probability β . We now illustrate Problem 12.9 through an example, highlighting the role of these different parameters.

12.2.3 Illustrative example

We reinterpret the problem sketched in Fig. 12.1 (where we considered the extinction probability of a disease over time) in the context of Problem 12.9. The solution function for this example problem is analogous to Example 12.6, but then with more than six time bounds. How can we solve Problem 12.9 for this example?

A change in perspective | Suppose that we are given the set $\{u_1, \dots, u_N\}$ of $N \in \mathbb{N}_{>0}$ values for the parameters, drawn according to \mathbb{P} . Each sample u_i leads to a solution vector $\text{sol}(u_i)$, as shown in Fig. 12.1c (recall that each curve is, in fact, an interpolation of finitely many points). To solve Problem 12.9, we want to compute a prediction region that overapproximates this set of N solution vectors.

To explain how we achieve this, let us change the perspective and look at only the solution vectors for two horizons, namely t_1 and t_2 (which are also indicated in Fig. 12.1c). We represent the solution vectors for these two horizons as in Fig. 12.2, where each point is a pair of probabilities that the disease becomes extinct before time t_1 and before t_2 . That is, each point in Fig. 12.2 is a vector for an instantiation $u \in \mathcal{V}_{CV}$ defined as

$$\left[\Pr^{CV[u]}(s_I \models \diamond^{\leq t_1} S_T), \Pr^{CV[u]}(s_I \models \diamond^{\leq t_2} S_T) \right] \in \mathbb{R}^2.$$

The boxes in Fig. 12.2 suggest how we can overapproximate the solution vector as a rectangle. However, when we map this rectangle back to the perspective in Fig. 12.1, we obtain a prediction region that is only defined at the time points at which the measure is computed. Instead, we want to obtain a prediction region that contains the curves in Fig. 12.1c for all time points $t \in [100, 200]$. In other words, we want a prediction region defined as a set-valued function over time.

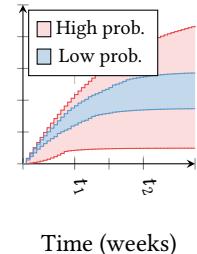
Definition 12.10 (Prediction region over time) A prediction region $R \subset \mathbb{R}^m$ over time is given by two curves $\underline{c}, \bar{c}: \mathbb{Q}_{\geq 0} \rightarrow \mathbb{R}$ as the area in-between:

$$R = \{(t, y) \in \mathbb{Q} \times \mathbb{R} \mid \underline{c}(t) \leq y \leq \bar{c}(t)\}.$$

Restricting the domain of t in Def. 12.10 to the rational numbers $\mathbb{Q}_{\geq 0}$ ensures that CSL model checking problems over these time bounds are decidable [ASSB00]. We need this restriction because our approach for solving Problem 12.9 is based on solving a set of CSL model checking problems over the time bounds of a prediction region.

In practice, we exploit the monotonicity² of the measure with respect to the time bound. Thus, the functions \underline{c}, \bar{c} in Def. 12.10 are obtained as two step functions. If desired, we can smoothen the resulting prediction region by taking an upper and lower bound on these step functions.

Problem 12.9 solved | Using Def. 12.10, we lift the two prediction regions in Fig. 12.2 to the prediction regions in Fig. 12.1d (copied here for convenience), which are defined over the continuous time interval $t \in [100, 200]$. When looking closely at Fig. 12.1d, one can indeed see the step functions that make up the prediction regions. These prediction regions provide a solution to Problem 12.9. For this specific example, using a confidence level of $\beta = 99\%$ and considering $N = 100$ curves, we conclude that $\mu = 79.4\%$ for the red region and $\mu = 7.5\%$ for the blue region. For a higher confidence level of $\beta = 99.9\%$, we would obtain slightly more conservative bounds on the containment probability. This concludes our intuitive explanation of how we solve Problem 12.9.



Time (weeks)

12.2.4 Our approach

To actually compute prediction regions, we use techniques from the *scenario approach* (also called *scenario optimization*), a data-driven methodology for solving stochastic optimization problems [CG18a; CCG21]. Starting with a set of solution vectors $N \in \mathbb{N}_{>0}$, denoted by $\{\text{sol}(u_i)\}_{i=1}^N$, we construct a prediction region based on the solution to a convex optimization problem. Our method can balance the size of the prediction region with the containment probability, as illustrated by the two boxes in Fig. 12.1d.

scenario approach

Extensions | Our approach offers more than prediction regions on probability curves as in Fig. 12.1. For example, we also allow for solution vectors that represent *multiple objectives*, such as the reachability with respect to different goal states, expected rewards, or even the probability mass of paths satisfying more complex temporal specifications. In our experiments, we show that this multi-objective approach—also on probability curves—yields much tighter bounds on the containment probability than an approach that analyzes each objective independently. We can also produce prediction regions as other shapes than boxes, as, for example, shown in Fig. 12.3.

²In this example, only the upper limit on the time bound is varied, so the resulting measure is monotonically increasing with the time bound.

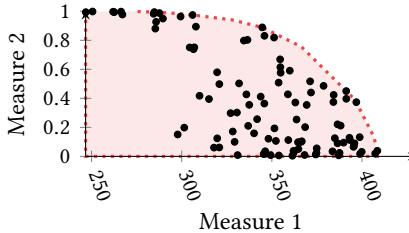


Figure 12.3: With our approach, we can also compute prediction regions of different shapes, such as this Pareto front for two measures.

12.3 Precise Sampling-Based Prediction Regions

In this section, we use scenario optimization [CG08; CG18a] to compute a prediction region that solves Problem 12.9. First, in Sect. 12.3.1, we describe how to compute a prediction region using the solution vectors $\{\text{sol}(u_i)\}_{i=1}^N$ for a given set of $N \in \mathbb{N}_{>0}$ parameter samples. In Sect. 12.3.2, we clarify how to compute a lower bound on the containment probability with respect to this prediction region. In Sect. 12.3.3, we present an algorithm based on those results that solves Problem 12.9.

12.3.1 Constructing prediction regions

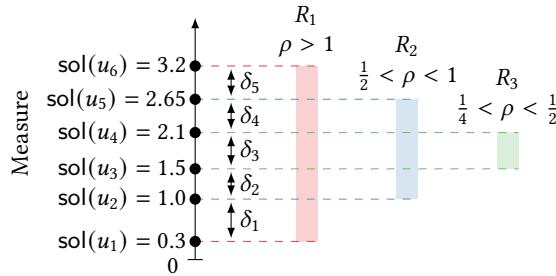
For conciseness, we restrict ourselves to R being a hyperrectangle in \mathbb{R}^m , with m the dimension of the solution function sol , and briefly describe extensions in Remark 12.15 below. Thus, we can represent R using two vectors $\underline{x}, \bar{x} \in \mathbb{R}^m$ such that, using pointwise inequalities, $R = \{x \in \mathbb{R}^m : \underline{x} \leq x \leq \bar{x}\}$. An example of such a rectangular prediction region is depicted earlier by Fig. 12.2.

Relaxation of solutions | As shown in Fig. 12.2, we do *not* require R to contain all solutions $\{\text{sol}(u_i)\}_{i=1}^N$. Instead, we have two orthogonal goals: We aim to minimize the size of R , while also minimizing the L^1 distance (also known as Manhattan distance) of the solutions $\{\text{sol}(u_i)\}_{i=1}^N$ to R . We call the solutions that are not contained in R *relaxed*. These goals define a *multi-objective problem*, which we solve by weighting the two objectives using a fixed parameter $\rho > 0$, called the *cost of relaxation*, that is used to scale the distance to R . Then, $\rho \rightarrow \infty$ enforces $\text{sol}(u_i) \subseteq R$ for all $i = 1, \dots, N$, as in the outer box in Fig. 12.2, while for $\rho \rightarrow 0$, R is reduced to a point. Thus, the cost of relaxation ρ is a tuning parameter that determines the size of the prediction region R and hence the fraction of the solution vectors that is contained in R (see [CG18b; CCG21] for details).

Scenario optimization problem | We capture the problem described above in the following convex *scenario optimization problem* $\mathfrak{L}_{\mathcal{U}}^\rho$:

$$\mathfrak{L}_{\mathcal{U}}^\rho : \underset{\substack{\underline{x} \in \mathbb{R}^m, \bar{x} \in \mathbb{R}^m, \\ \xi_i \in \mathbb{R}_{\geq 0}^m \forall i=1,\dots,N}}{\text{minimize}} \quad \|\bar{x} - \underline{x}\|_1 + \rho \sum_{i=1}^n \|\xi_i\|_1 \quad (12.1a)$$

$$\text{subject to } \underline{x} - \xi_i \leq \text{sol}(u_i) \leq \bar{x} + \xi_i \quad \forall i = 1, \dots, N. \quad (12.1b)$$

Figure 12.4: The prediction region changes with the cost of relaxation ρ .

The decision variables $\underline{x}, \bar{x} \in \mathbb{R}^m$ represent the lower and upper bound of the prediction region, respectively. Furthermore, for every $i = 1, \dots, N$, the decision variable $\xi_i \in \mathbb{R}_{\geq 0}^m$ is a slack variable representing the distance to R . The objective function in Eq. (12.1a) minimizes the size of R (by minimizing the sum of the width of the prediction region in all dimensions) plus ρ times the distances of the sampled solutions to R . We denote the (arguments of the) optimal solution to problem \mathfrak{L}_U^ρ for a given ρ by R_ρ^*, ξ_ρ^* , where $R_\rho^* = [\underline{x}_\rho^*, \bar{x}_\rho^*]$ for the rectangular case.

Proposition 12.11 A finite optimal solution R_ρ^*, ξ_ρ^* to \mathfrak{L}_U^ρ exists.

Proof. The solution vectors $\{\text{sol}(u_i)\}_{i=1}^N$ are finite-valued by definition. Thus, the constraints of problem \mathfrak{L}_U^ρ form a non-empty compact set on the decision variables $\underline{x}, \bar{x}, \xi$. Minimizing a linear objective over a compact set always has a finite optimum, and thus, a finite optimal solution to \mathfrak{L}_U^ρ exists. ■

While an optimal solution to \mathfrak{L}_U^ρ exists, it may not be unique, as illustrated by the following example.

Example 12.12 Fig. 12.4 shows a set of (1-dimensional) solutions for six sampled parameter instantiations, labeled $\text{sol}(u_1), \dots, \text{sol}(u_6)$, which we ordered without loss of generality. The figure also shows three prediction regions, R_1 , R_2 , and R_3 , which are obtained by solving \mathfrak{L}_U^ρ for different values of ρ . The values $\delta_1, \dots, \delta_5$ are the distances between the solutions, as shown in Fig. 12.4.

First consider prediction region $R_1 = [\text{sol}(u_1), \text{sol}(u_6)] = [0.3, 3.2]$. The corresponding objective value Eq. (12.1a) is

$$\|\bar{x} - \underline{x}\| + \rho \cdot \sum \xi_i = \|3.2 - 0.3\| = \delta_1 + \delta_2 + \delta_3 + \delta_4 + \delta_5,$$

as $\xi_i = 0$ for all $i = 1, \dots, 6$. Similarly, for region $R_2 = [\text{sol}(u_2), \text{sol}(u_5)] = [1.0, 2.65]$, where $\text{sol}(u_1)$ and $\text{sol}(u_6)$ are relaxed, the objective value is

$$\begin{aligned} \|\bar{x} - \underline{x}\| + \rho \cdot \sum \xi_i &= \|2.65 - 1.0\| + \rho(3.2 - 2.65) + \rho(1.0 - 0.3) \\ &= \delta_2 + \delta_3 + \delta_4 + \rho \cdot \delta_1 + \rho \cdot \delta_5. \end{aligned}$$

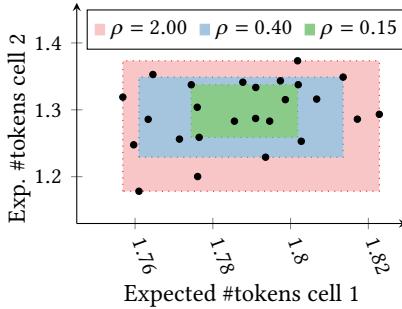


Figure 12.5: Prediction regions as boxes for different costs of relaxations ρ .

For $\rho > 1$, the objective value of R_1 is smaller than that of R_2 , so solving $\mathfrak{L}_{\mathcal{U}}^{\rho}$ yields the prediction region R_1 . By contrast, for $\rho < 1$ the opposite holds, so R_2 is optimal over R_1 . For exactly $\rho = 1$, the objective values of R_1 and R_2 are equal, so the optimal solution to $\mathfrak{L}_{\mathcal{U}}^{\rho}$ is not unique. Similarly, for $\rho < \frac{1}{2}$, relaxing samples $\text{sol}(u_1), \text{sol}(u_2), \text{sol}(u_5)$, and $\text{sol}(u_6)$ is cost-optimal, resulting in the prediction region $R_3 = [\text{sol}(u_3), \text{sol}(u_4)]$ (and for $\rho = \frac{1}{2}$ both R_2 and R_3 are optimal).

To ensure the uniqueness of optimal solutions, we make the following assumption.

Assumption 12.13 The optimal solution $R_{\rho}^{\star}, \xi_{\rho}^{\star}$ to $\mathfrak{L}_{\mathcal{U}}^{\rho}$ is unique.

If the solution to Eq. (12.1) is not unique, we apply a suitable tie-break rule that selects one solution of the optimal set (e.g., the solution with a minimum Euclidean norm, see [CG08]). Doing so, we can always satisfy Assumption 12.13 in practice.

Properties of prediction regions | Let us explore the geometrical properties of prediction regions obtained from solving problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$. Recall that the parameter ρ , called the cost of relaxation, is a tuning parameter that determines the size of the prediction region R . This effect is illustrated by the following example.

Example 12.14 We consider the Kanban manufacturing system benchmark from [CT96] with a Gaussian distribution over the parameters. In Fig. 12.5, we present $N = 25$ solution vectors for two expected cost measures. We use these solutions in problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ in Eq. (12.1) and solve for $\rho = 2, 0.4$, and 0.15 , resulting in the three prediction regions in Fig. 12.5. For $\rho = 2$, the prediction region contains all vectors, while for a lower cost of relaxation ρ , more vectors are left outside.

Interestingly, any prediction region obtained as the optimal solution to problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ in Eq. (12.1) has at least one sample exactly on its boundary. This property is also observed in Fig. 12.5 and can be explained based on the objective function $\|\bar{x} - \underline{x}\|_1 + \rho \sum_{i=1}^n \|\xi_i\|_1$. If, for a feasible solution $\underline{x}, \bar{x}, \xi$, there is no sample on the boundary of the box $[\underline{x}, \bar{x}]$, then we can reduce the size of the prediction region (thus decreasing the term $\|\bar{x} - \underline{x}\|_1$) without having to increase the term $\rho \sum_{i=1}^n \|\xi_i\|_1$.

Remark 12.15 (Shape of prediction region) While problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ in Eq. (12.1) yields a rectangular prediction region, we can also produce other shapes. We may, e.g., construct a Pareto front as in Fig. 12.3, by adding additional affine constraints [BV14]. In fact, our only requirement is that the objective function is convex, and the constraints are convex in the decision variables (the dependence of the constraints on u may be arbitrary) [CCG21].

12.3.2 Bounding the containment probability

The previous section shows how we compute a prediction region based on convex optimization. We now characterize a sound statistical lower bound on the containment probability with respect to the prediction region given by the optimal solution to this optimization problem. Toward that result, we introduce the so-called *complexity* of a solution to problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ in Eq. (12.1), a concept used in [CCG21] that is related to the *compressibility* of the solution vectors $\{\text{sol}(u_i)\}_{i=1}^N$. Roughly, the idea of this compressibility is that only some of the solution vectors are needed to retain the same optimal solution to problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ (and thus the same prediction region).

Definition 12.16 (Complexity) For $\mathfrak{L}_{\mathcal{U}}^{\rho}$ with optimal solution $R_{\rho}^{\star}, \xi_{\rho}^{\star}$, consider a subset of samples $\mathcal{W} \subseteq \mathcal{U}_N$ and the associated problem $\mathfrak{L}_{\mathcal{W}}^{\rho}$ with optimal solution $\tilde{R}_{\rho}, \tilde{\xi}_{\rho}$. The set \mathcal{W} is called *critical* if

$$\tilde{R}_{\rho} = R_{\rho}^{\star} \quad \text{and} \quad \left\{ u_i \in \mathcal{U}_N : \xi_{\rho,i}^{\star} > 0 \right\} \subseteq \mathcal{W}.$$

The *complexity* c_{ρ}^{\star} of $R_{\rho}^{\star}, \xi_{\rho}^{\star}$ is the cardinality of the smallest critical set. For brevity, we also call c_{ρ}^{\star} the complexity of $\mathfrak{L}_{\mathcal{U}}^{\rho}$.

complexity

Intuitively, the complexity c_{ρ}^{\star} is the number of samples $i = 1, \dots, N$ not contained in the prediction region R_{ρ}^{\star} (because for these samples, we have $\xi_{\rho,i}^{\star} > 0$), plus the minimum number of samples needed on the boundary of the region to keep the solution unchanged (that is, $\tilde{R}_{\rho} = R_{\rho}^{\star}$). We describe in Sect. 12.3.3 how we algorithmically determine (an upper bound on) the complexity.

Example 12.17 In Fig. 12.5, the prediction region for $\rho = 2$ contains all solution vectors, so $\xi_{2,i}^{\star} = 0 \forall i$. Moreover, if we remove *all but four* solutions (the ones on the boundary of the region), the optimal solution to problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ remains unchanged, so the complexity is $c_{2,0}^{\star} = 0 + 4$. Similarly, the complexity for $\rho = 0.4$ is $c_{0.4}^{\star} = 8 + 2 = 10$ (8 solutions outside the region and 2 on the boundary).

Recall that Def. 12.8 defines the containment probability of a generic region R . Here, we consider the containment probability $\Pr_{\mathcal{V}_{CV}}(R_{\rho}^{\star})$ of the optimal to $\mathfrak{L}_{\mathcal{U}}^{\rho}$, which is a random variable in the product space \mathcal{V}_{CV}^N with probability measure \mathbb{P}^N .

We adapt the following theorem from [CCG21], which gives a lower bound on the containment probability $\Pr_{\mathcal{V}_{CV}}(R_{\rho}^{\star})$ of an optimal solution to $\mathfrak{L}_{\mathcal{U}}^{\rho}$ for a predefined value of ρ . This lower bound is correct with a user-defined confidence level of $\beta \in (0, 1)$, which we typically choose close to one (e.g., $\beta = 0.99$).

Theorem 12.18 (Solution to Problem 12.9) Let \mathcal{U}_N be a set of $N \in \mathbb{N}_{>0}$ samples, and let c^* be the complexity of problem $\mathfrak{L}_{\mathcal{U}}^\rho$. For any confidence level $\beta \in (0, 1)$ and any upper bound $d^* \geq c^*$, it holds that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{CV}^N : \Pr_{\mathcal{V}_{CV}}(R_\rho^*) \geq \eta(d^*) \right\} \geq \beta, \quad (12.2)$$

where R_ρ^* is the prediction region for $\mathfrak{L}_{\mathcal{U}}^\rho$. Moreover, $\eta: \mathbb{N} \rightarrow \mathbb{R}_{\geq 0}$ is a function defined as $\eta(N) = 0$, and otherwise, $\eta(c)$ is the smallest positive real-valued solution to the following polynomial in the t variable for a complexity of c :

$$\binom{N}{c} t^{N-c} - \frac{1-\beta}{2N} \sum_{i=c}^{N-1} \binom{i}{c} t^{i-c} - \frac{1-\beta}{6N} \sum_{i=N+1}^{4N} \binom{i}{c} t^{i-c} = 0. \quad (12.3)$$

Proof. The proof of Theorem 12.18 is based on [CCG21], which states that for a complexity c_ρ^* and for $\eta(c_\rho^*)$ the smallest positive solution to Eq. (12.3), it holds that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{CV}^N : V(R_\rho^*) \leq 1 - \eta(c_\rho^*) \right\} \geq \beta, \quad (12.4)$$

where $V(R_\rho^*)$ is the so-called *violation probability*, which is defined as

$$V(R_\rho^*) = \mathbb{P}\{u \in \mathcal{V}_{CV} : \text{sol}(u) \notin R_\rho^*\}.$$

Observe that $\Pr_{\mathcal{V}_{CV}}(R_\rho^*) + V(R_\rho^*) = 1$. Thus, we rewrite Eq. (12.4) as

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{CV}^N : \Pr_{\mathcal{V}_{CV}}(R_\rho^*) \geq \eta(c_\rho^*) \right\} \geq \beta. \quad (12.5)$$

It is shown by [GC22] that $\eta(c)$ is monotonically decreasing in c [GC22], so for any $d_\rho^* \geq c_\rho^*$, we have $\eta(d_\rho^*) \leq \eta(c_\rho^*)$. Hence, Eq. (12.5) also implies Eq. (12.2), which concludes the proof. ■

With a probability of at least β , Theorem 12.18 yields a correct lower bound on the containment probability. That is, if we solve $\mathfrak{L}_{\mathcal{U}}^\rho$ for many more sets of N parameter samples (recall that, as the samples are i.i.d., these sets are drawn according to the product probability measure \mathbb{P}^N), the inequality in Eq. (12.2) is incorrect for *at most* a $1 - \beta$ fraction of the cases. We plot the lower bound $\eta(c)$ as a function of the complexity $c = 0, \dots, N$ in Fig. 12.6, for different sample sizes N and confidence levels β . These figures show that an increased complexity leads to a lower η , while increasing the sample size leads to a tighter bound.

Example 12.19 We continue Example 12.17. Recall that the complexity for the outer region in Fig. 12.5 is $c_{2,0}^* = 4$. With Theorem 12.18, we compute that, for a confidence level of $\beta = 0.9$, the containment probability for this prediction region is at least $\eta = 0.615$ (cf. Fig. 12.6a). For a stronger confidence level of $\beta = 0.999$, we obtain a more conservative lower bound of $\eta = 0.455$.

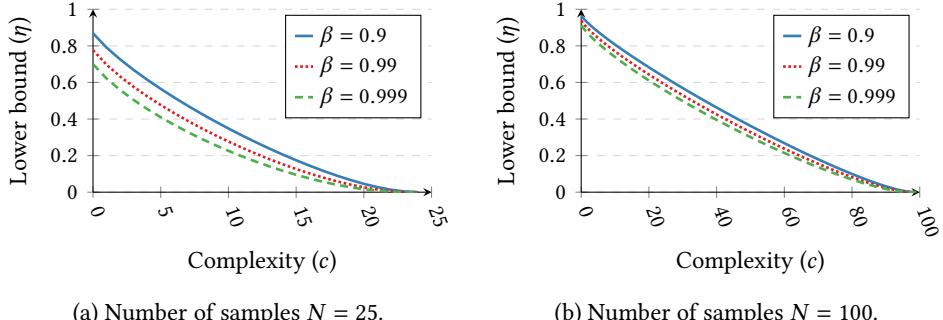


Figure 12.6: Lower bounds η on the containment probability as a function of the complexity c , obtained from Theorem 12.18 for different confidences β .

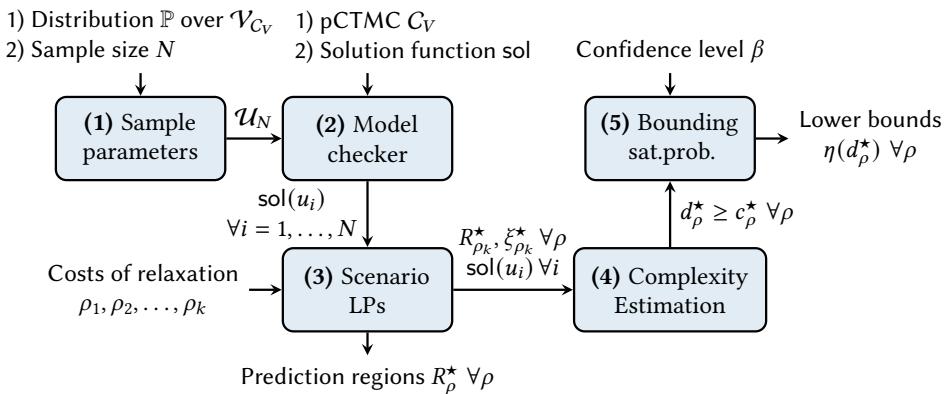


Figure 12.7: Overview of our approach for solving the problem statement.

12.3.3 Algorithm for computing prediction regions

We combine the previous results in the algorithm outlined in Fig. 12.7. Our algorithm produces a set of $k \in \mathbb{N}_{>0}$ prediction regions as in Fig. 12.5 and their associated lower bounds. To strictly solve the problem statement (which asks for a single prediction region), one can set $k = 1$ in the exposition below. We first outline the complete procedure before describing steps 3–5 in more detail.

Recall that Problem 12.9 assumes that we are given a upCTMC (C_V, \mathbb{P}), a solution function $\text{sol}: \mathcal{V}_{CV} \rightarrow \mathbb{R}^m$, and a confidence level $\beta \in (0, 1)$. As preprocessing steps, we first sample a set \mathcal{U}_N of N parameter values (step 1 in Fig. 12.7). In step 2, we feed the given pCTMC C_V and the solution function to a model checking algorithm that computes the solution vector $\text{sol}(u)$ for each $u \in \mathcal{U}_N$, yielding the set of solutions $\{\text{sol}(u_i)\}_{i=1}^N$ (step 2). In step 3, we use the solution vectors $\{\text{sol}(u_i)\}_{i=1}^N$ in the scenario problem $\Omega_{\mathcal{U}}^\rho$ in Eq. (12.1), which we solve for k predefined values ρ_1, \dots, ρ_k , yielding k prediction regions $R_{\rho_1}^*, \dots, R_{\rho_k}^*$. Step 4 is to compute an upper bound d_ρ^* on the complexity c_ρ^* for all values of ρ . Finally, in step 5, we use Theorem 12.18 for a given confidence β to compute the lower bound on the containment probability $\eta(d_\rho^*)$ of each R_ρ^* . Using Def. 12.10, we can postprocess this region to a prediction region over the probability curves.

Step (3): Choosing values for ρ | Example 12.12 shows that relaxation of additional solution vectors (and thus a change in the prediction region) only occurs at *critical* values of $\rho = \frac{1}{N}$, for $N \in \mathbb{N}_{>0}$. In our experiments in Sect. 12.6, we will use $\rho = \frac{1}{N+0.5}$ for ± 10 values of $N \in \mathbb{N}_{>0}$ to obtain gradients of prediction regions.

Step (4): Upper bounding the complexity | Computing the complexity c_ρ^* is a combinatorial problem in general [GC22], because we must consider the removal of all combinations of the solutions on the boundary of the prediction region R_ρ^* . In practice, we compute an upper bound $d_\rho^* \geq c_\rho^*$ on the complexity via a greedy algorithm. Specifically, we iteratively solve \mathfrak{L}_U^ρ in Eq. (12.1) with *one more sample on the boundary removed*. If the optimal solution is unchanged, we conclude that this sample does not contribute to the complexity. If the optimal solution is changed, we put the sample back and proceed by removing a different sample. This greedy algorithm terminates when we have tried removing all solutions on the boundary.

Step (5): Lower bounding the containment probability | Theorem 12.18 characterizes a computable function $B(d^*, N, \beta)$ that returns zero for $d^* = N$ (i.e., all samples are critical), and otherwise uses the polynomial Eq. (12.3) to obtain η , which we solve with an approximate root finding method in practice (see [GC22] for details on how to ensure that we find the smallest root). For every upper bound on the complexity d^* and any requested confidence, we obtain the lower bound $\eta = B(d^*, N, \beta)$ for the containment probability with respect to the prediction region R_ρ^* .

12.4 Imprecise Sampling-Based Prediction Regions

Thus far, we have solved Problem 12.9 under the assumption that we are able to compute the solution vectors precisely (up to numerics). For some models, however, computing precise solutions is expensive or even practically infeasible. In such a case, we may choose to compute an approximation, given as an *interval* on each entry of the solution function. In this section, we deal with such *imprecise solutions*.

Setting | An imprecise solution is described by the lower bound $\text{sol}_-(u) \in \mathbb{R}^m$ and upper bound $\text{sol}_+(u) \in \mathbb{R}^m$, such that $\text{sol}_-(u) \leq \text{sol}(u) \leq \text{sol}_+(u)$ holds with pointwise inequalities. Our goal is to compute a prediction region R and a (high-confidence) lower bound μ such that $\Pr_{\mathcal{V}_{CV}}(R) \geq \mu$, i.e., a lower bound on the probability that any *precise solution* $\text{sol}(u)$ is contained in R . However, we must now compute R and $\Pr_{\mathcal{V}_{CV}}(R)$ from the imprecise solutions sol_- and sol_+ . Thus, we aim to provide a guarantee with respect to the *precise* solution $\text{sol}(u)$, based on *imprecise* solutions.

Challenge | Intuitively, if we increase the (unknown) prediction region R^* from problem \mathfrak{L}_U^ρ (for the unknown precise solutions) while also overapproximating the complexity of \mathfrak{L}_U^ρ , we obtain sound bounds. We formalize this idea as follows.

Lemma 12.20 (Imprecise solutions) Let R_ρ^* be the prediction region and c_ρ^* be the complexity that result from solving \mathfrak{L}_U^ρ for the precise (unknown) solutions $\{\text{sol}(u_i)\}_{i=1}^N$. Given a set $R \in \mathbb{R}^N$ and $d \in \mathbb{N}$, for any confidence level $\beta \in (0, 1)$, the following implication holds:

$$R_\rho^\star \subseteq R \text{ and } c_\rho^\star \leq d \implies \mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{CV}^N : \Pr_{\mathcal{V}_{CV}}(R) \geq \eta(d) \right\} \geq \beta,$$

where $\eta(n) = 0$, and otherwise, $\eta(d)$ is the smallest positive real-valued solution to the polynomial equality in Eq. (12.3).

Proof. Recall from the proof of Theorem 12.18 that for $\eta(c_\rho^\star)$ the solution to Eq. (12.3), where c_ρ^\star is the true complexity of problem \mathfrak{L}_U^ρ , it holds that

$$\mathbb{P}^N \left\{ \{u_1, \dots, u_N\} \in \mathcal{V}_{CV}^N : \Pr_{\mathcal{V}_{CV}}(R_\rho^\star) \geq \eta(c_\rho^\star) \right\} \geq \beta. \quad (12.6)$$

Observe that for any two sets $R_\rho^\star \subseteq R$, we have $\Pr_{\mathcal{V}_{CV}}(R) \geq \Pr_{\mathcal{V}_{CV}}(R_\rho^\star)$. Moreover, recall that $\eta(c)$ is monotonically decreasing in c (as also observed visually from Fig. 12.6), and thus, the condition $c_\rho^\star \leq d$ implies that $\eta(d) \leq \eta(c_\rho^\star)$. Hence, under the proposed conditions, we rewrite Eq. (12.6) as the right-hand side of the implication in Lemma 12.20, which concludes the proof. ■

What is left is to compute the appropriate R and d in Lemma 12.20, which we will deal with in Sects. 12.4.1 and 12.4.3, respectively. As we will see, in contrast to Sect. 12.3, these results do *not* carry over to non-rectangular prediction regions. Thus, we will exclusively consider rectangular regions in the remainder of this section.

12.4.1 Prediction regions on imprecise solutions

In this section, we show how to compute $R \supseteq R_\rho^\star$, satisfying the first term in the premise of Lemma 12.20. We construct a *conservative box* around the imprecise solutions as in Fig. 12.8, containing both $\text{sol}_-(u)$ and $\text{sol}_+(u)$. We compute this box by solving the following problem \mathfrak{G}_U^ρ as a modified version of \mathfrak{L}_U^ρ in Eq. (12.1):

$$\mathfrak{G}_U^\rho : \underset{\substack{\underline{x} \in \mathbb{R}^m, \bar{x} \in \mathbb{R}^m, \\ \xi_i \in \mathbb{R}_{\geq 0}^m \forall i=1, \dots, N}}{\text{minimize}} \quad \|\bar{x} - \underline{x}\|_1 + \rho \sum_{i=1}^n \|\xi_i\|_1 \quad (12.7a)$$

$$\text{subject to} \quad \underline{x} - \xi_i \leq \text{sol}_-(u_i) \quad \forall i = 1, \dots, N \quad (12.7b)$$

$$\text{sol}_+(u_i) \leq \bar{x} + \xi_i \quad \forall i = 1, \dots, N. \quad (12.7c)$$

We denote the (arguments of the) optimal solution of \mathfrak{G}_U^ρ by $[\underline{x}'_\rho, \bar{x}'_\rho], \xi'_\rho$. For comparison, recall that the optimal solution to \mathfrak{L}_U^ρ is written as $[\underline{x}^\star_\rho, \bar{x}^\star_\rho], \xi^\star_\rho$.³ If a sample $u_i \in \mathcal{V}_{CV}$ in problem \mathfrak{G}_U^ρ is relaxed (i.e., has a non-zero ξ_i), part of the set $[\text{sol}_-(u_i), \text{sol}_+(u_i)]$ is not contained in the prediction region. The following result relates \mathfrak{L}_U^ρ and \mathfrak{G}_U^ρ , showing that we can use $[\underline{x}'_\rho, \bar{x}'_\rho]$ as R in Lemma 12.20.

Theorem 12.21 (Solving Problem 12.9 with imprecise solutions) Let the cost of relaxation $\rho > 0$ and the sample set $\mathcal{U}_N = \{u_1, \dots, u_N\}$ be given. For the prediction region $[\underline{x}'_\rho, \bar{x}'_\rho]$ as the optimal solution to problem \mathfrak{G}_U^ρ , it holds that $[\underline{x}^\star_\rho, \bar{x}^\star_\rho] \subseteq [\underline{x}'_\rho, \bar{x}'_\rho]$, with $[\underline{x}^\star_\rho, \bar{x}^\star_\rho]$ the optimal solution to \mathfrak{L}_U^ρ .

³We write $[\underline{x}^\star_\rho, \bar{x}^\star_\rho]$ and $[\underline{x}'_\rho, \bar{x}'_\rho]$, because the results in Sect. 12.4 apply only to rectangular regions.

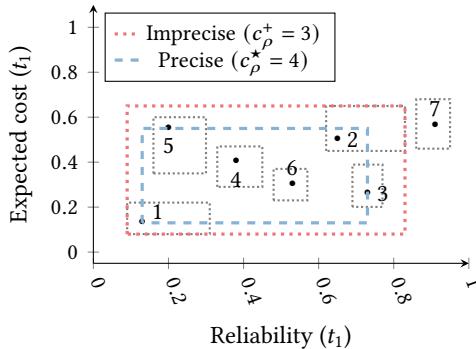


Figure 12.8: Complexity of imprecise solution versus that of the precise solution.

Since the proof of Theorem 12.21 is nontrivial, we first explain the result intuitively and thereafter provide the formal proof in Sect. 12.4.2. In particular, the entries ξ_i from the optimization problems \mathfrak{G}_U^ρ and \mathfrak{L}_U^ρ are incomparable, as are their objective functions. Instead, Theorem 12.21 relies on two observations:

1. Due to the use of the 1-norm, the optimization problem \mathfrak{G}_U^ρ can be decomposed into N smaller problems, whose results combine into a solution to the original problem. This allows us to consider individual dimensions of the solution vectors.
2. The solution vectors that are relaxed depend on the value of ρ and on their *relative order*, but not on the *precise position* within that order, which is also illustrated by Example 12.12. In combination with the observation from Example 12.12 that the *outermost* samples are relaxed at the (relatively) highest ρ , we can provide conservative guarantees on which samples are (or are surely not) relaxed.

12.4.2 *Proof of Theorem 12.21

The one-dimensional case | Let us first consider the case for solution vectors in \mathbb{R}^1 , i.e., those modeling one measure. Recall from Example 12.12 that for precise solutions in one dimension, the *outermost* two samples (labeled A and F) are relaxed under the (unique) optimal solution if $\rho < 1$, samples B and E if $\rho < \frac{1}{2}$, etc. Denote by $\text{sol}_-(u)_r \in \mathbb{R}$ and $\text{sol}_+(u)_r \in \mathbb{R}$ the r^{th} entries of the respective imprecise solution vectors for sample $u \in \mathcal{V}_{CV}$. In formalizing the relationship between the value of ρ and whether a sample is relaxed, we state the following definition:

Definition 12.22 (Counting samples) Let $r \in \{1, \dots, m\}$, where m is the dimension of the solution vector $\text{sol}(u) \in \mathbb{R}^m$. For any $(\text{sol}_-(u_i), \text{sol}_+(u_i))$, we define $\text{num}_+^\geq(u_i)_r \in \{1, \dots, N\}$ and $\text{num}_-^\leq(u_i)_r \in \{1, \dots, N\}$ as the number of samples whose upper bound is at least $\text{sol}_+(u_i)$ (or at most $\text{sol}_-(u_i)$), when projected to dimension r :

$$\begin{aligned}\text{num}_+^\geq(u_i)_r &= |\{u_j \in \mathcal{V}_{CV} : \text{sol}_+(u_j)_r \geq \text{sol}_+(u_i)_r\}| \\ \text{num}_-^\leq(u_i)_r &= |\{u_j \in \mathcal{V}_{CV} : \text{sol}_-(u_j)_r \leq \text{sol}_-(u_i)_r\}|.\end{aligned}$$

The intuition of Def. 12.22 is illustrated by Fig. 12.9, which shows the values of $\text{num}_+^\leq(u_i)$ for the case where $r = 1$ (thus, the subscripts r are omitted).

*Section with details that can be skipped safely

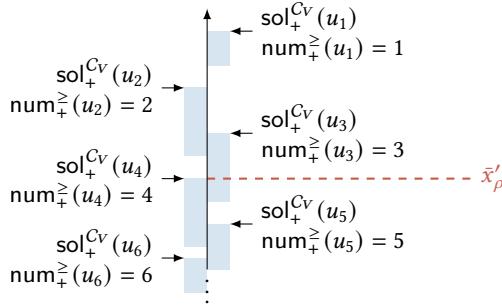


Figure 12.9: The upper bounds of five imprecise solutions and their values of $\text{num}_+^{\geq}(u)$.

The following lemma then characterizes the relaxed solutions:

Lemma 12.23 (Containment of imprecise solution) An imprecise solution $(\text{sol}_-(u), \text{sol}_+(u))$, $u \in \mathcal{V}_{CV}$ is not contained in prediction region $[\underline{x}'_\rho, \bar{x}'_\rho]$ of $\mathfrak{G}_{\mathcal{U}'}^\rho$ projected to dimension $r \in \{1, \dots, m\}$ of the solution function, if

$$\rho < \min \left\{ \text{num}_+^{\geq}(u)_r, \text{num}_-^{\leq}(u)_r \right\}^{-1}. \quad (12.8)$$

Note that, for precise solutions, there is no distinction between $\text{sol}_-(u_i)$ and $\text{sol}_+(u_i)$. Nevertheless, Lemma 12.23 still applies.

Proof (Proof of Theorem 12.21; part I). We prove the theorem by contradiction for the 1-dimensional case, and we generalize afterward. In a single dimension, $[\underline{x}'_\rho, \bar{x}'_\rho] \not\subseteq [\underline{x}'_\rho, \bar{x}'_\rho]$ requires that either $\underline{x}'_\rho < \underline{x}'_\rho$ or $\bar{x}'_\rho > \bar{x}'_\rho$. First consider the upper bounds \bar{x}'_ρ and \bar{x}'_ρ , which we can make explicit using Lemma 12.23:

$$\begin{aligned} \bar{x}'_\rho &= \max \left\{ \text{sol}_+(u), u \in \mathcal{V}_{CV} : \text{num}_+^{\geq}(u) > \rho^{-1} \right\} \\ \bar{x}'_\rho &= \max \left\{ \text{sol}(u), u \in \mathcal{V}_{CV} : \text{num}_\geq(u) > \rho^{-1} \right\}. \end{aligned}$$

For $\bar{x}'_\rho > \bar{x}'_\rho$ to hold, the maximum $\text{num}_\geq(u) > \rho^{-1}$ (i.e., the highest *precise* solution for which there are more than ρ^{-1} solutions at least as high) must exceed $\text{num}_+^{\geq}(u) > \rho^{-1}$ (the highest *imprecise upper bound* solution for which there are more than ρ^{-1} imprecise upper bound solutions at least as high). This can only be true if the number of samples for which $\text{sol}(u) > \bar{x}'$ is higher than the number for which $\text{sol}_+(u) > \bar{x}'$. However, by construction, $\text{sol}(u) \leq \text{sol}_+$, so this is impossible, and thus, it holds that $\bar{x}'_\rho \leq \bar{x}'_\rho$. While omitted for brevity, the proof that the lower bound $\underline{x}'_\rho \geq \bar{x}'_\rho$ follows analogous to the upper bound. ■

The multi-dimensional case | We now generalize the results above to the multi-dimensional case, i.e., with multiple measures.

Lemma 12.24 Problem $\mathfrak{G}_{\mathcal{U}}^{\rho}$ can be decomposed into an independent problem for every of the $m \in \mathbb{N}_{>0}$ dimensions of the solution function $\text{sol}: \mathcal{V}_{CV} \rightarrow \mathbb{R}^m$.

Lemma 12.24 holds because the objective Eq. (12.1a) is additive and all constraints for all measures Eq. (12.1b) are independent. Thus, we can equivalently solve problem $\mathfrak{L}_{\mathcal{U}}^{\rho}$ for all m measures modeled by the solution function separately.

Proof (Proof of Theorem 12.21; part II). We now generalize the result to multiple dimensions. Lemma 12.24 states that for rectangular prediction regions, problems $\mathfrak{L}_{\mathcal{U}}^{\rho}$ and $\mathfrak{G}_{\mathcal{U}}^{\rho}$ can be solved for each dimension separately. As such, we obtain an element-wise inequality $\underline{x}'_{\rho} \leq \underline{x}_{\rho}^* \leq \bar{x}_{\rho}^* \leq \bar{x}'_{\rho}$, which also implies that $[\underline{x}_{\rho}^*, \bar{x}_{\rho}^*] \subseteq [\underline{x}'_{\rho}, \bar{x}'_{\rho}]$, so the claim follows. ■

12.4.3 Computing the complexity

To satisfy the second term of the premise in Lemma 12.20, we compute an upper bound on the complexity. We first present a negative result. Let the complexity c'_{ρ} of problem $\mathfrak{G}_{\mathcal{U}}^{\rho}$ be defined analogous to Def. 12.16, but with $[\underline{x}'_{\rho}, \bar{x}'_{\rho}]$ as the region.

Lemma 12.25 In general, $c_{\rho}^* \leq c'_{\rho}$ does not hold.

Proof. In Fig. 12.8, the smallest critical set for the imprecise solutions are those labeled $\{1, 2, 7\}$, while this set is $\{1, 3, 5, 7\}$ under precise solutions, so $c_{\rho}^* > c'_{\rho}$. ■

Thus, we cannot upper bound the complexity directly from the result to $\mathfrak{G}_{\mathcal{U}}^{\rho}$. We can, however, determine the samples that are certainly *not* in any critical set (recall Def. 12.16). Intuitively, a sample is *surely noncritical* if its (imprecise) solution is strictly within the prediction region and does not overlap with any solution on the region's boundary. In Fig. 12.9, sample u_6 is surely noncritical, but sample u_5 is not (whether u_5 is critical depends on its precise solution). Formally, let δR be the boundary (i.e., the set of all limit points of R , as we defined in the preliminaries in Sect. 2.1) of region $[\underline{x}'_{\rho}, \bar{x}'_{\rho}]$, and let \mathcal{B} be the set of samples whose solutions overlap with δR , which is

$$\mathcal{B} = \{u \in \mathcal{U}_N : [\text{sol}_-(u), \text{sol}_+(u)] \cap \delta R \neq \emptyset\}.$$

We then define a surely noncritical sample as follows.

Definition 12.26 (Surely noncritical samples) For a region $[\underline{x}'_{\rho}, \bar{x}'_{\rho}]$, let $\mathcal{I} \subset [\underline{x}'_{\rho}, \bar{x}'_{\rho}]$ be the rectangle of largest volume, such that $\mathcal{I} \cap [\text{sol}_-(u), \text{sol}_+(u)] = \emptyset$ for any $u \in \mathcal{B}$. A sample $u_i \in \mathcal{V}_{CV}$ is *surely noncritical* if $[\text{sol}_-(u_i), \text{sol}_+(u_i)] \subseteq \mathcal{I}$. The set of all surely noncritical samples with respect to the (unknown) prediction region $[\underline{x}_{\rho}^*, \bar{x}_{\rho}^*]$ is denoted by $\mathcal{X} \subset \mathcal{U}_N$.

Before we state the last technical result of this section, we provide the following useful lemma about surely noncritical samples.

Lemma 12.27 Any surely noncritical sample cannot be in the (smallest) critical set, defined in Def. 12.16.

Proof. Recall from Def. 12.16 that a sample may (potentially) be critical if it is either outside or on the boundary of the prediction region $[x_\rho^*, \bar{x}_\rho^*]$. While the boundary of the prediction region $[x_\rho^*, \bar{x}_\rho^*]$ is unknown, it cannot be smaller than the inner rectangle \mathcal{I} defined in Def. 12.26. By construction, any surely noncritical sample is a subset of this set \mathcal{I} . Hence, any surely noncritical sample cannot be in the (smallest) critical set, and the claim follows. ■

As a worst case, any sample not surely noncritical can be in the smallest critical set, leading to the following bound on the complexity as required by Lemma 12.20.

Theorem 12.28 Let \mathcal{X} be the set of surely noncritical samples. Then, it holds that $c_\rho^* \leq |\mathcal{U}_N \setminus \mathcal{X}|$.

Proof. The proof follows almost directly from Lemma 12.27. The complexity is the cardinality of the smallest critical set, which cannot contain any surely noncritical sample, as stated by Lemma 12.27. Hence, it follows that $N - |\mathcal{X}| = |\mathcal{U}_N \setminus \mathcal{X}|$, where \mathcal{X} is the set of surely noncritical samples, which concludes the proof. ■

For imprecise solutions, the bound in Theorem 12.28 is conservative but can potentially be improved, as discussed in the following.

12.4.4 Solution refinement scheme

Practical model checking algorithms for CTMCs can compute upper and lower bounds on solutions with an arbitrary precision. However, the higher the requested precision, the higher the computation time. Thus, there is a trade-off between the tightness of the bounds on the solution and the computation time. This trade-off naturally leads to a *refinement* scheme for the imprecise solutions used to compute prediction regions.

Refining solutions | If the prediction region for a given set of imprecise solution vectors leads to an unsatisfactory solution to Problem 12.9, we refine these imprecise solutions by increasing the requested precision in the model checking algorithm. By doing so, we obtain imprecise solutions that are closer to their true (but unknown) value, which leads to improved prediction regions and upper bound on the complexity, which in turn improves the computed bound on the containment probability.

Refinement scheme | Specifically, we propose the following rule for refining solutions. After solving $\mathfrak{G}_{\mathcal{U}}^\rho$ for a given set of imprecise solutions, we refine the solutions on the boundary of the obtained prediction region. We then resolve problem $\mathfrak{G}_{\mathcal{U}}^\rho$, thus adding a loop back from step (4) to (2) in Fig. 12.7. In our experiments, we show that with this refinement scheme, we iteratively improve our upper bound $d \geq c_\rho^*$ and the smallest overapproximation $R \supseteq R_\rho^*$ of the prediction region.

12.5 Batch Verification for CTMCs

One practical bottleneck in our method is to obtain the necessary number of solution vectors $\{\text{sol}(u_i)\}_{i=1}^N$ by model checking. The following improvements, while mild, are essential in our implementation and, therefore, deserve a brief discussion.

In general, computing $\text{sol}(u)$ via model checking consists of two parts. First, the

high-level representation of the upCTMC—given in Prism [KNP11], in JANI [BDHH⁺17], or as a dynamic *fault tree*⁴—is translated into a concrete CTMC $C_V[u]$. Then, from $C_V[u]$ we construct $\text{sol}(u)$ using off-the-shelf algorithms [BHHK03]. We adapt the pipeline by tailoring the translation and the approximate analysis as outlined below.

Our implementation supports two methods for building the concrete CTMC for a parameter sample: (1) by first instantiating the valuation in the specification and then building the resulting concrete CTMC, or (2) by first building the pCTMC C_V (only once) and then instantiating it for each parameter sample to obtain the concrete CTMC $C_V[u]$. Which method is faster depends on the specific model (we only report results for the fastest method in Sect. 12.6 for brevity).

Partial models | To accelerate the time-consuming computation of solution vectors for large models, it is natural to abstract the models into smaller models that can be analyzed faster. Similar to ideas used for dynamic fault trees [VJK18], infinite CTMCs [RNBM⁺22], and continuous-time Markov decision processes (CTMDPs) [ABHK18], we employ an abstraction which only keeps the most relevant parts of a model, i.e., states with a sufficiently large probability to be reached from the initial state(s). Analysis of this partial model then yields best- and worst-case results for each measure, by assuming that all removed states are either target states (best case) or are not (worst case). This method returns imprecise solution vectors as used in Sect. 12.4, which can be refined up to arbitrary precision by retaining more states of the original model.

Similar to building the complete models, two approaches are possible to create the partial models: (1) fixing the valuation and directly abstracting the concrete CTMC, or (2) first building the complete pCTMC and then abstracting the concrete CTMC. We reuse partial models for similar valuations to avoid costly computations. We cluster parameter valuations that are close to each other (in Euclidean distance). For parameter valuations within one cluster, we reuse the same partial model (in terms of the states), albeit instantiating it according to the precise valuation.

12.6 Numerical Experiments

We answer three questions about (a prototype implementation of) our approach:

- Q1. Can we verify CTMCs taking into account the uncertainty about the rates?
- Q2. How well does our approach scale with respect to the number of measures (modeled by the solution function) and samples?
- Q3. How does our approach compare to naïve baselines (to be defined below)?

Setup | We implement our approach in Python, using the explicit engine of Storm [HJKQ⁺22] and the improvements of Sect. 12.5 to sample from upCTMCs. Our current implementation is limited to pCTMC instantiations that are *graph-preserving*, i.e., for any pair $s, s' \in S$ either $R(s, s')[u] = 0$ or $R(s, s')[u] > 0$ for all u . We solve optimization problems using the ECOS solver [DCB13]. All experiments ran single-threaded on a computer with 32 3.7 GHz cores and 64 GB RAM. We show the effectiveness of our method on a large number of publicly available pCTMC [HKPQ⁺19] and fault tree benchmarks [RBNS⁺19] across domains.

⁴Recall from Chapter 11 that a fault tree can be represented as a CTMC.

Reproducibility | All source code, benchmarks, and logfiles used to produce the experimental results are archived at <https://doi.org/10.5281/zenodo.6523863>.

12.6.1 Converting pCTMCs into upCTMCs

We illustrate on two specific benchmarks how we convert a pCTMC into a upCTMC by equipping its parameters with a probability distribution. We omit details on the other benchmarks for brevity.

Epidemic modeling | We consider the classical *SIR infection model* [AB12] of an infectious disease spreading through a population. The factored state $s = (\bar{S}, \bar{I}, \bar{R})$ of this CTMC counts the susceptible, infected, and recovered populations, denoted by $\bar{S}, \bar{I}, \bar{R} \geq 0$. Infections and recoveries (we assume that immunization is permanent) occur based on the following parametric rules that depend on parameters λ_i and λ_r :

$$\text{Infection: } S + I \xrightarrow{\lambda_i \cdot \bar{S} \cdot \bar{I}} I + I, \quad \text{Recovery: } I \xrightarrow{\lambda_r \cdot \bar{I}} R.$$

In the classical SIR model, λ_i and λ_r are assumed to be known precisely, while we consider the parameters $\lambda_i = \mathcal{N}(0.05, 0.002)$ and $\lambda_r = \mathcal{N}(0.04, 0.002)$ to be normally distributed (recall that we only use samples from these distributions). We define a solution function $\text{sol}: \mathcal{V}_{CV} \rightarrow \mathbb{R}^m$, which is defined for all $u \in \mathcal{V}_{CV}$ as

$$\text{sol}(u) = \left[\Pr^{CV[u]}(s_I \models \varphi_{t_1}), \dots, \Pr^{CV[u]}(s_I \models \varphi_{t_m}) \right]^\top,$$

where φ_t is a CSL path formula defined as

$$\varphi_t = (\bar{I} > 0) \mathbf{U}_{[100, t]}(\bar{I} = 0).$$

Thus, the solution function models the probability of the disease becoming extinct ($\bar{I} = 0$) within the time interval $[100, t]$, for different values of t . In this experiment, we define these time points to be evenly spaced over the time interval $[100, 200]$:

$$t_i = 100 + \frac{100 \cdot i}{m} \quad \text{for all } i = 1, \dots, m.$$

Buffer system | We augment the *producer-consumer buffering* system considered by [CCGK⁺18]. We equip the six parameters of this pCTMC by uniform probability distributions over their respective domains. This pCTMC models the transfer of requests from a producer (at a rate of $\lambda_g \in [32, 38]$) to consumers who consume them (at a rate of $\lambda_c \in [27, 33]$). The requests are sent at a rate of $\lambda_t \in [27, 33]$, via either a slow or a fast buffer, with probabilities of 0.6 and 0.4, respectively. While being faster, the fast buffer is less reliable than the slow buffer (it loses requests with a probability $\lambda_{\text{loss}} \in [0.025, 0.075]$), and has a smaller capacity. Requests from the slow buffer are transferred to the fast buffer with a probability proportional to the occupancy. The transmission rate of the slow buffer is $\lambda_{\text{slow}} \in [5, 15]$; the rate of the fast buffer is $\lambda_{\delta} \in [5, 15]$ higher. We consider a solution function $\text{sol}: \mathcal{V}_{CV} \rightarrow \mathbb{R}^2$ modeling two measures: (1) the expected transferred requests until time 25, and (2) the probability that the utilization of both buffers is above 75% within the time $[20, 25]$.

Table 12.1: Excerpt of the benchmark statistics (sampling time is per 100 CTMCs; m is the dimension of the solution function).

benchmark	m	Model size			Storm run time [s]		Scen.opt. time [s]	
		#pars	#states	#trans	Init.	Sample ($\times 100$)	$N = 100$	$N = 200$
SIR (140)	26	2	9 996	19 716	0.29	2947.29	18.26	63.27
SIR (140) ^a	26	2	9 996	19 716	0.29	544.27	25.11	129.66
KANBAN (3)	4	13	58 400	446 400	4.42	46.95	2.28	6.69
KANBAN (5)	4	13	2 546 432	24 460 016	253.39	4363.63	2.03	5.94
POLLING (9)	2	2	6 912	36 864	0.64	22.92	2.13	6.66
BUFFER	2	6	5 632	21 968	0.48	20.70	1.21	4.15
TANDEM (31)	2	5	2 016	6 819	0.11	862.41	5.19	24.30
RBC	40	6	2 269	12 930	0.01	1.40	5.27	16.88
Rc (1,1)	25	21	8 401	49 446	27.20	74.90	5.75	20.34
Rc (1,1) ^a	25	21	n/a ^b	n/a ^b	0.02	2.35	29.23	150.61
Rc (2,2) ^a	25	29	n/a ^b	n/a ^b	0.03	27.77	24.86	132.63
HECS (2,1) ^a	25	5	n/a ^b	n/a ^b	0.02	9.83	26.78	145.77
HECS (2,2) ^a	25	24	n/a ^b	n/a ^b	0.02	194.25	33.06	184.32

^a Computed using approximate model checking up to a relative gap between upper bound $\text{sol}_+^{\mathcal{C}_V}(u)$ and lower bound $\text{sol}_-^{\mathcal{C}_V}(u)$ below 1% for every sample $u \in \mathcal{V}_{\mathcal{C}_V}$.

^b Model size is unknown, as the approximation does not build the full state-space.

12.6.2 Applicability

An excerpt of the benchmark statistics is shown in Table 12.1 (for the full table, we refer to [3, Appendix C]). For all but the smallest benchmarks, sampling and computing the solution vectors by model checking is more expensive than solving the scenario problems. In the following, we illustrate that 100 samples are sufficient to provide qualitatively good prediction regions and associated lower bounds.

Plotting prediction regions | Fig. 12.10 shows prediction regions on the extinction probability of the disease in the SIR model and is analogous to the tubes in Fig. 12.1d.⁵ These regions are obtained by applying our algorithm with varying values for the cost of relaxation ρ . For a confidence level of $\beta = 99\%$, the widest (smallest) tube in Fig. 12.10 corresponds to a lower bound probability of $\mu = 91.1\%$ ($\mu = 23.9\%$). Thus, we conclude that, with a confidence of at least 99%, the curve created by the CTMC for a parameter value sampled according to \mathbb{P} will lie within the outermost region in Fig. 12.10 with a probability of at least 91.1%.

Similarly, Fig. 12.11 shows $N = 200$ solution vectors for the BUFFER benchmark, with a prediction region as a Pareto front on two measures, thus highlighting that our approach also supports regions of different shapes. For a confidence level of $\beta = 99\%$, the outer and inner prediction regions are associated with a lower bound probability of $\mu = 91.1\%$ and $\mu = 66.2\%$, respectively.

Tightness of the solution | In Table 12.2, we investigate the tightness of our results. For the experiment, we set $\rho = 1.1$ and solve $\Omega_{\mathcal{U}}^\rho$ for different values of N , repeating every experiment 10 times, resulting in the average bounds $\bar{\mu}$. Then, we sample 1 000

⁵We present the analogous plots for various other benchmarks in [3, Appendix C].

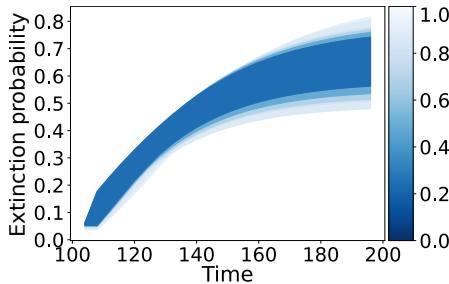


Figure 12.10: The prediction regions for the SIR (60) benchmark with $N = 400$ samples.

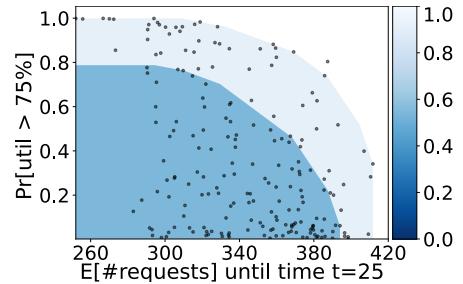


Figure 12.11: The Pareto front for the BUFFER benchmark with $N = 200$ samples.

Table 12.2: Lower bounds $\bar{\mu}$ and standard deviation (SD), vs. the observed number of 1 000 additional solutions that indeed lie within the obtained regions.

(a) KANBAN (3).					
	$\beta = 0.9$	$\beta = 0.999$	Frequentist		
N	$\bar{\mu}$	SD	$\bar{\mu}$	SD	Observed
100	0.862	0.000	0.798	0.000	959 ± 22.7
200	0.930	0.000	0.895	0.000	967 ± 17.4
400	0.965	0.001	0.947	0.001	984 ± 8.6
800	0.982	0.000	0.973	0.000	994 ± 3.2

(b) Railway crossing, Rc (1,1,hc).					
	$\beta = 0.9$	$\beta = 0.999$	Frequentist		
N	$\bar{\mu}$	SD	$\bar{\mu}$	SD	Observed
100	0.895	0.018	0.835	0.020	954 ± 26.8
200	0.945	0.007	0.912	0.008	980 ± 12.8
400	0.975	0.004	0.958	0.005	990 ± 8.3
800	0.986	0.002	0.977	0.003	995 ± 4.3

solutions and count the *observed* number of solutions contained in every prediction region, resulting in an empirical approximation of the containment probability. Recall that for $\rho > 1$, we obtain a prediction region that contains all solutions, so this observed count grows toward N . The lower bounds grow toward the empirical count for an increased N , with the smallest difference (Rc, $N = 800$, $\beta = 0.9$) being as small as 0.9%. Similar observations hold for other values of ρ .

Handling imprecise solutions | The approximate model checker is significantly faster (see Table 12.1 for SIR (140) and Rc), at the cost of obtaining imprecise solution vectors.⁶ For SIR (140), the sampling time is reduced from 49 to 9 min, while the scenario optimization time is slightly higher at 129 s. This difference only grows larger with the size of the CTMC. For the larger instances of Rc and HECS, computing exact solutions is infeasible at all (one HECS (2,2) sample alone takes 15 min).

While the bounds on the containment probability under imprecise solutions may initially be poor (for example, the setting in Fig. 12.12a results in $\mu = 2.1\%$), we can improve the results significantly using the refinement scheme proposed in Sect. 12.4.4. For example, Fig. 12.12c shows the prediction region after refining 31 of the 100 solutions, which yields $\mu = 74.7\%$. Thus, by iteratively refining only the imprecise solutions on the boundary of the prediction regions, we significantly tighten the obtained bounds on the containment probability.

⁶We terminate at a relative gap between upper/lower bound of the solution below 1%.

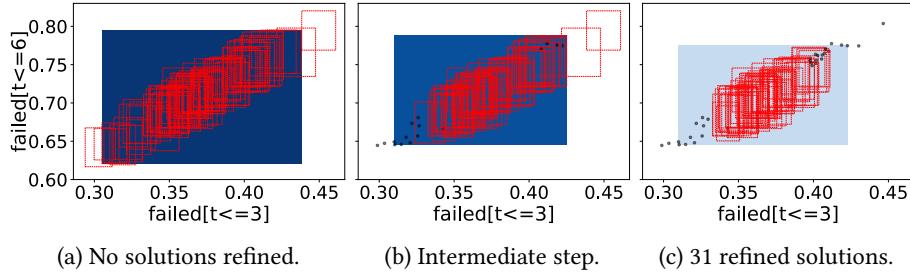


Figure 12.12: Refining imprecise solution vectors (red boxes) for Rc (2,2), $N = 100$.

Table 12.3: Run times in [s] for solving the scenario problems for SIR and Rc with $\rho = 0.1$ (timeout (TO) of 1 hour) for different sample sizes n and numbers of measures m modeled by the solution function.

(a) SIR (population 20).						(b) Railway crossing, Rc (1,1,hc).				
N / m	50	100	200	400	800	N / m	50	100	200	400
100	0.97	1.59	3.36	9.17	25.41	100	1.84	3.40	8.18	24.14
200	3.69	7.30	22.91	59.45	131.78	200	6.35	14.56	45.09	113.09
400	29.43	76.13	153.03	310.67	640.70	400	34.74	96.68	203.77	427.80
800	261.97	491.73	955.77	1924.15	TO	800	292.32	579.09	1215.67	2553.98

12.6.3 Scalability

In Table 12.3, we report the run times for steps (3)–(5) of our algorithm in Fig. 12.7 (i.e., for solving the scenario problems, but not for computing the solution vectors in Storm). We solve problem \mathfrak{L}_U^ρ for $\rho = 0.1$, with different numbers of samples N and numbers of measures m modeled by the solution function. Our approach scales well to realistic numbers of samples (up to 800) and measures (up to 400). The computational complexity of the scenario problems is largely *independent of the size of the CTMC*, and hence, similar run times are observed across the benchmarks (as also shown by Table 12.1).

12.6.4 Comparison to baselines

We compare against two baselines: (1) Scenario optimization to analyze each measure independently, yielding a separate probabilistic guarantee on each measure. (2) A frequentist (Monte Carlo) baseline, which samples a large number of parameter values and counts the number of associated solutions within a region.

Analyzing measures independently | To show that analyzing a full *set of measures* at once, e.g., the complete probability curve, is essential, we compare our method to the baseline that analyzes *each measure independently* and combines the obtained bounds on each measure afterward. We consider the PCS (BUFFER) benchmark with precise samples and solve \mathfrak{L}_U^ρ for $\rho = 2$.

Table 12.4 presents the full comparison on the PCS benchmark between our approach and the baseline scenario approach that analyzes each measure independently. In this table, we report the average lower bounds (over 10 iterations) on the containment probability for different sample sizes $N = 100, \dots, 800$, and confidence levels β . For

Table 12.4: Obtained bounds (for the PCS fault tree) on the containment probability for our approach and the baseline that analyzes each measure independently.

Method	$N = 100$		$N = 200$		$N = 400$		$N = 800$	
	$\beta = 0.9$	$\beta = 0.999$						
Our approach	0.908	0.848	0.937	0.903	0.976	0.960	0.984	0.975
Baseline	0.045	0.010	0.212	0.103	0.461	0.322	0.679	0.567

$N = 100$ samples and $\beta = 99.9\%$, our approach returns a lower bound probability of $\mu = 84.8\%$. By contrast, for this value of $\beta = 99.9\%$, the naïve baseline yields a lower bound of only 1.0%. While the difference between both methods decreases with an increasing number of samples, our approach consistently outperforms the baseline for different values of β and N . There are two reasons for this large difference. First, the baseline applies Theorem 12.28 once for each of the 25 measures, so it must use a more conservative confidence level of $\tilde{\beta} = 1 - \frac{1-\beta}{25} = 0.9996$. Second, the baseline takes the conjunction over the 25 independent lower bounds, which drastically reduces the obtained bound.

Frequentist baseline | The comparison to the frequentist baseline on the KANBAN and Rc benchmarks yields the previously discussed results in Table 12.2. The results in Tables 12.1 and 12.3 show that *the time spent for sampling is (for most benchmarks) significantly higher than for scenario optimization*. Thus, our scenario-based approach has a relatively low cost while resulting in valuable guarantees which the baseline does not give. To still obtain a high confidence in the result, a much larger sample size is needed for the frequentist baseline than for our approach.

12.7 Related Work

Several verification approaches exist to handle uncertain Markov models. We briefly discuss the approaches most related to our contributions from this chapter.

For (discrete-time) *interval* Markov chains or Markov decision processes, several approaches verify against all probabilities within the intervals [JL91; GLD00; SVA06; Sku09; PLSS13]. Lumpability of interval CTMCs is considered in [CGLT⁺21]. In contrast to upCTMCs, interval Markov chains have no dependencies between transition uncertainties, and no distributions are attached to the intervals.

Parametric Markov models generally define probabilities or rates via functions over the parameters. The standard parameter synthesis problem for discrete-time models is to find all valuations of parameters that satisfy a specification. Techniques range from computing a solution function over the parameters, to directly solving the underlying optimization problems [Daw04; HHZ11b; JÁHJ⁺24; CJJK⁺22]. Parametric CTMCs are investigated in [HKM08; CDPK⁺17], but are generally restricted to a few parameters. The work [CCGK⁺18] aims to find a robust parameter valuation in pCTMCs.

For all approaches listed so far, the results may be rather conservative, as no prior information on the uncertainties (the intervals) is used. That is, the uncertainty is not quantified, and all probabilities or rates are treated equally as likely. In our approach, we do not compute solution functions, as the underlying methods are computationally expensive and usually restricted to a few parameters.

Quantified uncertainty is studied in [MMAG14]. Similarly to our work, the approach draws parameter values from a probability distribution over the model parameters and analyzes the induced model via model checking. However, [MMAG14] studies DTMCs and performs a frequentist (Monte Carlo) approach, cf. Sect. 12.6, to compute estimates for a single measure, without prediction regions. Moreover, our approach requires significantly fewer samples, cf. the comparison in Sect. 12.6.4.

The work in [BMS16; BS18] takes a sampling-driven Bayesian approach for pCTMCs. In particular, they take a prior on the solution function over a single measure and update it based on samples (potentially obtained via statistical model checking). We assume no prior on the solution function and, as mentioned before, do not compute the solution function due to the expensive underlying computations.

Statistical model checking (SMC) [AP18; LDB10; DHKP17; LLTY⁺19] samples paths in stochastic models to perform model checking. This technique has been applied to numerous models [DLLM⁺15; DLLM⁺11; DHS18; RGRK⁺09], including CTMCs [SVA05; YS06]. SMC analyzes a *concrete* CTMC by sampling from the known transition rates, whereas for upCTMC, these rates are parametric.

The scenario approach [CG08; CCG21] is widely used in control theory [CC06] and recently in machine learning [CG20] and reliability engineering [RC21]. Within a verification context, closest to our work is our paper [2], which we discussed in Chapter 9 and considers the verification of single measures for uncertain parametric MDPs. [2] relies on the so-called sampling-and-discriminating approach [CG11], while we use the risk-and-complexity perspective [GC22], yielding better results for problems with many decision variables.

Summary

- ⇒ Uncertain parametric CTMCs (upCTMCs) extend parametric CTMCs (pCTMCs) with a probability distribution over the parameter space.
- ⇒ We presented a novel sampling-based approach for analyzing upCTMCs.
- ⇒ Our method provides statistical guarantees on the typical performance characteristics from a finite set of parameter samples.
- ⇒ As demonstrated by our experiments, using a few hundred samples is sufficient to obtain high-confidence statistical guarantees.

13 CTMCs With Imprecisely Timed Observations

Summary | We consider runtime monitoring for continuous-time Markov chains (CTMCs). In such applications, we must incorporate past observations, whose timings may be uncertain. Thus, we consider a setting in which we are given a sequence of imprecisely timed labels called the evidence. The problem is to compute reachability probabilities, which we condition on this evidence. Our key contribution is a method that solves this problem by unfolding the CTMC states over all possible timings for the evidence. We formalize this unfolding as a Markov decision process (MDP) in which each timing for the evidence is reflected by a scheduler. This MDP has infinitely many states and actions in general, making a direct analysis infeasible. Thus, we abstract this continuous MDP into a finite interval Markov decision process (IMDP) and develop an iterative refinement scheme to upper-bound conditional probabilities in the CTMC. We showcase the feasibility of our method on several numerical benchmarks.

Origins | This chapter is based on

- [11] Badings, Volk, Junges, Stoelinga and Jansen (2024) ‘*CTMCs with Imprecisely Timed Observations*’. TACAS.

In this paper, we present the first method to compute weighted conditional reachability probabilities in CTMCs with imprecisely timed observations.

Background | Familiarity with CTMCs and their analysis as discussed in Chapter 11 is assumed. Furthermore, the reader is assumed to be familiar with (i)MDPs and computing reachability probabilities for these models, as discussed in Chapter 3.



13.1 Introduction

Standard model checking algorithms for continuous-time Markov chains (CTMCs) assume a static and known initial state [BHHK03; ASSB00]. However, in applications such as *runtime monitoring* [SSAB⁺19; BDDF⁺18], we need to analyze an already running system without a static initial state. As no initial state is known, we must incorporate past observations instead. These observations are given as a sequence of CTMC labels, each of which is observed at a specific time. We call this sequence of timed labels the *evidence*. We want to incorporate this evidence by conditioning the state of the CTMC on the evidence. For example, we want to answer the question: “*What is the probability of a failure for a production machine (modeled as a CTMC) before time T, given that we have observed particular labels at earlier times t_1, t_2, \dots, t_n ?*”

Imprecise observation times | These conditional probabilities depend on the *exact time* at which each label was observed. However, in realistic scenarios, the times for the labels in the evidence may not be known precisely. For example, suppose that inspections are always done in the first week of a month, but the precise moment of inspection may be unknown. Intuitively, we can interpret such *imprecisely timed evidence* as a potentially infinite set of (precisely timed) *instances* of the evidence that vary only in the observation times. For example, an inspection done on “*January 2 exactly at noon*” is an instance of the imprecise observation time of “*the first week of January*.” This perspective motivates a robust version of the previous question: “Given the imprecisely timed evidence, what is the *maximal* probability of a failure before time T over all instances of the evidence?”

Overall goal | In this chapter, we consider a setting in which we are given a labeled CTMC together with imprecisely timed evidence. For each instance of the evidence, we can define the probability of reaching a set of target states, conditioned on that evidence. The problem is to compute the supremum over these conditional probabilities for all instances of the evidence. We generalize this problem by considering *weighted* conditional reachability probabilities (or simply the *weighted reachability*), where we assign to each state a nonnegative weight. Standard conditional reachability is then a special case with a weight of one for the target states and zero elsewhere. Overall, our main contribution is the first method to compute weighted conditional reachability probabilities in CTMCs with imprecisely timed evidence.

Outline | This chapter is structured as follows. First, in Sect. 13.2, we formalize the problem statement of conditional reachability probabilities for CTMCs. Second, in Sect. 13.3, we present our approach to solving this problem, which is based on an unfolding of the CTMC into an Markov decision process (MDP) with (potentially) infinitely many states and actions. Analyzing this infinite MDP is infeasible, so in Sect. 13.4 we create an abstraction into a finite interval Markov decision process (IMDP). In Sect. 13.5, we use this abstraction to compute bounds on conditional reachability probabilities. In Sect. 13.6, we show the feasibility of our method on several numerical benchmarks. We discuss related work in Sect. 13.7. Finally, we provide mathematical proofs in Sect. 13.8 and discuss open challenges in Sect. 13.9.

13.2 The CTMC Monitoring Problem

The key problem we want to solve is to compute reachability probabilities for the CTMC conditioned on a timed sequence of labels. We call such a timed sequence of labels the *evidence*. In this section, we define the concept of evidence more formally and state the formal *CTMC monitoring* problem that we solve.

CTMC
monitoring

Labeled CTMCs | In this chapter, we consider labeled CTMCs as in Def. 11.1. For simplicity, we assume that the labeling function maps every state $s \in S$ to exactly one atomic proposition, i.e., the labeling function is of the form $L: S \rightarrow AP$. This restriction is without loss of generality, and our results can be extended to CTMCs with a labeling

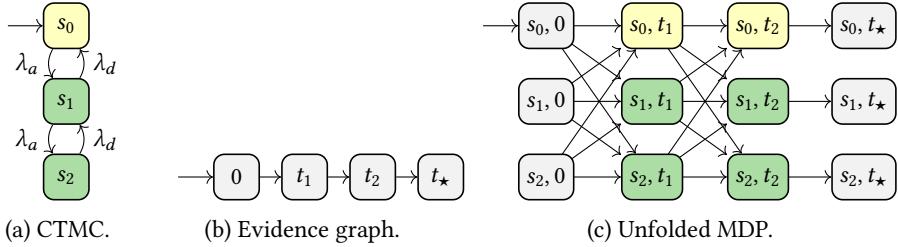


Figure 13.1: The CTMC (a) for Example 13.1, (b) the graph for the precise evidence $\rho = (t_1, o_1), (t_2, o_2)$, and (c) the states of the MDP unfolding from Def. 13.7.

function of the form $L: S \rightarrow 2^{AP}$ as in Def. 11.1.¹ We will also refer to the atomic propositions as *colors*.

colors

Example 13.1 (Inventory system) Consider a simple, single-product inventory where the number of items in stock ranges from 0 to 2, but we can only observe whether the inventory is empty or not. The inventory is initially empty, and products arrive and deplete with a fixed arrival and depletion rate, respectively. This system is modeled by the CTMC shown in Fig. 13.1a with states $S = \{s_0, s_1, s_2\}$ (modeling the stock level) and labels shown by the two colors (\circlearrowleft for empty and \bullet for nonempty). The rates at which items arrive and deplete are $R(s_0, s_1) = R(s_1, s_2) = \lambda_a = 3$ and $R(s_1, s_0) = R(s_2, s_1) = \lambda_d = 2$, respectively.

Evidence | The *evidence* $\rho = (t_1, o_1), \dots, (t_d, o_d) \in (\mathbb{R}_{>0} \times C)^d$ is a sequence of d times and labels such that $t_i < t_{i+1}$ for all $i \in \{1, \dots, d-1\}$. A timed label (t_i, o_i) means that at time t_i , the CTMC was in a state $s \in S$ for which $L(s) = o_i$. Since each time $t \in \mathbb{R}_{>0}$ can only occur once in ρ , we overload ρ and denote the evidence at time $t \in \{t_1, \dots, t_d\}$ by $\rho(t) = o \in C$, such that $(t, o) \in \rho$.

evidence

While a timed path of a CTMC describes the state at every continuous point in time, the evidence only contains the observations at d points in time. We say that a path π is *consistent* with evidence ρ , written as $\pi \models \rho$, if each timed label in ρ matches the label of path π at time t , i.e., if $L(\pi(t)) = \rho(t) \forall t \in \{t_1, \dots, t_d\}$.

path consistent with evidence

Example 13.2 An example of evidence for the CTMC in Fig. 13.1a is $\rho = (0.4, \circlearrowleft), (1.9, \bullet)$. The path fragment $\pi = s_0 1 s_1 0.5 s_2 0.8 s_1$ is consistent with evidence ρ , i.e., $\pi \models \rho$, because the CTMC was in state s_0 (with color \circlearrowleft) after time 0.4, and in state s_2 after time 1.9.

Conditional probabilities | We want to compute the conditional probability $\Pr^C(\pi(t_d) = s) \mid [\pi \models \rho]$ that the CTMC C with initial state s_I generates a path being in state s at time t_d , conditioned on the evidence ρ . Using *Bayes' rule*, we can

Bayes' rule

¹Specifically, we would define a larger set of atomic propositions AP' , such that each $c \in AP'$ represents a possible subset of atomic propositions $C \subseteq AP$ in the original CTMC.

characterize this conditional probability as follows (assuming $\frac{0}{0} = 0$, for brevity):

$$\Pr^C(\pi(t_d) = s \mid [\pi \models \rho]) = \frac{\Pr^C([\pi(t_d) = s] \cap [\pi \models \rho])}{\Pr^C(\pi \models \rho)}. \quad (13.1)$$

Imprecise timings | We extend evidence with uncertainty in the timing of each label. The *imprecisely timed evidence* (or *imprecise evidence*) $\Omega = (\mathcal{T}_1, o_1), \dots, (\mathcal{T}_d, o_d)$ is a sequence of d labels and uncertain timings $\mathcal{T}_i = \cup_{j=1}^q [\underline{t}_j, \bar{t}_j]$, with $\underline{t}_j \leq \bar{t}_j$, $q \in \mathbb{N}$.

Note that the uncertain timing \mathcal{T}_i can model both singletons ($\mathcal{T}_i = \{1, 2, 3\}$) and unions of intervals ($\mathcal{T}_i = [1, 1.5] \cup [2, 2.5]$). We require that $\max_{t \in \mathcal{T}_i}(t) < \min_{t' \in \mathcal{T}_{i+1}}(t')$ for all $i \in \{1, \dots, d-1\}$, which means that the order of the labels is known, despite the uncertainty in the observation times. Again, we overload notation and denote the evidence at time t by $\Omega(t) = o$, such that $\exists (\mathcal{T}, o) \in \Omega$ with $t \in \mathcal{T}$.

Imprecise evidence induces a set of *instances* of the evidence that only differ in the label times. This set of instances is uncountably infinite whenever one of the imprecise timings \mathcal{T} is a continuous set. Formally, the evidence $\rho = (t_1, o_1), \dots, (t_d, o_d)$ is an instance of the imprecise evidence Ω , written as $\rho \in \Omega$, if $t_i \in \mathcal{T}_i$ for all $i = 1, \dots, d$.

Example 13.3 An example of imprecise evidence for the CTMC in Example 13.1 is $\Omega = ([0.2, 0.8], \textcolor{yellow}{\circ}), ([1.4, 2.1], \textcolor{green}{\circ})$. The precise evidence $\rho = (0.4, \textcolor{yellow}{\circ}), (1.9, \textcolor{green}{\circ})$ described in Example 13.2 is an instance of Ω , i.e., $\rho \in \Omega$. However, $\rho' = (0.1, \textcolor{yellow}{\circ}), (1.9, \textcolor{green}{\circ})$ and $\rho'' = (0.4, \textcolor{green}{\circ}), (1.9, \textcolor{yellow}{\circ})$ are not, i.e., $\rho' \notin \Omega$, $\rho'' \notin \Omega$, as the timings and labels do not match, respectively.

13.2.1 Problem statement

Let $w: S \rightarrow \mathbb{R}_{\geq 0}$ be a so-called *state-weight function*, which assigns to each CTMC state $s \in S$ a nonnegative weight. The weight $w(s)$ represents a general measure of risk associated with each state $s \in S$, as used in [JTS21]. For example, $w(s)$ may represent the probability of reaching a set of target states S_G from $s \in S$ within some time horizon $h \geq 0$. We then consider the following problem.

Problem 13.4 (Weighted conditional reachability) Given a CTMC C , a state-weight function w , and the imprecisely timed evidence Ω , compute the (maximal) weighted conditional reachability probability $W(\Omega)$:

$$W(\Omega) = \sup_{\rho \in \Omega} \sum_{s \in S} \Pr^C(\pi(t_d) = s \mid [\pi \models \rho]) \cdot w(s).$$

Let us illustrate Problem 13.4 by means of the following example.

Example 13.5 For the CTMC in Example 13.1, consider the state-weight function that assigns to each state the probability of reaching state s_0 within time $t = 0.1$. Then, the problem above is interpreted as follows: “Given the imprecisely timed evidence Ω , compute the probability (conditioned on Ω) of reaching state s_0 within time $t = 0.1$ (after the final time point of the evidence).”

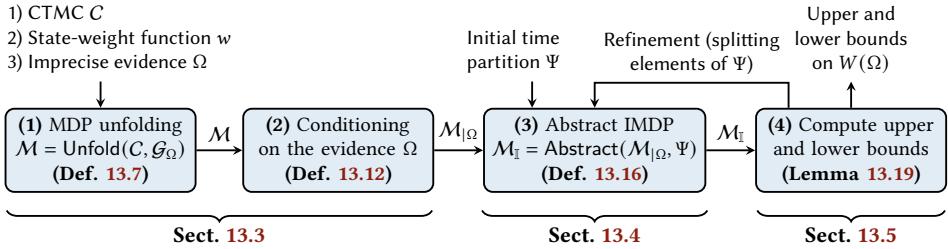


Figure 13.2: Conceptual workflow of our approach for solving Problem 13.4.

Variations | For ease of notation, we focus on computing the *maximal* weighted conditional reachability as formulated in Problem 13.4. To instead compute the *minimal* weighted conditional reachability, we would swap every inf and sup (and max and min) in our approach, but our general approach remains the same. Furthermore, by setting $w(s) = 1$ for all $s \in S_G$ and $w(s) = 0$ otherwise, we can also compute the probability of being in a state in S_G *immediately* after the evidence. Finally, we remark that Problem 13.4 only considers events *after* the end of the evidence. This setting is motivated by applications where the exact system state is not observable, but where actual system failures can be observed. Thus, one can typically assume that the system has not failed yet, and the problem, as formalized in Problem 13.4, is to predict the conditional probability of a future system failure.

13.2.2 Our approach

Our overall workflow to solve Problem 13.4 is summarized in Fig. 13.2 and consists of four blocks. We briefly introduce our approach from a high-level perspective, after which we discuss each block in more detail in Sects. 13.3 to 13.5.

1) Unfolding | The first step is to *unfold* the CTMC over all possible timings of the imprecisely timed evidence. We formalize this unfolding as an MDP [Put94], in which the timing imprecision is reflected by nondeterminism.

2) Conditioning | The weighted reachability $W(\Omega)$ in Problem 13.4 can be computed via (unconditional) reachability probabilities on a transformed version of this MDP. Inspired by ideas from [BKKM14; JTS21], we *refute* paths through the model that are *inconsistent* with all evidence instances $\rho \in \Omega$, and instead loop these paths back to the initial state. For the special case of evidence with precise observation times, we obtain a precise solution to the problem that we can directly compute. By contrast, imprecisely timed evidence yields an unfolded MDP with infinitely many states and actions.

3) Abstraction | Analyzing the infinite-state/action MDP directly is infeasible. Thus, we propose an abstraction of this infinite MDP as a finite IMDP, similar to game-based abstractions [KKNP10]. A robust analysis of the IMDP yields upper and lower bounds on the weighted reachability for the CTMC. Moreover, we propose an iterative refinement scheme that converges to the weighted reachability in the limit.

4) Computing bounds in practice | Finally, we use the IMDP abstraction and refinement to obtain sound upper and lower bounds on the weighted reachability in practice. In Sect. 13.6, we show the feasibility of our method across several numerical benchmarks. Concretely, we show that we obtain reasonably tight bounds on the weighted reachability within a reasonable time.

13.3 Conditional Reachability With Imprecise Evidence

In this section, we treat the first two blocks of Fig. 13.2. In Sect. 13.3.1, we *unfold* the CTMC over the times in the imprecise evidence into an MDP. The main result of this section, Theorem 13.11, states that the conditional reachability on the CTMC in Problem 13.4 is equal to the *maximal* conditional reachability probabilities in the MDP over a *subset of schedulers* (those that we call *consistent*; see Def. 13.8). In Sect. 13.3.2, we use results from [BKKM14] to determine these conditional probabilities via unconditional reachability probabilities on a transformed version of the MDP.

13.3.1 Unfolding the CTMC into an MDP

We interpret the (precisely timed) evidence $\rho = (t_1, o_1), \dots, (t_d, o_d)$ as a directed graph that encodes the trivial progression over the time steps t_1, \dots, t_d . While this perspective is trivial for precisely timed evidence, it will be useful when dealing with imprecisely timed evidence later on.

evidence graph

Definition 13.6 (Evidence graph) An *evidence graph* $\mathcal{G} = (\mathcal{N}, \mathcal{E})$ is a directed graph where each node $t \in \mathcal{N} \subseteq \mathbb{R}_{>0} \cup \{0, \dots, t_\star\}$ is a point in time (or the initial time 0 or terminal time t_\star), and with directed edges $\mathcal{E} \subset \{t \rightarrow t' : t, t' \in \mathcal{N}\}$, such that $t' > t$ for all $t \rightarrow t' \in \mathcal{E}$.

The evidence graph $\mathcal{G}_\rho = (\mathcal{N}_\rho, \mathcal{E}_\rho)$ for the precise evidence ρ has nodes $\mathcal{N}_\rho = \{0, t_1, \dots, t_d, t_\star\}$ and edges $\mathcal{E}_\rho = \{t_{i-1} \rightarrow t_i : i = 2, \dots, d\} \cup \{0 \rightarrow t_1, t_d \rightarrow t_\star\}$. As illustrated in Fig. 13.1b, the graph \mathcal{G}_ρ has exactly one path, which follows the time points t_1, \dots, t_d of the evidence ρ itself. Likewise, we model the imprecise evidence Ω as a graph \mathcal{G}_Ω which is the union of all graphs \mathcal{G}_ρ for all instances $\rho \in \Omega$, i.e.,

$$\mathcal{G}_\Omega = (\mathcal{N}_\Omega, \mathcal{E}_\Omega) = \cup_{\rho \in \Omega} (\mathcal{G}_\rho) = (\cup_{\rho \in \Omega} (\mathcal{N}_\rho), \cup_{\rho \in \Omega} (\mathcal{E}_\rho)).$$

If Ω has infinitely many instances, then \mathcal{G}_Ω has infinite branching. Importantly, we observe that every path $t_0 t_1 \dots t_d t_\star$ through graph \mathcal{G}_Ω corresponds to the time points of the precise evidence $\rho = (t_1, o_1), \dots, (t_d, o_d) \in \Omega$ (and vice versa).

We denote the successor nodes of $t \in \mathcal{N}$ by $\text{post}(t) = \{t' \in \mathcal{N} : t \rightarrow t' \in \mathcal{E}\}$. For example, the graph in Fig. 13.1b has $\text{post}(0) = t_1$, $\text{post}(t_1) = t_2$ and $\text{post}(t_2) = t_\star$.

MDP unfolding | We introduce the *unfolding operator* $\mathcal{M} = \text{Unfold}(C, \mathcal{G})$, which takes a CTMC C and a graph \mathcal{G} , and returns the *unfolded MDP* \mathcal{M} defined as follows. Recall from Sect. 11.2 that $P_{\delta_s}(t) \in \text{Distr}(S)$ denotes the transient probability distribution over states at time t , when starting in the fixed state $s \in S$ (at time zero).

Definition 13.7 (Unfolded MDP) For a CTMC $C = (S, s_I, R, L)$ and a graph $\mathcal{G} = (\mathcal{N}, \mathcal{E})$, the **unfolded MDP** $\text{Unfold}(C, \mathcal{G}) = (Q, A, q_I, P)$ has states $Q = S \times \mathcal{N}$, actions $A = \mathcal{N}$, initial state $q_I = (s_I, 0)$, and transition function P , which is defined for all $(s, t) \in Q$, $t' \in \text{post}(t)$, $s' \in S$ as

$$P((s, t), t', (s', t')) = \begin{cases} \mathbf{P}_{\delta_s}(t' - t)(s') & \text{if } t' \neq t_\star, \\ \mathbb{1}_{(s=s')} & \text{if } t' = t_\star, \end{cases} \quad (13.2)$$

unfolded
MDP

As an example, the unfolding of the CTMC in Fig. 13.1a over the graph in Fig. 13.1b is shown in Fig. 13.1c. A state $(s, t) \in Q$ in the unfolded MDP is interpreted as being in CTMC state $s \in S$ at time t .

In state (s, t) , the set of enabled actions is $A((s, t)) = \text{post}(t) \subset \mathcal{N}$, and taking an action $t' \in \text{post}(t)$ corresponds to *deterministically* jumping to time t' . In other words, the set of enabled actions coincides with the successor states in the graph \mathcal{G} . Thus, for precisely timed evidence, each state in the unfolded MDP has exactly one enabled action, i.e., the model is, in fact, a discrete-time Markov chain (DTMC). For imprecisely timed evidence, however, we obtain an unfolded MDP with infinitely many actions enabled in each state in general.

The effect of each action is *stochastic* and determines the next CTMC state. The transition probability $P((s, t), t', (s', t'))$ for $t' \neq t_\star$ models the probability of starting in CTMC state $s \in S$ and being in state $s' \in S$ after time $t' - t$ has elapsed, which is precisely the transient probability $\mathbf{P}_{\delta_s}(t' - t)(s')$ defined in Sect. 11.2. Finally, the (terminal) states (s, t_\star) for all $s \in S$ are absorbing.

Interpretation of schedulers | Every instance $\rho \in \Omega$ of the imprecise evidence $\Omega = (\mathcal{T}_1, o_1), \dots, (\mathcal{T}_d, o_d)$ corresponds to fixing a precise time $t_i \in \mathcal{T}_i$ for all $i = 1, \dots, d$. For each such evidence instance $\rho \in \Omega$, there exists a (memoryless deterministic; see Sect. 3.1.1) scheduler $\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}$ for MDP $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$ that induces a Markov chain which only visits those times t_1, \dots, t_d . We call such a scheduler σ *consistent* with the evidence ρ .

Definition 13.8 (Consistent scheduler) A scheduler $\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}}$ is *consistent* with $\rho = (t_1, o_1), \dots, (t_d, o_d) \in \Omega$, written as $\sigma \sim \rho$, if for all CTMC states $s \in S$, we have

$$\sigma((s, 0)) = t_1, \quad \sigma((s, t_i)) = t_{i+1} \quad \forall i \in \{0, \dots, d-1\}, \quad \sigma((s, t_d)) = t_\star.$$

We denote the set of all consistent schedulers by $\widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}} \subseteq \mathfrak{S}_{\text{stat}}^{\mathcal{M}}$.

consistent
scheduler

A consistent scheduler chooses the same action $\sigma((s, t)) = \sigma((s', t'))$ in any two MDP states $(s, t), (s', t') \in Q$ for which $t = t'$. There is a one-to-one correspondence between choices $\rho \in \Omega$ and consistent schedulers: For every $\rho \in \Omega$, there exists a scheduler $\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}}$ such that $\sigma \sim \rho$, and vice versa.

13

Example 13.9 Consider imprecise evidence $\Omega = ([0.2, 0.8], \text{O}), ([1.4, 2.1], \text{O})$ for the CTMC in Example 13.1. A scheduler with $\sigma((s_0, 0.4)) = 1.5$, $\sigma((s_1, 0.4)) = 1.8$ is inconsistent as it chooses different actions in MDP states with the same time.

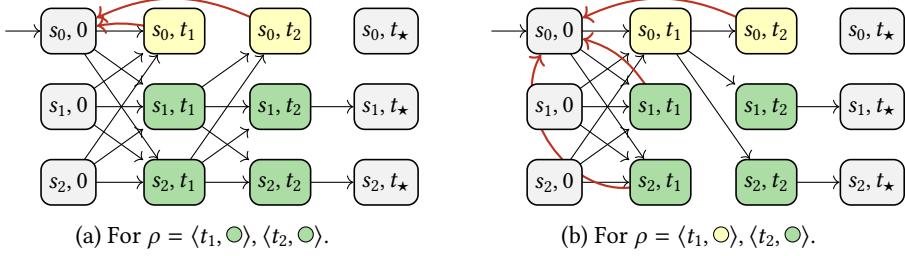


Figure 13.3: The unfolded MDP from Fig. 13.1c conditioned on different precise evidences. States that do not agree with the evidence are looped back to the initial state (depicted as red transitions).

Remark 13.10 The unfolded MDP $\mathcal{M}' = \text{Unfold}(C, \mathcal{G}_\rho)$ for the precise evidence ρ has only a single action enabled in every state. Hence, \mathcal{M}' has only one scheduler and directly reduces to a discrete-time Markov chain, such that we have $\widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}'} = \mathfrak{S}_{\text{stat}}^{\mathcal{M}'}$.

Conditional reachability on unfolded MDP | As a main result, we show that $W(\Omega)$ in Problem 13.4 can be expressed as maximizing conditional reachability probabilities in the unfolded MDP \mathcal{M} over the consistent schedulers $\widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}} \subset \mathfrak{S}_{\text{stat}}^{\mathcal{M}}$.

Theorem 13.11 (Conditional reachability in unfolded MDP) For a CTMC C and the imprecise evidence Ω with graph \mathcal{G}_Ω , let $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$ be the unfolded MDP. Then, using the notation from Sect. 3.2.2 (for the probability measure $\Pr_\sigma^{\mathcal{M}}$ over MDP paths $\xi \in \Pi_{\text{fin}}^{\mathcal{M}}$), Problem 13.4 is equivalently written as

$$W(\Omega) = \sup_{\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}}} \sum_{s \in S} \Pr_\sigma^{\mathcal{M}}(q_I \models \diamond(s, t_\star) \mid [\xi \models \rho, \sigma \sim \rho]) \cdot w(s).$$

Proof. The proof is slightly tedious and is, therefore, presented in Sect. 13.8 at the end of this chapter. Intuitively, the proof shows that for every instance $\rho \in \Omega$, the conditional transient probabilities in the CTMC are equivalent to conditional reachability probabilities in the unfolded MDP under a $\sigma \sim \rho$ consistent to ρ .

13.3.2 Computing conditional probabilities in MDPs

We describe a transformation of the unfolded MDP to compute the conditional reachability probabilities in Theorem 13.11. Intuitively, we *refute* all paths through the MDP that do not agree with the labels in the evidence. Specifically, we find the subset of MDP states $Q_{\text{reset}}(\Omega) \subset Q$ that disagree with the evidence, defined as

$$Q_{\text{reset}}(\Omega) = \{(s, t) \in Q : L(s) \neq \Omega(t)\} \subset Q. \quad (13.3)$$

We *reset* all states in $Q_{\text{reset}}(\Omega)$ by adding transitions back to the initial state with probability one. Formally, we define the *conditioned MDP* $\mathcal{M}|_\Omega$ as follows.

Definition 13.12 (Conditioned MDP) For $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega) = (Q, A, q_I, P)$, the *conditioned MDP* $\mathcal{M}|_\Omega = (Q, A, q_I, P|_\Omega)$ has the same states and actions, but the transition function is defined for all $(s, t) \in Q, t' \in \text{post}(t), s' \in S$ as

$$P|_\Omega((s, t), t', (s', t')) = \begin{cases} P((s, t), t', (s', t')) & \text{if } (s, t) \notin Q_{\text{reset}}(\Omega), \\ \mathbb{1}_{(s'=s)} & \text{if } (s, t) \in Q_{\text{reset}}(\Omega). \end{cases}$$

Two examples of conditioning on precise evidence are shown in Fig. 13.3. Compared to Fig. 13.1c, we removed all probability mass over paths that are not consistent with the evidence and normalized the probabilities for all other paths. The following result from [BKKM14] shows that conditional reachabilities in the unfolded MDP are equal to *unconditional* reachabilities in the conditioned MDP.

Lemma 13.13 (Thm. 1 in [BKKM14]) Let Ω be imprecisely timed evidence, C a CTMC, and $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$ an unfolded MDP. Furthermore, let $\mathcal{M}|_\Omega$ be the conditioned MDP defined by Def. 13.12. For all consistent schedulers $\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}}$ and for all CTMC states $s \in S$, it holds that

$$\Pr_{\sigma}^{\mathcal{M}}(q_I \models \diamond(s, t_\star) \mid [\xi \models \rho, \sigma \sim \rho]) = \Pr_{\sigma}^{\mathcal{M}|_\Omega}(q_I \models \diamond(s, t_\star)).$$

Importantly, observe that Lemma 13.13 holds only for *consistent schedulers* (for inconsistent schedulers, $\sigma \sim \rho$ does not exist).

Combining Lemma 13.13 with Theorem 13.11 directly expresses the conditional reachability $W(\Omega)$ in terms of reachability probabilities on the conditioned MDP.

Theorem 13.14 (Solution to Problem 13.4 via conditioned MDP) Given a CTMC C , a state-weight function w , and the imprecisely timed evidence Ω , let $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$. Then, it holds that

$$W(\Omega) = \sup_{\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}}} \sum_{s \in S} \Pr_{\sigma}^{\mathcal{M}|_\Omega}(q_I \models \diamond(s, t_\star)) \cdot w(s).$$

Solving Problem 1 with precisely timed evidence is now straightforward by solving a finite DTMC, see Remark 13.10. Furthermore, if the imprecise evidence has finitely many instances, then the MDP is finite. A naïve approach to optimize over the consistent schedulers is by enumeration, which we discuss in more detail in Sect. 13.5.

13.3.3 Computing evidence probability

We discuss a variation of Problem 13.4, which computes the probability for observing a given precise evidence ρ . We show that, with minor modifications to our unfolding procedure, we can compute the probability that a CTMC generates the given (precise) evidence ρ . Instead of looping all states $(s, t) \in Q_{\text{reset}}$ inconsistent with the evidence (defined in Eq. (13.3)) back to the initial state, we now create self-loops for those states. Formally, given an unfolded MDP (or in fact, DTMC) $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\rho) = (Q, A, q_I, P)$ for precise evidence ρ , we define the modified MDP $\mathcal{M}_\rho = (Q, A, q_I, P_\rho)$ with transition

conditioned
MDP

function P_ρ defined for all $(s, t), (s', t') \in Q$ as

$$P_\rho((s, t), t', (s', t')) = \begin{cases} P((s, t), t', (s', t')) & \text{if } (s, t) \notin Q_{\text{reset}}(\Omega), \\ 1_{((s, t)=(s', t'))} & \text{if } (s, t) \in Q_{\text{reset}}(\Omega), \end{cases}$$

with Q_{reset} defined by Eq. (13.3). This transformation of the unfolded MDP is illustrated by Fig. 13.4 for two different precisely timed evidences. Then, the probability $\Pr^C(\pi \models \rho)$ that CTMC C generates the evidence ρ is the probability that the modified MDP \mathcal{M}_ρ reaches a state (s, t_\star) for time t_\star and any CTMC state $s \in S$ (since \mathcal{M}_ρ is, in fact, a DTMC, we omit the scheduler):

$$\Pr^C(\pi \models \rho) = \sum_{s \in S} \Pr^{\mathcal{M}_\rho}(q_I \models \diamond(s, t_\star)). \quad (13.4)$$

Intuitively, the right-hand side of Eq. (13.4) is the probability of ever reaching a terminal state at time t_\star . Because all paths inconsistent with the evidence ρ are trapped by the self-loops (in non-terminal states), Eq. (13.4) thus is the probability that the CTMC generates a path that is consistent with ρ . For imprecise evidence Ω , we can also ask for the *worst-case* probability to obtain any instance $\rho \in \Omega$, by modifying the unfolded MDP $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$ in an analogous manner to Def. 13.12.

13.4 Abstraction of Conditioned MDPs

For imprecisely timed evidence with *infinitely many instances* (e.g., imprecise timings over intervals), the conditioned MDP from Sect. 13.3 has infinitely many states and actions. In this section, we treat block (3) of Fig. 13.2 and propose an *abstraction* of this continuous MDP into a finite IMDP. Similar to game-based abstractions [KKNP10; HHWZ10; HNPW⁺11], we capture abstraction errors as nondeterminism in the transition function of the IMDP. We will show that robust reachability probabilities in the IMDP yield sound bounds on the conditional reachability $W(\Omega)$.

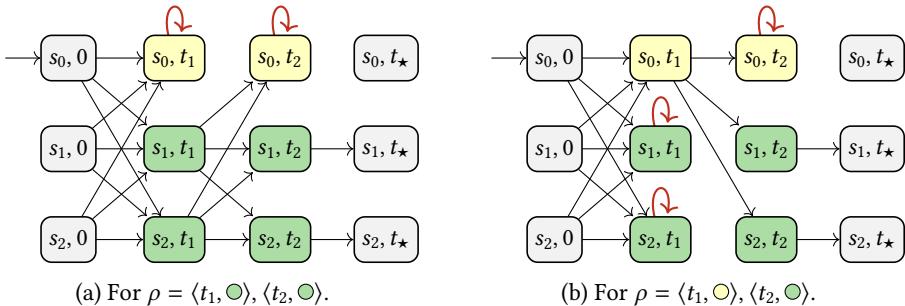


Figure 13.4: Using the unfolded MDP to compute the probability for two precisely timed evidences. States that do not agree with the evidence are made absorbing (depicted as red transitions).

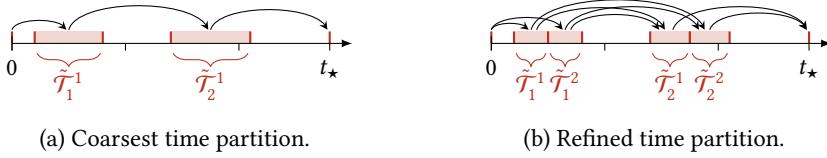


Figure 13.5: Two partitions of imprecise evidence $\Omega = ([0.2, 0.8], o_1), ([1.4, 2.1], o_2)$. The partition in (a) consists of two elements, such that $\tilde{\mathcal{T}}_1^1 = [0.2, 0.8]$ and $\tilde{\mathcal{T}}_2^1 = [1.4, 2.1]$, where (b) refines this to $\tilde{\mathcal{T}}_1^1 \cup \tilde{\mathcal{T}}_1^2 = [0.2, 0.8]$ and $\tilde{\mathcal{T}}_2^1 \cup \tilde{\mathcal{T}}_2^2 = [1.4, 2.1]$.

13.4.1 Abstracting evidence times

The crux of our abstraction is to create a finite partition of the (infinite) sets of uncertain timings in the evidence, as illustrated by Fig. 13.5. A partition of a set X is a collection of sets $\text{partition}(X) = (X_1, \dots, X_n)$ that cover X (i.e., $X = \bigcup_{i=1}^n X_i$) and the interior of each element is disjoint (i.e., $\text{int}(X_i) \cap \text{int}(X_j) = \emptyset$, $i, j \in \{1, \dots, n\}$, $i \neq j$).

Definition 13.15 (Time partition) A *time partition* Ψ of the imprecise evidence $\Omega = (\mathcal{T}_1, o_1), \dots, (\mathcal{T}_d, o_d)$ is a set $\Psi = \cup_{i=1}^d \text{partition}(\mathcal{T}_i) \cup \{0, t_★\}$, where each $\text{partition}(\mathcal{T}_i) = \{\mathcal{T}_i^1, \dots, \mathcal{T}_i^{n_i}\}$ is a finite partition of \mathcal{T}_i into $n_i \in \mathbb{N}_{>0}$ elements.

With abuse of notation, the element of Ψ containing time t is $\Psi(t) \in \Psi$, and $\Psi^{-1}(\psi) = \{t : \Psi(t) = \psi\}$ is the set of times mapping to $\psi \in \Psi$. As shown by Fig. 13.5, for each $i \in \{1, \dots, d\}$, the sets $\tilde{\mathcal{T}}_i^1, \dots, \tilde{\mathcal{T}}_i^{n_i}$ are a partition of the set \mathcal{T}_i .

To illustrate the abstraction, let $(s, t) \xrightarrow{t':P'} (s', t')$ denote the MDP transition from state $(s, t) \in Q$, under action $t' \in A((s, t))$ to state $(s', t') \in Q$, which has probability P' . Using this notation, we can express any MDP path as

$$(s_I, 0) \xrightarrow{t:P} (s, t) \xrightarrow{t':P'} (s', t') \xrightarrow{t'':P''} \dots \xrightarrow{t'''':P'''} (s, t_★). \quad (13.5)$$

For every element $\psi \in \Psi$ of partition Ψ , the abstraction merges all MDP states $(s, t) \in Q$ for which the time t belongs to the element ψ , that is, for which $t \in \Psi^{-1}(\psi)$. Thus, we merge infinitely many MDP states into finitely many abstract states. The MDP path in Eq. (13.5) matches the next path in the abstraction:

$$(s_I, 0) \xrightarrow{\mathcal{T}:P} (s, \mathcal{T}) \xrightarrow{\mathcal{T}':P'} (s', \mathcal{T}') \xrightarrow{\mathcal{T}'':P''} \dots \xrightarrow{\mathcal{T}''':P'''} (s, t_★), \quad (13.6)$$

where each $t \in \mathcal{T}$, and each P is a *set of probabilities*. The abstraction contains the behavior of the continuous MDP if $P \in \mathcal{P}$ at every step in Eqs. (13.5) and (13.6), see, e.g., [JL91]. The following IMDP abstraction satisfies these requirements.

Definition 13.16 (IMDP abstraction) For conditioned MDP $\mathcal{M}_{|\Omega} = (Q, A, q_I, P)$ and time partition Ψ of Ω , the IMDP $\mathcal{M}_{\mathbb{I}} = \text{Abstract}(\mathcal{M}_{|\Omega}, \Psi) = (\tilde{Q}, \tilde{A}, q_I, \mathcal{P})$, with states $\tilde{Q} = \{(s, \Psi(t)) : (s, t) \in Q\}$, actions $\tilde{A} = \{\Psi(t) : t \in A\}$, and uncertain

transition function \mathcal{P} defined for all $(s, \mathcal{T}), (s', \mathcal{T}') \in \tilde{\mathcal{Q}}$ as

$$\mathcal{P}((s, \mathcal{T}), \mathcal{T}', (s', \mathcal{T}')) = \text{closure}\left(\bigcup_{\substack{t \in \Psi^{-1}(\mathcal{T}), \\ t' \in \Psi^{-1}(\mathcal{T}')}} P((s, t), t', (s', t'))\right),$$

where $\text{closure}(x) = [\min(x), \max(x)]$ is the interval closure of x .

An abstraction under the coarse time partition from Fig. 13.5a is shown in Fig. 13.6a. The transition probabilities for each MDP state are defined by transient probabilities for the CTMC. Thus, the uncertain transition function \mathcal{P} of the IMDP overapproximates these transient probabilities over a *range of times* (as shown in Fig. 13.6b), yielding the probability intervals for $P(q)$ and $P(q'')$ shown in Fig. 13.6c.

Conditional reachability on IMDP | We show that the IMDP abstraction can be used to obtain sound upper and lower bounds on the conditional reachability $W(\Omega)$. Let $W_{\mathcal{M}_I}(\tilde{P}, \sigma) \geq 0$ denote the value for the MDP $\mathcal{M}_I[\tilde{P}]$ induced by IMDP \mathcal{M}_I under transition function \tilde{P} , and with scheduler $\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_I}$:

$$W_{\mathcal{M}_I}(\tilde{P}, \sigma) := \sum_{s \in S} \Pr_{\sigma}^{\mathcal{M}_I[\tilde{P}]}(q_I \models \diamond(s, t_{\star})) \cdot w(s). \quad (13.7)$$

The next theorem, which we will prove in Sect. 13.8, is the main result of this section.

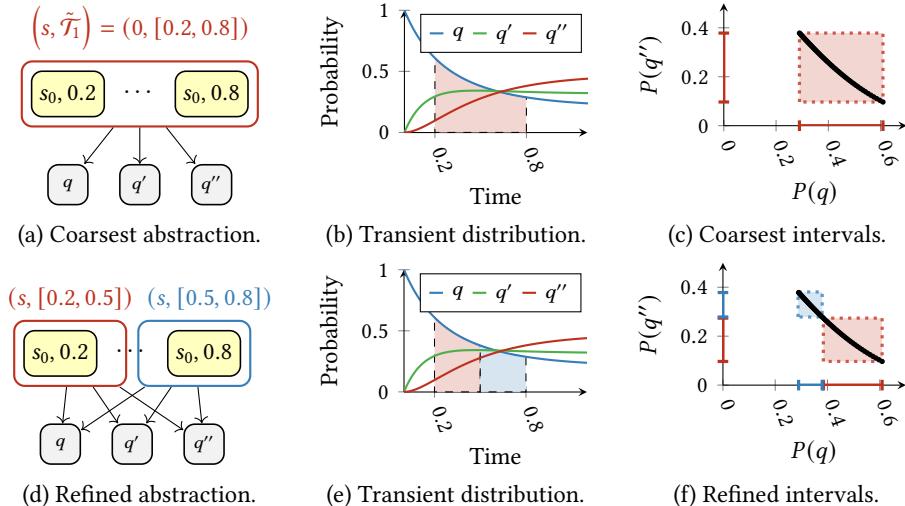


Figure 13.6: Abstraction of an infinite set of MDP states for all times $t \in [0.2, 0.8]$ into (a) a single IMDP state $(s, [0.2, 0.8])$ with probability intervals that overapproximate the transient distribution (b) as the rectangular set in (c), where the line shows the MDP transition probabilities for all $t \in [0.2, 0.8]$. The refinement (d) into two IMDP states $(s, [0.2, 0.5])$ and $(s, [0.5, 0.8])$ splits the approximation of the transient (e) into the two (less conservative) rectangular sets in (f).

Theorem 13.17 (Bounding conditional reachability) For a given conditioned MDP $\mathcal{M}_{|\Omega}$ and a time partition Ψ of Ω , let $\mathcal{M}_{\mathbb{I}} = \text{Abstract}(\mathcal{M}_{|\Omega}, \Psi)$ be the corresponding IMDP abstraction. Then, it holds that

$$\max_{\sigma \in \tilde{\Xi}_{\text{stat}}^{\mathcal{M}_{\mathbb{I}}} \tilde{P} \in \mathcal{P}} W_{\mathcal{M}_{\mathbb{I}}}(\tilde{P}, \sigma) \leq W(\Omega) \leq \max_{\sigma \in \tilde{\Xi}_{\text{stat}}^{\mathcal{M}_{\mathbb{I}}} \tilde{P} \in \mathcal{P}} \max W_{\mathcal{M}_{\mathbb{I}}}(\tilde{P}, \sigma). \quad (13.8)$$

Construction of the IMDP | We want to construct the abstract IMDP directly from the CTMC without first constructing the continuous MDP $\mathcal{M}_{|\Omega}$. Consider computing the probability interval $\mathcal{P}((s, \mathcal{T}), \mathcal{T}', (s', \mathcal{T}'))$ for the IMDP transition from state (s, \mathcal{T}) to (s', \mathcal{T}') . This interval is given by the minimum and maximum transient probabilities $P_{\delta_s}(t' - t)(s')$ over all $t \in \mathcal{T}$ and $t' \in \mathcal{T}'$. However, the problem is that the transient probabilities are not monotonic over time in general (as also illustrated by Figs. 13.6b and 13.6f), so it is unclear how to compute this interval.

Instead, we compute upper and lower bounds for the transient probabilities. Let $\underline{t} = \min(\mathcal{T})$ and $\bar{t} = \max(\mathcal{T})$. An upper bound on the transient probability is given by the probability to reach s' from s at *some* time $t' - t$, $t \in \mathcal{T}$, $t' \in \mathcal{T}'$:

$$\begin{aligned} \sup_{t \in \mathcal{T}, t' \in \mathcal{T}'} P_{\delta_s}(t' - t)(s') &\leq \sup_{t \in \mathcal{T}, t' \in \mathcal{T}'} \Pr^C(s \models \diamond^{[\underline{t}, t']} s') \\ &= \Pr^C(s \models \diamond^{[\underline{t}, \bar{t}']} s'), \end{aligned} \quad (13.9)$$

where we remark that $\bar{t}' - \underline{t}$ is the maximal time difference. A lower bound is given symmetrically by the transient probability to reach s' in the CTMC at the *earliest* possible time $\underline{t}' - \bar{t}$ and staying there for the *full* remaining time $(\bar{t}' - \underline{t}) - (\underline{t}' - \bar{t})$:

$$\inf_{t \in \mathcal{T}, t' \in \mathcal{T}'} P_{\delta_s}(t' - t)(s') \geq P_{\delta_s}(\underline{t}' - \bar{t})(s') \cdot \Pr^C(s \models \square^{[0, (\bar{t}' - \underline{t}) - (\underline{t}' - \bar{t})]} s'). \quad (13.10)$$

13.4.2 Abstraction refinement

To improve the tightness of the bounds in Theorem 13.17, we propose a refinement step that splits elements of the time partition Ψ . For example, we may split the single abstract state in Fig. 13.6a into the two states in Fig. 13.6d.

Definition 13.18 (Refinement of time partition) Let Ψ and Ψ' be partitions as per Def. 13.15, for which $|\Psi'| > |\Psi|$. We call Ψ' a *refinement* of Ψ if for all $\psi' \in \Psi'$, there exists a $\psi \in \Psi$ such that $\psi' \subseteq \psi$.

refinement

Any refinement Ψ' of partition Ψ can be constructed by finitely many splits.

Splitting the partition refines the IMDP's probability intervals, as also shown by Figs. 13.6c and 13.6f. The refined IMDP $\mathcal{M}_{\mathbb{I}}' = \text{Abstract}(\mathcal{M}_{|\Omega}, \Psi')$ has more states and actions but also smaller probability intervals and thus, in general, tighter bounds in Theorem 13.17. Repeatedly refining every element of the partition yields an IMDP with arbitrarily many states and actions and with arbitrarily small probability intervals. Hence, in the limit, we may recover the original continuous MDP by refinements, which also implies that the bounds in Theorem 13.17 on the refined IMDP converge.

Refinement strategy | By splitting every element of the partition Ψ , the number of IMDP states and actions double per iteration, and the number of transitions grows exponentially. Thus, we employ the following *guided refinement strategy*. At each iteration, we extract the scheduler σ^* that attains the upper bound in Theorem 13.17 and determine the set $\tilde{Q}_{\text{reach}}^{\sigma^*} \subset \tilde{Q}$ of reachable IMDP states. We only refine the reachable elements $\psi \in \Psi$, that is, for which there exists a $t \in \psi$ and $s \in S$ such that $(s, t) \in \tilde{Q}_{\text{reach}}^{\sigma^*}$. Using this guided strategy, we iteratively shrink only the relevant probability intervals, resulting in the same convergence behavior as the naïve strategy but without the severe increase in abstraction size.

13.5 Bounding the Conditional Reachability

Theorem 13.17 provides bounds on the conditional reachability $W(\Omega)$ in Problem 13.4, but computing these bounds involves optimizing over the subset of consistent schedulers. Recall from Def. 13.8 that a consistent scheduler chooses the same actions in different states.² As we are not aware of any efficient algorithm to optimize over the consistent schedulers, we compute the following straightforward bounds:

Lemma 13.19 (Bounds on Problem 13.4) Let $\mathcal{M}_I = \text{Abstract}(\mathcal{M}|_\Omega, \Psi)$ be the IMDP abstraction for unfolded MDP $\mathcal{M}|_\Omega$ and a time partition Ψ . It holds that

$$W(\Omega) \leq \max_{\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}_I}} \max_{\tilde{P} \in \mathcal{P}} W_{\mathcal{M}_I}(\tilde{P}, \sigma) \leq \max_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_I}} \max_{\tilde{P} \in \mathcal{P}} W_{\mathcal{M}_I}(\tilde{P}, \sigma).$$

Moreover, any consistent scheduler $\hat{\sigma} \in \widehat{\mathfrak{S}}_{\text{stat}}^{\mathcal{M}_I}$ results in a lower bound.

Obtaining lower bounds | While we can use *any* consistent scheduler in Lemma 13.19 to compute a lower bound on $W(\Omega)$, we obtain better bounds by modifying a (potentially non-consistent) optimal scheduler σ^- under the worst-case choice of probabilities, i.e., $\sigma^- = \operatorname{argmax}_{\sigma \in \mathfrak{S}_{\text{stat}}^{\mathcal{M}_I}} \min_{\tilde{P} \in \mathcal{P}} W_{\mathcal{M}_I}(\tilde{P}, \sigma)$. First, we check the following condition in all pairs of states $(s, t), (s', t') \in \tilde{Q}_{\text{reach}}^{\sigma^-} \subset \tilde{Q}$ reachable under σ^- :

$$t = t' \implies \sigma((s, t)) = \sigma((s', t)) \quad \forall (s, t), (s', t') \in \tilde{Q}_{\text{reach}}^{\sigma^-}.$$

Violation of this condition indicates an inconsistency in the scheduler, in which case we change the action in one of the states to match the others. We take a greedy approach and always adapt to the action chosen most often across all IMDP states $(s, t) \in \tilde{Q}$ for the same time t . For example, if $\sigma((s, t)) = \sigma((s', t)) \neq \sigma((s'', t))$, then we only modify $\sigma((s'', t))$ to match the other actions. Because the set $\tilde{Q}_{\text{reach}}^{\sigma^-}$ is finite by construction, a finite number of modifications suffices to render any scheduler consistent.

Obtaining upper bounds | The set of consistent schedulers is finite but prohibitively large, so enumerating over all consistent schedulers is infeasible. For a sound upper bound, we instead optimize over all schedulers. The experiments in Sect. 13.6 show that we obtain (relatively) tight bounds. To further refine these upper bounds, the

²In fact, consistent schedulers are similar to (memoryless) schedulers in partially observable MDPs (POMDPs) that choose the same action in states with the same observation label.

Table 13.1: Overview of considered benchmarks.

Example		CTMC size		State-weight function
Name	Evid. len. ($ \Omega $)	States	Transit.	Specification
INVENT	3-14	3	4	“Prob. empty inventory within time 0.1”
AHRS	4	74	196	“Prob. system failure within time 50”
PHIL	4	34	89	“Prob. deadlock within time 1”
TANDEM	2	120	363	“Prob. both queues full within time 10”
POLLING	3	576	2208	“Prob. all stations empty within time 10”

literature suggests another abstraction refinement loop, which can be formulated either directly on the imprecise evidence [CJJK19] or on the consistent schedulers [WJWJ⁺21]. The latter approach leverages the fact that consistent schedulers can also be modeled as searching for (memoryless) schedulers in partially observable MDPs, where the schedulers would only observe the time but not the state. Finally, the hardness of optimizing over consistent schedulers in the IMDP remains open: Classical NP-hardness results for the problems above do not carry over.

13.6 Numerical Experiments

We implemented our approach in a prototypical Python tool building on top of Storm [HJKQ⁺22] for the analysis of CTMCs and IMDPs. It takes as input a CTMC C , a specification defining the state-weight function w , and imprecisely timed evidence Ω . The tool constructs the abstract IMDP for the coarsest time partition, computing the probability intervals as per Eqs. (13.9) and (13.10). The bounds on the conditional reachability in Lemma 13.19 are computed using robust value iteration. Then, the tool applies guided refinements, as in Sect. 13.4.2, and starts a new iteration with the refined partition. After a predefined time limit, the tool returns the lower bound $\underline{W}(\Omega)$ and upper bound $\overline{W}(\Omega)$ on the conditional reachability $W(\Omega)$:

$$\underline{W}(\Omega) = \min_{\tilde{P} \in \mathcal{P}} W_{\mathcal{M}_\Omega}(\tilde{P}, \hat{\sigma}) \leq W(\Omega) \leq \max_{\sigma \in \mathcal{S}_{\text{stat}}^{\mathcal{M}_\Omega}} \max_{\tilde{P} \in \mathcal{P}} W_{\mathcal{M}_\Omega}(\tilde{P}, \sigma) = \overline{W}(\Omega), \quad (13.11)$$

where the consistent scheduler $\hat{\sigma}$ for the lower bound is obtained by fixing all inconsistencies in the scheduler σ^- defined in Sect. 13.5. The tool can also compute minimal conditional reachabilities (by swapping all min and max operators).

Benchmarks | We evaluate our approach on several CTMCs from the literature. For each CTMC, we create multiple imprecisely timed evidences. Table 13.1 lists the evidence length (i.e., the number of observed times and labels), the number of CTMC states and transitions, and the specification of the state-weight function. More specifically, we perform experiments on the following benchmarks:

1. INVENT is the inventory model from Fig. 13.1a with the label `empty` if the inventory is empty and $\neg\text{empty}$ otherwise. The state-weight function is defined by the probability of reaching an empty inventory within a time bound of 0.1.

fault tree

2. AHRS is a dynamic *fault tree* model of an Active Heat Rejection System [BD05]. The model was taken from the FFORT fault tree collection [RBNS⁺19] and converted into a CTMC using Storm [VJK18]. The evidence is given by observations of the failures of sub-systems and components, for instance, $A1_f$ and $Spare_f$. The state-weight function is given by the probability of system failure within the next 50 time units.
3. PHIL models a variant of the dining philosophers [Dij71] and was taken from the QCOMP benchmark collection [HKPQ⁺19]. As evidence, we can observe for each fork whether it is currently in use ($\neg fork_i$) or available. The state-weight function is given by the probability of reaching a deadlock within 1 time unit.
4. TANDEM models a tandem queuing network consisting of a Coxian distribution with two phases sequentially composed with a M/M/1-queue [HMS99]. As evidence, we observe if any of the two queues is full. The state-weight function is given by the probability that both queues are full within 10 time units.
5. POLLING models a cyclic server polling system [CT92]. One polling server handles six stations, processing their jobs at a given rate. As evidence, we observe whether stations are empty, i.e., have no jobs. The state-weight function is given by the probability that all stations have no jobs within 10 time units.

Reproducibility | All experiments run on an Intel Core i5 with 8GB RAM, using a time limit of 10 minutes. The source code, benchmarks, and logfiles to produce the experimental results are archived at <https://doi.org/10.5281/zenodo.10438984>.

13.6.1 Feasibility

We investigate if our approach yields tight bounds on the weighted reachability. Fig. 13.7 shows the results for each example with different imprecise evidences. The gray area shows the weighted reachabilities (as per Theorem 13.14) for 500 precisely timed instances $\rho \in \Omega$ sampled from the imprecise evidence. Recall that the weighted reachability $W(\Omega)$ is an upper bound to the weighted reachability for each precisely timed evidence $\rho \in \Omega$. Thus, the upper bound of the gray areas in Fig. 13.7, indicated as $W(\Omega)'$, is a lower bound of the actual (but unknown) value $W(\Omega)$. The blue lines are the upper bound $\overline{W(\Omega)}$ (solid) and lower bound $\underline{W(\Omega)}$ (dashed) on $W(\Omega)$ returned by our approach over the runtime (note the log-scale). Similarly, the red lines are the bounds obtained for *minimizing* the minimal weighted reachability.

Tightness of bounds | Fig. 13.7 shows that we obtain reasonably tight bounds within a minute. In all examples, the lower bound converges close to the maximum of the samples. The improvement is steepest at the start, indicating that the bounds can be quickly improved by only a few refinement steps. In the long run, the improvement of the bounds diminishes, both because each refinement takes longer, and the improvement in each iteration gets smaller.

While not clearly visible in Fig. 13.7a, the lower bound $\underline{W(\Omega)}$ (dashed blue line) slightly exceeds the maximal sampled value $W(\Omega)'$ (gray area) in the end. Thus, the lower bound $W(\Omega)$ is closer to the actual weighted reachability $W(\Omega)$ than the maximal lower bound obtained by sampling. We observed the same results when increasing the number of samples used to compute $W(\Omega)'$ to 10 000.

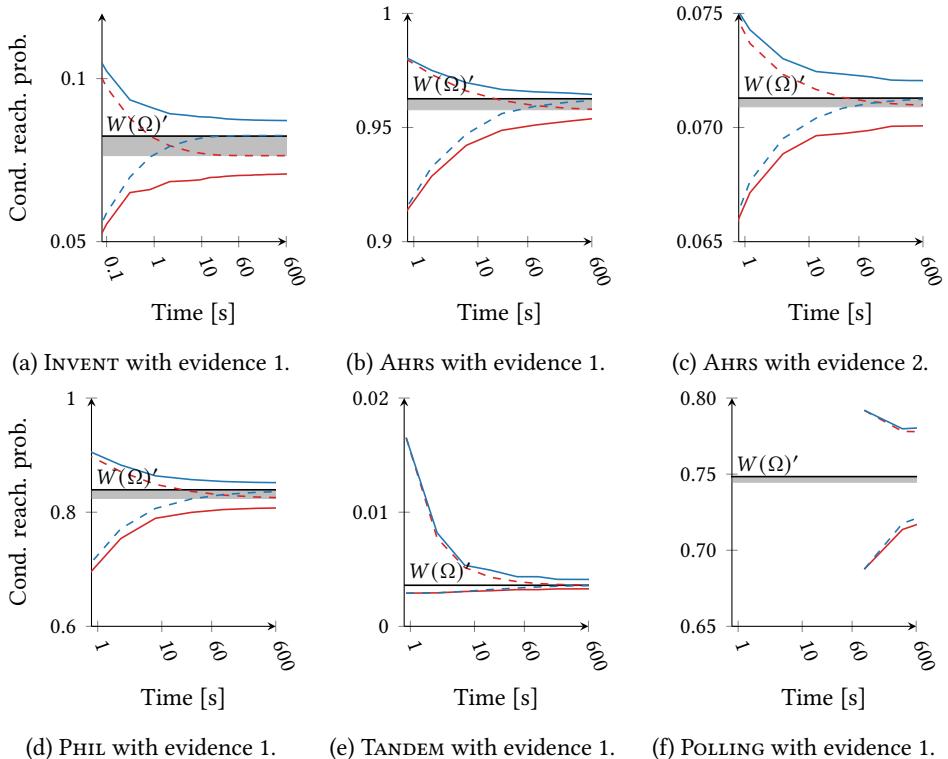


Figure 13.7: Results for different CTMCs and different imprecisely timed evidence. The blue lines are the upper bound $\overline{W(\Omega)}$ (solid) and lower bound $\underline{W(\Omega)}$ (dashed) on $W(\Omega)$; red lines show the analogous lower bounds.

Figs. 13.7b and 13.7c show the general benefit of conditioning on evidence. While evidence 1 for AHRS results in a state in which a system failure within the next 50 time units is very likely, a failure conditioned on evidence 2 is very unlikely.

13.6.2 Scalability

We investigate the scalability of our approach. Table 13.2 provides the refinement statistics, bounds, model sizes, and runtimes for all benchmarks. The refinement statistics show the number of iterations (Iter.) and the total number of splits made in the partition. The bounds on $W(\Omega)$ (which are the solid and dashed blue lines in Fig. 13.7) and the IMDP sizes are both given for the final iteration. For the timings, we provide the total time (over all iterations) and distinguish between the time spent on unfolding the model, i.e., constructing the IMDP and analyzing it. Our approach terminates if, after an iteration, the total run time so far exceeds the time limit of 10 minutes. The total runtime can, therefore, be significantly longer than 10 minutes.

CTMC size | The size of the CTMC has a large impact on the total runtime. For example, for evidence with 4 labels, we can perform up to 27 iterations for INVENT (3 CTMC states) but only 6-8 for AHRS (74 CTMC states). For POLLING (576 states) with

Table 13.2: Results for all benchmarks (evidence length $|\Omega|$ is given after the name).

Example	Refine		Results		IMDP size			Timings [s]		
	Name ($ \Omega $)	Iter.	#split	Bounds on $W(\Omega)$	States	Actions	Transit.	Unfold	Analysis	Total
INVENT-1 (4)	25	555	[0.082536, 0.087138]	898	128307	278163	537.51	100.28	637.81	
INVENT-2 (4)	27	585	[0.071768, 0.078328]	1180	167917	503537	606.91	43.85	650.74	
INVENT-3 (9)	14	1176	[0.071757, 0.078577]	2372	369329	1107877	658.77	127.83	786.57	
INVENT-4 (15)	7	528	[0.070924, 0.080409]	1016	39927	115119	42.63	974.89	1017.50	
AHRS-1 (4)	6	177	[0.962041, 0.964306]	6283	282538	1415346	620.75	179.65	800.39	
AHRS-2 (4)	8	154	[0.071239, 0.072057]	727	20626	81362	577.64	69.19	646.85	
AHRS-3 (4)	6	176	[0.964936, 0.969535]	6112	280954	1334231	749.38	152.61	902.00	
AHRS-4 (4)	7	300	[0.209591, 0.213820]	7179	535763	3618439	1801.81	111.39	1913.18	
PHIL-1 (5)	7	339	[0.836695, 0.851548]	4122	370091	3887339	851.92	60.32	912.23	
PHIL-2 (5)	6	209	[0.236734, 0.246067]	4050	203549	3669721	419.97	376.73	796.70	
TANDEM-1 (2)	9	77	[0.003577, 0.004009]	1203	24561	362657	917.29	3.11	920.42	
TANDEM-2 (2)	7	80	[0.130187, 0.162762]	587	25096	75548	549.03	327.93	876.96	
POLLING-1 (3)	2	9	[0.731410, 0.781912]	3267	9798	2379462	348.83	2603.08	2951.89	

evidence of length 2, performing 2 iterations takes nearly 50 minutes. The CTMC size affects the unfolding, which requires computing the transient probabilities from all states in one layer to all states in the next one. A clear example is TANDEM-1 (120 CTMC states), where nearly all time is spent on the unfolding. Furthermore, a larger CTMC also leads to more transitions in the IMDP and, thus, can increase the analysis time. An example is POLLING-1 (576 CTMC states), where most runtime is spent in the analysis.

Length of evidence | The time per refinement step increases with the length of the evidence. For example, for INVENT-4 (with 15 labels), only 7 iterations are performed because the resulting IMDP has 15 layers, so the value iteration becomes the bottleneck (nearly 96% of the runtime for this example is spent on analyzing the IMDP). This observation is consistent with experiments on unfolded MDPs in [JTS21; HJQW23], where policy iteration-based methods lead to better results.

Caching improves performance | To reduce runtimes, we implemented caching in our tool, which allows reusing transient probability computations. For example, if all labels in the evidence have a time interval of the same width (which is the case for AHRS-1), transient probabilities are the same between layers of the unfolding. Table 13.1 shows that the unfolding times for AHRS-1 are indeed lower than for, e.g., AHRS-3, which has time intervals of different widths.

Likelihood of evidence | The size of the IMDP is influenced by the number of CTMC states corresponding to the observed labels. Less likely observations can, therefore, mean that fewer CTMC states need to be considered in each layer. For example, the evidence in AHRS-2 is 17 times less likely (probability of 0.01, with 569 states) than AHRS-4 (probability of 0.17, with 4007 states), and as a result, the total runtime of AHRS-2 is less than for AHRS-4.

13.7 Related Work

Closest to our problem is work on model checking CTMCs against deterministic timed automaton (DTA) [CHKM11; AD18; FKLX⁺18]. Evidence can be expressed as a single-

clock DTA, and tools such as MC4CSL [AD10] can calculate the weighted reachability for precise timings. However, for imprecisely timed evidence, checking CTMCs against DTAs yields the *sum of probabilities* over all instances of the evidence, whereas we are interested in the *maximal probability* over all instances.

Our setting is also similar to synthesizing timeouts in CTMCs with fixed-delay transitions [BKKN⁺15; KKR16; BDKK⁺19]. Finding optimal timeouts is similar to our objective of finding an instance of the imprecisely timed evidence such that the weighted reachability is maximized. While timeouts can model the time *between* observations, we consider *global* observation times, i.e., the time between observations depends on the previous time of observation—which cannot be modeled with timeouts.

Imprecisely timed evidence can also be expressed via multiphase timed until formulas in continuous-time linear logic [GY22]. However, similar to DTA, conditioning and computing the maximal weighted reachability are not supported.

Conditional probabilities naturally appear in runtime monitoring [SSAB⁺19; BDDF⁺18] and speech recognition [GY07], and is, e.g., studied for hidden Markov models [SBSG⁺11] and MDPs [BKKM14; JTS21]. Approximate model checking of conditional continuous stochastic logic for CTMCs is studied in [GXZZ13; GHZZ13] by means of a product construction formalized as CTMC, but their algorithm is incompatible with imprecise observation times. Conditional sampling in CTMCs is studied by [HS09], and maximum likelihood inference of paths in CTMCs by [Per09].

The abstraction of continuous stochastic models into IMDPs is well-studied [LSAZ22]. Similar to the abstractions we developed in Part II, various papers study abstractions of stochastic hybrid and dynamical systems into IMDPs [CA19] and relate to early work in [JL91]. Our abstraction in Sect. 13.4 is similar to a game-based abstraction, in which the (possibly infinite-state) model is abstracted into a two-player stochastic game [KKNP10; HHWZ10; HNPW⁺11]. In particular, IMDPs are a special case of a stochastic game in which the actions of the second player in each state only differ in transition probabilities [NG05; Iye05]. An interesting extension is to combine our approach with the CTMCs with uncertain transition rates in Chapter 12.

13.8 *Proofs

We provide the proofs of Theorems 13.11 and 13.17 presented earlier in this chapter.

13.8.1 Proof of Theorem 13.11

The proof of Theorem 13.11 is based on Lemma 13.20 below, which states that, for every evidence instance $\rho \in \Omega$ with consistent scheduler $\sigma \sim \rho$, it holds that

$$\Pr^C(\pi(t_d) = s \mid [\pi \models \rho]) = \Pr_{\sigma}^{\mathcal{M}}(q_I \models \diamond(s, t_{\star}) \mid [\xi \models \rho]). \quad (13.12)$$

That is, the conditional transient probability $\Pr^C(\pi(t_d) = s \mid [\pi \models \rho])$ equals the conditional reachability probabilities in Eq. (13.12) for the unfolded MDP \mathcal{M} , under a scheduler $\sigma \sim \rho$ consistent to ρ . We then use Eq. (13.12) to rewrite Problem 13.4 as

$$W(\Omega) = \sup_{\rho \in \Omega} \sum_{s \in S} \Pr_{\sigma}^{\mathcal{M}}(q_I \models \diamond(s, t_{\star}) \mid [\xi \models \rho]) \cdot w(s),$$

*Section with details that can be skipped safely

where $\sigma \sim \rho$, as per Def. 13.8. Due to the one-to-one correspondence between choices $\rho \in \Omega$ and consistent schedulers, we can replace the supremum over $\rho \in \Omega$ by the supremum over consistent schedulers, which yields the expression in Theorem 13.11.

Next, we formalize the lemma that shows Eq. (13.12). In the proof of this lemma, we use the notion of the *state-trace* $s\text{Tr}_\rho(\pi) \in S^d$ of a CTMC path π onto the time points t_1, \dots, t_d of the precisely timed evidence ρ , which is defined as follows:

$$s\text{Tr}_\rho(\pi) = (\pi(t_1), \pi(t_2), \dots, \pi(t_d)). \quad (13.13)$$

Conditional reachability probabilities in the CTMC and in the unfolded MDP are then related as follows.

Lemma 13.20 (CTMC vs. unfolded MDP) For a CTMC C and the imprecise evidence Ω , let $\mathcal{M} = \text{Unfold}(C, \mathcal{G}_\Omega)$ be the unfolded MDP. For every instance $\rho \in \Omega$ with corresponding consistent scheduler $\sigma \in \widehat{\mathfrak{S}}_{\text{stat}}^\mathcal{M}$, i.e., such that $\sigma \sim \rho$, it holds that

$$\Pr^C(\pi(t_d) = s \mid [\pi \models \rho]) = \Pr_\sigma^\mathcal{M}(q_I \models \diamond(s, t_\star) \mid [\xi \models \rho]). \quad (13.14)$$

First, let us use Bayes' rule to rewrite the right-hand side of Eq. (13.14) as

$$\Pr_\sigma^\mathcal{M}(q_I \models \diamond(s, t_\star) \mid [\xi \models \rho]) = \frac{\Pr_\sigma^\mathcal{M}(q_I \models \diamond(s, t_\star) \cap [\xi \models \rho])}{\Pr_\sigma^\mathcal{M}([\xi \models \rho])}. \quad (13.15)$$

We will prove Lemma 13.20 by showing that the numerator and denominator in Eq. (13.15) are equivalent to those in Eq. (13.1). In other words, we will show that

$$\Pr^C([\pi(t_d) = s] \cap [\pi \models \rho]) = \Pr_\sigma^\mathcal{M}(q_I \models \diamond(s, t_\star) \cap [\pi \models \rho]) \quad \forall s \in S \quad (13.16)$$

$$\Pr^C(\pi \models \rho) = \Pr_\sigma^\mathcal{M}([\xi \models \rho]), \quad (13.17)$$

where $\sigma \sim \rho$ are consistent as per Def. 13.8. We prove Eq. (13.17) first and then prove Eq. (13.16) in a largely analogous manner.

Proof of Eq. (13.17) | From Eq. (13.1), we have for every $\rho \in \Omega$ that

$$\Pr^C(\pi \models \rho) = \int_{\Pi^C} \mathbb{1}_{(\pi \models \rho)} \Pr^C(\pi) d\pi = \Pr^C(\pi \in \Pi_\rho^C), \quad (13.18)$$

where $\Pi_\rho^C = \{\pi \in \Pi^C : \pi \models \rho\} \subset \Pi^C$ is the subset of CTMC paths consistent with evidence ρ . Let Γ_ρ^C be the set of state-traces that are consistent with evidence ρ :

$$\Gamma_\rho^C = \bigcup_{\pi \in \Pi^C} \{s\text{Tr}_\rho(\pi) : \pi \models \rho\} \subseteq S^d.$$

Let us denote (x_1, \dots, x_k) by $x_{1:k}$ for brevity. Using this notation, the preimages $s\text{Tr}_\rho^{-1}(s_{1:d})$ for all $s_{1:d} \in \Gamma_\rho^C$ form a partition of Π_ρ^C , that is:

$$\Pi_\rho^C = \bigcup_{s_{1:d} \in \Gamma_\rho^C} s\text{Tr}_\rho^{-1}(s_{1:d}) \quad \text{and} \quad s\text{Tr}_\rho^{-1}(s_{1:d}) \cup s\text{Tr}_\rho^{-1}(s'_{1:d}) = \emptyset \quad \forall s_{1:d}, s'_{1:d} \in \Gamma_\rho^C.$$

Thus, we can rewrite Eq. (13.18) as a finite sum over all state-traces $s_{1:d} \in \Gamma_\rho^C$:

$$\Pr^C(\pi \models \rho) = \sum_{s_{1:d} \in \Gamma_\rho^C} \Pr^C(\pi \in \Pi^C : \text{sTr}_\rho(\pi) = s_{1:d}).$$

The term $\Pr^C(\pi \in \Pi^C : \text{sTr}_\rho(\pi) = s_{1:d})$ is the probability for a path π whose state-trace is $s_{1:d}$. This probability is equal to the product of the appropriate transient probabilities $\mathbf{P}_{\delta_{s_{i-1}}}(t_i - t_{i-1})(s_i)$ for all $s \in \{1, \dots, d\}$, as defined in Sect. 11.2:

$$\Pr^C(\pi \models \rho) = \sum_{s_{1:d} \in \Gamma_\rho^C} \prod_{i=1}^d \mathbf{P}_{\delta_{s_{i-1}}}(t_i - t_{i-1})(s_i)$$

where $s_0 = s_I$ and $t_0 = 0$. Recall from Def. 13.7 that the unfolded MDP has transition probabilities $P((s, t), t', (s', t')) = \mathbf{P}_{\delta_s}(t' - t)(s')$. Hence, we obtain

$$\begin{aligned} \Pr^C(\pi \models \rho) &= \sum_{s_{1:d} \in \Gamma_\rho^C} \prod_{i=1}^d P((s_{i-1}, t_{i-1}), t_i, (s_i, t_i)) \\ &= \sum_{s_{1:d} \in \Gamma_\rho^C} \Pr_\sigma^M(\xi \in \Pi^M : \xi = (s_I, 0), (s_1, t_1), \dots, (s_d, t_d)). \end{aligned} \tag{13.19}$$

A state-trace $s_{1:d}$ belongs to the set of state-traces Γ_ρ^C consistent with evidence ρ if and only if the associated MDP path $\xi = (s_I, 0), (s_1, t_1), \dots, (s_d, t_d) \in \Pi^M$ is consistent with ρ , i.e., $\xi \models \rho$. Thus, we can rewrite Eq. (13.19) as the desired expression:

$$\begin{aligned} \Pr^C(\pi \models \rho) &= \sum_{\xi \in \Pi^M} \Pr_\sigma^M(\xi) \cdot \mathbb{1}_{(\xi \models \rho)} \\ &= \Pr_\sigma^M(\xi \models \rho). \end{aligned}$$

Proof of Eq. (13.16) | Again, using the fact that the preimages $\text{sTr}^{-1}(s_{1:d})$ for all $s_{1:d} \in \Gamma_\rho^C$ form a partition of Π_ρ^C (where sTr is defined by Eq. (13.13)), we obtain

$$\begin{aligned} \Pr^C([\pi(t_d) = s] \cap [\pi \models \rho]) &= \\ &\sum_{s_{1:d} \in \Gamma_\rho^C} \Pr^C(\pi \in \Pi^C : [\pi(t_d) = s] \cap [\text{sTr}_\rho(\pi) = s_{1:d}]). \end{aligned}$$

Compared to Eq. (13.16), we additionally require that $\pi(t_d) = s$, which corresponds with reaching the terminal state $(s, t_\star) \in Q$ in the unfolded MDP M corresponding with

CTMC state $s \in S$. As a result, we have that

$$\begin{aligned} \Pr^C([\pi(t_d) = s] \cap [\pi \models \rho]) &= \sum_{s_{1:d} \in \Gamma_\rho^C} \prod_{i=1}^d P((s_{i-1}, t_{i-1}), t_i, (s_i, t_i)) \cdot \mathbb{1}_{(s_d=s)} \\ &= \sum_{s_{1:d} \in \Gamma_\rho^C} \Pr_\sigma^M((s_1, 0), (s_1, t_1), \dots, (s_d, t_d)) \cdot \mathbb{1}_{(s_d=s)} \\ &= \sum_{\xi \in \Pi^M} \Pr_\sigma^M(\xi) \cdot \mathbb{1}_{(\xi \models \rho)} \cdot \mathbb{1}_{(\xi \models \diamond(s, t_\star))} \\ &= \Pr_\sigma^M(q_I \models \diamond(s, t_\star) \cap [\pi \models \rho]), \end{aligned}$$

which is the desired expression in Eq. (13.16), so we conclude the proof.

13.8.2 Proof of Theorem 13.17

Let $H: Q \rightarrow \tilde{Q}$ be a function that maps every state of MDP $\mathcal{M}|_\Omega$ to a state of IMDP \mathcal{M}_I , such that $H((s, t)) = (s, \mathcal{T}) \in \tilde{Q}$, where $t \in \mathcal{T}$. The mapping H is well-defined as \tilde{Q} represents a proper partition of Q . We prove Theorem 13.17 by showing that for every MDP state $(s, t) \in Q$, the corresponding IMDP state $H((s, t)) = (s, \mathcal{T}) \in \tilde{Q}$ overapproximates its behavior. Formally, for the conditioned MDP, take any transition from state $(s, t) \in Q$ via (enabled) action $t' \in A((s, t))$ to state $(s', t') \in Q$. For any such transition, there exists an IMDP transition $(s, \mathcal{T}) \in \tilde{Q}$ via $\mathcal{T}' \in A((s, \mathcal{T}))$ to state $(s', \mathcal{T}') \in \tilde{Q}$ such that

1. there exists $\tilde{P} \in \mathcal{P}$ such that $P((s, t), t', (s', t')) = \tilde{P}((s, \mathcal{T}), \mathcal{T}', (s', \mathcal{T}'))$, and
2. it holds that $H((s, t)) = (s, \mathcal{T}) \in \tilde{Q}$ and $H((s', t')) = (s', \mathcal{T}') \in \tilde{Q}$.

Observe that the converse also holds: For any IMDP transition, there exists a corresponding MDP transition such that the conditions above hold. These conditions formalize that there always exists a transition function $\tilde{P} \in \mathcal{P}$ such that the induced MDP $\mathcal{M}_I[\tilde{P}]$ is a *probabilistic bisimulation* of the conditioned MDP $\mathcal{M}|_\Omega$, similar as in [JL91]. Hence, there exists a $\tilde{P} \in \mathcal{P}$ such that

$$\max_{\sigma \in \widehat{\mathcal{S}}_{\text{stat}}} W_{\mathcal{M}_I}(\tilde{P}, \sigma) = W(\Omega),$$

leading to the upper and lower bounds in Eq. (13.8). Thus, we conclude the proof.

13.9 Discussion

We close this chapter with a brief discussion of open directions and questions about the method presented in this chapter. In particular, a natural next step is to embed our method in a predictive runtime monitoring framework, which introduces the challenge of running our algorithm in real-time. Toward this goal, we identify the following concrete challenges in further improving the applicability and scalability of the method presented in this chapter.

1. In our IMDP abstraction, we used methods to overapproximate the union over MDP probabilities in Def. 13.16. As a result, the performance of our method may

be improved significantly by finding better methods to overapproximate these probabilities. In particular, doing so may lead to tighter bounds on the weighted reachability in Problem 13.4.

2. Recall that our theoretical results require optimizing over the subset of consistent schedulers only. Since we are not aware of any efficient algorithm to achieve this, we chose to optimize over all schedulers instead. As a result, the tightness of our bounds presented in the numerical experiments in Sect. 13.6 may be improved by optimizing over the consistent schedulers only. One promising direction to optimize over consistent schedulers only is to leverage techniques recently developed by [ACJK⁺21] for the synthesis of probabilistic programs.
3. The refinement strategies that we used for the IMDP are very simple and potentially suboptimal. As such, investigating more sophisticated refinement strategies that only refine parts of the abstraction where beneficial could lead to significant reductions in the complexity of our approach.
4. Finally, the computational performance of our implementation can still be improved significantly. One promising option to improve performance is to adapt symbolic policy iteration [BDKK⁺19], which only considers small sets of candidate actions instead of all actions.

Summary

- ⇒ We have presented a method for computing reachability probabilities in CTMCs, conditioned on evidence with imprecise observation times.
- ⇒ We have shown that these conditional probabilities can be computed as unconditional probabilities on an unfolded MDP.
- ⇒ Since the unfolded MDP generally has infinitely many states and actions, we have presented an effective and iterative abstraction into a finite IMDP.

Part V

Outlook

14 Tool Support

We give a brief overview of the prototypical tools we have developed as part of this thesis. All of these tools are open-source and available as Zenodo archive and/or Git repository. For most tools, we also provide a Docker container that can be used to reproduce the experiments from the respective chapters.

14.1 Probabilistic model checkers

As discussed in Chapter 1, we prominently use probabilistic model checkers to analyze (finite-state) Markov models. In particular, the two model checkers that we use are PRISM¹ and Storm.² Both PRISM and Storm are open-source tools for the verification of various Markov models, such as discrete-time Markov chains (DTMCs), Markov decision processes (MDPs), and continuous-time Markov chains (CTMCs). Both tools also support parametric and interval DTMCs and MDPs. PRISM and Storm can solve quantitative verification problems for these Markov models against specifications expressed in a variety of temporal logics, including probabilistic computation tree logic (PCTL) and continuous stochastic logic (CSL). PRISM is implemented in Java, whereas Storm is implemented in C++. Python bindings for Storm exist through the Python package Stormpy.³ For more details about the functionalities and capabilities of both tools, we refer to the respective documentation on the websites.

14.2 DynAbs

Our abstraction methods for discrete-time stochastic systems (DTSSs) presented in Part II are implemented in a prototypical Python tool called DynAbs. The source code is available on GitHub at <https://github.com/LAVA-LAB/DynAbs> and on Zenodo at <https://doi.org/10.5281/zenodo.13348782>. The benchmark files provided with DynAbs can be used to set up new benchmarks with different DTSSs and reach-avoid specifications. DynAbs supports reach-avoid specifications over finite and infinite horizons, and stochastic linear dynamical systems with set-bounded parameter uncertainty.

Given a DTSS (possibly with set-bounded parameter uncertainty) and a reach-avoid specification, DynAbs generates an interval Markov decision process (IMDP) abstraction, roughly by performing the following steps:

1. Create a partition of the state space of the DTSS as the discrete IMDP states;
2. Define a set of target points on the state space of the DTSS, each of which corresponds to an IMDP action;
3. Determine which IMDP actions are enabled in which IMDP states;

¹<https://prismmodelchecker.org/>

²<https://stormchecker.org/>

³<https://moves-rwth.github.io/stormpy/>

4. Compute the transition probability intervals for each transition;
5. Export the IMDP to the input format used by PRISM.⁴

For a more detailed description of each step, we refer to the algorithms presented in Chapters 6 and 7, and to the documentation of DynAbs. For detailed numerical experiments and benchmarking, we also refer to Chapters 6 and 7.

DynAbs uses PRISM for model checking the IMDP abstractions it generates. Since PRISM is written in Java and DynAbs in Python, communication between the tools is currently quite inefficient. In particular, DynAbs first has to export the IMDPs in the PRISM language and then call PRISM to perform the model checking. Thereafter, PRISM writes the results into text files, which are loaded back into DynAbs. We remark that this exporting and importing of files can be a significant overhead, and thus, our implementation can still be improved significantly. For example, Storm with its Python bindings in Stormpy could be used instead of PRISM, which would avoid the need for writing and reading (potentially very large) model files. However, when we started developing DynAbs, Storm did not yet support IMDPs, and hence, we built on top of PRISM instead.

Another aspect to improve the scalability of DynAbs is by parallelizing the abstraction generation procedure. Most of the computations in generating the IMDP abstraction involve matrix-vector multiplications (for computing backward reachable sets), solving linear inequalities (for computing transition probability intervals), and other linear algebra operations. Such operations can be parallelized efficiently, and thus, the abstraction generation can be sped up significantly by running the code on a GPU cluster. Parallelized construction of finite-state abstractions has previously been investigated and implemented by the tools AMYTISS [LKSZ20] and IMPaCT [WL24]. We leave extensions of DynAbs with parallelization and GPU support for future work.

14.3 Scenario Approach for pMDPs

We implemented our method presented in Chapter 9 for sampling-based verification of parametric Markov decision processes (pMDPs) with a distribution over the parameters in Python. Our implementation is available on Zenodo at <https://doi.org/10.5281/zenodo.6674059>. This archive contains a Docker container to reproduce the experiments from Chapter 9, as well as the Python source code. We provide a Docker container to reproduce the experiments from Chapter 9.

Our implementation builds on top of Storm to verify MDPs induced by the sampled parameter values. While the derivation of our method from Chapter 9 relies on formulating scenario optimization problems, we do not need to solve these optimization problems explicitly in practice. Instead, as discussed in Chapter 9, the solutions to these scenario optimization problems can be derived analytically. As a result, the implementation of our method does not require a convex optimization solver, and the main computational costs lie in the model checking using Storm.

⁴More specifically, we export the IMDPs in *explicit PRISM format*, which stores the states, labels, and transitions as explicit lists in three separate files. For more details, we refer to the PRISM documentation on <https://prismmodelchecker.org/manual/Appendices/ExplicitModelFiles>.

14.4 Differentiation of pRMCs

A Python implementation of our methods from Chapter 10 for sensitivity analysis of parametric (robust) Markov chains is available on GitHub at <https://github.com/LAVA-LAB/prmc-sensitivity>. A Docker container to directly reproduce the experimental results presented in Chapter 10 is available as a Zenodo archive at <https://doi.org/10.5281/zenodo.7864260>.

We (only) use Storm to parse the benchmark models but not for the verification of these models. The computation of partial derivatives involves solving (potentially large) linear systems of equations and convex optimization problems. We use the SciPy sparse solver to solve equation systems, and we use Gurobi and its Python API via the package GurobiPy⁵ for solving linear programs. While open-source optimization packages such as CVXPY⁶ can also be used to solve these linear programs, we experienced better performance when using GurobiPy directly.

One main restriction of our implementation is that we cannot handle cases where solutions are not unique, because derivatives for these models are not well-defined. As a practical workaround for such corner cases, we slightly perturb the parameter values to ensure that derivatives are well-defined. However, this workaround is still prone to numerical stability issues. Such numerical issues could be prevented by using exact arithmetic for solving linear equation systems; however, this would restrict our implementation to small models only.

14.5 SLURF

In Chapter 12, we presented a sampling-based verification method for parametric CTMCs (pCTMCs) with a distribution over the parameters. We implemented this method in a Python tool called SLURF, which is available on GitHub at <https://github.com/LAVA-LAB/slurf> and on Zenodo at <https://doi.org/10.5281/zenodo.6523863>. The Zenodo archive contains a Docker container to directly reproduce our numerical experiments. Roughly, given a pCTMC, a distribution over the parameters, and a set of measures, SLURF performs the following steps:

1. Sample a set of $N > 0$ values for the parameters of the pCTMC;
2. For each of the N parameter samples, build the induced CTMC and compute the corresponding solution vector;
3. Construct and solve the scenario optimization problem for the obtained N solution vectors;
4. Return the probably approximately correct (PAC) verification result using the theoretical results from Chapter 12.

As discussed in Chapter 12, SLURF also supports *imprecise* solution vectors, which are obtained by only constructing partial CTMC models. Using imprecise solution vectors can significantly reduce the computational costs of verifying CTMCs, at the cost of obtaining more conservative results.

SLURF uses Storm with its Python bindings via Stormpy to verify CTMCs, and the

⁵<https://pypi.org/project/gurobipy/>

⁶<https://cvxpy.org/>

convex optimization package CVXPY to solve scenario optimization problems. CVXPY can be used with a variety of solvers, such as the open-source solvers ECOS and OSQP, but also commercial solvers such as Gurobi (which we used in our experiments). By building on top of Storm, SLURF can handle both CTMCs and dynamic fault trees (recall that fault trees can be represented as a CTMC), which can be given in either PRISM [KNP11] or JANI [BDHH⁺17] format.

14.6 Conditional Reachability in CTMCs

Finally, in Chapter 13, we presented a method for computing conditional reachability probabilities in CTMCs. A Python implementation of this method is available as a Zenodo archive at <https://doi.org/10.5281/zenodo.10438984>. This Zenodo archive contains a Docker container to reproduce the experiments from Chapter 13, as well as the Python source code. As discussed in more detail in Sect. 13.6, our implementation builds on top of Storm for the analysis of CTMCs and IMDPs. For several concrete challenges in further improving the applicability and scalability of our implementation, we refer back to Sect. 13.9.

15 Conclusion and Future Work

In this thesis, we have studied quantitative verification problems for Markov models with uncertainty. We have developed verification methods that can be used to provide rigorous guarantees that are robust against the model's uncertainty. In this final chapter, we reflect on the challenges posed in the introduction in Chapter 1: Have our contributions addressed these challenges, and to what extent? Thereafter, we zoom out and condense our contributions into a set of guidelines for the verification of Markov models under uncertainty. Based on the contributions and limitations of our work, we conclude the thesis by discussing several directions for future research.

15.1 Summary of Contributions

In Chapter 1, we formulated four challenges that we aimed to address in this thesis:

1. Robust policy synthesis for uncertain Markov models with continuous state and action spaces;
2. Data-driven verification of uncertain Markov models with prior knowledge;
3. Robust verification for continuous-time Markov chains with uncertainty;
4. Tool support for analyzing Markov models with uncertainty.

We discussed challenges 1–3 in Part II, III, and IV of this thesis, respectively. Furthermore, challenge 4 was covered throughout the whole thesis, and in particular in Chapter 14. We now discuss to what extent our contributions have addressed these challenges.

Challenge 1 | We start by reflecting on Part II, where we studied robust policy synthesis for discrete-time stochastic systems (DTSSs). These DTSSs are indeed Markov models with continuous state and action spaces. Furthermore, we considered two prime sources of uncertainty for these models: (1) stochastic uncertainty modeled as noise affecting the state transitions, and (2) set-bounded uncertainty in the model parameters. For both settings, we have developed a tractable abstraction algorithm, which can be used to synthesize robust policies that are guaranteed to satisfy a reach-avoid specification with at least a certain probability. We formalized abstractions as interval Markov decision processes (IMDPs) and showed the effectiveness of our approach on several benchmarks. Our abstraction algorithm in Chapter 6 is, in first instance, exact and thus provides *formal guarantees* on the synthesized policies. However, we thereafter turned to a sample-based method for computing intervals on the transition probabilities, thus leading to *statistical guarantees* instead.

While our contributions cover the first challenge, we must also acknowledge the inherent limitations of our approach. In particular, abstractions are computationally expensive to generate and their size increases exponentially with the dimension of the state space. We discuss ideas for mitigating these limitations in Sect. 15.3.

Challenge 2 | Second, in Part III, we investigated parametric Markov decision processes (pMDPs) as models for encoding prior knowledge about transition probabilities. In particular, we have seen that pMDPs allow for modeling dependencies between different probabilities of the model.

In Chapter 9, we considered a setting where, in addition to a pMDP, we have access to a probability distribution over the parameter space. We particularly used this distribution to encode prior knowledge about the parameters.

In Chapter 10, we proposed to use partial derivatives of the polynomial function that describes the pMDP (i.e., the solution function), as a measure of sensitivity for each of the parameters. We showed that these partial derivatives can be used to improve the sample efficiency of an iterative learning scheme.

Both of these methods are data-driven and lead to robust verification results with *statistical guarantees*. While our contributions cover the second challenge, our methods only apply to very particular settings. Especially the setting in Chapter 10 is quite restrictive because we only considered models without nondeterministic action choices, namely parametric (robust) Markov chains. As we have discussed in Sect. 10.8, extensions to the more general setting of pMDPs are interesting yet challenging.

Challenge 3 | Third, in Part IV, we shifted focus to models evolving in continuous time and studied continuous-time Markov chains (CTMCs) and their parametric extension. Again, we considered two sources of uncertainty for these models.

In Chapter 12, we investigated parametric CTMCs (pCTMCs) together with a probability distribution over the parameter space (similar to the setting in Chapter 9). We used data-driven methods from the scenario approach to obtain *statistical guarantees* on the model’s behavior when sampling a random value for the parameters from the distribution and plugging it into the model.

In Chapter 13, we considered a setting where the uncertainty is not part of the CTMC but instead arises from imprecise observations of the model’s state. We proposed a method to compute reachability probabilities (and other measures) for the CTMC, conditioned on a set of imprecisely timed observations. These imprecisely timed observations give rise to a *set of belief distributions* over the states. With our method, we can reason *robustly* over this set of belief distributions.

Thus, we can conclude that our contributions in Part IV have indeed covered the third challenge for two particular problem settings.

Challenge 4 | The final challenge we identified was to develop (prototypical) tool support for analyzing Markov models with uncertainty. As discussed in Chapter 14, we have developed several prototypical tools that implement the algorithms presented in the respective chapters. Within these tools, we make use of the probabilistic model checkers Storm [HJKQ⁺22] and PRISM [KNP11], as well as tools for convex optimization, such as Gurobi [Gur23]. Our main goal was to provide fellow researchers with a basis for further research, and we believe that we have achieved this goal. On the other hand, our goal was not to develop full-fledged tools with, for example, graphical user interfaces. While such more mature tools could be instrumental in the adoption of our methods in practice, developing these tools would be very time-consuming and requires a different set of skills. Thus, we defer the development of more mature tools to future work, as we discuss in Sect. 15.3.

15.2 A Guide to Robust Verification Under Uncertainty

We now zoom out from the technical contributions of this thesis. Instead, we provide a set of five guiding questions that practitioners can ask when aiming to verify a Markov model with uncertainty. While not an exhaustive list, we hope that these guidelines can help practitioners navigate some of the important questions that one should ask when facing a verification problem (or a more general decision-making problem) with uncertainty. We remark that the order of the questions is not fixed, and the importance of each question may vary depending on the application.

1. Do you actually need guarantees?

The first and perhaps most important question to ask is whether you need guarantees at all. In this thesis, we essentially skipped over this question, by focusing on applications where safety and/or performance guarantees are essential. However, in many applications, it might be sufficient to have a “*good enough*” solution, which performs well in practice but does not carry any guarantees. One argument is that the guarantees we can obtain are only as good as the models of the system, the uncertainty, and the system’s requirements. If too little is known about these models, then using a robust approach may lead to overly conservative or even trivial solutions. In such a case, there is little hope of obtaining any meaningful guarantees, so it might be better to forget about formal or statistical guarantees and instead aim to find a solution that empirically works well in practice.

2. What do you know about the model and its uncertainty?

The second question aims to decide the type of model one can use to represent the system and its uncertainty. In particular, one can ask questions such as:

1. Are continuous states and/or actions needed to faithfully represent the system?
2. Do you need a continuous-time model, or does it suffice to make predictions over discrete time steps?
3. What prior knowledge do you have about the uncertainty? (e.g., do you have likelihoods of outcomes, or do you only know the set of possible values?)

If a continuous model is needed to capture the system dynamics, then one should consider modeling the system as a DTSS and using our methods from Part II. Similarly, if a continuous-time model is needed, then one can look into CTMCs and our methods from Part IV. The combination of continuous-time and continuous-state/action models can be modeled as a continuous-time Markov decision processes (CTMDPs) [Kat16; GH09; BHKH05], which are not covered by this thesis. We emphasize, however, that verifying CTMDP is extremely challenging, and existing algorithms resort to discretization to obtain approximate optimal schedulers in practice [Kat16; BFKK⁺13]. Finally, depending on the prior knowledge about the uncertainty, one can either look into set-based or stochastic uncertainty models. In particular, if a probability distribution over the values of the uncertain variable can be derived, then one can use our sampling-based methods from Chapter 9 (for pMDPs) or Chapter 12 (for pCTMCs).

3. What are the system requirements?

In this thesis, we focused on requirements that can be expressed as *temporal logic specifications*. For Markov decision processes (MDPs), we considered *probabilistic computation tree logic (PCTL)*, while for CTMCs, we expressed requirements in *continuous stochastic logic (CSL)*. Furthermore, we also expressed objectives in terms of *expected cumulative rewards*. However, in some applications, the requirements are much more simple. For example, many practical control problems simply require stabilizing a system around a given point. For such simpler requirements, the methods presented in this thesis could be overkill and thus be unnecessarily complex. Before selecting a suitable verification method, it is, therefore, important to ask what the system requirements are and in what formalism you can model these requirements.

4. What type of guarantees do you need?

Besides asking whether you need guarantees at all, another important question is what type of guarantees you need. Do you need *formal* (probabilistic) guarantees, or do *statistical* guarantees suffice? Generally, sampling-based methods such as the scenario approach (as in Chapters 9 and 12) can provide statistical guarantees (i.e., guarantees that hold with a certain confidence probability). On the other hand, methods such as abstraction (as in Chapters 6 and 13) and robust optimization (as in Chapter 10) can often provide formal guarantees. Formal guarantees are often more conservative than statistical guarantees and can be more computationally expensive to obtain. At the same time, statistical guarantees leave a nonzero probability that the system actually does not meet the requirements. Thus, statistical guarantees can lead to a significantly higher variance in the performance of a system than formal guarantees.

5. Does your algorithm need to run online or offline?

Whether you need to apply an algorithm *online* or *offline* has great consequences for the type of methods that are applicable. In particular, online methods (such as for runtime monitoring or low-level control of an autonomous drone) are typically required to run at a high frequency and thus need to be computationally inexpensive. On the other hand, offline methods can usually be more computationally expensive because they are applied *before* running a system. Generally, abstraction-based methods (such as our abstractions of DTSSs) are computationally expensive and thus mainly suited for offline applications. On the other hand, sample-based methods (such as the derivative-guided learning scheme from Chapter 10) are often less expensive and thus better suited for online applications. Thus, it is important to decide the allowed computational expenses before choosing a suitable verification methodology.

15.3 Limitations, Challenges, and Perspectives

While the contributions of this thesis have covered the four challenges we posed, our methods certainly have significant limitations as well. In this section, we highlight some of these limitations and discuss potential opportunities for future work.

15.3.1 Combining learning and verification

The abstraction-based methods presented in this thesis (and most abstraction methods in general) can be computationally expensive. At the same time, these abstraction methods are often quite naïve. Consider, for example, again our IMDP abstraction algorithm for DTSSs, where we partitioned the state space and tried to model as much of the information from the DTSS in the abstraction. However, the synthesized Markov policy we obtain will (with very high probability) typically only visit a *fraction of the states* of the IMDP abstraction. This observation raises the natural question: “*Can we use learning to narrow the search space we cover with the abstraction?*”

Such combinations of learning and verification are not novel in general, such as done in *counterexample-guided inductive synthesis (CEGIS)*, also known as learner-verifier approaches [AAEG⁺21; AEG22; CGJL⁺03; LZCH22; ZLHC23b]. Furthermore, the use of machine learning algorithms to compute upper and lower bounds on reachability probabilities in MDPs is studied by [BCCF⁺14; BCCF⁺24], and statistical model checking with deep neural networks is explored by [GHHK⁺20].

To the best of our knowledge, the combination of learning and verification has not been studied in the scope of abstractions for stochastic dynamical control systems. Our abstraction techniques could naturally fit in an iterative framework, where a *learner* tries to determine what part of the DTSS state/action spaces to focus on, and where a *verifier* uses our techniques to generate an abstraction of those parts only. The main challenge in such an approach lies in the interplay/interaction between the learner and verifier. For example, if the verifier is not able to compute a Markov policy that solves the problem at hand, what information (such as counterexamples or other diagnostic information) should the verifier give back to the learner? We see such a combination of learning and verification as a very promising direction for future research.

15.3.2 Integration with reinforcement learning

Several of the methods we proposed in this thesis involve reasoning over models that are learned from data obtained by interacting with an (unknown) environment. As such, these methods are closely related to *model-based reinforcement learning (RL)* [MBPJ23; PN17; SLO22]. Following the terminology from [MBPJ23], the goal in model-based RL is to learn a global solution to a problem (e.g., an optimal policy for an MDP) based on either a learned or known model of the environment.¹ In this context, several of our methods can be seen as *model-based reinforcement learning with a partially known model*. For example, in Chapter 7, we assumed a model given in the form of a DTSS (known part of the model) with parameters that are only known up to a given set (unknown part of the model), which we called a robust DTSS (RDTSS). Similarly, in Chapters 9 and 12, we respectively assumed that we are given a pMDP and pCTMC (known part of the model), where the parameters are described by a fixed but unknown probability distribution (unknown part of the model). Despite these similarities, our methods differ from standard model-based RL in at least two conceptual aspects.

First, in this thesis, we have approached uncertainty in models from a relatively

¹Whether learning an optimal policy on a known model (e.g., an MDP) should be considered as model-based RL is a matter of debate. In our opinion, since such an optimal policy may be computed using linear programming, we would not regard this class of problems to involve *learning* at all. However, a proper discussion of this terminology is beyond our scope, so we follow [MBPJ23] for simplicity.

static perspective, often assuming that a model or representation of the uncertainty is available as a starting point. Going back to the examples above, we did not update our belief of the uncertain parameters of a RDTSS, and we did not update the probability distributions over the parameters of a pMDP or pCTMC. By contrast, model-based RL commonly takes a more *dynamic* perspective to uncertainty, where the obtained policy is also used to further improve the model (and vice versa) [SB98]. From this perspective, we may argue that most of our techniques are *one-shot approaches*, where we aim to provide a solution to a verification problem under uncertainty, but where we did not update the model or the uncertainty representation based on this solution. That said, this restriction is not inherent to our methods, and a natural next step would be to integrate our methods in a more dynamic uncertainty setting, more similar to model-based RL.

Second, while techniques from (model-based) RL have shown impressive performance on a wide range of complex control tasks, these methods often lack the ability to provide formal guarantees on the obtained solutions. By contrast, our methods are designed to provide such rigorous guarantees, at the cost of being more computationally expensive and requiring more stringent assumptions. In response to the lack of guarantees in RL, the field of learning-based control with formal guarantees has recently received much attention [GF15; AOSC⁺16; BTSK17]. One particularly relevant technique is that of *shielding* in reinforcement learning [ABEK⁺18; CJT23], where the idea is to constrain the exploration of an RL agent to a subset of safe actions generated by a reactive monitor (called a shield). We believe that several of our techniques may naturally fit in such a reactive monitoring framework, which can be used to provide formal guarantees in learning-based settings.

15.3.3 Partial observability

With the exception of Chapter 13, we have assumed that all models are *fully observable*. That is, we as decision-makers, or the scheduler/policy, have access to the *exact* state of the model. This assumption is often unrealistic in practice. For example, an autonomous car navigating in the real world does not know its *exact* position on the map; instead, it uses sensor measurements to *estimate* its state [TBF05; NPCN⁺21]. Filtering techniques such as the Kalman filter [Kal60; WB01] or the particle filter [TBF05; VT24; KFM04] can be used to compute these state estimates.

Incorporating such partial knowledge of states led to the development of the partially observable Markov decision process (POMDP) [CKL94; KLC98]. However, because several problems for POMDPs are proven to be undecidable [MHC99; CCT16], research has largely focused on approximate methods [PGT03; SV05; SPK13; YSHL17; KOA17]. An orthogonal research direction focuses on over/under-approximating values in POMDPs [Lov91; BG09; BJKQ20; BKQ22]. We believe that incorporating partial observability in our verification methods, while retaining the rigorous guarantees we provide, can be an interesting yet challenging avenue for further research.

15.3.4 Exploiting structure in AI

In Chapter 10, we have seen how knowledge about the underlying parametric structure of a model can be used to improve the data efficiency of learning methods. In a much broader context, we strongly believe that using such *structural information* can immensely improve the efficiency of AI methods. This can involve information about the

model's structure (as we did in Chapter 10), but also about the underlying uncertainty or the system requirements. For example, such structure in the model and its policy are made explicit in factored and decentralized (PO)MDPs [GKP01; GKP03; OSV08; OA16].

Structure in the system requirements can also be exploited. The focus in reinforcement learning has recently shifted from considering expected reward objectives only, to considering objectives as *reward machines* [IKVM18; GB20; XGAM⁺20; JBBA21; IKVM22]. These reward machines allow for expressing temporal logic specifications, which thus encode a *logical structure* of the underlying task at hand.

We believe that exploiting structure is needed to increase the (data-)efficiency of AI methods. While we have not focused much on such approaches in this thesis, we still see this as an interesting direction with high potential for future research.

15.3.5 Continuous-time models with nondeterminism

We have studied models with nondeterministic action choices, as well as models that evolve in continuous time. However, we have not investigated their combination, which creates a continuous-time Markov decision process (CTMDP) [Kat16; GH09; BHKH05]. Schedulers for CTMDPs not only decide which action to play based on the states previously visited but also on the elapsed time in every state [NSK09; Mil68]. As a result, a CTMDP has uncountable many deterministic schedulers (compared to countably many for MDPs; see Chapter 3 for details), and existing algorithms resort to discretization to obtain ϵ -optimal schedulers in practice [Kat16]. Thus, developing more sophisticated algorithms for CTMDPs poses an interesting direction for future work.

15.3.6 Mature tool support

As we already discussed, our focus in this thesis was on developing prototypical tools that can be used and extended by fellow researchers. However, to make an actual impact in the real world, more accessible, mature, and user-friendly software is needed. As an example, the model checker PRISM started out as a research tool that implemented algorithms for the analysis of Markov models [KNP02]. Nowadays, the PRISM website states that “*the tool is described in over 850 research papers and has been used in industrial projects and labs.*” We see the development of such mature tool support as a key factor that determines the societal impact of research. Hence, turning our prototypical implementations into more mature software is an important further research direction.

15.4 Final Remarks

As a final note, we expect that rigorous reasoning about uncertainty will become increasingly important in the field of artificial intelligence. As artificial intelligence models get bigger and bigger, it becomes more and more difficult to verify that their behavior is reliable, safe, and fair. We believe that understanding when and how we can verify these models while obtaining rigorous and robust guarantees remains an important field of research. Thus, much future research remains to be done. We hope that this thesis can be a valuable source for future researchers who aspire to solve the intriguing challenges of robust verification under uncertainty.

Part VI

Back Matter

A Bibliography

- [AAEG⁺21] A. Abate, D. Ahmed, A. Edwards, M. Giacobbe and A. Peruffo. ‘*FOSSIL: a software tool for the formal synthesis of lyapunov functions and barrier certificates using neural networks*’. *HSCC*. ACM, 2021, 24:1–24:11. doi: [10.1145/3447928.3456646](https://doi.org/10.1145/3447928.3456646).
- [AB12] H. Andersson and T. Britton. ‘*Stochastic epidemic models and their statistical analysis*’. Volume 151. Springer Science & Business Media, 2012. doi: [10.1007/978-1-4612-1158-7](https://doi.org/10.1007/978-1-4612-1158-7).
- [ABCK⁺18] S. Arming, E. Bartocci, K. Chatterjee, J. Katoen and A. Sokolova. ‘*Parameter-Independent Strategies for pMDPs via POMDPs*’. *QEST*. Volume 11024. Lecture Notes in Computer Science. Springer, 2018, pages 53–70. doi: [10.1007/978-3-319-99154-2_4](https://doi.org/10.1007/978-3-319-99154-2_4).
- [ABEK⁺18] M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer, S. Niekum and U. Topcu. ‘*Safe Reinforcement Learning via Shielding*’. *AAAI*. AAAI Press, 2018, pages 2669–2678. doi: [10.1609/AAAI.V32I1.11797](https://doi.org/10.1609/AAAI.V32I1.11797).
- [ABH16] K. Abbas, J. Berkhou and B. Heidergott. ‘*A critical account of perturbation analysis of Markov chains*’. *arXiv preprint* (2016). doi: [10.48550/arXiv.1609.04138](https://doi.org/10.48550/arXiv.1609.04138).
- [ABHK18] P. Ashok, Y. Butkova, H. Hermanns and J. Kretínský. ‘*Continuous-Time Markov Decisions Based on Partial Exploration*’. *ATVA*. Volume 11138. Lecture Notes in Computer Science. Springer, 2018, pages 317–334. doi: [10.1007/978-3-030-01090-4_19](https://doi.org/10.1007/978-3-030-01090-4_19).
- [ABRD⁺20] Z. Akata, D. Balliet, M. de Rijke, F. Dignum, V. Dignum, G. Eiben, A. Fokkens, D. Grossi, K. V. Hindriks, H. H. Hoos, H. Hung, C. M. Jonker, C. Monz, M. A. Neerincx, F. A. Oliehoek, H. Prakken, S. Schlobach, L. C. van der Gaag, F. van Harmelen, H. van Hoof, B. van Riemsdijk, A. van Wynsberghe, R. Verbrugge, B. Verheij, P. Vossen and M. Welling. ‘*A Research Agenda for Hybrid Intelligence: Augmenting Human Intellect With Collaborative, Adaptive, Responsible, and Explainable Artificial Intelligence*’. *Computer* 53.8 (2020), pages 18–28. doi: [10.1109/MC.2020.2996587](https://doi.org/10.1109/MC.2020.2996587).
- [ACDK⁺17] P. Ashok, K. Chatterjee, P. Daca, J. Kretínský and T. Meggendorfer. ‘*Value Iteration for Long-Run Average Reward in Markov Decision Processes*’. *CAV* (1). Volume 10426. Lecture Notes in Computer Science. Springer, 2017, pages 201–221. doi: [10.1007/978-3-319-63387-9_10](https://doi.org/10.1007/978-3-319-63387-9_10).
- [ACJK⁺21] R. Andriushchenko, M. Ceska, S. Junges, J. Katoen and S. Stupinský. ‘*PAYNT: A Tool for Inductive Synthesis of Probabilistic Programs*’. *CAV* (1). Volume 12759. Lecture Notes in Computer Science. Springer, 2021, pages 856–869. doi: [10.1007/978-3-030-81685-8_40](https://doi.org/10.1007/978-3-030-81685-8_40).

- [AD10] E. G. Amparore and S. Donatelli. ‘MC4CSLTA: An Efficient Model Checking Tool for CSLTA’. *QEST*. IEEE Computer Society, 2010, pages 153–154. doi: [10.1109/QEST.2010.26](https://doi.org/10.1109/QEST.2010.26).
- [AD18] E. G. Amparore and S. Donatelli. ‘Efficient model checking of the stochastic logic CSL^{TA}’. *Perform. Evaluation* 123-124 (2018), pages 1–34. doi: [10.1016/J.PEVA.2018.03.002](https://doi.org/10.1016/J.PEVA.2018.03.002).
- [AEG22] A. Abate, A. Edwards and M. Giacobbe. ‘Neural Abstractions’. *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*. 2022. doi: [10.48550/arXiv.2301.11683](https://doi.org/10.48550/arXiv.2301.11683).
- [AGKM22] C. Agarwal, S. Guha, J. Kretínský and P. Muruganandham. ‘PAC Statistical Model Checking of Mean Payoff in Discrete- and Continuous-Time MDP’. *CAV (2)*. Volume 13372. Lecture Notes in Computer Science. Springer, 2022, pages 3–25. doi: [10.1007/978-3-031-13188-2_1](https://doi.org/10.1007/978-3-031-13188-2_1).
- [AGR24] A. Abate, M. Giacobbe and D. Roy. ‘Stochastic Omega-Regular Verification and Control with Supermartingales’. *CAV (3)*. Volume 14683. LNCS. Springer, 2024, pages 395–419. doi: [10.1007/978-3-031-65633-0__18](https://doi.org/10.1007/978-3-031-65633-0__18).
- [AH90] J. Aspnes and M. Herlihy. ‘Fast Randomized Consensus Using Shared Memory’. *J. Algorithms* 11.3 (1990), pages 441–461. doi: [10.1016/0196-6774\(90\)90021-6](https://doi.org/10.1016/0196-6774(90)90021-6).
- [AHKV98] R. Alur, T. A. Henzinger, O. Kupferman and M. Y. Vardi. ‘Alternating Refinement Relations’. *CONCUR*. Volume 1466. Lecture Notes in Computer Science. Springer, 1998, pages 163–178. doi: [10.1007/BFb0055622](https://doi.org/10.1007/BFb0055622).
- [AHLPO0] R. Alur, T. A. Henzinger, G. Lafferriere and G. J. Pappas. ‘Discrete abstractions of hybrid systems’. *Proc. IEEE* 88.7 (2000), pages 971–984. doi: [10.1109/5.871304](https://doi.org/10.1109/5.871304).
- [AKW19] P. Ashok, J. Kretínský and M. Weininger. ‘PAC Statistical Model Checking for Markov Decision Processes and Stochastic Games’. *CAV (1)*. Volume 11561. Lecture Notes in Computer Science. Springer, 2019, pages 497–519. doi: [10.1007/978-3-030-25540-4_29](https://doi.org/10.1007/978-3-030-25540-4_29).
- [All10] L. J. Allen. ‘An introduction to stochastic processes with applications to biology’. CRC press, 2010. doi: [10.1201/b12537](https://doi.org/10.1201/b12537).
- [All17] L. J. Allen. ‘A primer on stochastic epidemic models: Formulation, numerical simulation, and analysis’. *Infectious Disease Modelling* 2.2 (2017), pages 128–142. doi: [10.1016/j.idm.2017.03.001](https://doi.org/10.1016/j.idm.2017.03.001).
- [ÅM10] K. J. Åström and R. M. Murray. ‘Feedback systems: an introduction for scientists and engineers’. Princeton university press, 2010. doi: [10.1515/9781400828739](https://doi.org/10.1515/9781400828739).
- [AM90] B. D. O. Anderson and J. B. Moore. ‘Optimal control: linear quadratic methods’. Prentice-Hall, Inc., 1990. doi: [10.5555/79089](https://doi.org/10.5555/79089).

- [AOSC⁺16] D. Amodei, C. Olah, J. Steinhardt, P. F. Christiano, J. Schulman and D. Mané. ‘Concrete Problems in AI Safety’. *CoRR* abs/1606.06565 (2016). doi: [10.48550/arXiv.1606.06565](https://doi.org/10.48550/arXiv.1606.06565).
- [AP18] G. Agha and K. Palmskog. ‘A Survey of Statistical Model Checking’. *ACM Trans. Model. Comput. Simul.* 28.1 (2018), 6:1–6:39. doi: [10.1145/3158668](https://doi.org/10.1145/3158668).
- [APLS08] A. Abate, M. Prandini, J. Lygeros and S. Sastry. ‘Probabilistic reachability and safety for controlled discrete time stochastic hybrid systems’. *Autom.* 44.11 (2008), pages 2724–2734. doi: [10.1016/J.AUTOMATICA.2008.03.027](https://doi.org/10.1016/J.AUTOMATICA.2008.03.027).
- [ASSB00] A. Aziz, K. Sanwal, V. Singhal and R. Brayton. ‘Model-checking continuous-time Markov chains’. *ACM Transactions on Computational Logic* 1.1 (2000), pages 162–170. doi: [10.1145/343369.343402](https://doi.org/10.1145/343369.343402).
- [Åst12] K. J. Åström. ‘Introduction to stochastic control theory’. Courier Corporation, 2012. url: <https://portal.research.lu.se/en/publications/introduction-to-stochastic-control-theory-2>.
- [AVLA⁺17] A. Ahmed, P. Varakantham, M. Lowalekar, Y. Adulyasak and P. Jaillet. ‘Sampling Based Approaches for Minimizing Regret in Uncertain Markov Decision Processes (MDPs)’. *J. Artif. Intell. Res.* 59 (2017), pages 229–264. doi: [10.1613/JAIR.5242](https://doi.org/10.1613/JAIR.5242).
- [Bar18] S. Barratt. ‘On the differentiability of the solution to convex optimization problems’. *arXiv preprint* (2018). doi: [10.48550/arXiv.1804.05098](https://doi.org/10.48550/arXiv.1804.05098).
- [BBC11] D. Bertsimas, D. B. Brown and C. Caramanis. ‘Theory and Applications of Robust Optimization’. *SIAM Rev.* 53.3 (2011), pages 464–501. doi: [10.1137/080734510](https://doi.org/10.1137/080734510).
- [BCCF⁺14] T. Brázdil, K. Chatterjee, M. Chmelík, V. Forejt, J. Kretínský, M. Z. Kwiatkowska, D. Parker and M. Ujma. ‘Verification of Markov Decision Processes Using Learning Algorithms’. *ATVA*. Volume 8837. Lecture Notes in Computer Science. Springer, 2014, pages 98–114. doi: [10.1007/978-3-319-11936-6_8](https://doi.org/10.1007/978-3-319-11936-6_8).
- [BCCF⁺24] T. Brázdil, K. Chatterjee, M. Chmelík, V. Forejt, J. Kretínský, M. Kwiatkowska, T. Meggendorfer, D. Parker and M. Ujma. ‘Learning Algorithms for Verification of Markov Decision Processes’. *CoRR* abs/2403.09184 (2024). doi: [10.48550/ARXIV.2403.09184](https://doi.org/10.48550/ARXIV.2403.09184).
- [BCHT17] S. Bansal, M. Chen, S. L. Herbert and C. J. Tomlin. ‘Hamilton-Jacobi reachability: A brief overview and recent advances’. *CDC*. IEEE, 2017, pages 2242–2253. doi: [10.1109/CDC.2017.8263977](https://doi.org/10.1109/CDC.2017.8263977).
- [BCMJJ19] E. Bøhn, E. M. Coates, S. Moe and T. A. Johansen. ‘Deep Reinforcement Learning Attitude Control of Fixed-Wing UAVs Using Proximal Policy optimization’ (2019), pages 523–533. doi: [10.1109/ICUAS.2019.8798254](https://doi.org/10.1109/ICUAS.2019.8798254).

- [BD05] H. Boudali and J. B. Dugan. ‘*A new Bayesian network approach to solve dynamic fault trees*’. *Proc. of RAMS. IEEE*. 2005, pages 451–456. doi: [10.1109/RAMS.2005.1408404](https://doi.org/10.1109/RAMS.2005.1408404).
- [BDDF⁺18] E. Bartocci, J. V. Deshmukh, A. Donzé, G. Fainekos, O. Maler, D. Nickovic and S. Sankaranarayanan. ‘*Specification-Based Monitoring of Cyber-Physical Systems: A Survey on Theory, Tools and Applications*’. *Lectures on Runtime Verification*. Volume 10457. Lecture Notes in Computer Science. Springer, 2018, pages 135–175. doi: [10.1007/978-3-319-75632-5_5](https://doi.org/10.1007/978-3-319-75632-5_5).
- [BDFL⁺18] A. Bart, B. Delahaye, P. Fournier, D. Lime, É. Monfroy and C. Truchet. ‘*Reachability in parametric Interval Markov Chains using constraints*’. *Theor. Comput. Sci.* 747 (2018), pages 48–74. doi: [10.1016/J.TCS.2018.06.016](https://doi.org/10.1016/J.TCS.2018.06.016).
- [BDHH⁺17] C. E. Budde, C. Dehnert, E. M. Hahn, A. Hartmanns, S. Junges and A. Turrini. ‘*JANI: Quantitative Model and Tool Interaction*’. *TACAS (2)*. Volume 10206. Lecture Notes in Computer Science. 2017, pages 151–168. doi: [10.1007/978-3-662-54580-5_9](https://doi.org/10.1007/978-3-662-54580-5_9).
- [BDKK⁺19] C. Baier, C. Dubslaff, L. Koreniak, A. Kucera and V. Rehák. ‘*Mean-payoff Optimization in Continuous-time Markov Chains with Parametric Alarms*’. *ACM Trans. Model. Comput. Simul.* 29.4 (2019), 28:1–28:26. doi: [10.1145/3310225](https://doi.org/10.1145/3310225).
- [Bel66] R. Bellman. ‘*Dynamic programming*’. *Science* 153.3731 (1966), pages 34–37. doi: [10.1126/science.153.3731.34](https://doi.org/10.1126/science.153.3731.34).
- [BFKK⁺13] T. Brázdil, V. Forejt, J. Krcál, J. Kretínský and A. Kucera. ‘*Continuous-time stochastic games with time-bounded reachability*’. *Inf. Comput.* 224 (2013), pages 46–70. doi: [10.1016/J.IC.2013.01.001](https://doi.org/10.1016/J.IC.2013.01.001).
- [BG09] B. Bonet and H. Geffner. ‘*Solving POMDPs: RTDP-Bel vs. Point-based Algorithms*’. *IJCAI*. 2009, pages 1641–1646. URL: <https://dl.acm.org/doi/10.5555/1661445.1661709>.
- [BGHY⁺22] L. Brunke, M. Greeff, A. W. Hall, Z. Yuan, S. Zhou, J. Panerati and A. P. Schoellig. ‘*Safe Learning in Robotics: From Learning-Based Control to Safe Reinforcement Learning*’. *Annu. Rev. Control. Robotics Auton. Syst.* 5 (2022), pages 411–444. doi: [10.1146/annurev-control-042920-020211](https://doi.org/10.1146/annurev-control-042920-020211).
- [BGN09] A. Ben-Tal, L. E. Ghaoui and A. Nemirovski. ‘*Robust Optimization*’. Volume 28. Princeton Series in Applied Mathematics. Princeton University Press, 2009. doi: [10.1515/9781400831050](https://doi.org/10.1515/9781400831050).
- [BH22] D. Bertsimas and D. den Hertog. ‘*Robust and Adaptive Optimization*’. Dynamic Ideas LLC, 2022. URL: <https://www.dynamic-ideas.com/books/robust-and-adaptive-optimization>.
- [BH97] C. Baier and H. Hermanns. ‘*Weak Bisimulation for Fully Probabilistic Processes*’. *CAV*. Volume 1254. Lecture Notes in Computer Science. Springer, 1997, pages 119–130. doi: [10.1007/3-540-63166-6_14](https://doi.org/10.1007/BF03342514).

- [BHHJ⁺20] C. Baier, C. Hensel, L. Hutschenreiter, S. Junges, J. Katoen and J. Klein. ‘Parametric Markov chains: PCTL complexity and fraction-free Gaussian elimination’. *Inf. Comput.* 272 (2020), page 104504. doi: [10.1016/J.IC.2019.104504](https://doi.org/10.1016/J.IC.2019.104504).
- [BHHK03] C. Baier, B. R. Haverkort, H. Hermanns and J. Katoen. ‘Model-Checking Algorithms for Continuous-Time Markov Chains’. *IEEE Trans. Software Eng.* 29.6 (2003), pages 524–541. doi: [10.1109/TSE.2003.1205180](https://doi.org/10.1109/TSE.2003.1205180).
- [BHKH05] C. Baier, H. Hermanns, J. Katoen and B. R. Haverkort. ‘Efficient computation of time-bounded reachability probabilities in uniform continuous-time Markov decision processes’. *Theor. Comput. Sci.* 345.1 (2005), pages 2–26. doi: [10.1016/J.TCS.2005.07.022](https://doi.org/10.1016/J.TCS.2005.07.022).
- [BHL19] G. Bacci, M. Hansen and K. G. Larsen. ‘Model Checking Constrained Markov Reward Models with Uncertainties’. *QEST*. Volume 11785. Lecture Notes in Computer Science. Springer, 2019, pages 37–51. doi: [10.1007/978-3-030-30281-8_3](https://doi.org/10.1007/978-3-030-30281-8_3).
- [BHTB⁺18] J. Buckman, D. Hafner, G. Tucker, E. Brevdo and H. Lee. ‘Sample-Efficient Reinforcement Learning with Stochastic Ensemble Value Expansion’. *NeurIPS*. 2018, pages 8234–8244. doi: [10.48550/arXiv.1807.01675](https://doi.org/10.48550/arXiv.1807.01675).
- [BJKQ20] A. Bork, S. Junges, J. Katoen and T. Quatmann. ‘Verification of Indefinite-Horizon POMDPs’. *ATVA*. Volume 12302. Lecture Notes in Computer Science. Springer, 2020, pages 288–304. doi: [10.1007/978-3-030-59152-6_16](https://doi.org/10.1007/978-3-030-59152-6_16).
- [BJS11] M. S. Bazaraa, J. J. Jarvis and H. D. Sherali. ‘Linear programming and network flows’. John Wiley & Sons, 2011. doi: [10.1002/9780471703778](https://doi.org/10.1002/9780471703778).
- [BK08] C. Baier and J. Katoen. ‘Principles of model checking’. MIT Press, 2008. URL: <https://mitpress.mit.edu/9780262026499/principles-of-model-checking/>.
- [BKKM14] C. Baier, J. Klein, S. Klüppelholz and S. Märcker. ‘Computing Conditional Probabilities in Markovian Models Efficiently’. *TACAS*. Volume 8413. Lecture Notes in Computer Science. Springer, 2014, pages 515–530. doi: [10.1007/978-3-642-54862-8_43](https://doi.org/10.1007/978-3-642-54862-8_43).
- [BKKN⁺15] T. Brázdil, L. Korenciak, J. Krcál, P. Novotný and V. Rehák. ‘Optimizing Performance of Continuous-Time Stochastic Systems Using Timeout Synthesis’. *QEST*. Volume 9259. Lecture Notes in Computer Science. Springer, 2015, pages 141–159. doi: [10.1007/978-3-319-22264-6_10](https://doi.org/10.1007/978-3-319-22264-6_10).
- [BKQ22] A. Bork, J. Katoen and T. Quatmann. ‘Under-Approximating Expected Total Rewards in POMDPs’. *TACAS (2)*. Volume 13244. Lecture Notes in Computer Science. Springer, 2022, pages 22–40. doi: [10.1007/978-3-030-99527-0_2](https://doi.org/10.1007/978-3-030-99527-0_2).
- [BLDT⁺21] F. S. Barbosa, B. Lacerda, P. Duckworth, J. Tumova and N. Hawes. ‘Risk-Aware Motion Planning in Partially Known Environments’. *CDC*. IEEE, 2021, pages 5220–5226. doi: [10.1109/CDC45484.2021.9683744](https://doi.org/10.1109/CDC45484.2021.9683744).

- [BLM13] S. Boucheron, G. Lugosi and P. Massart. ‘*Concentration Inequalities - A Nonasymptotic Theory of Independence*’. Oxford University Press, 2013. doi: [10.1093/ACPROF:OSO/9780199535255.001.0001](https://doi.org/10.1093/ACPROF:OSO/9780199535255.001.0001).
- [BMS16] L. Bortolussi, D. Milios and G. Sanguinetti. ‘*Smoothed model checking for uncertain Continuous-Time Markov Chains*’. *Inf. Comput.* 247 (2016), pages 235–253. doi: [10.1016/J.IC.2016.01.004](https://doi.org/10.1016/J.IC.2016.01.004).
- [BOBW10] L. Blackmore, M. Ono, A. Bektassov and B. C. Williams. ‘*A Probabilistic Particle-Control Approximation of Chance-Constrained Stochastic Predictive Control*’. *IEEE Trans. Robotics* 26.3 (2010), pages 502–517. doi: [10.1109/TRO.2010.2044948](https://doi.org/10.1109/TRO.2010.2044948).
- [Box76] G. E. P. Box. ‘*Science and Statistics*’. *Journal of the American Statistical Association* 71.356 (1976), pages 791–799. doi: [10.1080/01621459.1976.10480949](https://doi.org/10.1080/01621459.1976.10480949).
- [BR07] V. I. Bogachev and M. A. S. Ruas. ‘*Measure theory*’. Volume 1. Springer, 2007. doi: [10.1007/978-3-540-34514-5](https://doi.org/10.1007/978-3-540-34514-5).
- [BS18] L. Bortolussi and S. Silvetti. ‘*Bayesian Statistical Parameter Synthesis for Linear Temporal Properties of Stochastic Models*’. TACAS (2). Volume 10806. Lecture Notes in Computer Science. Springer, 2018, pages 396–413. doi: [10.1007/978-3-319-89963-3_23](https://doi.org/10.1007/978-3-319-89963-3_23).
- [BS78] D. P. Bertsekas and S. E. Shreve. ‘*Stochastic Optimal Control: The Discrete-time Case*’. Athena Scientific, 1978, page 330. URL: <https://web.mit.edu/dimitrib/www/soc.html>.
- [BSJJ24] E. M. Bovy, M. Suilen, S. Junges and N. Jansen. ‘*Imprecise Probabilities Meet Partial Observability: Game Semantics for Robust POMDPs*’. CoRR abs/2405.04941 (2024). doi: [10.48550/arXiv.2405.04941](https://doi.org/10.48550/arXiv.2405.04941).
- [BT02] R. I. Brafman and M. Tennenholtz. ‘*R-MAX - A General Polynomial Time Algorithm for Near-Optimal Reinforcement Learning*’. *J. Mach. Learn. Res.* 3 (2002), pages 213–231. doi: [10.1162/153244303765208377](https://doi.org/10.1162/153244303765208377).
- [BTSK17] F. Berkenkamp, M. Turchetta, A. P. Schoellig and A. Krause. ‘*Safe Model-based Reinforcement Learning with Stability Guarantees*’. NIPS. 2017, pages 908–918. doi: [10.48550/arXiv.1705.08551](https://doi.org/10.48550/arXiv.1705.08551).
- [BV14] S. P. Boyd and L. Vandenberghe. ‘*Convex Optimization*’. Cambridge University Press, 2014. doi: [10.1017/CBO9780511804441](https://doi.org/10.1017/CBO9780511804441).
- [BYG17] C. Belta, B. Yordanov and E. A. Gol. ‘*Formal methods for discrete-time dynamical systems*’. Volume 15. Springer, 2017. doi: [10.1007/978-3-319-50763-7](https://doi.org/10.1007/978-3-319-50763-7).
- [CA19] N. Cauchi and A. Abate. ‘*StochHy: Automated Verification and Synthesis of Stochastic Processes*’. TACAS (2). Volume 11428. Lecture Notes in Computer Science. Springer, 2019, pages 247–264. doi: [10.1007/978-3-030-17465-1_14](https://doi.org/10.1007/978-3-030-17465-1_14).
- [CB21] G. Casella and R. L. Berger. ‘*Statistical inference*’. Cengage Learning, 2021. doi: [10.1201/9781003456285](https://doi.org/10.1201/9781003456285).

- [CBSN⁺16] R. C. Cavalcante, R. C. Brasileiro, V. L. F. Souza, J. P. Nóbrega and A. L. I. Oliveira. ‘Computational Intelligence and Financial Markets: A Survey and Future Directions’. *Expert Syst. Appl.* 55 (2016), pages 194–211. doi: [10.1016/J.ESWA.2016.02.006](https://doi.org/10.1016/J.ESWA.2016.02.006).
- [CC05] G. C. Calafiore and M. C. Campi. ‘Uncertain convex programs: randomized solutions and confidence levels’. *Math. Program.* 102.1 (2005), pages 25–46. doi: [10.1007/S10107-003-0499-Y](https://doi.org/10.1007/S10107-003-0499-Y).
- [CC06] G. C. Calafiore and M. C. Campi. ‘The scenario approach to robust control design’. *IEEE Trans. Autom. Control*. 51.5 (2006), pages 742–753. doi: [10.1109/TAC.2006.875041](https://doi.org/10.1109/TAC.2006.875041).
- [CC97] X. Cao and H. Chen. ‘Perturbation realization, potentials, and sensitivity analysis of Markov processes’. *IEEE Trans. Autom. Control*. 42.10 (1997), pages 1382–1393. doi: [10.1109/9.633827](https://doi.org/10.1109/9.633827).
- [CCG21] M. C. Campi, A. Carè and S. Garatti. ‘The scenario approach: A tool at the service of data-driven decision making’. *Annu. Rev. Control.* 52 (2021), pages 1–17. doi: [10.1016/J.ARCONTROL.2021.10.004](https://doi.org/10.1016/J.ARCONTROL.2021.10.004).
- [CCGK⁺18] R. Calinescu, M. Ceska, S. Gerasimou, M. Kwiatkowska and N. Paoletti. ‘Efficient synthesis of robust models for stochastic systems’. *J. Syst. Softw.* 143 (2018), pages 140–158. doi: [10.1016/J.JSS.2018.05.013](https://doi.org/10.1016/J.JSS.2018.05.013).
- [CCT16] K. Chatterjee, M. Chmelik and M. Tracol. ‘What is decidable about partially observable Markov decision processes with ω -regular objectives’. *J. Comput. Syst. Sci.* 82.5 (2016), pages 878–911. doi: [10.1016/J.JCSS.2016.02.009](https://doi.org/10.1016/J.JCSS.2016.02.009).
- [CDPK⁺17] M. Ceska, F. Dannerberg, N. Paoletti, M. Kwiatkowska and L. Brim. ‘Precise parameter synthesis for stochastic biochemical systems’. *Acta Informatica* 54.6 (2017), pages 589–623. doi: [10.1007/S00236-016-0265-2](https://doi.org/10.1007/S00236-016-0265-2).
- [CE15] S. N. Cohen and R. J. Elliott. ‘Stochastic calculus and applications’. Volume 2. Springer, 2015. doi: [10.1007/978-1-4939-2867-5](https://doi.org/10.1007/978-1-4939-2867-5).
- [CFRS14] T. Chen, Y. Feng, D. S. Rosenblum and G. Su. ‘Perturbation Analysis in Verification of Discrete-Time Markov Chains’. CONCUR. Volume 8704. Lecture Notes in Computer Science. Springer, 2014, pages 218–233. doi: [10.1007/978-3-662-44584-6_16](https://doi.org/10.1007/978-3-662-44584-6_16).
- [CG08] M. C. Campi and S. Garatti. ‘The Exact Feasibility of Randomized Solutions of Uncertain Convex Programs’. *SIAM J. Optim.* 19.3 (2008), pages 1211–1230. doi: [10.1137/07069821X](https://doi.org/10.1137/07069821X).
- [CG11] M. C. Campi and S. Garatti. ‘A Sampling-and-Discarding Approach to Chance-Constrained Optimization: Feasibility and Optimality’. *J. Optim. Theory Appl.* 148.2 (2011), pages 257–280. doi: [10.1007/S10957-010-9754-6](https://doi.org/10.1007/S10957-010-9754-6).
- [CG18a] M. C. Campi and S. Garatti. ‘Introduction to the scenario approach’. SIAM, 2018. doi: [10.1137/1.9781611975444](https://doi.org/10.1137/1.9781611975444).

- [CG18b] M. C. Campi and S. Garatti. ‘*Wait-and-judge scenario optimization*’. *Math. Program.* 167.1 (2018), pages 155–189. doi: [10.1007/S10107-016-1056-9](https://doi.org/10.1007/S10107-016-1056-9).
- [CG20] M. C. Campi and S. Garatti. ‘*Scenario optimization with relaxation: a new tool for design and application to machine learning problems*’. *CDC. IEEE*, 2020, pages 2463–2468. doi: [10.1109/CDC42340.2020.9303914](https://doi.org/10.1109/CDC42340.2020.9303914).
- [CGJL⁺03] E. M. Clarke, O. Grumberg, S. Jha, Y. Lu and H. Veith. ‘*Counterexample-guided abstraction refinement for symbolic model checking*’. *J. ACM* 50.5 (2003), pages 752–794. doi: [10.1145/876638.876643](https://doi.org/10.1145/876638.876643).
- [CGJP⁺16] R. Calinescu, C. Ghezzi, K. Johnson, M. Pezzè, Y. Rafiq and G. Tamburrelli. ‘*Formal Verification With Confidence Intervals to Establish Quality of Service Properties of Software Systems*’. *IEEE Trans. Reliab.* 65.1 (2016), pages 107–125. doi: [10.1109/TR.2015.2452931](https://doi.org/10.1109/TR.2015.2452931).
- [CGLT⁺21] L. Cardelli, R. Grosu, K. G. Larsen, M. Tribastone, M. Tschaikowski and A. Vandin. ‘*Lumpability for Uncertain Continuous-Time Markov Chains*’. *QEST. Volume 12846. Lecture Notes in Computer Science*. Springer, 2021, pages 391–409. doi: [10.1007/978-3-030-85172-9_21](https://doi.org/10.1007/978-3-030-85172-9_21).
- [CGP09] M. C. Campi, S. Garatti and M. Prandini. ‘*The scenario approach for systems and control design*’. *Annu. Rev. Control.* 33.2 (2009), pages 149–157. doi: [10.1016/J.ARCONTROL.2009.07.001](https://doi.org/10.1016/J.ARCONTROL.2009.07.001).
- [CH08] K. Chatterjee and T. A. Henzinger. ‘*Value Iteration*’. *25 Years of Model Checking. Volume 5000. Lecture Notes in Computer Science*. Springer, 2008, pages 107–138. doi: [10.1007/978-3-540-69850-0_7](https://doi.org/10.1007/978-3-540-69850-0_7).
- [CHHK⁺13] T. Chen, E. M. Hahn, T. Han, M. Z. Kwiatkowska, H. Qu and L. Zhang. ‘*Model Repair for Markov Decision Processes*’. *TASE. IEEE Computer Society*, 2013, pages 85–92. doi: [10.1109/TASE.2013.20](https://doi.org/10.1109/TASE.2013.20).
- [CHKM11] T. Chen, T. Han, J. Katoen and A. Mereacre. ‘*Model Checking of Continuous-Time Markov Chains Against Timed Automata Specifications*’. *Log. Methods Comput. Sci.* 7.1 (2011). doi: [10.2168/LMCS-7\(1:12\)2011](https://doi.org/10.2168/LMCS-7(1:12)2011).
- [Cho19] V. Chonev. ‘*Reachability in Augmented Interval Markov Chains*’. *RP. Volume 11674. Lecture Notes in Computer Science*. Springer, 2019, pages 79–92. doi: [10.1007/978-3-030-30806-3_7](https://doi.org/10.1007/978-3-030-30806-3_7).
- [CHQ08] G. D. Cooman, F. Hermans and E. Quaeghebeur. ‘*Sensitivity analysis for finite Markov chains in discrete time*’. *UAI. AUAI Press*, 2008, pages 129–136. doi: [10.1561/11638](https://doi.org/10.1561/11638).
- [CJJK⁺17] M. Cubuktepe, N. Jansen, S. Junges, J. Katoen, I. Papusha, H. A. Poonawala and U. Topcu. ‘*Sequential Convex Programming for the Efficient Verification of Parametric MDPs*’. *TACAS (2). Volume 10206. Lecture Notes in Computer Science*. 2017, pages 133–150. doi: [10.1007/978-3-662-54580-5_8](https://doi.org/10.1007/978-3-662-54580-5_8).

- [CJJK⁺18] M. Cubuktepe, N. Jansen, S. Junges, J. Katoen and U. Topcu. ‘*Synthesis in pMDPs: A Tale of 1001 Parameters*’. ATVA. Volume 11138. Lecture Notes in Computer Science. Springer, 2018, pages 160–176. doi: [10.1007/978-3-030-01090-4_10](https://doi.org/10.1007/978-3-030-01090-4_10).
- [CJJK⁺20] M. Cubuktepe, N. Jansen, S. Junges, J. Katoen and U. Topcu. ‘*Scenario-Based Verification of Uncertain MDPs*’. TACAS (1). Volume 12078. Lecture Notes in Computer Science. Springer, 2020, pages 287–305. doi: [10.1007/978-3-030-45190-5_16](https://doi.org/10.1007/978-3-030-45190-5_16).
- [CJJK⁺22] M. Cubuktepe, N. Jansen, S. Junges, J. Katoen and U. Topcu. ‘*Convex Optimization for Parameter Synthesis in MDPs*’. IEEE Trans. Autom. Control. 67.12 (2022), pages 6333–6348. doi: [10.1109/TAC.2021.3133265](https://doi.org/10.1109/TAC.2021.3133265).
- [CJJK19] M. Ceska, N. Jansen, S. Junges and J. Katoen. ‘*Shepherding Hordes of Markov Chains*’. TACAS (2). Volume 11428. Lecture Notes in Computer Science. Springer, 2019, pages 172–190. doi: [10.1007/978-3-030-17465-1_10](https://doi.org/10.1007/978-3-030-17465-1_10).
- [CJJM⁺21] M. Cubuktepe, N. Jansen, S. Junges, A. Marandi, M. Suilen and U. Topcu. ‘*Robust Finite-State Controllers for Uncertain POMDPs*’. AAAI. AAAI Press, 2021, pages 11792–11800. doi: [10.1609/AAAI.V35I13.17401](https://doi.org/10.1609/AAAI.V35I13.17401).
- [CJJT23] S. Carr, N. Jansen, S. Junges and U. Topcu. ‘*Safe Reinforcement Learning via Shielding under Partial Observability*’. AAAI. AAAI Press, 2023, pages 14748–14756. doi: [10.1609/AAAI.V37I12.26723](https://doi.org/10.1609/AAAI.V37I12.26723).
- [CKL94] A. R. Cassandra, L. P. Kaelbling and M. L. Littman. ‘*Acting Optimally in Partially Observable Stochastic Domains*’. AAAI. AAAI Press / The MIT Press, 1994, pages 1023–1028. url: <https://cdn.aaai.org/AAAI/1994/AAAI94-157.pdf>.
- [CLLA⁺19] N. Cauchi, L. Laurenti, M. Lahijanian, A. Abate, M. Kwiatkowska and L. Cardelli. ‘*Efficiency through uncertainty: scalable formal synthesis for stochastic hybrid systems*’. HSCC. ACM, 2019, pages 240–251. doi: [10.1145/3302504.3311805](https://doi.org/10.1145/3302504.3311805).
- [CLZ97] A. R. Cassandra, M. L. Littman and N. L. Zhang. ‘*Incremental Pruning: A Simple, Fast, Exact Method for Partially Observable Markov Decision Processes*’. UAI. Morgan Kaufmann, 1997, pages 54–61. doi: [10.48550/arXiv.1302.1525](https://arxiv.org/abs/1302.1525).
- [COB21] G. Chou, N. Ozay and D. Berenson. ‘*Model Error Propagation via Learned Contraction Metrics for Safe Feedback Motion Planning of Unknown Systems*’. CDC. IEEE, 2021, pages 3576–3583. doi: [10.1109/CDC45484.2021.9683354](https://doi.org/10.1109/CDC45484.2021.9683354).
- [ÇOK22] M. M. Çelikok, F. A. Oliehoek and S. Kaski. ‘*Best-Response Bayesian Reinforcement Learning with Bayes-adaptive POMDPs for Centaurs*’. AAMAS. International Foundation for Autonomous Agents and Multiagent Systems (IFAAMAS), 2022, pages 235–243. doi: [10.5555/3535878](https://doi.org/10.5555/3535878).

- [CP34] C. J. Clopper and E. S. Pearson. ‘*The use of confidence or fiducial limits illustrated in the case of the binomial*’. *Biometrika* 26.4 (1934), pages 404–413. DOI: [10.2307/2331986](https://doi.org/10.2307/2331986).
- [CPM23] R. Coppola, A. Peruffo and M. Mazo Jr. ‘*Data-Driven Abstractions for Verification of Linear Systems*’. *IEEE Control. Syst. Lett.* 7 (2023), pages 2737–2742. DOI: [10.1109/LCSYS.2023.3288731](https://doi.org/10.1109/LCSYS.2023.3288731).
- [CRLH23] C. Costen, M. Rigter, B. Lacerda and N. Hawes. ‘*Planning with Hidden Parameter Polynomial MDPs*’. *AAAI*. AAAI Press, 2023, pages 11963–11971. DOI: [10.1609/AAAI.V37I10.26411](https://doi.org/10.1609/AAAI.V37I10.26411).
- [CS07] T. R. Colburn and G. M. Shute. ‘*Abstraction in Computer Science*’. *Minds Mach.* 17.2 (2007), pages 169–184. DOI: [10.1007/S11023-007-9061-7](https://doi.org/10.1007/S11023-007-9061-7).
- [CSKG22] B. Charpentier, R. Senanayake, M. J. Kochenderfer and S. Günnemann. ‘*Disentangling Epistemic and Aleatoric Uncertainty in Reinforcement Learning*’. *CoRR* abs/2206.01558 (2022). DOI: [10.48550/arXiv.2206.01558](https://doi.org/10.48550/arXiv.2206.01558).
- [CT92] H. Choi and K. S. Trivedi. ‘*Approximate Performance Models of Polling Systems Using Stochastic Petri Nets*’. *INFOCOM*. IEEE Computer Society, 1992, pages 2306–2314. DOI: [10.1109/INFCOM.1992.263520](https://doi.org/10.1109/INFCOM.1992.263520).
- [CT96] G. Ciardo and M. Tilgner. ‘*On the use of Kronecker operators for the solution of generalized stochastic Petri nets*’ (1996). URL: <https://ntrs.nasa.gov/citations/20040110963>.
- [CW98] X. Cao and Y. Wan. ‘*Algorithms for sensitivity analysis of Markov systems through potentials and perturbation realization*’. *IEEE Trans. Control. Syst. Technol.* 6.4 (1998), pages 482–494. DOI: [10.1109/87.701341](https://doi.org/10.1109/87.701341).
- [Dav18] M. H. Davis. ‘*Markov models & optimization*’. Routledge, 2018. DOI: [10.1201/9780203748039](https://doi.org/10.1201/9780203748039).
- [Daw04] C. Daws. ‘*Symbolic and Parametric Model Checking of Discrete-Time Markov Chains*’. *ICTAC*. Volume 3407. Lecture Notes in Computer Science. Springer, 2004, pages 280–294. DOI: [10.1007/978-3-540-31862-0_21](https://doi.org/10.1007/978-3-540-31862-0_21).
- [DBB18] N. M. van Dijk, S. P. J. van Brummelen and R. J. Boucherie. ‘*Uniformization: Basics, extensions and applications*’. *Perform. Evaluation* 118 (2018), pages 8–32. DOI: [10.1016/J.PEVA.2017.09.008](https://doi.org/10.1016/J.PEVA.2017.09.008).
- [DCB13] A. Domahidi, E. Chu and S. P. Boyd. ‘*ECOS: An SOCP solver for embedded systems*’. *ECC*. IEEE, 2013, pages 3071–3076. DOI: [10.23919/ECC.2013.6669541](https://doi.org/10.23919/ECC.2013.6669541).
- [DDNZ00] M. S. De Queiroz, D. M. Dawson, S. P. Nagarkatti and F. Zhang. ‘*Lyapunov-based control of mechanical systems*’. Springer Science & Business Media, 2000. DOI: [10.1007/978-1-4612-1352-9](https://doi.org/10.1007/978-1-4612-1352-9).
- [Del15] B. Delahaye. ‘*Consistency for Parametric Interval Markov Chains*’. *SynCoP*. Volume 44. OASIcs. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2015, pages 17–32. DOI: [10.4230/OASIcs.SYNCP.2015.17](https://doi.org/10.4230/OASIcs.SYNCP.2015.17).

- [DEP02] J. Desharnais, A. Edalat and P. Panangaden. ‘*Bisimulation for Labelled Markov Processes*’. *Inf. Comput.* 179.2 (2002), pages 163–193. doi: [10.1006/INCO.2001.2962](https://doi.org/10.1006/INCO.2001.2962).
- [DGJP04] J. Desharnais, V. Gupta, R. Jagadeesan and P. Panangaden. ‘*Metrics for labelled Markov processes*’. *Theor. Comput. Sci.* 318.3 (2004), pages 323–354. doi: [10.1016/J.TCS.2003.09.013](https://doi.org/10.1016/J.TCS.2003.09.013).
- [DHDU18] S. Depeweg, J. M. Hernández-Lobato, F. Doshi-Velez and S. Udluft. ‘*Decomposition of Uncertainty in Bayesian Deep Learning for Efficient and Risk-sensitive Learning*’. *ICML*. Volume 80. Proceedings of Machine Learning Research. PMLR, 2018, pages 1192–1201. doi: [10.48550/arXiv.1710.07283](https://doi.org/10.48550/arXiv.1710.07283).
- [DHK15] F. Dannenberg, E. M. Hahn and M. Z. Kwiatkowska. ‘*Computing Cumulative Rewards Using Fast Adaptive Uniformization*’. *ACM Trans. Model. Comput. Simul.* 25.2 (2015), 9:1–9:23. doi: [10.1145/2688907](https://doi.org/10.1145/2688907).
- [DHKP17] P. Daca, T. A. Henzinger, J. Kretínský and T. Petrov. ‘*Faster Statistical Model Checking for Unbounded Temporal Properties*’. *ACM Trans. Comput. Log.* 18.2 (2017), 12:1–12:25. doi: [10.1145/3060139](https://doi.org/10.1145/3060139).
- [DHS18] P. R. D’Argenio, A. Hartmanns and S. Sedwards. ‘*Lightweight Statistical Model Checking in Nondeterministic Continuous Time*’. *ISoLA* (2). Volume 11245. Lecture Notes in Computer Science. Springer, 2018, pages 336–353. doi: [10.1007/978-3-030-03421-4_22](https://doi.org/10.1007/978-3-030-03421-4_22).
- [Die17] T. G. Dietterich. ‘*Steps Toward Robust Artificial Intelligence*’. *AI Mag.* 38.3 (2017), pages 3–24. doi: [10.1609/AIMAG.V38I3.2756](https://doi.org/10.1609/AIMAG.V38I3.2756).
- [Dij71] E. W. Dijkstra. ‘*Hierarchical Ordering of Sequential Processes*’. *Acta Informatica* 1 (1971), pages 115–138. doi: [10.1007/BF00289519](https://doi.org/10.1007/BF00289519).
- [DJJC⁺15] C. Dehnert, S. Junges, N. Jansen, F. Corzilius, M. Volk, H. Bruintjes, J. Katoen and E. Ábrahám. ‘*PROPhESY: A PRObabilistic ParamEter SYnthesis Tool*’. *CAV* (1). Volume 9206. Lecture Notes in Computer Science. Springer, 2015, pages 214–231. doi: [10.1007/978-3-319-21690-4_13](https://doi.org/10.1007/978-3-319-21690-4_13).
- [DJJL01] P. R. D’Argenio, B. Jeannet, H. E. Jensen and K. G. Larsen. ‘*Reachability Analysis of Probabilistic Systems by Successive Refinements*’. *PAPM-PROBMIV*. Volume 2165. Lecture Notes in Computer Science. Springer, 2001, pages 39–56. doi: [10.1007/3-540-44804-7_3](https://doi.org/10.1007/3-540-44804-7_3).
- [DJKV17] C. Dehnert, S. Junges, J. Katoen and M. Volk. ‘*A Storm is Coming: A Modern Probabilistic Model Checker*’. *CAV* (2). Volume 10427. Lecture Notes in Computer Science. Springer, 2017, pages 592–600. doi: [10.1007/978-3-319-63390-9_31](https://doi.org/10.1007/978-3-319-63390-9_31).
- [DLLM⁺11] A. David, K. G. Larsen, A. Legay, M. Mikucionis and Z. Wang. ‘*Time for Statistical Model Checking of Real-Time Systems*’. *CAV*. Volume 6806. Lecture Notes in Computer Science. Springer, 2011, pages 349–355. doi: [10.1007/978-3-642-22110-1_27](https://doi.org/10.1007/978-3-642-22110-1_27).

- [DLLM⁺15] A. David, K. G. Larsen, A. Legay, M. Mikucionis and D. B. Poulsen. ‘*Uppaal SMC tutorial*’. *Int. J. Softw. Tools Technol. Transf.* 17.4 (2015), pages 397–415. doi: [10.1007/S10009-014-0361-Y](https://doi.org/10.1007/S10009-014-0361-Y).
- [DLP16] B. Delahaye, D. Lime and L. Petrucci. ‘*Parameter Synthesis for Parametric Interval Markov Chains*’. VMCAI. Volume 9583. Lecture Notes in Computer Science. Springer, 2016, pages 372–390. doi: [10.1007/978-3-662-49122-5_18](https://doi.org/10.1007/978-3-662-49122-5_18).
- [Dry43] H. L. Dryden. ‘*A review of the statistical theory of turbulence*’. *Quarterly of Applied Mathematics* 1.1 (1943), pages 7–42. URL: <https://www.ams.org/journals/qam/1943-01-01/S0033-569X-1943-08209-8/S0033-569X-1943-08209-8.pdf>.
- [Dur10] R. Durrett. ‘*Probability: Theory and Examples, 4th Edition*’. Cambridge University Press, 2010. doi: [10.1017/CBO9780511779398](https://doi.org/10.1017/CBO9780511779398).
- [ECL11] P. M. Esfahani, D. Chatterjee and J. Lygeros. ‘*On a problem of stochastic reach-avoid set characterization*’. CDC/ECC. IEEE, 2011, pages 7069–7074. doi: [10.1109/CDC.2011.6160403](https://doi.org/10.1109/CDC.2011.6160403).
- [EK18] P. M. Esfahani and D. Kuhn. ‘*Data-driven distributionally robust optimization using the Wasserstein metric: performance guarantees and tractable reformulations*’. *Math. Program.* 171.1-2 (2018), pages 115–166. doi: [10.1007/S10107-017-1172-1](https://doi.org/10.1007/S10107-017-1172-1).
- [ESL15] P. M. Esfahani, T. Sutter and J. Lygeros. ‘*Performance Bounds for the Scenario Approach and an Extension to a Class of Non-Convex Programs*’. *IEEE Trans. Autom. Control*. 60.1 (2015), pages 46–58. doi: [10.1109/TAC.2014.2330702](https://doi.org/10.1109/TAC.2014.2330702).
- [FCGA21] X. Fang, R. Calinescu, S. Gerasimou and F. Alhwikem. ‘*Fast Parametric Model Checking through Model Fragmentation*’. ICSE. IEEE, 2021, pages 835–846. doi: [10.1109/ICSE43902.2021.00081](https://doi.org/10.1109/ICSE43902.2021.00081).
- [FCTS15] J. F. Fisac, M. Chen, C. J. Tomlin and S. S. Sastry. ‘*Reach-avoid problems with time-varying dynamics, targets and constraints*’. HSCC. ACM, 2015, pages 11–20. doi: [10.1145/2728606.2728612](https://doi.org/10.1145/2728606.2728612).
- [FH94] M. C. Fu and J. Hu. ‘*Smoothed perturbation analysis derivative estimation for Markov chains*’. *Oper. Res. Lett.* 15.5 (1994), pages 241–251. doi: [10.1016/0167-6377\(94\)90084-1](https://doi.org/10.1016/0167-6377(94)90084-1).
- [FHHW⁺11] M. Fränzle, E. M. Hahn, H. Hermanns, N. Wolovick and L. Zhang. ‘*Measurability and safety verification for stochastic hybrid systems*’. HSCC. ACM, 2011, pages 43–52. doi: [10.1145/1967701.1967710](https://doi.org/10.1145/1967701.1967710).
- [FKLX⁺18] Y. Feng, J. Katoen, H. Li, B. Xia and N. Zhan. ‘*Monitoring CTMCs by Multi-clock Timed Automata*’. CAV (1). Volume 10981. Lecture Notes in Computer Science. Springer, 2018, pages 507–526. doi: [10.1007/978-3-319-96145-3_27](https://doi.org/10.1007/978-3-319-96145-3_27).

- [FPE21] G. F. Franklin, J. D. Powell and A. Emami-Naeini. ‘*Feedback control of dynamic systems*’. Volume 8. Pearson, 2021. URL: <https://www.pearson.com/en-us/subject-catalog/p/feedback-control-of-dynamic-systems/P200000003343/9780137516834>.
- [FPK14] N. Ferns, D. Precup and S. Knight. ‘*Bisimulation for Markov Decision Processes through Families of Functional Expressions*’. *Horizons of the Mind*. Volume 8464. Lecture Notes in Computer Science. Springer, 2014, pages 319–342. doi: [10.1007/978-3-319-06880-0_17](https://doi.org/10.1007/978-3-319-06880-0_17).
- [FQMN⁺22] C. Fan, Z. Qin, U. Mathur, Q. Ning, S. Mitra and M. Viswanathan. ‘*Controller Synthesis for Linear System With Reach-Avoid Specifications*’. *IEEE Trans. Autom. Control*. 67.4 (2022), pages 1713–1727. doi: [10.1109/TAC.2021.3069723](https://doi.org/10.1109/TAC.2021.3069723).
- [FT14] J. Fu and U. Topcu. ‘*Probably Approximately Correct MDP Learning and Control With Temporal Logic Constraints*’. *Robotics: Science and Systems*. 2014. doi: [10.15607/RSS.2014.X.039](https://doi.org/10.15607/RSS.2014.X.039).
- [FTG16] A. Filieri, G. Tamburrelli and C. Ghezzi. ‘*Supporting Self-Adaptation via Quantitative Verification and Sensitivity Analysis at Run Time*’. *IEEE Trans. Software Eng.* 42.1 (2016), pages 75–99. doi: [10.1109/TSE.2015.2421318](https://doi.org/10.1109/TSE.2015.2421318).
- [FÜ11] C. R. Fox and G. Ülkümen. ‘*Distinguishing two dimensions of uncertainty*’. *Essays in Judgment and Decision Making* (2011). doi: [10.2139/ssrn.3695311](https://doi.org/10.2139/ssrn.3695311).
- [Fuk21] K. Fukuda. ‘*Cddlib reference manual*’. Report version 094m, McGill University, Montréal, Quebec, Canada (2021). URL: https://people.inf.ethz.ch/fukudak/cdd_home/cddlibman2021.pdf.
- [GB20] M. Gaon and R. I. Brafman. ‘*Reinforcement Learning with Non-Markovian Rewards*’. AAAI. AAAI Press, 2020, pages 3980–3987. doi: [10.1609/AAAI.V34I04.5814](https://doi.org/10.1609/AAAI.V34I04.5814).
- [GBLL24] I. Gracia, D. Boskos, L. Laurenti and M. Lahijanian. ‘*Data-driven strategy synthesis for stochastic systems with unknown nonlinear disturbances*’. L4DC. Volume 242. Proceedings of Machine Learning Research. PMLR, 2024, pages 1633–1645. doi: [10.48550/arXiv.2406.09704](https://doi.org/10.48550/arXiv.2406.09704).
- [GC06] J. C. Geromel and P. Colaneri. ‘*Robust stability of time varying polytopic systems*’. *Syst. Control. Lett.* 55.1 (2006), pages 81–85. doi: [10.1016/J.SYSCONLE.2004.11.016](https://doi.org/10.1016/J.SYSCONLE.2004.11.016).
- [GC22] S. Garatti and M. C. Campi. ‘*Risk and complexity in scenario optimization*’. *Math. Program.* 191.1 (2022), pages 243–279. doi: [10.1007/S10107-019-01446-4](https://doi.org/10.1007/S10107-019-01446-4).
- [GDG03] R. Givan, T. L. Dean and M. Greig. ‘*Equivalence notions and model minimization in Markov decision processes*’. *Artif. Intell.* 147.1-2 (2003), pages 163–223. doi: [10.1016/S0004-3702\(02\)00376-4](https://doi.org/10.1016/S0004-3702(02)00376-4).

- [GF15] J. García and F. Fernández. ‘*A comprehensive survey on safe reinforcement learning*’. *J. Mach. Learn. Res.* 16 (2015), pages 1437–1480. doi: [10.5555/2789272.2886795](https://doi.org/10.5555/2789272.2886795).
- [GG23] V. Goyal and J. Grand-Clément. ‘*Robust Markov Decision Processes: Beyond Rectangularity*’. *Math. Oper. Res.* 48.1 (2023), pages 203–226. doi: [10.1287/MOOR.2022.1259](https://doi.org/10.1287/MOOR.2022.1259).
- [GH09] X. Guo and O. Hernández-Lerma. ‘*Continuous-Time Markov Decision Processes*’. In *Continuous-Time Markov Decision Processes: Theory and Applications*. Springer Berlin Heidelberg, 2009, pages 9–18. doi: [10.1007/978-3-642-02547-1_2](https://doi.org/10.1007/978-3-642-02547-1_2).
- [GHHK⁺20] T. P. Gros, H. Hermanns, J. Hoffmann, M. Klauck and M. Steinmetz. ‘*Deep Statistical Model Checking*’. *FORTE*. Volume 12136. Lecture Notes in Computer Science. Springer, 2020, pages 96–114. doi: [10.1007/978-3-03-0086-3_6](https://doi.org/10.1007/978-3-03-0086-3_6).
- [GHS18] P. Gainer, E. M. Hahn and S. Schewe. ‘*Incremental Verification of Parametric and Reconfigurable Markov Chains*’. Lecture Notes in Computer Science 11024 (2018), pages 140–156. doi: [10.1007/978-3-319-99154-2_9](https://doi.org/10.1007/978-3-319-99154-2_9).
- [GHZZ13] Y. Gao, E. M. Hahn, N. Zhan and L. Zhang. ‘*CCMC: A Conditional CSL Model Checker for Continuous-Time Markov Chains*’. *ATVA*. Volume 8172. Lecture Notes in Computer Science. Springer, 2013, pages 464–468. doi: [10.1007/978-3-319-02444-8_36](https://doi.org/10.1007/978-3-319-02444-8_36).
- [GKM10] C. Goerzen, Z. Kong and B. Mettler. ‘*A Survey of Motion Planning Algorithms from the Perspective of Autonomous UAV Guidance*’. *J. Intell. Robotic Syst.* 57.1-4 (2010), pages 65–100. doi: [10.1007/S10846-009-9383-1](https://doi.org/10.1007/S10846-009-9383-1).
- [GKP01] C. Guestrin, D. Koller and R. Parr. ‘*Multiagent Planning with Factored MDPs*’. *NIPS*. MIT Press, 2001, pages 1523–1530. URL: https://proceedings.neurips.cc/paper_files/paper/2001/file/7af6266cc52234b5aa339b16695f7fc4-Paper.pdf.
- [GKPV03] C. Guestrin, D. Koller, R. Parr and S. Venkataraman. ‘*Efficient Solution Algorithms for Factored MDPs*’. *J. Artif. Intell. Res.* 19 (2003), pages 399–468. doi: [10.1613/JAIR.1000](https://doi.org/10.1613/JAIR.1000).
- [GLD00] R. Givan, S. M. Leach and T. L. Dean. ‘*Bounded-parameter Markov decision processes*’. *Artif. Intell.* 122.1-2 (2000), pages 71–109. doi: [10.1016/S0004-3702\(00\)00047-3](https://doi.org/10.1016/S0004-3702(00)00047-3).
- [GLMA⁺24] I. Gracia, L. Laurenti, M. Mazo Jr., A. Abate and M. Lahijanian. ‘*Temporal Logic Control for Nonlinear Stochastic Systems Under Unknown Disturbances*’ . 2024. doi: [10.48550/arXiv.2412.11343](https://doi.org/10.48550/arXiv.2412.11343).
- [GM84] D. Gross and D. R. Miller. ‘*The Randomization Technique as a Modeling Tool and Solution Procedure for Transient Markov Processes*’. *Oper. Res.* 32.2 (1984), pages 343–361. doi: [10.1287/OPRE.32.2.343](https://doi.org/10.1287/OPRE.32.2.343).

- [GMT14] V. Gabrel, C. Murat and A. Thiele. ‘Recent advances in robust optimization: An overview’. *Eur. J. Oper. Res.* 235.3 (2014), pages 471–483. doi: [10.1016/J.EJOR.2013.09.036](https://doi.org/10.1016/J.EJOR.2013.09.036).
- [GP09] A. Girard and G. J. Pappas. ‘Hierarchical control system design using approximate simulation’. *Autom.* 45.2 (2009), pages 566–571. doi: [10.1016/J.AUTOMATICA.2008.09.016](https://doi.org/10.1016/J.AUTOMATICA.2008.09.016).
- [GS10] J. Goh and M. Sim. ‘Distributionally Robust Optimization and Its Tractable Approximations’. *Oper. Res.* 58.4-Part-1 (2010), pages 902–917. doi: [10.1287/OPRE.1090.0795](https://doi.org/10.1287/OPRE.1090.0795).
- [GSD12] A. Guez, D. Silver and P. Dayan. ‘Efficient Bayes-Adaptive Reinforcement Learning using Sample-Based Search’. *NIPS*. 2012, pages 1034–1042. doi: [10.48550/arXiv.1205.3109](https://doi.org/10.48550/arXiv.1205.3109).
- [Gur23] Gurobi Optimization, LLC. ‘Gurobi Optimizer Reference Manual’. 2023. URL: <https://www.gurobi.com>.
- [GXZZ13] Y. Gao, M. Xu, N. Zhan and L. Zhang. ‘Model checking conditional CSL for continuous-time Markov chains’. *Inf. Process. Lett.* 113.1-2 (2013), pages 44–50. doi: [10.1016/J.IPL.2012.09.009](https://doi.org/10.1016/J.IPL.2012.09.009).
- [GY07] M. J. F. Gales and S. J. Young. ‘The Application of Hidden Markov Models in Speech Recognition’. *Found. Trends Signal Process.* 1.3 (2007), pages 195–304. doi: [10.1561/2000000004](https://doi.org/10.1561/2000000004).
- [GY22] J. Guan and N. Yu. ‘A Probabilistic Logic for Verifying Continuous-time Markov Chains’. *TACAS* (2). Volume 13244. Lecture Notes in Computer Science. Springer, 2022, pages 3–21. doi: [10.1007/978-3-030-99527-0_1](https://doi.org/10.1007/978-3-030-99527-0_1).
- [GZMG⁺16] S. Grammatico, X. Zhang, K. Margellos, P. J. Goulart and J. Lygeros. ‘A Scenario Approach for Non-Convex Control Design’. *IEEE Trans. Autom. Control*. 61.2 (2016), pages 334–345. doi: [10.1109/TAC.2015.2433591](https://doi.org/10.1109/TAC.2015.2433591).
- [HBK17] L. Hutschenreiter, C. Baier and J. Klein. ‘Parametric Markov Chains: PCTL Complexity and Fraction-free Gaussian Elimination’. *GandALF*. Volume 256. EPTCS. 2017, pages 16–30. doi: [10.4204/EPTCS.256.2](https://doi.org/10.4204/EPTCS.256.2).
- [HCHB⁺17] S. L. Herbert, M. Chen, S. Han, S. Bansal, J. F. Fisac and C. J. Tomlin. ‘FaSTrack: A modular framework for fast and guaranteed safe motion planning’. *CDC*. IEEE, 2017, pages 1517–1522. doi: [10.1109/CDC.2017.8263867](https://doi.org/10.1109/CDC.2017.8263867).
- [Her02] H. Hermanns. ‘Interactive Markov Chains: The Quest for Quantified Quality’. Volume 2428. Lecture Notes in Computer Science. Springer, 2002. doi: [10.1007/3-540-45804-2](https://doi.org/10.1007/3-540-45804-2).
- [HH12] H. Hatefi and H. Hermanns. ‘Model Checking Algorithms for Markov Automata’. *Electron. Commun. Eur. Assoc. Softw. Sci. Technol.* 53 (2012). doi: [10.14279/TUJ.ECEASST.53.783](https://doi.org/10.14279/TUJ.ECEASST.53.783).
- [HHA17] S. Haesaert, P. M. J. V. den Hof and A. Abate. ‘Data-driven and model-based verification via Bayesian identification and reachability analysis’. *Autom.* 79 (2017), pages 115–126. doi: [10.1016/J.AUTOMATICA.2017.01.037](https://doi.org/10.1016/J.AUTOMATICA.2017.01.037).

- [HHHT16] E. M. Hahn, V. Hashemi, H. Hermanns and A. Turrini. ‘*Exploiting Robust Optimization for Interval Probabilistic Bisimulation*’. *QEST*. Volume 9826. Lecture Notes in Computer Science. Springer, 2016, pages 55–71. doi: [10.1007/978-3-319-43425-4_4](https://doi.org/10.1007/978-3-319-43425-4_4).
- [HHK00] B. R. Haverkort, H. Hermanns and J. Katoen. ‘*On the Use of Model Checking Techniques for Dependability Evaluation*’. *SRDS*. IEEE Computer Society, 2000, pages 228–237. doi: [10.1109/RELDI.2000.885410](https://doi.org/10.1109/RELDI.2000.885410).
- [HHSS⁺16] V. Hashemi, H. Hermanns, L. Song, K. Subramani, A. Turrini and P. Wojciechowski. ‘*Compositional Bisimulation Minimization for Interval Markov Decision Processes*’. *LATA*. Volume 9618. Lecture Notes in Computer Science. Springer, 2016, pages 114–126. doi: [10.1007/978-3-319-30000-9_9](https://doi.org/10.1007/978-3-319-30000-9_9).
- [HHWZ10] E. M. Hahn, H. Hermanns, B. Wachter and L. Zhang. ‘*PASS: Abstraction Refinement for Infinite Probabilistic Models*’. *TACAS*. Volume 6015. Lecture Notes in Computer Science. Springer, 2010, pages 353–357. doi: [10.1007/978-3-642-12002-2_30](https://doi.org/10.1007/978-3-642-12002-2_30).
- [HHZ11a] E. M. Hahn, T. Han and L. Zhang. ‘*Synthesis for PCTL in Parametric Markov Decision Processes*’. *NASA Formal Methods*. Volume 6617. Lecture Notes in Computer Science. Springer, 2011, pages 146–161. doi: [10.1007/978-3-642-20398-5_12](https://doi.org/10.1007/978-3-642-20398-5_12).
- [HHZ11b] E. M. Hahn, H. Hermanns and L. Zhang. ‘*Probabilistic reachability for parametric Markov models*’. *Int. J. Softw. Tools Technol. Transf.* 13.1 (2011), pages 3–19. doi: [10.1007/S10009-010-0146-X](https://doi.org/10.1007/S10009-010-0146-X).
- [HJ02] J. Han and P. Jonker. ‘*A System Architecture Solution for Unreliable Nanoelectronic Devices*’. *IEEE Transactions on Nanotechnology* 1 (2002), pages 201–208. doi: [10.1109/TNANO.2002.807393](https://doi.org/10.1109/TNANO.2002.807393).
- [HJ94] H. Hansson and B. Jonsson. ‘*A Logic for Reasoning about Time and Reliability*’. *Formal Aspects Comput.* 6.5 (1994), pages 512–535. doi: [10.1007/BF01211866](https://doi.org/10.1007/BF01211866).
- [HJKQ⁺22] C. Hensel, S. Junges, J. Katoen, T. Quatmann and M. Volk. ‘*The probabilistic model checker Storm*’. *Int. J. Softw. Tools Technol. Transf.* 24.4 (2022), pages 589–610. doi: [10.1007/S10009-021-00633-Z](https://doi.org/10.1007/S10009-021-00633-Z).
- [HJQW23] A. Hartmanns, S. Junges, T. Quatmann and M. Weininger. ‘*A Practitioner’s Guide to MDP Model Checking Algorithms*’. *TACAS (1)*. Volume 13993. Lecture Notes in Computer Science. Springer, 2023, pages 469–488. doi: [10.1007/978-3-031-30823-9_24](https://doi.org/10.1007/978-3-031-30823-9_24).
- [HK10] A. J. Hoffman and J. B. Kruskal. ‘*Integral Boundary Points of Convex Polyhedra*’. *50 Years of Integer Programming*. Springer, 2010, pages 49–76. doi: [10.1007/978-3-540-68279-0_3](https://doi.org/10.1007/978-3-540-68279-0_3).
- [HKK14] H. Hermanns, J. Krcál and J. Kretínský. ‘*Probabilistic Bisimulation: Naturally on Distributions*’. *CONCUR*. Volume 8704. Lecture Notes in Computer Science. Springer, 2014, pages 249–265. doi: [10.1007/978-3-662-44584-6_18](https://doi.org/10.1007/978-3-662-44584-6_18).

- [HKM08] T. Han, J. Katoen and A. Mereacre. ‘*Approximate Parameter Synthesis for Probabilistic Time-Bounded Reachability*’. RTSS. IEEE Computer Society, 2008, pages 173–182. doi: [10.1109/RTSS.2008.19](https://doi.org/10.1109/RTSS.2008.19).
- [HKPQ⁺19] A. Hartmanns, M. Klauck, D. Parker, T. Quatmann and E. Ruijters. ‘*The Quantitative Verification Benchmark Set*’. TACAS (1). Volume 11427. Lecture Notes in Computer Science. Springer, 2019, pages 344–350. doi: [10.1007/978-3-030-17462-0_20](https://doi.org/10.1007/978-3-030-17462-0_20).
- [HMS99] H. Hermanns, J. Meyer-Kayser and M. Siegle. ‘*Multi terminal binary decision diagrams to represent and analyse continuous time Markov chains*’. 3rd Int. Workshop on the Numerical Solution of Markov Chains. Citeseer. 1999, pages 188–207. URL: <https://www.unibw.de/technische-informatik/mitarbeiter/professoren/siegle/publikationen/nsmc99.pdf>.
- [HNPW⁺11] E. M. Hahn, G. Norman, D. Parker, B. Wachter and L. Zhang. ‘*Game-based Abstraction and Controller Synthesis for Probabilistic Hybrid Systems*’. QEST. IEEE Computer Society, 2011, pages 69–78. doi: [10.1109/QEST.2011.17](https://doi.org/10.1109/QEST.2011.17).
- [HNVT⁺18] S. Haesaert, P. Nilsson, C. I. Vasile, R. Thakker, A. Agha-mohammadi, A. D. Ames and R. M. Murray. ‘*Temporal Logic Control of POMDPs via Label-based Stochastic Simulation Relations*’. ADHS. Volume 51. IFAC-PapersOnLine 16. Elsevier, 2018, pages 271–276. doi: [10.1016/J.IFACOL.2018.08.046](https://doi.org/10.1016/J.IFACOL.2018.08.046).
- [HPSW⁺11] H. Hermanns, A. Parma, R. Segala, B. Wachter and L. Zhang. ‘*Probabilistic Logical Characterization*’. Inf. Comput. 209.2 (2011), pages 154–172. doi: [10.1016/J.IC.2010.11.024](https://doi.org/10.1016/J.IC.2010.11.024).
- [HPW18] C. P. Ho, M. Petrik and W. Wiesemann. ‘*Fast Bellman Updates for Robust MDPs*’. ICML. Volume 80. Proceedings of Machine Learning Research. PMLR, 2018, pages 1984–1993. URL: <https://proceedings.mlr.press/v80/ho18a.html>.
- [HRW12] J. Humpherys, P. Redd and J. M. West. ‘*A Fresh Look at the Kalman Filter*’. SIAM Rev. 54.4 (2012), pages 801–823. doi: [10.1137/100799666](https://doi.org/10.1137/100799666).
- [HS09] A. Hobolth and E. A. Stone. ‘*Simulation from endpoint-conditioned, continuous-time Markov chains on a finite state space, with applications to molecular evolution*’. The annals of applied statistics 3.3 (2009), page 1204. doi: [10.1214/09-AOAS247](https://doi.org/10.1214/09-AOAS247).
- [HSA17] S. Haesaert, S. E. Z. Soudjani and A. Abate. ‘*Verification of General Markov Decision Processes by Approximate Similarity Relations and Policy Refinement*’. SIAM J. Control. Optim. 55.4 (2017), pages 2333–2367. doi: [10.1137/16M1079397](https://doi.org/10.1137/16M1079397).
- [HSJM⁺22] L. Heck, J. Spel, S. Junges, J. Moerman and J. Katoen. ‘*Gradient-Descent for Randomized Controllers Under Partial Observability*’. VMCAI. Volume 13182. Lecture Notes in Computer Science. Springer, 2022, pages 127–150. doi: [10.1007/978-3-030-94583-1_7](https://doi.org/10.1007/978-3-030-94583-1_7).

- [HSV93] L. Helmink, M. P. A. Sellink and F. W. Vaandrager. ‘*Proof-Checking a Data Link Protocol*’. *TYPES*. Volume 806. Lecture Notes in Computer Science. Springer, 1993, pages 127–165. doi: [10.1007/3-540-58085-9_75](https://doi.org/10.1007/3-540-58085-9_75).
- [HW21] E. Hüllermeier and W. Waegeman. ‘*Aleatoric and epistemic uncertainty in machine learning: an introduction to concepts and methods*’. *Mach. Learn.* 110.3 (2021), pages 457–506. doi: [10.1007/S10994-021-05946-3](https://doi.org/10.1007/S10994-021-05946-3).
- [HW93] D. Haussler and M. Warmuth. ‘*The Probably Approximately Correct (PAC) and Other Learning Models*’. *Foundations of Knowledge Acquisition: Machine Learning*. Edited by A. L. Meyrowitz and S. Chipman. Springer US, 1993, pages 291–312. doi: [10.1007/978-0-585-27366-2_9](https://doi.org/10.1007/978-0-585-27366-2_9).
- [IKVM18] R. T. Icarte, T. Q. Klassen, R. A. Valenzano and S. A. McIlraith. ‘*Using Reward Machines for High-Level Task Specification and Decomposition in Reinforcement Learning*’. *ICML*. Volume 80. Proceedings of Machine Learning Research. PMLR, 2018, pages 2112–2121. url: <http://proceedings.mlr.press/v80/icarte18a.html>.
- [IKVM22] R. T. Icarte, T. Q. Klassen, R. A. Valenzano and S. A. McIlraith. ‘*Reward Machines: Exploiting Reward Function Structure in Reinforcement Learning*’. *J. Artif. Intell. Res.* 73 (2022), pages 173–208. doi: [10.1613/JAIR.1.12440](https://doi.org/10.1613/JAIR.1.12440).
- [Iye05] G. N. Iyengar. ‘*Robust Dynamic Programming*’. *Math. Oper. Res.* 30.2 (2005), pages 257–280. doi: [10.1287/MOOR.1040.0129](https://doi.org/10.1287/MOOR.1040.0129).
- [JÁHJ⁺24] S. Junges, E. Ábrahám, C. Hensel, N. Jansen, J. Katoen, T. Quatmann and M. Volk. ‘*Parameter synthesis for Markov models: covering the parameter space*’. *Formal Methods Syst. Des.* 62.1 (2024), pages 181–259. doi: [10.1007/S10703-023-00442-X](https://doi.org/10.1007/S10703-023-00442-X).
- [JBBA21] K. Jothimurugan, S. Bansal, O. Bastani and R. Alur. ‘*Compositional Reinforcement Learning from Logical Specifications*’. *NeurIPS*. 2021, pages 10026–10039. doi: [10.48550/arXiv.2106.13906](https://doi.org/10.48550/arXiv.2106.13906).
- [Jen53] A. Jensen. ‘*Markoff chains as an aid in the study of Markoff processes*’. *Scandinavian Actuarial Journal* 3 (1953), pages 87–91. doi: [10.1080/03461238.1953.10419459](https://doi.org/10.1080/03461238.1953.10419459).
- [JJK22] N. Jansen, S. Junges and J. Katoen. ‘*Parameter Synthesis in Markov Models: A Gentle Survey*’. *Principles of Systems Design*. Volume 13660. Lecture Notes in Computer Science. Springer, 2022, pages 407–437. doi: [10.1007/978-3-031-22337-2_20](https://doi.org/10.1007/978-3-031-22337-2_20).
- [JKPW21] S. Junges, J. Katoen, G. A. Pérez and T. Winkler. ‘*The complexity of reachability in parametric Markov decision processes*’. *J. Comput. Syst. Sci.* 119 (2021), pages 183–210. doi: [10.1016/J.JCSS.2021.02.006](https://doi.org/10.1016/J.JCSS.2021.02.006).
- [JL91] B. Jonsson and K. G. Larsen. ‘*Specification and Refinement of Probabilistic Processes*’. *LICS*. IEEE Computer Society, 1991, pages 266–277. doi: [10.1109/LICS.1991.151651](https://doi.org/10.1109/LICS.1991.151651).

- [JLFL20] J. Jackson, L. Laurenti, E. W. Frew and M. Lahijanian. ‘*Safety Verification of Unknown Dynamical Systems via Gaussian Process Regression*’. *CDC*. IEEE, 2020, pages 860–866. doi: [10.1109/CDC42340.2020.9303814](https://doi.org/10.1109/CDC42340.2020.9303814).
- [JTS21] S. Junges, H. Torfah and S. A. Seshia. ‘*Runtime Monitors for Markov Decision Processes*’. *CAV (2)*. Volume 12760. Lecture Notes in Computer Science. Springer, 2021, pages 553–576. doi: [10.1007/978-3-030-81688-9_26](https://doi.org/10.1007/978-3-030-81688-9_26).
- [Jun20] S. Junges. ‘*Parameter synthesis in Markov models*’. PhD thesis. RWTH Aachen University, Germany, 2020. doi: [10.18154/RWTH-2020-02348](https://doi.org/10.18154/RWTH-2020-02348).
- [Kak03] S. M. Kakade. ‘*On the sample complexity of reinforcement learning*’. PhD thesis. University of London, University College London (United Kingdom), 2003. URL: https://homes.cs.washington.edu/~sham/papers/thesis/sham_thesis.pdf.
- [Kal02] O. Kallenberg. ‘*Foundations of modern probability*’. Probability and its applications. Springer, 2002. doi: [10.1007/978-3-030-61871-1](https://doi.org/10.1007/978-3-030-61871-1).
- [Kal60] R. E. Kalman. ‘*A new approach to linear filtering and prediction problems*’. *Journal of Fluids Engineering, Transactions of the ASME* 82.1 (1960), pages 35–45. doi: [10.1115/1.3662552](https://doi.org/10.1115/1.3662552).
- [Kat16] J. Katoen. ‘*The Probabilistic Model Checking Landscape*’. *LICS*. ACM, 2016, pages 31–45. doi: [10.1145/2933575.2934574](https://doi.org/10.1145/2933575.2934574).
- [KBJT19] J. Kenanian, A. Balkan, R. M. Jungers and P. Tabuada. ‘*Data driven stability analysis of black-box switched linear systems*’. *Autom.* 109 (2019). doi: [10.1016/J.AUTOMATICA.2019.108533](https://doi.org/10.1016/J.AUTOMATICA.2019.108533).
- [KCOB21] C. Knuth, G. Chou, N. Ozay and D. Berenson. ‘*Planning With Learned Dynamics: Probabilistic Guarantees on Safety and Reachability via Lipschitz Constants*’. *IEEE Robotics Autom. Lett.* 6.3 (2021), pages 5129–5136. doi: [10.1109/LRA.2021.3068889](https://doi.org/10.1109/LRA.2021.3068889).
- [KFM04] C. C. T. Kwok, D. Fox and M. Meila. ‘*Real-time particle filters*’. *Proc. IEEE* 92.3 (2004), pages 469–484. doi: [10.1109/JPROC.2003.823144](https://doi.org/10.1109/JPROC.2003.823144).
- [KG02] H. K. Khalil and J. W. Grizzle. ‘*Nonlinear systems*’. Volume 3. Prentice hall Upper Saddle River, NJ, 2002.
- [KGS07] B. T. Kulakowski, J. F. Gardner and J. L. Shearer. ‘*Dynamic modeling and control of engineering systems*’. Cambridge University Press, 2007. doi: [10.1017/CBO9780511805417](https://doi.org/10.1017/CBO9780511805417).
- [KIKF22] M. Koegel, M. Ibrahim, C. Kallies and R. Findeisen. ‘*Safe Hierarchical Model Predictive Control and Planning for Autonomous Systems*’. *CoRR* abs/2203.14269 (2022). doi: [10.48550/arXiv.2203.14269](https://doi.org/10.48550/arXiv.2203.14269).
- [KKNP01] J. Katoen, M. Z. Kwiatkowska, G. Norman and D. Parker. ‘*Faster and Symbolic CTMC Model Checking*’. *PAPM-PROBMIV*. Volume 2165. Lecture Notes in Computer Science. Springer, 2001, pages 23–38. doi: [10.1007/3-540-44804-7_2](https://doi.org/10.1007/3-540-44804-7_2).

- [KKNP10] M. Kattenbelt, M. Z. Kwiatkowska, G. Norman and D. Parker. ‘*A game-based abstraction-refinement framework for Markov decision processes*’. *Formal Methods Syst. Des.* 36.3 (2010), pages 246–280. doi: [10.1007/S10703-010-0097-6](https://doi.org/10.1007/S10703-010-0097-6).
- [KKR16] L. Koreniak, A. Kucera and V. Rehák. ‘*Efficient Timeout Synthesis in Fixed-Delay CTMC Using Policy Iteration*’. *MASCOTS*. IEEE Computer Society, 2016, pages 367–372. doi: [10.1109/MASCOTS.2016.34](https://doi.org/10.1109/MASCOTS.2016.34).
- [KLC98] L. P. Kaelbling, M. L. Littman and A. R. Cassandra. ‘*Planning and Acting in Partially Observable Stochastic Domains*’. *Artif. Intell.* 101.1-2 (1998), pages 99–134. doi: [10.1016/S0004-3702\(98\)00023-X](https://doi.org/10.1016/S0004-3702(98)00023-X).
- [KLM96] L. P. Kaelbling, M. L. Littman and A. W. Moore. ‘*Reinforcement Learning: A Survey*’. *J. Artif. Intell. Res.* 4 (1996), pages 237–285. doi: [10.1613/JAIR.301](https://doi.org/10.1613/JAIR.301).
- [KNP02] M. Z. Kwiatkowska, G. Norman and D. Parker. ‘*PRISM: Probabilistic Symbolic Model Checker*’. *Computer Performance Evaluation / TOOLS*. Volume 2324. Lecture Notes in Computer Science. Springer, 2002, pages 200–204. doi: [10.1007/3-540-46029-2_13](https://doi.org/10.1007/3-540-46029-2_13).
- [KNP06] M. Z. Kwiatkowska, G. Norman and A. Pacheco. ‘*Model checking expected time and expected reward formulae with random time bounds*’. *Comput. Math. Appl.* 51.2 (2006), pages 305–316. doi: [10.1016/J.CAMWA.2005.11.016](https://doi.org/10.1016/J.CAMWA.2005.11.016).
- [KNP11] M. Z. Kwiatkowska, G. Norman and D. Parker. ‘*PRISM 4.0: Verification of Probabilistic Real-Time Systems*’. *CAV*. Volume 6806. Lecture Notes in Computer Science. Springer, 2011, pages 585–591. doi: [10.1007/978-3-642-22110-1_47](https://doi.org/10.1007/978-3-642-22110-1_47).
- [KNP12] M. Z. Kwiatkowska, G. Norman and D. Parker. ‘*The PRISM Benchmark Suite*’. *QEST*. IEEE Computer Society, 2012, pages 203–204. doi: [10.1109/QEST.2012.14](https://doi.org/10.1109/QEST.2012.14).
- [KOA17] S. Katt, F. A. Oliehoek and C. Amato. ‘*Learning in POMDPs with Monte Carlo Tree Search*’. *ICML*. Volume 70. Proceedings of Machine Learning Research. PMLR, 2017, pages 1819–1827. url: <http://proceedings.mlr.press/v70/katt17a.html>.
- [Kra07] J. Kramer. ‘*Is abstraction the key to computing?*’ *Commun. ACM* 50.4 (2007), pages 36–42. doi: [10.1145/1232743.1232745](https://doi.org/10.1145/1232743.1232745).
- [KS02] M. J. Kearns and S. Singh. ‘*Near-Optimal Reinforcement Learning in Polynomial Time*’. *Mach. Learn.* 49.2-3 (2002), pages 209–232. doi: [10.1023/A:1017984413808](https://doi.org/10.1023/A:1017984413808).
- [LAB15] M. Lahijanian, S. B. Andersson and C. Belta. ‘*Formal Verification and Synthesis for Discrete-Time Stochastic Systems*’. *IEEE Trans. Autom. Control*. 60.8 (2015), pages 2031–2045. doi: [10.1109/TAC.2015.2398883](https://doi.org/10.1109/TAC.2015.2398883).
- [LaV06] S. M. LaValle. ‘*Planning Algorithms*’. Cambridge University Press, 2006. doi: [10.1017/CBO9780511546877](https://doi.org/10.1017/CBO9780511546877).

- [LB02] C. M. Lagoa and B. R. Barmish. ‘*Distributionally robust Monte Carlo simulation: A tutorial survey*’. *IFAC Proceedings Volumes* 35.1 (2002), pages 151–162. doi: [10.3182/20020721-6-ES-1901.00360](https://doi.org/10.3182/20020721-6-ES-1901.00360).
- [LB16] T. Lipp and S. Boyd. ‘*Variations and extension of the convex-concave procedure*’. *Optimization and Engineering* 17 (2016), pages 263–287. doi: [10.1007/s11081-015-9294-x](https://doi.org/10.1007/s11081-015-9294-x).
- [LBBS⁺18] Y. R. S. Llerena, M. Böhme, M. Brünink, G. Su and D. S. Rosenblum. ‘*Verifying the long-run behavior of probabilistic system models in the presence of uncertainty*’. *ESEC/SIGSOFT FSE*. ACM, 2018, pages 587–597. doi: [10.1145/3236024.3236078](https://doi.org/10.1145/3236024.3236078).
- [LDB10] A. Legay, B. Delahaye and S. Bensalem. ‘*Statistical Model Checking: An Overview*’. *RV*. Volume 6418. Lecture Notes in Computer Science. Springer, 2010, pages 122–135. doi: [10.1007/978-3-642-16612-9_11](https://doi.org/10.1007/978-3-642-16612-9_11).
- [LKSZ20] A. Lavaei, M. Khaled, S. Soudjani and M. Zamani. ‘*AMYTISS: Parallelized Automated Controller Synthesis for Large-Scale Stochastic Systems*’. *CAV* (2). Volume 12225. Lecture Notes in Computer Science. Springer, 2020, pages 461–474. doi: [10.1007/978-3-030-53291-8_24](https://doi.org/10.1007/978-3-030-53291-8_24).
- [LLTY⁺19] A. Legay, A. Lukina, L. Traonouez, J. Yang, S. A. Smolka and R. Grosu. ‘*Statistical Model Checking*’. *Computing and Software Science*. Volume 10000. Lecture Notes in Computer Science. Springer, 2019, pages 478–504. doi: [10.1007/978-3-319-91908-9_23](https://doi.org/10.1007/978-3-319-91908-9_23).
- [LMT07] R. Lanotte, A. Maggiolo-Schettini and A. Troina. ‘*Parametric probabilistic transition systems for system design and analysis*’. *Formal Aspects Comput.* 19.1 (2007), pages 93–109. doi: [10.1007/S00165-006-0015-2](https://doi.org/10.1007/S00165-006-0015-2).
- [LOE13] K. Lesser, M. M. K. Oishi and R. S. Erwin. ‘*Stochastic reachability for control of spacecraft relative motion*’. *CDC*. IEEE, 2013, pages 4705–4712. doi: [10.1109/CDC.2013.6760626](https://doi.org/10.1109/CDC.2013.6760626).
- [Lov91] W. S. Lovejoy. ‘*Computationally Feasible Bounds for Partially Observed Markov Decision Processes*’. *Oper. Res.* 39.1 (1991), pages 162–175. doi: [10.1287/OPRE.39.1.162](https://doi.org/10.1287/OPRE.39.1.162).
- [LS91] K. G. Larsen and A. Skou. ‘*Bisimulation through Probabilistic Testing*’. *Inf. Comput.* 94.1 (1991), pages 1–28. doi: [10.1016/0890-5401\(91\)90030-6](https://doi.org/10.1016/0890-5401(91)90030-6).
- [LSAZ22] A. Lavaei, S. Soudjani, A. Abate and M. Zamani. ‘*Automated verification and synthesis of stochastic hybrid systems: A survey*’. *Autom.* 146 (2022), page 110617. doi: [10.1016/J.AUTOMATICA.2022.110617](https://doi.org/10.1016/J.AUTOMATICA.2022.110617).
- [LSFZ23] A. Lavaei, S. Soudjani, E. Frazzoli and M. Zamani. ‘*Constructing MDP Abstractions Using Data With Formal Guarantees*’. *IEEE Control. Syst. Lett.* 7 (2023), pages 460–465. doi: [10.1109/LCSYS.2022.3188535](https://doi.org/10.1109/LCSYS.2022.3188535).
- [LSS20] A. Loquercio, M. Segù and D. Scaramuzza. ‘*A General Framework for Uncertainty Estimation in Deep Learning*’. *IEEE Robotics Autom. Lett.* 5.2 (2020), pages 3153–3160. doi: [10.1109/LRA.2020.2974682](https://doi.org/10.1109/LRA.2020.2974682).

- [LV95] N. A. Lynch and F. W. Vaandrager. ‘Forward and Backward Simulations: I. Untimed Systems’. *Inf. Comput.* 121.2 (1995), pages 214–233. doi: [10.1006/INCO.1995.1134](#).
- [LWL06] L. Li, T. J. Walsh and M. L. Littman. ‘Towards a Unified Theory of State Abstraction for MDPs’. *AI&M*. 2006. URL: <http://anytime.cs.umass.edu/aimath06/proceedings/P21.pdf>.
- [LZCH22] M. Lechner, D. Zikelic, K. Chatterjee and T. A. Henzinger. ‘Stability Verification in Stochastic Control Systems via Neural Network Supermartingales’. *AAAI*. AAAI Press, 2022, pages 7326–7336. doi: [10.1609/AAAI.V36I7.20695](#).
- [Mar20] G. Marcus. ‘The Next Decade in AI: Four Steps Towards Robust Artificial Intelligence’. *CoRR* abs/2002.06177 (2020). doi: [10.48550/arXiv.2002.06177](#).
- [MBPJ23] T. M. Moerland, J. Broekens, A. Plaat and C. M. Jonker. ‘Model-based Reinforcement Learning: A Survey’. *Found. Trends Mach. Learn.* 16.1 (2023), pages 1–118. doi: [10.1561/2200000086](#).
- [MCL23] F. B. Mathiesen, S. C. Calvert and L. Laurenti. ‘Safety Certification for Stochastic Systems via Neural Barrier Functions’. *IEEE Control. Syst. Lett.* 7 (2023), pages 973–978. doi: [10.1109/LCSYS.2022.3229865](#).
- [MG07] J. Matousek and B. Gärtner. ‘Integer Programming and LP Relaxation’. *Understanding and Using Linear Programming*. Springer Berlin Heidelberg, 2007, pages 29–40. doi: [10.1007/978-3-540-30717-4_3](#).
- [MGF21] A. Makdesi, A. Girard and L. Fribourg. ‘Efficient Data-Driven Abstraction of Monotone Systems with Disturbances’. *ADHS*. Volume 54. IFAC-PapersOnLine 5. Elsevier, 2021, pages 49–54. doi: [10.1016/J.IFACOL.2021.08.473](#).
- [MGL14] K. Margellos, P. Goulart and J. Lygeros. ‘On the Road Between Robust Optimization and the Scenario Approach for Chance Constrained Optimization Problems’. *IEEE Trans. Autom. Control.* 59.8 (2014), pages 2258–2263. doi: [10.1109/TAC.2014.2303232](#).
- [MHC99] O. Madani, S. Hanks and A. Condon. ‘On the Undecidability of Probabilistic Planning and Infinite-Horizon Partially Observable Markov Decision Problems’. *AAAI/IAAI*. AAAI Press / The MIT Press, 1999, pages 541–548. doi: [10.1016/S0004-3702\(02\)00378-8](#).
- [Mil68] B. L. Miller. ‘Finite state continuous time Markov decision processes with a finite planning horizon’. *SIAM Journal on Control* 6.2 (1968), pages 266–280. doi: [10.1016/0022-247X\(68\)90194-7](#).
- [Mil71] R. Milner. ‘An Algebraic Definition of Simulation Between Programs’. *IJCAI*. William Kaufmann, 1971, pages 481–489. URL: <https://dl.acm.org/doi/abs/10.5555/1622876.1622926>.
- [Mil89] R. Milner. ‘Communication and concurrency’. PHI Series in computer science. Prentice Hall, 1989. URL: <https://dl.acm.org/doi/book/10.5555/534666>.

- [ML03] C. B. Moler and C. V. Loan. ‘*Nineteen Dubious Ways to Compute the Exponential of a Matrix, Twenty-Five Years Later*’. *SIAM Rev.* 45.1 (2003), pages 3–49. doi: [10.1137/S00361445024180](https://doi.org/10.1137/S00361445024180).
- [MMAG14] I. Meedeniya, I. Moser, A. Aleti and L. Grunske. ‘*Evaluating probabilistic models with uncertain model parameters*’. *Softw. Syst. Model.* 13.4 (2014), pages 1395–1415. doi: [10.1007/S10270-012-0277-5](https://doi.org/10.1007/S10270-012-0277-5).
- [Mob02] R. K. Mobley. ‘*An introduction to predictive maintenance*’. Elsevier, 2002. doi: [10.1016/B978-0-7506-7531-4.X5000-3](https://doi.org/10.1016/B978-0-7506-7531-4.X5000-3).
- [Mod22] H. Modares. ‘*Data-driven Safe Control of Linear Systems Under Epistemic and Aleatory Uncertainties*’. *CoRR* abs/2202.04495 (2022). doi: [10.48550/arXiv.2202.04495](https://doi.org/10.48550/arXiv.2202.04495).
- [MRTT53] T. S. Motzkin, H. Raiffa, G. L. Thompson and R. M. Thrall. ‘*The double description method*’. *Contributions to the Theory of Games* 2.28 (1953), pages 51–73. doi: [10.1515/9781400881970-004](https://doi.org/10.1515/9781400881970-004).
- [MSXA18] Y. Mao, M. Szmuk, X. Xu and B. Açikmese. ‘*Successive convexification: A superlinearly convergent algorithm for non-convex optimal control problems*’. *arXiv preprint* (2018). doi: [10.48550/arXiv.1804.06539](https://doi.org/10.48550/arXiv.1804.06539).
- [MT17] A. Majumdar and R. Tedrake. ‘*Funnel libraries for real-time robust feedback motion planning*’. *Int. J. Robotics Res.* 36.8 (2017), pages 947–982. doi: [10.1177/0278364917712421](https://doi.org/10.1177/0278364917712421).
- [Mun14] R. Munos. ‘*From Bandits to Monte-Carlo Tree Search: The Optimistic Principle Applied to Optimization and Planning*’. *Found. Trends Mach. Learn.* 7.1 (2014), pages 1–129. doi: [10.1561/2200000038](https://doi.org/10.1561/2200000038).
- [MWW24] T. Meggendorfer, M. Weininger and P. Wienhöft. ‘*What Are the Odds? Improving the foundations of Statistical Model Checking*’. *CoRR* abs/2404.05424 (2024). doi: [10.48550/arXiv.2404.05424](https://doi.org/10.48550/arXiv.2404.05424).
- [New98] R. G. Newcombe. ‘*Two-sided confidence intervals for the single proportion: comparison of seven methods*’. *Statistics in medicine* 17.8 (1998), pages 857–872. doi: [10.1002/\(SICI\)1097-0258\(19980430\)17:8<857::AID-SIM777>3.0.CO;2-E](https://doi.org/10.1002/(SICI)1097-0258(19980430)17:8<857::AID-SIM777>3.0.CO;2-E).
- [NG05] A. Nilim and L. E. Ghaoui. ‘*Robust Control of Markov Decision Processes with Uncertain Transition Matrices*’. *Oper. Res.* 53.5 (2005), pages 780–798. doi: [10.1287/OPRE.1050.0216](https://doi.org/10.1287/OPRE.1050.0216).
- [NPCN⁺21] T. Nyberg, C. Pek, L. D. Col, C. Norén and J. Tumova. ‘*Risk-aware Motion Planning for Autonomous Vehicles with Safety Specifications*’. IV. IEEE, 2021, pages 1016–1023. doi: [10.1109/IV48863.2021.9575928](https://doi.org/10.1109/IV48863.2021.9575928).
- [NSK09] M. R. Neuhausser, M. Stoelinga and J. Katoen. ‘*Delayed Nondeterminism in Continuous-Time Markov Decision Processes*’. *FoSSaCS*. Volume 5504. Lecture Notes in Computer Science. Springer, 2009, pages 364–379. doi: [10.1007/978-3-642-00596-1_26](https://doi.org/10.1007/978-3-642-00596-1_26).
- [OA16] F. A. Oliehoek and C. Amato. ‘*A Concise Introduction to Decentralized POMDPs*’. Springer Briefs in Intelligent Systems. Springer, 2016. doi: [10.1007/978-3-319-28929-8](https://doi.org/10.1007/978-3-319-28929-8).

- [OSV08] F. A. Oliehoek, M. T. J. Spaan and N. Vlassis. ‘Optimal and Approximate Q -value Functions for Decentralized POMDPs’. *J. Artif. Intell. Res.* 32 (2008), pages 289–353. doi: [10.1613/JAIR.2447](https://doi.org/10.1613/JAIR.2447).
- [PAQM18] Y. V. Pant, H. Abbas, R. A. Quaye and R. Mangharam. ‘Fly-by-logic: control of multi-drone fleets with temporal logic objectives’. *ICCPs*. IEEE Computer Society / ACM, 2018, pages 186–197. doi: [10.1109/ICCPs.2018.00026](https://doi.org/10.1109/ICCPs.2018.00026).
- [Per09] T. J. Perkins. ‘Maximum likelihood trajectories for continuous-time Markov chains’. *NIPS*. Curran Associates, Inc., 2009, pages 1437–1445. URL: https://proceedings.neurips.cc/paper_files/paper/2009/file/afda332245e2af431fb7b672a68b659d-Paper.pdf.
- [PGT03] J. Pineau, G. J. Gordon and S. Thrun. ‘Point-based value iteration: An anytime algorithm for POMDPs’. *IJCAI*. Morgan Kaufmann, 2003, pages 1025–1032. URL: <https://dl.acm.org/doi/10.5555/1630659.1630806>.
- [PK08] G. Papaefthymiou and B. Klöckl. ‘MCMC for wind power simulation’. *IEEE Transactions on Energy Conversion* 23.1 (2008), pages 234–240. doi: [10.1109/TEC.2007.914174](https://doi.org/10.1109/TEC.2007.914174).
- [PKOW20] E. van der Pol, T. Kipf, F. A. Oliehoek and M. Welling. ‘Plannable Approximations to MDP Homomorphisms: Equivariance under Actions’. *AAMAS*. International Foundation for Autonomous Agents and Multiagent Systems, 2020, pages 1431–1439. doi: [10.5555/3398761.3398926](https://doi.org/10.5555/3398761.3398926).
- [PLSS13] A. Puggelli, W. Li, A. L. Sangiovanni-Vincentelli and S. A. Seshia. ‘Polynomial-Time Verification of PCTL Properties of MDPs with Convex Uncertainties’. *CAV*. Volume 8044. Lecture Notes in Computer Science. Springer, 2013, pages 527–542. doi: [10.1007/978-3-642-39799-8_35](https://doi.org/10.1007/978-3-642-39799-8_35).
- [PM23] A. Peruffo and M. Mazo Jr. ‘Data-Driven Abstractions With Probabilistic Guarantees for Linear PETC Systems’. *IEEE Control. Syst. Lett.* 7 (2023), pages 115–120. doi: [10.1109/LCSYS.2022.3186187](https://doi.org/10.1109/LCSYS.2022.3186187).
- [PN17] A. S. Polydoros and L. Nalpantidis. ‘Survey of Model-Based Reinforcement Learning: Applications on Robotics’. *J. Intell. Robotic Syst.* 86.2 (2017), pages 153–173. doi: [10.1007/S10846-017-0468-Y](https://doi.org/10.1007/S10846-017-0468-Y).
- [Pnu77] A. Pnueli. ‘The Temporal Logic of Programs’. *FOCS*. IEEE Computer Society, 1977, pages 46–57. doi: [10.1109/SFCS.1977.32](https://doi.org/10.1109/SFCS.1977.32).
- [PP18] L. Petrucci and J. van de Pol. ‘Parameter Synthesis Algorithms for Parametric Interval Markov Chains’. *FORTE*. Volume 10854. Lecture Notes in Computer Science. Springer, 2018, pages 121–140. doi: [10.1007/978-3-319-92612-4_7](https://doi.org/10.1007/978-3-319-92612-4_7).
- [Pré03] A. Prékopa. ‘Probabilistic Programming’. *Stochastic Programming*. Volume 10. Handbooks in Operations Research and Management Science. Elsevier, 2003, pages 267–351. doi: [https://doi.org/10.1016/S0927-0507\(03\)10005-9](https://doi.org/10.1016/S0927-0507(03)10005-9).
- [Pré13] A. Prékopa. ‘Stochastic programming’. Volume 324. Springer Science & Business Media, 2013. doi: [10.1007/978-94-017-3087-7](https://doi.org/10.1007/978-94-017-3087-7).

- [PSQ13] S. Park, E. Serpedin and K. A. Qaraqe. ‘*Gaussian Assumption: The Least Favorable but the Most Useful [Lecture Notes]*’. *IEEE Signal Process. Mag.* 30.3 (2013), pages 183–186. doi: [10.1109/MSP.2013.2238691](https://doi.org/10.1109/MSP.2013.2238691).
- [Put94] M. L. Puterman. ‘*Markov Decision Processes: Discrete Stochastic Dynamic Programming*’. Wiley Series in Probability and Statistics. Wiley, 1994. doi: [10.1002/9780470316887](https://doi.org/10.1002/9780470316887).
- [PWHA16] E. Polgreen, V. B. Wijesuriya, S. Haesaert and A. Abate. ‘*Data-Efficient Bayesian Verification of Parametric Markov Chains*’. *QEST*. Volume 9826. Lecture Notes in Computer Science. Springer, 2016, pages 35–51. doi: [10.1007/978-3-319-43425-4_3](https://doi.org/10.1007/978-3-319-43425-4_3).
- [PWHA17] E. Polgreen, V. B. Wijesuriya, S. Haesaert and A. Abate. ‘*Automated Experiment Design for Data-Efficient Verification of Parametric Markov Decision Processes*’. *QEST*. Volume 10503. Lecture Notes in Computer Science. Springer, 2017, pages 259–274. doi: [10.1007/978-3-319-66335-7_16](https://doi.org/10.1007/978-3-319-66335-7_16).
- [PWHO⁺20] E. van der Pol, D. E. Worrall, H. van Hoof, F. A. Oliehoek and M. Welling. ‘*MDP Homomorphic Networks: Group Symmetries in Reinforcement Learning*’. *NeurIPS*. 2020. doi: [10.48550/arXiv.2006.16908](https://doi.org/10.48550/arXiv.2006.16908).
- [QDJJ⁺16] T. Quatmann, C. Dehnert, N. Jansen, S. Junges and J. Katoen. ‘*Parameter Synthesis for Markov Models: Faster Than Ever*’. *ATVA*. Volume 9938. Lecture Notes in Computer Science. 2016, pages 50–67. doi: [10.1007/978-3-319-46520-3_4](https://doi.org/10.1007/978-3-319-46520-3_4).
- [RAM23] L. Rickard, A. Abate and K. Margellos. ‘*Learning Robust Policies for Uncertain Parametric Markov Decision Processes*’. *CoRR* abs/2312.06344 (2023). doi: [10.48550/arXiv.2312.06344](https://doi.org/10.48550/arXiv.2312.06344).
- [RB01] B. Ravindran and A. G. Barto. ‘*Symmetries and Model Minimization in Markov Decision Processes*’. Technical report. 2001. URL: <https://dl.acm.org/doi/10.5555/897533>.
- [RB03] B. Ravindran and A. G. Barto. ‘*SMDP Homomorphisms: An Algebraic Approach to Abstraction in Semi-Markov Decision Processes*’. *IJCAI*. Morgan Kaufmann, 2003, pages 1011–1018. URL: <http://ijcai.org/Proceedings/03/Papers/145.pdf>.
- [RBNS⁺19] E. J. J. Ruijters, C. E. Budde, M. C. Nakhaee, M. I. A. Stoelinga, D. Bucur, D. Hiemstra and S. Schivo. ‘*FFORT: A Benchmark Suite for Fault Tree Analysis*’. *ESREL*. Research Publishing, 2019, pages 878–885. doi: [10.3850/978-981-11-2724-3_0641-cd](https://doi.org/10.3850/978-981-11-2724-3_0641-cd).
- [RC21] R. Rocchetta and L. G. Crespo. ‘*A scenario optimization approach to reliability-based and risk-based design: Soft-constrained modulation of failure probability bounds*’. *Reliab. Eng. Syst. Saf.* 216 (2021), page 107900. doi: [10.1016/J.RESS.2021.107900](https://doi.org/10.1016/J.RESS.2021.107900).

- [RGRK⁺09] K. D. Rao, V. Gopika, V. V. S. S. Rao, H. S. Kushwaha, A. K. Verma and A. Srividya. ‘Dynamic fault tree analysis using Monte Carlo simulation in probabilistic safety assessment’. *Reliab. Eng. Syst. Saf.* 94.4 (2009), pages 872–883. doi: [10.1016/J.RESS.2008.09.007](https://doi.org/10.1016/J.RESS.2008.09.007).
- [RLH21] M. Rigter, B. Lacerda and N. Hawes. ‘Risk-Averse Bayes-Adaptive Reinforcement Learning’. *NeurIPS*. 2021, pages 1142–1154. doi: [10.48550/arXiv.2102.05762](https://doi.org/10.48550/arXiv.2102.05762).
- [RM19] H. Rahimian and S. Mehrotra. ‘Distributionally Robust Optimization: A Review’. *CoRR* abs/1908.05659 (2019). doi: [10.48550/arXiv.1908.05659](https://doi.org/10.48550/arXiv.1908.05659).
- [RN10] S. J. Russell and P. Norvig. ‘Artificial Intelligence - A Modern Approach, Third International Edition’. Pearson Education, 2010. doi: https://people.engr.tamu.edu/guni/csce421/files/AI_Russell_Norvig.pdf.
- [RNBM⁺22] R. Roberts, T. Neupane, L. Buecherl, C. J. Myers and Z. Zhang. ‘STAMINA 2.0: Improving Scalability of Infinite-State Stochastic Model Checking’. *VMA-CAI*. Volume 13182. Lecture Notes in Computer Science. Springer, 2022, pages 319–331. doi: [10.1007/978-3-030-94583-1_16](https://doi.org/10.1007/978-3-030-94583-1_16).
- [RPM23] L. Romao, A. Papachristodoulou and K. Margellos. ‘On the Exact Feasibility of Convex Scenario Programs With Discarded Constraints’. *IEEE Trans. Autom. Control*. 68.4 (2023), pages 1986–2001. doi: [10.1109/TAC.2022.3165320](https://doi.org/10.1109/TAC.2022.3165320).
- [RPT16] P. Reist, P. Preiswerk and R. Tedrake. ‘Feedback-motion-planning with simulation-based LQR-trees’. *Int. J. Robotics Res.* 35.11 (2016), pages 1393–1416. doi: [10.1177/0278364916647192](https://doi.org/10.1177/0278364916647192).
- [RRBS19] E. Ruijters, D. Reijsbergen, P. de Boer and M. Stoelinga. ‘Rare event simulation for dynamic fault trees’. *Reliab. Eng. Syst. Saf.* 186 (2019), pages 220–231. doi: [10.1016/J.RESS.2019.02.004](https://doi.org/10.1016/J.RESS.2019.02.004).
- [RS03] A. Ruszczyński and A. Shapiro. ‘Stochastic Programming Models’. *Stochastic Programming*. Volume 10. Handbooks in Operations Research and Management Science. Elsevier, 2003, pages 1–64. doi: [https://doi.org/10.1016/S0927-0507\(03\)10001-1](https://doi.org/10.1016/S0927-0507(03)10001-1).
- [RS15] E. Ruijters and M. Stoelinga. ‘Fault tree analysis: A survey of the state-of-the-art in modeling, analysis and tools’. *Comput. Sci. Rev.* 15 (2015), pages 29–62. doi: [10.1016/J.COSREV.2015.03.001](https://doi.org/10.1016/J.COSREV.2015.03.001).
- [RSA22] U. Rosolia, A. Singletary and A. D. Ames. ‘Unified Multirate Control: From Low-Level Actuation to High-Level Planning’. *IEEE Trans. Autom. Control*. 67.12 (2022), pages 6627–6640. doi: [10.1109/TAC.2022.3184664](https://doi.org/10.1109/TAC.2022.3184664).
- [Rud⁺64] W. Rudin et al. ‘Principles of mathematical analysis’. Volume 3. McGraw-hill New York, 1964. URL: <https://www.mheducation.com/highered/product/principles-mathematical-analysis-rudin/M9780070542358.html>.

- [RWR17] G. Reissig, A. Weber and M. Rungger. ‘*Feedback Refinement Relations for the Synthesis of Symbolic Controllers*’. *IEEE Trans. Autom. Control.* 62.4 (2017), pages 1781–1796. doi: [10.1109/TAC.2016.2593947](https://doi.org/10.1109/TAC.2016.2593947).
- [SA13] S. E. Z. Soudjani and A. Abate. ‘*Adaptive and Sequential Gridding Procedures for the Abstraction and Verification of Stochastic Processes*’. *SIAM J. Appl. Dyn. Syst.* 12.2 (2013), pages 921–956. doi: [10.1137/120871456](https://doi.org/10.1137/120871456).
- [SAP24] Y. Schnitzer, A. Abate and D. Parker. ‘*Certifiably Robust Policies for Uncertain Parametric Environments*’. *CoRR* abs/2408.03093 (2024). doi: [10.48550/ARXIV.2408.03093](https://doi.org/10.48550/ARXIV.2408.03093).
- [SB98] R. S. Sutton and A. G. Barto. ‘*Reinforcement learning - an introduction*’. Adaptive computation and machine learning. MIT Press, 1998. URL: <https://mitpress.mit.edu/9780262039246/reinforcement-learning/>.
- [SBHH17] D. Scheftelowitsch, P. Buchholz, V. Hashemi and H. Hermanns. ‘*Multi-Objective Approaches to Markov Decision Processes with Uncertain Transition Parameters*’. *VALUETOOLS*. ACM, 2017, pages 44–51. doi: [10.1145/3150928.3150945](https://doi.org/10.1145/3150928.3150945).
- [SBSG⁺11] S. D. Stoller, E. Bartocci, J. Seyster, R. Grosu, K. Havelund, S. A. Smolka and E. Zadok. ‘*Runtime Verification with State Estimation*’. *RV*. Volume 7186. Lecture Notes in Computer Science. Springer, 2011, pages 193–207. doi: [10.1007/978-3-642-29860-8_15](https://doi.org/10.1007/978-3-642-29860-8_15).
- [Seg95] R. Segala. ‘*Modeling and verification of randomized distributed real-time systems*’. PhD thesis. Massachusetts Institute of Technology, Cambridge, MA, USA, 1995. URL: <https://groups.csail.mit.edu/tds/papers/Segala/phd1.pdf>.
- [SFCR16] G. Su, Y. Feng, T. Chen and D. S. Rosenblum. ‘*Asymptotic Perturbation Bounds for Probabilistic Model Checking with Empirically Determined Probability Parameters*’. *IEEE Trans. Software Eng.* 42.7 (2016), pages 623–639. doi: [10.1109/TSE.2015.2508444](https://doi.org/10.1109/TSE.2015.2508444).
- [Shm04] V. Shmatikov. ‘*Probabilistic analysis of an anonymity system*’. *J. Comput. Secur.* 12.3-4 (2004), pages 355–377. doi: [10.3233/JCS-2004-123-403](https://doi.org/10.3233/JCS-2004-123-403).
- [SJCT20] M. Suilen, N. Jansen, M. Cubuktepe and U. Topcu. ‘*Robust Policy Synthesis for Uncertain POMDPs via Convex Optimization*’. *IJCAI*. ijcai.org, 2020, pages 4113–4120. doi: [10.24963/IJCAI.2020/569](https://doi.org/10.24963/IJCAI.2020/569).
- [SJK21] J. Spel, S. Junges and J. Katoen. ‘*Finding Provably Optimal Markov Chains*’. *TACAS (1)*. Volume 12651. Lecture Notes in Computer Science. Springer, 2021, pages 173–190. doi: [10.1007/978-3-030-72016-2_10](https://doi.org/10.1007/978-3-030-72016-2_10).
- [SKCC⁺15] M. Svorenová, J. Kretínský, M. Chmelík, K. Chatterjee, I. Cerná and C. Belta. ‘*Temporal logic control for stochastic linear systems using abstraction refinement of probabilistic games*’. *HSCC*. ACM, 2015, pages 259–268. doi: [10.1145/2728606.2728608](https://doi.org/10.1145/2728606.2728608).

- [SKD21] L. N. Steimle, D. L. Kaufman and B. T. Denton. ‘*Multi-model Markov decision processes*’. *IIS Trans.* 53.10 (2021), pages 1124–1139. doi: [10.1080/24725854.2021.1895454](https://doi.org/10.1080/24725854.2021.1895454).
- [SKLT11] S. Summers, M. Kamgarpour, J. Lygeros and C. J. Tomlin. ‘*A stochastic reach-avoid problem with random obstacles*’. *HSCC*. ACM, 2011, pages 251–260. doi: [10.1145/1967701.1967738](https://doi.org/10.1145/1967701.1967738).
- [Sku09] D. Skulj. ‘*Discrete time Markov chains with interval probabilities*’. *Int. J. Approx. Reason.* 50.8 (2009), pages 1314–1329. doi: [10.1016/J.IJAR.2009.06.007](https://doi.org/10.1016/J.IJAR.2009.06.007).
- [SL10] S. Summers and J. Lygeros. ‘*Verification of discrete time stochastic hybrid systems: A stochastic reach-avoid decision problem*’. *Autom.* 46.12 (2010), pages 1951–1961. doi: [10.1016/J.AUTOMATICA.2010.08.006](https://doi.org/10.1016/J.AUTOMATICA.2010.08.006).
- [SL95] R. Segala and N. A. Lynch. ‘*Probabilistic Simulations for Probabilistic Processes*’. *Nord. J. Comput.* 2.2 (1995), pages 250–273. doi: [10.1007/978-3-540-48654-1_35](https://doi.org/10.1007/978-3-540-48654-1_35).
- [SLO22] R. A. N. Starre, M. Loog and F. A. Oliehoek. ‘*Model-Based Reinforcement Learning with State Abstraction: A Survey*’. *BNAIC/BENELEARN*. Volume 1805. Communications in Computer and Information Science. Springer, 2022, pages 133–148. doi: [10.1007/978-3-031-39144-6\9](https://doi.org/10.1007/978-3-031-39144-6\9).
- [Smi13] A. Smith. ‘*Sequential Monte Carlo methods in practice*’. Springer Science & Business Media, 2013. doi: [10.1007/978-1-4757-3437-9](https://doi.org/10.1007/978-1-4757-3437-9).
- [Smi14] R. C. Smith. ‘*Uncertainty Quantification - Theory, Implementation, and Applications*’. Computational science and engineering. SIAM, 2014. doi: [10.1137/1.9781611973228](https://doi.org/10.1137/1.9781611973228).
- [Söd02] T. Söderström. ‘*Discrete-time stochastic systems: estimation and control*’. Springer Science & Business Media, 2002. doi: [10.1007/978-1-4471-0101-7](https://doi.org/10.1007/978-1-4471-0101-7).
- [SPK13] G. Shani, J. Pineau and R. Kaplow. ‘*A survey of point-based POMDP solvers*’. *Auton. Agents Multi Agent Syst.* 27.1 (2013), pages 1–51. doi: [10.1007/S10458-012-9200-2](https://doi.org/10.1007/S10458-012-9200-2).
- [SSAB⁺19] C. Sánchez, G. Schneider, W. Ahrendt, E. Bartocci, D. Bianculli, C. Colombo, Y. Falcone, A. Francalanza, S. Krstic, J. M. Lourenço, D. Nickovic, G. J. Pace, J. Rufino, J. Signoles, D. Traytel and A. Weiss. ‘*A survey of challenges for runtime verification from advanced application domains (beyond software)*’. *Formal Methods Syst. Des.* 54.3 (2019), pages 279–335. doi: [10.1007/S10703-019-00337-W](https://doi.org/10.1007/S10703-019-00337-W).
- [SSPJ22] M. Suilen, T. D. Simão, D. Parker and N. Jansen. ‘*Robust Anytime Learning of Markov Decision Processes*’. *NeurIPS*. 2022. doi: [10.48550/arXiv.2205.15827](https://doi.org/10.48550/arXiv.2205.15827).
- [STBR11] S. L. Smith, J. Tumova, C. Belta and D. Rus. ‘*Optimal path planning for surveillance with temporal-logic constraints*’. *Int. J. Robotics Res.* 30.14 (2011), pages 1695–1708. doi: [10.1177/0278364911417911](https://doi.org/10.1177/0278364911417911).

- [Ste94] W. J. Stewart. ‘*Introduction to the numerical solution of Markov Chains*’. Princeton University Press, 1994. doi: [10.2307/j.ctv182jsw5](https://doi.org/10.2307/j.ctv182jsw5).
- [Sto02] M. Stoelinga. ‘*An Introduction to Probabilistic Automata*’. *Bull. EATCS* 78 (2002), pages 176–198. url: https://web.archive.org/web/20170829004616id_/http://wwwhome.ewi.utwente.nl/~marielle/papers/pa.pdf.
- [Str23] G. Strang. ‘*Introduction to linear algebra*’. 6th. SIAM, 2023. url: <https://math.mit.edu/~gs/linearalgebra/ila6/indexila6.html>.
- [Sul15] T. J. Sullivan. ‘*Introduction to uncertainty quantification*’. Volume 63. Springer, 2015. doi: [10.1007/978-3-319-23395-6](https://doi.org/10.1007/978-3-319-23395-6).
- [SV05] M. T. J. Spaan and N. Vlassis. ‘*Perseus: Randomized Point-based Value Iteration for POMDPs*’. *J. Artif. Intell. Res.* 24 (2005), pages 195–220. doi: [10.1613/JAIR.1659](https://doi.org/10.1613/JAIR.1659).
- [SVA05] K. Sen, M. Viswanathan and G. Agha. ‘*On Statistical Model Checking of Stochastic Systems*’. *CAV*. Volume 3576. Lecture Notes in Computer Science. Springer, 2005, pages 266–280. doi: [10.1007/11513988_26](https://doi.org/10.1007/11513988_26).
- [SVA06] K. Sen, M. Viswanathan and G. Agha. ‘*Model-Checking Markov Chains in the Presence of Uncertainties*’. *TACAS*. Volume 3920. Lecture Notes in Computer Science. Springer, 2006, pages 394–410. doi: [10.1007/11691372_26](https://doi.org/10.1007/11691372_26).
- [SVAO19] H. Sar tipizadeh, A. P. Vinod, B. A cikmese and M. Oishi. ‘*Voronoi Partition-based Scenario Reduction for Fast Sampling-based Stochastic Reachability Computation of Linear Systems*’. *ACC*. IEEE, 2019, pages 37–44. doi: [10.23919/ACC.2019.8814354](https://doi.org/10.23919/ACC.2019.8814354).
- [SZ15] F. Shmarov and P. Zuliani. ‘*ProbReach: verified probabilistic delta-reachability for stochastic hybrid systems*’. *HSCC*. ACM, 2015, pages 134–139. doi: [10.1145/2728606.2728625](https://doi.org/10.1145/2728606.2728625).
- [SZ23] A. Salamati and M. Zamani. ‘*Safety Verification of Stochastic Systems: A Repetitive Scenario Approach*’. *IEEE Control. Syst. Lett.* 7 (2023), pages 448–453. doi: [10.1109/LCSYS.2022.3186932](https://doi.org/10.1109/LCSYS.2022.3186932).
- [Tab09] P. Tabuada. ‘*Verification and Control of Hybrid Systems - A Symbolic Approach*’. Springer, 2009. doi: [10.1007/978-1-4419-0224-5](https://doi.org/10.1007/978-1-4419-0224-5).
- [TBBB⁺20] B. Ton, R. Basten, J. Bolte, J. Braaksma, A. Di Buccianico, P. van de Calseyde, F. Groteman, T. Heskes, N. Jansen, W. Teeuw, T. Tinga and M. Stoelinga. ‘*PrimaVera: Synergising Predictive Maintenance*’. *Applied Sciences* 10.23 (2020). doi: [10.3390/app10238348](https://doi.org/10.3390/app10238348).
- [TBF05] S. Thrun, W. Burgard and D. Fox. ‘*Probabilistic robotics*’. Intelligent robotics and autonomous agents. MIT Press, 2005. url: <https://mitpress.mit.edu/9780262201629/probabilistic-robotics/>.
- [TLS21] S. Thiebes, S. Lins and A. Sunyaev. ‘*Trustworthy artificial intelligence*’. *Electron. Mark.* 31.2 (2021), pages 447–464. doi: [10.1007/S12525-020-00441-4](https://doi.org/10.1007/S12525-020-00441-4).

- [TMTR10] R. Tedrake, I. R. Manchester, M. M. Tobenkin and J. W. Roberts. ‘*LQR-trees: Feedback Motion Planning via Sums-of-Squares Verification*’. *Int. J. Robotics Res.* 29.8 (2010), pages 1038–1052. doi: [10.1177/0278364910369189](https://doi.org/10.1177/0278364910369189).
- [TP19] A. Tsiamis and G. J. Pappas. ‘*Finite Sample Analysis of Stochastic System Identification*’. *CDC. IEEE*, 2019, pages 3648–3654. doi: [10.1109/CDC40024.2019.9029499](https://doi.org/10.1109/CDC40024.2019.9029499).
- [TSYA20] A. J. Taylor, A. Singletary, Y. Yue and A. D. Ames. ‘*Learning for Safety-Critical Control with Control Barrier Functions*’. *L4DC*. Volume 120. Proceedings of Machine Learning Research. PMLR, 2020, pages 708–717. doi: [10.48550/arXiv.1912.10099](https://doi.org/10.48550/arXiv.1912.10099).
- [Vaa17] F. W. Vaandrager. ‘*Model learning*’. *Commun. ACM* 60.2 (2017), pages 86–95. doi: [10.1145/2967606](https://doi.org/10.1145/2967606).
- [VGO19] A. P. Vinod, J. D. Gleason and M. M. K. Oishi. ‘*SReachTools: a MATLAB stochastic reachability toolbox*’. *HSCC. ACM*, 2019, pages 33–38. doi: [10.1145/3302504.3311809](https://doi.org/10.1145/3302504.3311809).
- [VGOH⁺19] P. Virtanen, R. Gommers, T. E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, E. Burovski, P. Peterson, W. Weckesser, J. Bright, S. van der Walt, M. Brett, J. Wilson, K. J. Millman, N. Mayorov, A. R. J. Nelson, E. Jones, R. Kern, E. Larson, C. Carey, I. Polat, Y. Feng, E. W. Moore, J. VanderPlas, D. Laxalde, J. Perktold, R. Cimrman, I. Henriksen, E. A. Quintero, C. R. Harris, A. M. Archibald, A. H. Ribeiro, F. Pedregosa, P. van Mulbregt and SciPy. ‘*SciPy 1.0-Fundamental Algorithms for Scientific Computing in Python*’. *CoRR* abs/1907.10121 (2019). doi: [10.1038/s41592-019-0686-2](https://doi.org/10.1038/s41592-019-0686-2).
- [VIT22] A. P. Vinod, A. Israel and U. Topcu. ‘*On-the-Fly Control of Unknown Nonlinear Systems With Sublinear Regret*’. *IEEE Transactions on Automatic Control* (2022), pages 1–13. doi: [10.1109/TAC.2022.3186425](https://doi.org/10.1109/TAC.2022.3186425).
- [VJK18] M. Volk, S. Junges and J. Katoen. ‘*Fast Dynamic Fault Tree Analysis by Model Checking Techniques*’. *IEEE Trans. Ind. Informatics* 14.1 (2018), pages 370–379. doi: [10.1109/TII.2017.2710316](https://doi.org/10.1109/TII.2017.2710316).
- [Vol22] M. Volk. ‘*Dynamic fault trees: semantics, analysis and applications*’. PhD thesis. Dissertation, RWTH Aachen University, 2022, 2022. doi: [10.18154/RWTH-2023-04092](https://doi.org/10.18154/RWTH-2023-04092).
- [VT24] M. Vahs and J. Tumova. ‘*Risk-aware Control for Robots with Non-Gaussian Belief Spaces*’. *ICRA. IEEE*, 2024, pages 11661–11667. doi: [10.1109/ICRA57147.2024.10611412](https://doi.org/10.1109/ICRA57147.2024.10611412).
- [WA19] V. B. Wijesuriya and A. Abate. ‘*Bayes-Adaptive Planning for Data-Efficient Verification of Uncertain Markov Decision Processes*’. *QEST*. Volume 11785. Lecture Notes in Computer Science. Springer, 2019, pages 91–108. doi: [10.1007/978-3-030-30281-8_6](https://doi.org/10.1007/978-3-030-30281-8_6).
- [WB01] G. Welch and G. Bishop. ‘*An introduction to the Kalman filter*’. *Proc of SIGGRAPH, Course* 8.27599-23175 (2001), page 41. url: https://www.cs.unc.edu/~welch/media/pdf/kalman_intro.pdf.

- [Wil91] D. Williams. ‘*Probability with Martingales*’. Cambridge mathematical textbooks. Cambridge University Press, 1991. doi: [10.1017/CBO9780511813658](https://doi.org/10.1017/CBO9780511813658).
- [WJPK19] T. Winkler, S. Junges, G. A. Pérez and J. Katoen. ‘*On the Complexity of Reachability in Parametric Markov Decision Processes*’. CONCUR. Volume 140. LIPIcs. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2019, 14:1–14:17. doi: [10.4230/LIPICS.CONCUR.2019.14](https://doi.org/10.4230/LIPICS.CONCUR.2019.14).
- [WJWJ⁺21] L. Winterer, S. Junges, R. Wimmer, N. Jansen, U. Topcu, J. Katoen and B. Becker. ‘*Strategy Synthesis for POMDPs in Robot Planning via Game-Based Abstractions*’. IEEE Trans. Autom. Control. 66.3 (2021), pages 1040–1054. doi: [10.1109/TAC.2020.2990140](https://doi.org/10.1109/TAC.2020.2990140).
- [WKR13] W. Wiesemann, D. Kuhn and B. Rustem. ‘*Robust Markov Decision Processes*’. Math. Oper. Res. 38.1 (2013), pages 153–183. doi: [10.1287/MOOR.1120.0566](https://doi.org/10.1287/MOOR.1120.0566).
- [WKS14] W. Wiesemann, D. Kuhn and M. Sim. ‘*Distributionally Robust Convex Optimization*’. Oper. Res. 62.6 (2014), pages 1358–1376. doi: [10.1287/OPRE.2014.1314](https://doi.org/10.1287/OPRE.2014.1314).
- [WL24] B. Wooding and A. Lavaei. ‘*IMPACT: A Parallelized Software Tool for IMDP Construction and Controller Synthesis with Convergence Guarantees*’. HSCC. ACM, 2024, 30:1–30:2. doi: [10.1145/3641513.3652532](https://doi.org/10.1145/3641513.3652532).
- [Wol20] L. A. Wolsey. ‘*Integer programming*’. John Wiley & Sons, 2020. doi: [10.1002/9781119606475](https://doi.org/10.1002/9781119606475).
- [WTM12] E. M. Wolff, U. Topcu and R. M. Murray. ‘*Robust control of uncertain Markov Decision Processes with temporal logic specifications*’. CDC. IEEE, 2012, pages 3372–3379. doi: [10.1109/CDC.2012.6426174](https://doi.org/10.1109/CDC.2012.6426174).
- [WW99] D. S. Wilks and R. L. Wilby. ‘*The weather generation game: a review of stochastic weather models*’. Progress in Physical Geography: Earth and Environment 23.3 (1999), pages 329–357. doi: [10.1177/030913339902300302](https://doi.org/10.1177/030913339902300302).
- [XGAM⁺20] Z. Xu, I. Gavran, Y. Ahmad, R. Majumdar, D. Neider, U. Topcu and B. Wu. ‘*Joint Inference of Reward Machines and Policies for Reinforcement Learning*’. ICAPS. AAAI Press, 2020, pages 590–598. doi: [10.48550/arXiv.1909.05912](https://doi.org/10.48550/arXiv.1909.05912).
- [XM10] H. Xu and S. Mannor. ‘*Distributionally Robust Markov Decision Processes*’. NIPS. Curran Associates, Inc., 2010, pages 2505–2513. URL: https://proceedings.neurips.cc/paper_files/paper/2010/file/19f3cd308f1455b3fa09a282e0d496f4-Paper.pdf.
- [XM12] H. Xu and S. Mannor. ‘*Distributionally Robust Markov Decision Processes*’. Math. Oper. Res. 37.2 (2012), pages 288–300. doi: [10.1287/MOOR.1120.0540](https://doi.org/10.1287/MOOR.1120.0540).
- [Yed14] R. K. Yedavalli. ‘*Robust control of uncertain dynamic systems*’. AMC 10 (2014), page 12. doi: [10.1007/978-1-4614-9132-3](https://doi.org/10.1007/978-1-4614-9132-3).

- [YS06] H. L. S. Younes and R. G. Simmons. ‘*Statistical probabilistic model checking with a focus on time-bounded properties*’. *Inf. Comput.* 204.9 (2006), pages 1368–1409. doi: [10.1016/J.IC.2006.05.002](https://doi.org/10.1016/J.IC.2006.05.002).
- [YSHL17] N. Ye, A. Soman, D. Hsu and W. S. Lee. ‘*DESPOT: Online POMDP Planning with Regularization*’. *J. Artif. Intell. Res.* 58 (2017), pages 231–266. doi: [10.1613/JAIR.5328](https://doi.org/10.1613/JAIR.5328).
- [YTCB⁺12] B. Yordanov, J. Tumova, I. Cerna, J. Barnat and C. Belta. ‘*Temporal Logic Control of Discrete-Time Piecewise Affine Systems*’. *IEEE Trans. Autom. Control* 57.6 (2012), pages 1491–1504. doi: [10.1109/TAC.2011.2178328](https://doi.org/10.1109/TAC.2011.2178328).
- [ZD98] K. Zhou and J. C. Doyle. ‘*Essentials of robust control*’. Volume 104. Prentice hall Upper Saddle River, NJ, 1998. url: <https://www.ece.lsu.edu/kemin/essentials.htm>.
- [ZG21] M. Zanon and S. Gros. ‘*Safe Reinforcement Learning Using Robust MPC*’. *IEEE Trans. Autom. Control* 66.8 (2021), pages 3638–3652. doi: [10.1109/TAC.2020.3024161](https://doi.org/10.1109/TAC.2020.3024161).
- [ZLHC23b] D. Zikelic, M. Lechner, T. A. Henzinger and K. Chatterjee. ‘*Learning Control Policies for Stochastic Systems with Reach-Avoid Guarantees*’. AAAI. AAAI Press, 2023, pages 11926–11935. doi: [10.1609/AAAI.V37I10.26407](https://doi.org/10.1609/AAAI.V37I10.26407).
- [ZSRH⁺12] L. Zhang, Z. She, S. Ratschan, H. Hermanns and E. M. Hahn. ‘*Safety Verification for Probabilistic Hybrid Systems*’. *Eur. J. Control* 18.6 (2012), pages 572–587. doi: [10.3166/EJC.18.572-587](https://doi.org/10.3166/EJC.18.572-587).

B Index

- reach-avoid specification, 55
- absorbing state, 30, 85
- abstraction, 12, 85, 244, 245
 - refinement, 247
- atomic proposition, 29
- backward reachable set, 86, 122
- Bayes' rule, 237
- binary relation, 65
 - lifted relation, 67
 - single-valued, 65
 - strict, 65
- Borel space, 51
- boundary, 23
- chance-constrained optimization
 - problem, 155
- closure, 23
- consistent scheduler, 241
- continuous stochastic logic, 201
 - path formula, 201
 - state formula, 201
- continuous-time Markov chain, 198
 - colors, 237
 - exit rates, 198
 - path, 200
 - probability measure, 200
 - residence time, 198
 - transient distribution, 199
 - uniformization, 204
- convex function, 25
- convex hull, 24
- convex polytope, 24, 176
 - parametric, 177
- CTMC monitoring, 236
 - conditioned MDP, 243
 - evidence, 237
 - imprecise evidence, 238
- path consistent with evidence, 237
- unfolded MDP, 241
- deadlock, 30
- diagonal matrix, 24
- discount factor, 37
- discrete-time Markov chain, 30
 - induced, 32
 - robust, 41
- discrete-time stochastic system, 51
 - execution, 54
 - induced, 118
 - linear, 82
 - Markov policy, 53
 - probability measure, 54
 - process noise, 52
 - robust, 118
 - sample path, 54
 - time-invariant, 51
- epistemic error, 121
- evaluation problem, 4
- evidence graph, 240
- fault tree, 201, 228, 250
- filtration, 27
 - natural, 27
- gradient, 178
- Hoeffding's inequality, 99, 129, 190
- identity matrix, 24
- indicator function, 24
- interface function, 72
 - robust, 77
- interior, 23
- interval MDP, 41
- Markov decision process, 29
 - optimal scheduler, 37

- path, 31
 probability measure, 33
 measurable function, 26
 measurable space, 26
 measure (for CTMC), 203
 expected cumulative reward, 204
 steady-state probability, 203
 transient probability, 203
 measure (for MDP), 35
 reach-avoid probability, 36
 expected cumulative reward, 36
 reachability probability, 36
 nature, 42
 Markov, 43
 optimistic, 44
 pessimistic, 44
 stationary, 43
 optimization problem, 24
 convex, 25
 parameter synthesis, 148
 parametric CTMC, 205
 parameter space, 206
 solution function, 207
 vector-valued solution function, 212
 parametric MDP, 144
 parameter instantiation, 145
 parameter space, 145
 parametric Markov chain, 144
 solution function, 146
 parametric robust Markov chain, 177
 robust solution function, 178
 partial derivative, 178
 partial map, 24
 partial transition function, 29
 partition, 65, 85, 245
 prediction region, 213
 containment probability, 213
 probabilistic computation tree logic, 34
 path formula, 34
 state formula, 34
 probabilistic simulation relation, 67, 89
 alternating, 75, 102
 probability distribution
 Dirac, 26
 discrete, 26
 probability space, 26
 pseudoinverse, 24
 random variable, 26
 reach-avoid specification
 consistent, 66
 equivalent, 66
 robust MDP, 41
 rectangularity, 41
 robust value, 43
 uncertainty set, 41
 robust optimization, 13
 satisfaction probability, 35
 scenario approach, 14, 91, 128, 156, 215
 complexity, 219
 scenario optimization problem, 14, 156, 216
 scheduler, 31
 deterministic, 32
 Markov scheduler, 32
 randomized, 31
 stationary, 32
 self-loop, 30
 set-bounded parameter uncertainty, 119
 state-weight function, 238
 stochastic kernel, 54
 stochastic process, 2, 27
 adapted, 27
 stochasticity, 2
 synthesis problem, 4
 target point, 86
 target set, 121
 transition, 30, 198
 uncertain parametric CTMC, 211
 uncertain parametric MDP, 153
 satisfaction probability, 154
 uncertainty set, 7
 value iteration, 38

C Research Data Management

C

This thesis research has been carried out under the research data management policy of the Institute for Computing and Information Science of Radboud University, located in Nijmegen, The Netherlands.¹

The following code repositories have been produced during this Ph.D. research:

- Chapters 6 and 7: Thom Badings, Alessandro Abate, Nils Jansen, David Parker, Hasan Poonawala, Licio Romao and Marielle Stoelinga; <https://doi.org/10.5281/zenodo.13348782> (Python source code)
- Chapter 9: Thom Badings, Murat Cubuktepe, Nils Jansen, Sebastian Junges, Joost-Pieter Katoen and Ufuk Topcu; <https://doi.org/10.5281/zenodo.6674059> (Docker container and Python source code)
- Chapter 10: Thom Badings, Sebastian Junges, Ahmadreza Marandi, Ufuk Topcu and Nils Jansen; <https://doi.org/10.5281/zenodo.7864260> (Docker container and Python source code)
- Chapter 12: Thom Badings, Nils Jansen, Sebastian Junges, Marielle Stoelinga and Matthias Volk; <https://doi.org/10.5281/zenodo.6523863> (Docker container and Python source code)
- Chapter 13: Thom Badings, Matthias Volk, Sebastian Junges, Marielle Stoelinga and Nils Jansen; <https://doi.org/10.5281/zenodo.10438984> (Docker container and Python source code)

For a more detailed overview of these code repositories and the tools and implementations they contain, we refer to Chapter 14.

¹ru.nl/icis/research-data-management/, last accessed August 19th, 2024.

D About the Author

Thom Badings was born on July 17, 1995 in Zwolle, the Netherlands. After his secondary education at the Meander College in Zwolle, he moved to Groningen to study Industrial Engineering and Management at the University of Groningen. In 2017, he obtained his Bachelor of Science degree with a double specialization, as well as the Honours College Bachelor's program. Thom continued with a master's in Industrial Engineering and Management at the same university, which he completed in 2019 with a *cum laude* distinction. He followed a specialization in smart systems in control and automation, and his master's thesis was on the topic of model predictive control for power systems with uncertain electricity generation.



Thom started his Ph.D. research in 2020 at Radboud University in Nijmegen, the Netherlands, under the supervision of Prof. Nils Jansen and Prof. Marielle Stoelinga. He is part of PrimaVera, an academic research consortium on the topic of predictive maintenance. Thom's research interests are broadly on the intersection between control theory, artificial intelligence, and formal methods. His Ph.D. research focused on developing novel verification methods for proving properties about the behavior of stochastic systems in the presence of uncertainty. With his research, Thom aims to develop methods that can be used to provide rigorous mathematical guarantees about properties of safety, reliability, and performance of complex systems, even if no perfect model of the system is available.

In 2022, he received a distinguished paper award at AAAI, one of the leading conferences in artificial intelligence. During his Ph.D., Thom has obtained the University Teaching Qualification (UTQ) and has taught in several courses at both bachelor's and master's level. Since November 2024, Thom is a postdoctoral research associate with the Department of Computer Science at the University of Oxford, where he works with Prof. Alessandro Abate on research about understanding decision making under uncertainty in artificial intelligence.

Titles in the IPA Dissertation Series since 2022

A. Fedotov. *Verification Techniques for xMAS.* Faculty of Mathematics and Computer Science, TU/e. 2022-01

M.O. Mahmoud. *GPU Enabled Automated Reasoning.* Faculty of Mathematics and Computer Science, TU/e. 2022-02

M. Safari. *Correct Optimized GPU Programs.* Faculty of Electrical Engineering, Mathematics & Computer Science, UT. 2022-03

M. Verano Merino. *Engineering Language-Parametric End-User Programming Environments for DSLs.* Faculty of Mathematics and Computer Science, TU/e. 2022-04

G.F.C. Dupont. *Network Security Monitoring in Environments where Digital and Physical Safety are Critical.* Faculty of Mathematics and Computer Science, TU/e. 2022-05

T.M. Soethout. *Banking on Domain Knowledge for Faster Transactions.* Faculty of Mathematics and Computer Science, TU/e. 2022-06

P. Vukmirović. *Implementation of Higher-Order Superposition.* Faculty of Sciences, Department of Computer Science, VU. 2022-07

J. Wagemaker. *Extensions of (Concurrent) Kleene Algebra.* Faculty of Science, Mathematics and Computer Science, RU. 2022-08

R. Janssen. *Refinement and Partiality for Model-Based Testing.* Faculty of Science, Mathematics and Computer Science, RU. 2022-09

M. Laveaux. *Accelerated Verification of Concurrent Systems.* Faculty of Mathematics and Computer Science, TU/e. 2022-10

S. Kochanthara. *A Changing Landscape: On Safety & Open Source in Automated and Connected Driving.* Faculty of Mathematics and Computer Science, TU/e. 2023-01

L.M. Ochoa Venegas. *Break the Code? Breaking Changes and Their Impact on Software Evolution.* Faculty of Mathematics and Computer Science, TU/e. 2023-02

N. Yang. *Logs and models in engineering complex embedded production software systems.* Faculty of Mathematics and Computer Science, TU/e. 2023-03

J. Cao. *An Independent Timing Analysis for Credit-Based Shaping in Ethernet TSN.* Faculty of Mathematics and Computer Science, TU/e. 2023-04

K. Dokter. *Scheduled Protocol Programming.* Faculty of Mathematics and Natural Sciences, UL. 2023-05

J. Smits. *Strategic Language Workbench Improvements.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2023-06

A. Arslanagić. *Minimal Structures for Program Analysis and Verification.* Faculty of Science and Engineering, RUG. 2023-07

M.S. Bouwman. *Supporting Railway Standardisation with Formal Verification.* Faculty of Mathematics and Computer Science, TU/e. 2023-08

S.A.M. Lathouwers. *Exploring Annotations for Deductive Verification.* Faculty of Electrical Engineering, Mathematics & Computer Science, UT. 2023-09

J.H. Stoel. *Solving the Bank, Lightweight Specification and Verification Techniques for Enterprise Software.* Faculty of Mathematics and Computer Science, TU/e. 2023-10

- D.M. Groenewegen.** *WebDSL: Linguistic Abstractions for Web Programming.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2023-11
- D.R. do Vale.** *On Semantical Methods for Higher-Order Complexity Analysis.* Faculty of Science, Mathematics and Computer Science, RU. 2024-01
- M.J.G. Olsthoorn.** *More Effective Test Case Generation with Multiple Tribes of AI.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2024-02
- B. van den Heuvel.** *Correctly Communicating Software: Distributed, Asynchronous, and Beyond.* Faculty of Science and Engineering, RUG. 2024-03
- H.A. Hiep.** *New Foundations for Separation Logic.* Faculty of Mathematics and Natural Sciences, UL. 2024-04
- C.E. Brandt.** *Test Amplification For and With Developers.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2024-05
- J.I. Hejderup.** *Fine-Grained Analysis of Software Supply Chains.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2024-06
- J. Jacobs.** *Guarantees by construction.* Faculty of Science, Mathematics and Computer Science, RU. 2024-07
- O. Bunte.** *Cracking OIL: A Formal Perspective on an Industrial DSL for Modelling Control Software.* Faculty of Mathematics and Computer Science, TU/e. 2024-08
- R.J.A. Erkens.** *Automaton-based Techniques for Optimized Term Rewriting.* Faculty of Mathematics and Computer Science, TU/e. 2024-09
- J.J.M. Martens.** *The Complexity of Bisimilarity by Partition Refinement.* Faculty of Mathematics and Computer Science, TU/e. 2024-10
- L.J. Edixhoven.** *Expressive Specification and Verification of Choreographies.* Faculty of Science, OU. 2024-11
- J.W.N. Paulus.** *On the Expressivity of Typed Concurrent Calculi.* Faculty of Science and Engineering, RUG. 2024-12
- J. Denkers.** *Domain-Specific Languages for Digital Printing Systems.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2024-13
- L.H. Applis.** *Tool-Driven Quality Assurance for Functional Programming and Machine Learning.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2024-14
- P. Karkhanis.** *Driving the Future: Facilitating C-ITS Service Deployment for Connected and Smart Roadways.* Faculty of Mathematics and Computer Science, TU/e. 2024-15
- N.W. Cassee.** *Sentiment in Software Engineering.* Faculty of Mathematics and Computer Science, TU/e. 2024-16
- H. van Antwerpen.** *Declarative Name Binding for Type System Specifications.* Faculty of Electrical Engineering, Mathematics, and Computer Science, TUD. 2025-01
- I.N. Mulder.** *Proof Automation for Fine-Grained Concurrent Separation Logic.* Faculty of Science, Mathematics and Computer Science, RU. 2025-02
- T.S. Badings.** *Robust Verification of Stochastic Systems: Guarantees in the Presence of Uncertainty.* Faculty of Science, Mathematics and Computer Science, RU. 2025-03

